Making the Elite: Caste, Class, and Cultural Capital in Multinational Hiring

Soumitra Shukla*
Harvard Business School

March 10, 2026

Abstract

Class disparities persist in elite hiring, but where do they arise and why? India's codified caste categories provide a tractable proxy for cumulative social class, allowing me to identify where disparities emerge and to evaluate whether cultural-capital screening is a central channel through which caste-linked disadvantage operates via class-coded biographical cues. Using screening data from over 1,000 jobs at multinationals recruiting at an elite Indian college, I follow candidates across four rounds and show that subjective personal interviews assessing "cultural fit" account for over 95% of the caste success-rate penalty. Two tests triangulate the mechanism. First, large language model analysis of nearly one million words from publicly available mock-preparation transcripts shows that personal interviews emphasize biographical markers of cultural capital—hometown, schooling, hobbies, family background, and parents' occupation—while earlier subjective rounds such as group debates focus on abstract policy reasoning. Second, surname timing shows penalties at application screening for the 15% of students with caste-identifiable surnames, but only after fit interviews for the 85% with initially ambiguous surnames, consistent with a pattern in which, for the majority, disparities emerge when interview content makes cultural-capital signals legible to evaluators rather than from in-person cues alone. Disadvantaged-caste hires show lower attrition and higher promotion rates across job selectivity tiers and role types, patterns difficult to reconcile with simple customer or accurate statistical discrimination. Interviews with senior HR heads at five multinationals point to organizational frictions—screener-supervisor separation and weak performance feedback—that may sustain biased fit screening.

*I would like to thank Joseph Altonji, Costas Meghir, and John Eric Humphries for serving on my committee and providing guidance and feedback. I would also like to thank Jason Abaluck, Peter Arcidiacono, Michael Dinerstein, Arpit Gupta, Patrick Kline, Scott Kominers, Eliza McNay, Suanna Oh, Evan Rose, Kamalini Ramdas, Jung Sakong, Rob Valletta, Jon Williams, seminar participants at the Yale labor workshop, the Chicago FED, Chicago Harris, Wharton (BEPP), Penn State, UCSD, SOLE, Southeastern Micro Labor Workshop, System Equitable Growth Conference at the New York Fed, the "Discrimination in the 21st Century Conference" organized by the Becker Friedman Institute and Chicago Booth, MIT Organizational Economics Lunch, and the NBER Race and Stratification Working Group for helpful feedback and support. I used OpenAI's ChatGPT and Anthropic's Claude for light copy-editing.

"Class distinctions do not die; they merely learn new ways of expressing themselves."
—R. Hoggart, introduction to George Orwell's *The Road to Wigan Pier* (1989)

"Every established order tends to produce ... the naturalization of its own arbitrariness."
—P. Bourdieu, *Outline of a Theory of Practice* (1977)

1 Introduction

Despite his wealth and polish, Jay Gatsby is dismissed in Fitzgerald's eponymous novel as "Mr. Nobody from Nowhere" (Fitzgerald, 1925). The line captures a familiar dynamic: biographical origins can shape evaluation even when credentials are strong. Elite hiring can work similarly when evaluators read cultural-capital markers—family background, linguistic polish, and social ease—as signals of merit. This paper quantifies how strongly those inferences map into hiring penalties and when those penalties are consistent with bias.

Class disparities persist in elite career outcomes (Folke and Rickne, 2025; Laurison and Friedman, 2019; Reeves and Friedman, 2024; Stansbury, 2024), yet basic questions remain unanswered: How large are class-based hiring penalties, where in multi-stage hiring processes do they emerge, and why? Even within discrimination research, we rarely observe real-world hiring pipelines closely enough to pin down magnitude, location, and mechanism (Altonji and Blank, 1999; Quillian et al., 2020). This gap is especially acute for class, where quantitative evidence remains much thinner than for race and gender (Stansbury, 2024).

This paper answers these questions using end-to-end campus recruitment data from an elite Indian college whose selectivity and global comparisons place it alongside leading Ivy League institutions (Martellini et al., 2024). I trace candidates through four sequential screening rounds and show that final-round personal interviews[1] assessing "cultural fit" account for over 95% of the caste hiring success-rate penalty. The estimated objects are caste-category penalties; I interpret them as informative about class-linked cultural capital in this setting. Converging evidence from surname timing, post-hire outcomes, large language model (LLM) analysis, and firm-side interviews indicates that these interviews screen biographical markers of cultural capital in ways that are difficult to reconcile with simple customer-discrimination predictions or with a broad class of accurate statistical-discrimination models, and are consistent with disadvantaged candidates facing a higher effective interview threshold.

[1] In this paper, "personal interview," "HR interview," and "HR-style fit interview" are used interchangeably to denote the same final open-ended, fit-oriented interview round. These terms do not refer to technical, coding, or case-study interviews.

Measuring class discrimination is difficult because class reflects intergenerationally transmitted advantage that often persists even as current income and credentials change (Clark, 2014; Erikson and Goldthorpe, 1992; Reeves and Friedman, 2024). Unlike socioeconomic status (SES), which is a point-in-time snapshot of income and education, class captures *cumulative* economic, social, and cultural capital transmitted across generations (Bourdieu, 1984; Khan, 2011; Lareau, 2003; Weber, 1978). Because this bundle is only partially observed, researchers and evaluators typically infer class from combinations of cues rather than any single SES indicator. Many of those cues are interactional and embodied.[2] Bourdieu argued that class reproduces itself not only through economic resources but through *cultural capital*[3]—markers such as accent, school pedigree, parental occupation, hobbies, and family background—that rarely map cleanly to a single variable and often surface through social interaction.

This creates a *double problem* for researchers: in most settings, class itself is measured imprecisely—should we use income, education, or parental occupation?—and even when a penalty is detected, the specific cue that triggered it often remains unclear (Lamont, 1992; Rivera, 2012, 2015). India's caste system helps resolve both problems. Caste categories are administratively codified (Bayly, 2008; Dirks, 2001; Mandal Commission, 1980) and provide a high-signal proxy for cumulative historical advantage and disadvantage, while remaining highly salient—98% of Indians identify with a caste group (Pew Research Center, 2021). This substantially reduces ambiguity both for evaluators, who may use caste as a proxy for status, and for researchers trying to identify caste discrimination empirically. Ethnographic work on elite hiring further suggests that subjective personal interviews probe biographical markers of class origin (Laurison and Friedman, 2019; Rivera, 2011), making interviews that assess "cultural fit" a plausible site where caste-linked disadvantage can operate through class-coded biographical cues.

The setting offers unusual leverage. I study recruitment by multinational consulting, technology, and financial-services firms for their India operations at an elite Indian public college, where hiring is centralized through the placement office and follows structured screening protocols. Because the institution operates under constitutionally mandated reservations, disadvantaged castes comprise roughly half of enrolled students (Madheswaran, 2008). Yet disadvantaged castes—who constitute nearly 70% of India's population—remain underrepresented in elite private-sector employment (Deshpande, 2011; Jodhka, 2015). Since these first jobs shape career trajectories for decades in the world's most populous country (Subramanian,

[2] I use *embodied* to mean dispositions and styles acquired through socialization and expressed in speech, tastes, comportment, and interaction (Bourdieu, 1986; Goffman, 1959; Khan, 2011; Lareau, 2003).

[3] For a synthetic account of Bourdieu's framework, see Swartz (1997).

2019), identifying how much of this representation gap reflects discrimination—and where and why it arises—is a first-order question for long-run inequality.

I observe over 1,000 jobs, each with four sequential rounds that move from relatively standardized to more subjective assessment: application reading, written aptitude tests, group debates, and personal interviews assessing candidate "fit." Over 95% of the roughly 30-percentage-point caste penalty in hiring success rates emerges after personal interviews. Earlier rounds show negligible caste gaps after conditioning on observables (GPA, entrance exam scores, internships, major, and degree), though Section 6.4.2 documents an earlier penalty for the surname-identifiable subgroup. To address unobserved heterogeneity in candidate quality and hiring thresholds, I examine the full distribution of job selectivity, from jobs that cut few applicants before offers to jobs that cut many. The post-personal-interview penalty remains sizeable across this distribution and shows little systematic attenuation, including in the most selective jobs where estimates are less precise, indicating that disparities are not confined to low-threshold positions or weaker candidate pools.

The claim is not that fit interviews reveal no valuable information, but that whatever they reveal is difficult to reconcile with the magnitude and pattern of disparities. Disadvantaged-caste hires have approximately 20% higher promotion rates and 28% lower attrition rates. A natural objection is inframarginality (Ayres, 2002): conditional on being hired, disadvantaged candidates may simply be unusually strong. If that were the full explanation, performance gaps should narrow in the most selective jobs, where hires are closer to common margins. Instead, promotion and attrition advantages remain broadly stable across selectivity tiers, including the most selective jobs. This pattern is more consistent with disadvantaged candidates facing a higher effective interview threshold.

Three broad explanations matter for interpreting the hiring penalty: accurate statistical discrimination, customer discrimination, and bias. Point estimates for post-hire outcomes move in the same direction in both client-facing and non-client-facing roles, making a standalone customer-discrimination account less plausible. The pattern is also hard to reconcile with accurate statistical discrimination: if employers held correct productivity beliefs (in means and/or variances), post-hire advantages should tend to narrow as jobs become more selective.[4] Instead, disadvantaged workers outperform on retention and promotion across the selectivity distribution.

To triangulate this quantitative evidence, I draw on interviews with senior HR heads who oversee campus recruitment in India at Tata Consultancy Services, IRIS Software Group,

[4]This attenuation prediction assumes monotone screening on expected productivity, reasonably well-behaved productivity distributions, and broadly stable task composition across selectivity tiers.

Tata Motors, Aditya Birla Group, and Jindal South West Steel. These firms span IT services, enterprise software, manufacturing, steel, and diversified conglomerate businesses, with substantial operations in India and internationally. Across them, interviewees emphasize retention and promotion as key metrics signaling hire quality. I use this limited qualitative sample for institutional triangulation rather than stand-alone identification, but it strengthens confidence that the post-hire patterns are difficult to reconcile with accurate statistical discrimination.

Taken together, the body of evidence is consistent with disadvantaged candidates facing a higher effective interview threshold, which can reflect taste-based discrimination, inaccurate beliefs about group productivity, or both. It does not by itself mechanically rule out every interview-specific latent trait first revealed in one-on-one interviews. But any such account must explain a demanding bundle of facts jointly: near-zero pre-interview gaps even in highly selective jobs, a sharp discontinuity at the personal-interview stage, immediate penalties when caste is already legible from surnames, later penalties when surnames are ambiguous and biographical probing begins, broadly similar patterns in client-facing and non-client-facing roles, and higher promotion with lower first-job leaving among disadvantaged hires.

An important mechanism question is what interviews are screening: directly observable caste labels, or class-like inference from biographical markers of cumulative social position. Put differently, do disparities for most candidates arise from overt labels, or when interviews make cultural-capital cues legible? I use two complementary pieces of evidence to argue that, in this setting, cultural-capital screening is an important channel through which caste becomes legible to evaluators, especially for candidates whose surnames are not initially caste-identifiable. First, because observing actual interview content at scale is rarely feasible (Quillian et al., 2020), I analyze nearly 1 million words from public mock HR-interview and group-debate transcripts. These transcripts are used to characterize widely consumed preparation materials, not as direct transcripts of the focal employers' realized interviews. These YouTube materials are widely used by candidates preparing for similar entry-level jobs in India and have cumulatively garnered hundreds of millions of views. At the *dominant-content level*, HR interviews are overwhelmingly about the candidate (roughly 90%, versus under 1% for group debates), whereas group debates are overwhelmingly about policy and society. At the *topic level*, HR interviews repeatedly elicit biographical markers of cultural capital—hometown, schooling, hobbies, family background, and parental occupation—while group debates instead center on policy topics such as economic policy, education policy, governance and democracy, technology, and artificial intelligence.

The same interviews I conducted with senior HR heads provide a direct triangulation

check for whether the transcript evidence reflects actual screening content. Across all five firms, interviewees explicitly describe final rounds as fit assessments ("fit," "cultural fit," or "organizational fit"). Within that shared framing, a subset explicitly reports biographical probes on family background or schooling, while others emphasize rapport, communication, and motivation checks. As one interviewee put it, "what's the fit in the organization, that's the basic." Another stated, "In the HR round, we do ask for family backgrounds and other things," and a third confirmed, "[Do things like hobbies and schooling come up in conversation?] Yes, it does."[5] While fit assessments may in principle aim to improve matching, triangulating these interviews with linked post-hire outcomes shows that disadvantaged-caste hires have lower attrition and higher promotion—the same downstream metrics interviewees emphasize as indicators of hire quality. This pattern is also consistent with evidence on caste-linked screening and exclusion in Indian corporate labor markets (Jodhka, 2008; Jodhka and Newman, 2007; Jodhka, 2015; Shakthi, 2023), and is less consistent with a pure productivity-screening interpretation.

The second piece of evidence ties interview content to caste revelation. While some surnames are caste-identifying, most are ambiguous in urban-educated settings because migration, "Sanskritization,"[6] and regional naming conventions blur name-based inference (e.g., Jodhka, 2015; Mamidi, 2011; Srinivas, 1966; Wiltshire, 2020). In this setting, only about 15% of students have caste-identifiable surnames. Among this 15%, penalties emerge at application reading; for the 85% with caste-ambiguous surnames, disparities emerge only after personal interviews. This timing pattern is consistent with penalties arising when interview interactions make caste legible. Combined with the transcript evidence, it is most consistent with cultural-capital screening as a central channel through which evaluators discern caste in "fit" interviews.

Why might such bias persist? My interviews with senior HR heads point to two organizational frictions. First, campus hires are recruited into central pools rather than matched to specific managers, creating weak interviewer-supervisor alignment ("the placements are so random, that it's almost like lottery"; "there is no line manager to student alignment"). Second, recruiters are evaluated mainly on process metrics (targets, turnaround time, join rates), not downstream hire performance. Interviewees also noted that post-hire outcomes are difficult to trace back to recruitment because team assignment, onboarding, and workplace

[5]The phrase "you can't expect them to read encyclopedias and Enid Blyton" in one interview similarly signals how social background can enter communication judgments; Enid Blyton was a widely read twentieth-century British children's author.

[6]"Sanskritization" refers to mobility processes in which lower-status *jāti* (localized endogamous caste) groups adopt names, customs, and practices associated with higher-status groups to claim upward status mobility (Srinivas, 1966).

dynamics intervene. This separation and weak feedback may reduce incentives to correct biased fit screening.

Scholars disagree on whether caste is best understood as analogous to race, ethnicity, class, or as a distinct institution that partially overlaps all three.[7] In this paper, I use the class framework in a *measurement-focused, not ontological* sense. Three observations support that choice. First, caste penalties are most pronounced at interview stages that elicit biographical markers—especially for the roughly 85% of candidates with caste-ambiguous surnames—rather than at first visual exposure in group debates; these markers—hometown, schooling, and parental occupation—are precisely those that leading theories of stratification treat as constitutive of class (Blau and Duncan, 1967; Bourdieu, 1986; Erikson and Goldthorpe, 1992; Lareau, 2003; Weber, 1978). Second, in the Supreme Court of India's landmark decision in *Indra Sawhney v. Union of India* (1992), the Court states that "[a] caste can be and quite often is a social class in India" (Supreme Court of India, 1992). Third, the *Mandal Commission* defined disadvantage through bundled indicators of occupation, education, housing, marriage patterns, and social standing designed to capture cumulative social position, and used a weighted scoring rule over social, educational, and economic indicators to classify groups as backward (Galanter, 1984; Mandal Commission, 1980). Together, these observations justify treating caste as a high-signal measure of cumulative class position in this setting, clarifying what the empirical results identify and grounding the contribution claims that follow. Stated narrowly, the paper identifies caste-based disadvantage operating through class-coded biographical signals in elite hiring, rather than class discrimination in a fully general sense.

Related literature. This paper contributes to discrimination research by providing, to my knowledge, some of the first large-scale evidence from major multinationals that localizes a caste penalty within a real hiring pipeline and links it to cultural-capital screening in subjective interviews. Discrimination research rarely observes multi-stage hiring deeply enough to recover mechanism salience or adjudicate whether this channel reflects efficient screening or bias (Bertrand and Duflo, 2017; Small and Pager, 2020). Using stage-by-stage outcomes, surname timing, mock-transcript evidence on interview content, and post-hire comparisons, I show that interview-stage penalties are difficult to reconcile with simple

[7] This debate is longstanding. Ghurye observed that "we do not possess a real general definition of caste," and that any single definition is difficult because of the phenomenon's complexity (Ghurye, 1969). Inden similarly notes the absence of a universally accepted definition (Inden, 2001). Competing traditions emphasize ritual hierarchy (Dumont, 1966), colonial-administrative reification (Bayly, 2008; Dirks, 2001), ethnicity-like *jāti* boundaries (Béteille, 1996), and socioeconomic stratification in labor and status terms (Bourdieu, 1986; Deshpande, 2011; Weber, 1978).

customer-discrimination predictions or broad accurate statistical-discrimination accounts, and are consistent with a higher effective interview threshold for disadvantaged candidates. Over 95% of penalties emerge at personal interviews—a margin canonical audit and correspondence designs cannot observe (Quillian et al., 2020). Much of the evidence closest to the present question quantifies caste discrimination through names-based audit or correspondence designs, which are well suited to detecting differential treatment at first contact (Thorat and Attewell, 2007). The contribution here is to show how caste-linked disadvantage can instead arise in high-stakes hiring when names are often ambiguous and "fit" assessments make caste legible through class-coded biographical markers, consistent with a long literature on the persistence and modern transformation of caste in India (Béteille, 1965; Srinivas, 1966; Deshpande, 2011; Munshi, 2019; Jodhka, 2015; Mosse, 2020) and with more recent evidence that surname-based identifiability is often incomplete in urban settings (Wiltshire, 2020; Bhupatiraju et al., 2024). This implies that upstream anonymization policies, such as CV blinding, may reduce initial gaps yet still miss substantial downstream discrimination generated in interaction-heavy interviews. Relative to existing caste labor-market research (Asher et al., 2024; Cassan et al., 2021; Deshpande, 2011; Munshi, 2019), the contribution here is granular pipeline evidence that quantifies where penalties arise within hiring and links penalties in subjective interviews to cultural-capital screening as a quantitatively salient mechanism.

Second, the paper contributes to class-discrimination research on hiring and screening, where evidence remains thinner than in race and gender (Stansbury, 2024). A central challenge is that cultural-capital screening is difficult to identify and quantify: the relevant cues are interactional, multidimensional, and typically unobserved in administrative hiring data (Goffman, 1959; Savage, 2015). Researchers therefore rely on point-in-time proxies—income, education, and parental occupation—that imperfectly capture cumulative intergenerational position (Blau and Duncan, 1967; Erikson and Goldthorpe, 1992). Canonical sociological scholarship conceptualizes cultural capital as embodied, interactional, and relational, and links it to class reproduction and evaluative boundary-drawing (Bourdieu, 1984, 1986; Lamont and Lareau, 1988; Lareau, 2003; Savage et al., 2015; Swidler, 1986). Building on this literature, I combine caste as a measurement device for cumulative social position (Bayly, 2008; Mandal Commission, 1980), stage-wise decomposition of where penalties emerge, and LLM analysis of interview content. This design identifies when class-coded biographical markers become legible, quantifies how strongly penalties concentrate within the pipeline, and connects those patterns to downstream outcomes.

Third, the paper contributes to personnel and organizational economics by providing field evidence on the allocative consequences of discretionary fit interviews in high-stakes hiring.

Research on person-organization fit and enculturation shows that organizations value cultural matching (Chatman, 1991; Goldberg et al., 2016; O'Reilly et al., 1991), while personnel research continues to debate the predictive value of unstructured interviews relative to structured or mechanical approaches (Benson and Shaw, 2024; Highhouse, 2008; Hoffman and Stanton, 2024; Kuncel et al., 2013; Sackett et al., 2022). This creates a core tension: enculturation can facilitate coordination, but discretionary fit screening—including referral-based channels that reproduce the social composition of incumbent workforces (Friebel et al., 2023)—can also reward social similarity rather than job-relevant capability (Srivastava et al., 2018). The evidence here speaks directly to that tension. In this setting, the stage with the greatest discretionary fit evaluation is also where disparities concentrate, even though post-hire outcomes do not indicate weaker performance among disadvantaged hires. This pattern is consistent with fit-intensive screening placing weight on social-similarity criteria that are imperfectly aligned with productivity-relevant allocation, consistent with broader evidence that subjective evaluation systems can distort allocation and advancement (Cowgill, 2020; Hoffman et al., 2018; Sarsons, 2017).

Fourth, the paper contributes to mobility and stratification debates by showing how downstream hiring discretion can offset upstream equalization. Admissions-side policies can broaden access to elite colleges (Karabel, 2005; Stevens, 2009), including affirmative action in this setting (Bagde et al., 2016; Madheswaran, 2008; Yenmez and Sönmez, 2022), but evidence from elite labor markets shows that the conversion of elite education into top jobs and incomes remains unequal by social background (Jodhka and Naudet, 2019; Zimmerman, 2019; Michelman et al., 2022)—a pattern potentially reinforced by exclusionary preferences that forgo gains from inclusion to preserve in-group advantage (Imas et al., 2025). In these data, that disconnect is concrete: representation is broadened at college entry through affirmative action, yet disparities remain concentrated at the final interview stage of the hiring pipeline. The pattern is consistent with frameworks in which cultural transmission and group-based social interactions can sustain inequality despite formal access expansion (Bisin and Verdier, 2025; Cassan et al., 2021; Durlauf, 2006). The implication is that admissions-side equalization alone may not translate into durable mobility without changes in downstream hiring.

2 Background: India's Administrative Caste Categories as a Measure of Social Class

Scholars disagree on whether caste is best understood as race, ethnicity, class, or a distinct institution that overlaps all three (Bayly, 2008; Dirks, 2001; Dumont, 1966; Ghurye, 1969). For this paper, I use a measurement-focused rather than ontological framework. The empirical object is the modern administrative classification observed in the data. That is the right layer for this design because the mechanism of interest is interview-stage biographical screening, and those categories record cumulative social position in a form the state has already codified. Employers do not observe the admissions classification directly; they infer status through names, CVs, and interview answers. The design therefore asks when caste-linked status becomes legible through those cues rather than assuming direct employer knowledge of the administrative category. Appendix A develops the broader conceptual discussion; Table A.1 provides a comparative map of caste, class, and race across mechanism-relevant dimensions.

2.1 India's Administrative and Legal Framework

The English word "caste" bundles distinctions that Indian scholarship keeps separate. *Jāti* refers to localized endogamous communities, while *varna* denotes a broader classificatory order (Basham, 1954; Béteille, 1996; Cohn, 1987). Modern policy categories are administrative aggregations, not literal translations of the classical four-*varna* schema. Many communities now coded as SC or ST lie outside that canonical frame, and OBC coding is heterogeneous rather than reducible to a single *varna* bin (Bayly, 2008; Cohn, 1987; Dirks, 2001; Galanter, 1984).

For this paper, the key point is administrative aggregation. The observed layer in the data is the modern state classification—SC, ST, OBC, and General/Unreserved—rather than community-level *jāti* coding or inferred *varna* status. The bridge from historical hierarchy to modern administration is explicit: the 1936 Scheduled Castes Order and later constitutional and statutory arrangements converted social hierarchy into state-recognized policy categories (His Majesty in Council, 1936; Bayly, 2008; Galanter, 1984; Miller, 1919; Rajya Sabha Committee on Personnel, Public Grievances, Law and Justice, 2005). Table 1 summarizes the distinction between conceptual layers and the measurement layer used in the empirical analysis. Appendix A, Figures A.1–A.4, provide archival schedule evidence for this terminology.

The constitutional and policy architecture reinforces why the administrative layer is the

Table 1: **Conceptual Layers, Administrative Coding, and Empirical Use**

Layer	**What it denotes**	**Observed in data?**	**Role in this paper**
Jati	Localized endogamous communities tied to region, occupation, and local history	No	Underlying social identities. Not coded community-by-community in the hiring data.
Varna	Broad classical/status discourse (*Brahmin*, *Kshatriya*, *Vaishya*, *Shudra*)	No	Conceptual backdrop only. Not a one-to-one map to SC/ST/OBC/-General; SC and ST often fall outside the canonical four-*varna* schema, and OBC is heterogeneous.
Administrative categories	State-verified policy categories (SC, ST, OBC, General/Unreserved)	Yes	Legally verified classification observed in admissions records and linked to hiring outcomes. This is the main measurement layer used in the empirical analysis.
Empirical coding used here	*Disadvantaged*= SC, ST, or OBC; *Advantaged*= General/Unreserved	Yes	Regression coding used to identify where disparities emerge in the hiring pipeline. This is a measurement choice, not an ontological claim that caste is reducible to class.

Notes: For the *jati*/*varna* distinction, see Basham (1954); Béteille (1996); Cohn (1987). For administrative codification and aggregation, see Bayly (2008); Dirks (2001); Galanter (1984); Mandal Commission (1980). For the judicial statement that caste can function as social class in India, see Supreme Court of India (1992). The empirical design uses the administrative layer because it is the state-verified classification observed in the data.

relevant one for empirical work. Scheduled Castes, Scheduled Tribes, and Other Backward Classes are administered through a common framework of social and educational disadvantage rather than a literal ritual taxonomy (Bayly, 2008; Cohn, 1987; Dirks, 2001; Galanter, 1984; Mandal Commission, 1980). The Mandal Commission operationalized that framework using a composite of social, educational, and economic indicators, with heavier weight on the first two (Galanter, 1984; Mandal Commission, 1980). *Indra Sawhney* then supplied the judicial articulation most relevant for this paper's measurement strategy: caste can function as social class in India (Supreme Court of India, 1992). The claim here is narrower than caste and class being identical in every sense. It is that these administratively verified categories are a disciplined proxy for cumulative disadvantage in a hiring setting where biography and social ease are screened.[8]

2.2 Academic Scholarship on Caste and Class

No single definition of caste commands the literature (Inden, 2001; Ghurye, 1969). What matters for this paper is that three strands are especially relevant. One emphasizes ritual hierarchy and endogamy (Dumont, 1966; Srinivas, 1966). A second emphasizes administrative legibility and state codification (Bayly, 2008; Cohn, 1987; Dirks, 2001). A third emphasizes materially grounded stratification through occupation, education, marriage, and mobility (Béteille, 1965, 1996; Cassan et al., 2021; Deshpande, 2011; Jodhka, 2015). This paper draws mainly on the third strand because the mechanism under study is interview-stage screening on biography and social ease rather than ritual status in the abstract.

That move also fits broader comparative theory. Béteille links *varna* to class-like ordering (Béteille, 1996). Weber emphasizes status closure (Weber, 1978). And Bourdieu's account of embodied cultural capital points directly to the kinds of cues elite interviews can surface—accent, conversational ease, familiarity with elite settings, and ways of presenting self (Bourdieu, 1986; Khan, 2011; Lareau, 2003; Lamont, 1992). This is a narrow mechanism claim, not a universal theory of caste. Appendix A develops the broader conceptual discussion.

2.3 Operational Coding of Caste in Affirmative Action Framework

Operationally, I code SC/ST/OBC candidates as *Disadvantaged* and General/Unreserved candidates as *Advantaged*. That matches the administrative categories observed in the data

[8] "Creamy layer" excludes relatively better-off OBC households from reservation benefits in public-sector admissions and jobs while group classification remains OBC (Galanter, 1984; Supreme Court of India, 1992). In this private-sector setting, OBC creamy-layer candidates still face substantial interview penalties; see Section 2.3.

while using them as proxies for cumulative social position in private-sector hiring. This is a measurement choice, not an ontological claim that caste is reducible to class. Because affirmative action in India applies to public education and government jobs, not private-sector multinational hiring (Galanter, 1984), the mechanism documented here operates in a domain where reservation rules do not directly constrain employer screening. Disaggregated estimates support this coding: even OBC candidates above creamy-layer thresholds face substantial interview penalties, which is hard to square with an income-alone interpretation. Appendix B reports the disaggregated results.

A complementary institutional fact points in the same direction. Appendix C shows that 81% of recruiting firms have no substantive DEI page, and among the firms that do, caste is rarely mentioned. That silence does not prove the absence of informal norms. But it fits qualitative scholarship documenting employer resistance to caste-conscious policy and the continued use of merit and fit language in corporate hiring discourse (Jodhka, 2008; Jodhka and Newman, 2007; Thorat and Newman, 2010).

3 Key Terms and Institutional Setting

This section defines the terms used throughout the paper and describes the institutional setting: an on-campus job fair at an elite Indian college, anonymized under the data agreement. The setting is representative of elite hiring in India. Top-ranked colleges use similar standardized placement procedures, elite employers recruit from the same set of campuses, graduates predominantly pursue elite private-sector positions, and student participation is near-universal. Appendix D documents this representativeness evidence.

3.1 Definitions

Five terms are central.

a. **Disadvantaged castes.** Scheduled Castes (SCs), Scheduled Tribes (STs), and Other Backward Classes (OBCs), as defined in Section 2.3.

b. **Advantaged castes.** Castes not classified as SC, ST, or OBC, commonly referred to as the "General Category" in India's administrative framework.

c. **Job.** A firm-designation pair (e.g., Google software engineer).

d. **Elite jobs.** Entry-level positions at multinationals, major national corporations, or unicorn startups (privately held startups valued at over $1 billion).

e. **Caste-based affirmative action.** Constitutional reservations requiring government jobs and educational institutions to allocate up to 50% of seats or positions to SC, ST, and OBC candidates combined at central institutions (some states exceed this ceiling); these reservations do not apply to private-sector positions.

3.2 The Job Fair

The elite college's job fair occurs entirely on campus and has five phases.

a. **Pre-Placement.** Firms deliver pre-placement talks between June and mid-August to gauge interest and finalize roles.

b. **Return offers from internships.** By late August, firms extend pre-placement offers to summer interns. Unlike in the U.S., these play a minor role: fewer than 5% of students accept and exit the placement process.

c. **Student and employer registration.** Students register by late August; employers submit job roles, compensation, and vacancies by early September. Compensation is non-negotiable: the placement office verifies that offer letters match advertised pay, and non-compliance can restrict future campus visits. Comparable rules and timelines are publicly documented across elite Indian colleges (for example, IIT Delhi Office of Career Services and Delhi University Placement Cell).

d. **Placement.** Starting in mid-September, students apply for jobs. Employers then screen applications, administer written aptitude tests and group debates, and conduct final-round personal interviews.

e. **Post-Placement ("aftermarket").** From late January until graduation, the placement office runs an aftermarket for students without offers. These employers are typically lower-paying startups using less standardized recruitment.

4 Data and Sample

The main dataset comes from the elite college's centralized job portal and spans four placement cycles (exact years anonymized). It covers 4,164 students pursuing technical degrees and competing for more than 1,000 positions in software engineering, quantitative consulting, data analytics, and manufacturing at multinational firms.

Data Source. Students upload standardized profiles to the portal, including GPA, entrance-exam scores, coursework, and internship experience. Employers upload shortlists after each screening stage, generating end-to-end documentation from application through final offer. Registered students must participate in all four screening rounds, and the placement office verifies that offer-letter salaries match the non-negotiable compensation schedules filed by firms.

Student Sample. The 4,164 students pursue four technical degrees: Bachelor of Technology (B.Tech.), Dual Degree, Master of Technology (M.Tech.), and Master of Science (M.S.). Admissions are based exclusively on standardized national entrance examinations with caste- and major-specific cutoffs; there is no subjective admissions component. Approximately 90% of students are male, which is typical of elite technical colleges in India (Datta, 2017). Appendix O, Table O.1 reports the degree breakdown.

Caste Measurement and Representation. Caste-based affirmative action in admissions yields approximate parity in the overall sample: 48% disadvantaged and 52% advantaged. Reservation status is not inferred; it comes from administrative records and government certificates used to claim reserved seats (Galanter, 1984). OBC creamy-layer status is documented separately through income verification, and I refer to those students as OBC-CL. I still treat OBC-CL students as disadvantaged because the empirical question concerns cumulative social position rather than current income alone. Appendix B shows that OBC-CL candidates continue to face substantial interview penalties.

Researcher Observes State-Verified Caste Categories; Recruiters Do Not. The placement records include state-verified caste categories, but recruiters are not directly shown that information during screening. For candidates with caste-ambiguous surnames, recruiters must infer caste from other biographical markers rather than observe it directly. The 15% surname-identifiable share should therefore be read as a lower bound on pre-interview caste legibility, because recruiters may also infer caste from first name, hometown, schooling language, accent, or communication style.

Summary Statistics for Students. Disadvantaged-caste students score lower on GPA and entrance exams, but both groups span the full performance distribution (Appendix N, Figure N.1). All empirical analyses condition on these differences. Caste is also not strongly associated with internship sector: among undergraduates, sector distributions are nearly

identical across caste groups (about 25% technology, 37% consulting, and 38% manufacturing), with similar patterns for postgraduate degrees. Appendix O, Table O.2 reports the full summary statistics by caste and degree.

5 Institutional Features of the Hiring Process

Three institutional features are central for interpretation: applications are low-cost, the recruitment sequence is standardized, and compensation is fixed within job. Together, these features make it possible to localize where disparities arise and to interpret job-quality gaps as differences in access rather than wage negotiation.

Applications. Applying to a job requires only clicking on the employer's listing in the centralized portal, so students apply broadly within eligibility constraints.

Four-Stage Recruitment Pipeline. Once applications are submitted, jobs follow the same four-stage sequence: (1) *application screening*, where employers typically use GPA and profile thresholds to filter resumes; (2) *written aptitude tests*, standardized assessments of technical and analytical skills; (3) *group debates*, structured discussions involving 20–30 students divided into two teams with common prompts; and (4) *personal interviews*, roughly 30-minute final-round evaluations emphasizing interpersonal skills, organizational fit, and cultural fit.

Recruiters observe standardized profile fields from the application stage onward, including name, GPA, entrance-exam rank, school marks, and internships. Group debates add real-time communication cues and evaluator discretion, which makes their near-zero caste gap substantively informative: the late break is not just a generic artifact of in-person or subjective assessment. Personal interviews add direct biographical probing—for example, hometown, schooling trajectory, family background, and parents' occupation—which is precisely where caste salience can intensify for surname-ambiguous candidates.

Selectivity is substantial throughout the process. On average, jobs receive about 209 applications, 125 candidates advance to written tests, 102 to group debates, 26 to interviews, and about 6 receive offers. Appendix O, Table O.4 reports the stage-by-stage counts.

Standardized, Non-Negotiable Compensation. Salaries and benefits are standardized, non-negotiable, and uniform within job. As a result, any earnings gaps in the data reflect differential access to jobs rather than idiosyncratic within-job wage setting.

Internship Return Offer Policy. Students who accept internship return offers are withdrawn from the placement cycle by the placement office and are excluded from the main analysis.

India-Based Positions and Indian Recruiters. About 99% of positions are India-based. Recruitment and interviewing are therefore conducted almost entirely by Indian hiring managers.

6 Caste Penalties Emerge at "Fit" Interviews

The central empirical result of the paper is simple: the caste penalty is generated overwhelmingly at personal interviews, not at earlier standardized screens. This section establishes that claim in five steps. Section 6.1 shows that the pooled caste gap opens sharply at interviews. Section 6.2 then shows that this late-stage penalty does not systematically attenuate across the full distribution of job selectivity, from low-cutoff to high-cutoff jobs. Section 6.3 shows that disadvantaged hires are subsequently promoted more and leave first jobs less, with similar qualitative patterns in client-facing and non-client-facing roles. Section 6.4 then ties the timing of the penalty to biographical probing and caste revelation using transcript evidence and surname timing. Section 6.5 brings those facts together to assess accurate statistical discrimination, customer discrimination, and bias, and uses firm-side interviews to document why biased fit screening can persist. Throughout, the qualitative interviews with senior HR heads are used for institutional triangulation; full theme-ordered excerpts appear in Appendix F.

Identification interpretation. The selectivity and post-hire exercises rely on three maintained assumptions, formalized in Appendix I and Appendix H. First, conditional on observed credentials and fixed effects, higher cumulative cuts proxy tighter screening thresholds rather than fundamentally different disadvantaged and advantaged applicant pools across jobs; Section 6.2 supports that reading by showing near-zero pre-interview effects even in the highest-cutoff jobs. Second, promotion and first-job retention are informative post-hire outcomes for match quality, conditional on controls and fixed effects. Third, the evidence rules out broad pre-interview cumulative-deficit stories that should surface in high-cutoff jobs, while not claiming to mechanically rule out every interview-specific trait first observed in round 4.

6.1 Caste Penalties in Hiring

Success Rates. The success-rate analysis uses the 4,005 students—96% of the 4,164 graduating cohort—who entered the centralized placement system. For each stage, the outcome is whether a candidate advances to the next round for a given job.

To estimate stage-specific caste gaps, I use:

$$\text{Advance}_{ijt} = \alpha + \beta \cdot \text{Disadvantaged Caste}_i + \mathbf{X}_i'\gamma + \delta_j + \theta_t + \epsilon_{ijt}, \tag{1}$$

where Advance_{ijt} equals 1 if student i advances to the next stage for job j in cohort-year t; δ_j are job fixed effects; θ_t are cohort-year fixed effects; and $\mathbf{X}_i$ includes GPA, entrance-exam percentile, 10th/12th marks, internship sector, and department/degree fixed effects. The coefficient β is the within-job caste gap in advancement probability at that stage.

Table 2 reports the results. Pre-interview differences are economically negligible: application screening is -0.25pp ($p < 0.10$), written tests are -0.02pp (n.s.), and group debates are -0.05pp (n.s.). The sharp break appears only at personal interviews, where the disadvantaged-caste penalty is -29.38pp ($p < 0.01$). In absolute value, the interview-stage coefficient accounts for essentially all of the stage-level penalty mass.[9]

Table 2: **Caste Differences in the Success Rate at Each Job Search Stage**

	App → Test (1)	Test → Debate (2)	Debate → Int (3)	Int → Offer (4)
Disadv. Caste (OLS)	-0.0025	-0.0002	-0.0005	-0.2938***
	(0.0018)	(0.0023)	(0.0051)	(0.0107)
Mean Success Rate	0.5990	0.8110	0.2580	0.2610
Observations	305,541	183,150	148,489	38,378
Adjusted R^2	0.61069	0.34360	0.01858	0.24450
Year FE	✓	✓	✓	✓
Job FE	✓	✓	✓	✓
Other Controls	✓	✓	✓	✓

Notes: Table 2 reports the disadvantaged-caste coefficient from Equation 1, estimated separately for each stage. Each observation is a student-job-year combination for candidates still in contention at that stage, and the dependent variable is an indicator for advancing to the next stage. Controls include college GPA, entrance exam score (JEE rank percentile), 10th and 12th grade marks, prior internship sector indicators, job fixed effects, department fixed effects, degree-stream fixed effects, and cohort-year fixed effects; continuous scores enter linearly. When firms conduct common screening rounds for multiple positions, each position is treated as a separate job. Standard errors are two-way clustered at the student and job levels. *$p < 0.1$; **$p < 0.05$; ***$p < 0.01$.

[9] I operationalize the "interview share" as the absolute value of the interview-to-offer coefficient divided by the sum of absolute values of all four stage-specific coefficients (App→Test, Test→GD, GD→Int, Int→Offer). Each coefficient is the disadvantaged-caste estimate from the stage-specific linear probability model. This mechanical decomposition is what the text summarizes as "over 95%."

Robustness and sample selection. The interview penalty is nearly identical in client-facing and non-client-facing jobs (Appendix O, Table O.5), which is hard to square with customer-discrimination stories that predict materially larger penalties in client-facing roles. The result is also robust to category disaggregation (Appendix B, Table B.1), flexible controls (Appendix G, Table G.1), and Oster bounds (Appendix H, Table H.1).

Selection into the success-rate sample is more likely to attenuate the estimated interview penalty than to exaggerate it. Appendix O, Table O.6, Panel A, shows that disadvantaged students who never enter the job fair are negatively selected on GPA, while disadvantaged hires within the same job have slightly *higher* GPAs. Excluding non-participants should therefore bias the estimated interview penalty toward zero, not away from it.

Job Quality. Success rates localize where screening gaps arise, but they weight each job equally. To add an economic margin, I measure job quality as the average posted salary of the jobs each candidate remains in contention for at each stage. Because salaries are standardized and non-negotiable within jobs (Section 3.2), these gaps reflect differential access to jobs rather than bargaining.

To estimate stage-specific gaps in job quality, I use:

$$\log(\overline{\text{Salary}}_{it}) = \alpha + \beta \cdot \text{Disadvantaged Caste}_i + \mathbf{X}_i'\boldsymbol{\gamma} + \theta_t + \epsilon_{it}, \tag{2}$$

where $\overline{\text{Salary}}_{it}$ is the average posted salary across the jobs student i remains eligible for at a given stage in cohort-year t; θ_t are cohort-year fixed effects; and $\mathbf{X}_i$ includes the same controls as Equation 1. Job fixed effects are omitted because the dependent variable is an average across jobs.

Figure 1 mirrors the success-rate pattern: job-quality gaps are negligible through tests and debates and become sharply negative only at interviews and offers. Appendix O, Table O.7 reports the full stage coefficients, and Appendix O, Table O.8 shows similar controlled offer-stage gaps in client- and non-client-facing roles. The estimates are again robust to category disaggregation, flexible controls, and Oster bounds (Appendix B, Table B.2; Appendix G, Table G.2; Appendix H, Table H.2). Selection into the job-quality sample also appears more likely to attenuate estimated gaps than to inflate them: Appendix O, Table O.6, Panel B, shows that excluded disadvantaged students are negatively selected on GPA.

Figure 1: **Job Quality Penalty by Hiring Stage**

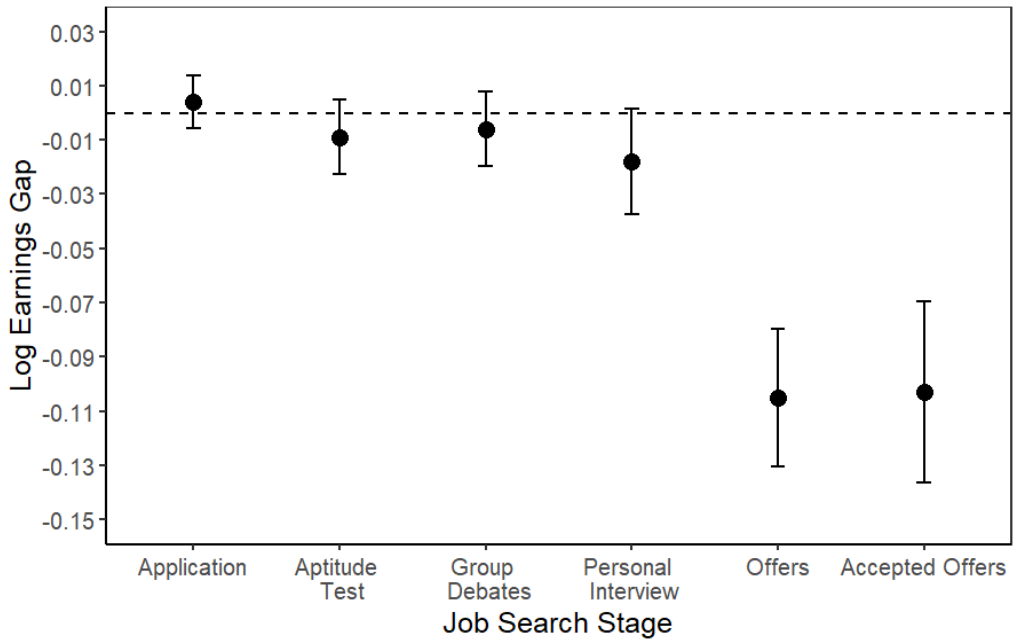

Notes: Figure 1 reports the disadvantaged-caste coefficient from Equation 2 at each hiring stage. Job quality is measured as the average posted salary of the jobs a candidate remains in contention for at that stage. Because salaries are standardized and non-negotiable within jobs, these gaps reflect differential job access rather than within-job pay variation. Black dots show 95% confidence intervals. Controls include college GPA, entrance exam score (JEE rank percentile), 10th and 12th grade marks, prior internship sector indicators, department fixed effects, degree-stream fixed effects, and cohort-year fixed effects. Standard errors are HC1 robust for App/Test/Debate/Interview/Offer and job-clustered for Choice. Job fixed effects are omitted because the dependent variable is a student-stage average across jobs.

6.2 Heterogeneity: Penalties Show Little Attenuation Across Job Selectivity

Section 6.1 documented the average stage pattern. This section asks whether the interview-stage result could instead reflect broad cumulative ability deficits rather than discriminatory screening. The selectivity data say no. Pre-interview penalties remain near zero across the full selectivity distribution, while a large interview penalty appears and shows little systematic attenuation. Section 6.4.2 discusses the key subgroup exception: candidates with caste-identifiable surnames can face earlier penalties.

We observe advance/not-advance decisions rather than stage-level scores. Jobs vary widely in selectivity (Appendix O, Table O.9), so higher-cutoff jobs probe deeper into the applicant pool. That is enough for the interpretation here: the argument does not require equality of latent ability across caste groups, only that conditional on observables and fixed

effects the pre-interview pool not diverge systematically across jobs in a way that would mechanically generate a late-stage pattern. Under broad cumulative-deficit stories, some pre-interview caste gap should emerge as cutoffs rise. It does not. The common-pool interpretation used below is narrower than equality of talent across caste groups: conditional on observed credentials and fixed effects, firms do not appear to be fishing from systematically different disadvantaged and advantaged pools across jobs before interviews. The absence of any pre-interview gap, even in high-cutoff jobs, is what supports that interpretation.[10]

To examine heterogeneity across job selectivity, I estimate separately for each job-stage:

$$\text{Advance}_{ijt} = \alpha_j + \beta_j \cdot \text{Disadvantaged Caste}_i + \mathbf{X}_i'\gamma_j + \theta_t + \epsilon_{ijt}, \tag{3}$$

where β_j is the within-job caste gap in advancement probability at that stage, conditional on observables. I then plot each β_j against the job's cumulative cut rate—the fraction of applicants eliminated by that stage.

Figure 2 shows two facts. First, application, written-test, and group-debate coefficients stay near zero across the selectivity distribution. Second, the interview penalty is roughly 30 percentage points and remains large across the observed selectivity range. Appendix N, Figure N.2 shows the same qualitative pattern in raw gaps: modest early gaps shrink through written tests and debates, then reopen sharply at interviews. That is difficult to reconcile with a broad cumulative ability deficit, because the candidates who survive to interviews in more selective jobs are positively selected on earlier screens.

The model intuition is correspondingly simple. Conditional on observables, more selective jobs can be read as applying stricter pre-interview cutoffs to candidates drawn from the same underlying applicant pools. If disadvantaged candidates had a broad pre-interview ability deficit, stricter cutoffs should eventually expose it, so some negative pre-interview caste gap should emerge as selectivity rises. It does not. Appendix I formalizes this logic and shows how a higher effective interview threshold for disadvantaged candidates can instead generate the observed late-stage penalty together with positive selection among those ultimately hired.

[10]This evidence is also hard to reconcile with early-round compensatory favoritism toward disadvantaged candidates: when caste is directly legible from surnames, penalties already appear at application reading (Section 6.4.2), not as early-stage advantages.

Figure 2: **Caste Penalty in the Conditional Success Rate At Each Screening Round**

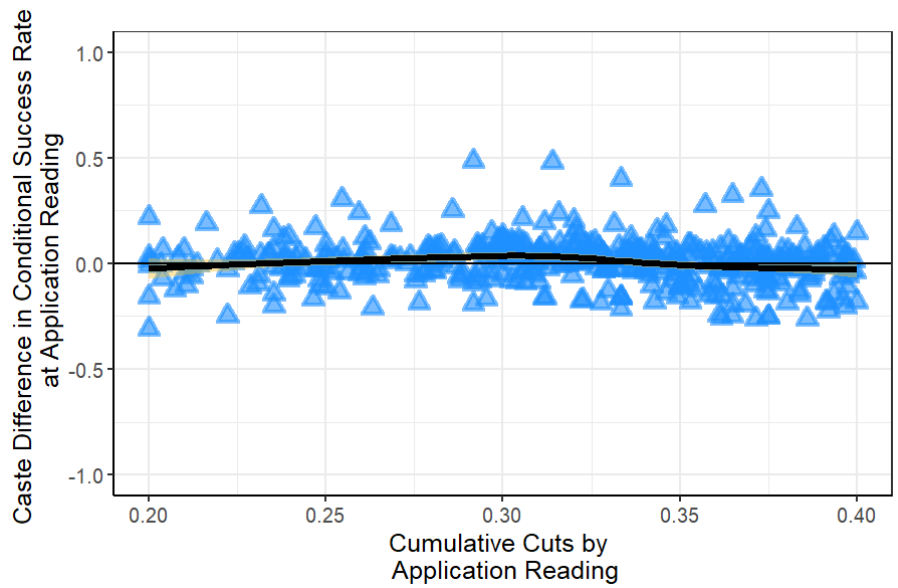

(a) Caste Penalty in Conditional Success Rates versus Cumulative Cuts by Application Reading

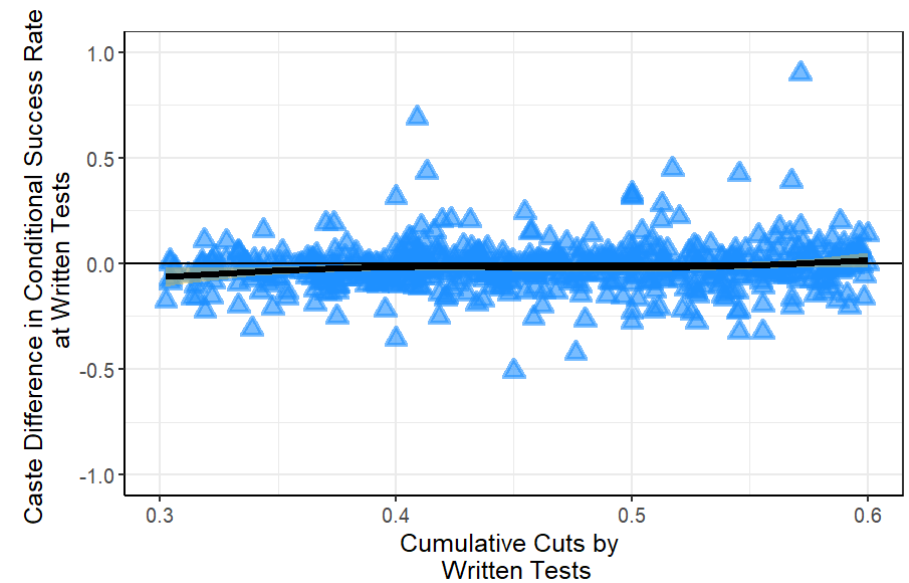

(b) Caste Penalty in Conditional Success Rates versus Cumulative Cuts by Written Tests

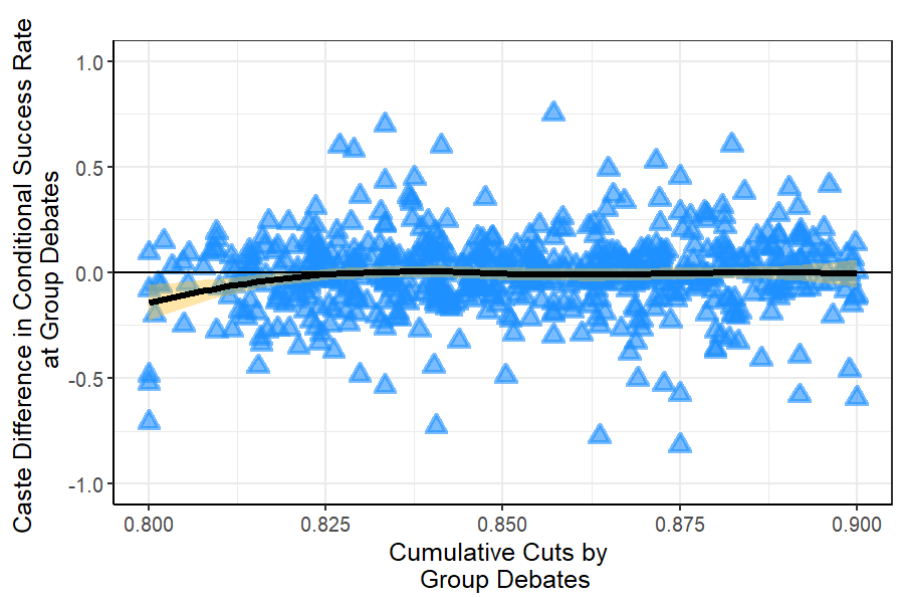

(c) Caste Penalty in Conditional Success Rates versus Cumulative Cuts by Group Debates

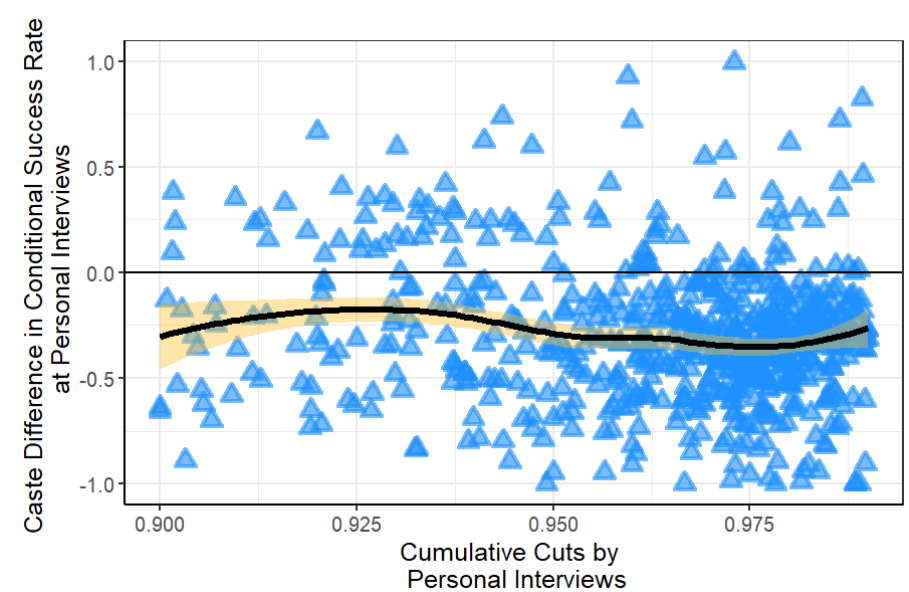

(d) Caste Penalty in Conditional Success Rates versus Cumulative Cuts by Personal Interviews

Notes: Figure 2 plots the disadvantaged-caste coefficient from Equation 3 against the stage-specific cumulative cut rate for each job. Each blue triangle is a single job; negative values indicate lower advancement rates for disadvantaged castes. The x-axis is the cumulative proportion of applicants eliminated by the end of that stage (not the stage-conditional cut). Black lines show LOESS smooths with 95% confidence bands. Controls include college GPA, entrance exam score (JEE rank percentile), 10th and 12th grade marks, prior internship sector, department fixed effects, degree-stream fixed effects, and cohort-year fixed effects; collinear dummies are dropped in single-year jobs. The key pattern is near-zero caste penalties at all pre-interview stages across the selectivity distribution, but a large interview-stage penalty with little systematic attenuation across the observed selectivity range. This is also the empirical pattern that motivates the common-pool interpretation used later in Appendix I: conditional on observables, firms do not appear to be fishing from systematically different disadvantaged and advantaged applicant pools before interviews.

6.3 Caste Differences in Promotion and Attrition

Disadvantaged hires are more likely to be promoted and less likely to leave their first job than advantaged peers, measured within the first employer over up to five years post-graduation. Read together, those two margins matter more than either one alone: the post-hire evidence does not rest on attrition alone, because promotion and retention move in the same direction. The claim is not that interviews reveal nothing useful. It is that these downstream patterns are difficult to reconcile with disadvantaged hires being weaker on productivity-relevant traits once the localization of the penalty is taken into account.

Why promotion and first-job retention remain informative. Promotion and first-job retention are imperfect proxies for realized worker quality, so I use them as triangulation rather than stand-alone identification. Their force here comes from four features. First, inference does not rest on attrition alone: both margins move in the same direction, with disadvantaged hires promoted more and leaving less. Second, both outcomes are measured within the first employer, which reduces confounding from cross-firm mobility. Third, promotion is identified from within-firm title progression and validated using salary changes and benchmark tenure-to-title ladders (Wilmers et al., 2025). Fourth, the HR interviews reported below indicate that firms themselves treat retention, progression, and promotion as central indicators of campus-hire quality. Relative to most discrimination studies, which cannot follow hiring to downstream firm outcomes at all, this is unusually informative post-hire evidence (e.g., Hoffman et al., 2018; Quillian et al., 2020). Turnover is also a firm-relevant margin because replacement, onboarding, and training are costly (Bloom et al., 2012; Doeringer and Piore, 1971; Minni, 2024).

Specification. To examine post-hire performance, I estimate:

$$Y_{ijt} = \alpha + \beta \cdot \text{DisadvCaste}_i + \mathbf{X}_i'\gamma + \delta_j + \theta_t + \epsilon_{ijt} \tag{4}$$

where Y_{ijt} is an indicator for promotion (or leaving first job) for individual i in job j and cohort t; δ_j are job fixed effects; θ_t are cohort-year fixed effects; and $\mathbf{X}_i$ includes GPA, entrance exam scores, prior internship sector, and department and degree fixed effects. For promotion outcomes, $\mathbf{X}_i$ also includes job tenure, so the tenure-controlled estimate captures the promotion advantage conditional on equal time at the firm.

Promotions and Tenure. Table 3, Panel A, shows that disadvantaged hires have 14–20% higher promotion rates than advantaged peers in the pooled sample. Point estimates remain

positive across selectivity tiers and across client-facing and non-client-facing jobs, with no systematic attenuation as selectivity rises. Panel B shows that disadvantaged hires are 17.0 percentage points less likely to leave their first job. That retention advantage is also directionally stable across selectivity tiers and job types. Results are robust to functional-form checks (Appendix G, Table G.3) and Oster bounds (Appendix H, Table H.3).

Differential Attrition Is Unlikely to Inflate the Estimated Gap. A natural concern is that selective exit among advantaged workers could mechanically inflate the promotion gap. Appendix O, Table O.10 addresses this directly. Within the same job, hired disadvantaged candidates have higher GPAs, leavers have slightly lower GPAs than stayers, and the leaver-stayer GPA gradient does not differ meaningfully by caste. Together, these estimates make it unlikely that differential attrition is generating the promotion advantage.

Interview Triangulation on Quality Metrics. To triangulate interpretation, I use semi-structured interviews with senior HR heads from five recruiting firms: Tata Consultancy Services, IRIS Software Group, Tata Motors, Aditya Birla Group, and Jindal South West Steel. This is qualitative triangulation rather than stand-alone identification; theme-ordered excerpts appear in Appendix F. Across firms, interviewees treat retention, performance progression, and promotion as central indicators of campus-hire quality. One senior HR head described the campus team as being evaluated on "the quality of the overall campus hires and their retention and their cultural fit and their growth in the organization." The firms differ in vocabulary, but not in the underlying point: these are quality-relevant outcomes rather than automatic post-hire transitions. That matches the quantitative evidence here.

Stylized Model: Inframarginality, Latent Productivity, and Higher Interview Thresholds. An important identification concern is inframarginality: even with common interview thresholds, post-hire differences can arise from who survives selection at the margin. Appendix I develops a stylized threshold model (with fuller derivations in Appendix I.1) to make this logic explicit. Under equal thresholds, any latent productivity disadvantage should attenuate as job selectivity rises, because both groups' hires are drawn from increasingly similar right tails. By contrast, with a higher effective interview threshold for disadvantaged candidates, stable post-hire advantages can persist across selectivity tiers because disadvantaged hires are selected from deeper in the right tail. The empirical pattern matches this higher-threshold case more closely than an equal-threshold latent-productivity account.

Table 3: **Caste Differences in Promotion and Job Leaving**

	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)	(9)	(10)	(11)
	All	All	Top 75%	Top 50%	Top 40%	Top 30%	Top 20%	Top 15%	Top 10%	Client	Non-Client
					By Job Selectivity					*By Job Type*	
Panel A: Outcome = Promoted											
Disadv. Caste	0.137***	0.096***	0.104***	0.112**	0.136***	0.147*	0.190**	0.204**	0.084	0.107**	0.082**
	(0.026)	(0.030)	(0.033)	(0.049)	(0.051)	(0.075)	(0.076)	(0.093)	(0.122)	(0.053)	(0.037)
Mean Y	0.684	0.684	0.667	0.668	0.676	0.711	0.711	0.706	0.73	0.687	0.682
Observations	2,489	2,489	2,012	1,246	997	748	502	374	252	835	1,654
Adjusted R^2	0.183	0.205	0.203	0.216	0.211	0.182	0.192	0.189	0.127	0.194	0.214
Panel B: Outcome = Left First Job											
Disadv. Caste	-0.186***	-0.170***	-0.130***	-0.160***	-0.210***	-0.187**	-0.207**	-0.145	-0.132	-0.078	-0.215***
	(0.044)	(0.046)	(0.048)	(0.053)	(0.060)	(0.088)	(0.099)	(0.122)	(0.156)	(0.085)	(0.046)
Mean Y	0.605	0.605	0.616	0.648	0.655	0.662	0.675	0.668	0.659	0.592	0.611
Observations	2,489	2,489	2,012	1,246	997	748	502	374	252	835	1,654
Adjusted R^2	0.215	0.216	0.24	0.247	0.214	0.179	0.085	0.084	0.054	0.236	0.207
Job FE	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
Year FE	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
Other Controls		✓	✓	✓	✓	✓	✓	✓	✓	✓	✓

Notes: Table 3 reports disadvantaged-caste coefficients from Equation 4 for graduates placed through the main placement process. Panel A uses promotion within five years as the outcome; Panel B uses leaving the first job. Each observation is one graduate's first placed job. Column (1) includes job and cohort-year fixed effects. Column (2) adds controls: college GPA, entrance exam score (JEE rank percentile), 10th and 12th grade marks, prior internship sector indicators, department fixed effects, and degree-stream fixed effects; Panel A also controls for job tenure, so the estimate captures the promotion gap conditional on equal time at the firm. Columns (3)–(9) split by job selectivity; columns (10)–(11) split by client-facing status. Standard errors are clustered at the job level. *$p < 0.1$; **$p < 0.05$; ***$p < 0.01$.

6.4 The Caste Penalty Tracks Cultural-Capital Screening and Caste Revelation

Sections 6.1–6.3 establish a higher effective bar at personal interviews rather than generic subjectivity. That distinction matters because group debates are already live, interpersonal, and partly discretionary, yet they do not generate a comparable caste penalty. The remaining mechanism question is what that bar is screening. I use two tests. First, I compare what personal interviews and group debates actually probe. Second, I use surname timing to examine when caste becomes legible. If personal interviews disproportionately extract biography, and penalties appear earlier when surnames already reveal caste but later when they do not, the combined timing supports caste revelation through class-coded biography.

6.4.1 LLM Analysis of Mock Interview and Group Debate Transcripts

Group debates and personal interview rounds are both subjective and in-person, but only personal interviews produce large caste penalties in the hiring data. Because the administrative records do not contain firms' interview scripts, I use mock transcripts as a characterization of widely consumed preparation material rather than as direct transcripts of the focal employers' realized interviews. I therefore treat this evidence as comparative triangulation, not stand-alone identification of the firms' exact scripts. I analyze 274 live mock transcripts (126 personal interviews, 148 group debates) totaling more than 870,000 words, after excluding coaching-only and purely technical sessions.[11] Appendix J documents the prompt design, taxonomy, transcript construction, sampling rules, and robustness checks. The HR interviews discussed below provide a direct firm-side plausibility check.

Dominant-content and topic-level evidence. Each transcript is classified as "about speaker," "about society," or "mixed," with biography coded only when it is the topic of exchange rather than rhetorical context in a policy discussion. Personal interviews are overwhelmingly about the candidate (89.7% coded "about speaker"), whereas group debates are overwhelmingly about society (99.3% coded "about society"). The same split replicates under Claude Sonnet 4 and `gpt-5.2` (Appendix J, Tables J.4 and J.5). Figure 3 visualizes this contrast.

[11] Earlier working drafts used a smaller live subset (183 transcripts: 129 personal interviews and 54 group debates). The current live corpus (274: 126 personal interviews, 148 group debates) reflects expanded group-debate ingestion, stricter live-versus-tips reclassification, and exclusion of personally conducted HR-interview files from the analysis dataset.

Figure 3: **Content Focus: Personal Interviews vs. Group Debates**

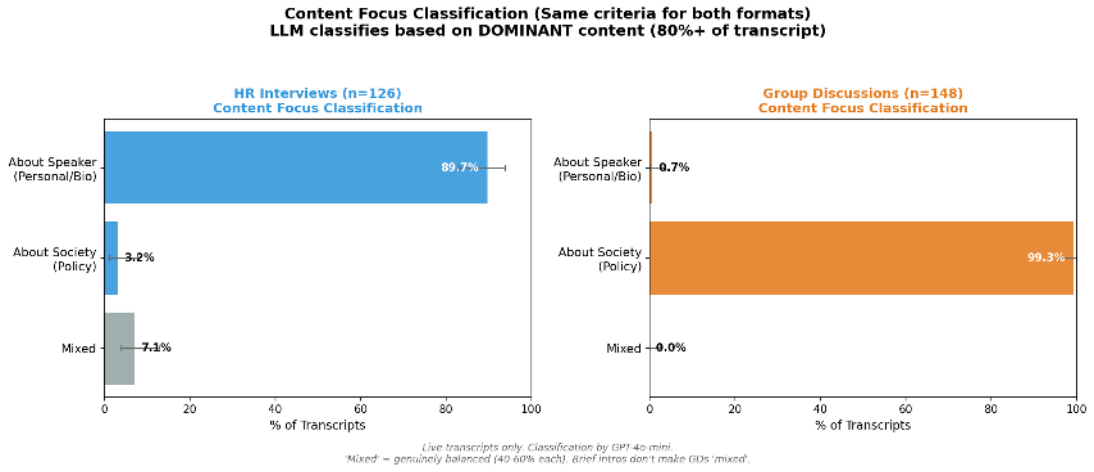

Notes: Figure 3 classifies 274 live transcripts (126 personal interviews, 148 group debates) from YouTube mock-interview preparation videos by dominant content: "About Speaker," "About Society," or "Mixed." Classification was performed using GPT-4o-mini with structured JSON outputs; the prompt distinguishes biography *as the topic* from personal anecdotes used rhetorically in policy debate. Personal interviews overwhelmingly focus on the candidate, whereas group debates overwhelmingly focus on policy and society. Error bars show 95% Wilson confidence intervals. Results are robust to alternative sampling strategies and to cross-LLM validation using Claude Sonnet 4 and `gpt-5.2`. See Appendix J for full methodology, prompts, and robustness figures.

Figure 4 applies a common 20-topic taxonomy and shows the same asymmetry at finer resolution. Personal interviews repeatedly elicit hometown, schooling, hobbies, family structure, and parents' occupation. Group debates concentrate on policy topics such as economic policy, education policy, governance and democracy, and technology and artificial intelligence. To keep the comparison conservative, biography is coded only when it is itself the topic of questioning, not when personal experience is used as context in debate. I also exclude "candidate's college/university education" from the biographical-marker interpretation because every student has already cleared elite-college admission, even though that topic appears in 84.9% of personal interviews and 6.8% of group debates.

Taken together, the dominant-content and topic-level evidence point in the same direction: personal interviews are the round that systematically extracts biography, while group debates are mostly policy-debate settings. That is exactly the kind of asymmetry one would expect if caste becomes legible primarily when interviewers can probe class-coded biography among candidates whose surnames are not initially informative.

Figure 4: **Topics Probed in Personal Interviews vs. Group Debates**

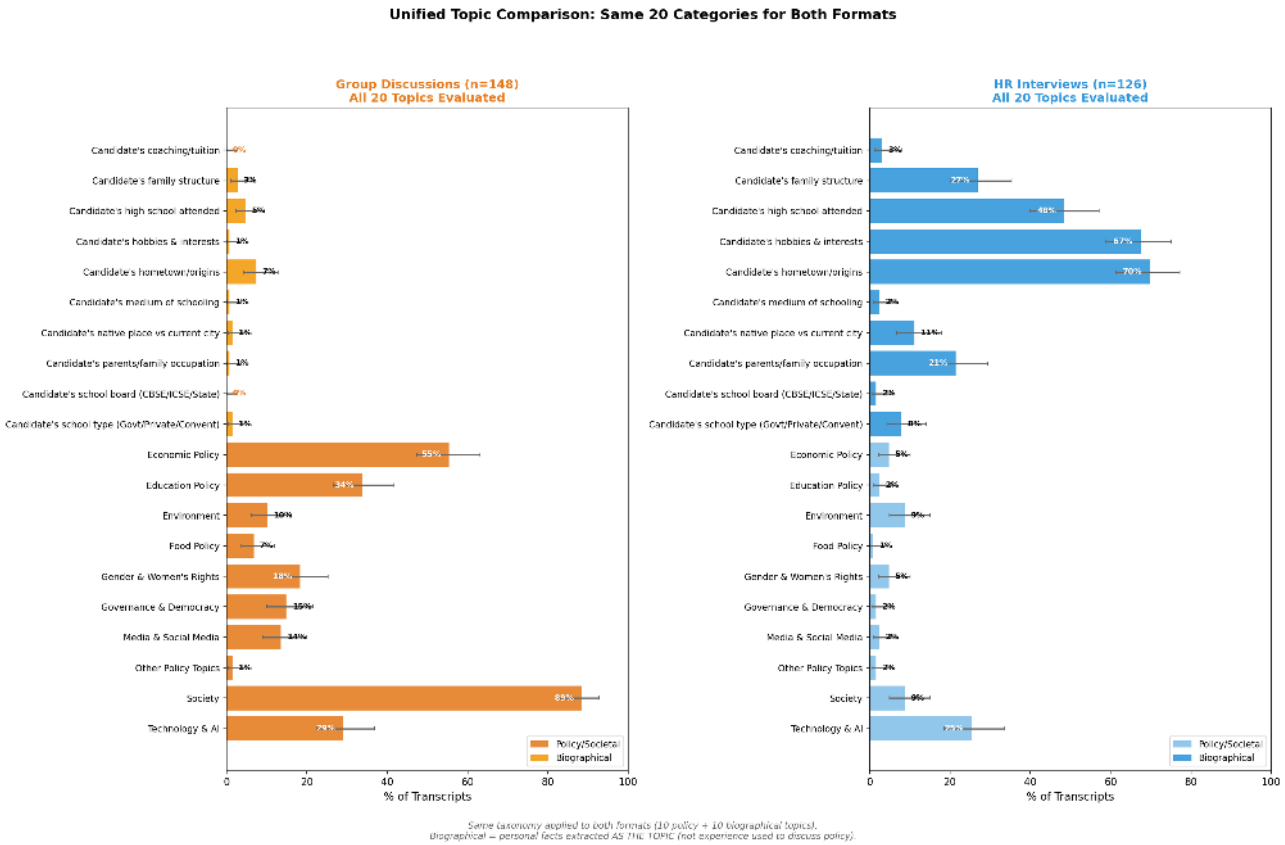

Notes: Figure 4 shows the share of transcripts in which each topic appears, using an identical 20-category taxonomy for personal interviews and group debates. The sample comprises 274 live transcripts (126 personal interviews, 148 group debates) from YouTube preparation videos. The prompt distinguishes biography *as the topic* (e.g., "Where are you from?") from personal experience used as rhetorical *context* in policy debate. This avoids false positives in group debates, where anecdotes often support abstract arguments rather than reveal biography. "Candidate's college/university education" is excluded from the biographical-marker interpretation because all students in this sample already attend an elite college; for reference, it appears in 84.9% of personal interviews and 6.8% of group debates. Error bars show 95% Wilson confidence intervals. See Appendix J for the full taxonomy, prompts, sampling strategies, and cross-LLM robustness checks using Claude Sonnet 4 and `gpt-5.2`.

Firm-side triangulation and bounded interpretation. Interviews I conducted with senior HR heads at Tata Consultancy Services, IRIS Software Group, Tata Motors, Aditya Birla Group, and Jindal South West Steel provide a direct institutional check. Across firms, final rounds are described as fit-oriented, and several interviewees explicitly mention family background, schooling, hobbies, or related probes. One senior HR head at Tata Motors put it bluntly: "Definitely in the HR round, we do ask for family backgrounds... what basically is the cultural fit of this person." This pattern aligns with scholarship on Indian corporate labor markets, where caste-linked screening persists within meritocratic organizational language (Jodhka, 2008; Jodhka and Newman, 2007; Jodhka, 2015; Shakthi, 2023), and with evidence

from elite U.S. firms, where final-round fit interviews are often biographical and rapport-based rather than purely technical (Rivera, 2012, p. 1006, 1008); see also Lamont (2009) and Turco (2010).

The transcript evidence is nevertheless indirect. These are candidate-consumed mock preparation sessions rather than proprietary transcripts of the focal employers' realized interviews, and LLM coding cannot recover every in-room signal such as posture or tone. Its role in the paper is therefore comparative and triangulatory, not stand-alone identification. That bounded use is still informative. If generic in-person cues in group debates were sufficient to reveal caste among surname-ambiguous candidates, penalties should appear there too. They do not. By contrast, when caste is already legible from surnames, penalties appear at application reading. In a setting with weak formal caste-DEI signaling (Appendix C), the later emergence of penalties when biographical probing intensifies is more consistent with caste revelation than with generic subjectivity.

6.4.2 Surname Identifiability Analysis

This second test asks when caste penalties appear as caste becomes legible. If penalties emerge at application reading for caste-identifiable surnames but only after personal interviews for non-identifiable surnames, timing identifies caste revelation rather than generic late-stage subjectivity. Most surnames are ambiguous in urban settings (e.g., Jodhka, 2015; Parmar, 2020).[12] Under a conservative administrative-list coding, about 15% of students have surnames that unambiguously identify caste status.

Coding and main timing result. A surname is coded as *caste-identifiable* only when it provides an unambiguous advantaged/disadvantaged signal based on official administrative lists and corroborating sources. This conservative protocol likely understates employers' true ability to infer caste from names. Appendix K, Table K.1 gives illustrative mappings, and Appendix K reports the full source hierarchy and rules.

[12]See also Bhupatiraju et al. (2024)—a study of more than one million High Court cases from the urban center of Patna—which shows that roughly half of petitioners, respondents, and judges use *caste-neutral* surnames, underscoring how common name-based ambiguity is in urban India.

Figure 5: **The Caste Job Quality Penalty by Surname Identifiability**

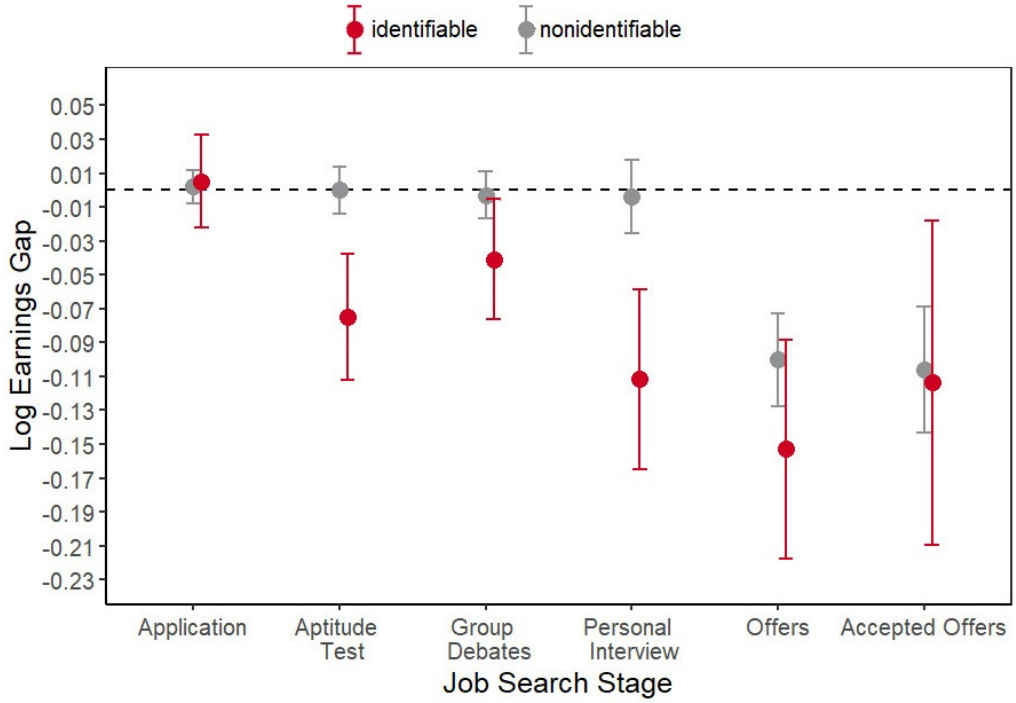

Notes: Figure 5 reports the disadvantaged-caste coefficient from Equation 2, estimated separately for students with caste-identifiable surnames (red, n=619) and non-identifiable surnames (gray, n=3,386). Each dot shows the caste gap in log average salary of jobs in contention at that stage, with 95% confidence intervals. Because salaries are standardized and non-negotiable within jobs, gaps reflect differences in job access rather than within-job pay variation. All specifications include the baseline controls from Equation 2. For identifiable surnames, penalties emerge by the written-test stage—before any in-person interaction. For non-identifiable surnames, no penalty appears until after personal interviews. Full regression results are in Appendix O, Table O.11, Panels A and B; Oster bounds in Appendix H, Table H.4, Panels A and B; and functional-form robustness in Appendix G, Table G.4, Panels A and B, and Table G.5, Panels A and B.

Figure 5 shows the main timing split in job quality: identifiable surnames diverge by the written-test stage, while non-identifiable surnames diverge only after personal interviews. The success-rate splits tell the same story even more sharply: in the identifiable sample, penalties appear already at application reading; in the non-identifiable sample, they appear only after interviews. Appendix O, Tables O.11 and O.12 report the full coefficients.

Selection and interpretation. A natural concern is that the identifiable-surname sample may differ in baseline productivity. Appendix L, Figures L.1 and L.2 address that directly by comparing raw and controlled application-stage gaps. Controls remove the early gap in the non-identifiable sample but not in the identifiable sample. Appendix L, Table L.1 then shows that the caste–GPA slope is nearly identical across surname-identifiability groups, so the early divergence for identifiable surnames is not just a GPA-composition artifact.

For identifiable surnames, later-round coefficients are often more negative but imprecise; the key result is early emergence, not strict monotonic deepening. For non-identifiable

surnames, the absence of a group-debate penalty is equally informative. Group debates may reveal class-linked style, accent, or confidence, but in this setting those cues are not sufficiently unambiguous and actionable to generate systematic penalties for the surname-ambiguous majority. Personal interviews, by contrast, allow one-on-one probing across multiple linked dimensions—schooling, hometown, family background, hobbies, and communication style—which makes caste revelation much more plausible. Together with the transcript evidence, this timing pattern is most consistent with cultural-capital screening that makes caste legible.

6.5 Discrimination Channels: Statistical, Customer, and Bias

Given the interview-stage penalty, the selectivity pattern, and post-hire outcomes, three explanations are central: accurate statistical discrimination, customer discrimination, and bias. Customer-discrimination accounts predict stronger penalties in client-facing roles; instead, the interview penalty and post-hire patterns are directionally similar across client-facing and non-client-facing jobs (Sections 6.1 and 6.3). Accurate statistical-discrimination accounts predict attenuation of post-hire gaps as selectivity rises; instead, promotion and retention advantages remain broadly stable across the selectivity distribution (Sections 6.2 and 6.3). The stylized model in Appendix I formalizes this selectivity logic for a broad class of mean- and variance-based screening stories; it narrows those interpretations sharply, while not mechanically ruling out every conceivable higher-moment model without additional structure. Section 6.4 then shows that penalties track caste visibility via surnames and biographical probing in personal interview rounds, rather than generic productivity signals observed at earlier stages. Taken together, these patterns are difficult to reconcile with simple customer-discrimination predictions and with a broad class of accurate statistical-discrimination accounts. They are consistent with a higher effective interview threshold for disadvantaged candidates; that threshold can reflect taste-based discrimination, reservation-stigma or other stereotype-based inference, inaccurate beliefs about group productivity, or some combination of these channels. The evidence does not mechanically rule out every interview-specific latent trait, but any such trait would need to operate in a very particular way: it would have to matter sharply in one-on-one fit interviews but not earlier standardized screens or group debates, align with the surname-timing evidence on caste legibility, appear similarly in client-facing and non-client-facing jobs, and still leave disadvantaged hires with higher promotion and lower first-job leaving.

Qualitative Triangulation on Why Bias Can Persist. To triangulate why this bias channel can persist organizationally, I use the same interviews with senior HR heads at Tata Consul-

tancy Services, IRIS Software Group, Tata Motors, Aditya Birla Group, and Jindal South West Steel. These interviews are a limited qualitative sample and not stand-alone identification; full theme-ordered excerpts appear in Appendix F. Both frictions appear across all five firms, though they manifest differently: some firms describe random pool-based assignment, others panel-based hiring with post-hoc function matching; some cite turnaround-time metrics, others offer-drop ratios.

First, interview screeners are often separated from eventual supervisors because campus hires enter central pools and are assigned later. As one interviewee put it,

> *"So it's a pool which is hired... which person will go to which location is not very firmed up... it's not the direct manager who is interviewing them."* (Senior HR head, Tata Motors)

Another described the same structure more starkly:

> *"The placements are so random, that it's almost like lottery... there's no interviewer matching done when placement happens."* (Senior HR head, Tata Consultancy Services)

Second, recruiters face weak accountability for downstream performance. One interviewee stated,

> *"For talent acquisition [teams], it's mainly the TAT [turnaround time]... [t]heir performance is not linked to the performance of the hired candidates."* (Senior HR head, Tata Motors)

and another noted,

> *"It is almost impossible to track back their performance, their appraisal or their promotion to the recruitment branch."* (Senior HR head, Tata Consultancy Services)

These interview accounts do not identify causal effects on their own, but they are consistent with an agency wedge between interview screening and downstream performance feedback, which can allow biased fit assessments to persist.

7 Alternative Explanations and External Validity

This section addresses the main alternative explanations and then considers how far the findings should travel beyond one elite Indian institution.

Alternative explanations. No single diagnostic settles interpretation on its own. The paper's claim rests on the coherence of the joint pattern. If disadvantaged candidates had broad cumulative ability deficits, pre-interview gaps should persist or widen in more selective jobs. Instead, pooled pre-interview gaps remain near zero even at high cutoffs (Section 6.2). Nor does the pattern look like generic discretion or in-person subjectivity: group debates are already live, interpersonal, and partly discretionary, yet they do not generate a comparable caste penalty. If a broad class of accurate statistical-discrimination models explained the interview penalty, post-hire differences should attenuate as thresholds rise. Instead, promotion and first-job-retention advantages remain broadly stable across selectivity tiers (Section 6.3; Appendix M). If customer discrimination were primary, penalties should be materially larger in client-facing jobs. Instead, the key hiring and post-hire patterns appear in both client-facing and non-client-facing roles (Section 6.5). And any remaining selection-on-unobservables story must still explain the timing evidence: earlier penalties when surnames reveal caste and later penalties when surnames are ambiguous until interviews begin extracting biography.

Residual interview-specific latent traits. A reasonable skeptic may still ask whether one-on-one interviews reveal some genuine latent trait not well proxied by GPA, entrance scores, or group debates. The paper's cleanest claim is therefore not a full decomposition of whether interviewers respond to taste, reservation stigma, or mistaken beliefs. It is that caste-based disadvantage operates through class-coded biographical cues at the fit-interview stage. But a benign latent-trait account has to explain more than a generic late-stage break. It must explain why pre-interview gaps remain near zero even in highly selective jobs, why penalties appear immediately when caste is already legible from surnames, why they appear later when surnames are ambiguous and interviews begin eliciting biography, why the pattern looks similar in client-facing and non-client-facing roles, and why disadvantaged hires then exhibit both higher promotion and lower first-job leaving. That bundle is increasingly difficult to reconcile with a standard productivity-screening story and is more consistent with a higher effective interview threshold for disadvantaged candidates.

A related possibility is reservation backlash: evaluators may react negatively to quota discourse even where quotas do not govern these private-sector jobs (Jodhka, 2008; Jodhka and Newman, 2007). That possibility is not inconsistent with the paper's interpretation, but the timing evidence still points to information revelation rather than blanket hostility from the outset. Penalties emerge immediately when surnames reveal caste and otherwise emerge mainly when interview content makes caste legible. Institutional context fits that interpretation. Appendix C shows that 81% of sampled recruiting firms (826/1,017) have no

substantive DEI page, and among firms with substantive DEI pages, caste appears in only 5.1% of India-headquartered pages and 1.0% of foreign-headquartered pages (Figure C.1). This formal-policy silence is consistent with weak caste salience in corporate DEI language and may partly reflect exclusionary preferences that resist broadening diversity categories beyond those with established institutional legitimacy (Imas et al., 2025). Appendix M reports additional discussion and checks.

External validity. The data come from one institution, but several features support broader relevance. First, the institution is globally elite in STEM (Martellini et al., 2024). Second, Section 5 documents that the campus placement architecture is standardized across top Indian colleges, and Appendix D provides supporting evidence. Third, the stage sequence—resume screening, standardized tests, and discretionary final interviews—closely matches recruitment pipelines used by multinationals elsewhere (Bloom et al., 2012; Rivera, 2015). Fourth, evidence from the U.S. and U.K. likewise finds class-linked penalties concentrated at subjective fit gates (Laurison and Friedman, 2019; Stansbury and Rodriguez, 2024), while India's civil service shows analogous personality-interview penalties paired with post-hire outperformance (Bhavnani and Lee, 2021). Finally, Folke and Rickne (2025) document an analogous "class ceiling" in politics: working-class candidates face discretionary promotion barriers throughout the pipeline even after controlling for ability, motivation, and voter support. Because this paper studies technical roles with relatively standardized assessments, the estimates here are plausibly conservative relative to less standardized hiring environments.

8 Conclusion

This paper opens the hiring pipeline at multinational recruiters in India and shows that the caste penalty is generated overwhelmingly at one place: the final personal interview. Earlier standardized screens account for little of the gap. The same localization appears in job quality, not just in success rates. The central fact of the paper is therefore not merely that disadvantaged candidates receive worse outcomes, but that the penalty is concentrated in the fit-oriented stage of elite hiring.

The rest of the evidence points to a coherent interpretation. The interview penalty shows little systematic attenuation across selectivity tiers. Disadvantaged hires then exhibit higher promotion rates and lower first-job leaving, with similar qualitative patterns in client-facing and non-client-facing roles. And the mechanism evidence lines up with the timing of the penalty: personal interviews systematically elicit biography, while surname-identifiable

candidates are penalized earlier and surname-ambiguous candidates are penalized later, when biography is probed. Taken together, these facts support a higher effective interview threshold for disadvantaged candidates operating through class-coded biographical cues.

The paper does not recover the precise primitive inside that threshold. The evidence cannot separately identify taste, reservation stigma, or inaccurate beliefs about group productivity. But it does rule out much broader alternatives. The pattern is not what one would expect from a broad pre-interview productivity deficit, from a simple customer-discrimination story concentrated in client-facing roles, or from an accurate-statistical-discrimination account whose implications should weaken as selectivity rises. The contribution is therefore to localize the penalty, show how it becomes legible, and narrow the class of interpretations that fit the data.

The organizational implication is immediate. These interviews are not simply neutral tiebreakers after objective screening. In this setting they are the main site at which disparities are created. Interviews with senior HR heads suggest one reason such screening can persist: the people conducting campus interviews are often separated from eventual supervisors and weakly accountable for downstream performance. That diagnosis also points to reform. Structured interviews with standardized questions and scoring rubrics can reduce the scope for discretionary bias while preserving valid assessment (Campion et al., 1997; Schmidt and Hunter, 1998; Levashina et al., 2014). Linking interviewer assessments to later outcomes can strengthen accountability (Mero and Motowidlo, 1995; Castilla, 2015).

India's caste system makes this process unusually observable, but the broader lesson is not India-specific. The heterogeneity analysis in Section 6.2 also suggests a natural next step: linking variation in firm-specific penalties to concrete organizational practices. Wherever elite hiring relies on subjective assessments of who "fits," evaluators can screen on biographical signals that track cumulative social advantage—accent, schooling trajectory, family background, cultural references, and social ease (Bourdieu, 1986; Durlauf, 2006; Laurison and Friedman, 2019; Reeves and Friedman, 2024; Rivera, 2015; Stansbury and Rodriguez, 2024). What caste adds here is measurement: it bundles those advantages into administratively legible categories, making visible a mechanism that is harder to study elsewhere. In that sense, the paper's broader implication is straightforward: late-stage fit assessments can convert inherited social position into outcomes interpreted as merit. "Every established order tends to produce ... the naturalization of its own arbitrariness" (Bourdieu, 1977, p. 164). Elite hiring's language of "fit" is one such naturalization.

References

Altonji, J. G. and R. M. Blank (1999). Race and gender in the labor market. *Handbook of Labor Economics 3*, 3143–3259.

Anderson, R. N. (1938). The work of the placement officer. *Teachers College Record 40*(1), 1–7.

Asher, S., P. Novosad, and C. Rafkin (2024). Intergenerational mobility in india: New measures and estimates across time and social groups. *American Economic Journal: Applied Economics 16*(2), 1–25.

Ayres, I. (2002). Outcome Tests of Racial Disparities in Police Practices. *Justice Research and Policy 4*(1-2), 131–142.

Bagde, S., D. Epple, and L. Taylor (2016). Does affirmative action work? caste, gender, college quality, and academic success in india. *American Economic Review 106*(6), 1495–1521.

Banerjee, A., E. Duflo, M. Ghatak, and J. Lafortune (2013). Marry for what? caste and mate selection in modern india. *American Economic Journal: Microeconomics 5*(2), 33–72.

Bapuji, H., K. Gupta, S. Chrispal, and T. Roulet (2023). What managers everywhere must know about caste. *MIT Sloan Management Review Winter*. Accessed: 2024-04-22.

Basham, A. L. (1954). *The Wonder That Was India: A Survey of the Culture of the Indian Sub-Continent Before the Coming of the Muslims*. London: Sidgwick and Jackson. Classic survey of pre-Islamic Indian civilization.

Bayly, S. (2008). *Caste, Society and Politics in India from the Eighteenth Century to the Modern Age*. The New Cambridge History of India. Cambridge: Cambridge University Press.

Benson, A. and K. Shaw (2024). What do managers do? *Annual Review of Economics*. Forthcoming.

Bertrand, M. and E. Duflo (2017). Field experiments on discrimination. In A. V. Banerjee and E. Duflo (Eds.), *Handbook of Economic Field Experiments*, Volume 1, pp. 309–393. Amsterdam: North-Holland.

Bertrand, M. and S. Mullainathan (2004). Are Emily and Greg more employable than Lakisha and Jamal? a field experiment on labor market discrimination. *American Economic Review 94*(4), 991–1013.

Béteille, A. (1965). *Caste, Class, and Power: Changing Patterns of Stratification in a Tanjore Village*. Berkeley: University of California Press.

Béteille, A. (1996). Caste in contemporary india. *Caste Today*, 150–179. Argues jati corresponds to ethnicity rather than caste.

Bhavnani, R. R. and A. Lee (2021). Does affirmative action worsen bureaucratic performance? Evidence from the indian administrative service. *American Journal of Political Science 65*(1), 5–20.

Bhupatiraju, S., D. L. Chen, S. Joshi, and P. Neis (2024). Caste aside? Names, networks and justice in the courts of Bihar, India. *European Journal of Empirical Legal Studies 1*(2), 151–178.

Bisin, A. and T. Verdier (2025). Economic models of cultural transmission. CEPR Discussion Paper DP20340, CEPR Press, Paris and London.

Blau, P. M. and O. D. Duncan (1967). *The American Occupational Structure*. New York: Wiley.

Bloom, N., R. Sadun, and J. Van Reenen (2012). The organization of firms across countries. *The Quarterly Journal of Economics 127*(4), 1663–1705.

Bourdieu, P. (1977). *Outline of a Theory of Practice*. Cambridge: Cambridge University Press.

Bourdieu, P. (1984). *Distinction: A Social Critique of the Judgement of Taste*. Cambridge, MA: Harvard University Press.

Bourdieu, P. (1986). The forms of capital. In J. G. Richardson (Ed.), *Handbook of Theory and*

Research for the Sociology of Education, pp. 241–258. New York: Greenwood.
BusinessLine (2018). An Inclusive India Inc. Business Line, January 13, 2018.
Campion, M. A., D. K. Palmer, and J. E. Campion (1997). A review of structure in the selection interview. *Personnel Psychology 50*(3), 655–702.
Cassan, G., D. Keniston, and T. Kleineberg (2021). A division of laborers: Identity and efficiency in india. Working Paper 28462, National Bureau of Economic Research. R&R at Econometrica.
Castilla, E. J. (2015). Accounting for the gap: A firm study manipulating organizational accountability and transparency in pay decisions. *Organization Science 26*(2), 311–333.
Chatman, J. A. (1991). Matching people and organizations: Selection and socialization in public accounting firms. *Administrative Science Quarterly 36*(3), 459–484.
Clark, G. (2014). *The Son Also Rises: Surnames and the History of Social Mobility*. Princeton, NJ: Princeton University Press.
Cohn, B. S. (1987). The census, social structure and objectification in South Asia. In *An Anthropologist Among the Historians and Other Essays*, pp. 224–254. Delhi: Oxford University Press.
Cowgill, B. (2020). Bias and productivity in humans and algorithms: Theory and evidence from résumé screening. Working Paper, Columbia University.
Datta, S. (2017). *A History of the Indian University System: Emerging from the Shadows of the Past*. Palgrave Macmillan.
Deshpande, A. (2011). *The Grammar of Caste: Economic Discrimination in Contemporary India*. New Delhi: Oxford University Press.
Dirks, N. B. (2001). *Castes of Mind: Colonialism and the Making of Modern India*. Princeton, NJ: Princeton University Press.
Doeringer, P. B. and M. J. Piore (1971). *Internal Labor Markets and Manpower Analysis*. Lexington, MA: Heath Lexington Books.
Dumont, L. (1966). *Homo Hierarchicus: The Caste System and Its Implications*. Chicago: University of Chicago Press. English translation 1970.
Durlauf, S. N. (2006). Groups, social influences, and inequality: A memberships theory perspective on poverty traps. In S. Bowles, S. N. Durlauf, and K. Hoff (Eds.), *Poverty Traps*. Princeton: Princeton University Press.
Erikson, R. and J. H. Goldthorpe (1992). *The Constant Flux: A Study of Class Mobility in Industrial Societies*. Oxford: Oxford University Press.
Fitzgerald, F. S. (1925). *The Great Gatsby*. New York: Charles Scribner's Sons.
Folke, O. and J. Rickne (2025). The class ceiling in politics. *American Political Science Review 119*(3), 1375–1392.
Friebel, G., M. Heinz, M. Hoffman, and N. Zubanov (2023). What do employee referral programs do? measuring the direct and overall effects of a management practice. *Journal of Political Economy 131*(3), 633–686.
Galanter, M. (1984). *Competing Equalities: Law and the Backward Classes in India*. Berkeley: University of California Press.
Ghurye, G. S. (1969). *Caste and Race in India* (5th ed.). Bombay: Popular Prakashan.
Goffman, E. (1959). *The Presentation of Self in Everyday Life*. Garden City, NY: Doubleday.
Goldberg, A., S. B. Srivastava, V. G. Manian, W. Monroe, and C. Potts (2016). Fitting in or standing out? The tradeoffs of structural and cultural embeddedness. *American Sociological Review 81*(6), 1190–1222.
Government of India (1950). The constitution of india (as amended). Relevant provisions:

Articles 15(4), 16(4), 340, 341, 342, and 338B.
Highhouse, S. (2008). Stubborn reliance on intuition and subjectivity in employee selection. *Industrial and Organizational Psychology 1*(3), 333–342.
His Majesty in Council (1936). The government of india (scheduled castes) order, 1936. Order in Council under the Government of India Act, 1935.
Hoffman, M., L. B. Kahn, and D. Li (2018). Discretion in Hiring. *The Quarterly Journal of Economics 133*(2), 765–800.
Hoffman, M. and C. T. Stanton (2024). People, practices, and productivity: A review of new advances in personnel economics. NBER Working Paper 32849.
Imas, A., K. Madarász, and H. Sarsons (2025). Jealousy of trade: Exclusionary preferences and economic nationalism. Working Paper 34351, National Bureau of Economic Research.
Inden, R. (2001). *Imagining India* (2nd ed.). Bloomington: Indiana University Press. Notes no universally accepted definition of caste.
Jodhka, S. S. (2008, October). Caste & the corporate sector. *Indian Journal of Industrial Relations 44*(2), 185–193. Invited article; JSTOR stable/27768188.
Jodhka, S. S. (2015). *Caste in Contemporary India*. Routledge India.
Jodhka, S. S. and J. Naudet (2019). *Mapping the Elite: Power, Privilege, and Inequality*. Oxford University Press.
Jodhka, S. S. and K. Newman (2007). In the name of globalisation: Meritocracy, productivity and the hidden language of caste. *Economic and Political Weekly 42*(41), 4125–4132.
Karabel, J. (2005). *The Chosen: The Hidden History of Admission and Exclusion at Harvard, Yale, and Princeton*. Boston: Houghton Mifflin.
Khan, S. R. (2011). *Privilege: The Making of an Adolescent Elite at St. Paul's School*. Princeton, NJ: Princeton University Press.
Kuncel, N. R., D. M. Klieger, B. S. Connelly, and D. S. Ones (2013). Mechanical versus clinical data combination in selection and admissions decisions: A meta-analysis. *Journal of Applied Psychology 98*(6), 1060–1072.
Lamont, M. (1992). *Money, Morals, and Manners: The Culture of the French and the American Upper-Middle Class*. Chicago: University of Chicago Press.
Lamont, M. (2009). *How Professors Think: Inside the Curious World of Academic Judgment*. Cambridge, MA: Harvard University Press.
Lamont, M. and A. Lareau (1988). Cultural capital: Allusions, gaps and glissandos in recent theoretical developments. *Sociological Theory 6*(2), 153–168.
Lareau, A. (2003). *Unequal Childhoods: Class, Race, and Family Life*. Berkeley, CA: University of California Press.
Laurison, D. and S. Friedman (2019). *The Class Ceiling: Why it Pays to be Privileged*. Bristol, UK: Bristol University Press, Policy Press.
Levashina, J., C. J. Hartwell, F. P. Morgeson, and M. A. Campion (2014). The structured employment interview: Narrative and quantitative review of the research literature. *Personnel Psychology 67*(1), 241–293.
Madheswaran, S. (2008). Is affirmative action policy for private sector necessary? *Indian Journal of Industrial Relations 44*(2), 164–184.
Mamgain, R. P. (2019). Growth, poverty and non-farm employment in india. *Indian Journal of Labour Economics 62*(2), 197–214.
Mamidi, P. (2011). Strategic manipulation of caste identity signals. Doctoral Dissertation, Oxford University.
Mandal Commission (1980). The Mandal Commission Report. Prime Minister's High Level

Committee, Cabinet Secretariat, Government of India.
Martellini, P., T. Schoellman, and J. Sockin (2024, February). The global distribution of college graduate quality. *Journal of Political Economy 132*(2), 434–483.
Mero, N. P. and S. J. Motowidlo (1995). The effects of rater accountability on the accuracy and impression-management bias of performance appraisal. *Journal of Applied Psychology 80*(4), 517–524.
Michelman, V., J. Price, and S. D. Zimmerman (2022). Old boys' clubs and upward mobility among the educational elite. *The Quarterly Journal of Economics 137*(2), 845–909.
Miller, L. C. (1919). Report of the committee appointed to consider steps necessary for the adequate representation of communities in the public service.
Minni, V. (2024). Global managers, local workers: Wage setting inside a multinational firm. *AEA Papers & Proceedings 114*, 586–91.
Mosse, D. (2020). The modernity of caste and the market economy. *Modern Asian Studies 54*(4), 1225–1271.
Munshi, K. (2019). Caste and the indian economy. *Journal of Economic Literature 57*(4), 781–834.
Omi, M. and H. Winant (2014). *Racial Formation in the United States* (3rd ed.). New York: Routledge.
O'Reilly, C. A., J. Chatman, and D. F. Caldwell (1991). People and organizational culture: A profile comparison approach to assessing person–organization fit. *Academy of Management Journal 34*(3), 487–516.
Oster, E. (2019). Unobservable selection and coefficient stability: Theory and evidence. *Journal of Business & Economic Statistics 37*(2), 187–204.
Parkin, F. (1979). *Marxism and Class Theory: A Bourgeois Critique*. London: Tavistock.
Parmar, D. (2020). Caste and the politics of perception in india's labor markets. *World Development 135*, 105066.
Pew Research Center (2021). Religion in india: Tolerance and segregation. Technical report, Pew Research Center.
Quillian, L., J. J Lee, and M. Oliver (2020). Evidence from field experiments in hiring shows substantial additional racial discrimination after the callback. *Social Forces 99*(2), 732–759.
Rajya Sabha Committee on Personnel, Public Grievances, Law and Justice (2005). Eighth report on the scheduled castes, scheduled tribes and other backward classes (reservation in posts and services) bill, 2004. Technical report, Rajya Sabha. Historical background on Mysore's 1918–21 policy and the Miller Committee.
Reeves, A. and S. Friedman (2024). *Born to Rule: The Making and Remaking of the British Elite*. Cambridge, MA: Belknap Press of Harvard University Press.
Rivera, L. A. (2011). Ivies, extracurriculars, and exclusion: Elite employers' use of educational credentials. *Research in Social Stratification and Mobility 29*(1), 71–90.
Rivera, L. A. (2012). Hiring as cultural matching: The case of elite professional service firms. *American Sociological Review 77*(6), 999–1022.
Rivera, L. A. (2015). *Pedigree: How Elite Students Get Elite Jobs*. Princeton, NJ: Princeton University Press.
Sackett, P. R., C. Zhang, C. M. Berry, and F. Lievens (2022). Revisiting meta-analytic estimates of validity in personnel selection: Addressing systematic overcorrection for restriction of range. *Journal of Applied Psychology 107*(12), 2040–2068.
Sarsons, H. (2017). Recognition for group work: Gender differences in academia. *AEA Papers and Proceedings 107*(5), 141–145.
Savage, M. (2015). *Social Class in the 21st Century*. Pelican Books.

Savage, M., F. Devine, N. Cunningham, M. Taylor, Y. Li, J. Hjellbrekke, B. Le Roux, S. Friedman, and A. Miles (2015). On social class, anno 2014. *Sociology 49*(6), 1011–1030.

Schmidt, F. L. and J. E. Hunter (1998). The validity and utility of selection methods in personnel psychology: Practical and theoretical implications of 85 years of research findings. *Psychological Bulletin 124*(2), 262–274.

Shakthi, S. (2023). Corporate Brahminism and tech work: Caste in a modern Indian profession. *South Asia: Journal of South Asian Studies 46*(5), 920–933.

Small, M. L. and D. Pager (2020). Sociological perspectives on racial discrimination. *Journal of Economic Perspectives 34*(2), 49–67.

Srinivas, M. N. (1966). *Social Change in Modern India*. Berkeley: University of California Press.

Srivastava, S. B., A. Goldberg, V. G. Manian, and C. Potts (2018). Enculturation trajectories: Language, cultural adaptation, and individual outcomes in organizations. *Management Science 64*(3), 1348–1364.

Stansbury, A. (2024). Social class and career progression in elite professions. *American Economic Journal: Applied Economics Forthcoming*.

Stansbury, A. and K. Rodriguez (2024). The class gap in career progression: Evidence from US academia. R&R, *Econometrica*.

Stevens, M. P. (2009). *Creating a Class: College Admissions and the Education of Elites*. Harvard University Press.

Subramanian, A. (2019). *The Caste Of Merit: Engineering Education in India*. Harvard University Press.

Supreme Court of India (1992). Indra sawhney v. union of india. AIR 1993 SC 477.

Swartz, D. L. (1997). *Culture and Power: The Sociology of Pierre Bourdieu*. Chicago: University of Chicago Press.

Swidler, A. (1986). Culture in action: Symbols and strategies. *American Sociological Review 51*(2), 273–286.

Tambe, P. and L. M. Hitt (2012). The productivity of information technology investments: New evidence from IT labor data. *Information Systems Research 23*(3), 599–617.

Thorat, S. and P. Attewell (2007). The legacy of social exclusion: A correspondence study of job discrimination in india. *Economic and Political Weekly 42*(41), 4141–4145.

Thorat, S. and K. S. Newman (Eds.) (2010). *Blocked by Caste: Economic Discrimination in Modern India*. New Delhi: Oxford University Press.

Turco, C. J. (2010). Cultural foundations of tokenism: Evidence from the leveraged buyout industry. *American Sociological Review 75*(6), 894–913.

Weber, M. (1978). *Economy and Society: An Outline of Interpretive Sociology*. Berkeley: University of California Press.

Wilmers, N., S. Wang, and L. Zhang (2025). Remote work and careers. SSRN working paper.

Wiltshire, C. R. (2020). *Uniformity and Variability in the Indian English Accent*. Cambridge University Press.

Yenmez, M. B. and T. Sönmez (2022). Affirmative action in India via vertical, horizontal, and overlapping reservations. *Econometrica 90*(3), 1143–1176.

Zimmerman, S. D. (2019). Elite colleges and upward mobility to top jobs and top incomes. *American Economic Review 109*(1), 1–47.

Online Appendix (Not for Publication)

Table of Contents

A Legal and Scholarly Foundations on Caste and Class **A-3**

A.1 Legal Architecture of Administrative Coding A-3

A.2 Scholarly Foundations for a Class-Relevant Mechanism A-10

A.3 Comparative Implications for This Paper A-11

B Robustness to Category Aggregation **A-13**

C Caste in Corporate DEI Policy **A-15**

D Representativeness of the Job Fair at the Elite College **A-17**

D.1 Systematic Evidence of Standardization Across Elite Institutions A-17

D.2 Corroborating Evidence A-17

E Employment Platform Coverage Benchmarks and Profile Matching Method **A-18**

E.1 Career Outcome Data Collection and Professional Profile Matching Methodology A-18

F HR Professional Interviews: Theme-Ordered Excerpts **A-18**

G Robustness to Functional Form **A-20**

G.1 Success Rates A-20

G.2 Job Quality A-21

G.3 Post-Hire Outcomes A-22

G.4 Success Rates by Surname Identifiability A-23

G.5 Job Quality by Surname Identifiability A-25

H Robustness to Omitted Variable Bias: Oster Bounds **A-26**

H.1 Success Rates A-26

H.2 Job Quality A-26

H.3 Promotion and Tenure Rates A-27

H.4 Job Quality by Surname Identifiability A-29

I A Model Exploiting Plausibly Exogenous Variation in Selectivity Across Jobs to Identify Discrimination **A-30**

I.1 Extending the Simple Model of Late-Stage Discrimination with Imperfect Information A-34

I.2 Derivation of Expected Performance Above Threshold A-35

J LLM-Based Transcript Analysis Methodology **A-36**

J.1 Key Design Principle: Unified Taxonomy A-36

J.2 Data Collection A-37

J.3 Classification Prompts A-37

J.4 Classification Process A-39

J.5 Results Summary A-39

J.6 Sampling-Based Sensitivity Analysis A-40

K Detailed Methodology for Classifying Caste-Identifiable Surnames **A-45**

K.1 Hierarchy of Data Sources A-45

K.2 Administrative Framework and the Definition of Disadvantage A-46

K.3 Operational Mapping and Conservative Criteria A-46

L Differential Selection Across Surname Samples **A-48**

M More Details on Alternative Explanations **A-51**

N Additional Figures **A-53**

O Additional Tables **A-55**

A Legal and Scholarly Foundations on Caste and Class

This appendix makes one narrow point. For the mechanism studied in the paper, India's SC/ST/OBC categories matter not because caste and class are identical in every domain, but because the state has standardized caste-linked disadvantage into legally legible administrative categories. That administrative structure is what makes it possible to study when status becomes legible in private hiring.

A.1 Legal Architecture of Administrative Coding

Pre-independence bridge: "Depressed Classes" to "Scheduled Castes." Long before independence, representative politics and colonial administration treated caste disadvantage as an administrable category rather than a purely ritual abstraction (Miller, 1919; Rajya Sabha Committee on Personnel, Public Grievances, Law and Justice, 2005). The Government of India Act, 1935 created the schedule-based machinery, and the Government of India (Scheduled Castes) Order, 1936 implemented it by listing the relevant communities (His Majesty in Council, 1936). "Scheduled Castes" is therefore literal: it denotes communities placed on an official schedule.

Constitutional continuation and post-independence categories. Independent India retained and extended this logic. Articles 341 and 342 govern Scheduled Castes and Scheduled Tribes through Presidential Orders, and Article 340 authorizes commissions on socially and educationally backward classes (Government of India, 1950). Figures A.1 and A.2 show that continuity directly. For this paper, the key implication is measurement: disadvantage is repeatedly codified through legally legible group categories. Those categories are coarse, but they are the categories that courts, commissions, and public institutions actually use (Galanter, 1984; Bayly, 2008; Dirks, 2001; Cohn, 1987).

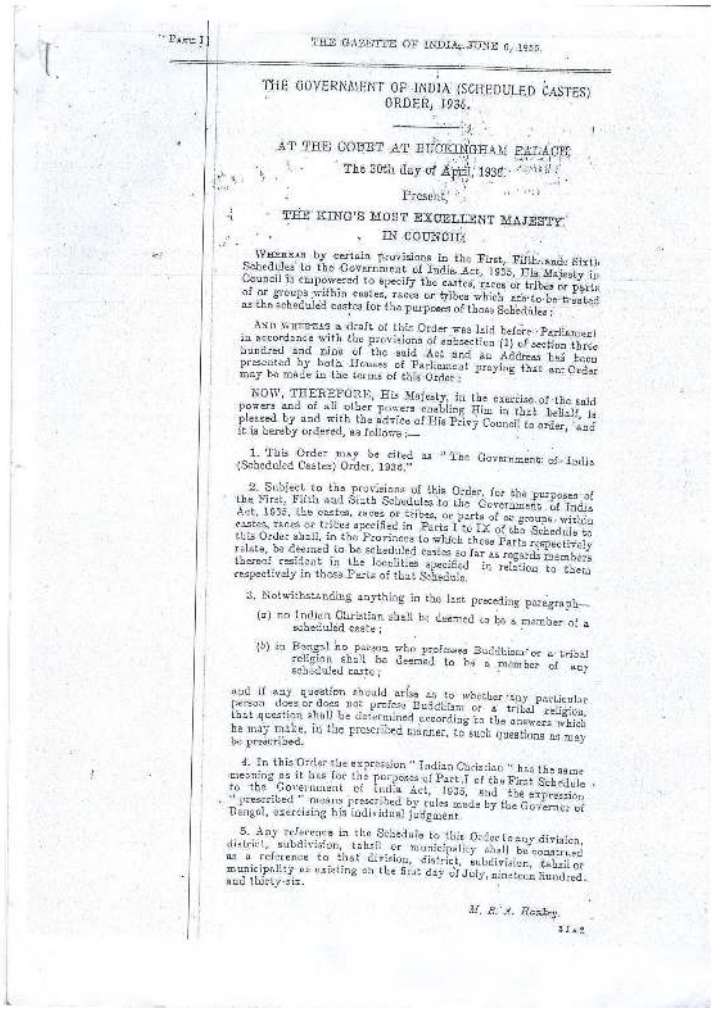

PART I] THE GAZETTE OF INDIA, JUNE 6, 1936.

THE GOVERNMENT OF INDIA (SCHEDULED CASTES) ORDER, 1936.

AT THE COURT AT BUCKINGHAM PALACE

The 30th day of April, 1936.

Present,

THE KING'S MOST EXCELLENT MAJESTY IN COUNCIL

WHEREAS by certain provisions in the First, Fifth and Sixth Schedules to the Government of India Act, 1935, His Majesty in Council is empowered to specify the castes, races or tribes or parts of or groups within castes, races or tribes which are to be treated as the scheduled castes for the purposes of those Schedules;

AND WHEREAS a draft of this Order was laid before Parliament in accordance with the provisions of subsection (1) of section three hundred and nine of the said Act and an Address has been presented by both Houses of Parliament praying that an Order may be made in the terms of this Order;

NOW, THEREFORE, His Majesty, in the exercise of the said powers and of all other powers enabling Him in that behalf, is pleased by and with the advice of His Privy Council to order, and it is hereby ordered, as follows:—

1. This Order may be cited as "The Government of India (Scheduled Castes) Order, 1936."

2. Subject to the provisions of this Order, for the purposes of the First, Fifth and Sixth Schedules to the Government of India Act, 1935, the castes, races or tribes, or parts of or groups within castes, races or tribes specified in Parts I to IX of the Schedule to this Order shall, in the Provinces to which those Parts respectively relate, be deemed to be scheduled castes so far as regards members thereof resident in the localities specified in relation to them respectively in those Parts of that Schedule.

3. Notwithstanding anything in the last preceding paragraph—

(a) no Indian Christian shall be deemed to be a member of a scheduled caste;

(b) in Bengal no person who professes Buddhism or a tribal religion shall be deemed to be a member of any scheduled caste;

and if any question should arise as to whether any particular person does or does not profess Buddhism or a tribal religion, that question shall be determined according to the answers which he may make, in the prescribed manner, to such questions as may be prescribed.

4. In this Order the expression "Indian Christian" has the same meaning as it has for the purposes of Part I of the First Schedule to the Government of India Act, 1935, and the expression "prescribed" means prescribed by rules made by the Governor of Bengal, exercising his individual judgment.

5. Any reference in the Schedule to this Order to any division, district, subdivision, tahsil or municipality shall be construed as a reference to that division, district, subdivision, tahsil or municipality as existing on the first day of July, nineteen hundred and thirty-six.

M. E. A. Realey.

3 I a 2

Government of India (Scheduled Castes) Order, 1936: order page

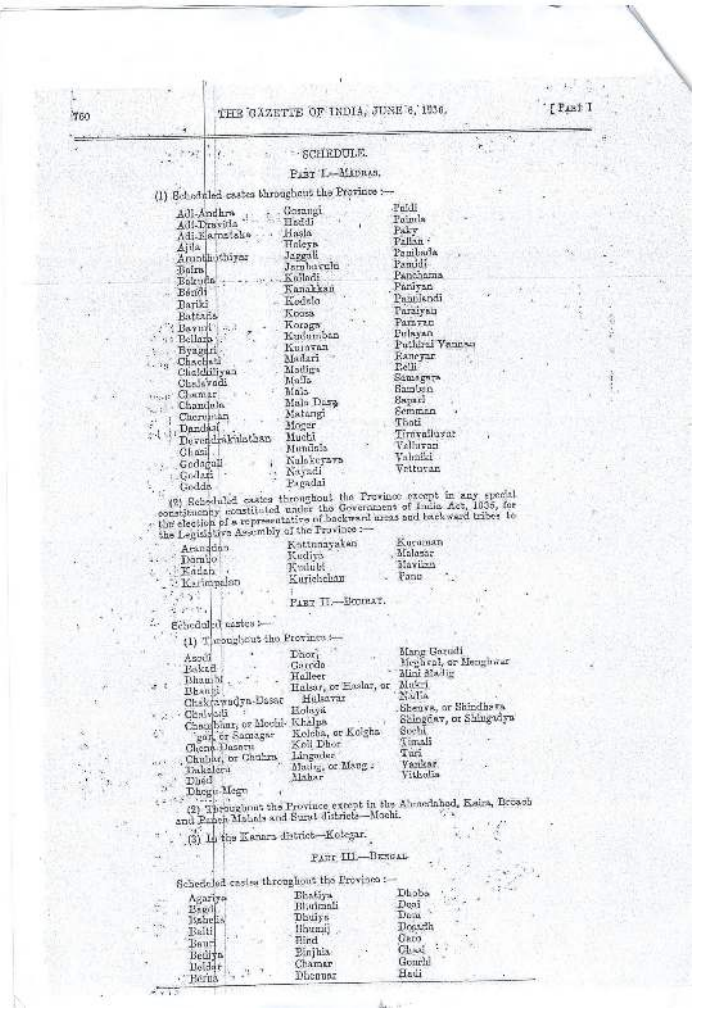

760 THE GAZETTE OF INDIA, JUNE 6, 1936. [PART I

SCHEDULE.

PART I.—MADRAS.

(1) Scheduled castes throughout the Province:—

Adi-Andhra	Gosangi	Relli
Adi-Dravida	Haddi	Tolmla
Adi-Karnataka	Hasla	Paky
Ajila	Holeya	Pallan
Arunthathiyar	Jaggali	Pambada
Baira	Jambavulu	Panadi
Baluta	Kalladi	Panchama
Bandi	Kanakkan	Paniyan
Barki	Kodelo	Pannandi
Battada	Koosa	Paraiyan
Bayadi	Koraga	Parayan
Bellara	Kudumban	Pulayan
Byagari	Kuravan	Pudhiri Vannan
Chachati	Madari	Raneyar
Chakkiliyan	Madiga	Relli
Chalavadi	Malla	Samagara
Chamar	Mala	Sambon
Chandala	Mala Dasa	Sapari
Cherumban	Malangi	Semman
Dandasi	Moger	Thoti
Devendrakulathan	Mochi	Tiruvalluvar
Chani	Mundala	Valluvan
Godagali	Nalakeyava	Valmiki
Godari	Nayadi	Vettuvan
Godda	Pagadai	

(2) Scheduled castes throughout the Province except in any special constituency constituted under the Government of India Act, 1935, for the election of a representative of backward areas and backward tribes to the Legislative Assembly of the Province:—

Arandan	Kattunayakan	Kuruman
Dambu	Kodiya	Malasar
Kadan	Kudubi	Maviian
Kalladan	Kurichchan	Pano
Kanikkalan		

PART II.—BOMBAY.

Scheduled castes:—

(1) Throughout the Province:—

Asodi	Dhor	Mang Garudi
Bakad	Garoda	Meghwal, or Menghwar
Bhambi	Halleer	Mini Madig
Bhangi	Halsar, or Halalar, or	Mukri
Chakrawadya-Dasar	Halsavar	Nalia
Chalwadi	Holeya	Shenva, or Shindhava
Chambhar, or Mochi-gar, or Samagar	Khalpa	Shingaiar, or Shingadya
Cheno-Dasaru	Koliba, or Kolgha	Sochi
Chuhar, or Ghuhra	Koli Dhor	Timali
Dakalera	Lingadar	Turi
Dhed	Madig, or Mang	Vankar
Dhegu-Megh	Mahar	Vithoba

(2) Throughout the Province except in the Ahmedabad, Kaira, Broach and Panch Mahals and Surat districts—Mochi.

(3) In the Kanara district—Kolegar.

PART III.—BENGAL.

Scheduled castes throughout the Province:—

Agariya	Bhatiya	Dhoba
Bagdi	Bhumiali	Dosi
Bahelia	Dhulya	Dom
Baiti	Bhumij	Doseth
Bauri	Bind	Garo
Beliya	Binjhia	Ghasi
Bolder	Chamar	Gonchi
Berua	Dhanuar	Hadi

Government of India (Scheduled Castes) Order, 1936: schedule page

Figure A.1: **Scheduled Castes, 1936: Late-Colonial Origin of Schedule-Based Classification**

Notes: Panel (a) reproduces the operative first page of the Government of India (Scheduled Castes) Order, 1936. Panel (b) reproduces a schedule page listing Scheduled Castes by province, making the schedule structure itself visible rather than merely stated. Read together, the two pages show that the Government of India Act, 1935 created the legal machinery and that the 1936 order implemented it by naming the communities to be treated as Scheduled Castes in official schedules. The schedule page makes clear that the category was constituted through official listing, not simply through descriptive language. Source: Ministry of Social Justice and Empowerment archival PDF.

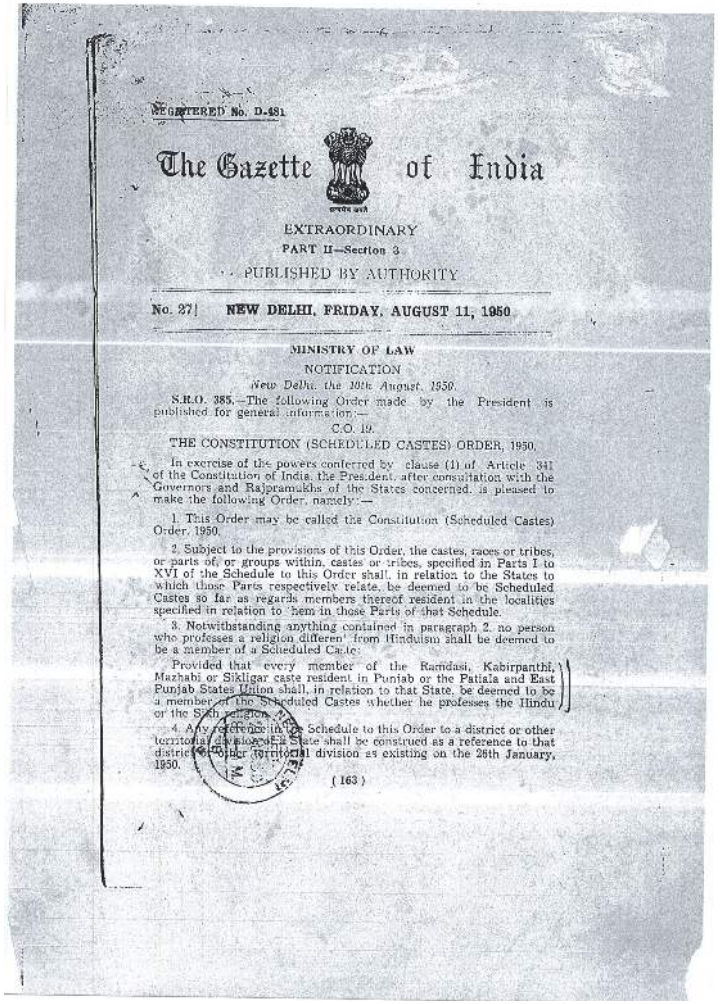

REGISTERED No. D-481

The Gazette of India

EXTRAORDINARY

PART II—Section 3

PUBLISHED BY AUTHORITY

No. 27] NEW DELHI, FRIDAY, AUGUST 11, 1950

MINISTRY OF LAW

NOTIFICATION

New Delhi, the 10th August, 1950.

S.R.O. 385.—The following Order made by the President is published for general information:—

C.O. 19.

THE CONSTITUTION (SCHEDULED CASTES) ORDER, 1950.

In exercise of the powers conferred by clause (1) of Article 341 of the Constitution of India, the President, after consultation with the Governors and Rajpramukhs of the States concerned, is pleased to make the following Order, namely:—

1. This Order may be called the Constitution (Scheduled Castes) Order, 1950.

2. Subject to the provisions of this Order, the castes, races or tribes, or parts of, or groups within, castes or tribes, specified in Parts I to XVI of the Schedule to this Order shall, in relation to the States to which those Parts respectively relate, be deemed to be Scheduled Castes so far as regards members thereof resident in the localities specified in relation to them in those Parts of that Schedule.

3. Notwithstanding anything contained in paragraph 2, no person who professes a religion different from Hinduism shall be deemed to be a member of a Scheduled Caste:

Provided that every member of the Ramdasi, Kabirpanthi, Mazhabi or Sikligar caste resident in Punjab or the Patiala and East Punjab States Union shall, in relation to that State, be deemed to be a member of the Scheduled Castes whether he professes the Hindu or the Sikh religion.

4. Any reference in the Schedule to this Order to a district or other territorial division of a State shall be construed as a reference to that district or other territorial division as existing on the 26th January, 1950.

(163)

Constitution (Scheduled Castes) Order, 1950: order page

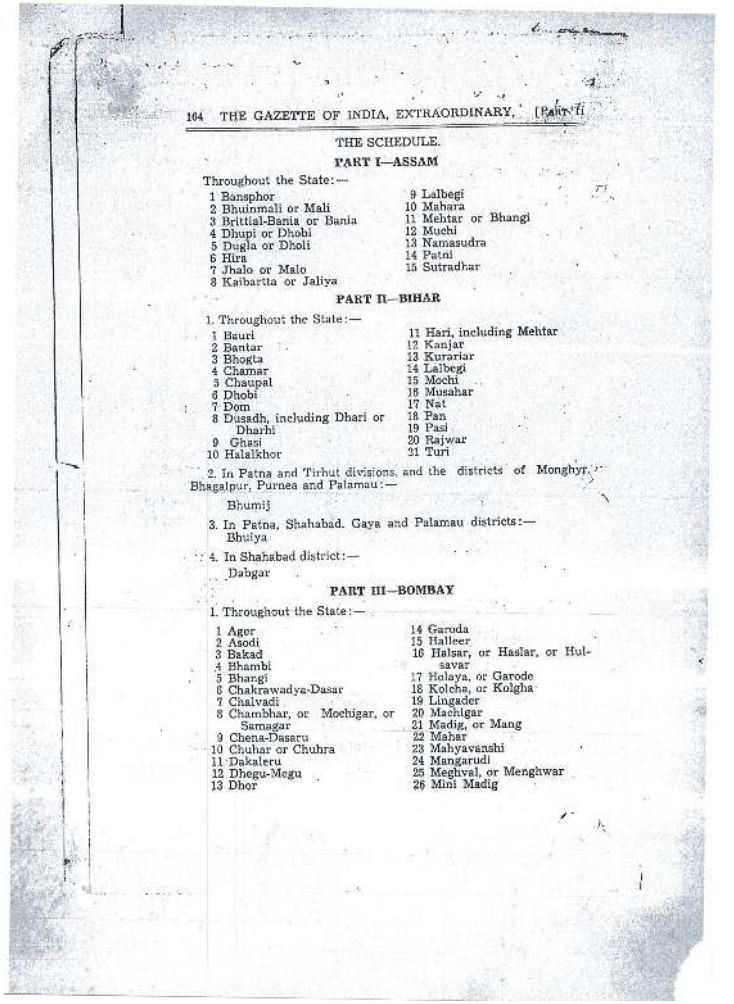

164 THE GAZETTE OF INDIA, EXTRAORDINARY, [PART II

THE SCHEDULE.

PART I—ASSAM

Throughout the State:—

1 Bansphor
2 Bhuinmali or Mali
3 Brittial-Bania or Bania
4 Dhupi or Dhobi
5 Dugla or Dholi
6 Hira
7 Jhalo or Malo
8 Kaibartta or Jaliya
9 Lalbegi
10 Mahara
11 Mehtar or Bhangi
12 Muchi
13 Namasudra
14 Patni
15 Sutradhar

PART II—BIHAR

1. Throughout the State:—

1 Bauri
2 Bantar
3 Bhogta
4 Chamar
5 Chaupal
6 Dhobi
7 Dom
8 Dusadh, including Dhari or Dharhi
9 Ghasi
10 Halalkhor
11 Hari, including Mehtar
12 Kanjar
13 Kurariar
14 Lalbegi
15 Mochi
16 Musahar
17 Nat
18 Pan
19 Pasi
20 Rajwar
21 Turi

2. In Patna and Tirhut divisions, and the districts of Monghyr, Bhagalpur, Purnea and Palamau:—

Bhumij

3. In Patna, Shahabad, Gaya and Palamau districts:—

Bhuiya

4. In Shahabad district:—

Dabgar

PART III—BOMBAY

1. Throughout the State:—

1 Ager
2 Asodi
3 Bakad
4 Bhambi
5 Bhangi
6 Chakrawadya-Dasar
7 Chalvadi
8 Chambhar, or Mochigar, or Samagar
9 Chena-Dasaru
10 Chuhar or Chuhra
11 Dakaleru
12 Dhegu-Megu
13 Dhor
14 Garoda
15 Halleer
16 Halsar, or Haslar, or Hulsavar
17 Holaya, or Garode
18 Kolcha, or Kolgha
19 Lingader
20 Machigar
21 Madig, or Mang
22 Mahar
23 Mahyavanshi
24 Mangarudi
25 Meghval, or Menghwar
26 Mini Madig

Constitution (Scheduled Castes) Order, 1950: schedule page

Figure A.2: **Scheduled Castes, 1950: Constitutional Coding and Official State-Wise Schedules**

Notes: Panel (a) reproduces the first page of the Constitution (Scheduled Castes) Order, 1950. Panel (b) reproduces the beginning of the constitutional Schedule, listing Scheduled Castes state by state. The pair shows that independent India retained the schedule-based coding logic of the late-colonial period, but now embedded it in the constitutional order under Article 341 and associated Presidential Orders. The state-wise schedule makes the constitutional continuity of official caste listing visually explicit. Source: Ministry of Social Justice and Empowerment official PDF.

Scheduled Tribes and the breadth of the schedule-based framework. The same constitutional machinery also governs Scheduled Tribes, though ST classification has a distinct historical genealogy tied to tribal communities, frontier governance, and uneven state incorporation rather than untouchability alone (Bayly, 2008; Dirks, 2001; Cohn, 1987; Galanter, 1984). For the paper, the relevant point is institutional rather than ethnographic: the Indian state codified multiple forms of durable disadvantage through official schedules. Figure A.3 therefore reinforces the measurement claim in the main text: SC/ST/OBC are policy categories designed to capture persistent social disadvantage, not point-in-time income.

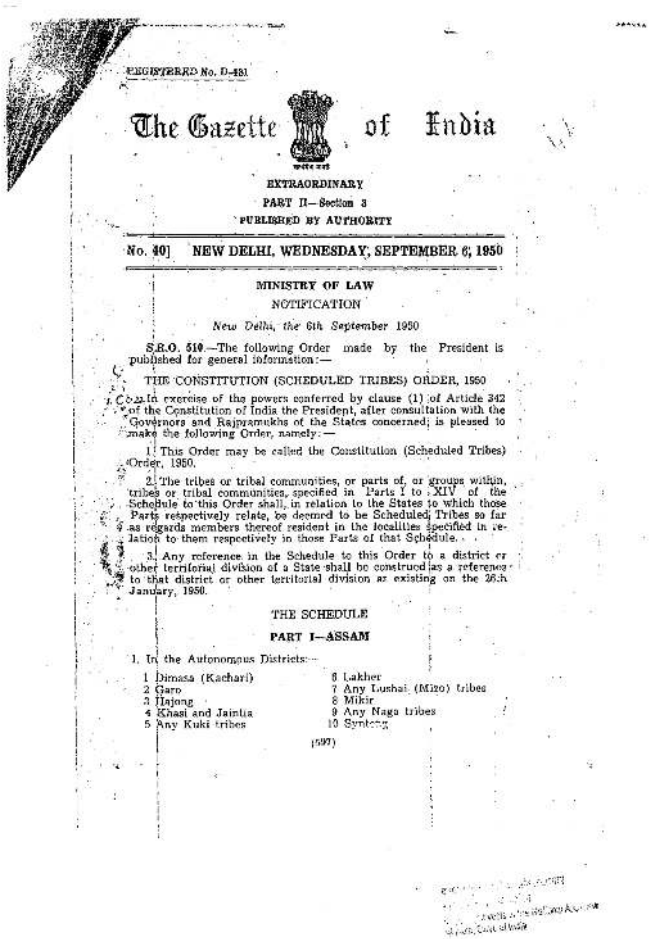

REGISTERED No. D-481

The Gazette of India

EXTRAORDINARY

PART II—Section 3

PUBLISHED BY AUTHORITY

No. 40] NEW DELHI, WEDNESDAY, SEPTEMBER 6, 1950

MINISTRY OF LAW

NOTIFICATION

New Delhi, the 6th September 1950

S.R.O. 510.—The following Order made by the President is published for general information:—

THE CONSTITUTION (SCHEDULED TRIBES) ORDER, 1950

In exercise of the powers conferred by clause (1) of Article 342 of the Constitution of India the President, after consultation with the Governors and Rajpramukhs of the States concerned, is pleased to make the following Order, namely:—

1. This Order may be called the Constitution (Scheduled Tribes) Order, 1950.

2. The tribes or tribal communities, or parts of, or groups within, tribes or tribal communities, specified in Parts I to XIV of the Schedule to this Order shall, in relation to the States to which those Parts respectively relate, be deemed to be Scheduled Tribes so far as regards members thereof resident in the localities specified in relation to them respectively in those Parts of that Schedule.

3. Any reference in the Schedule to this Order to a district or other territorial division of a State shall be construed as a reference to that district or other territorial division as existing on the 26th January, 1950.

THE SCHEDULE

PART I—ASSAM

1. In the Autonomous Districts:—

1 Dimasa (Kachari)
2 Garo
3 Hajong
4 Khasi and Jaintia
5 Any Kuki tribes
6 Lakher
7 Any Lushai (Mizo) tribes
8 Mikir
9 Any Naga tribes
10 Synteng

(597)

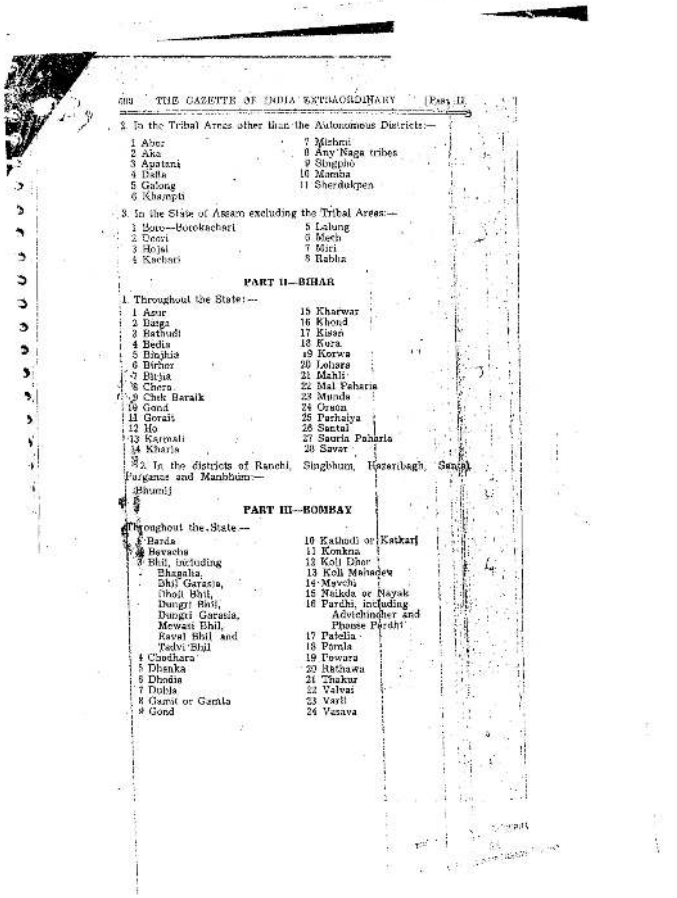

THE GAZETTE OF INDIA EXTRAORDINARY [PART II

2. In the Tribal Areas other than the Autonomous Districts:—

1 Abor
2 Aka
3 Apatani
4 Daffa
5 Galong
6 Khampti
7 Mishmi
8 Any Naga tribes
9 Singpho
10 Miroha
11 Sherdukpen

3. In the State of Assam excluding the Tribal Areas:—

1 Boro—Borokachari
2 Deori
3 Hojai
4 Kachari
5 Lalung
6 Mech
7 Miri
8 Rabha

PART II—BIHAR

1. Throughout the State:—

1 Asur
2 Birjia
3 Bathudi
4 Bedia
5 Binjhia
6 Birhor
7 Birjia
8 Chero
9 Chik Baraik
10 Gond
11 Gorait
12 Ho
13 Kareali
14 Kharia
15 Kharwar
16 Khond
17 Kisan
18 Kora
19 Korwa
20 Lohara
21 Mahli
22 Mal Paharia
23 Munda
24 Oraon
25 Parhaiya
26 Santal
27 Sauria Paharia
28 Savar

2. In the districts of Ranchi, Singbhum, Hazaribagh, Santal Parganas and Manbhum:—

Bhumij

PART III—BOMBAY

Throughout the State:—

1 Barda
2 Bevacha
3 Bhil, including Bhagalia, Bhil Garasia, Dholi Bhil, Dungri Bhil, Dungri Garasia, Mewasi Bhil, Rawal Bhil and Tadvi Bhil
4 Chodhara
5 Dhanka
6 Dhodia
7 Dubla
8 Gamit or Gamta
9 Gond
10 Kathodi or Katkari
11 Konkna
12 Koli Dhor
13 Koli Mahadev
14 Mavchi
15 Naikda or Nayak
16 Pardhi, including Advichinchar and Phanse Pardhi
17 Patelia
18 Porma
19 Powara
20 Rathawa
21 Thakur
22 Valvai
23 Varli
24 Vasava

Constitution (Scheduled Tribes) Order, 1950: front page

Constitution (Scheduled Tribes) Order, 1950: schedule page

Figure A.3: **Scheduled Tribes: Constitutional Order and State Lists**

Notes: Panel (a) reproduces the front page of the Constitution (Scheduled Tribes) Order, 1950. Panel (b) reproduces a schedule page listing Scheduled Tribes across multiple states, making the territorial organization of the ST schedule explicit. The pair matters for the paper's measurement argument because it shows that the same schedule-based constitutional framework extended beyond Scheduled Castes to Scheduled Tribes under Article 342. The state listing also makes clear that tribal coding, like SC coding, was operationalized through official schedules rather than abstract category labels alone. Source: Official Constitution (Scheduled Tribes) Order, 1950 PDF.

Backward classes and the administrative logic of cumulative disadvantage. Article 340 was operationalized most clearly by the Mandal Commission, which defined backwardness using eleven indicators spanning social, educational, and economic dimensions rather than current income alone (Government of India, 1950; Mandal Commission, 1980; Galanter, 1984). For this paper, that is the crucial measurement fact: the administrative category is explicitly a bundled social position. Figure A.4 is included because the indicator page is the clearest archival statement of that logic.

Judicial consolidation: *Indra Sawhney* and the class-relevant reading of caste. *Indra Sawhney v. Union of India* (1992) consolidated this framework by treating caste as a legally usable indicator of persistent social and educational disadvantage in public policy, while also introducing creamy-layer exclusion within OBCs (Supreme Court of India, 1992). That distinction matters here. Reservation eligibility can change with current income or status, but the underlying administrative group classification remains. OBC-CL outcomes are therefore still informative if interviewers respond to cumulative social position rather than income alone (Supreme Court of India, 1992; Galanter, 1984).

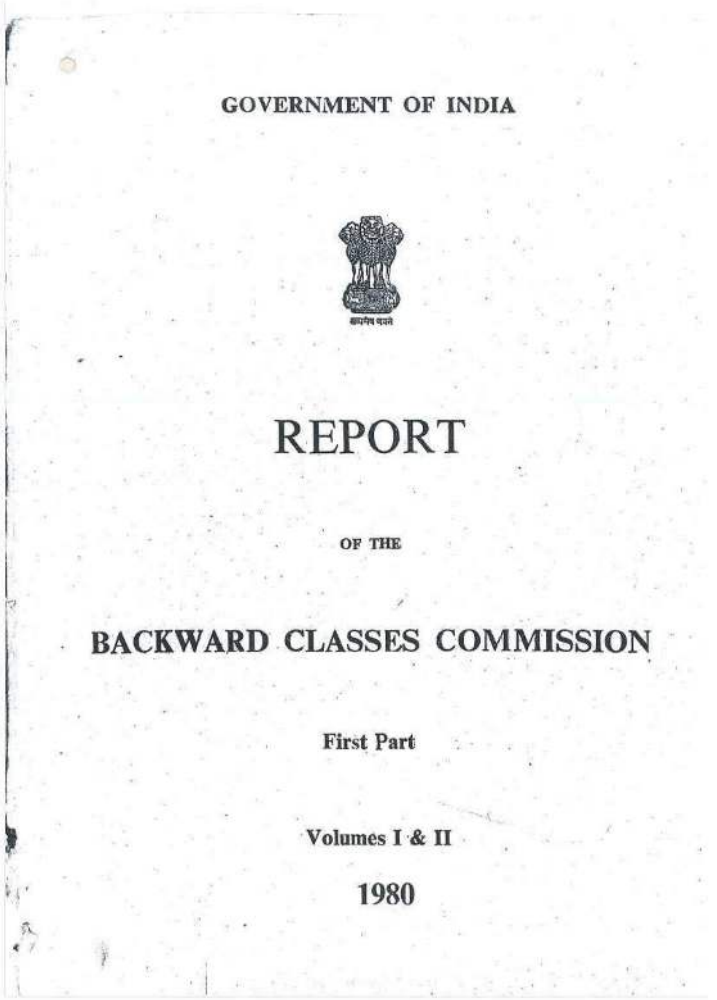

GOVERNMENT OF INDIA

सत्यमेव जयते

REPORT

OF THE

BACKWARD CLASSES COMMISSION

First Part

Volumes I & II

1980

Backward Classes Commission report cover

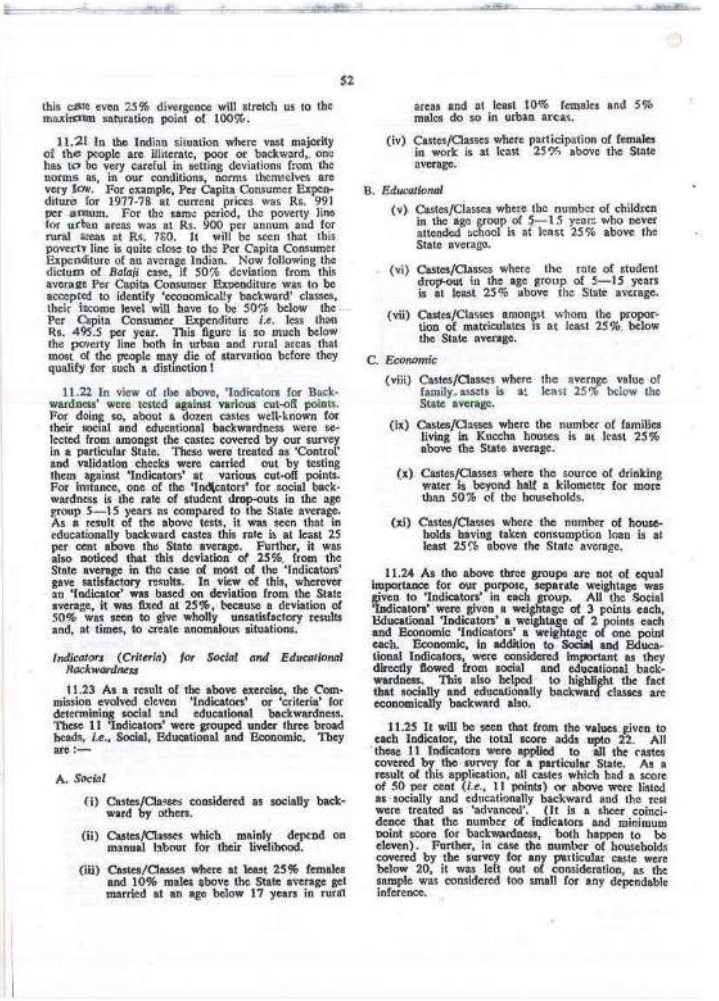

this caste even 25% divergence will stretch us to the maximum saturation point of 100%.

11.21 In the Indian situation where vast majority of the people are illiterate, poor or backward, one has to be very careful in setting deviations from the norms as, in our conditions, norms themselves are very low. For example, Per Capita Consumer Expenditure for 1977-78 at current prices was Rs. 991 per annum. For the same period, the poverty line for urban areas was at Rs. 900 per annum and for rural areas at Rs. 780. It will be seen that this poverty line is quite close to the Per Capita Consumer Expenditure of an average Indian. Now following the dictum of *Balaji* case, if 50% deviation from this average Per Capita Consumer Expenditure was to be accepted to identify 'economically backward' classes, their income level will have to be 50% below the Per Capita Consumer Expenditure *i.e.* less than Rs. 495.5 per year. This figure is so much below the poverty line both in urban and rural areas that most of the people may die of starvation before they qualify for such a distinction!

11.22 In view of the above, 'Indicators for Backwardness' were tested against various cut-off points. For doing so, about a dozen castes well-known for their social and educational backwardness were selected from amongst the castes covered by our survey in a particular State. These were treated as 'Control' and validation checks were carried out by testing them against 'Indicators' at various cut-off points. For instance, one of the 'Indicators' for social backwardness is the rate of student drop-outs in the age group 5—15 years as compared to the State average. As a result of the above tests, it was seen that in educationally backward castes this rate is at least 25 per cent above the State average. Further, it was also noticed that this deviation of 25% from the State average in the case of most of the 'Indicators' gave satisfactory results. In view of this, wherever an 'Indicator' was based on deviation from the State average, it was fixed at 25%, because a deviation of 50% was seen to give wholly unsatisfactory results and, at times, to create anomalous situations.

Indicators (Criteria) for Social and Educational Backwardness

11.23 As a result of the above exercise, the Commission evolved eleven 'Indicators' or 'criteria' for determining social and educational backwardness. These 11 'Indicators' were grouped under three broad heads, *i.e.*, Social, Educational and Economic. They are :—

A. *Social*

(i) Castes/Classes considered as socially backward by others.

(ii) Castes/Classes which mainly depend on manual labour for their livelihood.

(iii) Castes/Classes where at least 25% females and 10% males above the State average get married at an age below 17 years in rural areas and at least 10% females and 5% males do so in urban areas.

(iv) Castes/Classes where participation of females in work is at least 25% above the State average.

B. *Educational*

(v) Castes/Classes where the number of children in the age group of 5—15 years who never attended school is at least 25% above the State average.

(vi) Castes/Classes where the rate of student drop-out in the age group of 5—15 years is at least 25% above the State average.

(vii) Castes/Classes amongst whom the proportion of matriculates is at least 25% below the State average.

C. *Economic*

(viii) Castes/Classes where the average value of family assets is at least 25% below the State average.

(ix) Castes/Classes where the number of families living in Kuccha houses is at least 25% above the State average.

(x) Castes/Classes where the source of drinking water is beyond half a kilometer for more than 50% of the households.

(xi) Castes/Classes where the number of households having taken consumption loan is at least 25% above the State average.

11.24 As the above three groups are not of equal importance for our purpose, separate weightage was given to 'Indicators' in each group. All the Social 'Indicators' were given a weightage of 3 points each, Educational 'Indicators' a weightage of 2 points each and Economic 'Indicators' a weightage of one point each. Economic, in addition to Social and Educational Indicators, were considered important as they directly flowed from social and educational backwardness. This also helped to highlight the fact that socially and educationally backward classes are economically backward also.

11.25 It will be seen that from the values given to each Indicator, the total score adds upto 22. All these 11 Indicators were applied to all the castes covered by the survey for a particular State. As a result of this application, all castes which had a score of 50 per cent (*i.e.*, 11 points) or above were listed as socially and educationally backward and the rest were treated as 'advanced'. (It is a sheer coincidence that the number of indicators and minimum point score for backwardness, both happen to be eleven). Further, in case the number of households covered by the survey for any particular caste were below 20, it was left out of consideration, as the sample was considered too small for any dependable inference.

Indicator framework for identifying backward classes

Figure A.4: **Mandal Commission: Report and Indicator Framework**

Notes: Panel (a) reproduces the front page of the Backward Classes Commission (Mandal) Report, 1980. Panel (b) reproduces the page setting out the indicator framework for identifying backward classes: eleven indicators grouped under social, educational, and economic heads, with unequal weights summing to 22 points. This is the most relevant archival evidence for the paper's measurement argument because it shows that backwardness was operationalized as a bundled social position rather than income alone. Sources: National Commission for Backward Classes archive; Ministry of Social Justice and Empowerment official excerpt.

A.2 Scholarly Foundations for a Class-Relevant Mechanism

Administrative categories are not *varna* labels. Indian scholarship has long distinguished broad *varna* schemas from localized *jāti* groupings (Basham, 1954; Béteille, 1996). The empirical object in this paper is not a canonical four-*varna* map. It is the modern administrative aggregation SC/ST/OBC, which bundles many local groups into policy categories defined around durable social disadvantage.

Why the interview mechanism is class-relevant. Class theory predicts that evaluators infer social position from bundles of capital revealed in interaction. In Weberian and Bourdieusian terms, market position, status closure, and embodied cultural capital are distinct but tightly linked in practice (Weber, 1978; Parkin, 1979; Bourdieu, 1986, 1984; Lareau, 2003; Lamont and Lareau, 1988). Rivera's work on elite hiring makes the same point in organizational language: final-round "fit" assessments often reward ease, polish, and biographical familiarity (Rivera, 2012). Those frameworks match the timing in this paper. Penalties should emerge where evaluators can probe biography and rapport, not merely where candidates are visible.

Historical accounts likewise caution against treating caste as primarily phenotypic. Colonial and postcolonial institutions stabilized categories through censuses, schedules, and administrative practice (Dirks, 2001; Bayly, 2008; Cohn, 1987). That is why discrimination can operate through interactional inference even when phenotype is weakly informative. The same logic also clarifies OBC-CL results: current income thresholds do not erase inherited differences in schooling, language socialization, or networks (Clark, 2014; Deshpande, 2011; Munshi, 2019; Jodhka, 2015).

A.3 Comparative Implications for This Paper

Relative to many race settings where assignment is more immediate at first exposure (Omi and Winant, 2014), the mechanism here is more biographical and interactional. That helps explain why group debates and personal interviews diverge even though both are subjective rounds: group debates can reveal some social cues, but personal interviews create much denser opportunities to elicit family background, schooling trajectories, cultural references, and rapport. The measurement advantage is also unusual. In many class settings researchers observe only noisy proxies such as neighborhood or parental education; India's SC/ST/OBC architecture supplies high-signal administrative coding of cumulative disadvantage.

Institutional boundaries and measurement validity. Private multinational hiring in this setting is not governed by public-sector reservation rules (Galanter, 1984; Jodhka, 2008). That boundary matters for identification: late-stage caste penalties are not mechanical compliance effects. Administrative categories remain coarse relative to lived heterogeneity, but they are informative for group-threshold questions, and Appendix B shows that the main pattern survives disaggregation, including for OBC-CL.

The legal and sociological evidence therefore points to a specific interpretation of the main results. In a setting where caste is often inferred from biography rather than read off phenotype, disparities should appear when interviews make status legible. The state did not determine that timing; it made the timing test observable. Table A.1 summarizes the comparison.

Table A.1: Comparing Caste, Class, and Race Across Nine Dimensions Relevant to Discrimination Research

Dimension	Race	Class	Caste
Primary identification	Phenotypic (skin color, facial features)—visible at a glance (Omi and Winant, 2014)	Inferred from credentials, consumption, accent (in extended conversation), networks (Bourdieu, 1986)	Inferred from biographical markers (surname, hometown, parental occupation, educational pedigree) (Thorat and Attewell, 2007)
Intergenerational persistence	High (but boundaries shift historically) (Ignatiev, 1995)	High (0.7–0.8 correlation across 10–15 generations) (Clark, 2014)	High (administratively codified) (Mandal Commission, 1980)
Formal mutability	Generally immutable (though "passing" documented) (Hobbs, 2014)	Formally mutable (though empirically persistent)	Formally immutable (though Sanskritization documented) (Srinivas, 1966)
Endogamy	Historically enforced; declining	Assortative mating by education/status (Schwartz, 2013)	Strong; marriage markets sort explicitly on caste (Banerjee et al., 2013)
State classification	Census categories; affirmative action	Rarely codified administratively	Explicit constitutional categories (SC/ST/OBC) (Galanter, 1984)
Basis of classification	Phenotype + ancestry	Occupation, income, education, credentials	"Socially and educationally backward" (11 composite indicators) (Mandal Commission, 1980)
Genetic distinctiveness	Population structure exists but continuous	N/A	Historic gene flow across boundaries; not phenotypically distinguishable (Moorjani et al., 2013)
Capitals bundled	Some correlation with economic/social capital	Economic, cultural, social, symbolic (Bourdieu, 1986)	Economic, cultural, social, symbolic (administratively codified) (Deshpande, 2011)
Discrimination mechanism	Name-based identification → taste-based or statistical (Bertrand and Mullainathan, 2004)	Biographical inference → "fit" screening (Rivera, 2012)	Biographical inference → "fit" screening

Note: This table is a comparative map for discrimination mechanisms, not an exhaustive theory of stratification. For this paper, three dimensions matter most: identification mechanism, degree of state codification, and whether the system bundles multiple forms of capital. On those dimensions, caste aligns with class for the mechanism studied here: both rely more on biographical inference than immediate phenotype, but India's administrative coding provides a measurement advantage unavailable for class in most Western settings. See Section 2 for the main-text discussion and Appendix A for the extended treatment.

B Robustness to Category Aggregation

Motivation and takeaway. This section tests whether the pooled disadvantaged-category result is driven by composition across Scheduled Castes (SC), Scheduled Tribes (ST), Other Backward Classes non-creamy-layer (OBC-NCL), and Other Backward Classes creamy-layer (OBC-CL). It is not. Interview-stage penalties remain large across disaggregated categories, including OBC-CL, which is reservation-ineligible. If reservation eligibility or current income were the main channel, OBC-CL penalties should attenuate sharply. Instead, across hiring, job quality, and post-hire outcomes, disaggregation supports the same interpretation as the pooled estimates: the mechanism loads on cumulative social position rather than current income alone.

Results: Hiring Success Rates. Table B.1 reports stage-by-stage success-rate gaps. Column (1) reproduces the pooled disadvantaged-category result; Columns (2)–(5) disaggregate to SC, ST, OBC-NCL, and OBC-CL. The main pattern is unchanged: pre-interview gaps are small, while interview-stage penalties remain large across disadvantaged categories, including OBC-CL.

Table B.1: Success Rate Differences by Disaggregated Caste Category

		Coefficient on Category Indicator (Relative to General)				
Stage	N	Disadv. Caste	SC	ST	OBC–NCL	OBC–CL
App → Test	305,541	-0.0025	-0.0254***	-0.0396***	0.0112***	-0.0511***
		(0.0018)	(0.0018)	(0.0023)	(0.0015)	(0.0045)
Test → GD	183,150	-0.0002	-0.0189***	-0.0584***	0.0088***	-0.0526***
		(0.0023)	(0.0027)	(0.0043)	(0.0019)	(0.0072)
GD → Int	148,489	-0.0005	0.0042	-0.0389***	-0.0002	0.0074
		(0.0051)	(0.0043)	(0.0073)	(0.0028)	(0.0104)
Int → Offer	38,378	-0.2938***	-0.2799***	-0.2704***	-0.2958***	-0.3634***
		(0.0107)	(0.0074)	(0.0135)	(0.0050)	(0.0186)
Year FE		✓	✓	✓	✓	✓
Job FE		✓	✓	✓	✓	✓
Other Controls		✓	✓	✓	✓	✓

Notes: This table reports disaggregated versions of Equation 1, replacing the pooled disadvantaged-caste indicator with separate indicators for SC, ST, OBC-NCL, and OBC-CL; General/Unreserved is the omitted reference group. Rows correspond to the four stage transitions (App→Test, Test→GD, GD→Int, Int→Offer), each estimated on candidates still in contention at that stage. The N column reports stage-specific estimation samples. All specifications include the baseline controls and fixed effects from the main success-rate regressions. Standard errors are two-way clustered at the student and job levels. Significance levels: * p<0.10, ** p<0.05, *** p<0.01.

Results: Job Quality. Table B.2 reports job-quality differences (log mean salary in USD PPP) using the same category split. The salary patterns mirror the hiring-success results, again including sizable OBC-CL penalties.

Results: Post-Hire Outcomes. Table B.3 extends the same disaggregation to promotions and first-job exit. The post-hire pattern remains consistent with a higher-bar mechanism: disadvantaged hires, including OBC-CL, are more likely to be promoted and less likely to leave conditional on being hired.

Table B.2: Job Quality Differences by Disaggregated Caste Category

		Coefficient on Category Indicator (Relative to General)				
Outcome	N	Disadv. Caste	SC	ST	OBC–NCL	OBC–CL
Log Salary (App)	4,005	0.0036	0.0118	0.0195*	-0.0013	0.0081
		(0.0049)	(0.0075)	(0.0101)	(0.0054)	(0.0195)
Log Salary (Test)	3,618	-0.0086	-0.0150	-0.0177	-0.0071	0.0163
		(0.0070)	(0.0110)	(0.0161)	(0.0076)	(0.0283)
Log Salary (Debate)	3,266	-0.0063	-0.0063	-0.0038	-0.0071	0.0087
		(0.0067)	(0.0114)	(0.0189)	(0.0072)	(0.0306)
Log Salary (Interview)	3,031	-0.0184*	-0.0224	-0.0085	-0.0181*	-0.0163
		(0.0099)	(0.0177)	(0.0314)	(0.0107)	(0.0446)
Log Salary (Offer)	2,527	-0.1049***	-0.0954***	-0.1150**	-0.1039***	-0.1582***
		(0.0132)	(0.0266)	(0.0513)	(0.0140)	(0.0587)
Log Salary (Choice)	2,527	-0.1035***	-0.0682*	-0.1280*	-0.1075***	-0.1377*
		(0.0174)	(0.0351)	(0.0677)	(0.0185)	(0.0775)
Year FE		✓	✓	✓	✓	✓
Department FE		✓	✓	✓	✓	✓
Other Controls		✓	✓	✓	✓	✓

Notes: This table reports disaggregated versions of Equation 2, with General/Unreserved as the omitted reference group. For each stage, the dependent variable is the log mean posted salary across jobs for which the candidate remains eligible (and for Choice, the accepted-job salary). The N column reports stage-specific estimation samples. Coefficient columns report the pooled disadvantaged-category estimate and disaggregated coefficients for SC, ST, OBC-NCL, and OBC-CL. All specifications include the baseline controls and fixed effects from the main job-quality regressions. For App/Test/Debate/Interview/Offer stages, HC1 robust standard errors are reported; for Choice, standard errors are clustered at the job level. * $p<0.10$, ** $p<0.05$, *** $p<0.01$.

Table B.3: Post-Hire Outcomes by Disaggregated Caste Category

		Coefficient on Category Indicator (Relative to General)				
Outcome	N	Disadv. Caste	SC	ST	OBC–NCL	OBC–CL
Promotion	2,489	0.0962***	0.0805*	0.0349	0.0967***	0.1760*
		(0.0300)	(0.0478)	(0.0850)	(0.0270)	(0.0960)
Left First Job	2,489	-0.1700***	-0.1374**	-0.1811*	-0.1273***	-0.3000***
		(0.0460)	(0.0537)	(0.0963)	(0.0301)	(0.1088)
Year FE		✓	✓	✓	✓	✓
Job FE		✓	✓	✓	✓	✓
Other Controls		✓	✓	✓	✓	✓

Notes: This table reports disaggregated versions of Equation 4, with General/Unreserved as the omitted reference group. "Promoted" is an indicator for receiving a promotion; "Left First Job" is an indicator for exiting the first employer. The N column reports outcome-specific estimation samples. Coefficient columns report the pooled disadvantaged-category estimate and disaggregated coefficients for SC, ST, OBC-NCL, and OBC-CL. All specifications include the baseline controls and fixed effects from the main post-hire regressions; the promotion specification additionally controls for tenure duration. Standard errors are clustered at the job level. * $p<0.10$, ** $p<0.05$, *** $p<0.01$.

C Caste in Corporate DEI Policy

This appendix asks a simple question: do firms recruiting in this market publicly treat caste as a DEI category? I combine the employer-attitudes literature with a systematic scrape of DEI pages from recruiting firms.

Qualitative backdrop. India-specific scholarship documents strong employer resistance to caste-conscious private-sector affirmative action, usually framed in the language of merit and efficiency (Jodhka, 2008; Jodhka and Newman, 2007; Thorat and Newman, 2010). As Jodhka (2008) summarizes:

> *"The response of the Indian private sector to the proposal was unequivocally negative... Reservations, they repeatedly argued, would work against 'merit' and 'efficiency'."*

Related audit and labor-market evidence is consistent with a gap between universalist public language and caste-linked treatment in practice (Thorat and Attewell, 2007; Cassan et al., 2021; Munshi, 2019). The quantitative exercise below asks whether caste is even visible in the formal DEI language of firms active in elite recruitment.

Data and collection protocol. The sample includes 1,017 firms recruiting at elite Indian engineering colleges (377 India-headquartered and 640 foreign-headquartered). For each firm, I queried standard DEI paths (for example, `/diversity`, `/dei`, `/inclusion`) and India-specific alternatives when available, parsed the resulting pages, and retained only substantive text exceeding 200 words. This yields 191 qualifying DEI pages; 161 have clean headquarters metadata (59 India-headquartered, 102 foreign-headquartered), while 30 enter aggregate counts but are excluded from headquarters splits.

Coding dictionary. Each DEI page is scanned for demographic-category mentions using a fixed keyword dictionary. Detection is permissive (any keyword hit); interpretation is conservative, treating a hit as evidence of public salience rather than substantive policy depth.

Note: A page is coded as mentioning a category if any corresponding keyword appears in the extracted DEI text.

Category	Keywords
Gender	women, gender, female, male, non-binary, gender equality, gender parity
Race/Ethnicity	race, racial, ethnicity, ethnic, minorities, black, asian, hispanic, latino
LGBTQ+	lgbtq, lgbt, lesbian, gay, bisexual, transgender, queer, sexual orientation
Caste	**caste, scheduled caste, scheduled tribe, sc/st, dalit, obc, other backward class**
Socioeconomic	socioeconomic, economic background, class, low-income, first-generation
Disability	disability, disabled, differently abled, accessibility, pwd, handicapped
Age	age, generational, ageism, older workers, age diversity, multigenerational
Religion	religion, religious, faith, beliefs, religious diversity

Persistence and interpretation. Wayback snapshots for 2017–2024 (262 archived pages) show that caste mentions are rare over time: only 1.9% of archived DEI pages mention caste, with zero mentions in 2017, 2018, 2023, and 2024.

Figure C.1 shows the main cross-sectional result. Among firms with substantive DEI pages and known headquarters metadata, caste appears in only 5.1% of India-headquartered pages and 1.0% of foreign-headquartered pages. Combined with the fact that 81% of sampled firms have no qualifying DEI page at all, this indicates very low caste salience in formal DEI language. That makes explicit caste-targeted promotion or hiring initiatives an unlikely explanation for the post-hire patterns in the main text (Jodhka, 2008; Jodhka and Newman, 2007; Thorat and Newman, 2010). The evidence is about public policy language, not internal practice, so it narrows but does not mechanically rule out informal channels.

Figure C.1: **Demographic Mentions in Corporate DEI Statements by Company Headquarters**

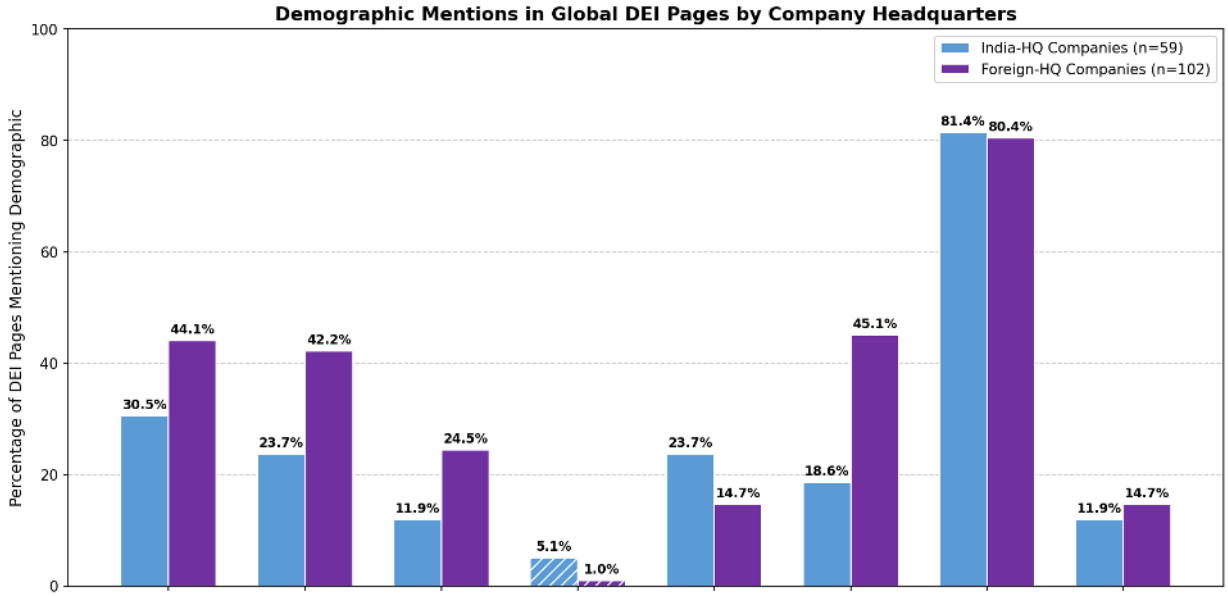

Notes: Figure C.1 shows the percentage of corporate DEI pages that mention each demographic category, separately by headquarters location. We scraped 1,017 companies recruiting at elite Indian engineering colleges; 191 had DEI pages with substantive content (>200 words), of which 161 could be matched to headquarters (59 India-headquartered, 102 foreign-headquartered). Each bar is the share of qualifying DEI pages mentioning the category. Caste keywords searched: caste, scheduled caste, scheduled tribe, SC/ST, dalit, OBC, other backward class. The 826 companies (81%) with no qualifying DEI page are not shown.

D Representativeness of the Job Fair at the Elite College

This appendix shows that the studied job fair closely matches elite campus recruitment elsewhere in India.

D.1 Systematic Evidence of Standardization Across Elite Institutions

I examined placement procedures at 108 elite institutions—all 23 IITs, all 31 NITs, 20 IIMs, and 34 other elite institutions—using automated collection from official placement-office websites in November 2025. Twenty-eight institutions (25.9%) published detailed procedures. Among those that did, the same four-stage pipeline documented in this paper repeatedly appears: application screening, written tests, group debates, and personal interviews. That uniformity spans IITs, NITs, IIMs, and private engineering colleges, suggesting that the focal institution is a high-information example of a broader elite-placement architecture rather than an outlier.

Interpretation and limit. The transparent institutions are not a random sample, so the representativeness claim rests on two facts: near-identical procedures among institutions that disclose them, and strong incentives for standardization in campus recruitment (shared recruiters, timelines, and placement-office templates). This does not imply perfect homogeneity, but it strongly supports external validity.

D.2 Corroborating Evidence

Four additional observations point in the same direction:

1. **Historical roots of standardization.** India's premier technical institutions were built with explicit reference to elite Western universities and inherited placement-office structures that were already well developed in the U.S. by the mid-twentieth century (Datta, 2017; Anderson, 1938).

2. **Graduates predominantly pursue elite private-sector jobs.** Graduates from elite colleges overwhelmingly target multinational and top Indian private-sector employers, with public-sector roles playing a secondary part (Subramanian, 2019).

3. **Elite firms recruit from the same campuses.** The same employers repeatedly target the same high-status campuses, generating a similar candidate pool across institutions and cohorts (Jodhka and Naudet, 2019).

4. **Student withdrawal is minimal.** Participation in campus placements is near-universal, and voluntary deregistration is low, consistent with evidence from comparable institutions (e.g., Mamgain, 2019).

Taken together, the institutional audit and corroborating evidence indicate that the four-stage screening pipeline documented here is standard practice in elite Indian campus recruitment.

E Employment Platform Coverage Benchmarks and Profile Matching Method

This appendix benchmarks the professional-profile matching used for post-hire outcomes and summarizes the matching protocol.

External evidence suggests LinkedIn reaches roughly 80% of U.S. bachelor's-degree holders, with even higher penetration in STEM and elite-college populations (Wilmers et al., 2025; Tambe and Hitt, 2012). LinkedIn's own STEM report and cross-country evidence point to strong representation among Indian STEM workers as well (Baird et al., 2023; Kashyap and Verkroost, 2021). Industry and elite-alumni audits are consistent with very high coverage in the relevant labor market (Tambe et al., 2021; Zhang and Vucetic, 2016). In our data, profile matches are checked against placement records, manually audited for ambiguous names, and de-duplicated, making a 98% match rate plausible for an elite STEM cohort.

E.1 Career Outcome Data Collection and Professional Profile Matching Methodology

Post-hire outcomes are measured by matching alumni records to professional profiles on a major employment-network platform. The institutional data make this unusually feasible because names, cohorts, majors, and initial placements are all observed.

Matching proceeds in three passes: exact anchors on name+cohort+firm, probabilistic expansion for common names using degree-year-major compatibility, and manual adjudication using job history, title progression, and employer history. Placement-office records provide quasi-unique identifiers, so the final 98% match rate is achieved with high confidence in identity resolution.

F HR Professional Interviews: Theme-Ordered Excerpts

To triangulate the administrative evidence, I conducted semi-structured interviews in 2024 with five senior HR professionals involved in campus recruitment at Tata Consultancy Services, Tata Motors, Aditya Birla Group, IRIS Software Group, and Jindal South West Steel.[13] Together they represent more than 70 years of recruiting experience across IT services, enterprise software, manufacturing, steel/metals, and diversified conglomerates. The interviews covered pool hiring, who conducts personal interviews, what those interviews probe, whether interviewers later supervise hires, and how recruiting teams are evaluated. This is a small convenience sample, so I use it for institutional context rather than stand-alone identification.

Theme 1: Retention and promotion are treated as quality metrics. Across firms, interviewees described retention, growth, and non-automatic promotion as key quality metrics.

> *"It's evaluated on the quality of the overall campus hires and their retention and their cultural fit and their growth in the organization."* (HR head, Aditya Birla Group)

[13] Names are anonymized for confidentiality; firm names are retained for institutional context.

"Performance, it is basically the retention part... we invest a lot of time and money on them... if we lose them, then the entire thing is lost." (HR head, Tata Motors)

"[A]fter six months... we do a thorough evaluation and then we decide who are those top performers... [f]or promotion... some will get, some will not get." (HR head, IRIS Software Group)

"[Past] performance rating... feeds into your promotion... [t]he first promotion kicks in minimum three years after." (HR head, Tata Consultancy Services)

Theme 2: Central pool hiring separates interviewers from eventual supervisors. Interviewees consistently described campus recruitment as central-pool hiring, with weak alignment between interviewers and final supervisors.

"Predominantly campus hiring will be for a central pool. It is not for any specific business." (HR head, IRIS Software Group)

"So it's a pool which is hired... which person will go to which location is not very firmed up... it's not the direct manager who is interviewing them." (HR head, Tata Motors)

"The placements are so random, that it's almost like lottery. Accidentally, you can get one person you may have interviewed, but there's no interviewer matching done when placement happens." (HR head, Tata Consultancy Services)

Theme 3: Process metrics dominate, with weak downstream feedback. Interviewees emphasized turnaround time and conversion metrics, with little interviewer-level accountability for later performance.

"For talent acquisition team, it's mainly the TAT [turnaround time]... how much time they took to hire... and how many accepted offers really joined. Their performance is not linked to the performance of the hired candidates." (HR head, Tata Motors)

"It is almost impossible to track back their performance, their appraisal or their promotion to the recruitment branch." (HR head, Tata Consultancy Services)

"After six months of probation we do an evaluation... I do not think too much of data goes to recruitment after six months... beyond that, the ownership will be with the business." (HR head, IRIS Software Group)

"There is no qualitative parameters to assess the hiring process itself, except looking at the offer drop ratio." (HR head, Jindal South West Steel)

Theme 4: Biographical probing is embedded in fit assessment. All five interviewees described final personal rounds as fit assessments. Some reported direct questions on family background and schooling; others emphasized rapport, communication, and motivation.

> *"In the HR round, we do ask for family backgrounds... we ask about their education backgrounds, where they have studied... and what kind of background they come from."* (HR head, Tata Motors)

> *"[When discussing communication expectations, some candidates] came from backgrounds where their parents were working as domestic help or driver... you can't expect them to read encyclopedias and Enid Blyton."* (HR head, Tata Consultancy Services)

> *"[Do things like hobbies and schooling come up in conversation?] Yes, it does. It does for sure."* (HR head, IRIS Software Group)

> *"These are more aptitude and attitude and aspiration related questions... what's the fit in the organization, that's the basic... [and when asked if interviewers can ask about family background and schooling:] there are no issues on any questions."* (HR head, Aditya Birla Group)

Taken together, these interviews line up with the transcript evidence in Section 6.4: personal interviews are a setting in which fit is central and biographical markers can become salient, even if firms do not ask identical questions.

Comparative literature anchor. These excerpts align with Indian evidence that caste-linked evaluative norms can persist in ostensibly meritocratic corporate settings (Jodhka, 2008; Jodhka and Newman, 2007; Jodhka, 2015; Shakthi, 2023), and with broader evidence that elite employers often weight rapport and cultural matching heavily in final rounds (Rivera, 2012, 2015; Laurison and Friedman, 2019). The external literature does not identify caste-specific channels; it helps explain why biographical probing can become decisive in discretionary interviews.

G Robustness to Functional Form

This section checks whether the results depend on linear controls for GPA, entrance scores, and school marks. I compare the baseline specification to three more flexible alternatives: natural splines, decile indicators, and—where feasible—saturated department $\times$ year $\times$ stream fixed effects. I report the results for success rates (Section G.1), job quality (Section G.2), post-hire outcomes (Section G.3), and the surname-split analyses (Sections G.4 and G.5). The coefficients are stable throughout.

G.1 Success Rates

Table G.1 reports the coefficient on disadvantaged caste from Equation 1 under alternative functional forms. The natural-spline specification uses four degrees of freedom for each

continuous ability control, a standard flexible choice (Harrell, 2015); the decile specification replaces each continuous measure with decile indicators.

The estimates barely move. At the interview-to-offer stage, the coefficient ranges from -0.293 to -0.298, compared with -0.294 in the baseline. Early-stage effects remain economically small. Functional-form assumptions therefore do not drive the success-rate results.

Table G.1: **Success Rate Robustness: Functional Form Specifications**

	App → Test	Test → Debate	Debate → Int	Int → Offer
	(1)	(2)	(3)	(4)
Panel A: Main Specification (Linear Controls)				
Disadv. Caste	-0.003	-0.000	-0.000	-0.294***
	(0.002)	(0.002)	(0.005)	(0.011)
N	305,541	183,150	148,489	38,378
Panel B: Natural Splines (df=4)				
Disadv. Caste	0.018***	0.015***	-0.003	-0.295***
	(0.002)	(0.002)	(0.006)	(0.011)
N	305,541	183,150	148,489	38,378
Panel C: Decile Bins				
Disadv. Caste	0.012***	0.011***	-0.001	-0.293***
	(0.002)	(0.002)	(0.006)	(0.011)
N	305,541	183,150	148,489	38,378
Panel D: Saturated FE (Dept × Year × Stream)				
Disadv. Caste	-0.002	-0.001	-0.001	-0.298***
	(0.002)	(0.002)	(0.005)	(0.010)
N	305,541	183,150	148,489	38,378

Notes: Table G.1 reports the coefficient on disadvantaged caste (SC/ST/OBC indicator) from linear probability models of advancing to the next hiring stage, estimated separately for each stage. The dependent variable is an indicator for advancing. Panel A uses the main specification with linear controls: college GPA, entrance exam score (JEE rank percentile), 10th and 12th grade marks, prior internship sector, department fixed effects, degree stream fixed effects, and cohort-year fixed effects. Panel B replaces the four ability measures with natural splines (df=4). Panel C replaces them with decile indicators. Panel D omits the four ability measures entirely and replaces the separate department, stream, and year fixed effects with saturated department × year × stream fixed effects (i.e., a full set of triple-interaction dummies), so that all ability variation is absorbed nonparametrically within cells. All specifications include job fixed effects. Standard errors are two-way clustered at the student and job levels. *** p<0.01, ** p<0.05, * p<0.1.

G.2 Job Quality

Table G.2 applies the same exercise to job quality. Because salaries are standardized within jobs, these regressions omit job fixed effects and compare the baseline specification to splines and decile bins only.

The offer-stage coefficient ranges from -0.089 to -0.105, with decile bins at -0.103. The job-quality results are therefore likewise insensitive to functional form.

Table G.2: **Job Quality Robustness: Functional Form Specifications**

	App (1)	Test (2)	Debate (3)	Interview (4)	Offer (5)	Choice (6)
Panel A: Main Specification (Linear Controls)						
Disadv. Caste	0.004	−0.009	−0.006	−0.018*	−0.105***	−0.103***
	(0.005)	(0.007)	(0.007)	(0.010)	(0.013)	(0.017)
N	4,005	3,618	3,266	3,031	2,527	2,527
Panel B: Natural Splines (df=4)						
Disadv. Caste	0.005	0.000	0.004	−0.008	−0.089***	−0.085***
	(0.005)	(0.007)	(0.007)	(0.010)	(0.013)	(0.017)
N	4,005	3,618	3,266	3,031	2,527	2,527
Panel C: Decile Bins						
Disadv. Caste	0.003	−0.003	0.001	−0.017*	−0.103***	−0.103***
	(0.005)	(0.007)	(0.007)	(0.010)	(0.013)	(0.017)
N	4,005	3,618	3,266	3,031	2,527	2,527

Notes: Table G.2 reports the coefficient on disadvantaged caste (SC/ST/OBC indicator) from regressions of log average salary at the offer stage on caste and controls, following Equation 2. Panel A uses the main specification with linear controls: college GPA, entrance exam score (JEE rank percentile), 10th and 12th grade marks, prior internship sector, department fixed effects, degree stream fixed effects, and cohort-year fixed effects. Panel B replaces the four ability measures with natural splines (df=4). Panel C replaces them with decile indicators. Job fixed effects are omitted because the dependent variable is at the student level (an average across jobs), making job identifiers undefined. Standard errors are HC1 robust. *** $p<0.01$, ** $p<0.05$, * $p<0.1$.

G.3 Post-Hire Outcomes

Table G.3 extends the same exercise to promotions and first-job exit. The coefficients are stable: promotion ranges from 0.096 to 0.117, and leaving first job ranges from −0.170 to −0.177. The post-hire advantage is therefore not a functional-form artifact.

Table G.3: **Post-Hire Robustness: Functional Form Specifications**

	Promoted (1)	Left Job (2)
Panel A: Main Specification (Linear Controls)		
Disadv. Caste	0.096***	-0.170***
	(0.030)	(0.046)
N	2,489	2,489
Panel B: Natural Splines (df=4)		
Disadv. Caste	0.104***	-0.177***
	(0.032)	(0.046)
N	2,489	2,489
Panel C: Decile Bins		
Disadv. Caste	0.117***	-0.177***
	(0.031)	(0.045)
N	2,489	2,489

Notes: Table G.3 reports the coefficient on disadvantaged caste (SC/ST/OBC indicator) from post-hire outcome regressions. Column (1), Promoted (Panel A), and Column (2), Left Job (Panel B), are estimated on the same sample. Row groups A–C apply the main linear-controls specification, natural splines (df=4), and decile-bin controls, respectively. All specifications include job fixed effects, department fixed effects, degree stream fixed effects, and cohort-year fixed effects. Standard errors are clustered at the job level. $^{*}p<0.1$, $^{**}p<0.05$, $^{***}p<0.01$.

G.4 Success Rates by Surname Identifiability

Table G.4 repeats the success-rate analysis separately for students with caste-identifiable surnames (Panel A) and non-identifiable surnames (Panel B), corresponding to the surname-identifiability analysis in Section 6.4.2. The timing pattern is stable: in Panel A, early penalties remain large and the interview coefficient ranges from -0.224 to -0.246; in Panel B, early coefficients remain near zero while the interview penalty stays around -0.30. The surname-timing result is therefore not driven by functional form.

Table G.4: **Success Rate Robustness by Surname Identifiability**

Panel A: Identifiable Surnames

	App→Test (1)	Test→GD (2)	GD→Int (3)	Int→Offer (4)
Panel A: Main Specification (Linear Controls)				
Disadv. Caste	-0.184***	0.110***	-0.063***	-0.233***
	(0.006)	(0.006)	(0.008)	(0.022)
N	46,163	21,407	18,899	3,932
Panel B: Natural Splines (df=4)				
Disadv. Caste	-0.166***	0.108***	-0.065***	-0.240***
	(0.005)	(0.006)	(0.008)	(0.024)
N	46,163	21,407	18,899	3,932
Panel C: Decile Bins				
Disadv. Caste	-0.179***	0.096***	-0.061***	-0.246***
	(0.005)	(0.007)	(0.010)	(0.024)
N	46,163	21,407	18,899	3,932
Panel D: Saturated FE (Dept × Year × Stream)				
Disadv. Caste	-0.177***	0.122***	-0.063***	-0.224***
	(0.006)	(0.006)	(0.010)	(0.029)
N	46,163	21,407	18,899	3,932

Panel B: Non-Identifiable Surnames

	App→Test (1)	Test→GD (2)	GD→Int (3)	Int→Offer (4)
Panel A: Main Specification (Linear Controls)				
Disadv. Caste	0.026***	-0.011***	0.008	-0.298***
	(0.002)	(0.002)	(0.005)	(0.011)
N	259,378	161,743	129,590	34,446
Panel B: Natural Splines (df=4)				
Disadv. Caste	0.048***	0.006**	0.005	-0.300***
	(0.002)	(0.002)	(0.006)	(0.011)
N	259,378	161,743	129,590	34,446
Panel C: Decile Bins				
Disadv. Caste	0.041***	0.002	0.006	-0.297***
	(0.002)	(0.002)	(0.006)	(0.012)
N	259,378	161,743	129,590	34,446
Panel D: Saturated FE (Dept × Year × Stream)				
Disadv. Caste	0.028***	-0.012***	0.005	-0.302***
	(0.002)	(0.002)	(0.005)	(0.011)
N	259,378	161,743	129,590	34,446

Notes: Table G.4 reports the coefficient on disadvantaged caste from linear probability models of advancing to the next hiring stage, estimated separately for each stage. Panel A restricts the sample to students with caste-identifiable surnames (n=619), while Panel B restricts the sample to students with non-identifiable surnames (n=3,386). Within each panel, row group A uses the main specification with linear controls, and row groups B–D apply alternative functional forms (natural splines, decile bins, saturated fixed effects). All specifications include job fixed effects. Standard errors are two-way clustered at the student and job levels. *** $p<0.01$, ** $p<0.05$, * $p<0.1$.

G.5 Job Quality by Surname Identifiability

Table G.5 shows the same robustness for job quality, again corresponding to Section 6.4.2. Offer-stage penalties remain stable for both identifiable surnames (-0.143 to -0.146) and non-identifiable surnames (-0.082 to -0.100), confirming that the salary patterns are not driven by functional form.

Table G.5: **Job Quality Robustness by Surname Identifiability**

Panel A: Identifiable Surnames

	Offer (1)	Choice (2)
Panel A: Main Specification (Linear Controls)		
Disadv. Caste	-0.146***	-0.107**
	(0.033)	(0.049)
N	378	378
Panel B: Natural Splines (df=4)		
Disadv. Caste	-0.145***	-0.103**
	(0.033)	(0.050)
N	378	378
Panel C: Decile Bins		
Disadv. Caste	-0.143***	-0.098*
	(0.035)	(0.051)
N	378	378

Panel B: Non-Identifiable Surnames

	Offer (1)	Choice (2)
Panel A: Main Specification (Linear Controls)		
Disadv. Caste	-0.100***	-0.106***
	(0.014)	(0.019)
N	2,149	2,149
Panel B: Natural Splines (df=4)		
Disadv. Caste	-0.082***	-0.085***
	(0.014)	(0.018)
N	2,149	2,149
Panel C: Decile Bins		
Disadv. Caste	-0.092***	-0.099***
	(0.014)	(0.019)
N	2,149	2,149

Notes: Table G.5 reports the coefficient on disadvantaged caste from regressions of log average salary on caste and controls. Panel A restricts the sample to students with caste-identifiable surnames (n=378 job-getters); Panel B restricts to students with non-identifiable surnames (n=2,149 job-getters). Column (1) reports offer-stage salary and Column (2) reports final job-choice salary. Within each panel, row group A uses the main specification with linear controls, and row groups B–C apply alternative functional forms (natural splines and decile bins). Offer-stage standard errors are HC1 robust; Choice-stage standard errors are clustered at the job level. *** $p<0.01$, ** $p<0.05$, * $p<0.1$.

H Robustness to Omitted Variable Bias: Oster Bounds

This section asks whether omitted variables could explain the estimated caste gaps. I use the Oster (2019) bounds methodology to compare baseline estimates to bias-adjusted estimates under a demanding calibration: $\delta = 1$ and $R^*_{\max} = 1.3\tilde{R}$ (Stansbury and Rodriguez, 2024). I do not report δ^* row-by-row in the tables because, outside the load-bearing coefficients, it is often mechanically unstable when estimates are near zero and its sign is easy to over-interpret. The main object here is whether the adjusted coefficients materially differ from the baseline ones. They generally do not.

H.1 Success Rates

Table H.1 compares OLS and Oster-adjusted estimates for stage-specific success rates. The interview coefficient barely moves, from -0.294 to -0.296. By contrast, the small pre-interview coefficients are mechanically sensitive and are not substantively interpreted. The omitted-variable result is therefore simple: under the stated calibration, the interview-stage penalty remains essentially unchanged.

Table H.1: **OLS vs. Oster-Adjusted Estimates: Success Rates**

		Disadvantaged Caste	
Stage	N	OLS	Oster-Adjusted
App → Test	305,541	-0.003^{*}	0.009
		(0.001)	(0.001)
Test → Debate	183,150	-0.001	0.010^{***}
		(0.002)	(0.002)
Debate → Int	148,489	-0.001	-0.003
		(0.003)	(0.003)
Int → Offer	38,378	-0.294^{***}	-0.296^{***}
		(0.005)	(0.005)

Notes: Table H.1 compares OLS estimates to bias-adjusted estimates using the Oster (2019) methodology. Each row reports the coefficient on disadvantaged caste (SC/ST/OBC indicator) for the probability of advancing from one hiring stage to the next. For each stage, Oster's short specification is the fixed-effects baseline (job and cohort-year FE only), and the full specification adds observed controls. The Oster adjustment uses $\delta = 1$ (selection on unobservables equals selection on observables) and $R^*_{\max} = 1.3\tilde{R}$ (hypothetical R-squared 30% higher than observed). All specifications include job fixed effects, cohort-year fixed effects, and linear controls: college GPA, entrance exam score (JEE rank percentile), 10th and 12th grade marks, prior internship sector indicators, department fixed effects, and degree stream fixed effects. Standard errors are conventional OLS (IID) and are reported in parentheses. The implied δ^* for the interview stage is -131.9; for the earlier near-zero coefficients, implied δ^* values are unstable or not substantively informative and are therefore omitted from the table. *** p<0.01, ** p<0.05, * p<0.1.

H.2 Job Quality

Table H.2 applies the same exercise to job quality. The offer-stage coefficient moves from -0.105 to -0.079, and the choice-stage coefficient from -0.103 to -0.076, so the late-stage job-quality penalty remains economically large. Early-stage adjusted coefficients are small and not load-bearing. For the central late-stage coefficients, the implied δ^* values exceed one in magnitude (not shown), so omitted variables do not plausibly explain away the job-quality penalty that appears after interviews.

Table H.2: **OLS vs. Oster-Adjusted Estimates: Job Quality**

		Disadvantaged Caste	
Stage	N	OLS	Oster-Adjusted
Application Stage	4,005	0.004	0.005
		(0.005)	(0.005)
Written Test Stage	3,618	−0.009	0.010
		(0.007)	(0.008)
Group Debate Stage	3,266	−0.006	0.011
		(0.007)	(0.008)
Interview Stage	3,031	−0.018*	0.012
		(0.010)	(0.011)
Offer Stage	2,527	−0.105***	−0.079***
		(0.013)	(0.014)
Choice Stage	2,527	−0.103***	−0.076***
		(0.017)	(0.019)

Notes: Table H.2 compares OLS estimates to bias-adjusted estimates using the Oster (2019) methodology. Each row reports the coefficient on disadvantaged caste (SC/ST/OBC indicator) for log average salary of jobs in contention at each hiring stage. For each stage, Oster's short specification includes cohort-year fixed effects only, and the full specification adds observed controls. The Oster adjustment uses $\delta = 1$ and $R^*_{\max} = 1.3\tilde{R}$. All specifications include cohort-year fixed effects and linear controls: college GPA, entrance exam score (JEE rank percentile), 10th and 12th grade marks, prior internship sector indicators, department fixed effects, and degree stream fixed effects. Standard errors are conventional OLS (IID) and are reported in parentheses. The implied δ^* values for the offer and choice stages are 4.0 and 3.7, respectively; earlier-stage values are omitted from the table because those coefficients are small and not central to the paper's interpretation. *** p<0.01, ** p<0.05, * p<0.1.

H.3 Promotion and Tenure Rates

Table H.3 reports Oster bounds for post-hire outcomes. The full-controls promotion coefficient is essentially unchanged (0.096 to 0.093), and the full-controls first-job-exit coefficient is likewise stable (−0.170 to −0.180). The same qualitative pattern holds across selectivity tiers and client-facing status. For the full-controls rows, the implied δ^* values are large in magnitude (29.6 for promotion; −17.9 for first-job exit), so omitted variables do not provide a convincing explanation for the post-hire advantage of disadvantaged hires.

Table H.3: **OLS vs. Oster-Adjusted Estimates: Post-Hire Outcomes**

Panel A: Promotion Rates

		Disadvantaged Caste	
Outcome	N	OLS	Oster-Adjusted
Promoted (All)	2,489	0.137***	0.142***
		(0.025)	(0.026)
Promoted (w/ controls)	2,489	0.096***	0.093***
		(0.026)	(0.028)
Promoted at Top 75%	2,012	0.104***	0.108***
		(0.029)	(0.031)
Promoted at Top 50%	1,246	0.112***	0.122***
		(0.038)	(0.037)
Promoted at Top 40%	997	0.136***	0.141***
		(0.043)	(0.041)
Promoted at Top 30%	748	0.147***	0.157**
		(0.054)	(0.058)
Promoted at Top 20%	502	0.190**	0.196**
		(0.073)	(0.082)
Promoted at Top 15%	374	0.204**	0.206**
		(0.084)	(0.082)
Promoted at Top 10%	252	0.084	0.074
		(0.109)	(0.131)
Promoted (Client Facing)	835	0.107**	0.106***
		(0.044)	(0.039)
Promoted (Non-Client Facing)	1,654	0.082**	0.075**
		(0.032)	(0.036)

Panel B: Job Leaving Rates (Tenure)

		Disadvantaged Caste	
Outcome	N	OLS	Oster-Adjusted
Overall Specifications			
Left Job (Job FE)	2,489	−0.188***	−0.198
Left Job (Job + Year FE)	2,489	−0.186***	−0.197
Left Job (Full Controls)	2,489	−0.170***	−0.180
By Job Selectivity			
Top 75% Selectivity	2,012	−0.124***	−0.139
Top 50% Selectivity	1,246	−0.135***	−0.147
Top 40% Selectivity	997	−0.210***	−0.213
Top 30% Selectivity	748	−0.189***	−0.200
Top 25% Selectivity	623	−0.158**	−0.157
Top 20% Selectivity	502	−0.152*	−0.128
Top 15% Selectivity	374	−0.191**	−0.170
Top 10% Selectivity	252	−0.181	−0.161
By Client-Facing Status			
Consumer Facing	835	−0.067	−0.070
Non-Consumer Facing	1,654	−0.171***	−0.188

Notes: Table H.3 compares OLS estimates to Oster-adjusted estimates for post-hire outcomes. Panel A reports promotion outcomes by selectivity and client-facing status; Panel B reports first-job leaving outcomes by selectivity and client-facing status. For each row, Oster's short specification uses fixed effects only, and the full specification adds observed controls. The Oster adjustment uses $\delta = 1$ and $R^*_{\max} = 1.3\tilde{R}$. Across both panels, coefficients remain stable after adjustment, supporting robustness of post-hire results to omitted-variable bias. Standard errors are conventional OLS (IID) and are reported in parentheses where shown. For the full-controls rows, the implied δ^* values are 29.6 for promotion and −17.9 for first-job exit; the remaining non-zero rows show the same qualitative pattern, so δ^* is omitted from the table to save space and avoid over-emphasizing sign changes that are not central to interpretation. *** p<0.01, ** p<0.05, * p<0.1.

H.4 Job Quality by Surname Identifiability

Table H.4 applies the same exercise separately by surname-identifiability sample. The timing pattern remains intact after adjustment: identifiable surnames show earlier penalties, while non-identifiable surnames show near-zero pre-interview coefficients and a robust offer-stage penalty. The omitted δ^* values at the offer stage are 5.6 and 3.8, respectively. The Oster results therefore reinforce the caste-visibility interpretation.

Table H.4: **OLS vs. Oster-Adjusted Estimates: Job Quality by Surname Identifiability**

Panel A: Identifiable Surnames

		Disadvantaged Caste	
Stage	N	OLS	Oster-Adjusted
Application	619	0.012	0.013
		(0.014)	(0.016)
Written Test	552	-0.068^{***}	-0.049^{***}
		(0.019)	(0.021)
Group Debate	500	-0.034^{*}	-0.016
		(0.018)	(0.022)
Interview	449	-0.105^{***}	-0.078^{**}
		(0.027)	(0.032)
Offer	378	-0.146^{***}	-0.119^{***}
		(0.033)	(0.042)

Panel B: Non-Identifiable Surnames

		Disadvantaged Caste	
Stage	N	OLS	Oster-Adjusted
Application	3,386	0.002	0.004
		(0.005)	(0.006)
Written Test	3,066	0.0004	0.019^{***}
		(0.007)	(0.007)
Group Debate	2,766	-0.003	0.013
		(0.007)	(0.008)
Interview	2,582	-0.004	0.024^{**}
		(0.011)	(0.011)
Offer	2,149	-0.100^{***}	-0.075^{***}
		(0.014)	(0.015)

Notes: Table H.4 compares OLS estimates to Oster-adjusted estimates separately by surname-identifiability sample. Panel A (identifiable surnames, n=619) shows penalties emerging at written tests and persisting through offers. Panel B (non-identifiable surnames, n=3,386) shows near-zero pre-interview effects and a robust offer-stage penalty. All specifications include cohort-year fixed effects and linear controls. The Oster adjustment uses $\delta = 1$ and $R^*_{\max} = 1.3\tilde{R}$. Standard errors are conventional OLS (IID) and are reported in parentheses. The implied δ^* values at the offer stage are 5.6 in Panel A and 3.8 in Panel B; elsewhere δ^* is either unstable because coefficients are near zero or adds little beyond the displayed adjusted estimates, so it is omitted from the table. *** p<0.01, ** p<0.05, * p<0.1.

I A Model Exploiting Plausibly Exogenous Variation in Selectivity Across Jobs to Identify Discrimination

This section formalizes the selectivity argument used in the main text. The empirical problem is to distinguish a higher interview threshold from unobserved ability differences. The key facts are the absence of early-stage gaps and the stable post-hire advantage of disadvantaged hires across selectivity tiers. The stylized model shows why those facts are hard to reconcile with equal interview thresholds.

Consider jobs $j \in \mathcal{J}$ with four sequential screening rounds $s \in \{1, 2, 3, 4\}$, where $s = 1, 2, 3$ are pre-interview stages and $s = 4$ is the personal interview. There are two caste groups $g \in \{0, 1\}$, where $g = 1$ denotes disadvantaged castes and $g = 0$ denotes advantaged castes. Candidates draw stage-specific screening scores $M_{s,i,g}$ from distributions $F_{s,g}$, common across jobs within each stage after residualizing on observables. Jobs set thresholds $\tau_{s,j}$ for each round, except that interview thresholds may differ by caste, $\tau_{4,j}^{(g)}$. A candidate passes round s if $M_{s,i,g} \geq \tau_{s,j}$ and they have passed all preceding rounds. A caste penalty at interviews is indicative of discrimination if $\tau_{4,j}^{(1)} > \tau_{4,j}^{(0)}$.

I use five assumptions.

Assumption 1 (Perfect information). Ability is perfectly proxied by observed screening scores at each stage. This simplifies exposition; the main threshold-comparison logic continues to hold under noisy signals, as shown in Appendix I.1.

Assumption 2 (Regularity conditions). The performance distributions $F_{s,g}$ are continuous with smooth densities and finite moments.

Assumption 3 (Monotonicity of thresholds). Thresholds are weakly increasing across rounds:

$$\tau_{4,j} \;\geq\; \tau_{3,j} \;\geq\; \tau_{2,j} \;\geq\; \tau_{1,j}.$$

This comparison is coherent because stage scores are implicitly mapped to a common latent ability scale.

Assumption 4 (Common talent pool across jobs). For each screening stage $s \in \{1, 2, 3, 4\}$ and caste group $g \in \{0, 1\}$, let

$$U_{s,i,g,j} \equiv M_{s,i,g,j} - m_s(X_i),$$

where X_i denotes the observed controls. There exist job-invariant distributions $\{F_{s,g}\}$ such that

$$U_{s,i,g,j} \sim F_{s,g} \qquad \text{for all eligible jobs } j,$$

with support independent of j.

This assumption is plausible in our setting because the centralized portal, low marginal application costs, and binding eligibility rules sharply limit caste-based sorting across jobs (Section 5). It does *not* require equal group means or variances. The narrower claim is that, after conditioning on observables, more selective jobs are not systematically drawing from weaker disadvantaged pools and stronger advantaged pools. Section 6.2 supports that reading: pre-interview caste gaps remain near zero even in the most selective jobs.

Assumption 5 (Light tails). For each caste g and stage s, there exist constants $C_g > 0$ and $\beta > 0$ such that

$$\Pr\big(M_{s,i,g} \geq m\big) \;\leq\; C_g\, e^{-\beta m}.$$

If support is unbounded, assume the hazard rate $h_{s,g}(m) \equiv f_{s,g}(m)/[1 - F_{s,g}(m)]$ diverges as $m \to \infty$. If support is finite with upper endpoint $\bar{m}_{s,g}$, interpret this as $h_{s,g}(m) \to \infty$ as $m \uparrow \bar{m}_{s,g}$. Because screening scores are bounded in practice, the finite-support case is the empirically relevant one here.

These assumptions imply two corollaries.

Corollary 1 (Exogeneity of cross-job threshold variation). Under Assumptions 1–5, the institutional features documented in Section 5, and the fact that caste is typically not salient prior to interviews (Section 6.4), cross-job variation in pre-interview thresholds is orthogonal to group identity and caste-specific unobservables conditional on observables:

$$\tau_{s,j} \perp G_i \mid X_i \quad \text{and} \quad \mathbb{E}\big[U_{s,i,g,j} \mid \tau_{s,j}, X_i, G_i\big] = \mathbb{E}\big[U_{s,i,g,j} \mid X_i, G_i\big], \qquad s < 4.$$

The first condition is the identifying restriction motivated by early-stage caste invisibility; the second is the mean-independence implication of Assumption 4. Mean independence is sufficient for the propositions below.

Corollary 2 (Interview cutoff as a valid instrument). Under the same assumptions and organizational separation between hiring and promotion (HR sets $\tau_{4,j}$; line managers decide promotions later), cross-job variation in the interview cutoff $\tau_{4,j}$ is a valid instrument for the composition of hires: *(i) relevance* $\partial P(\text{Hired}_{j,g} \mid \tau_{4,j})/\partial \tau_{4,j} \neq 0$; *(ii) exclusion* for any job j, group g, and ability index a, the promotion mapping $\pi_{j,g}(a)$ does not depend on $\tau_{4,j}$ once (j, a, g) are fixed.

Interpreting the Main Empirical Patterns

The logic is simple. If disadvantaged candidates had a meaningful pre-interview right-tail deficit, that deficit should appear in high-cutoff jobs before interviews. If they instead face a higher interview threshold, stable post-hire outperformance can persist even at high selectivity.

Proposition 1 (No Early Gap Implies No Right-Tail Ability Deficit Before Interviews). Consider a job j with four rounds of screening, where $s = 1, 2, 3$ are pre-interview stages. Under Assumptions 1–5, if disadvantaged castes ($g = 1$) were genuinely weaker in the right tail before interviews, then for a sufficiently high threshold $\tau_{s,j}$ at some pre-interview stage $s < 4$, the within-job conditional pass-rate gap

$$\Delta_{s,j} \equiv P(\text{pass}_{s,j} \mid g = 1) - P(\text{pass}_{s,j} \mid g = 0)$$

would become negative. Conversely, observing $\Delta_{s,j} \geq 0$ even in the most selective jobs rules out any meaningful right-tail deficit for disadvantaged castes before interviews.

Sketch of Proof. Figure I.1 illustrates the argument. Under Corollary 1, cross-job variation in pre-interview thresholds probes progressively deeper into the common right tail of both groups' score distributions.

- **Panel (a): Right-tail deficit.** If disadvantaged castes had a thinner pre-interview right tail, higher cutoffs would eventually generate negative pass-rate gaps.
- **Panel (b): No right-tail deficit.** Section 6.2 instead shows near-zero pre-interview gaps even in the most selective jobs.

Figure I.1: **Testing for Pre-Interview Performance Differences**

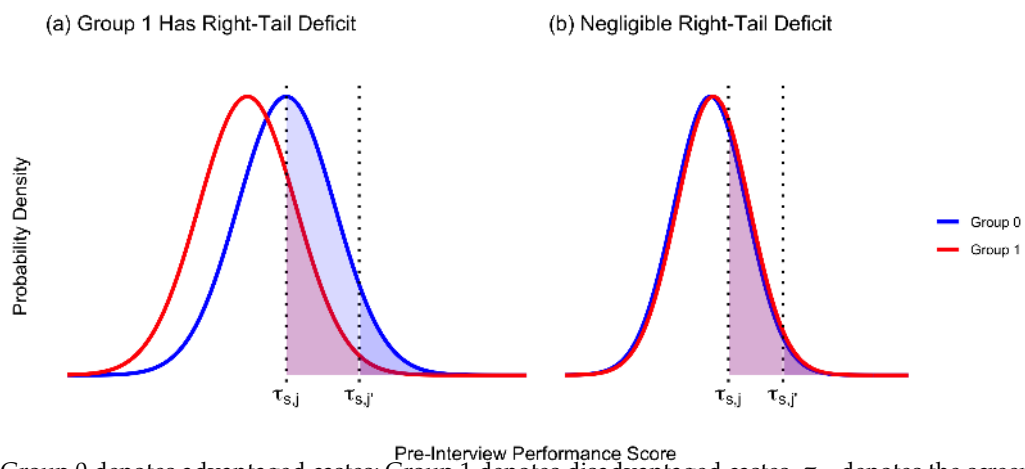

Notes: Group 0 denotes advantaged castes; Group 1 denotes disadvantaged castes. $\tau_{s,j}$ denotes the screening threshold at stage $s < 4$ for job j; j' represents a more selective job. Shaded regions show candidates who pass the thresholds (blue for Group 0, red for Group 1, purple where they overlap). Panel (a) shows how a right-tail deficit for Group 1 would become more visible at higher cutoffs. Panel (b) shows the pattern implied by the data in Section 6.2: no meaningful pre-interview right-tail deficit.

□

Ruling out Cumulative Performance Deficits

The same logic rules out cumulative pre-interview deficits. Any broad disadvantage that accumulates across rounds must appear somewhere before the interview stage, especially in high-cutoff jobs. Section 6.2 shows that it does not.

Do Interviews Reveal Unobserved Ability Differences That Justify Disparities?

Once Proposition 1 rules out meaningful pre-interview deficits, the late-stage gap must reflect either interview-stage information revelation or a higher interview threshold. The test below uses post-hire performance across job selectivity to distinguish between them.

Let $\bar{M}_{i,g}$ denote a monotone aggregate of a candidate's four stage scores; for exposition I use their average. Assume a single-index structure,

$$M_{s,i,g} = f_s(A_{i,g}) + \varepsilon_{s,i,g},$$

where $A_{i,g}$ is a common latent ability draw and $\varepsilon_{s,i,g}$ is small stage-specific noise. Candidate i is promoted if $\bar{M}_{i,g} \geq \kappa_j$, where κ_j is a job-specific promotion cutoff. For jobs in selectivity tier q, let $\tau^{q}_{4,j}$ denote the common interview cutoff absent discrimination and $\tau^{(g)}_{4,j}$ the actual group-specific cutoff. Define

$$\bar{Y}^{(g)}_{j} = \mathbb{E}[\bar{M}_{i,g} \mid M_{4,i,g} \geq \tau^{(g)}_{4,j}],$$

and let the promotion gap in selectivity tier q be

$$\Delta^{\text{Prom}}(q) = \bar{Y}_j^{(1)} - \bar{Y}_j^{(0)}.$$

Under any monotone promotion rule, a positive $\Delta^{\text{Prom}}(q)$ implies a corresponding promotion-rate advantage.

Proposition 2 (Stable Promotion Gaps Across Selectivity Tiers Indicate Discrimination). Fix a selectivity tier q. Under Assumptions 1–5:

1. *No discrimination.* If $\tau_{4,j}^{(1)} = \tau_{4,j}^{(0)} = \tau_{4,j}^{q}$ for all jobs in tier q, then

$$\Delta^{\text{Prom}}(q) \to 0 \qquad \text{as } q \to 1.$$

2. *Discrimination.* If $\tau_{4,j}^{(1)} = \tau_{4,j}^{(0)} + \delta$ with $\delta > 0$, then

$$\Delta^{\text{Prom}}(q) \to \delta > 0 \qquad \text{as } q \to 1,$$

under the tail normalization described above.

Thus, a positive and stable promotion gap as job selectivity increases is more consistent with discrimination than with equal interview thresholds plus latent ability differences.

Sketch of Proof. The key idea is average-to-marginal convergence. By Corollary 2, the interview cutoff varies across jobs and affects promotions only through who gets hired.

- **Panel (a): No discrimination.** If both groups face the same cutoff, selected means converge to the same marginal threshold as selectivity rises, following the mean-residual-life result in Proposition 5, so the gap vanishes.
- **Panel (b): Discrimination.** If disadvantaged castes face a higher interview threshold, their selected mean converges to a higher margin and the post-hire performance gap remains positive.

Figure I.2 illustrates the two cases. Assumption 4 is what makes this informative across jobs: higher cutoffs index a common margin rather than different applicant pools. Combined with the organizational separation between hiring and later promotion (Section 6.3), cross-job selectivity is therefore informative about discrimination rather than composition.

Figure I.2: **Testing for Discrimination**

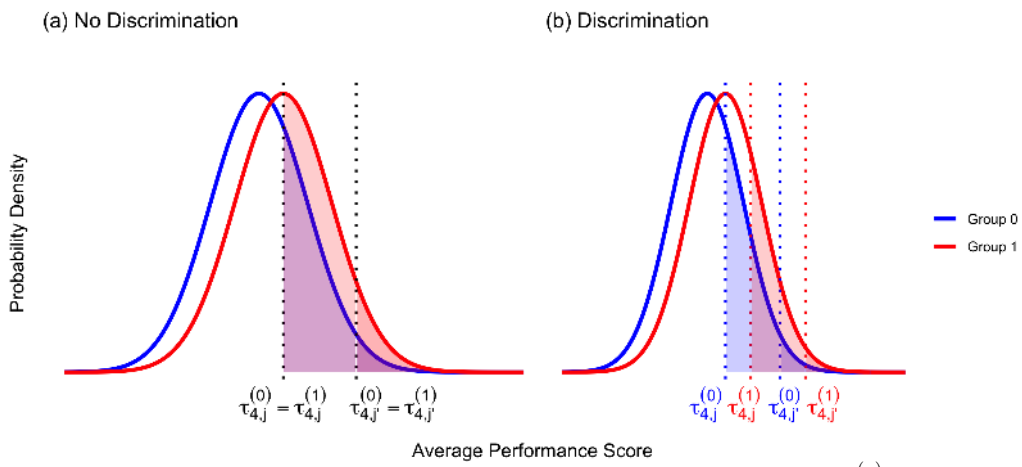

Notes: Group 0 denotes advantaged castes; Group 1 denotes disadvantaged castes. $\tau^{(g)}_{4,j}$ denotes group g's interview threshold; j' represents a more selective job than j. Shaded regions show candidates who pass the thresholds (blue for Group 0, red for Group 1, purple where they overlap). Panel (a) shows that equal thresholds imply vanishing performance gaps as selectivity increases. Panel (b) shows that a constant threshold wedge $\delta > 0$ preserves a stable performance advantage for Group 1.

□

I.1 Extending the Simple Model of Late-Stage Discrimination with Imperfect Information

I now relax the perfect-information assumption. The intuition is unchanged.

Setup. As before, there are two groups $g \in \{0, 1\}$, a continuum of jobs $j \in \mathcal{J}$, and four sequential screening rounds. Let $A_{i,g}$ denote latent ability. At each stage,

$$M_{s,i,g} = A_{i,g} + \eta_{s,i,g},$$

where $\eta_{s,i,g}$ is a stage-s innovation. Let $\mathcal{I}_s$ denote the employer's information set after observing signals through round s. Passing stage s requires

$$\mathbb{E}[A_{i,g} \mid \mathcal{I}_s] \geq \bar{\tau}_{s,j},$$

which, under symmetric mean-zero noise, is equivalent to thresholding on the realized score $M_{s,i,g} \geq \tau_{s,j}$.

Assume:

(i) $\mathbb{E}[\eta_{s,i,g} \mid \mathcal{I}_{s-1}] = 0$ for all s, i, g;

(ii) $\eta_{s,i,1} \overset{d}{=} \eta_{s,i,0}$ for every stage s.

Noise may be serially correlated in raw performance, but it does not create systematic group-level bias. I maintain Assumptions 2–5 from the main model.

Proposition 3 (No Early Gap Implies No Right-Tail Deficit Before Interviews). Consider any pre-interview stage $s < 4$. Under Assumptions 2–5, if disadvantaged castes truly have fewer high-ability candidates in the right tail, then for sufficiently high $\tau_{s,j}$,

$$P(\text{pass}_{s,j} \mid g = 1) - P(\text{pass}_{s,j} \mid g = 0) < 0.$$

Therefore, observing no such gap even among highly selective jobs rules out a meaningful right-tail deficit before interviews.

Sketch of Proof. The argument is unchanged. Under the common-pool assumption, higher pre-interview thresholds probe deeper into the same group-specific tails across jobs. Symmetric noise cannot systematically hide a large latent right-tail deficit at high cutoffs, so the empirical absence of early gaps still rules out a meaningful pre-interview deficit. □

Given Proposition 3, any late-stage gap must again come from interview-stage information revelation or from discrimination. To distinguish between them, consider post-hire performance.

Proposition 4 (Stable Promotion Gaps Identify Discrimination). Under Assumptions 2–5:

1. *No discrimination.* If $\tau_{4,j}^{(1)} = \tau_{4,j}^{(0)}$, then promotion gaps vanish as selectivity increases.

2. *Discrimination.* If $\tau_{4,j}^{(1)} = \tau_{4,j}^{(0)} + \delta$ with $\delta > 0$, then the group-1 minus group-0 average-performance gap converges to a positive limit proportional to δ (equal to δ under the normalization in Proposition 2).

Hence, persistent promotion gaps at high selectivity are difficult to reconcile with an unobserved-ability explanation and are more consistent with discrimination under the maintained assumptions.

Sketch of Proof. The same average-to-marginal logic applies under imperfect information. If both groups face the same threshold, selected average ability converges to the same margin and the promotion gap vanishes. If group 1 faces a threshold wedge δ, the selected average remains higher by a positive amount proportional to δ. Symmetric noise alone cannot generate a stable positive gap at high selectivity. □

I.2 Derivation of Expected Performance Above Threshold

This appendix establishes the tail result used in Proposition 2,

$$\lim_{\tau \to \infty} \big(\mathbb{E}[M \mid M \geq \tau] - \tau \big) = 0.$$

Setup. Let M be a real-valued random variable with CDF F, PDF f, survival function $S(a) = 1 - F(a)$, and hazard rate $h(a) = f(a)/S(a)$ wherever $S(a) > 0$.

Assumption A. M has a continuously differentiable density, finite mean, and diverging hazard rate:

$$\lim_{a\to\infty} h(a) = \infty.$$

This covers familiar light-tailed families such as the Normal, Logistic, and Weibull with shape parameter > 1, but excludes constant-hazard cases such as the Exponential.

Proposition 5 (Mean of a Right-Truncated Distribution)**.** For any τ with $S(\tau) > 0$, define the mean residual life

$$m(\tau) \equiv \mathbb{E}[M - \tau \mid M \geq \tau] = \frac{\int_\tau^\infty S(u)\,du}{S(\tau)}.$$

Under Assumption A,

$$\lim_{\tau\to\infty} m(\tau) = 0, \qquad \text{hence} \qquad \mathbb{E}[M \mid M \geq \tau] = \tau + m(\tau) \to \tau.$$

Proof. Since

$$m(\tau) = \frac{\int_\tau^\infty S(u)\,du}{S(\tau)},$$

L'Hôpital's rule gives

$$\lim_{\tau\to\infty} m(\tau) = \lim_{\tau\to\infty} \frac{\int_\tau^\infty S(u)\,du}{S(\tau)} = \lim_{\tau\to\infty} \frac{-S(\tau)}{-f(\tau)} = \lim_{\tau\to\infty} \frac{S(\tau)}{f(\tau)} = \lim_{\tau\to\infty} \frac{1}{h(\tau)} = 0$$

under Assumption A. □

Discussion. The key requirement is a diverging hazard rate. Heavy-tailed or constant-hazard cases such as the Exponential do not satisfy it. In our setting, screening scores are effectively bounded, so the required tail regularity is the empirically relevant case.

J LLM-Based Transcript Analysis Methodology

This appendix documents the transcript analysis used to triangulate the interview mechanism. The corpus is candidate-consumed mock-preparation material rather than the focal firms' realized interviews, so the exercise is not stand-alone identification. Its value is comparative: the same taxonomy is applied to HR interviews and group debates, making it possible to test whether the two formats systematically differ in whether they elicit biography or policy discussion.

J.1 Key Design Principle: Unified Taxonomy

The core design choice is to apply the *same* classification scheme to both transcript types. The classifier records 21 raw topic dimensions—10 policy topics and 11 biographical topics, including college education. For the main comparison, college education is excluded, leaving a 20-topic analytic set. The key coding rule distinguishes biography *as the topic* from biography

used only as rhetorical context for a policy claim; this is what prevents false positives in group debates.

J.2 Data Collection

I collected mock interview and group-debate transcripts from YouTube, the dominant interview-preparation platform in India. Search queries combined general terms with company-specific queries for firms in the recruiting sample. To focus on material candidates actually consume, I retained only videos with at least 1,000 views.

Transcripts were extracted using commercial transcription services, checked manually, and de-duplicated using content-based hashing. The full corpus contains 390 transcripts. After filtering to live mock sessions and excluding tips/meta-content, the analysis sample contains 274 live transcripts: 148 group debates (622,322 words) and 126 HR interviews (247,956 words), for a total of roughly 870,000 words.

Filtering to Live Transcripts

Many preparation videos are advice or commentary rather than mock interactions. After transcript-type classification (Section J.3), I retain only live mock sessions. All main-text results use this live subset.

J.3 Classification Prompts

The classification uses GPT-4o-mini via the OpenAI API with structured JSON outputs and temperature 0.1. Four features matter: a unified taxonomy for HR and GD transcripts, an explicit biography-versus-context distinction, descriptive rather than evaluative coding, and structured outputs that can be aggregated systematically.

System Prompt

```
You are analyzing interview/discussion transcripts to identify what topics are
covered.

Your task is to objectively identify:  1.  HIGH-LEVEL FOCUS: Is this primarily
about the speaker/candidate personally, or about society/policy?  2.  POLICY
TOPICS: What societal/policy topics are discussed (if any)?  3.  BIOGRAPHICAL
TOPICS: What personal facts about the candidate are extracted (if any)?
```

```
CRITICAL DISTINCTION FOR BIOGRAPHICAL TOPICS:

BIOGRAPHICAL = The conversation EXTRACTS personal facts about the candidate
AS THE MAIN TOPIC. - ‘‘Where are you from?’’  → ‘‘I’m from Jaipur’’ = BIOGRAPHICAL
(hometown is the topic) - ‘‘What does your father do?’’  → ‘‘He’s a farmer’’ =
BIOGRAPHICAL (parents’ occupation is the topic) - ‘‘Tell me about your family’’ →
‘‘I have two siblings’’ = BIOGRAPHICAL (family structure is the topic)
```

NOT BIOGRAPHICAL = Personal experience used to discuss a SOCIETAL topic. - ‘‘In my experience at college, reservation policy needs reform’’ = NOT biographical (policy is the topic) - ‘‘When I worked at Tata, I saw how AI is transforming manufacturing’’ = NOT biographical (technology is the topic) - ‘‘As someone from a rural area, I think infrastructure investment is crucial’’ = NOT biographical (economic policy is the topic)

The key question: Is the TOPIC the person’s biography, or is personal experience just CONTEXT for discussing society/policy?

Apply the SAME analysis to all transcripts regardless of format. Be thorough and objective.

Unified Classification Prompt (Same for HR and GD)

Analyze this transcript and identify ALL topics covered.

IMPORTANT: Apply the same analysis regardless of transcript type (GD or HR interview).

STEP 1 - TRANSCRIPT TYPE: - live_gd: A live group debate with participants debating - live_hr: A live HR/personal interview with Q&A - tips: Coaching/tips content about interviews or GDs - meta: Discussion about interviews/GDs without actual content - other: Something else

STEP 2 - HIGH-LEVEL CLASSIFICATION: Based on the DOMINANT content (80%+ of transcript), is this: - about_speaker: Primarily extracts/discusses personal information about the candidate - about_society: Primarily discusses societal/policy topics - mixed: Genuinely balanced (40-60% each) - use sparingly - If one side is 60-80%, classify to that dominant side (about_speaker or about_society) and mark confidence as medium.

STEP 3 - DETAILED TOPIC ANALYSIS: Identify which specific topics are covered. Use the SAME criteria for both GD and HR.

CRITICAL: For biographical topics, mark TRUE only if personal info is EXTRACTED AS THE TOPIC. Do NOT mark TRUE if someone merely mentions personal experience to argue about policy.

Return ONLY a valid JSON object with these top-level keys: - transcript_type: live_gd | live_hr | tips | meta | other - content_focus: about_speaker | about_society | mixed - content_focus_confidence: high | medium | low - dominant_content_pct: estimated share statement - policy_topics: nested object with {discussed, detail} by topic - biographical_topics: nested object with

```
{discussed, detail} by topic
```

```
Policy-topic keys:  economic_policy, education_policy, technology_ai,
gender_women, environment, governance_democracy, media_social_media,
society_culture, food_policy, other_policy.

Biographical-topic keys:  hometown_origins, hobbies_interests, high_school,
family_structure, native_vs_current, parents_occupation, school_type,
coaching_tuition, medium_of_schooling, school_board, college_education.

For reported ‘‘biographical markers’’ figures/tables, college_education is
excluded, leaving a 10-marker set used in the main interpretation.
```

J.4 Classification Process

The classification has two layers. Each transcript first receives a dominant-content label ("about_speaker," "about_society," or "mixed"), and the classifier then records which of the 20 analytic topics appear.

J.5 Results Summary

All results are based on the 274 live transcripts. The contrast is informative because both formats are communication-intensive, but only one systematically makes personal biography the topic of evaluation.

High-Level Content Classification

Table J.1: Content Focus Classification

Content Focus	**HR Interviews** (n=126)	**Group Debates** (n=148)
About Speaker (Personal/Biographical)	89.7%	0.7%
About Society (Policy/Debate)	3.2%	99.3%
Mixed	7.1%	0.0%

Notes: Table J.1 shows transcript-level dominant-content classification for 274 live transcripts. "About Speaker" indicates that biography or personal extraction is the dominant topic; "About Society" indicates that policy or societal discussion is dominant; "Mixed" is used only for genuinely balanced cases. Classification uses GPT-4o-mini and the unified prompt described in Section J. Cross-LLM validation appears below.

Detailed Topic Analysis: All 20 Dimensions

Tables J.2 and J.3 report the full topic-level breakdown. Read together, they show a structural difference: group debates concentrate on policy discussion, whereas personal interviews disproportionately extract personal background information.

Table J.2: Policy Topics (10 Categories)

Policy Topic	**HR Interviews** (n=126)	**Group Debates** (n=148)
Economic Policy	4.8%	55.4%
Education Policy	2.4%	33.8%
Technology & AI	25.4%	29.1%
Gender & Women's Rights	4.8%	18.2%
Environment	8.7%	10.1%
Governance & Democracy	1.6%	14.9%
Media & Social Media	2.4%	13.5%
Society & Culture	8.7%	88.5%
Food Policy	0.8%	6.8%
Other Policy Topics	1.6%	1.4%

Notes: Table J.2 reports the share of transcripts in which each policy topic appears. A topic is marked as present when the classifier detects substantive discussion of that topic. The same unified taxonomy is applied to both HR and GD transcripts.

Table J.3: Biographical Topics (10 Categories)

Biographical Topic	**HR Interviews** (n=126)	**Group Debates** (n=148)
Hometown/Geographic Origins	69.8%	7.4%
Hobbies & Interests	67.5%	0.7%
High School Attended	48.4%	4.7%
Family Structure	27.0%	2.7%
Parents/Family Occupation	21.4%	0.7%
Native Place vs Current City	11.1%	1.4%
School Type (Govt/Private/Convent)	7.9%	1.4%
Coaching/Tuition	3.2%	0.0%
Medium of Schooling	2.4%	0.7%
School Board (CBSE/ICSE/State)	1.6%	0.0%

Notes: Table J.3 reports the share of transcripts in which each biographical topic is directly probed. A topic is coded as biographical only when personal information is extracted *as the topic*, not when personal experience is used rhetorically in a policy discussion.

J.6 Sampling-Based Sensitivity Analysis

To check that the findings are not an artifact of one classification strategy, I compare three approaches that vary how much of each transcript the model sees: the full transcript (the main analysis), paragraph-by-paragraph aggregation, and a random sample of paragraphs. Reduced-context strategies necessarily attenuate measured prevalence. The relevant question is whether the qualitative HR-versus-GD contrast survives. It does.

Results Across Sampling Strategies

Figure J.1 shows that HR interviews remain primarily about the candidate while group debates remain primarily about society across all three strategies. Figures J.2 and J.3 show the same robustness for the full topic set and the biographical-marker subset.

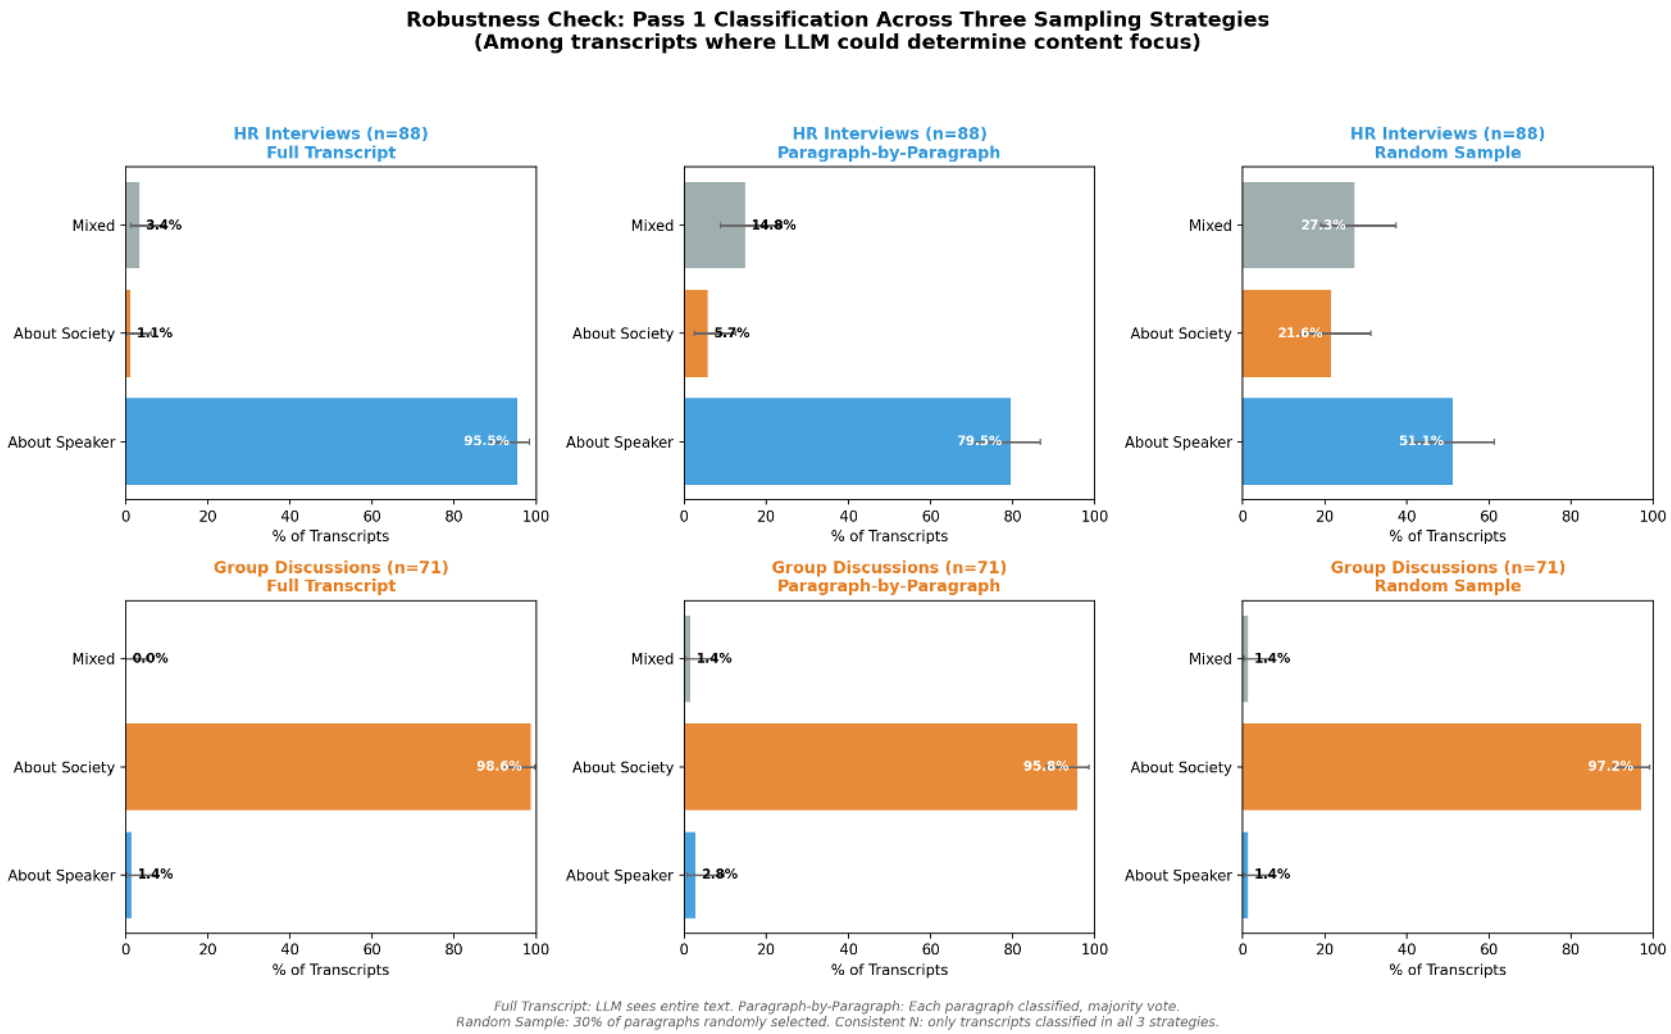

Figure J.1: **Robustness: Content Focus Classification Across Sampling Strategies.**

Notes: Figure J.1 compares transcript-level content-focus classification across the full-transcript, paragraph-by-paragraph, and random-sample strategies. The sample includes only transcripts successfully classified in all three approaches. Lower "about speaker" rates under reduced-context strategies are expected; the key point is that the HR-versus-GD contrast remains.

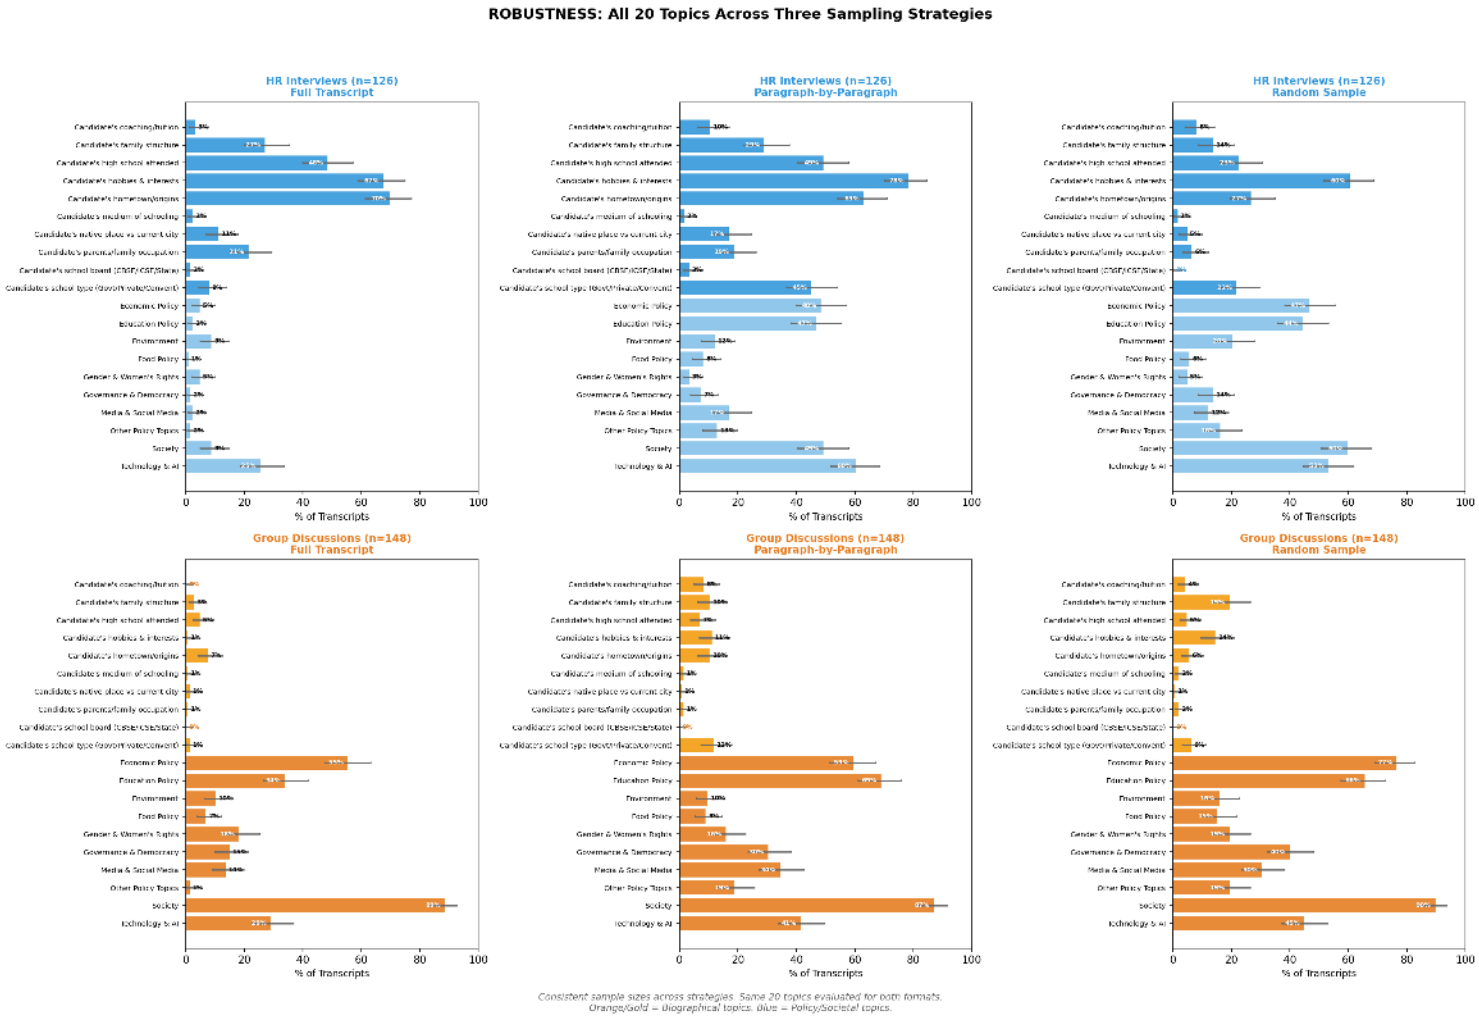

Figure J.2: **Robustness: All 20 Topic Dimensions Across Sampling Strategies.**

Notes: Figure J.2 reports prevalence of all 20 analytic topics across the three sampling strategies. The alternative strategies produce lower measured prevalence because they provide less context, but the relative ordering is preserved: HR remains higher on biographical topics and GD remains higher on policy topics.

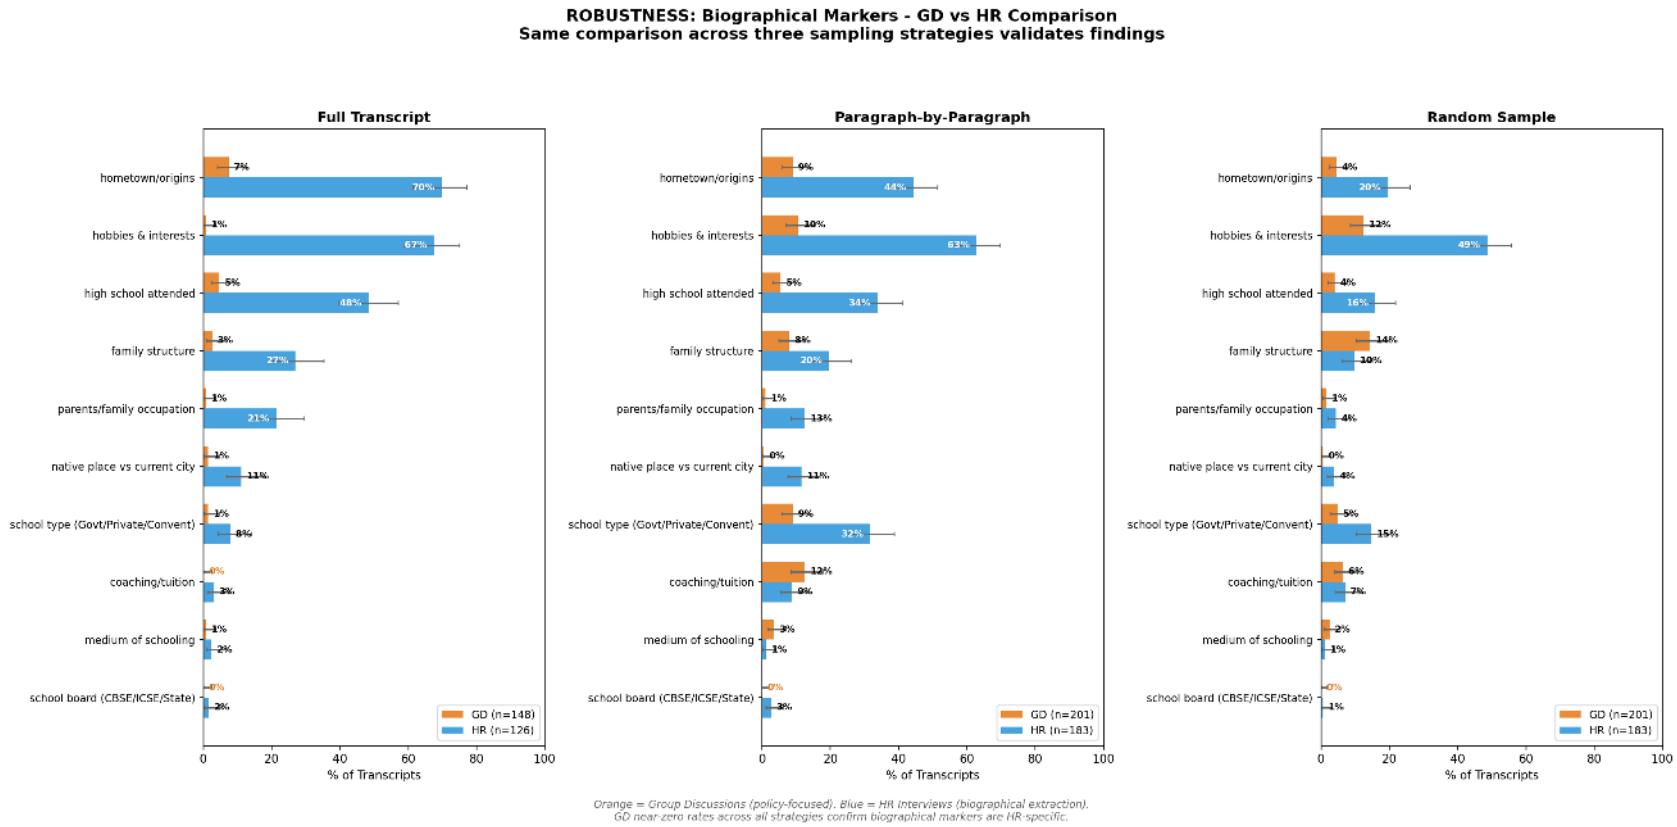

Figure J.3: **Robustness: Biographical Markers Across Sampling Strategies.**

Notes: Figure J.3 reports prevalence of the 10 biographical-marker dimensions across the three sampling strategies. Reduced-context strategies attenuate absolute prevalence, but the qualitative result is unchanged: HR interviews probe biography much more than GDs.

Cross-LLM Robustness: Claude and GPT-5.2 Validation

A separate concern is that the findings might reflect GPT-4o-mini rather than genuine content differences. To address this, I replicated the analysis using Anthropic's Claude Sonnet 4 and OpenAI's `gpt-5.2` (run in February 2026). All three models received identical prompts and the same taxonomy.

Table J.4 shows that the core pattern is robust across models: group debates are classified as "about society," HR interviews as "about speaker," and the prevalence of key biographical topics is similar. Table J.5 reports formal concordance on dominant-content labels for the overlap of transcripts classified as live by each model.

Table J.4: **Cross-LLM Validation: GPT-4o-mini vs. Claude Sonnet 4 vs. GPT-5.2**

	GPT-4o-mini		Claude Sonnet 4		GPT-5.2	
	HR	GD	HR	GD	HR	GD
High-Level Focus						
About Speaker	89.7%	0.7%	93.5%	0.0%	94.9%	0.0%
About Society	3.2%	99.3%	4.7%	100.0%	3.1%	99.1%
Biographical Topics						
Hometown/origins	69.8%	7.4%	75.7%	10.2%	74.5%	7.5%
Hobbies & interests	67.5%	0.7%	72.9%	0.8%	70.4%	0.9%
High school attended	48.4%	4.7%	46.7%	1.6%	48.0%	0.0%
Family structure	27.0%	2.7%	30.8%	0.8%	26.5%	0.9%
Parents' occupation	21.4%	0.7%	23.4%	0.0%	25.5%	0.0%
Sample size (live transcripts)	126	148	107	127	98	107

Notes: Table J.4 reports results under identical prompts and the same unified taxonomy across GPT-4o-mini, Claude Sonnet 4, and `gpt-5.2`. Sample-size differences reflect model-specific live-versus-tips filtering. The fact that the main HR-versus-GD asymmetry persists despite those differences supports robustness of both the live-filtering and topic-classification steps.

Table J.5: **Cross-LLM Dominant-Content Concordance on Overlapping Live Transcripts**

Model Pair	Overlap N	Agreement (%)	Cohen's κ
GPT-4o-mini vs. Claude Sonnet 4	234	96.2	0.925
GPT-4o-mini vs. GPT-5.2	205	98.1	0.962
Claude Sonnet 4 vs. GPT-5.2	199	98.5	0.970
All three models (exact-match agreement)	199	97.0	–

Notes: This table reports concordance for dominant-content classification (about speaker / about society / mixed) on transcript overlaps keyed by filename and source type after the standard exclusions used in the analysis pipeline. "Agreement" is raw percent agreement; κ is Cohen's kappa. Metrics are generated from `scripts/compute_cross_llm_agreement.py`.

K Detailed Methodology for Classifying Caste-Identifiable Surnames

This appendix describes how student surnames are coded as *caste-identifiable*. A surname is treated as identifiable only when it maps unambiguously to an administratively notified community recognized in the reservation system. Official government lists are binding; ethnographic sources are corroborative only.

K.1 Hierarchy of Data Sources

I rely on the following sources, in descending order of legal authority:

1. **Official central/state lists that define eligibility.** These are the binding sources for disadvantaged status.

   - ***OBC (Other Backward Classes):*** I use the National Commission for Backward Classes' state-wise *Central List of OBCs* (NCBC portal). For each State/Union Territory, the portal reports a *Caste/Community* entry together with the relevant *Gazette Resolution* numbers (e.g., Uttar Pradesh list).[14]
   - ***SC/ST (Scheduled Castes and Scheduled Tribes):*** Scheduled Castes (SC) and Scheduled Tribes (ST) are defined by Presidential Orders issued under Articles 341 and 342 of the Constitution and amended by Parliament. I use the Government of India's consolidated *List of Scheduled Castes* with the underlying constitutional orders (Ministry of Social Justice & Empowerment) and the *Constitution (Scheduled Tribes) Order, 1950* (Ministry of Tribal Affairs) together with the updated state-wise ST list (MoTA updated list, 2022).[15]

2. **State gazette notifications and implementing circulars.** Where necessary to harmonize nomenclature (synonyms, regional titles, subgroup labels) with the central lists, I consult the underlying State/Centre *Gazette Resolutions* linked from NCBC entries. These notices are the primary implementation documents.

3. **Ethnographic cross-walks for name variants (corroborative only).** To avoid false positives, I do *not* assign categories from ethnography alone. Instead, I use the Anthropological Survey of India's *People of India* series to verify widely used eponyms, aliases, and regional spellings for communities already recognized in the official lists (AnSI homepage).

[14] *Illustrative mapping (UP):* The NCBC *Central List* for Uttar Pradesh lists *"Ahir, Yadav"* as OBC, with Gazette resolutions `12011/88/98-BCC` (06/12/1999) and `12011/68/93-BCC(C)` (10/09/1993). See the NCBC Uttar Pradesh page. "OBC" is the official term "Other Backward *Classes*," used in central policy and overlapping with "Socially and Educationally Backward Classes" (SEBC) in constitutional usage.

[15] The Legislative Department also hosts the constitutional orders, including the *Constitution (Scheduled Castes) Order, 1950*: link.

K.2 Administrative Framework and the Definition of Disadvantage

What does "Caste/Community" on the NCBC pages denote? On the NCBC *Central List* pages, the *Caste/Community* column lists officially notified communities eligible for OBC benefits in a given State or Union Territory. For coding purposes, a surname is identifiable only if it is a widely recognized eponym for a notified community and that mapping is unambiguous across the official sources above. Surnames shared across multiple communities, states, or both sides of the reserved/unreserved split are coded *non-identifiable*.

How do community lists map to advantaged versus disadvantaged status? The administrative architecture itself defines this binary. Communities appearing in SC, ST, or OBC lists are treated as disadvantaged; those outside these lists are coded as General/Unreserved. The paper therefore follows the state's own reserved-versus-unreserved distinction.

K.3 Operational Mapping and Conservative Criteria

Mapping surnames to notified communities. The coding follows two steps: map a surname to a widely recognized community label, then verify that the community is notified as SC, ST, or OBC in the relevant official lists. A surname is coded as *identifiable* only if this chain—Surname $\rightarrow$ Community $\rightarrow$ Administrative Category—is unique and category-consistent.

The rule is intentionally conservative. A surname must be both pan-regional, mapping to the same side of the reserved/unreserved split across states, and category-consistent, mapping uniquely or only to communities on the same side of that split. Any cross-category or state-dependent mapping is coded *non-identifiable*. The identifiable-surname subsample therefore contains both advantaged and disadvantaged candidates; identification comes from comparing them *within* that subset. Appendix K, Table K.1 provides illustrative mappings.

Accounting for regional variation without home-state information. Because the data do not record students' home states, the coding rule must be valid without state-specific priors. That is why pan-regional clarity is required. Surnames such as "Patel" or "Verma" may map to different communities across states and are therefore coded as non-identifiable. This conservative rule explains why only about 15% of the sample is classified as having caste-identifiable surnames.

Salience over completeness. The method prioritizes salience over exhaustiveness. Official lists define communities, not complete surname dictionaries, so the coding captures only the most recognizable and administratively anchored surname signals. Ethnographic sources are used only to corroborate that a surname is a recognized label for a community already present in the official lists. The fact that only 15% of students meet this strict standard underscores how limited surname-based caste identification is in practice, especially without home-state information.

Table K.1: **Illustrative Surname-to-Category Mapping Under the Conservative Rule**

Surname (Illustrative)	Administrative Link Used for Coding	Coded as Identifiable?	Reason
Yadav	Linked to communities notified as OBC in central/state lists (e.g., NCBC state entries)	Yes	Pan-regional and category-consistent disadvantaged signal.
Chamar/Jatav	Linked to communities notified as SC under constitutional schedules	Yes	Unambiguous notified disadvantaged category mapping.
Patel	Used by multiple communities across states with non-unique category mapping	No	State-dependent and cross-community ambiguity.
Verma	Title used across multiple communities and categories in different regions	No	Cross-category ambiguity prevents a unique administrative mapping.

Notes: Table K.1 is illustrative only. Coding in the data follows the full two-step procedure described above and uses official lists as the binding source. A surname is coded as identifiable only when mapping is pan-regional and category-consistent.

L Differential Selection Across Surname Samples

A natural concern with the surname-identifiability analysis is composition: perhaps disadvantaged candidates with identifiable surnames are simply more negatively selected than their peers. This appendix addresses that concern in two steps. Table L.1 asks whether the caste–GPA relationship differs across surname-identifiability groups. Figures L.1 and L.2 then compare raw and conditional application-stage gaps across the full selectivity distribution. Taken together, the evidence weakens a differential-selection explanation.

Table L.1: **Differential Selection Test: GPA by Caste and Surname Identifiability**

	log(GPA)
Disadv. Caste	-0.115***
	(0.010)
Non-Identifiable	-0.004
	(0.006)
Disadv. Caste × Non-Identifiable	0.001
	(0.011)
Observations	4,005

Notes: Table L.1 reports OLS estimates of log GPA regressed on caste status, surname identifiability, and their interaction. The sample includes all students who applied to at least one job ($n = 4{,}005$). "Disadv. Caste" indicates Scheduled Caste, Scheduled Tribe, or Other Backward Classes status; "Non-Identifiable" indicates students whose surnames do not reveal caste affiliation under the conservative mapping rule. The interaction coefficient is near zero (0.001), indicating that the caste–GPA relationship is essentially the same across surname-identifiability groups. Controls include department, degree, and cohort-year fixed effects; heteroskedasticity-robust standard errors are in parentheses. $^{*}p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

Table L.1 is the mean-level check. The near-zero interaction indicates that disadvantaged candidates in the identifiable sample are not more negatively selected on GPA than disadvantaged candidates in the non-identifiable sample, relative to their advantaged counterparts.

Figures L.1 and L.2 extend the same logic across the application-stage selectivity distribution. Raw gaps are negative in both surname samples, which is unsurprising given average caste differences in observable credentials. After conditioning on the same controls, however, the samples diverge sharply: the application-stage gap is approximately zero for non-identifiable surnames but remains clearly negative for identifiable surnames.

That comparison is informative because the control set is identical in both samples. If the earlier penalty in the identifiable sample were mainly stronger negative selection on observables, the controls should attenuate it in the same way that they attenuate the non-identifiable gap. They do not. This does not mechanically rule out all unobservables, but any remaining selection account must explain three facts jointly: the near-zero caste–GPA interaction, the differential response to the same controls, and the fact that the residual penalty in the identifiable sample appears already at application reading. The more parsimonious interpretation is that recruiters respond to caste information carried by identifiable surnames.

Figure L.1: **Raw Caste Gap in Success Rates by Surname Identifiability (Without Controls)**

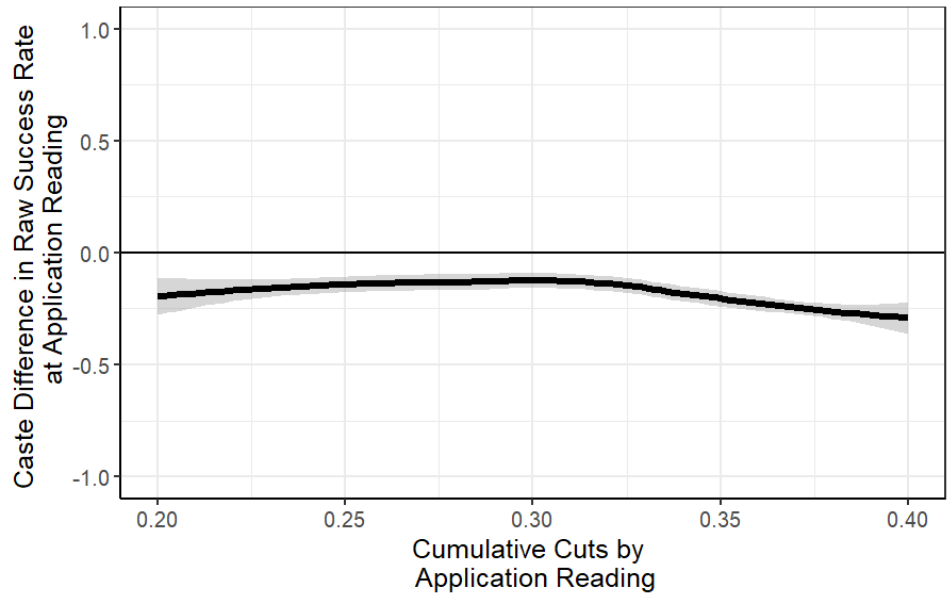

(a) Full Sample: Raw Caste Gap in Success Rates versus Cumulative Cuts

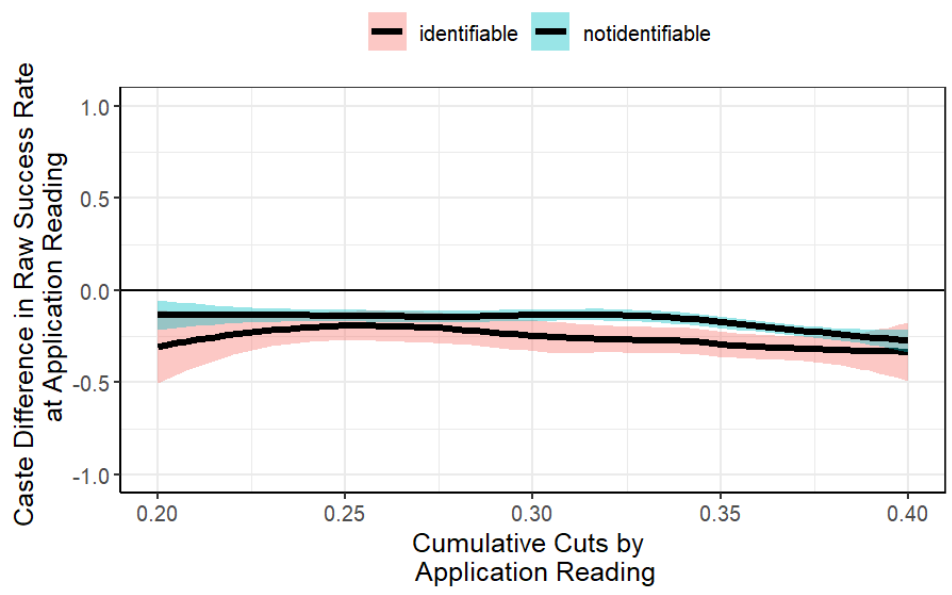

(b) Identifiable vs. Non-Identifiable Samples: Raw Caste Gap in Success Rates versus Cumulative Cuts

Notes: Figure L.1 shows *raw* caste differences in application-stage success rates. Panel (a) shows the full sample; Panel (b) separates students with identifiable (red, $n = 619$) and non-identifiable (cyan, $n = 3{,}386$) surnames. The x-axis shows the cumulative proportion of applicants cut by each job. Both surname samples exhibit sizable negative raw gaps, reflecting baseline caste differences in observable credentials. The comparison with Figure L.2 is the key selection check.

Figure L.2: **Conditional Caste Gap in Success Rates by Surname Identifiability (With Controls)**

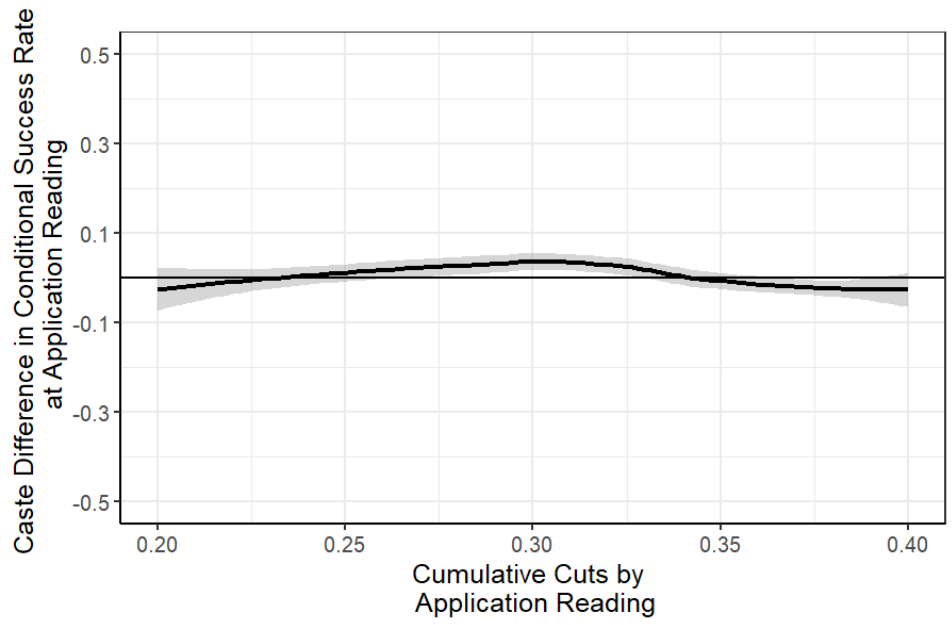

(a) Full Sample: Conditional Caste Gap in Success Rates versus Cumulative Cuts

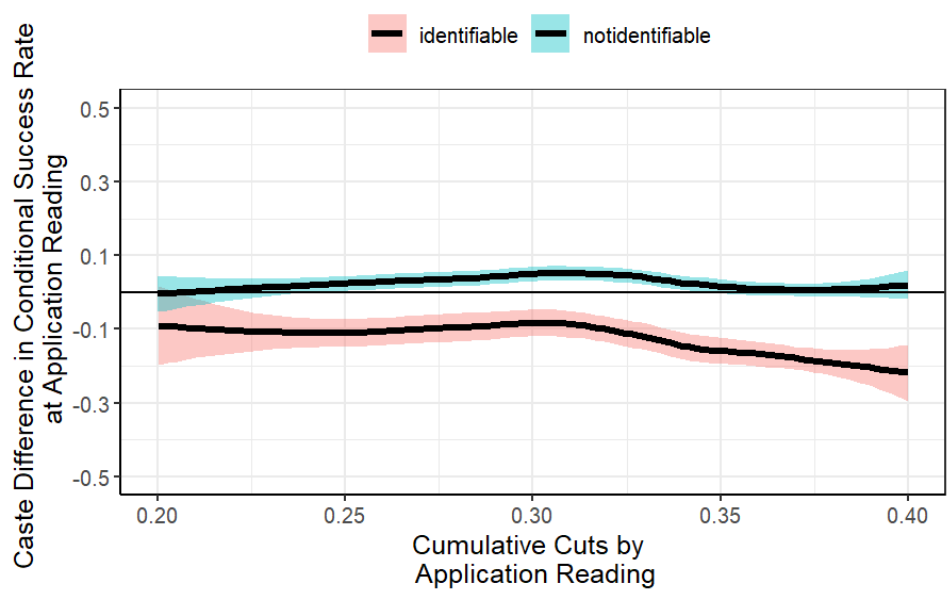

(b) Identifiable vs. Non-Identifiable Samples: Conditional Caste Gap in Success Rates versus Cumulative Cuts

Notes: Figure L.2 shows *conditional* caste differences in application-stage success rates after controlling for GPA, entrance exam score (JEE rank percentile), 10th and 12th grade marks, department, degree, and cohort-year fixed effects. Panel (a) shows the full sample; Panel (b) separates students with identifiable (red, $n = 619$) and non-identifiable (cyan, $n = 3{,}386$) surnames. The x-axis shows the cumulative proportion of applicants cut by each job. Once observables are controlled, the non-identifiable gap rises to approximately zero, while the identifiable-surname gap remains persistently negative. This is the key evidence against a purely compositional explanation for the early identifiable-surname penalty.

M More Details on Alternative Explanations

Several alternative explanations are plausible *a priori*: cumulative ability deficits, accurate statistical discrimination, customer discrimination, strategic delay of discrimination, preferential post-hire treatment, or selection on unobservables correlated with caste. No single diagnostic is dispositive, and the evidence does not mechanically rule out every interview-specific latent trait. But taken together the facts are difficult to reconcile with these alternatives: pre-interview gaps remain near zero even in highly selective jobs; penalties appear immediately when caste is legible from surnames and otherwise emerge mainly when interview content makes caste legible; post-hire promotion and first-job-retention advantages remain broadly stable across selectivity tiers; and patterns look similar in client-facing and non-client-facing roles. Functional-form checks and Oster-style bounds further limit omitted-variable interpretations (Oster, 2019).

1. **Early-round compensatory "fairness" or inclusivity.** If firms were lowering early-stage cutoffs to appear inclusive, we should observe some early-round advantage for disadvantaged candidates, especially where caste is already legible. We do not. In the pooled sample, pre-interview gaps are close to zero rather than positive, and in the surname-identifiable subgroup penalties already appear at application reading (Figure L.2). This does not rule out isolated cases of early compensatory behavior, but it cuts against generalized early-round inclusivity as an explanation for the observed timing pattern.

2. **Strategic delay of discrimination for optics.** A delay story predicts that employers avoid acting on caste information when it is already available and instead wait for a later discretionary stage. That is not what the data show. When caste is legible from surnames, penalties appear immediately at application reading; when surnames are not informative, penalties appear primarily when interview content makes caste legible. Group debates are also subjective and in-person, so if firms wanted simply to postpone discrimination until after objective screens, they could already do so there. The fact that they generally do not is more consistent with information revelation than with strategic timing. The weak formal salience of caste in corporate DEI language (Appendix C) and the qualitative literature on employer resistance to caste-conscious policy (Jodhka, 2008; Jodhka and Newman, 2007) also make a strong optics/compliance story less compelling.

3. **Hidden cumulative or right-tail ability deficits.** If disadvantaged candidates had weaker right-tail ability before interviews, the deficit should become visible in highly selective jobs *before* the interview stage. Under the common-pool and selectivity logic developed in Section I, higher pre-interview cutoffs probe deeper into the same right tail; any broad latent deficit should therefore surface there first. Instead, pooled pre-interview gaps remain near zero even in the most selective jobs (Section 6.2). This sharply weakens cumulative-deficit explanations for the interview penalty, though it does not by itself rule out every interview-specific trait first observed only in round 4.

4. **Accurate statistical discrimination based on noisy signals.** A broad class of accurate statistical-discrimination stories—whether based on believed group means or variances—predicts that post-hire differences should attenuate as job selectivity rises,

because both groups' hires are drawn from increasingly similar margins. Instead, disadvantaged hires show broadly stable promotion and first-job-retention advantages across selectivity tiers (Section 6.3; Appendix H). This does not rule out every conceivable higher-moment model without additional structure, but it makes standard accurate-screening rationalizations substantially less plausible.

5. **Broad socio-emotional skill deficits.** Group debates are not identical to personal interviews, so this test is not exhaustive. But they are the earlier stage most explicitly designed to evaluate communication, teamwork, and performance under direct social interaction. If disadvantaged candidates had a broad deficit in socio-emotional skills large enough to justify the interview penalty, some penalty should usually appear there as well. Instead, caste gaps in group-debate outcomes are near zero, while large penalties appear only in personal interviews. This weakens broad interpersonal-skill explanations for the full interview-stage gap.

6. **Customer discrimination.** If employers were screening to satisfy client preferences, penalties should be materially larger in client-facing roles. They are not. The interview penalty appears in both client-facing and non-client-facing jobs, and the post-hire promotion advantage also appears in both (Sections 6.1 and 6.3). This does not rule out modest customer effects at some firms, but it is difficult to reconcile with customer discrimination as the main explanation.

7. **Selection on unobservables or differential composition across surname samples.** A natural concern is that the identifiable- and non-identifiable-surname samples differ in unobserved productivity in ways that mechanically generate different timing patterns. Appendix L weakens that account in two ways. First, the interaction between disadvantaged caste and non-identifiable surnames in the GPA regression is near zero, indicating that the caste–GPA slope is essentially the same across surname samples. Second, the same observable controls that eliminate the early raw gap in the non-identifiable sample do *not* eliminate the early gap in the identifiable sample. This does not rule out all unobservables, but it makes a simple compositional explanation less persuasive.

8. **Outside options, missing data, or unobserved attrition within the pipeline.** Once students register for the placement process, the placement office enforces participation, prohibits skipping rounds, and records exit decisions through the centralized portal (Sections 4 and 6.1). Exit surveys corroborate these records. These features sharply limit the scope for silent attrition or missing stages to generate the observed hiring pattern. Outside options may matter for post-hire retention at the margin, but they cannot explain why penalties emerge when caste becomes legible in the hiring pipeline.

9. **Preferential promotion for disadvantaged hires.** Explicit caste-focused DEI initiatives in the private sector are rare. Appendix C shows caste is almost absent from formal DEI language, and qualitative work documents resistance to caste-based affirmative action (Bapuji et al., 2023; BusinessLine, 2018; Jodhka, 2008; Jodhka and Newman, 2007). This makes preferential promotion an unlikely explanation for the post-hire advantage. Even if isolated firms favored disadvantaged workers after hiring, that would not explain the earlier timing pattern—especially the immediate application-stage penalty when caste is legible from surnames.

N Additional Figures

Figure N.1: **Common Support Within Each College GPA Decile**

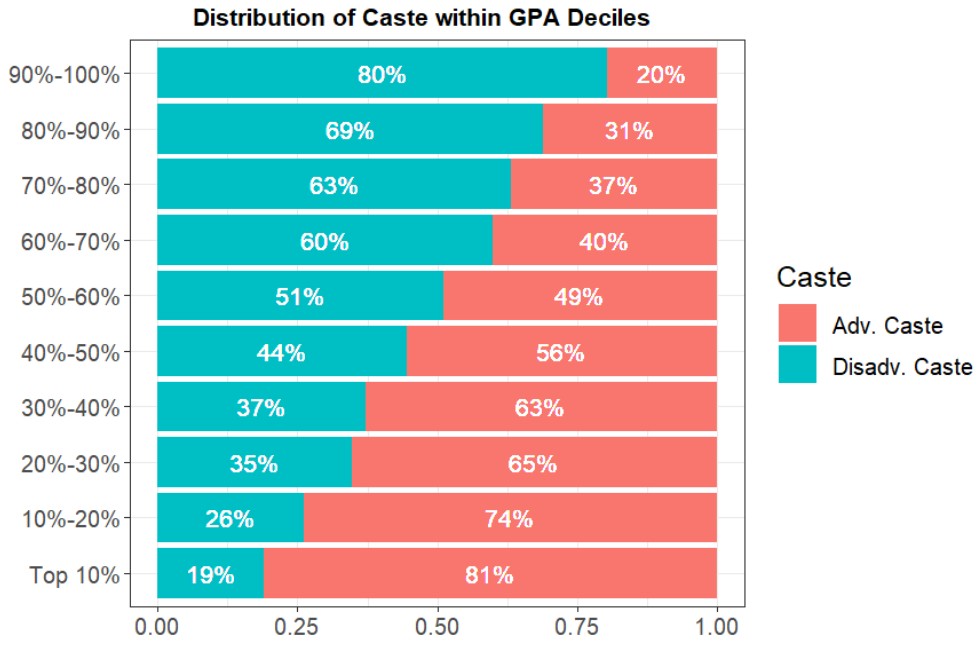

Notes: Figure N.1 demonstrates "common support"—the presence of both advantaged and disadvantaged caste students at every level of academic performance. Each bar shows the count of students within a GPA decile, with advantaged caste students in one color and disadvantaged caste students in another. GPAs are normalized to zero mean and unit standard deviation within each degree-year. The overlap across all deciles ensures that comparisons between caste groups are not confounded by differences in the support of the GPA distribution—i.e., both groups are represented throughout the ability range, including the right tail. This overlap is important for the paper's identification strategy, which relies on comparing outcomes across jobs with varying selectivity thresholds.

Figure N.2: **Caste Penalty in the Uncontrolled Success Rate At Each Screening Round**

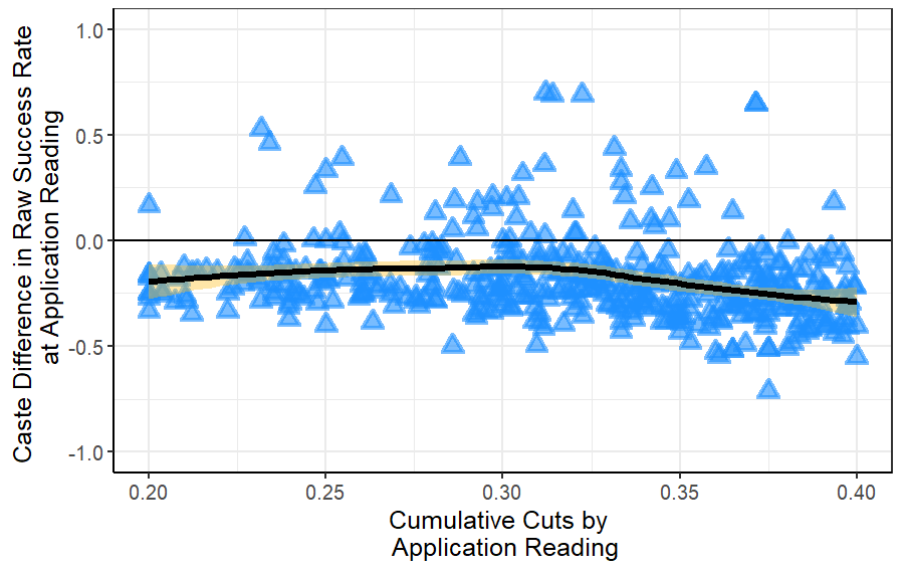

(a) Caste Penalty in Raw Success Rates versus Cumulative Cuts by Application Reading

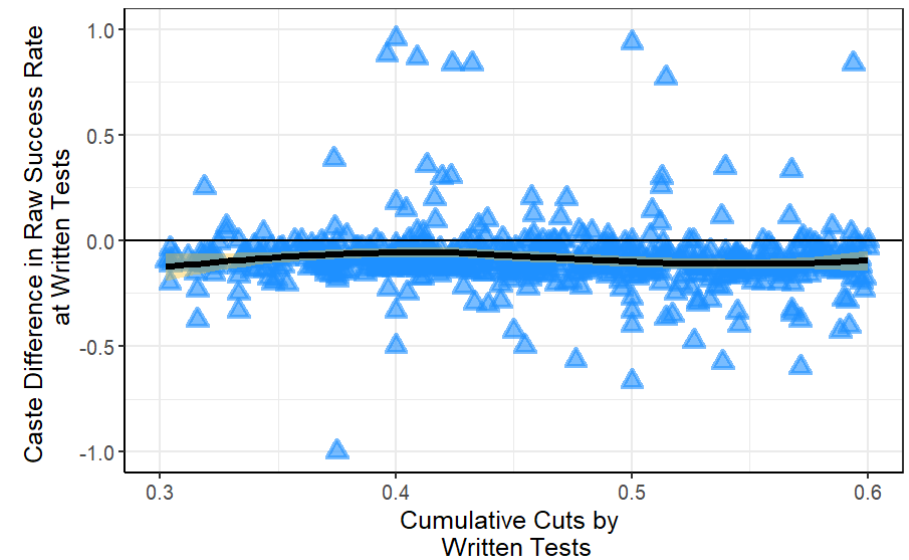

(b) Caste Penalty in Raw Success Rates versus Cumulative Cuts by Written Tests

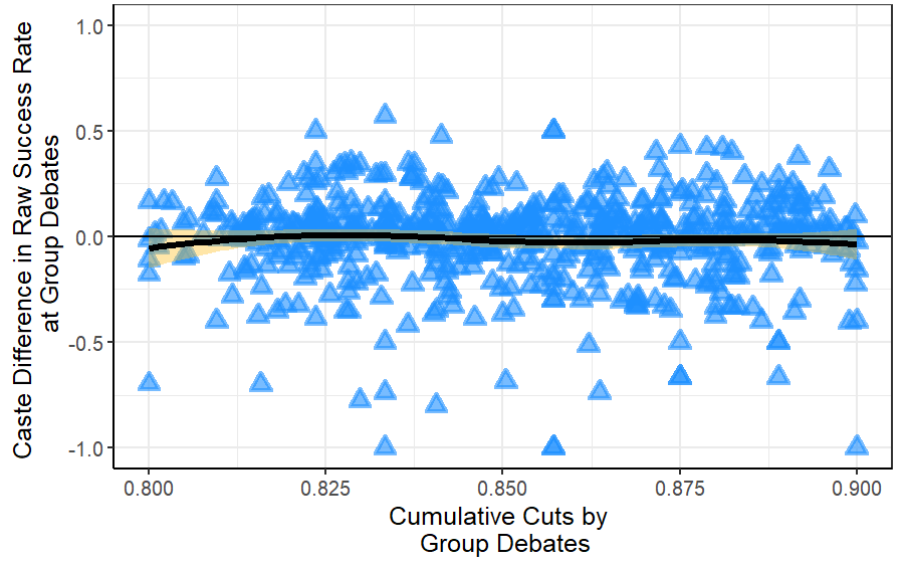

(c) Caste Penalty in Raw Success Rates versus Cumulative Cuts by Group Debates

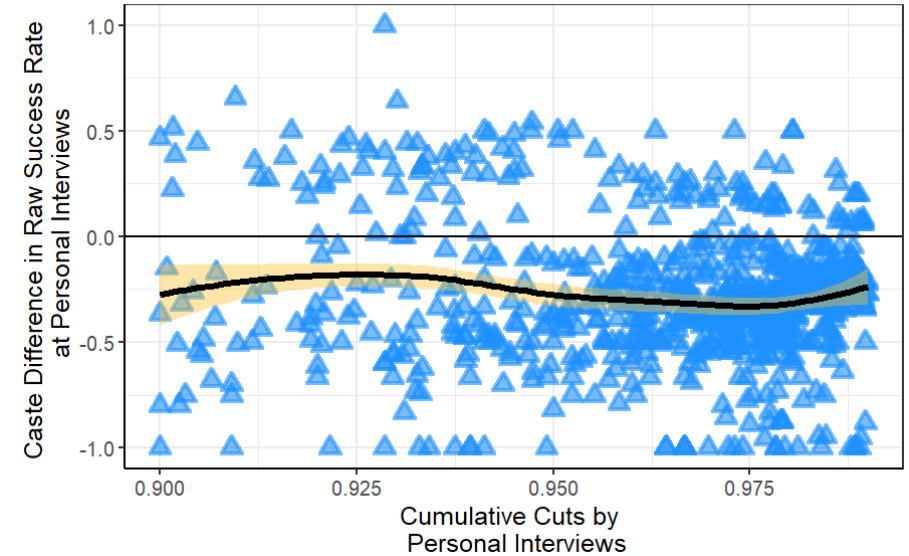

(d) Caste Penalty in Raw Success Rates versus Cumulative Cuts by Personal Interviews

Notes: Figure N.2 plots *raw* (uncontrolled) caste differences in success rates against job selectivity at each of four screening stages: application reading, written tests, group debates, and personal interviews. Each blue triangle represents a single job. Unlike Figure 2, these are simple differences in mean advancement rates between disadvantaged castes (SC/ST/OBC) and advantaged castes—with no individual-level controls and no fixed effects. The y-axis shows the raw caste gap for each job; negative values indicate lower advancement rates for disadvantaged castes. The x-axis shows job selectivity, measured as the cumulative proportion of applicants eliminated by that stage. Black lines show smoothed conditional means (LOESS); shaded regions show 95% confidence intervals. The horizontal line at zero indicates no caste difference. Key finding: a modest raw caste gap exists at the application stage, but this gap *shrinks toward zero* as selectivity increases through written tests and group debates. The gap then *re-emerges sharply* at interviews. This shrinking-then-reappearing pattern—observed even without controls—is inconsistent with cumulative ability deficits and suggests positive selection of disadvantaged candidates reaching interviews.

O Additional Tables

Table O.1: **Total Students by Caste and Degree**

	Total Students		
	Advantaged Caste	Disadvantaged Caste	Degree Total
B.Tech.	579	710	1289
Dual Degree	622	617	1239
M.Tech.	601	566	1167
M.S.	344	125	469
Caste Total	2146	2018	4164
Fraction	0.52	0.48	1.00

Notes: Table O.1 shows student counts by caste and college degree. B.Tech. stands for Bachelor of Technology, M.Tech. for Master of Technology, and M.S. for Master of Science. A Dual degree integrates undergraduate and post-graduate studies and is completed one year after the four-year B.Tech.

Table O.2: **Summary Statistics (Students)**

	Adv. Caste	Disadv. Caste	Diff (SD)
		B.Tech.	
Avg. Entrance Exam Score	0.41	-0.33	0.74***
Avg. 10th Grade Score	0.24	-0.19	0.43***
Avg. 12th Grade Score	0.24	-0.20	0.44***
Avg. Overall College GPA	0.51	-0.42	0.93***
		Dual Degree	
Avg. Entrance Exam Score	0.35	-0.35	0.70***
Avg. 10th Grade Score	0.20	-0.20	0.40***
Avg. 12th Grade Score	0.14	-0.14	0.28***
Avg. Overall College GPA	0.43	-0.43	0.86***
		M.Tech.	
Avg. Entrance Exam Score	0.27	-0.30	0.57***
Avg. 10th Grade Score	0.23	-0.15	0.38***
Avg. 12th Grade Score	0.17	-0.10	0.27***
Avg. Overall College GPA	0.34	-0.33	0.67***
		M.S.	
Avg. Entrance Exam Score	-0.02	0.02	-0.04
Avg. 10th Grade Score	0.03	-0.07	0.10
Avg. 12th Grade Score	0.02	-0.06	0.08
Avg. Overall College GPA	0.04	-0.14	0.18*

Notes: Table O.2 reports pre-college skills (entrance exam scores, 10th and 12th standard national exit exam scores) and college GPA by caste and degree. All scores normalized to zero mean and unit standard deviation within each degree and year. Entrance exam scores (originally ranks) renormalized so higher numbers are better. The "Caste Diff." column shows the difference (disadvantaged minus advantaged) in standard deviation units, with statistical significance indicated by asterisks. $^{*}p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$. Adv. Caste and Disadv. Caste denote advantaged and disadvantaged castes.

Table O.3: **Summary Statistics (Jobs)**

	Count	Fraction	Avg. Salary ($)	Std. Dev. ($)
Private	1011	0.98	60341.38	32591.51
Government	19	0.02	56202.23	20240.01
Headquartered Outside India	620	0.60	64562.74	36038.04
Headquartered in India	410	0.40	53766.04	23946.16
Client Facing	246	0.24	57553.23	28288.23
Non-Client Facing	784	0.76	61115.92	33361.87
All Jobs	1030	1.00	60265.02	32391.89

Notes: Table O.3 reports the distribution of jobs between public and private sectors, domestic and foreign-headquartered (primarily US or Europe) firms, and client-facing versus non-client-facing roles. Average salaries (USD PPP) are also shown for each group. The firms in the sample, regardless of headquarters location, are predominantly multinational corporations.

Table O.4: **Summary Statistics (Candidate Screening)**

Round	Mean	Median
Application Reading	208.85	140.17
Written Tests	124.91	87.00
Group Debates	101.54	70.00
Personal Interviews	26.07	18.00
Offers	6.49	4.00

Notes: Table O.4 reports mean and median candidate counts per job at each stage of the four-round screening process: applications received, candidates advancing to written tests, candidates advancing to group debates, candidates advancing to personal interviews, and final offers extended. Values calculated by averaging student counts across years per job, then averaging across jobs. The declining counts illustrate the substantial selectivity at each screening stage.

Table O.5: **Caste Differences in the Success Rate at Each Job Search Stage by Job Type**

Panel A. Client-Facing Jobs

	App → Test (1)	Test → Debate (2)	Debate → Int (3)	Int → Offer (4)
Disadv. Caste (OLS)	-0.0066**	-0.0007	-0.0136	-0.3062***
	(0.0028)	(0.0048)	(0.0087)	(0.0182)
Disadv. Caste (Oster)	0.002	0.006	-0.015***	-0.314***
	(0.002)	(0.004)	(0.005)	(0.009)
Mean Success Rate	0.5770	0.7860	0.2590	0.2600
Observations	117,770	67,927	53,422	13,842
Adjusted R^2	0.61588	0.37467	0.01189	0.24298
Year FE	✓	✓	✓	✓
Job FE	✓	✓	✓	✓
Other Controls	✓	✓	✓	✓

Panel B. Non-Client-Facing Jobs

	App → Test (1)	Test → Debate (2)	Debate → Int (3)	Int → Offer (4)
Disadv. Caste (OLS)	0.0024	0.0011	0.0070	-0.2867***
	(0.0022)	(0.0024)	(0.0064)	(0.0130)
Disadv. Caste (Oster)	0.008***	0.014***	0.007**	-0.286***
	(0.002)	(0.002)	(0.003)	(0.007)
Mean Success Rate	0.6140	0.8250	0.2580	0.2620
Observations	187,771	115,223	95,067	24,536
Adjusted R^2	0.60783	0.32634	0.02270	0.24723
Year FE	✓	✓	✓	✓
Job FE	✓	✓	✓	✓
Other Controls	✓	✓	✓	✓

Notes: Table O.5 reports within-job linear probability models of advancing to the next hiring stage by job type. Panel A restricts to client-facing jobs; Panel B restricts to non-client-facing jobs. Coefficients are shown for both baseline OLS and Oster-adjusted specifications at each stage (application to test, test to debate, debate to interview, interview to offer). The dependent variable is an indicator for advancing, with observation at the student-job-year level conditional on remaining in contention. Controls include college GPA, entrance exam score (JEE rank percentile), 10th and 12th grade marks, and prior internship sector, with job fixed effects, department fixed effects, degree stream fixed effects, and cohort-year fixed effects. Bias-adjusted coefficients follow Oster (2019), setting $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$. For the baseline OLS rows, standard errors are two-way clustered at the student and job levels. $^*p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

Table O.6: **Sample Selection into Success-Rate and Job-Quality Analyses: GPA Comparisons**

Panel A. Success-Rate Sample Selection (Participants vs. Non-Participants)

	Dependent variable: log (GPA)	
	(1)	(2)
Disadv. Caste	-0.114^{***}	0.009^{***}
	(0.004)	(0.002)
Did Not Participate in Job Fair	0.011	
	(0.016)	
Did Not Participate in Job Fair * Disadv. Caste	-0.044^{**}	
	(0.022)	
Constant	2.071^{***}	2.199^{***}
	(0.010)	(0.028)
Observations	4,164	2,527
Adjusted R^2	0.231	0.822
Year FE	✓	✓
Major	✓	✓
Degree	✓	✓
Job FE		✓
Note:	*p<0.1; **p<0.05; ***p<0.01	

Panel B. Job-Quality Sample Selection (Placed vs. Not in Job-Quality Regressions)

	Dependent variable: log (GPA)	
	(1)	(2)
Disadv. Caste	-0.032^{***}	0.009^{***}
	(0.004)	(0.002)
Not in Job Quality Regression	-0.158^{***}	
	(0.005)	
Not in Job Quality Regression * Disadv. Caste	-0.070^{***}	
	(0.007)	
Constant	2.127^{***}	2.199^{***}
	(0.008)	(0.028)
Observations	4,164	2,527
Adjusted R^2	0.559	0.822
Year FE	✓	✓
Major	✓	✓
Degree	✓	✓
Job FE		✓
Note:	*p<0.1; **p<0.05; ***p<0.01	

Notes: Table O.6 reports GPA-based selection checks for the two core outcomes. Panel A examines selection into the success-rate analysis: Column (1) compares participants in the centralized job fair with non-participants, and Column (2) benchmarks within-job GPA among hires. Panel B applies the same structure for selection into the job-quality analysis by comparing students in versus out of job-quality regressions. Across panels, interaction terms show disadvantaged candidates excluded from the main regressions have lower GPAs, while disadvantaged hires within the same job have slightly higher GPAs. This pattern implies that exclusion is unlikely to mechanically inflate estimated caste penalties. Year, major, and degree fixed effects are included in both columns; job fixed effects are included in Column (2) only. Standard errors are conventional OLS (IID) and are reported in parentheses. *p<0.1; **p<0.05; ***p<0.01.

Table O.7: **Caste Differences in Job Quality at Each Job Search Stage in the Full Sample**

	Dependent variable:					
	log (Avg. Salary)					
	App	Test	Debate	Interview	Offer	Choice
	(1)	(2)	(3)	(4)	(5)	(6)
Disadv. Caste (OLS)	0.004	−0.009	−0.006	−0.018*	−0.105***	−0.103***
	(0.005)	(0.007)	(0.007)	(0.010)	(0.013)	(0.017)
Disadv. Caste (Oster)	0.005	0.010	0.011	0.012	−0.079***	−0.076***
	(0.005)	(0.008)	(0.008)	(0.011)	(0.014)	(0.019)
Constant	11.032***	10.904***	10.848***	10.758***	10.672***	10.863***
	(0.011)	(0.015)	(0.015)	(0.023)	(0.030)	(0.040)
Observations	4,005	3,618	3,266	3,031	2,527	2,527
Adjusted R^2	0.578	0.491	0.547	0.549	0.488	0.412
Year FE	✓	✓	✓	✓	✓	✓
Other Controls	✓	✓	✓	✓	✓	✓
Job FE						

Notes: Table O.7 reports the coefficient on disadvantaged caste (SC/ST/OBC indicator) from OLS regressions of log(average salary) at each job search stage. Each column represents a different stage: (1) application, (2) written test, and subsequent rounds through job choice. The dependent variable is log(average salary), where the average is taken over the posted salaries of all jobs a student remains eligible for at that stage; for the "Choice" column this collapses to the salary of the accepted job. All regressions include linear controls: college GPA, entrance exam score (JEE rank percentile), 10th and 12th grade marks, prior internship sector, department fixed effects, degree stream fixed effects, and cohort-year fixed effects. Job fixed effects are omitted because salaries are standardized and identical within jobs. Bias-adjusted coefficients follow Oster (2019), setting $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$. For OLS rows, App/Test/Debate/Interview/Offer stages use HC1 robust standard errors, while Choice-stage standard errors are clustered at the job level. $^*p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

Table O.8: **Caste Differences in Mean Salary Among Graduates**

	Dependent variable:			
	log (Salary)			
	Without Controls	With Controls	Client Facing	Non-Client Facing
	(1)	(2)	(3)	(4)
Disadv. Caste (OLS)	−0.231***	−0.103***	−0.107***	−0.089***
	(0.020)	(0.017)	(0.027)	(0.022)
Disadv. Caste (Oster)		−0.076***	−0.095***	−0.072***
		(0.019)	(0.026)	(0.022)
Constant	11.235***	10.863***	10.883***	10.848***
	(0.012)	(0.040)	(0.057)	(0.052)
Observations	2,527	2,527	849	1,678
Adjusted R^2	0.049	0.412	0.549	0.382
Year FE	✓	✓	✓	✓
Other Controls	✓	✓	✓	✓
Job FE				

Notes: Table O.8 reports the coefficient on disadvantaged caste (SC/ST/OBC indicator) from OLS regressions of log(salary) of the accepted job, for graduates who obtained jobs through the college's job fair. Column (1) shows raw estimates without controls. Column (2) adds linear controls: college GPA, entrance exam score (JEE rank percentile), 10th and 12th grade marks, prior internship sector, department fixed effects, degree stream fixed effects, and cohort-year fixed effects. Columns (3)-(4) show separate estimates for client-facing and non-client-facing jobs with full controls. Job fixed effects are omitted because salaries are standardized and identical within jobs regardless of candidate characteristics. Bias-adjusted coefficients follow Oster (2019), setting $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$. Standard errors are clustered at the job level. $^*p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

Table O.9: **Variation in Conditional and Cumulative Cuts Across Jobs**

Cut Type	Round	Q1	Mean	Q2	Q3
Conditional Cut	App. Reading	0.30	0.39	0.39	0.48
Cumulative Cut	App. Reading	0.30	0.39	0.39	0.48
Conditional Cut	Written Test	0.10	0.18	0.15	0.26
Cumulative Cut	Written Test	0.38	0.49	0.49	0.61
Conditional Cut	Group Debate	0.69	0.74	0.74	0.79
Cumulative Cut	Group Debate	0.83	0.87	0.87	0.91
Conditional Cut	Personal Interview	0.61	0.75	0.78	0.90
Cumulative Cut	Personal Interview	0.95	0.96	0.97	0.99

Notes: Table O.9 shows the distribution of job selectivity across the four-round screening process. *Conditional cuts* report the fraction of candidates eliminated at each stage, given that they reached that stage. *Cumulative cuts* report the total fraction of original applicants eliminated by that stage. For example, by the personal interview stage, the median job has cumulatively eliminated 97% of original applicants (Q2 = 0.97 for Cumulative Cut at Personal Interview). Q1, Q2 (median), and Q3 denote quartiles across all jobs in the sample. This variation in cumulative selectivity is exploited in Section 6.2 to test whether caste penalties vary with job selectivity.

Table O.10: **Caste Differences in GPA Between Those Who Left the First Job Versus Those Who Did Not**

	Dependent variable:
	log (GPA)
Disadv. Caste	0.007***
	(0.002)
Left Job	−0.003*
	(0.002)
Left Job × Disadv. Caste	−0.004
	(0.003)
Observations	2,489
Adjusted R^2	0.832
Year FE	✓
Major	✓
Degree	✓
Job FE	✓
Note:	*p<0.1; **p<0.05; ***p<0.01

Notes: Table O.10 tests for differential selection on GPA between job leavers and stayers by caste. The dependent variable is log(GPA). The coefficient on Disadv. Caste indicates that disadvantaged hires have higher GPAs than advantaged peers within the same job. The coefficient on Left Job indicates that among advantaged workers, those who leave have slightly lower GPAs than stayers. The interaction term (Left Job × Disadv. Caste) tests whether the GPA gap between leavers and stayers differs by caste—the near-zero and insignificant coefficient indicates no differential selection. The regression includes job fixed effects, year fixed effects, and controls for major and degree. Standard errors are conventional OLS (IID) and are reported in parentheses. $^{*}p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

Table O.11: **Caste Differences in Job Quality at Each Job Search Stage by Surname Identifiability**

Panel A. Identifiable Sample

	Dependent variable: log (Avg. Salary)					
	App (1)	Test (2)	Debate (3)	Interview (4)	Offer (5)	Choice (6)
Disadv. Caste (OLS)	0.012	-0.068^{***}	-0.034^{*}	-0.105^{***}	-0.146^{***}	-0.107^{**}
	(0.014)	(0.019)	(0.018)	(0.027)	(0.033)	(0.049)
Disadv. Caste (Oster)	0.013	-0.049^{***}	-0.016	-0.078^{**}	-0.119^{***}	-0.082
	(0.016)	(0.021)	(0.022)	(0.032)	(0.042)	(0.055)
Constant	10.991^{***}	10.869^{***}	10.807^{***}	10.734^{***}	10.601^{***}	10.834^{***}
	(0.030)	(0.042)	(0.041)	(0.064)	(0.078)	(0.114)
Observations	619	552	500	449	378	378
Adjusted R^2	0.546	0.470	0.537	0.535	0.509	0.362
Year FE	✓	✓	✓	✓	✓	✓
Other Controls	✓	✓	✓	✓	✓	✓
Job FE						

Panel B. Non-Identifiable Sample

	Dependent variable: log (Avg. Salary)					
	App (1)	Test (2)	Debate (3)	Interview (4)	Offer (5)	Choice (6)
Disadv. Caste (OLS)	0.002	0.0004	-0.003	-0.004	-0.100^{***}	-0.106^{***}
	(0.005)	(0.007)	(0.007)	(0.011)	(0.014)	(0.019)
Disadv. Caste (Oster)	0.004	0.019^{***}	0.013	0.024^{**}	-0.075^{***}	-0.076^{***}
	(0.006)	(0.007)	(0.008)	(0.011)	(0.015)	(0.018)
Constant	11.041^{***}	10.913^{***}	10.857^{***}	10.766^{***}	10.685^{***}	10.869^{***}
	(0.012)	(0.017)	(0.016)	(0.024)	(0.033)	(0.043)
Observations	3,386	3,066	2,766	2,582	2,149	2,149
Adjusted R^2	0.584	0.500	0.552	0.556	0.486	0.421
Year FE	✓	✓	✓	✓	✓	✓
Other Controls	✓	✓	✓	✓	✓	✓
Job FE						

Notes: Table O.11 reports stage-specific job-quality regressions by surname-identifiability sample. Panel A uses students with identifiable caste surnames; Panel B uses students with non-identifiable surnames. In both panels, columns track the job-search pipeline from application through job choice, and the dependent variable is log average posted salary across jobs still in contention (or accepted salary in the "Choice" column). All regressions include the main controls and cohort-year fixed effects; job fixed effects are omitted because salary is standardized within jobs. Bias-adjusted coefficients follow Oster (2019), with $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$. For OLS rows, App/Test/Debate/Interview/Offer stages use HC1 robust standard errors, while Choice-stage standard errors are clustered at the job level. $^{*}p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

Table O.12: **Caste Differences in the Success Rate at Each Job Search Stage by Surname Identifiability**

Panel A. Identifiable Sample

	App → Test	Test → Debate	Debate → Int	Int → Offer
	(1)	(2)	(3)	(4)
Disadv. Caste	-0.1839***	0.1099***	-0.0629***	-0.2328***
	(0.0056)	(0.0061)	(0.0087)	(0.0252)
Mean Success Rate	0.4640	0.8830	0.2080	0.4070
Observations	46,163	21,407	18,899	3,932
Adjusted R^2	0.62095	0.29942	0.04530	0.18969

Panel B. Non-Identifiable Sample

	App → Test	Test → Debate	Debate → Int	Int → Offer
	(1)	(2)	(3)	(4)
Disadv. Caste	0.0260***	-0.0112***	0.0078	-0.2979***
	(0.0017)	(0.0024)	(0.0054)	(0.0111)
Mean Success Rate	0.6240	0.8010	0.2660	0.2450
Observations	259,378	161,743	129,590	34,446
Adjusted R^2	0.63019	0.35001	0.01739	0.24959

Notes: Table O.12 reports stage-specific linear probability models of advancing to the next hiring stage by surname-identifiability sample. Panel A uses students with identifiable caste surnames; Panel B uses students with non-identifiable surnames. Columns track the hiring pipeline from application reading through final interviews, and each coefficient reports the disadvantaged-caste gap in conditional progression at that stage. All specifications include the main linear controls, job fixed effects, department fixed effects, degree stream fixed effects, and cohort-year fixed effects. Standard errors are two-way clustered at the student and job levels. *$p < 0.1$; **$p < 0.05$; ***$p < 0.01$.

Supplementary References

Baird, M., N. Gahlawat, R. Hood, S. Lara, and P. Ko (2023). Pathways into STEM jobs. Technical report. LinkedIn Economic Graph, November 2023.

Harrell, F. E. (2015). *Regression Modeling Strategies: With Applications to Linear Models, Logistic and Ordinal Regression, and Survival Analysis* (2nd ed.). Cham, Switzerland: Springer.

Hobbs, A. (2014). *A Chosen Exile: A History of Racial Passing in American Life*. Cambridge, MA: Harvard University Press.

Ignatiev, N. (1995). *How the Irish Became White*. New York: Routledge.

Kashyap, R. and F. Verkroost (2021). Analysing global professional gender gaps using LinkedIn advertising data. *EPJ Data Science 10*(39), 1–32.

Moorjani, P., K. Thangaraj, N. Patterson, M. Lipson, P.-R. Loh, P. Govindaraj, B. Berger, D. Reich, and L. Singh (2013). Genetic evidence for recent population mixture in india. *American Journal of Human Genetics 93*(3), 422–438.

Schwartz, C. R. (2013). Trends and variation in assortative mating: Causes and consequences. *Annual Review of Sociology 39*, 451–470.

Subramanian, A. (2019). *The Caste Of Merit: Engineering Education in India*. Harvard University Press.

Tambe, P., L. Hitt, D. Rock, and E. Brynjolfsson (2021). Digital capital and superstar firms. R&R, *Management Science*.

Zhang, S. and S. Vucetic (2016). Sampling bias in Linkedin: A case study. In *Proceedings of the 25th International Conference Companion on World Wide Web*, pp. 145–146.