Making the Elite: Class Discrimination at Multinationals

Soumitra Shukla*

May 22, 2025

Abstract

What drives socioeconomic disparities in elite job access, and where in the hiring process do these inequalities emerge? Measuring class-based discrimination is challenging, especially in Western contexts where status is often inferred from subtle markers such as accent, education, or networks. India's caste system offers a uniquely explicit measure of social class, enabling precise identification of how social status affects access to elite positions. Analyzing detailed screening data from over 1,000 jobs at U.S. and European multinationals in India, this paper provides the first systematic quantification of class bias in elite recruitment, demonstrating that subjective personal interviews assessing candidate "fit" drive nearly 90% (0.16 standard deviations) of the caste earnings penalty. However, disadvantaged caste hires have around 20% higher promotion rates, even in highly selective roles, indicating these disparities cannot be justified by differences in candidate quality. To formally distinguish discrimination from unobserved ability differences, I develop a novel theoretical framework utilizing plausibly exogenous variation in screening selectivity across jobs, which I empirically test using promotion data. Consistently higher promotion rates among disadvantaged hires across selectivity tiers support the theory's prediction of stricter hiring thresholds rather than productivity differences. Given these persistent promotion advantages across both client-facing and non-client-facing roles, results suggest bias rather than statistical or customer-driven discrimination. These findings highlight a critical organizational challenge: preserving necessary screening discretion without reinforcing class biases that hinder diversity and efficient talent allocation.

*I would like to thank Joseph Altonji, Costas Meghir, and John Eric Humphries for serving on my committee and providing guidance and feedback. I would also like to thank Jason Abaluck, Peter Arcidiacono, Crémieux, Michael Dinerstein, Arpit Gupta, Patrick Kline, Scott Kominers, Eliza McNay, Suanna Oh, Evan Rose, Kamalini Ramadas, Jung Sakong, Rob Valletta, Jon Williams, seminar participants at the Yale labor workshop, the Chicago FED, Chicago Harris, Wharton (BEPP), Penn State, UCSD, SOLE, Southeastern Micro Labor Workshop, System Equitable Growth Conference at the New York Fed, the "Discrimination in the 21st Century Conference" organized by the Becker Friedman Institute and Chicago Booth, MIT Organizational Economics Lunch, and the NBER Race and Stratification Working Group for helpful feedback and support.

"People with advantages are loath to believe that they just happen to be people with advantages." —C. Wright Mills, *The Power Elite* (1956)

1 Introduction

Who gains access to top jobs, and how do hiring practices shape inequality in elite recruitment? Extensive research documents discrimination across labor markets (e.g., Altonji and Blank, 1999; Rivera, 2012, 2015; Friedman and Laurison, 2020), yet precise evidence remains limited on *where*, *how*, and the *magnitude* of hiring discrimination in high-stakes contexts like multinational corporations (MNCs).

Measuring discrimination is especially challenging when the relevant marker is social class. In many Western settings, status is often inferred from overlapping cues—accents, school pedigrees, inherited networks, income—complicating precise measurement (e.g., Lamont and Lareau, 1988; Durlauf, 1999; Ridgeway, 2019; Bisin and Verdier, 1998; Durlauf, 2008; Khan, 2011; Stansbury and Rodriguez, 2024a). Even seemingly objective indicators like income provide ambiguous signals: high-earning individuals with regional or "working-class" accents may still face status penalties, while those with elite cultural markers but modest incomes often retain class advantages through social networks and cultural capital (Bourdieu, 1986). For example, studies in the U.K. show that a high-earning executive who retains a regional working-class accent is often rated as lower status than a lower-paid professional speaking Received Pronunciation or listing a prestigious boarding school (Savage et al., 2013; Savage, 2015; Laurison and Friedman, 2019; Levon et al., 2021).[1] Studies of elite recruitment show that when skill differences among candidates are minimal, evaluators rely on subjective assessments of who "fits" with the organization's culture—referring to candidates' cultural alignment, interpersonal rapport, perceived similarity in values or backgrounds, expectations regarding team dynamics, and other job-specific competencies (Lamont, 1992).

Recent trends illustrate how well-intentioned policies inadvertently reinforce class advantages. Replacing objective tests with fit-focused "holistic" screens shifts opportunities toward higher-status applicants: test-optional policies reduce access for high-achieving, low-income students (Sacerdote et al., 2025), and legacy or athlete preferences at selective U.S. universities favor the wealthy despite weak links to later success (Chetty et al., 2023). India's civil-service hiring shows a parallel pattern: "personality" interviews penalize lower-status

[1] Received Pronunciation, sometimes called "BBC English" or "Queen's English," is the historically prestigious accent of English traditionally associated with the British upper classes and elite institutions (Savage et al., 2013).

candidates yet add little job-performance signal (Bhavnani and Lee, 2021). Across settings, adding subjective "fit" assessments to objective screens typically degrades predictive validity and systematically disadvantages lower-status applicants (Kuncel and Hezlett, 2007; Bol et al., 2014; Rivera, 2012; Laurison and Friedman, 2019; Stansbury and Rodriguez, 2024a). Such subjective judgments encode deep-seated cultural expectations (Bisin and Verdier, 2010) and predict hiring no better than chance (Carroll et al., 1982; Bergman and Lerner, 2025). For example, personal statements and reference letters used in college admissions often reflect reviewers' biases more than candidates' potential (Murphy et al., 2009; Kuncel et al., 2014). Decades of personnel-psychology evidence confirm that these assessments rarely out-predict standardized metrics, even among experts (e.g., Meehl, 1954, 1986; Ægisdóttir et al., 2006; Highhouse, 2008; Dana et al., 2013; Kuncel et al., 2013). Worldwide, fit-based evaluations entrench in-group advantage, disadvantaging lower-status candidates in ratings, promotions, and awards (e.g., Gorman, 2005; Sarsons, 2017; Sarsons et al., 2021; Sarsons, 2023; Benson et al., 2023), with privileged networks propagating these gaps even at elite institutions (Rivera, 2015; Laurison and Friedman, 2016, 2019; Stansbury and Schultz, 2023; Asher et al., 2024; Zimmerman, 2019; Barrios-Fernández et al., 2024). Thus, the very mechanism meant to refine selection often obscures merit.

These universal mechanisms of status discrimination operate across societies but remain difficult to measure precisely. This measurement challenge motivates my focus on the Indian caste system, which provides an unusually comprehensive marker of social stratification. While caste is uniquely Indian, it embodies the broader concept of social class that Bourdieu (1986) conceptualizes as a joint stock of four capital forms: economic (money, property rights), social (durable networks), cultural (embodied traits like accent and etiquette; objectified possessions; institutionalized credentials), and symbolic (prestige from public recognition of other capitals). Unlike the more fluid manifestations of class in Western contexts, India's caste system codifies these capital forms into a relatively stable, legally recognized hierarchy (Beteille, 1965, 1969). This formalization enables precise identification of social-class discrimination that would be methodologically challenging in other contexts, while illuminating mechanisms that operate universally in elite labor markets.

Leveraging this analytical advantage, I examine recruitment data from leading U.S. and European multinationals hiring in India. Despite affirmative action policies, disadvantaged caste individuals remain significantly underrepresented in elite jobs (Madheswaran, 2008), though the precise magnitude and locus of discrimination remain unclear. This analysis makes three key contributions: First, I systematically quantify social-status discrimination in elite recruitment. Second, I identify personal interviews that assess candidate "fit" as the primary mechanism revealing and penalizing social status—detailed data show most caste

disparities emerge here, with hired disadvantaged caste candidates subsequently exhibiting significantly higher promotion rates. Third, I demonstrate that class bias—rather than statistical or customer discrimination—drives entry-level disparities. By observing the full recruitment process through job offers and career outcomes, this analysis provides uniquely detailed evidence on discrimination's economic consequences compared to existing research focused on callback rates (e.g., Kline et al., 2022). Each of these three contributions is, to my knowledge, novel in the literature.

Elite entry-level positions significantly influence economic power, particularly in India, where disadvantaged castes constitute nearly 70% of the population (Bayly, 2008). The studied elite college job fair exemplifies globally widespread multinational hiring practices (cf. Bloom et al., 2012; Minni, 2023), with nearly half the students belonging to disadvantaged castes due to mandated affirmative action. Across the 1,050 jobs studied, each position involves four screening rounds progressing from standardized to subjective assessments: application reading, aptitude tests, group debates on socio-emotional skills, and final-round personal interviews assessing candidate "fit." The unconditional caste earnings gap is 23%, narrowing to 10% after controlling for pre-college skills, academic performance, and other employer-relevant factors. With salaries explicitly *identical across castes* for identical roles and *non-negotiable*, I decompose these gaps by counterfactually assigning identical salaries at each stage before offers are accepted. Personal interviews alone account for nearly 90% of the conditional earnings penalty (approximately 9% of average earnings, or 0.16 standard deviations) and about 95% of the nearly 30-percentage-point caste penalty in hiring success rates, highlighting their role in penalizing social status.

The caste penalty arises when caste is revealed. While some surnames indicate caste, most do not convey caste signals in urban-educated settings (e.g., elite colleges typically located in urban centers) due to migration and geographical variation in naming conventions (Mamidi, 2011; Jodkha, 2017; Wiltshire, 2020). Among the 15% of students with distinctly caste-identifiable surnames, penalties emerge immediately at the first screening round (application reading). For the remaining 85% with caste-ambiguous surnames, no early penalty appears—disparities emerge only after personal interviews. This differential timing strongly suggests that discrimination occurs precisely when caste becomes visible to employers, with personal interviews serving as the primary revelation mechanism for most students.

This paper does not argue that personal interviews are useless. Rather, it argues that any *additional diagnostic information* interviews provide cannot *justify* observed caste disparities. Disadvantaged caste hires consistently outperform advantaged peers by around 20% in promotion rates, even in highly selective roles, strongly indicating disparities stem from discrimination rather than unmeasured skill deficits.

To formalize this interpretation, I develop a theoretical framework that leverages plausibly exogenous variation in screening selectivity—the depth in the applicant pool from which each firm selects—across jobs. The framework yields two key implications: first, if disadvantaged candidates genuinely lacked skills, highly selective pre-interview screening stages would show a caste penalty, which they do not; second, if interviews accurately filtered weaker candidates, disadvantaged hires would not consistently outperform advantaged peers across selectivity tiers. The logic is straightforward. As selectivity increases, the average hired candidate's quality approaches marginal (cutoff) quality. If jobs were using the same cutoffs regardless of caste, promotion differences between castes should shrink as job selectivity increases. Thus, consistently higher promotion rates among disadvantaged hires across all selectivity levels suggest stricter hiring thresholds (discrimination) rather than productivity differences.

Analysis of distinctive versus non-distinctive surnames, combined with the selectivity evidence above, strongly indicates that hiring disparities reflect *disparate treatment*—explicit caste-based discrimination. This conclusion is reinforced by two additional findings: persistent promotion advantages for disadvantaged hires in both client-facing and non-client-facing roles (which comprise over 75% of positions). Together, these patterns suggest biases—either taste-based discrimination or biased beliefs—rather than statistical or customer-driven discrimination drive the caste hiring penalty (cf. Benson et al., 2024). Such discrimination explicitly violates Article 15 of the Indian Constitution (Ministry of Law and Justice, Government of India, 2021), severely limiting socioeconomic mobility and impeding efficient talent allocation.

This paper advances labor, personnel, and organizational economics by quantitatively pinning down social-status discrimination in elite hiring and unpacking its mechanisms. Tracking candidates through each stage of multinational recruitment, it shows that subjective "fit" interviews account for about 90 percent of the conditional hiring gap. The estimate enriches the still-nascent literature on class discrimination, which lags behind research on race and gender (e.g., Durlauf, 2006; Rivera, 2012, 2015). The paper is most closely related to Stansbury and Rodriguez (2024a), who document sizable class gaps in U.S. academic careers; Rivera (2012), who posits that "cultural matching" in elite-firm interviews could engender class disparities, yet stops short of quantifying their magnitude in determining hiring outcomes; Bhavnani and Lee (2021), who, in the context of recruitment in elite civil services, show that disadvantaged caste officers who would have qualified even without affirmative action quotas score markedly lower in "personality'' interviews yet later outperform advantaged officers by roughly 18–20 percent, effectively replicating the late-stage penalty documented here; and Zimmerman (2019), Michelman et al. (2022), and Barrios-Fernández

et al. (2024), who demonstrate how elite schooling reproduces advantage via peer-driven social and cultural capital, explicitly invoking Bourdieu's theory of social reproduction. While each study uncovers a critical piece of the puzzle, no paper jointly (i) quantifies the magnitude of class discrimination in elite hiring, (ii) identifies its location in the hiring pipeline, and (iii) isolates its underlying mechanisms, distinguishing bias from statistical or customer discrimination. Leveraging India's explicit caste identifiers and unusually granular recruitment data, the present paper is the first to accomplish all three, providing systematic evidence on the location, mechanisms, and magnitude of class bias—information that is indispensable for effective policy and organizational interventions.

As traditional social-status markers dilute across societies (Lenski, 1954; Grusky, 2005; Clark, 2015), class discrimination likely moves deeper into hiring processes, underscoring the need for detailed hiring-stage data to uncover how social status influences hiring decisions. This evolution mirrors broader patterns of how discrimination adapts to regulatory scrutiny and formal equality: as overt bias becomes socially and legally costly, it often shifts toward more subtle, difficult-to-detect forms (Pearson et al., 2009). My findings about interviews' role in perpetuating class disadvantage connect to research on how organizational practices can inadvertently preserve discrimination even while attempting to promote meritocracy (Castilla and Benard, 2010; Castilla, 2015). The results also suggest broader implications: subjective interviews may likewise disadvantage candidates with other less visible yet stigmatized identities—such as lower socioeconomic background (Rivera, 2012, 2015; Rivera and Tilcsik, 2016), LGBTQ+ status (Eames, 2024, 2025), or intersectional identities (Chow and Knowles, 2016). These patterns highlight how seemingly neutral evaluation practices can amplify subtle forms of disadvantage (Reskin, 2012), especially when criteria like "cultural fit" allow evaluators to act on implicit biases without explicit awareness (Rivera, 2012; Kang et al., 2016).

Another key contribution highlights managers' critical role in talent allocation through screening on social class. Recent literature documents substantial managerial influence on firm performance and culture (e.g., Turco, 2010; Lazear et al., 2015; Sadun et al., 2017; Hoffman and Tadelis, 2021; Espinosa and Stanton, 2022; Fenizia, 2022; Sadun, 2023; Benson and Shaw, 2024; Minni et al., 2025), particularly through assigning workers to teams and training investments (e.g., Minni, 2023; Sadun, 2023; Sadun et al., 2024). This paper uniquely integrates Bourdieu's concept of social class—especially its social and cultural capital dimensions like hobbies, networks, mannerisms, and conversational polish—to illuminate biases in managerial assessments of who "fits" (Bourdieu, 1986). Although such class markers might correlate with productivity, my empirical results show that subjective interview information—whether reflecting cultural similarity or other aspects of "fit"—cannot justify

observed disparities, as disadvantaged hires perform better. Thus, this study bridges cultural sociology and empirical personnel economics by demonstrating how managerial biases operate through cultural and social alignment.

Additionally, this paper highlights agency problems in hiring: the separation of hiring authority from performance accountability and managers' reliance on overconfident private judgments or productivity-irrelevant cues. Hiring managers in my setting rarely supervise recruits directly, creating incentives to prioritize perceived "fit" over long-term potential. Higher promotion rates for disadvantaged candidates, despite initial undervaluation, underscore this misalignment. Indeed, candidates rejected on subjective grounds frequently perform as well or better than initially accepted peers (e.g., DeVaul et al., 1987), and managers who override data-driven hiring-test recommendations often end up with lower-performing hires (Hoffman et al., 2018). Such subjective, short-term decisions systematically produce sub-optimal organizational outcomes, especially in the long-run (e.g., Gibbons and Roberts, 2013; Sadun et al., 2017; Benson and Shaw, 2024; Hoffman and Stanton, 2024).

This paper further contributes by examining discretionary evaluations—particularly personal interviews—in elite private-sector hiring, a context with limited empirical evidence despite significant implications. While economic studies extensively document discrimination at early stages like résumé screens (cf. Kline et al., 2022), few trace bias into the interview room where subjective evaluations occur. Meta-analyses of correspondence audits find sizable race and gender gaps at callbacks—and gaps of similar or larger size by the offer stage—suggesting additional discrimination emerges during face-to-face interactions (Quillian et al., 2020). Organizational-psychology research shows that subjective interviews invite snap judgments, inflate confidence, and promote stereotype-consistent "sense-making" (Highhouse, 2008; Dana et al., 2013), while meta-studies document that interview impressions add little predictive power beyond objective assessments (McDaniel et al., 1994; Levashina et al., 2014). Recent evidence also shows informal, subjective criteria consistently advantage privileged candidates yet correlate poorly with job performance (e.g., Sarsons, 2017; Benson et al., 2023; Cappelli, 2019; Hoffman et al., 2018). By pinpointing when class penalties crystallize and linking them to post-hire performance, this study bridges psychological insights with new field evidence, while highlighting limitations of interventions like résumé anonymization that may merely shift bias to later, subjective stages (cf. Doleac and Hansen, 2020).

This study also contributes to debates over standardized versus "holistic" evaluations common in elite hiring—such as college admissions, faculty hiring, and grant awards. These processes typically progress from objective metrics (GPA, test scores, publications) that reliably predict performance (Bol et al., 2014; Kuncel and Hezlett, 2007), toward subjective assessments (personal statements, interviews) with lower predictive validity (Murphy

et al., 2009; Borneman et al., 2007; Cooper et al., 2007; Kuncel et al., 2007; Alvero et al., 2021). In my context, screening moves from application reading and aptitude tests toward group debates and personal interviews, mirroring this general shift toward subjective evaluation. Recent research shows holistic evaluations disproportionately reflect evaluators' biases rather than candidates' potential (Chetty et al., 2023; Sacerdote et al., 2025; Bol et al., 2014; Checchi and van de Werfhorst, 2018; Barg, 2013). While emerging algorithmic screening tools aim to increase objectivity, they may inadvertently replicate biases without rigorous validation (Albright, 2019, 2024), though properly designed algorithms can enhance efficiency and equity (Bergman et al., 2024). By linking subjective evaluations to organizational inefficiencies—demonstrated by higher promotion rates among initially undervalued candidates—this research emphasizes the need to scrutinize and reform interview practices.

Finally, this paper provides insights into elite educational institutions' role in socioeconomic mobility and discriminatory hiring. Elite colleges like the studied Indian institution, with admission rates more competitive than many Ivy leagues, are gateways to prestigious careers (Jodhka and Naudet, 2019). Although affirmative action policies worldwide substantially increase disadvantaged students' access to elite colleges, research indicates elite education disproportionately benefits privileged students, perpetuating class inequalities (e.g., Zimmerman, 2019; Barrios-Fernández et al., 2024; Chetty et al., 2023).[2] This study shows how subjective interview practices can erode affirmative action gains, systematically disadvantaging these candidates (cf. Bhavnani and Lee, 2021). Given multinational firms' recruitment practices, findings highlight the need to integrate affirmative action with hiring reforms. Evaluating elite colleges separately from employer hiring overlooks critical complementarities in selection criteria and networks, emphasizing integrated approaches for socioeconomic mobility—particularly given how social class co-evolves with cultural and institutional factors shaping economic outcomes (e.g., Bisin et al., 2023).

2 Caste as Social Status Under Affirmative Action

India's affirmative-action framework formally treats caste as a unified marker of social status, bundling financial resources, social networks, and cultural capital into two administratively defined categories—"advantaged" and "disadvantaged." Because these policies apply only to public universities and government jobs, not to private multinationals, the labor market I study offers an unusually clean lens on where status biases emerge in hiring. For historical background on caste classification and reservation policy, see Online Appendix Section A.

[2]Complementary evidence from Ang (2021) and Ang (2023) shows that discriminatory environments beyond hiring—such as biased policing and media—likewise hinder upward mobility.

3 Key Definitions and Institutional Setting

This section defines key terms, describes the elite college's job fair studied in this paper, and clarifies why this job fair is representative of how elite Indian college graduates enter elite jobs.

3.1 Key Terms: Disadvantaged Castes, Elite Jobs, and Elite Colleges

Below, I define four key terms used throughout the analysis:

a. **Disadvantaged castes.** Following the Indian government's official classifications, the term "disadvantaged castes" encompasses individuals belonging to Scheduled Castes (SCs), Scheduled Tribes (STs), or Other Backward Classes (OBCs).

b. **Job.** A "job" refers specifically to a job designation within a firm. For example, if Google recruits both a product manager and a software engineer, these constitute two distinct jobs. Section 4.6 explains why a job typically corresponds to a firm-designation pair, rather than a particular office location.

c. **Elite jobs.** An "elite entry-level job" refers to a formal entry-level position at multinational corporations (MNCs), major national corporations, or unicorn startups (privately-held startups valued at over $1 billion). All elite jobs analyzed in this paper pay salaries within the top 1% of entry-level wages in India (Economic Advisory Council to the Prime Minister, 2022). Throughout this paper, the terms "elite entry-level jobs" and "elite jobs" are used interchangeably.

d. **Elite colleges.** "Elite colleges" include institutions consistently ranked among the top 30 in their respective fields—such as science, engineering, arts and humanities, social sciences, and business—by *India Today*, an Indian equivalent of the *US News & World Report*. Prominent examples include the Indian Institutes of Technology (IITs), Indian Institutes of Management (IIMs), National Institutes of Technology (NITs), and select Delhi University colleges such as St. Stephen's, Shri Ram College of Commerce, and Lady Shri Ram College. Several elite Indian colleges also place among the top colleges worldwide for graduate earnings, according to a cross-country analysis of college quality in Martellini et al. (2024), underscoring their international standing.

3.2 Structure of the On-Campus Job Fair

The elite college's job fair occurs *entirely on campus* and comprises three phases: (1) pre-placement, (2) placement, and (3) post-placement (or the "aftermarket").

a. **Pre-Placement Phase.** Firms are invited by the placement office prior to June. Between June and mid-August, recruiters deliver pre-placement talks, gauge student interest, and finalize roles. By late August, firms extend pre-placement offers (PPOs) to students who interned with them during the summer. Students accepting PPOs are withdrawn from the main placement process by the college's placement office.

   Significance of return offers from internships. Unlike in the US, internship return offers play a minor role. Fewer than 5% of students accept these offers and thus exit the placement process, leaving over 95% to actively participate in the main job fair (Section 5.1). Additionally, fewer than 2% of students choose public-sector employment (Section 4.6).

   Student and employer registration. Students register for placement by late August, while employers submit detailed registration forms by early September, specifying job roles, compensation, and vacancies, which become fixed thereafter. Compensation is *non-negotiable*, strictly enforced by the placement office through verification of offer letters. Non-compliance can result in restrictions on future campus visits by firms.[3]

b. **Placement Phase.** Starting in mid-September, students apply for jobs. Employers initially screen applications, followed by written aptitude tests and group debates assessing socio-emotional skills from September to early December. Final-round personal interviews occur between December and January, after which about two-thirds of the total student cohort secure full-time jobs through this standardized placement process.

c. **Post-Placement Phase ("Aftermarket").** From late January until graduation in July, the placement office operates an "aftermarket" for students without initial job offers, comprising roughly one-third of the cohort. These unplaced students typically pursue entrepreneurship, graduate studies, competitive exam preparation (such as civil services or elite MBAs), or find alternative employment.

[3]Such rules and timelines are common across elite Indian colleges and are available on their websites. For instance, see Central Placement Cell (Delhi University).

Excluding firms and students outside the standard recruitment pipeline. Employers in the aftermarket—typically smaller or newer startups that offer lower salaries and use less standardized recruitment—are excluded, as are any firms (and the students they eventually hire) whose hiring bypasses the four-stage pipeline (e.g., test-plus-interview or single-round interviews). Because the analysis explicitly targets standardized hiring practices, these exclusions are necessary; their implications are discussed in Section 5.4.

3.3 Why the Studied Job Fair is Representative

The studied job fair is representative of how elite Indian graduates transition into elite jobs due to four factors (further details and evidence are provided in Online Appendix Section B). First, nearly all top-ranked Indian colleges (defined in Section 3.1) use similar standardized procedures, making this fair typical of peer institutions. Second, elite employers—multinationals, major national firms, and unicorn startups—recruit consistently from the same top-tier colleges. These firms visit 30–40 colleges annually, collectively impacting around 50,000 students. Because these employers draw from a common institutional network, their recruitment strategies and job offerings remain consistent.

Third, elite entry-level roles almost exclusively recruit from top-ranked colleges, whose graduates predominantly fill these positions. This mutual exclusivity is documented in prior research (e.g., Subramanian, 2019; Jodhka and Naudet, 2019), creating an integrated elite labor market. Thus, the studied college's employment outcomes directly mirror national trends. Fourth, very few graduates opt out of the formal placement process. In this setting, fewer than 5% of students choose to forgo participation (Section 5.1). Comparable participation rates at peer institutions confirm the representativeness of the studied fair (e.g., Subramanian, 2019). Most students rely heavily on their college's placement office rather than pursuing off-campus opportunities, further reinforcing representativeness.

4 Detailed Data on Students, Jobs, and Recruitment Process

This section describes the data used in the analysis, detailing student and job characteristics, recruitment procedures, and career outcomes such as promotions and tenure. It concludes by discussing relevant data limitations.

4.1 Institutional Source and Completeness of Recruitment Data

Data are obtained directly from the placement office of an elite Indian university (anonymized for confidentiality), which meticulously records each recruitment stage—applications, screening rounds, interviews, and final job offers. The university's standardized recruitment practices benchmark placement processes at comparable institutions nationwide (cf. Subramanian, 2019).

4.2 Centralized Job Portal

Students register through a centralized online job portal managed by the placement office, creating standardized personal profiles including GPA, entrance exam scores, major coursework, and internship experiences. Firms similarly submit detailed profiles, specifying non-negotiable salary structures, job descriptions, and qualifications (Section 3.2). Both student and job profiles are mutually visible via the portal. Throughout recruitment, employers directly upload candidate shortlists from each screening round onto the portal, ensuring systematic documentation and data consistency.

4.3 Standardized Recruitment Rules and Mandatory Participation Ensuring Data Integrity

Recruitment at elite Indian colleges is tightly standardized: students who enter the campus placement fair are *mandated* by the placement office to participate in all four screening rounds—application screening, written aptitude tests, group debates, and personal interviews (Section 4.8)—prohibiting informal or incomplete hiring practices. Offer-letter salaries are likewise verified against the non-negotiable compensation schedules filed by firms, preventing unrecorded side deals and keeping wage data clean (Section 3.2). Such standardized rules, common across elite institutions, enhance transparency, comparability, and empirical rigor.[4]

4.4 Student Sample: Time Span, Degrees, Demographics, and Caste Representation

The dataset spans four placement cycles at an elite Indian college (exact years anonymized). It includes 4,164 students pursuing four degrees: Bachelor of Technology (B.Tech.), Dual Degree (integrating undergraduate and postgraduate coursework), Master of Technology (M.Tech.), and Master of Science (M.S.).

[4]Detailed placement guidelines are publicly accessible on elite colleges' websites, e.g., the Central Placement Cell (Delhi University).

Admissions based solely on standardized national entrance exams. Admissions are *exclusively* based on standardized national-level exams with caste- and major-specific cutoffs.

Near-equal caste representation via affirmative action in admissions. Due to caste-based affirmative action (Section 2), disadvantaged and advantaged castes are represented at near-parity *within each major*, and therefore within most degrees (Table 1).[5]

Caste self-reporting. Students self-report caste through official administrative documents, enabling accurate tracking by both the college and researcher.

Predominantly male sample. Approximately 90% of students are male, typical of elite technical colleges in India (cf. Datta, 2017).

Table 1: **Total Students by Caste and Degree**

	Total Students		
	Advantaged Caste	Disadvantaged Caste	Degree Total
B.Tech.	579	710	1289
Dual Degree	622	617	1239
M.Tech.	601	566	1167
M.S.	344	125	469
Caste Total	2146	2018	4164
Fraction	0.52	0.48	1.00

Notes: Table 1 shows student counts by caste and college degree. B.Tech. stands for Bachelor of Technology, M.Tech. for Master of Technology, and M.S. for Master of Science. A Dual degree integrates undergraduate and post-graduate studies, completed one year after the four-year B.Tech.

4.5 Student Characteristics and Caste-Based Differences

Student characteristics include college GPA, pre-college academic skills, and prior labor market experience, highlighting two key patterns:

1. **Significant caste differences in academic performance.** Large caste disparities exist in college GPA and entrance exam (EE) scores (Table 2), although students from both caste groups span all deciles of these metrics (Online Appendix Figure OA.1).

[5]In some degrees, particularly the M.S., disadvantaged-caste seats often remain partially unfilled due to eligibility constraints or fewer qualified applicants, consistent with patterns documented in Krishna and Frisancho Robles (2015).

2. **Weak correlation between caste and prior internships.** The dataset includes detailed internship information (e.g., duration, pay, employment sector). Caste is only weakly correlated with internship sector or experience. For instance, among B.Tech. and Dual Degree students, internships across castes show similar sector distributions (25% technology, 37% consulting, and 38% manufacturing). Similar patterns hold for postgraduate degrees (not shown). This likely arises because return internship offers are rare in elite Indian colleges (Section 3.2), making internship choices primarily driven by student sector preferences rather than caste-related pay or opportunities.

Table 2: **Summary Statistics (Students)**

	Adv. Caste	Disadv. Caste	Diff (SD)
		B.Tech.	
Avg. Entrance Exam Score	0.41	-0.33	0.74***
Avg. 10th Grade Score	0.24	-0.19	0.43***
Avg. 12th Grade Score	0.24	-0.20	0.44***
Avg. Overall College GPA	0.51	-0.42	0.93***
		Dual Degree	
Avg. Entrance Exam Score	0.35	-0.35	0.70***
Avg. 10th Grade Score	0.20	-0.20	0.40***
Avg. 12th Grade Score	0.14	-0.14	0.28***
Avg. Overall College GPA	0.43	-0.43	0.86***
		M.Tech.	
Avg. Entrance Exam Score	0.27	-0.30	0.57***
Avg. 10th Grade Score	0.23	-0.15	0.38***
Avg. 12th Grade Score	0.17	-0.10	0.27***
Avg. Overall College GPA	0.34	-0.33	0.67***
		M.S.	
Avg. Entrance Exam Score	-0.02	0.02	-0.04
Avg. 10th Grade Score	0.03	-0.07	0.10
Avg. 12th Grade Score	0.02	-0.06	0.08
Avg. Overall College GPA	0.04	-0.14	0.18*

Notes: Table 2 reports pre-college skills (entrance exam scores, 10th and 12th standard exit exams) and college GPA. All scores are normalized to zero mean and unit standard deviation within each degree and year. Entrance exam scores (originally ranks) are renormalized so higher numbers are better. Caste differences shown in standard deviation units. $^{*}p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$. Adv. Caste and Disadv. Caste denote advantaged and disadvantaged castes.

4.6 Characteristics of Multinational Private-Sector Jobs

This subsection summarizes key job characteristics, including sector distribution, salary structure, recruitment procedures, and managerial involvement.

Sector and employer overview. The dataset includes approximately 1,050 private-sector jobs, mostly at multinational corporations headquartered in the US or Europe, recruiting primarily for Indian offices (Table 3). Around 75% of roles are non-client-facing (e.g., software engineering), with 25% client-facing (e.g., consulting).

Exclusion of public-sector jobs. Public-sector roles (fewer than 2%; Table 3) are excluded due to distinct salary structures and stability.[6]

Standardized, non-negotiable salaries and benefits. Salaries and non-wage benefits (e.g., insurance, signing bonuses) are standardized, non-negotiable, and uniform for each job role, independent of individual characteristics like caste or gender. Firms typically adopt nationwide salary standards, recruiting identically across elite campuses.[7]

Table 3: **Summary Statistics (Jobs)**

	Count	Fraction	Avg. Salary ($)	Std. Dev. ($)
Private	1011	0.98	60341.38	32591.51
Government	19	0.02	56202.23	20240.01
Headquartered Outside India	620	0.60	64562.74	36038.04
Headquartered in India	410	0.40	53766.04	23946.16
Client Facing	246	0.24	57553.23	28288.23
Non-Client Facing	784	0.76	61115.92	33361.87
All Jobs	1030	1.00	60265.02	32391.89

Notes: Table 3 reports the distribution of jobs between public and private sectors, domestic and foreign headquarters (primarily US or Europe), and client-facing versus non-client-facing roles. Average salaries (USD PPP) are also shown for each group.

[6]For details, see Seventh Central Pay Commission, 2016.

[7]Large multinationals typically centralize core HR policies at corporate headquarters, codifying pay scales, incentives, and promotion criteria for all entry-level roles (e.g.,. Bloom et al., 2012; Sarsons et al., 2022; Minni, 2023, 2024; Sadun et al., 2024). Campus recruiters therefore have no discretion to vary compensation across locations or candidates, leading to nationally uniform, non-negotiable salary packages.

Job location and recruitment processes. Over 98% of jobs are India-based, employing unified national recruitment processes. Thus, jobs generally represent firm-designation pairs rather than specific office locations. Roles abroad (under 2%) are excluded due to different compensation and hiring structures.

Recruitment conducted by Indian hiring managers separately from direct supervisors. Because most jobs are based in India, recruitment is conducted by specialized Indian hiring managers. These managers typically remain distinct from eventual direct supervisors due to logistical constraints posed by campus visits across India.

Internship return offer policy. Students accepting internship return offers exit the main placement cycle and are excluded from analysis (Section 3.2).

4.7 Quantifying Job Selectivity at Each Recruitment Stage

This subsection summarizes the selectivity of jobs using two key concepts—*conditional cuts*, the fraction of applicants eliminated *at each round*, and *cumulative cuts*, the total fraction eliminated *up to a given round*. Table 4 reports these metrics. For instance, a job eliminating 50% of applicants each round has a 50% conditional cut at each stage, but a cumulative cut of 75% after two rounds. Typically, successful candidates rank within the top 5% of initial applicants. Occasionally, highly selective roles extend no offers.

Table 4: **Variation in Conditional and Cumulative Selectivity Across Jobs**

Cut Type	Round	Q1	Mean	Q2	Q3
Conditional Cut	App. Reading	0.30	0.39	0.39	0.48
Cumulative Cut	App. Reading	0.30	0.39	0.39	0.48
Conditional Cut	Written Test	0.10	0.18	0.15	0.26
Cumulative Cut	Written Test	0.38	0.49	0.49	0.61
Conditional Cut	Group Debate	0.69	0.74	0.74	0.79
Cumulative Cut	Group Debate	0.83	0.87	0.87	0.91
Conditional Cut	Personal Interview	0.61	0.75	0.78	0.90
Cumulative Cut	Personal Interview	0.95	0.96	0.97	0.99

Notes: Table 4 shows the distribution of conditional and cumulative selection cuts across jobs.

4.8 Four Recruitment Stages: From Objective to Subjective Evaluations

Recruitment for all jobs progresses uniformly from relatively objective assessments to subjective evaluations:

Low-cost applications via centralized portal. A single click on the centralized portal forwards the same standardized résumé to any additional job (Section 4.2), making the marginal cost of applying essentially zero; application volumes are therefore virtually identical across castes (Figure 1).

Eligibility criteria restrict applications. Applications are constrained by major, degree, and GPA. For example, software engineering roles often require Computer Science majors, whereas analyst roles typically allow all majors.

Four mandatory screening rounds. All jobs strictly follow these sequential stages:

1. **Application reading.** Employers use standardized criteria and GPA cutoffs to initially filter resumes, involving limited subjective assessment.
2. **Written aptitude tests.** Standardized online tests objectively evaluate technical and analytical skills.
3. **Group debates.** Structured debates involve 20–30 students divided into two teams, assessing communication, confidence, and teamwork under standardized conditions, including uniform prompts and equal speaking time per candidate.
4. **Personal interviews.** Final-stage interviews (30 minutes) heavily emphasize subjective evaluations of interpersonal skills, organizational fit, cultural alignment, and role-specific competencies.

Throughout each stage, employers maintain access to candidates' standardized profiles via the centralized portal, ensuring transparency and consistency. Additional details for screening rounds and candidate selection, including timeline, are shown in Online Appendix Section D.

Common screening rounds across multiple positions. Employers offering multiple distinct roles frequently administer common screening (e.g., same technical tests) simultaneously for all applicants.

4.9 Candidate Screening: Summary Statistics

Table 5 summarizes candidate progression through each recruitment stage. On average, jobs receive roughly 209 applications, with approximately 125 advancing to written tests, 102 reaching group debates, 26 participating in personal interviews, and about 6 candidates receiving job offers. This substantial reduction reflects high selectivity at each stage (Table 4). Eligibility constraints (major, degree, GPA) naturally limit the number of applications per student (Section 4.8).

Table 5: **Summary Statistics (Candidate Screening)**

Round	Mean	Median
Application Reading	208.85	140.17
Written Tests	124.91	87.00
Group Debates	101.54	70.00
Personal Interviews	26.07	18.00
Offers	6.49	4.00

Notes: Table 5 reports mean and median counts per job for applications, written tests, group debates, interviews, and offers. Values calculated by averaging student counts across years per job, then averaging across jobs.

4.10 Data on Job Tenure and Promotions

Career outcomes are measured using tenure and promotions, obtained for nearly all (over 98%) placed students using alumni records, also maintained by the placement office and largely curated from an employment-focused platform plus a few other sources (see Online Appendix C for coverage benchmarks). Tenure is calculated using reported employment start and end dates.

Classifying promotions. Promotions are defined as movements to higher job titles within the same firm validated by salary increases and typical career trajectories using external industry-standard benchmarks (such as Glassdoor, levels.fyi), effectively distinguishing them from lateral moves. Such classifications are particularly reliable for entry-level roles at multinational corporations with globally standardized career ladders and compensation structures (cf. Bloom et al., 2012; Minni, 2024).

4.11 Data Limitations and Mitigation

Two key limitations include the absence of detailed resumes and screening round performance scores. However, standardized portal data comprehensively capture relevant student characteristics (e.g., GPA, internships), ensuring robust empirical analysis. Strategies addressing the second limitation are discussed in Section 5.7.

5 Penalizing Social Status Through "Fit" Assessments

This section presents six core findings demonstrating how personal interviews penalize social status (caste) in entry-level hiring at multinational firms. First, roughly 90% of caste-based disparities in earnings and success rates arise specifically during subjective, "fit"-oriented personal interviews. Second, these disparities are likely conservative. Third, personal interviews are the primary mechanism through which social status is *revealed and penalized.* Fourth, caste penalties consistently emerge at personal interviews across nearly all job positions. Fifth, disadvantaged caste hires exhibit systematically higher promotion rates even in highly selective roles, strongly suggesting hiring disparities result from discrimination rather than unobserved skill differences. Finally, these caste penalties align most closely with *bias* rather than statistical or customer-driven discrimination.

5.1 Main Sample and Reliability of Initial Job Placements

As discussed in Section 3.2, approximately two-thirds of students secure full-time positions through the main placement process. This proportion aligns with findings from similar elite Indian institutions (e.g., Mamgain, 2019; Subramanian, 2019; Jodhka and Naudet, 2019). Specifically, fewer than 5% accept return offers from summer internships, while about 60% obtain employment through the standardized four-stage job fair. The remaining one-third participate in the subsequent *aftermarket* hiring phase.[8]

Offer-letter salaries are cross-checked by the placement office before graduation, reinforcing the accuracy of our wage data (Section 3.2). Nearly all placed students (>99%) begin employment in the roles offered. This is further verified through post-graduation surveys and cross-validated against external employment platforms, affirming the reliability of placement records as accurate indicators of initial job outcomes.

[8]Of the total 4,164 students, around 159 (4%) accept internship return offers, leaving 4,005 actively participating. Approximately 2,527 students (60% of the total cohort) secure jobs via the four-stage process, totaling 2,686 students (159 internship returns + 2,527 from placement rounds), or about 65% placed through the primary placement process. Consequently, roughly 1,478 students (one-third) enter the aftermarket. See Online Appendix Tables OA.1 and OA.2 for detailed statistics.

5.2 Caste Penalties in Earnings and Hiring Success Rates Emerge at Personal Interviews

Caste disparities in both earnings and hiring success rates predominantly arise at the subjective, fit-focused personal interview stage. Earlier recruitment rounds—application screening, written aptitude tests, and group debates—show minimal caste gaps after controlling for key student characteristics, including GPA, entrance exam scores, internships, major, and degree.[9]

Earnings Gap Decomposition. Disadvantaged castes earn significantly lower initial salaries overall: the unconditional earnings gap is approximately 23%, narrowing to roughly 10% after controlling for pre-college skills, college GPA, prior labor market experience, and other factors. This residual 10% gap—about 6,000 (PPP), or approximately 0.18 standard deviations of the private-sector salary distribution (Table 3)—reflects differences in job composition rather than pay disparities within identical roles, given standardized salaries (see also Online Appendix Table OA.1 for variations by job type).

Because jobs have *fixed, non-negotiable* salaries, caste earnings penalties across screening stages reflect differences in job access rather than candidate-specific salary variations (Section 4.6). This enables clear identification of where the caste earnings penalty first appears. Specifically, I quantify these earnings differences using the following regression specification, assigning posted salaries counterfactually to each caste group at each recruitment stage before final offers are accepted:

$$\log(\text{Avg. Job Salary}_{it}^{\text{Search Stage}}) = \alpha + \beta \times \text{Disadv. Caste}_i + \text{Controls}_i + \gamma_t + \epsilon_{it}, \quad (1)$$

where $i \in \mathcal{I}$ indexes students, and "Search Stage" $\in$ {Application, Written Tests, Group Debates, Personal Interviews, Offers, Accepted Offers}. Controls include GPA, entrance exam scores, prior internships, major, and degree; γ_t denotes cohort-year fixed effects. For each student, the "average job salary" at a given stage is the mean salary of all roles for which the student remains eligible.

The coefficient β, illustrated in Figure 1, quantifies the caste earnings gap at each recruitment stage. The application reading round shows no statistically significant earnings gap after controls, confirming that broad, low-cost applications (Section 4.8) lead to similar application behavior across castes. The written-test and group-debate rounds likewise reveal no meaningful gaps. Differential caste preferences over offered jobs also contribute negligibly to the total penalty. The earnings penalty emerges sharply at the personal interview stage,

[9] Even uncontrolled regressions show similar patterns. See Section 5.7.

accounting for roughly 90% of the total earnings gap observed among graduates in the sample. Full regression results with robustness checks using Oster (2019) bounds are presented in Online Appendix Table OA.4.[10]

Figure 1: **The Caste Earnings Penalty**

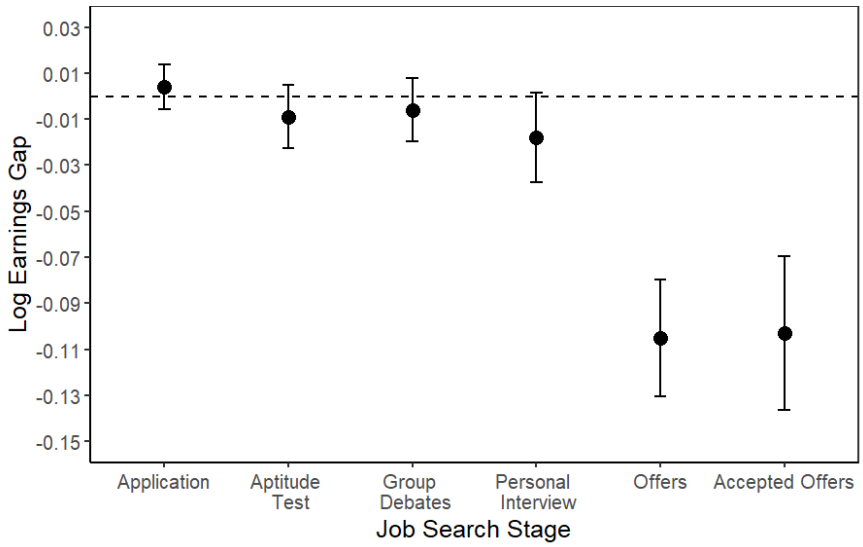

Notes: Figure 1 shows the log earnings gap across castes at each job search stage. Black dots represent coefficients showing the percentage difference in average salary between disadvantaged and advantaged castes, with 95% confidence intervals. Since salaries are identical and non-negotiable for identical jobs, the earnings gap reflects differences in job access at each stage. All regressions include controls.

Success Rate Gaps by Screening Stage. To provide complementary evidence to the earnings analysis, I examine success rates—the probability of progressing from one recruitment stage to the next *within* each job. Because salaries are fixed, a zero earnings gap at stage t+1 can occur only if both caste groups have similar odds of advancing at stage t; otherwise they would face different job sets and immediate pay differences. Success-rate regressions therefore furnish a direct consistency check on the earnings decomposition.

Unlike the earnings gap analysis, which cannot include job fixed effects, these linear-probability models exploit job and year fixed effects to isolate *within-job* caste disparities at each stage. Table 6 shows negligible caste differences at the application, written-test, and group-debate rounds. The gap opens only after personal interviews, where disadvantaged castes are about 30 percentage points less likely to advance, a drop that explains more than

[10]We apply Oster (2019)'s bounding methodology, setting $\delta = 1$ (assuming unobservables match observables in selection intensity) and $R^* = 1.3 \times \tilde{R}$, allowing the total explanatory power to exceed the observed R^2 by 30%. The main findings remain robust under these stringent assumptions and their variations.

95% of the overall success-rate disparity and aligns exactly with the interview-driven earnings penalty.[11]

Table 6: **Caste Differences in the Success Rate at Each Job Search Stage**

	Success App (1)	Success Test (2)	Success Debate (3)	Success Interview (4)
Disadv. Caste (OLS)	-0.0025*	-0.0002	-0.0005	-0.2938***
	(0.0014)	(0.0018)	(0.0027)	(0.0048)
Disadv. Caste (Oster)	0.009	0.010***	-0.002	-0.296***
	(0.0013)	(0.002)	(0.003)	(0.005)
Mean Success Rate	0.5990	0.8110	0.2580	0.2610
Observations	305,541	183,150	148,489	38,378
Adjusted R^2	0.61069	0.34360	0.01858	0.24450
Year FE	✓	✓	✓	✓
Job FE	✓	✓	✓	✓
Other Controls	✓	✓	✓	✓

Notes: Table 6 reports caste differences in success rates at each job search stage. Each observation is a student-job-year combination for candidates still in contention at that stage. All regressions include job fixed effects and controls. When firms conduct common screening rounds for multiple positions, each position is treated as a separate job. Bias-adjusted coefficients use Oster (2019) with $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$, assuming unobservables match observables in selection intensity and allowing total explanatory power to exceed observed R^2 by 30%. $^*p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

5.3 Robustness and Job-Level Heterogeneity in the Caste Penalty

These findings withstand extensive robustness checks, including a fine-grained *job-level* analysis that treats each employer–position pair as its own fixed effect. Because firms likely vary in written-test content, interview format, and caste composition, documenting a 30 percentage point caste penalty *within nearly every job* rules out the concern that unobserved résumé quality, firm-specific screening methods, or other forms of organizational heterogeneity explain the gap. Section 5.7 details these estimates and presents related job-performance evidence based on promotion outcomes.

5.4 Conservative Estimates Due to Exclusion of the Aftermarket

The earnings and success-rate gaps identified above likely represent conservative estimates. The main analysis excludes roughly one-third of students who enter the post-placement *aftermarket* hiring phase (Section 5.1). Disadvantaged caste students entering the aftermarket

[11] Based on the estimates in Table 6, disadvantaged caste candidates in our sample are about 70% less likely than their advantaged peers to clear the final interview stage (risk-ratio ≈ 0.28). Thorat and Attewell (2007) report an average 50% shortfall in callbacks for "disadvantaged" i.e., Dalit and Muslim names (risk-ratio ≈ 0.52). Although the point estimates differ, both imply that disadvantaged groups face roughly a one-half to two-thirds reduction in success probability, well within the same ballpark once sampling variation and design differences are taken into account.

generally exhibit much *lower* average GPAs than comparable advantaged castes (Online Appendix Table OA.2).[12] Because higher GPA typically correlates with higher-paying and more selective job placements, excluding a subset of lower-GPA disadvantaged students from the main analysis understates the true caste disparities.

5.5 Information Observed by Employers at Each Screening Round

Employers continuously access candidates' standardized academic profiles, including GPA, coursework, entrance exam scores, internships, and technical skills, through the centralized online job portal (Section 4.2). While early stages like application screening and written aptitude tests rely primarily on objective metrics (e.g., GPA cutoffs, standardized test scores), group debates and personal interviews introduce direct social interactions.

Group debates involve structured team discussions, allowing employers to assess communication, teamwork, and socio-emotional skills. However, these interactions remain standardized, with limited opportunity for individual background details or subtle indicators of caste to emerge.

In contrast, personal interviews lasting roughly 30 minutes each allow deeper, individualized conversations about candidates' interests, extracurricular activities, personal backgrounds, and interpersonal rapport. These social interactions significantly enhance employers' insights into candidates' cultural fit and organizational alignment, making subtle indicators of social status—particularly caste—more salient (cf. Bisin and Verdier, 2010; Benson et al., 2019; Rivera, 2015). Although we lack data on interviewers' own backgrounds, ethnographic evidence shows that evaluative norms are quickly socialized within elite firms, so managers with diverse origins converge on similar notions of fit, reinforcing status-based biases (e.g., Turco, 2010; Rivera, 2015). Section 5.6 provides quantitative evidence supporting this interpretation (qualitative evidence in Online Appendix Section E).

5.6 Interviews as the Main Mechanism for Revealing and Penalizing Social Status

When does caste become visible to employers? This subsection demonstrates quantitatively that subjective, fit-focused personal interviews are the primary stage at which caste is revealed and penalized.

[12] *Within* any offered job, disadvantaged hires have slightly higher GPAs than advantaged counterparts, consistent with the caste hiring penalty (Online Appendix Table OA.3).

5.6.1 Quantitative Evidence

Although many surnames are caste-indicative, most do not reliably signal caste in elite urban college settings, where students come from diverse regions with varied naming conventions (e.g., Jodkha, 2017; Parmar, 2020).[13] Consequently, only about 15% of students in my sample have clearly identifiable caste surnames. For this subgroup, caste penalties appear earlier—at the application reading round. However, for the remaining 85% with non-identifiable surnames, the caste penalty emerges predominantly during personal interviews where job-competencies, including candidate background and cultural fit, often become explicit topics of conversation (e.g., Rivera, 2012, 2015). Thus, personal interviews—particularly informal "fit" assessments—serve as the key mechanism for revealing and penalizing social status.

Caste earnings penalty emerges earlier when surnames reveal caste. Figure 2 replicates Figure 1, separately examining students with and without clearly identifiable caste surnames. I re-estimate Equation 1 separately for each subgroup. Students with identifiable surnames experience a significant penalty starting from application screening, while those without identifiable surnames see penalties emerge only at personal interviews.[14] Corresponding regression results, including Oster's bias-adjusted coefficients, are in Online Appendix Tables OA.7 and OA.8.

Similar patterns in success rates. The same story holds when we look at progression through the screening rounds. Students whose surnames clearly reveal caste suffer large penalties right from the application-reading stage across almost every job (Online Appendix Figure OA.2). By contrast, students whose surnames are *not* caste-identifiable track the overall pattern shown in Table 6: their success rates are essentially equal to advantaged peers until the personal interview, where the penalty appears and drives the final gap.

How are student names associated with caste? Surnames are coded as caste-identifiable versus not by utilizing official government sources, including regularly updated state-specific "schedules" of caste names (Scheduled Castes and Tribes, as described in Section 2). Common advantaged castes surname lists are compiled by the Anthropological Survey of India. Similarly, common disadvantaged castes (SC/ST/OBC) surnames are available through

[13] See also Bhupatiraju et al. (2024)—a study of more than one million High Court cases from the urban centre of Patna—which shows that roughly half of petitioners, respondents, and judges use *caste-neutral* surnames (49% of petitioners, 56% of respondents, and 50% of judges), underscoring how common name-based ambiguity is in urban India.

[14] For example, a purely income-based explanation would predict an early gap in both groups, contrary to what we observe.

Election Commission records of reserved constituencies and official state-maintained SC/ST/OBC lists.

Qualitative evidence. Summarized here (with full details in Online Appendix E), it shows that (i) in urban, highly educated settings, neither skin tone nor regional accent reliably signals caste; and (ii) the one-on-one, open-ended format of personal interviews both plausibly reveals new background cues—family, schooling, networks—and clarifies earlier, ambiguous signals, together making social-status distinctions much more salient.

Figure 2: **The Caste Earnings Penalty by Subsample**

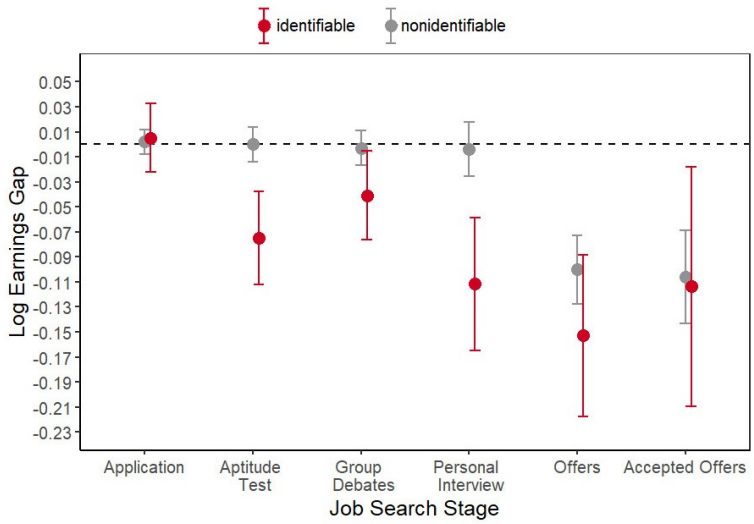

Notes: Figure 2 compares the log earnings gap across castes at each job search stage between students with identifiable (red) and non-identifiable (gray) caste surnames. Both gray and red dots show percentage differences in average salary between disadvantaged and advantaged castes, with 95% confidence intervals. Since salaries are identical and non-negotiable for identical jobs, gaps reflect differences in job access. All regressions include controls.

5.7 Heterogeneity in Caste Penalties, Selectivity, and Promotion Outcomes: Empirical Facts & Stylized Model

In this subsection, we examine *job-level heterogeneity* in the caste penalty, effectively netting out unobserved differences in screening tests, recruiting protocols, firm caste mix, and other organization-specific factors. We show that disadvantaged-caste applicants experience about a 30 percentage-point lower interview success rate *across nearly every job*, yet those who are hired exhibit roughly 20 percent higher promotion rates than their advantaged peers.

Since direct candidate performance measures at each stage are unavailable, unobserved ability differences might remain hidden until later rounds (Section 4.11). Thus, these headline findings raise two critical questions. First, might caste penalties appear earlier in *highly*

selective jobs, which recruit primarily from the right tail of the ability distribution? Second, could personal interviews reveal genuine unobserved ability differences—rather than discrimination—that justify hiring disparities across castes?

Variation in job-level selectivity offers an empirical test. If disadvantaged castes truly had lower pre-interview performance, such deficits should be particularly evident in highly selective jobs. Yet, pre-interview caste gaps remain negligible even among these selective positions. Furthermore, once hired, disadvantaged castes exhibit higher promotion rates. One might argue this simply reflects inframarginal productivity differences—specifically, a greater share of right-tail performers among disadvantaged hires. However, examining promotion patterns across jobs of increasing selectivity helps disentangle discrimination from ability differences. If hiring thresholds were identical across castes, disadvantaged hires' promotion advantage should shrink as jobs become more selective. Instead, we find this advantage remains stable across all selectivity tiers.

Below, I present two empirical facts leveraging job selectivity variation, followed by a formal model demonstrating that these patterns reflect a systematically stricter hiring threshold for disadvantaged castes. Fact 1 uses job-specific selectivity measures at each screening round, revealing no caste gap in pre-interview success rates across selectivity levels. Fact 2 categorizes jobs by selectivity and documents consistently higher promotion rates for disadvantaged hires throughout the selectivity distribution.

Fact 1 (Conditional Success Rates Reveal No Early Penalty, but a Large Final Interview Gap). Figure 5 plots, for each job, the conditional success-rate gap (disadvantaged minus advantaged) against the cumulative cut rate at each screening stage (defined in Section 4.7). Negative values indicate lower advancement rates for disadvantaged candidates relative to advantaged candidates. Formally, the gap for job j is the coefficient β_j from the regression:

$$\Pr(\text{Success}_{ijt}^{\text{Job Search Stage}}) = \alpha_j + \beta_j \times \text{Disadv. Caste}_i + \text{Controls}_i + \gamma_t + \epsilon_{ijt}, \tag{2}$$

where $i \in \mathcal{I}$ indexes students and γ_t denotes year fixed effects. Controls include GPA, entrance exam scores, prior internships, major, and degree.

Disadvantaged castes are about 30 percentage points less likely to advance than advantaged castes after personal interviews, consistently across nearly all jobs. Before personal interviews—during application reading, written tests, and group debates—the caste gaps remain negligible, *even* among highly selective jobs that eliminate most applicants. A distinct caste penalty emerges only after personal interviews, underscoring that late-stage subjective "fit" assessments, rather than pre-interview performance differences, drive observed caste dis-

parities. **Even uncontrolled regressions demonstrate the same pattern** (see Online Appendix Figure OA.3).

Fact 2 (Disadvantaged Hires Earn Comparably Higher Promotion Rates, Even in Highly Selective Jobs). To further test whether interviews reflect genuine ability or discrimination, I examine post-hire performance outcomes. Using data from alumni records, I successfully track career trajectories for all but approximately 2% of the students placed into full-time jobs through the standardized placement process, employing promotion rates as a measure of job performance.[15] If disadvantaged castes genuinely had lower unobserved abilities revealed during interviews, their relative performance should decline in more selective jobs, which impose higher—but equal—standards across castes. Instead, disadvantaged hires consistently maintain their promotion advantage even in highly selective positions. To demonstrate this, I estimate caste differences in promotion rates using the following regression across job selectivity tiers:

$$\text{Outcome}_{ijt} = \alpha + \beta \times \text{Disadv. Caste}_i + \text{Controls}_i + k_j + \gamma_t + \epsilon_{ijt}, \tag{3}$$

where $i \in \mathcal{I}$ indexes students, $j \in \mathcal{J}$ indexes jobs, γ_t denotes year fixed effects, and k_j represents job fixed effects. Controls include GPA, entrance exam scores, prior labor market experience, major, and degree. Table 7 presents these results:

- **Columns (1)–(2)** show that, pooling all jobs, disadvantaged hires have approximately a 20% *higher* probability of promotion relative to advantaged peers, despite initial hiring penalties.
- **Columns (3)–(8)** group jobs by *selectivity* (final rejection rates), revealing that disadvantaged hires consistently maintain a stable promotion advantage across the selectivity distribution. For most tiers, from the top 75% to the top 20%, the advantage is approximately 20-25% and statistically significant. In the most-selective tier (top 10 percent), the point estimate remains comparable (0.084) but is imprecise (unsurprising with only $N = 252$ hires) so it does not attain conventional significance. Taken together, the coefficients show that disadvantaged hires outperform advantaged peers even after the most stringent selection.
- **Columns (9)–(10)** demonstrate consistent patterns in both client-facing and non-client-facing roles, indicating that the promotion advantage is not specific to particular job functions.

[15] See Section 4.10 for more details and Section 5.1 for details on the placed sample.

India's elite civil service—which likewise screens candidates belonging to an equally elite talent pool (many aspirants come from top Indian universities) through a written exam followed by a "personality'' interview—replicates the late-stage caste penalty and its source: disadvantaged caste officers with similar or better written-test scores but markedly lower interview ratings outperform advantaged castes by about 18–20 percent (Table 2, Bhavnani and Lee 2021).[16]

Promotions as a performance metric. It is possible that promotions do not fully capture what firms are seeking to maximize, but it is hard to imagine that job performance is not captured by promotions. Promotions are also a particularly good measure of productivity in this setting because high-skilled, entry-level jobs at elite MNCs have rigid salary structures globally (e.g., Bloom et al., 2012; Minni, 2024).

Promotion-rate estimates are conservative. Global promotion grids at multinationals are formally standardized (e.g., Bloom et al., 2012), yet local managers still rely on discretionary ratings of "leadership," "communication," and "potential" (Benson et al., 2019)—subjective assessments that should *lower* disadvantaged candidates' odds of promotion, not raise them, because such criteria routinely privilege "cultural fit" and other class-coded signals valued by managers themselves (e.g., Rivera, 2012; Turco, 2010; Gorman, 2005). Explicit caste-conscious preferences in hiring or promotion are rare: even in the U.S. fewer than 10% of major firms mention social class in DEI statements (Stansbury and Rodriguez, 2024b), and caste references are virtually absent in Indian multinationals (Bapuji et al., 2023). Hence the 20% promotion edge we observe is, if anything, a lower bound.[17] Online Appendix Table OA.9 corroborates this conservatism: advantaged-caste job leavers (during the period for which tenure and promotion data were collected) typically have lower GPAs, whereas promotions correlate (albeit weakly) positively with GPA.

Key Takeaway. Facts 1 and 2 challenge the notion that personal interviews merely reveal latent skill deficits among disadvantaged castes. If disadvantaged candidates genuinely performed worse before interviews, we would observe significantly lower pre-interview success rates—especially in highly selective jobs. Similarly, if interviews accurately filtered out weaker candidates, promotion rates for disadvantaged hires should be comparable or lower.

[16] Such reverse "performance" gaps, observed only after selection, are consistent with outcome-based tests for discrimination in healthcare and algorithmic-fairness research, where higher conditional success rates for marginalized groups signal stricter screening thresholds (e.g., Chandra and Staiger, 2010; Abaluck et al., 2016; Chandra and Staiger, 2020).

[17] A 2018 U.S. survey found 67% of disadvantaged-caste employees reported workplace mistreatment (California v. Cisco Systems, Inc., 2022), reinforcing the bias concern.

Instead, disadvantaged hires consistently achieve *higher* promotion rates, even in the most selective jobs. Crucially, this promotion advantage does not diminish with job selectivity, contradicting explanations based solely on unobserved ability differences.

5.7.1 Explaining the Empirical Facts: A Simple Model of Late-Stage Discrimination

This section introduces a unified theoretical model demonstrating how (i) the absence of early-stage gaps for disadvantaged castes implies no genuine pre-interview ability deficit, and (ii) variation in hiring thresholds across jobs helps distinguish discrimination from genuine ability differences revealed at the interview stage.

Consider a set of jobs $j \in \mathcal{J}$. Each job has up to four sequential rounds $s \in \{1, 2, 3, 4\}$:

- Rounds $s = 1, 2, 3$ are pre-interview screening rounds (e.g. application, written tests, group debates)
- Round $s = 4$ is the personal interview

Groups and Screening. There are two castes (groups) $g \in \{0, 1\}$, where $g = 1$ denotes disadvantaged castes and $g = 0$ denotes advantaged castes. At each round s, employers observe a candidate-specific score $M_{s,i,g}$, drawn from a common across-jobs distribution F_g (either unconditional or residualized on observables; results are identical), for candidate i in group g.[18] A candidate passes round s if and only if their score satisfies $M_{s,i,g} \geq \tau_{s,j}$ and if they have successfully passed all preceding rounds. At the final interview stage ($s = 4$), the threshold $\tau_{4,j}^{(g)}$ may differ by group g, reflecting the fact that caste becomes salient for approximately 85% of candidates at this stage (Section 5.6.1). Thus, a final-stage caste gap is indicative of discrimination if $\tau_{4,j}^{(1)} > \tau_{4,j}^{(0)}$.

Perfect information. For simplicity, we assume employers observe these scores *perfectly* at each stage. This assumption simplifies exposition but is not essential for our key insights: even under imperfect information, variation in screening selectivity and hiring thresholds across jobs allows us to distinguish between discrimination and genuine ability differences (see Online Appendix Section F for details).

[18] We assume the performance distributions for each group are continuous with smooth densities, finite moments, and satisfy standard regularity conditions. These technical assumptions ensure well-behaved pass rates and expectations when calculating promotion probabilities. While these properties underpin our formal analysis, they are not critical to understanding the main conceptual results.

Monotonicity of thresholds. We assume that thresholds are monotonically increasing across screening rounds:

$$\tau_{4,j} \;\geq\; \tau_{3,j} \;\geq\; \tau_{2,j} \;\geq\; \tau_{1,j},$$

implying each successive round is at least as demanding as the previous one.

Assumption 1 (Identification from Exogenous Variation in Screening Thresholds Across Jobs). For identification, we assume the underlying group ability distributions are common across jobs. This assumption is essential because it ensures differences in hiring thresholds τ_j reflect genuine variations in job selectivity, rather than systematic differences in candidate quality.

To clarify this idea, imagine a large tub of ice cream representing a common talent pool. Each job selects candidates from this pool using a ladle of a specific size: a smaller ladle (higher cutoff) selects fewer candidates but with higher average quality; a larger ladle (lower cutoff) selects more candidates, including those with lower scores. Since all jobs draw from the same pool, ladle size captures selectivity alone. This assumption is justified empirically by several institutional and empirical features:

1. **Quota-driven representation and GPA overlap.** Caste-based admission quotas create almost equal representation *within each major*; there's also caste overlap within every GPA decile (Sections 4.4–4.5; Online Appendix Figure OA.1).

2. **Restricted pre-interview sorting.** Near-zero application costs, broad eligibility rules, and standardized recruitment protocols leave employers little scope to manipulate cutoffs or screen candidates by caste before interviews; caste is largely invisible at earlier stages (Section 4.8).

3. **No observed caste gap pre-interview.** Empirically, there is *no* caste difference in advancement through application reading, tests, or debates—even for highly selective jobs (Fact 1). If *cumulative* performance deficits existed before interviews, they would be most visible precisely in these high-cutoff positions, yet they are absent. This direct evidence is the strongest validation of the common-pool assumption.

Thus, *within a given job*, a higher cutoff naturally selects candidates of higher average quality. However, *across jobs*, variations in cutoffs reflect exogenous differences in selectivity, allowing us to use this variation as an instrument in our empirical analysis.

(A) Ruling Out Systematic Performance Differences Before Interviews

We first rule out systematic caste differences in candidate performance at pre-interview screening rounds.

Proposition 1 (No Early Gap Implies No Right-Tail Deficit Before Interviews)**.** Consider a job j with up to four rounds of screening, where the first three ($s = 1, 2, 3$) are pre-interview stages. Under the common-pool assumption (Assumption 1) that each job draws applicants from the same ability distributions F_g, increasing the cutoff $\tau_{s,j}$ across jobs is equivalent to looking farther into the shared right tail of F_g. Hence, if group 1 truly performed systematically worse, then for a sufficiently high threshold $\tau_{s,j}$ at some pre-interview stage $s < 4$, the pass-rate gap

$$\Delta_{s,j} \;=\; P(\text{Pass}_{s,i,j}|g = 1) \;-\; P(\text{Pass}_{s,i,j}|g = 0) \;<\; 0$$

would necessarily arise. Therefore, observing that $\Delta_{s,j} \geq 0$—even under high thresholds that specifically focus on the right tail—rules out any meaningful right-tail performance deficit for group 1 before the interview stage.

Sketch of Proof. We illustrate how variation in screening thresholds helps identify systematic performance differences at pre-interview rounds.

Figure 3 contrasts two scenarios. Each panel depicts two groups' performance distributions in pre-interview screening rounds, where blue (solid) represents advantaged candidates ($g = 0$) and red (dashed) represents disadvantaged candidates ($g = 1$). The shaded portion above each vertical dashed line represents candidates passing the screening cutoff for a specific job. The difference between the panels highlights how job selectivity can test for pre-interview performance deficits:

- **Panel (a): Hypothetical Right-Tail Deficit.** Suppose the screening score distribution for each group g, denoted by $\varphi_g(m)$, has a regularly decaying tail; specifically, assume $P(M_{s,i,g} \geq m) \sim C_g\, m^{-\alpha}$ for sufficiently large m, with constants $C_g > 0$ and $\alpha > 1$. This assumption implies that, as screening thresholds rise (reflecting more selective jobs), the pass rate shrinks at a rate determined by the tail density. If group 1's right tail is thinner than group 0 (i.e., a smaller C_1 or larger effective decay rate), the red "passing region" shrinks *more quickly* than the blue, resulting in increasingly negative pass-rate gaps as selectivity increases.

- **Panel (b): No Meaningful Deficit.** In contrast, Fact 1 shows no significant shortfall for disadvantaged candidates (Group 1) at any pre-interview screening stage. Even at high cutoffs for highly selective jobs, both groups maintain similar pass rates. This empirical result holds consistently across application reading, written tests, and group debates.

Key Takeaway. Fact 1 strongly supports Proposition 1. Thus, any observed late-stage gap at the interview stage ($s = 4$) must stem either from newly revealed ability signals

during interviews or from discrimination based explicitly on social status.

Figure 3: **Using Variation in Screening Thresholds to Test for Pre-Interview Performance Differences**

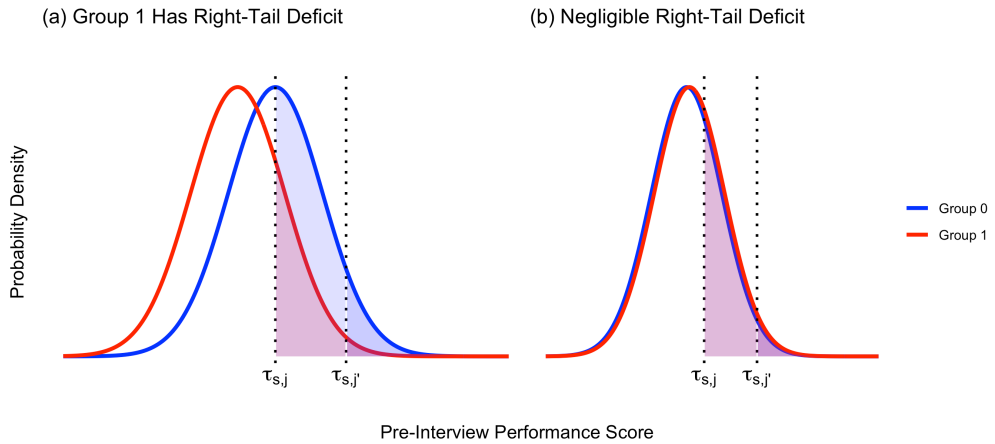

Notes: Figure 3 illustrates how selective screening would reveal systematic pre-interview performance differences. Panel (a) shows how a hypothetical right-tail deficit for Group 1 would become more apparent at higher cutoffs. Panel (b) shows the likely distribution of pre-interview performance scores implied by Fact 1, with no meaningful deficit. Shaded regions show candidates who pass the thresholds (blue for Group 0, red for Group 1, purple where they overlap). $\tau_{s,j}$ denotes screening threshold at stage $s < 4$ for job j (j' represents a more selective job).

□

(B) Using Variation in Hiring Thresholds to Test for Discrimination

Intuition. Variation in job-level hiring cutoffs provides a way to test whether the persistent 20% promotion advantage for disadvantaged hires arises from discriminatory hiring standards or genuine ability differences. If both castes faced identical hiring thresholds, this promotion advantage would diminish as jobs become increasingly selective. The persistence of this advantage instead points toward discrimination arising from systematically stricter hiring thresholds applied to disadvantaged candidates.

Framework. To formalize this intuition, we link pre-hire screening performance to post-hire job performance. For simplicity, assume job performance equals the average screening score across all four rounds. Specifically, let $\bar{M}_{i,g} = \frac{1}{4}\sum_{s=1}^{4} M_{s,i,g}$ denote candidate i's average screening score from group g.[19] Candidate i is promoted if their average score exceeds a job-specific promotion threshold κ_j.

We distinguish two types of hiring thresholds for jobs ranked by selectivity. For jobs within the top $(1-q)$ fraction by selectivity (e.g., if $q = 0.8$, the top 20% most selective jobs):

[19] Any monotonic link $f(.)$ between screening round performance and job performance works; we use the average for simplicity.

- The *selectivity-tier threshold* $\tau^{q}_{4,j}$ captures the common interview cutoff—*identical for both castes*—across jobs within this selectivity tier.
- The *group-specific threshold* $\tau^{(g)}_{4,j}$ reflects potentially caste-specific hiring standards within each job.

  Under discrimination, we expect $\tau^{(1)}_{4,j} > \tau^{(0)}_{4,j}$. Without discrimination, thresholds coincide: $\tau^{(1)}_{4,j} = \tau^{(0)}_{4,j} = \tau^{q}_{4,j}$.

Let

$$\bar{Y}^{(g)}_{j} \;=\; \mathbb{E}\big[\bar{M}_{i,g} \,\big|\, M_{4,i,g} \geq \tau^{(g)}_{4,j}\big]$$

represent the average screening score of hired candidates from group g. Define the promotion gap within selectivity tier q as

$$\Delta^{\text{Prom}}(q) \;=\; \bar{Y}^{(1)}_{j} - \bar{Y}^{(0)}_{j}.$$

Link to Assumption 1. Under the common–pool assumption, the job–specific interview cutoff $\tau_{4,j}$ acts like an *instrument*: an exogenous knob that tightens selection stringency without shifting underlying talent. As $\tau_{4,j}$ rises across increasingly selective jobs, the average ability of hires converges to the marginal cutoff. Hence, if both castes face the *same* threshold, the promotion gap $\Delta^{\text{Prom}}(q)$ must shrink toward 0. If the gap instead stabilises at a positive constant δ, the only explanation is that disadvantaged applicants face a uniformly higher hiring bar ($\tau^{(1)}_{4,j} = \tau^{(0)}_{4,j} + \delta$). We formalise and test this implication in Proposition 2.

Proposition 2 (Stable Promotion Gaps Across Selectivity Tiers Indicate Discrimination)**.** Consider a set of jobs J_q within the top $(1-q)$ fraction of selectivity, with associated (group-common) thresholds $\tau^{q}_{4,j}$. Suppose that the average screening score among hired candidates for group g in these jobs is given by:

$$\bar{Y}^{(g)}_{j} = \mathbb{E}\left[\bar{M}_{i,g} \mid M_{4,i,g} \geq \tau^{(g)}_{4,j}\right].$$

Under the common-pool assumption, $\tau_{4,j}$ is an *exogenous shifter* of the marginal hire's ability, so any limit behaviour of $\Delta^{\text{Prom}}(q)$ is driven solely by differences in hiring thresholds across groups.[20]

If no discrimination exists ($\tau^{(1)}_{4,j} = \tau^{(0)}_{4,j}$), then as $q \uparrow$ (as jobs become more selective), average scores of hired candidates converge to the common cutoff, and thus $\Delta^{\text{Prom}}(q) \to 0$.

[20]**IV analogy.** Instrument: $Z_j = \tau_{4,j}$ (varies widely across jobs). Relevance: $\text{Cov}(Z_j, \text{selection margin}) \neq 0$. Exclusion: Z_j affects promotion Y_{ij} *only* through the selection rule $H_{ij} = 1\{A_i \geq \tau^{(G_i)}_{4,j}\}$, which is valid because all jobs sample from the same $F(\cdot)$. Hence $\Delta^{\text{Prom}}(q)$ mimics a Wald estimator: if thresholds are equal, it tends to 0 as $q \to 1$; if disadvantaged candidates face $\delta > 0$ higher thresholds, it converges to that δ.

Conversely, if discrimination is present, with disadvantaged candidates facing a higher cutoff ($\tau_{4,j}^{(1)} = \tau_{4,j}^{(0)} + \delta$ for some $\delta > 0$), then average hired scores for disadvantaged candidates stabilize at approximately δ higher. Consequently, the promotion gap $\Delta^{\text{Prom}}(q)$ remains consistently positive across all levels of selectivity.

Sketch of Proof. Figure 4 illustrates the key intuition. In each scenario, average screening performance for the two groups is shown, with red (Group 1) consistently outperforming blue (Group 0) post-hire due to stricter selection.

- **Panel (a): No Discrimination.** Without discrimination, both groups face a common hiring cutoff: $\tau_{4,j}^{(1)} = \tau_{4,j}^{(0)} = \tau$. Under the regularity condition that each group's screening score distribution decays regularly—specifically, $P(M_{s,i,g} \geq m) \sim C_g\, m^{-\alpha}$ for large m, constants $C_g > 0$, and $\alpha > 1$, we have: $\lim_{\tau\to\infty} \big(\mathbb{E}[\bar{M}_{i,g} \mid M_{4,i,g} \geq \tau] - \tau\big) = 0$.

  Formally, with common CDF $F(a)$ and pdf $f(a)$,

$$\bar{Y}_j^{(g)} \;=\; \frac{\displaystyle\int_{\tau_{4,j}^{(g)}}^{\infty} a\, f(a)\, da}{1 - F\big(\tau_{4,j}^{(g)}\big)} \;=\; \tau_{4,j}^{(g)} + \frac{\displaystyle\int_{\tau_{4,j}^{(g)}}^{\infty} \big(1 - F(a)\big)\, da}{1 - F\big(\tau_{4,j}^{(g)}\big)} \;=\; \tau_{4,j}^{(g)} \;+\; \frac{1}{h\big(\tau_{4,j}^{(g)}\big)},$$

  where $h(\tau) \equiv f(\tau)\big/\big[1 - F(\tau)\big]$ is the hazard rate. As $\tau \uparrow \infty$ and the tails are regularly varying, $h(\tau) \to \infty$, so $\lim\limits_{\tau\to\infty} \big(\bar{Y}_j^{(g)} - \tau\big) = 0$.

  Consequently, as the common cutoff rises in more selective jobs, each group's average hired score converges to the marginal threshold, and the performance gap between groups shrinks to zero.

- **Panel (b): Discrimination.** With discrimination, disadvantaged candidates (Group 1) face systematically higher hiring thresholds than advantaged candidates (Group 0), such that $\tau_{4,j}^{(1)} = \tau_{4,j}^{(0)} + \delta$, $\delta > 0$. Under the same regularity condition stated above, as thresholds become increasingly selective:

$$\lim_{\tau\to\infty} \Big(\mathbb{E}[\bar{M}_{i,1} \mid M_{4,i,1} \geq \tau_{4,j}^{(1)}] - \mathbb{E}[\bar{M}_{i,0} \mid M_{4,i,0} \geq \tau_{4,j}^{(0)}]\Big) = \delta.$$

  Thus, the average hired score for disadvantaged candidates stabilizes at approximately δ points higher than for advantaged candidates, preserving a consistent performance gap across increasingly selective jobs.

Since a candidate is promoted if $\bar{M}_{i,g}$ exceeds a job-specific threshold κ_j, the asymptotic behavior of the screening score distributions implies that if both groups face the same thresh-

old, the promotion gap $\Delta^{\text{Prom}}(q)$ shrinks toward zero, whereas a constant gap in thresholds (indicative of discrimination) preserves a stable promotion gap.

Not every hire sits on the margin. The approach remains robust even when most hires comfortably exceed hiring thresholds. What matters for identification is how *thresholds* shift with selectivity, enabling clear discrimination tests irrespective of inframarginal hires.

Figure 4: **Using Variation in Hiring Thresholds To Test for Discrimination**

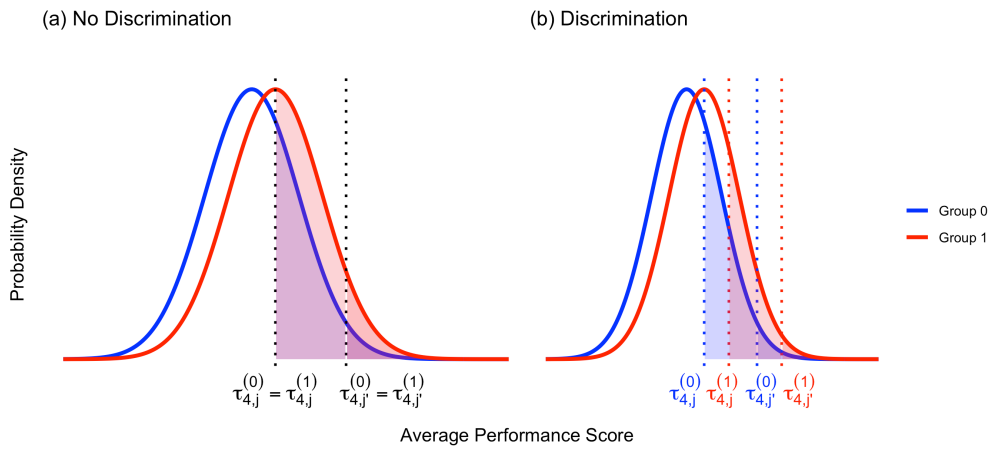

Notes: Figure 4 depicts two scenarios. In panel (a), a common cutoff for both groups shrinks their hired-performance gap at higher cutoffs. In panel (b), a constant δ gap in cutoffs preserves a stable performance advantage for Group 1. Shaded regions show candidates who pass the thresholds (blue for Group 0, red for Group 1, purple where they overlap). Here, $\tau^{(g)}_{4,j}$ denotes group g's interview threshold at job j, with j' representing a more selective job.

□

Key Takeaway. Fact 2 strongly supports Proposition 2. The persistent promotion advantage for disadvantaged hires across all selectivity tiers indicates that disadvantaged candidates face consistently higher hiring thresholds. Such a pattern is incompatible with explanations based solely on unobserved ability differences. Instead, it reinforces the assumption of latent group ability being common across jobs: once disadvantaged applicants meet the stricter bar, their post-hire performance is often stronger than that of advantaged peers, pointing squarely to discriminatory hiring standards rather than latent skill gaps.

Summary. Together, the negligible pre-interview caste gaps (Fact 1) and persistent post-interview promotion advantages (Fact 2) provide robust empirical evidence of discrimination at the hiring stage, particularly during personal interviews. These results are consistent with the model predictions of discriminatory hiring standards (Propositions 1 and 2), ruling out unobserved ability differences as the primary explanation.

Figure 5: **Caste Penalty in the Conditional Success Rate At Each Screening Round**

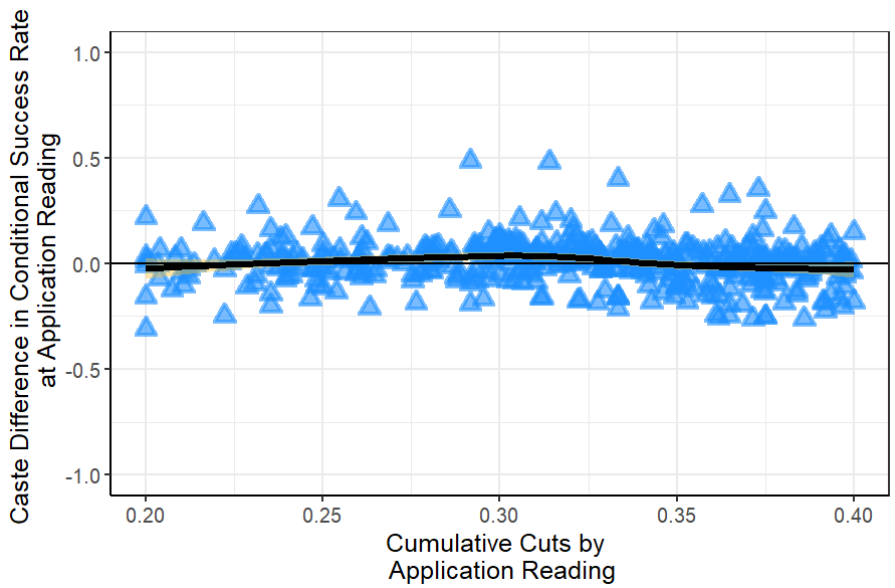

(a) Caste Penalty in Conditional Success Rates versus Cumulative Cuts by Application Reading

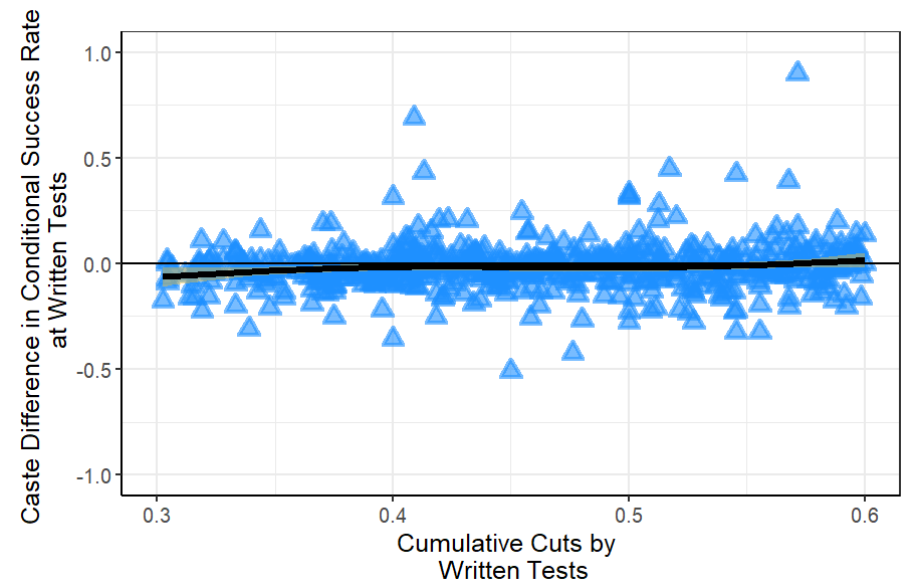

(b) Caste Penalty in Conditional Success Rates versus Cumulative Cuts by Written Tests

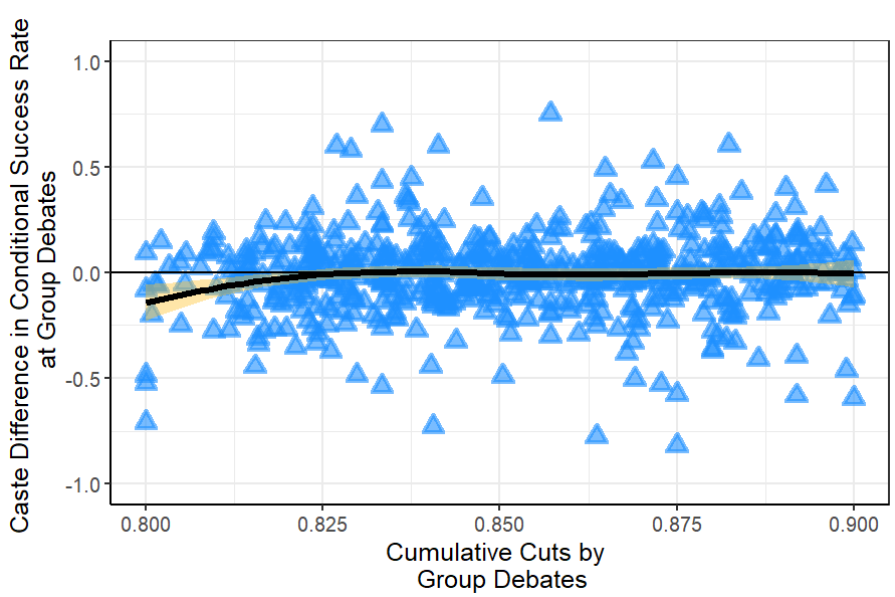

(c) Caste Penalty in Conditional Success Rates versus Cumulative Cuts by Group Debates

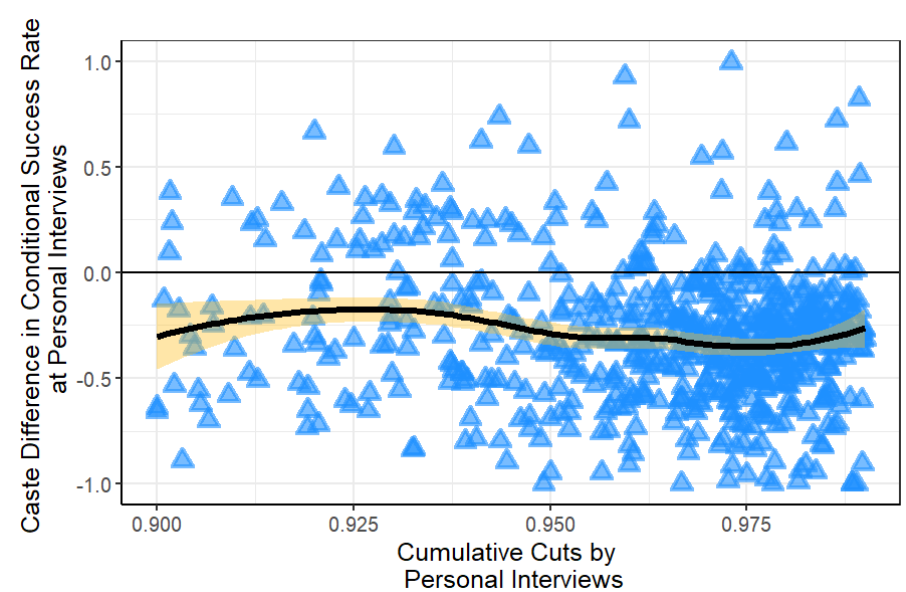

(d) Caste Penalty in Conditional Success Rates versus Cumulative Cuts by Personal Interviews

Notes: Figure 5 plots caste differences in conditional success rates against cumulative cuts at each screening round. Blue triangles represent individual jobs, with smoothed conditional means (black lines) and 95% confidence intervals (yellow). The x-axis shows the cumulative proportion of students eliminated by each round; the y-axis shows the caste coefficient from job-specific regressions with controls (degree, major, GPA, entrance exam scores) and cohort-year fixed effects. The horizontal line at zero indicates no caste difference.

Table 7: **Caste Differences in Promotion Rate**

	Dependent variable:										
	Promoted	Promoted	Promoted at Top 75%	Promoted at Top 50%	Promoted at Top 40%	Promoted at Top 30%	Promoted at Top 20%	Promoted at Top 15%	Promoted at Top 10%	Promoted at Client Facing	Promoted at Non-Client Facing
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)	(9)	(10)	(11)
Disadv. Caste (OLS)	0.137***	0.096***	0.104***	0.112***	0.136***	0.147***	0.190**	0.204**	0.084	0.107**	0.082**
	(0.025)	(0.026)	(0.029)	(0.038)	(0.043)	(0.054)	(0.073)	(0.084)	(0.109)	(0.044)	(0.032)
Disadv. Caste (Oster)	0.142***	0.093***	0.108***	0.122***	0.141***	0.157**	0.196**	0.206**	0.074	0.106***	0.075**
	(0.026)	(0.028)	(0.031)	(0.037)	(0.041)	(0.058)	(0.082)	(0.082)	(0.131)	(0.039)	(0.036)
Tenure Duration		0.084***	0.082***	0.076***	0.074***	0.068***	0.054*	0.079**	0.115**	0.065***	0.097***
		(0.012)	(0.014)	(0.019)	(0.021)	(0.025)	(0.030)	(0.037)	(0.049)	(0.021)	(0.015)
Mean Y	0.684	0.684	0.667	0.668	0.676	0.711	0.711	0.706	0.730	0.687	0.682
Observations	2,489	2,489	2,012	1,246	997	748	502	374	252	835	1,654
Adjusted R^2	0.183	0.205	0.203	0.216	0.211	0.182	0.192	0.189	0.127	0.194	0.214
Job FE	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
Year FE	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
Other Controls		✓	✓	✓	✓	✓	✓	✓	✓	✓	✓

Notes: Table 7 reports caste differences in promotion rates for graduates who obtained jobs through the main placement process. Column (1) includes only job and year fixed effects. Column (2) adds controls for GPA, major, degree, and job tenure. Columns (3)-(8) show results by job selectivity tiers, measured by cumulative applicant elimination rates. Columns (9)-(10) compare client-facing versus non-client-facing roles. Bias-adjusted coefficients use Oster (2019) with $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$, assuming unobservables match observables in selection intensity and allowing total explanatory power to exceed observed R^2 by 30%. $^*p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

5.8 Classification and Preferred Interpretation of the Caste Penalty: Distinguishing from Statistical or Customer Discrimination

Given the evidence that employers directly penalize caste (Section 5.6), the most plausible interpretation is *disparate treatment*—explicit discrimination based on caste. The disparities are also most consistent with *bias.* Four facts strongly support this view. First, disadvantaged castes are much less likely to get a job offer across the full distribution of jobs (Fact 1). Second, disadvantaged hires are consistently *more* likely to be promoted, even among the all selectivity tiers (Fact 2). Third, over three-quarters of the roles analyzed are non-customer-facing (Table 3). Fourth, disadvantaged hires outperform even within the minority of customer-facing roles (Table 7). Taken together, these findings are difficult to reconcile with explanations based purely on accurate statistical inference about productivity[21] or customer-driven discrimination; instead, they align closely with caste-based bias.

Explicit caste discrimination is prohibited under Article 15 of the Indian Constitution (Ministry of Law and Justice, Government of India, 2021). Although subjective criteria for assessing "fit" might inadvertently create a *disparate impact* (e.g., by favoring candidates with specific extracurricular backgrounds), the comprehensive evidence presented throughout Sections 5.2–5.7 provides the strongest support for the conclusion that observed hiring disparities primarily reflect *disparate treatment driven by bias.*

5.9 What Happens During Personal Interviews?

We do not directly observe interview transcripts, so we cannot parse the exact conversations. However, prior ethnographic research on elite hiring demonstrates that final-round interviews at multinational corporations place substantial weight on subjective markers of "cultural fit," which are closely tied to candidates' socioeconomic backgrounds rather than strictly job-relevant skills (e.g., Rivera, 2012, 2015; Lamont, 2009). Similar patterns have been documented in India's civil service examinations, where subjective "personality" interviews disproportionately penalize lower-status candidates despite having limited predictive power for job performance (Bhavnani and Lee, 2021). These external findings strongly support our interpretation that the significant late-stage caste penalty we document arises from discretionary, status-based evaluations rather than from objective assessments of ability. A richer qualitative account of how interviewers use cultural cues—complete with illustrative quotations from recruiters—is provided in Appendix G.

[21] An accurate statistical discrimination model typically assumes evaluators hold correct beliefs about average group-based productivity and that the underlying productivity distributions are relatively well-behaved (cf. Aigner and Cain, 1977).

6 Alternative Explanations

Below, I briefly summarize alternative explanations for the observed caste disparities emerging specifically during personal interviews.

1. **Early-stage inclusivity for optics.** Pre-interview pass rates differ by <1 percentage point (Table 6; Figure 5), and surname-identifiable applicants already face a substantial penalty at the initial application screening round (Figure 2; Online Appendix Figure OA.2). Moreover, there is weak institutional pressure on private employers to promote caste diversity: India lacks an EEOC-like regulator, caste or socioeconomic class is *not* a diversity metric that multinational firms typically monitor (Stansbury and Rodriguez, 2024a; Stansbury, 2024), and as recently as 2018, only 3% of India's top-100 listed firms even collected caste data (BusinessLine, 2018). In fact, many Indian MNCs explicitly view caste-targeted hiring as "anti-merit" (Jodhka and Newman, 2007).

2. **Delayed discrimination.** It remains unclear why employers would delay expressing caste preferences: when caste is identifiable through surnames, penalties emerge immediately at the application screening stage. If employers wished to exercise caste-based discrimination before interviews, the highly subjective group debate rounds—where employers have considerable discretion in candidate evaluation—provide ample opportunity. Yet no caste gaps emerge in these rounds. Interviews therefore matter primarily because they newly reveal caste for most candidates (Figure 2; Online Appendix Figure OA.2).

3. **Hidden ability gaps.** If disadvantaged castes had latent skill deficits, particularly in the right tail, we would observe lower pre-interview pass rates in highly selective jobs that target top performers. However, these rounds show no caste disadvantage (Fact 1). Moreover, hired disadvantaged candidates later outperform advantaged peers across job selectivity tiers (Fact 2)—further contradicting the notion of ability gaps.

4. **Cumulative performance deficits before interviews.** If disadvantaged candidates lagged on *aggregate* pre-interview performance, the gap would surface first—and most sharply—in the most selective jobs that have the highest performance cutoffs. Yet no caste penalty appears in these positions (Fact 1); interviews therefore cannot be "catch-up" screens that simply weed out lower cumulative performers.

5. **Attrition and external search.** Once students register and submit applications through the centralized portal, virtually no one exits the placement process for non-

elimination reasons—exits occur only through formal elimination decisions at each of the four strictly enforced screening rounds, with the placement office meticulously recording every advance-or-exit decision (Sections 4.1 and 4.3). Return offers from internships are rare (Section 3.2), and the placement office prohibits registered candidates from conducting parallel off-campus searches (cf. Central Placement Cell, Delhi University), leaving little scope for differential outside options to drive the observed gap.

6. **Socio-emotional skills.** Group debates—explicitly designed to evaluate communication, teamwork, and other socio-emotional skills—reveal no caste differences (Panels (c) of Figures 5–OA.3).

7. **Missing data or skipped rounds.** For students participating in the placement process, the placement office enforces mandatory attendance at every stage and strictly prohibits employers from skipping steps or altering procedures mid-process. Comprehensive exit surveys independently corroborate the completeness of official records (Sections 4.1 and 4.3; Section 5.1).

8. **Negotiation and salary differences.** Salaries and benefits are ex-ante standardized, strictly non-negotiable, and verified by the placement office. Exit surveys further confirm identical starting salaries (Section 3.2; Section 4.6; Section 5.1).

9. **Public-sector competition.** Quota-regulated government jobs account for only about 2% of all campus job listings and thus cannot meaningfully affect the observed caste disparities. All job listings are visible to both candidates and recruiters through a centralized placement portal, ensuring private-sector employers are fully aware of the limited public-sector competition. Moreover, only the most prestigious public-sector positions typically recruit at elite colleges (Table 3).

10. **Preferences over jobs.** Near-zero application costs lead students to apply broadly (Section 4.8). Moreover, preference-based sorting can at best explain only a small share of the initial earnings gap and none of the sharp interview-stage penalty (Figure 1).

11. **Preferential promotion for disadvantaged hires.** Given the rarity of caste-focused corporate initiatives (Jodkha, 2017; BusinessLine, 2018), and extensive survey evidence pointing to widespread bias *against* disadvantaged caste employees rather than in their favor (California v. Cisco Systems, Inc., 2022)[22], it is unlikely that the

[22]In a 2018 survey of South Asian employees in US workplaces, 67% of disadvantaged caste respondents reported experiencing unfair treatment (California v. Cisco Systems, Inc., 2022).

promotion results are confounded by policies explicitly favoring disadvantaged castes. Although a few multinational corporations initiated efforts to address caste discrimination in late 2022 and 2023 (Chakravartty and Subramanian, 2021; Bapuji et al., 2023), these changes occurred after our analysis period (Section 4.4).

12. **Alternative job-ladder explanation.** This explanation cannot account for our findings since promotion comparisons analyze identical positions (same job designation, same firm; see Section 3.1).

7 External Validity

How generalizable are findings from one elite Indian institution? While caste discrimination might initially appear uniquely Indian, five considerations underscore why these results hold broader implications, not only across India but also globally in elite hiring and other high-stakes selection contexts.

A uniquely comprehensive social status label. India's constitutionally recognized caste categories bundle economic, social, and cultural capital (Bourdieu, 1986) into a formally codified and well-recognized social hierarchy. This codification gives researchers precise measurement clarity typically unavailable in Western contexts (e.g., Beteille, 1969), where class-based discrimination must often be indirectly inferred from ambiguous signals such as educational background, accents, or informal networks (e.g., Rivera, 2015; Laurison and Friedman, 2019). Because the status cue is explicit rather than inferred, we can pinpoint *when* and *where* "fit"-based heuristics tip decisions against lower-status candidates in high-stakes settings, such as hiring at multinationals.

Mirroring global elite hiring pipelines. The hiring process we analyze is a textbook version of how top employers screen talent in India and abroad. Multinational corporations typically employ standardized salary bands and identical evaluation rubrics across their worldwide offices (cf. Bloom et al., 2012; Minni, 2024; Sadun et al., 2024). Within India, the same three-stage protocol—résumé screens, written tests, and a discretionary final interview—governs recruiting at virtually every premier engineering or management campus and even the Union civil-service exams (e.g., Subramanian, 2019; Bhavnani and Lee, 2021). That step-wise architecture also dominates on-campus hiring at Ivy-Plus universities and Oxbridge, where brief, rapport-heavy conversations decide who becomes a "culture carrier" (e.g., Rivera, 2012; Laurison and Friedman, 2019). By pinpointing class bias at precisely

this last gatekeeper stage, our evidence speaks directly to the mainstream hiring procedures used by global consultancies, tech giants, and investment banks (Rivera, 2015).

A globally competitive talent pool. The students in our sample resemble their Ivy-Plus counterparts on every external benchmark of human-capital quality. Using a harmonized alumni-earnings metric for 2,800 institutions worldwide, Martellini et al. (2024) place several elite Indian colleges alongside leading universities worldwide. That stature rests on an admissions bar arguably steeper than the Ivy League: fewer than 1% of the more than one million applicants win a seat at such institutions, and their alumni regularly helm Fortune 500 firms (Jodhka and Naudet, 2019). Because our focal college draws from a similar applicant pool, the labor-market segment we study is comparable in human capital quality to the most selective universities worldwide, bolstering the external validity of our findings.

Subjective "fit" screens reproduce status bias worldwide. Across settings, the discretionary final interview—or its admissions cousin, the "holistic" file read—systematically rewards the subtle cultural capital of higher-status candidates. As explicit status markers become less obvious in diverse urban labor markets worldwide (Jodkha, 2017; Clark, 2015), these subjective evaluations become increasingly critical gatekeepers. Evidence from U.S. and U.K. professional hiring shows that brief, rapport-heavy conversations often engender class gaps by privileging mannerisms, conversational style, and extracurricular touchstones that signal organizational likeness (Rivera, 2012; Laurison and Friedman, 2016, 2019; Stansbury and Rodriguez, 2024a). Cognitive research explains why: quick, intuitive assessments in interviews often rely on superficial impressions of rapport and familiarity that implicitly reflect socioeconomic background (Kahneman, 2011; Todorov, 2017). Similar dynamics pervade elite U.S. college admissions, where the shift from standardized metrics to "holistic" evaluations disadvantages high-achieving, lower-SES applicants by emphasizing subtle class-linked advantages over academic potential (Kuncel et al., 2013; Sacerdote et al., 2025). Legacy preferences and broad "character" reviews further tilt acceptances toward affluent applicants even when objective metrics are identical (Stevens, 2009). Indeed, evidence from elite public-sector hiring shows these subjective screens disadvantage lower-status candidates without predicting job performance (Bhavnani and Lee, 2021). Taken together, these findings underscore that the mechanism we document—status discrimination crystallizing at the final, subjective gatekeeper stage—is not idiosyncratic to caste or to India but a structural feature of elite selection systems globally.

Downstream erosion of colleges' upward-mobility policies. Policies that widen entry at the *front end* of elite pipelines often lose traction further *downstream*. In India, caste-targeted admissions quotas boost representation in top colleges, yet the current seat-allocation rule is riddled with design flaws that blunt their effect and spur costly litigation (Yenmez and Sönmez, 2022; Sönmez and Yenmez, 2024). As this study shows, discretionary hiring screens at multinational firms can erode the remaining gains. A parallel slippage appears in the United States: children from high-income families remain far likelier to parlay an Ivy-Plus degree into a top-tier career, aided by exclusive campus clubs and alumni networks that persist into professional life (Michelman et al., 2022). Those same peer ties transmit advantages to the next generation, tightening the intergenerational link between family pedigree and human capital (Barrios-Fernández et al., 2024). Even when low-SES or minority students clear the academic bar, ascent into executive roles still hinges on early-career social capital, leaving sizeable gaps in leadership and top-income ranks (Zimmerman, 2019). Together, these patterns indicate that admissions-focused reforms are necessary but not sufficient: without safeguards at the discretionary, reputation-driven stages of labor-market evaluation, elite pipelines can narrow precisely when candidates move from credentials to careers.

8 Conclusion

By opening the hiring "black box," this paper makes three novel contributions. First, by examining all recruitment steps—from initial applications through final job attainment— it provides the first, large-scale empirical evidence pinpointing exactly where class-based bias emerges during entry-level hiring at multinationals. Second, it quantifies this discrimination, showing that nearly 90% of disparities arise during subjective personal interviews evaluating candidate "fit." Third, disadvantaged candidates penalized by these interviews subsequently outperform advantaged peers by nearly 20% in promotion rates, demonstrating that subjective assessments often systematically undervalue talent from lower-status backgrounds and hinder efficient talent allocation.

The Indian setting provides unusual clarity for investigating status discrimination. Although caste is uniquely explicit, the underlying mechanism—subjective evaluations of who "belongs" or "fits"—is universal. As explicit class signals fade globally, subtle markers like conversational polish and extracurricular interests become increasingly influential, perpetuating socioeconomic disparities (e.g., Rivera, 2015; Jodkha, 2017; Laurison and Friedman, 2019; Stansbury and Rodriguez, 2024a). Rapid, intuitive judgments about rapport often overshadow deliberate, merit-based evaluations (e.g., Kahneman, 2011; Pinker, 2021), dis-

advantaging lower-status candidates worldwide.

These findings highlight a critical vulnerability in the affirmative action pipeline: successful educational interventions can be partially eroded by downstream subjective evaluations in hiring. Employers inadvertently re-privilege candidates from familiar socioeconomic backgrounds, prioritizing short-term cultural alignment over long-term performance potential (e.g., Benson and Shaw, 2024; Hoffman and Stanton, 2024).

What can firms do? Eliminating interviews entirely is impractical, but straightforward reforms—structured behavioral questions, blinded scoring, and linking interviewer ratings to candidate performance—can mitigate bias while preserving essential human judgment (cf. Highhouse, 2008; Dana et al., 2013), offering fruitful avenues for future research.

As remote and algorithmic screening expand, stage-by-stage audits—like those presented here—will be essential to designing fairer, more effective selection systems. Evidence from other high-stakes arenas shows that while simple rule-based decisions can yield large efficiency gains with modest equity costs (Albright, 2025), algorithmic defaults curb disparities only when overrides are limited and outcomes continuously monitored (Albright, 2019, 2024). Thus, whenever discretionary judgment creeps back in—whether in hiring, courts, or any algorithm-assisted decision—subtle cues linked to class, race, gender, or other stigmatized identities often resurface, reproducing inequality. Addressing this persistent barrier requires sustained policy and organizational reform, unlocking overlooked talent to the benefit of organizations and societies globally.

References

[1] Abaluck, J., L. Agha, C. Kabrhel, A. Raja, and A. Venkatesh (2016). The determinants of productivity in medical testing: Intensity and allocation. *American Economic Review 106*(12), 3730–3764.

[2] Ægisdóttir, S., M. J. White, P. M. Spengler, A. S. Maugherman, L. A. Anderson, R. S. Cook, C. N. Nichols, G. K. Lampropoulos, B. S. Walker, G. Cohen, and J. D. Rush (2006). The meta-analysis of clinical judgment project: Fifty-six years of accumulated research on clinical versus statistical prediction. *The Counseling Psychologist 34*(3), 341–382.

[3] Aigner, D. J. and G. G. Cain (1977). Statistical theories of discrimination in labor markets. *Industrial and Labor Relations Review 30*(2), 175–187.

[4] Albright, A. (2019). If you give a judge a risk score: Evidence from Kentucky bail decisions. Working paper.

[5] Albright, A. (2024). The hidden effects of algorithmic recommendations. Working paper, Federal Reserve Bank of Minneapolis.

[6] Albright, A. (2025, March). No money bail, no problems? Trade-offs in an automatic pretrial release program. Working paper.

[7] Altonji, J. G. and R. M. Blank (1999). *Race and Gender in the Labor Market. Handbook of Labor Economics*, ed. by O. C. Ashenfelter and D. Card, Elsevier, vol. 3C, chap. 48.

[8] Alvero, A., S. Giebel, B. Gebre-Medhin, A. L. Antonio, M. L. Stevens, and B. W. Domingue (2021). Essay content and style are strongly related to household income and sat scores: Evidence from 60,000 undergraduate applications. *Science Advances 7*(42), eabi9031.

[9] Ang, D. (2021). The effects of police violence on inner-city students. *Quarterly Journal of Economics 136*(1), 115–168.

[10] Ang, D. (2023). The birth of a nation: Media and racial hate. *American Economic Review 113*(6), 1424–60.

[11] Asher, S., C. Farquharson, E. Iljazi, and P. Novosad (2024). Access to opportunity in the sciences: Evidence from the nobel laureates. Working Paper.

[12] Baird, M., N. Gahlawat, R. Hood, S. Lara, and P. Ko (2023). Pathways into STEM jobs. Technical report. White paper, November 2023.

[13] Bapuji, H., K. Gupta, S. Chrispal, and T. Roulet (2023). What managers everywhere must know about caste. *MIT Sloan Management Review Winter*. Accessed: 2024-04-22.

[14] Barg, K. (2013). The influence of students' social background and parental involvement on teachers' school track choices: Reasons and consequences. *European Sociological Review 29*(3), 565–579.

[15] Barrios-Fernández, A., C. Neilson, and S. Zimmerman (2024). Elite universities and the intergenerational transmission of human and social capital. *American Economic Review*. conditionally accepted.

[16] Bayly, S. (2008). *Caste, Society and Politics in India from the Eighteenth Century to the Modern Age*. Chicago: Cambridge University Press.

[17] Benson, A., S. Board, and M. Meyer-ter-Vehn (2024). Discrimination in hiring: Evidence from retail. *Review of Economic Studies 91*(4), 1956–1987.

[18] Benson, A., D. Li, and K. Shue (2019). Promotions and the Peter principle. *The Quarterly Journal of Economics 134*(4), 2085–2134.

[19] Benson, A., D. Li, and K. Shue (2023). "Potential" and the gender promotion gap. R&R, *American Economic Review*.

[20] Benson, A. and K. Shaw (2024). What do managers do? *Annual Review of Economics*. Forthcoming.

[21] Bergman, P. and A. Lerner (2025). Are recruiters better than a coin flip at judging resumes? Here's the data. https://interviewing.io/blog/are-recruiters-better-than-a-coin-flip-at-judging-resumes. Accessed: 2025-04-21.

[22] Bergman, P., D. Li, and L. R. Raymond (2024). Hiring as exploration. *Review of Economic Studies (forthcoming)*. NBER Working Paper 27736, revised June 2024.

[23] Beteille, A. (1965). *Caste, Class and Power: Changing Patterns of Stratification in a Tanjore Village*. University of California Press.

[24] Beteille, A. (1969). *Castes: Old and New, Essays in Social Structure and Social Stratification*. Asia Publishing House.

[25] Bhavnani, R. R. and A. Lee (2021). Does affirmative action worsen bureaucratic performance? Evidence from the indian administrative service. *American Journal of Political Science 65*(1), 5–20.

[26] Bhupatiraju, S., D. L. Chen, S. Joshi, and P. Neis (2024). Caste aside? Names, networks and justice in the courts of bihar, india. *European Journal of Empirical Legal Studies 1*(2), 151–178.

[27] Bisin, A., J. Rubin, A. Seror, and T. Verdier (2023). Culture, institutions & the long divergence. *Journal of Economic Growth 28*, 1–40.

[28] Bisin, A. and T. Verdier (1998). On the cultural transmission of preferences for social status. *Journal of Public Economics 70*, 75–97.

[29] Bisin, A. and T. Verdier (2010). The economics of cultural transmission and socialization. In *Handbook of Social Economics*, pp. 339–416. Elsevier.

[30] Bloom, N., R. Sadun, and J. Van Reenan (2012). The organization of firms across countries. *The Quarterly Journal of Economics 127*(4), 1663–1705.

[31] Bol, T., J. Witschge, H. G. Van de Werfhorst, and J. Dronkers (2014). Curricular tracking and central examinations: Counterbalancing the impact of social background on student achievement in 36 countries. *Social Forces 92*(4), 1545–1572.

[32] Borneman, M. J., S. R. Cooper, D. M. Klieger, N. R. Kuncel, and N. Brenckman (2007). Admission interviews: Not so nice to meet you. In *Alternative Predictors of Academic Performance, Symposium conducted at the Annual Meeting of the Society for Industrial and Organizational Psychology, New York, NY.*

[33] Bourdieu, P. (1986). The forms of capital. In J. G. Richardson (Ed.), *Handbook of Theory and Research for the Sociology of Education*, pp. 241–258. Greenwood Press.

[34] BusinessLine (2018). An Inclusive India Inc. Business Line, January 13, 2018.

[35] California v. Cisco Systems, Inc. (2022). Dept. of Fair Employment and Housing v. Cisco Systems, Inc. Docket No. H048962.

[36] Cappelli, P. (2019). Your approach to hiring is all wrong. *Harvard Business Review*. Accessed: 2024-08-15.

[37] Carroll, J. S., R. L. Wiener, D. Coates, J. Galegher, and J. J. Alibrio (1982). Evaluation, diagnosis, and prediction in parole decision making. *Law & Society Review 17*, 199–228.

[38] Castilla, E. J. (2015). Accounting for the gap: A firm study manipulating organizational accountability and transparency in pay decisions. *Organization Science 26*(2), 311–333.

[39] Castilla, E. J. and S. Benard (2010). The paradox of meritocracy in organizations. *Administrative Science Quarterly 55*(4), 543–576.

[40] Chakravartty, P. and A. Subramanian (2021). Why is Caste Inequality Still Legal in America? *The New York Times*, May 25, 2021.

[41] Chandra, A. and D. O. Staiger (2010). Identifying provider prejudice in healthcare. Technical Report 16382, National Bureau of Economic Research.

[42] Chandra, A. and D. O. Staiger (2020). Identifying sources of inefficiency in health care. *Quarterly Journal of Economics 135*(2), 785–843.

[43] Checchi, D. and H. G. van de Werfhorst (2018). Policies, skills and earnings: How educational inequality affects earnings inequality. *Socio-Economic Review 16*(1), 137–160.

[44] Chetty, R., D. J. Deming, and J. N. Friedman (2023). Diversifying society's leaders? The determinants and causal effects of admission to highly selective private colleges. NBER Working Paper 31492.

[45] Chow, R. M. and E. D. Knowles (2016). Taking race off the table: Agenda setting and support for color-blind public policy. *Personality and Social Psychology Bulletin 42*(1), 25–39.

[46] Clark, G. (2015). *The Son Also Rises: Surnames and the History of Social Mobility*. The Princeton Economic History of the Western World.

[47] Cooper, S. R., D. M. Klieger, M. J. Borneman, N. R. Kuncel, and H. Mikeworth (2007). Personal statements: Please keep them to yourself. In *Alternative Predictors of Academic Performance, Symposium conducted at the Annual Meeting of the Society for Industrial and Organizational Psychology, New York, NY.*

[48] Dana, J., R. Dawes, and N. Peterson (2013). Belief in the unstructured interview: The persistence of an illusion. *Judgement and Decision Making 8*(5), 512–520.

[49] Datta, S. (2017). *A History of the Indian University System: Emerging from the Shadows of the Past*. Palgrave Macmillan.

[50] Deshpande, A. (2011). *The Grammar of Caste: Economic Discrimination in Contemporary India*. Oxford: Oxford University Press.

[51] DeVaul, R., F. Jervey, J. Chappell, P. Caver, B. Short, and S. O'Keefe (1987). Medical school performance of initially rejected students. *Journal of the American Medical Association 257*, 47–51.

[52] Doleac, J. L. and B. Hansen (2020). The unintended consequences of "Ban the Box": Statistical discrimination and employment outcomes when criminal histories are hidden. *Journal of Labor Economics 38*(2), 321–374.

[53] Durlauf, S. N. (1999). The case "against" social capital. In P. Dasgupta and I. Serageldin (Eds.), *Social Capital: A Multifaceted Perspective*, pp. 84–104. Washington, D.C.: World Bank.

[54] Durlauf, S. N. (2006). Groups, social influences, and inequality: A memberships theory perspective on poverty traps. In S. Bowles, S. N. Durlauf, and K. Hoff (Eds.), *Poverty Traps*. Princeton: Princeton University Press.

[55] Durlauf, S. N. (2008). Memberships and inequality. In D. Castiglione, J. van Deth, and G. Wolleb (Eds.), *Handbook of Social Capital*. Oxford: Oxford University Press.

[56] Eames, T. (2024). Taryn versus taryn (she/her) versus taryn (they/them): A field experiment on pronoun disclosure and hiring discrimination. Working Paper.

[57] Eames, T. (2025). How does sex-based occupational discrimination affect nonbinary applicants? *AEA Papers and Proceedings*.

[58] Economic Advisory Council to the Prime Minister (2022). The State of Inequality in India Report. Government of India.

[59] Erickson, B. H. (1996). Culture, class, and connections. *American Journal of Sociology 102*(1), 217–251.

[60] Espinosa, M. and C. T. Stanton (2022, July). Training, communications patterns, and spillovers inside organizations. Working Paper 30224, National Bureau of Economic Research.

[61] Fenizia, A. (2022). Managers and productivity in the public sector. *Econometrica 90*(3), 1063–1084.

[62] Friedman, S. and D. Laurison (2020). *The Class Ceiling: Why it Pays to be Privileged?* Bristol University Press, Policy Press.

[63] Gibbons, R. and J. Roberts (2013). Economic theories of incentives in organizations. *The Handbook of Organizational Economics*.

[64] Gorman, E. H. (2005). Gender stereotypes, same-gender preferences, and organizational variation in the hiring of women: Evidence from law firms. *American Sociological Review 70*(4), 702–728.

[65] Gottfried, J. (2024, January). Americans' social media use. Technical report, Pew Research Center.

[66] Grusky, D. B. (2005). Foundations of a neo-durkheimian class analysis. In E. O. Wright (Ed.), *Approaches to Class Analysis*, pp. 51–81. Cambridge University Press.

[67] Hampole, M., F. Truffa, and A. Wong (2024). Peer effects and the gender gap in corporate leadership: Evidence from MBA students. *Quarterly Journal of Economics*. Revise and Resubmit.

[68] Highet, K. (2023). "RIP English": Race, Class and 'Good English' in India. *Journal of Linguistic Anthropology 33*(2), 184–201.

[69] Highhouse, S. (2008). Stubborn Reliance on Intuition and Subjectivity in Employee Selection. *Industrial and Organizational Psychology: Perspectives on Science and Practice 1*(3), 333–342.

[70] Hoffman, M., L. B. Kahn, and D. Li (2018). Discretion in Hiring. *The Quarterly Journal of Economics 133*(2), 765–800.

[71] Hoffman, M. and C. T. Stanton (2024). People, practices, and productivity: A review of new advances in personnel economics. NBER Working Paper 32849.

[72] Hoffman, M. and S. Tadelis (2021). People management skills, employee attrition, and manager rewards: An empirical analysis. *Journal of Political Economy 129*(1), 243–285.

[73] Human Rights Watch (2001). Caste discrimination: A global concern. HRW Briefing Paper. See ch. 2, p. 2: "Lower-caste communities are almost invariably indistinguishable in physical appearance from higher-caste communities".

[74] Jodhka, S. S. and J. Naudet (2019). *Mapping the Elite: Power, Privilege, and Inequality*. Oxford University Press.

[75] Jodhka, S. S. and K. Newman (2007). In the name of globalisation: Meritocracy, productivity and the hidden language of caste. *Economic and Political Weekly 42*(41), 4125–4132.

[76] Jodkha, S. S. (2017). *Caste in Contemporary India*. Routledge India.

[77] Kahneman, D. (2011). *Thinking, Fast and Slow*. New York: Farrar, Straus and Giroux.

[78] Kang, S. K., K. A. DeCelles, A. Tilcsik, and S. Jun (2016). Whitened resumes: Race and self-presentation in the labor market. *Administrative Science Quarterly 61*(3), 469–502.

[79] Kashyap, R. and F. Verkroost (2021). Analysing global professional gender gaps using LinkedIn advertising data. *EPJ Data Science 10*(39), 1–32.

[80] Khan, S. R. (2011). *Privilege: The Making of an Adolescent Elite at St. Paul's School*. Princeton, NJ: Princeton University Press.

[81] Kline, P., E. Rose, and C. Walters (2022). Systemic discrimination among large U.S. employers. *Quarterly Journal of Economics 137*(4), 1963–2036.

[82] Krishna, K. and V. C. Frisancho Robles (2015). Affirmative Action in Higher Education in India: Targeting, Catch Up, and Mismatch. *Higher Education*.

[83] Kuncel, N. R., M. Credé, and L. L. Thomas (2007). A meta-analysis of the predictive validity of the graduate management admission test (GMAT) and undergraduate grade point average (UGPA) for graduate student academic performance. *Academy of Management Learning & Education 6*(1), 51–68.

[84] Kuncel, N. R. and S. A. Hezlett (2007). Standardized tests predict graduate students' success. *Science 315*(5815), 1080–1081.

[85] Kuncel, N. R., D. M. Klieger, B. S. Connelly, and D. S. Ones (2013). Mechanical versus clinical data combination in selection and admissions decisions: A meta-analysis. *Journal of Applied Psychology 98*(6), 1060–1072.

[86] Kuncel, N. R., R. J. Kochevar, and D. S. Ones (2014). A meta-analysis of letters of recommendation in college and graduate admissions: Reasons for hope. *International Journal of Selection and Assessment 22*(1), 101–107.

[87] Lamont, M. (1992). *Money, Morals, and Manners: The Culture of the French and the American Upper-Middle Class*. Chicago: University of Chicago Press.

[88] Lamont, M. (2009). *How Professors Think: Inside the Curious World of Academic Judgment*. Cambridge, MA: Harvard University Press.

[89] Lamont, M. and A. Lareau (1988). Cultural capital: Allusions, gaps and glissandos in recent theoretical developments. *Sociological Theory 6*(2), 153–168.

[90] Laurison, D. and S. Friedman (2016). The class pay gap in higher professional and managerial occupations. *American Sociological Review 81*(4), 668–695.

[91] Laurison, D. and S. Friedman (2019). *The Class Ceiling: Why it Pays to be Privileged*. Bristol, UK: Policy Press.

[92] Lazear, E. P., K. L. Shaw, and C. T. Stanton (2015). The value of bosses. *Journal of Labor Economics 33*(4), 823–861.

[93] Lenski, G. E. (1954). Status crystallization: A non-vertical dimension of social status. *American Sociological Review 19*(4), 405–413.

[94] Levashina, J., C. J. Hartwell, F. P. Morgeson, and M. A. Campion (2014). The structured employment interview: Narrative and quantitative review of the research literature. *Personnel Psychology 67*(1), 241–293.

[95] Levon, E., D. Sharma, D. J. L. Watt, A. Cardoso, and Y. Ye (2021). Accent bias and perceptions of professional competence in england. *Journal of English Linguistics 49*(4), 355–388.

[96] Madheswaran, S. (2008). Is affirmative action policy for private sector necessary? *Indian Journal of Industrial Relations 44*(2), 164–184.

[97] Mamgain, R. P. (2019). *Formal Labour Market in Urban India: Job Search, Hiring Practices and Discrimination*. SAGE Publications Pvt. Ltd.

[98] Mamidi, P. (2011). Strategic manipulation of caste identity signals. Doctoral Dissertation, Oxford University.

[99] Martellini, P., T. Schoellman, and J. Sockin (2024, February). The global distribution of college graduate quality. *Journal of Political Economy 132*(2), 434–483.

[100] McDaniel, M. A., D. T. Whetzel, F. L. Schmidt, and S. D. Maurer (1994). The validity of employment interviews: A comprehensive review and meta-analysis. *Journal of Applied Psychology 79*(4), 599–616.

[101] Meehl, P. E. (1954). *Clinical versus Statistical Prediction: A Theoretical Analysis and a Review of the Evidence*. University of Minnesota Press.

[102] Meehl, P. E. (1986). Causes and effects of my disturbing little book. *Journal of Personality Assessment 50*(3), 370–375.

[103] Michelman, V., J. Price, and S. D. Zimmerman (2022). Old boys' clubs and upward mobility among the educational elite. *The Quarterly Journal of Economics 137*(2), 845–909.

[104] Ministry of Law and Justice, Government of India (2021). The constitution of India. Article 15 explicitly prohibits discrimination on grounds of religion, race, caste, sex, or place of birth.

[105] Minni, V. (2023). Making the invisible hand visible: Managers and the allocation of workers to jobs. R&R, *Quarterly Journal of Economics*.

[106] Minni, V. (2024). Global managers, local workers: Wage setting inside a multinational firm. *AEA Papers & Proceedings 114*, 586–91.

[107] Minni, V., K.-T. Nguyen, and H. Sarsons (2025). Managers and the cultural transmission of gender norms. Working Paper.

[108] Mishra, N. (2015). India and colorism: The finer nuances. *Washington University Global Studies Law Review 14*, 725–750.

[109] Murphy, S. C., D. M. Klieger, M. J. Borneman, and N. R. Kuncel (2009). The predictive power of personal statements in admissions: A meta-analysis and cautionary tale. *Educational and Psychological Measurement 69*(6), 967–982.

[110] Neckerman, K. M. and J. Kirschenman (1991). Hiring strategies, racial bias, and inner-city workers. *Social Problems 38*(4), 433–447.

[111] Oster, E. (2019). Unobservable selection and coefficient stability: Theory and evidence. *Journal of Business & Economic Statistics 37*(2), 187–204.

[112] Parameswaran, R. and K. Cardoza (2015). Melanin on the margins: Advertising and the cultural politics of fair/light/white beauty in india. *Journalism & Communication Monographs 11*(3), 213–274.

[113] Parmar, R. (2020). Transacting caste in modern times: Changing social identity through surnames in urban Gujarat. *Contemporary Voice of Dalit 12*(2), 220–233.

[114] Pearson, A. R., J. F. Dovidio, and S. L. Gaertner (2009). The nature of contemporary prejudice: Insights from aversive racism. *Social and Personality Psychology Compass 3*(3), 314–338.

[115] Pinker, S. (2021). *Rationality: What It Is, Why It Seems Scarce, Why It Matters*. New York: Viking.

[116] Puri, S. (2025, February). Capitalism with Indian characteristics. Stigler Center Lecture Series, University of Chicago Booth School of Business. Accessed: 2025-02-11.

[117] Quillian, L., J. J Lee, and M. Oliver (2020). Evidence from Field Experiments in Hiring Shows Substantial Additional Racial Discrimination after the Callback. *Social Forces 99*(2), 732–759.

[118] Reskin, B. (2012). The race discrimination system. *Annual Review of Sociology 38*, 17–35.

[119] Ridgeway, C. L. (2019). *Status: Why Is It Everywhere? Why Does It Matter?* New York, NY: Russell Sage Foundation.

[120] Rivera, L. A. (2012). Hiring as cultural matching: The case of elite professional service firms. *American Sociological Review 77*(6), 999–1022.

[121] Rivera, L. A. (2015). *Pedigree: How Elite Students Get Elite Jobs*. Princeton University Press.

[122] Rivera, L. A. and A. Tilcsik (2016). Class advantage, commitment penalty: The gendered effect of social class signals in an elite labor market. *American Sociological Review 81*(6), 1097–1131.

[123] Sacerdote, B., D. O. Staiger, and M. Tine (2025, January). How test optional policies in college admissions disproportionately harm high achieving applicants from disadvantaged backgrounds. Working Paper 33389, National Bureau of Economic Research.

[124] Sadun, R. (2023). CEOs and firm performance. NBER Reporter, ISSN 0276-119X, National Bureau of Economic Research (NBER), Cambridge, MA, Iss. 1, pp. 12-16.

[125] Sadun, R., N. Bloom, and J. Van Reenen (2017, September-October). Why do we undervalue competent management? Neither great leadership nor brilliant strategy matters without operational excellence. *Harvard Business Review 95*(5), 120–127. Winner of 59th Annual HBR McKinsey Award.

[126] Sadun, R., J. Tamayo, Jorge A. Ramirez, Diaz, and A. B. S. Neyra-Nazarrett (2024, April). Training within firms. Working Paper 33670, National Bureau of Economic Research.

[127] Sarsons, H. (2017). Recognition for group work: Gender differences in academia. *AEA Papers and Proceedings 107*(5), 141–145.

[128] Sarsons, H. (2023). Interpreting Signals in the Labor Market: Evidence from Medical Referrals. R&R, *Review of Economic Studies*.

[129] Sarsons, H., K. Gërxhani, E. Reuben, and A. Schram (2021). Gender differences in recognition for group work. *Journal of Political Economy 129*(1), 101–147.

[130] Sarsons, H., J. Hjort, and X. Li (2022). Across-Country Wage Compression in Multinationals. R&R, *American Economic Review*.

[131] Savage, M. (2015). *Social Class in the 21st Century*. Pelican Books.

[132] Savage, M., F. Devine, N. Cunningham, M. Taylor, Y. Li, J. Hjellbrekke, B. Le Roux, S. Friedman, and A. Miles (2013). A new model of social class? Findings from the bbc's great british class survey experiment. *Sociology 47*(2), 219–250.

[133] Sebring, J. M. (1969). Caste indicators and caste identification of strangers. *Human Organization 28*(3), 199–207.

[134] Sharma, D. (2017). English in India. In J. Filipović and S. C. Deshors (Eds.), *World Englishes: New Theoretical and Methodological Considerations*, pp. 311–332. De Gruyter Mouton.

[135] Sönmez, T. and M. B. Yenmez (2024). Constitutional implementation of affirmative action policies in india. *The Review of Economics and Statistics*, 1–45. Published online 25 November 2024.

[136] Stansbury, A. (2024). Social class and career progression in elite professions. *American Economic Journal: Applied Economics Forthcoming*.

[137] Stansbury, A. and K. Rodriguez (2024a). The class gap in career progression: Evidence from US academia. R&R, *Econometrica*.

[138] Stansbury, A. and K. Rodriguez (2024b, December). The class gap in career progression: Evidence from US academia. Presented at the NBER Organizational Economics Working Group, Fall 2024.

[139] Stansbury, A. and R. Schultz (2023). The economics profession's socioeconomic diversity problem. *Journal of Economic Perspectives 37*(4), 207–230.

[140] Stevens, M. P. (2009). *Creating a Class: College Admissions and the Education of Elites*. Harvard University Press.

[141] Subramanian, S. (2019). India's elite colleges and the employment divide. *Economic and Political Weekly 54*(2), 42–48.

[142] Tambe, P., L. Hitt, D. Rock, and E. Brynjolfsson (2021). Digital capital and superstar firms. Technical report, Brookings Institution, Hutchins Center Working Paper #73.

[143] Thorat, S. and P. Attewell (2007, October). The legacy of social exclusion: A correspondence study of job discrimination in india. *Economic and Political Weekly 42*(41), 4141–4145.

[144] Todorov, A. (2017). *Face Value: The Irresistible Influence of First Impressions*. Princeton, NJ: Princeton University Press.

[145] Turco, C. J. (2010). Cultural foundations of tokenism: Evidence from the leveraged buyout industry. *American Sociological Review 75*(6), 894–913.

[146] Wiltshire, C. R. (2020). *Uniformity and Variability in the Indian English Accent*. Cambridge University Press.

[147] Yenmez, M. B. and T. Sönmez (2022). Affirmative action in India via vertical, horizontal, and overlapping reservations. *Econometrica 90*(3), 1143–1176.

[148] Zhang, L., S. Wang, and N. Wilmers (2025). Remote work and careers. SSRN working paper.

[149] Zhang, S. and S. Vucetic (2016). Sampling bias in Linkedin: A case study. In *Proceedings of the 25th International Conference Companion on World Wide Web*, pp. 145–146.

[150] Zimmerman, S. D. (2019). Elite colleges and upward mobility to top jobs and top incomes. *American Economic Review 109*(1), 1–47.

ONLINE APPENDIX

A Caste as Social Status Under Affirmative Action

India's caste system provides a uniquely precise measure of social status, offering clear administrative categories for analyzing class-based inequalities. Affirmative action policies formally recognize caste as a comprehensive marker of historical advantage and disadvantage, codifying how caste membership systematically shapes access to resources, education, and elite employment opportunities. This administrative clarity enables identification of status-based discrimination mechanisms that remain difficult to isolate in Western class analyses.

Jati, Varna, and Modern Caste Classifications. India's modern caste system originates from historical occupational hierarchies. At its most detailed level, the system comprises thousands of hereditary occupational groups known as *Jatis*, historically linked to trades such as merchants, priests, warriors, and laborers. These *Jatis* are grouped into broader categories called *Varnas*: Brahmins (priests), Kshatriyas (warriors), Vaishyas (merchants), Shudras (laborers), and historically marginalized "untouchables" (Dalits). The government maintains official lists classifying communities as "Scheduled Castes" (SCs), "Scheduled Tribes" (STs), and "Other Backward Classes" (OBCs)—groups defined as socially and educationally disadvantaged (Puri, 2025).

This paper adopts two broad categories aligned with India's affirmative action framework: advantaged ("upper") and disadvantaged ("lower") castes. By disadvantaged castes, I refer explicitly to Scheduled Castes, Scheduled Tribes, and Other Backward Classes. This administrative categorization recognizes caste as fundamentally intertwined with economic marginalization and social status, paralleling Western concepts of class (Puri, 2025).

Historical Development of Caste-based Affirmative Action Policies. Affirmative action for disadvantaged communities emerged during British colonial rule, beginning with the Government of India Act of 1919's reserved legislative seats for "depressed classes." The 1935 Government of India Act officially designated these groups as Scheduled Castes. Post-independence, India's 1949 Constitution institutionalized quota-based affirmative action in legislatures, public higher education, and government employment for Scheduled Castes and Scheduled Tribes. Following the Mandal Commission's recommendations in the 1980s, quotas expanded to include Other Backward Classes, raising reservation levels to nearly 50% (Bayly, 2008).

Caste as Proxy for Social Status. Although caste originated as a hereditary hierarchy, Indian policymakers now define caste categories as comprehensive indicators of socioeconomic "backwardness"—encompassing economic, educational, and social disadvantages. This approach recognizes caste as encapsulating multiple dimensions of social status: financial resources, educational credentials, social networks, and cultural capital (cf. Bourdieu, 1986). The Mandal Commission merged lineage-based identities with socioeconomic status, establishing caste as a consolidated marker of social status.

Limited Scope of Affirmative Action in the Private Sector. India's affirmative action policies do not extend to private-sector employment, including the multinational corporations studied here. Private-sector roles typically offer higher wages and career growth but lack reservation mandates, disproportionately excluding disadvantaged castes (Subramanian, 2019). This parallels Western contexts, where firms emphasize racial and gender diversity but rarely address socioeconomic diversity. Fewer than 10% of large US firms track socioeconomic diversity (Stansbury and Rodriguez, 2024b), while fewer than 3% of India's top firms maintain internal caste data (BusinessLine, 2018). Although a few multinational corporations have recently begun addressing caste discrimination (Chakravartty and Subramanian, 2021; Bapuji et al., 2023), these changes occurred in late 2022 or 2023, falling outside the timeframe of the data analyzed in this paper (Section 4.4).

B Representativeness of the Job Fair at the Elite College

The studied on-campus job fair provides a highly representative window into how India's elite-college graduates transition into elite private-sector jobs, supported by four key observations:

1. **Near-identical fairs across elite institutions.** Since their inception in the 1950s, India's premier institutions—including the Indian Institutes of Technology (IITs), Indian Institutes of Management (IIMs), and leading Delhi University colleges—have consciously modeled their practices after prominent Western universities such as MIT, Stanford, and Harvard. Consequently, these institutions share strikingly similar academic calendars, faculty-to-student ratios, curricula, and critically, their standardized placement office procedures. Today, the job fairs conducted by all 23 IITs, 20 IIMs, and top-tier Delhi University colleges closely mirror one another, featuring virtually identical recruitment stages, timelines, and employer pools (e.g., see Central Placement Cell, Delhi University).

2. **Graduates predominantly flow into elite private-sector jobs.** Graduates from India's elite colleges overwhelmingly pursue high-paying entry-level roles in multinational corporations and prominent Indian firms. Public-sector employment remains a distinctly secondary choice, reflecting a long-standing preference among top-tier college graduates for careers in consulting, technology, finance, and multinational management (e.g., Subramanian, 2019; Jodhka and Naudet, 2019).

3. **Elite firms recruit almost exclusively from elite campuses.** The multinational firms and prestigious Indian companies studied here systematically target the same elite Indian institutions year after year. Collectively, these recruiters visit around 30–40 top-ranked colleges annually, each institution typically enrolling more than 1,000 students. This concentrated recruiting strategy ensures a highly consistent candidate pool across years, facilitating a standardized recruitment process and comparable outcomes nationwide. Prior studies further corroborate that elite firms' hiring remains

strongly concentrated within these elite campuses (e.g., Subramanian, 2019; Jodhka and Naudet, 2019).

4. **Minimal student withdrawal from the job fair.** Participation in the on-campus job fair is near-universal among graduating students. Less than 5% of final-year students voluntarily deregister from the formal placement process, a proportion consistent across similar elite institutions (Subramanian, 2019). Such low withdrawal rates underscore the central role of these standardized recruitment fairs in determining employment outcomes for India's most academically competitive graduates.

Collectively, these facts demonstrate that the focal job fair closely mirrors elite hiring pathways across India's top universities.

C Benchmarking Employment-Platform Coverage: External Evidence

Large-scale validations suggest LinkedIn already reaches roughly four-fifths of U.S. bachelor's-degree holders overall (Zhang et al., 2025). Coverage is markedly higher in technical fields. LinkedIn's own "Pathways into STEM Employment" study reports that 92.7% of 2022 STEM master's graduates (and hence virtually all STEM users with active profiles) interacted with at least one STEM job on the site, a 12-point lift over the bachelor-level average (Baird et al., 2023). Cross-country work buttresses this. In their 92-nation ad-audience study, Kashyap and Verkroost (2021) show that for India the 'Information & Communication' sector accounts for 32% of female and 27% of male LinkedIn users but only about 1% of the professional labour force logged by the ILO, a more than 25-percentage-point enrichment that mirrors the global STEM premium they document.

Firm-level evidence points the same way: Tambe et al. (2021) find IT employees make up 5% of LinkedIn's U.S. workforce versus 2% in BLS data, implying near-census platform coverage of U.S. IT professionals once Pew's nearly 35% overall adoption rate is factored in (Gottfried, 2024). Even outside STEM, MBA follow-ups routinely match 85–90% of graduates from many cohorts to LinkedIn profiles (Hampole et al., 2024). And elite STEM-focused universities often comfortably clear 90%: Stanford's official count lists "more than 220,000 living alumni," yet LinkedIn shows about 314,000 profiles, which after de-duplication and removing mis-tags using standard methods (e.g., Zhang and Vucetic, 2016) suggest a more than 95% hit-rate. Given that our college's talent pool is comparable to many world-leading institutions (Martellini et al., 2024), a well over 90% in-sample match rate at our STEM-heavy campus is entirely plausible. Indeed, the third-party matched dataset achieves a nearly 98% match rate for graduates, aided by systematic institutional identifiers that follow predictable patterns within each major-degree-year, which greatly boosts match precision.

D Timeline and Screening Rounds

Every job in the sample follows the same four-stage sequence, eliminating candidates at each step:

1. **Application screening (September).** Recruiters conduct an initial screening, primarily based on GPA, major, degree, and standout résumé items (e.g., internships, competitions).

2. **Written aptitude tests (September–November).** Online technical assessments vary by role: coding exercises for tech/manufacturing jobs, and case-based tests for consulting positions. Firms advertising multiple posts usually administer one common test (e.g., Nvidia uses a single exam for all advertised profiles).

3. **Group debates (September–November).** A non-technical evaluation of soft skills:

   (a) Groups of 30–35 students, split into two teams, debate a general prompt (e.g., "Is social media a boon or a bane?"), while recruiters observe communication, teamwork, and confidence.

   (b) Firms often conduct a single group debate for all advertised roles—for example, Google may simultaneously assess applicants for program manager and software engineer positions.

4. **Personal interviews (December–January).** Final 30-minute "fit" conversations:

   (a) Interviews occur simultaneously in multiple parallel batches, each handled by different members of the recruiting team.

   (b) As with earlier rounds, a single interview panel commonly covers multiple job titles (e.g., Uber may interview data scientist and front-end engineer candidates within the same track).

   (c) Because more than 98% of roles are based in India, interviewers are almost certainly Indian (Section 4.6).

E Qualitative evidence supporting caste revelation at interviews

Qualitative evidence further clarifies why caste becomes most legible to recruiters during personal interviews. As discussed quantitatively in Section 5.6.1, surnames alone rarely disclose caste in urban Indian settings; instead, status cues emerge primarily through richer social interactions. Group debates provide limited scope for caste inference: candidates debate standardized prompts in large groups with minimal individualized questioning, and neither physical appearance nor linguistic cues reliably indicate caste. Anthropological research consistently finds phenotypic traits like skin color are unreliable caste indicators among urban, educated populations (Sebring, 1969; Mishra, 2015; Parameswaran and Cardoza, 2015). Human Rights Watch concurs, stating that "*Lower-caste communities are almost*

invariably indistinguishable in physical appearance from higher-caste communities" (Human Rights Watch, 2001, p. 2). Similarly, elite-college varieties of Indian English reflect regional or educational backgrounds rather than caste, as English-medium schooling has expanded significantly beyond upper castes due to affirmative action (Wiltshire, 2020; Sharma, 2017; Highet, 2023). Thus, group debates offer employers minimal caste-identifying information.

In contrast, personal interviews substantially shift the evaluative context. They amplify earlier signals such as schooling pedigree, resume details, and debate performances through deeper discussions. Interviews also elicit richer background information—personal interests, leisure activities, niche extracurricular involvement, and peer networks (e.g., Rivera, 2012, 2015), which are strongly correlated with caste-linked social and cultural capital, such as networks, resources, and cultural tastes (cf. Bourdieu, 1986). Such characteristics are well-documented caste correlates in elite Indian contexts (cf. Jodhka and Newman, 2007; Deshpande, 2011; Jodkha, 2017). Additionally, caste identities frequently become salient through social interactions in organizational settings, as managers learn subtle cues from conversations about personal background, residential origins, familial occupations, or shared acquaintances (Bapuji et al., 2023). Integrating these nuanced signals with earlier observations enables recruiters to more accurately infer caste, aligning precisely with the timing and intensity of caste penalties documented quantitatively (Section 5.6).

F Extending the Simple Model of Late-Stage Discrimination with Imperfect Information

In Section 5.7.1, I assumed that employers observe each candidate's true screening-round performance perfectly (i.e., $M_{s,i,g} = A_{i,g}$ at round s). Here, I show that *all* key insights—particularly how variation in job selectivity identifies discrimination versus ability deficits—remain valid even when employers observe *noisy* performance signals.

Setup. As in Section 5.7.1, there are two groups (castes) $g \in \{0, 1\}$, a continuum of jobs $j \in \mathcal{J}$, and four sequential screening rounds $\{1, 2, 3, 4\}$. Rounds $1, 2, 3$ constitute pre-interview filtering (e.g., applications, written tests, group debates), while round 4 is the final personal interview. Each candidate i from group g has a *latent ability* $A_{i,g}$. Passing round s requires

$$\mathbb{E}\big[M_{s,i,g} \,\big|\mathcal{I}_{s-1}\big] \;\geq\; \tau_{s,j}, \tag{OA.1}$$

where $\mathcal{I}_{s-1}$ denotes the employer's information after observing signals in rounds $1, \ldots, s-1$. Thresholds $\tau_{s,j}$ weakly increase with each stage and may vary across jobs.

Connection to Assumption 1. Throughout this appendix we maintain the common-pool assumption: each job still draws candidates from the same latent-ability distributions F_g. Screening noise $\eta_{s,i,g}$ is therefore orthogonal to job-level variation in $\tau_{s,j}$, exactly as in the main text.

Noisy Signals per Round. At each stage s, the employer observes

$$M_{s,i,g} \;=\; A_{i,g} \;+\; \eta_{s,i,g},$$

where $\eta_{s,i,g}$ is a zero-mean noise term. We relax the assumption of independence across candidates and rounds, requiring only that:

(i) $\mathbb{E}[\eta_{s,i,g}] = 0$ for all s, i, g, and

(ii) $\eta_{s,i,1} \stackrel{d}{=} \eta_{s,i,0}$, so that for any subset of candidates (including those who have passed earlier rounds), there is no systematic caste-level difference in noise distributions that favors one caste.

Thus, while noise might correlate across rounds or candidates (e.g., candidate performance may be interdependent across stages), it never introduces systematic group-level biases. Under these conditions, persistent group differences cannot result from noise alone.

Updating Beliefs and Passing Rounds. At each stage s, the employer updates belief about $A_{i,g}$ given $\mathcal{I}_s \equiv \{M_{1,i,g}, \ldots, M_{s,i,g}\}$. Let

$$\mathbb{E}[A_{i,g} \mid \mathcal{I}_s] \quad \text{and} \quad Var(A_{i,g} \mid \mathcal{I}_s)$$

denote the posterior mean and variance of $A_{i,g}$. With each successive stage, beliefs about candidate ability become more precise, and the passing condition can be evaluated using $\mathbb{E}[M_{s,i,g} \mid \mathcal{I}_{s-1}]$.

Final Interview Thresholds and Caste Revelation. As before, let $\tau_{4,j}^{(g)}$ be the *group-specific* final-round threshold at job j. If caste becomes salient at the interview for group 1, a higher threshold ($\tau_{4,j}^{(1)} > \tau_{4,j}^{(0)}$) indicates discrimination. The goal is to demonstrate how variation in thresholds across jobs disentangles true ability differences from discrimination—even with noisy signals.

F.0.1 No Early Caste Penalty Implies No Large Right-Tail Deficit Before Interviews

A substantial right-tail ability deficit for group 1 would yield a particularly large gap in pass rates at early stages for highly selective jobs (with high $\tau_{s,j}$). Formally:

Proposition 3 (No Early Gap $\implies$ No Right-Tail Deficit Under Imperfect Information)**.** Consider any pre-interview stage $s < 4$. Under Assumption 1 (common-pool ability distributions), If group 1 truly has fewer high-ability candidates in the top tail, then for sufficiently high $\tau_{s,j}$,

$$P\big(\text{Pass}_{s,i,j} \mid g = 1, \mathcal{I}_{s-1}\big) \;<\; P\big(\text{Pass}_{s,i,j} \mid g = 0, \mathcal{I}_{s-1}\big).$$

Therefore, empirically observing no such gap—even in highly selective jobs—rules out a meaningful right-tail deficit for group 1 before the interview stage.

Sketch. Under Assumption 1, all jobs sample from the same F_g. With noisy signals, passing round s depends on the posterior mean $\mathbb{E}[M_{s,i,g} \mid \mathcal{I}_{s-1}]$. For high $\tau_{s,j}$, employers require a sufficiently high $\mathbb{E}[A_{i,g} \mid \mathcal{I}_{s-1}]$. If group 1 has fewer high-ability individuals, as $\tau_{s,j}$ rises, the pass rate for group 1 falls faster than group 0. This mirrors the perfect-information

argument in Proposition 1 and further relies on the *regularly decaying tail* assumption (e.g., $P(M_{s,i,g} \geq m) \sim C_g\, m^{-\alpha}$ for large m), which ensures that higher cutoffs intensify right-tail gaps.

Because noise is symmetric across groups, it cannot conceal large latent ability gaps at high cutoffs. Thus, the absence of early-stage gaps even at very selective jobs implies no meaningful right-tail deficit for group 1. □

F.0.2 Stable Promotion Gaps Across Selectivity Still Imply Discrimination

Post-hire performance Y_i depends on latent ability $A_{i,g}$. Promotion requires $Y_i \geq \kappa_j$. Final-round passing demands $\mathbb{E}[M_{4,i,g} \mid \mathcal{I}_3] \geq \tau_{4,j}^{(g)}$. Suppose top $(1-q)$ fraction have cutoffs $\tau_{4,j}^{q}$. If both groups face identical thresholds, their hired performance converges, and the promotion gap vanishes. However, if, $\tau_{4,j}^{(1)} = \tau_{4,j}^{(0)} + \delta$ ($\delta > 0$), group 1's marginal hires possess higher posterior means, producing a stable advantage in realized performance.

Proposition 4 (Stable Promotion Gaps Under Imperfect Information Identify Discrimination)**.** Under Assumption 1 and the regularly decaying tail condition used in Proposition 2 (i.e., for large m, $P(M_{4,i,g} \geq m) \sim C_g\, m^{-\alpha}$ with $\alpha > 1$), if group 1 faces a threshold higher by a constant $\delta > 0$, then average hired performance for group 1 remains about δ units higher than group 0, even in highly selective jobs. Hence, persistent promotion-rate gaps at high selectivity levels reject an unobserved-ability explanation in favor of discrimination.

Sketch. The argument parallels Proposition 2. Under imperfect information, stage-4 passage requires $\mathbb{E}[M_{4,i,g} \mid \mathcal{I}_3] \geq \tau_{4,j}^{(g)}$. If $\tau_{4,j}^{(1)} = \tau_{4,j}^{(0)} + \delta$, typical group 1 hires satisfy:

$$\mathbb{E}[A_{i,1} \mid \mathcal{I}_3] \;\approx\; \tau_{4,j}^{(0)} + \delta \;-\; \mathbb{E}[\eta_{4,i,1} \mid \mathcal{I}_3].$$

Because noise is identically distributed across groups (mean-zero, no systematic bias), raising group 1's threshold by δ effectively shifts their final hires toward higher regions of the ability distribution. At high selectivity, each group's hired candidates converge around their respective thresholds. Since group 1's threshold is higher by δ, their average hired performance remains about δ above group 0's. Conversely, identical thresholds shrink the performance gap to zero. Regular tail decay ensures marginal hires dominate outcomes at high cutoffs, maintaining a stable δ-sized gap under discrimination.

□

Key Takeaway. Allowing for imperfect information preserves key identifying variations. The absence of pre-interview deficits at high selectivity, combined with stable or increasing performance gaps post-hire, strongly indicates stricter interview thresholds imposed on disadvantaged candidates after caste revelation, rather than hidden ability shortfalls.

G Cultural Fit in Final-Round Interviews at Elite Multinationals

Final-round interviews at elite multinational firms often hinge less on technical ability than on interviewers' discretionary judgements of *cultural fit*. Building on the concept of *cultural matching* (Rivera, 2012), evaluators favour candidates whose background, tastes, and self-presentation feel "similar to us.'' In Rivera's seminal study of U.S. professional service firms, more than half of interviewers ranked cultural fit as *the single most important* criterion in the final round, ahead of analytical skill or academic pedigree. A law-firm partner captured the prevailing logic: *"In our new associates we are, first and foremost, looking for cultural compatibility. Someone who will . . . fit in''* (Rivera, 2012, p. 1006).

"Would I Want to Be Stuck in an Airport With Them?"

Interviewers often invoke the *airport test*: would I enjoy being stranded with this candidate for hours? A senior banker explained, *"A lot of this job is attitude, not aptitude . . . So you can be the smartest guy ever, but I don't care. I need to be comfortable working everyday with you, then getting stuck in an airport with you, and then going for a beer after. You need chemistry. Not only that the person is smart, but that you like him.''* (Rivera, 2012, p. 1008). Such remarks show how late-stage evaluation often becomes an audition of social rapport rather than pure competence (Lamont, 2009).

To locate common ground quickly, interviewers scour the résumé for shared experiences—college sports, study-abroad destinations, niche hobbies—and open the interview with informal chat. If a spark appears, enthusiasm often translates into advocacy in the hiring committee. Conversely, when overlap is missing, evaluators may withdraw support even when the candidate excels technically (Rivera, 2012, pp. 1013–1019).

Class-Coded Signals

The traits that signal fit are deeply entwined with social class and cultural capital (Bourdieu, 1986). Playing squash or lacrosse can mark a candidate as privileged in one firm and "too snooty'' in another, depending on the organization's self-image. Rivera documents a consulting director who excluded a technically strong candidate because *"we're more rough and tumble''* (Rivera, 2012, p. 1009). Similar dynamics appear in other sectors: Turco's (2010) ethnography of a leveraged-buyout firm showed women struggling to meet an aggressive, sport-centred masculine ideal, while Erickson (1996) finds that shared leisure pursuits foster stronger workplace ties and, ultimately, advancement.

Crucially, professionalism and client-readiness are evaluated separately from fit. As one consultant noted, *"You need someone who speaks in a way that earns your trust, who presents their opinion respectfully but also convincingly . . . But in terms of "fit," it's someone that we want on our case team . . . You want someone that makes you feel comfortable, that you enjoy hanging out with, can maintain a cool head when times are tough and make tough times kind of fun''* (Rivera, 2012, p. 1007). Recruiters believed that while "polish," "presence," and "presentation style" could be taught, fit cannot. Because the finalist pool is already

academically and technically elite, these subtle cultural cues decisively influence the offer decision.

Implications for Caste

Our quantitative results show that caste penalties emerge almost exclusively at this final interview stage. The qualitative literature helps explain why. If lower-caste candidates have had less exposure to elite cultural codes—speech patterns, leisure interests, informal networks—they may appear a weaker "fit'' despite equal credentials. Interviewers' reliance on self-referential judgements thereby re-inscribes social-status hierarchies under the guise of meritocracy (Neckerman and Kirschenman, 1991; Deshpande, 2011). As Rivera (2012) concludes, employers often *"constructed merit in a manner that validated their own strengths and experiences and perceived similar candidates as better applicants''* (p. 1012), reproducing privilege while believing they are selecting the "best'' candidate.

H Tables and Figures

Table OA.1: **Caste Differences in Mean Salary Among Graduates**

	Dependent variable:			
	log (Salary)			
	Without Controls	With Controls	Client Facing	Non-Client Facing
	(1)	(2)	(3)	(4)
Disadv. Caste	-0.231^{***}	-0.103^{***}	-0.107^{***}	-0.089^{***}
	(0.020)	(0.017)	(0.027)	(0.022)
Disadv. Caste (Oster)		-0.076^{***}	-0.095^{***}	-0.072^{***}
		(0.019)	(0.026)	(0.022)
Constant	11.235^{***}	10.863^{***}	10.883^{***}	10.848^{***}
	(0.012)	(0.040)	(0.057)	(0.052)
Observations	2,527	2,527	849	1,678
Adjusted R^2	0.049	0.412	0.549	0.382
Year FE	✓	✓	✓	✓
Other Controls	✓	✓	✓	✓
Job FE				

Notes: This table reports estimates from earnings regressions for graduates who obtained jobs through the college's job fair. The dependent variable is log earnings. Controls include pre-college skills, GPA, prior labor market experience, major, and degree. Regressions include year fixed effects. Column (1) shows estimates without controls, Column (2) adds controls, and Columns (3)-(4) show separate estimates for client-facing and non-client-facing jobs with controls. Job fixed effects are omitted as salaries are identical within jobs regardless of caste. Bias-adjusted coefficients follow Oster (2019), setting $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$ (assuming unobservables match observables in selection intensity and allowing total explanatory power to exceed observed R^2 by 30%). $^{*}p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

Table OA.2: **Caste Differences in GPA Between Those Omitted From the Earnings Regression Versus Those Who Are Not**

	Dependent variable:	
	log (GPA)	
	(1)	(2)
Disadv. Caste	-0.033^{***}	-0.032^{***}
	(0.004)	(0.004)
Omitted	-0.167^{***}	-0.158^{***}
	(0.005)	(0.005)
Omitted * Disadv. Caste	-0.071^{***}	-0.070^{***}
	(0.007)	(0.007)
Constant	2.147^{***}	2.127^{***}
	(0.004)	(0.008)
Observations	4,164	4,164
Adjusted R^2	0.546	0.559
Year FE	✓	✓
Major		✓
Degree		✓
Job FE		
Note:	*p<0.1; **p<0.05; ***p<0.01	

Notes: This table shows caste differences in GPA between students included in and omitted from the earnings regression. "Omitted" comprises students who either deregistered from placement or received no job offers, typically pursuing aftermarket opportunities, higher education, or entrepreneurship (Sections 3.2 and 5.4). Column (1) includes only year fixed effects; Column (2) adds major and degree controls. Job fixed effects are excluded as they cannot be separately identified from the "omitted" coefficient for students without jobs. $^{*}p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

Table OA.3: **Caste Differences in GPA At Each Job Search Stage**

	Dependent variable:					
	log (GPA)					
	App	Test	Debate	Interview	Offer	Choice
	(1)	(2)	(3)	(4)	(5)	(6)
Disadv. Caste	-0.111^{***}	-0.049^{***}	-0.039^{***}	-0.019^{***}	0.010^{***}	0.009^{***}
	(0.005)	(0.003)	(0.002)	(0.001)	(0.002)	(0.002)
Disadv. Caste (Oster)	-0.104^{***}	-0.035^{***}	-0.034^{***}	-0.014^{***}	0.015^{***}	0.014^{***}
	(0.006)	(0.003)	(0.002)	(0.002)	(0.002)	(0.002)
Constant	2.057^{***}	2.131^{***}	2.163^{***}	2.212^{***}	2.214^{***}	2.199^{***}
	(0.028)	(0.024)	(0.023)	(0.006)	(0.006)	(0.028)
Mean GPA	7.776	7.967	8.17	8.26	8.419	8.419
Observations	4,005	3,618	3,266	3,031	2,527	2,527
Adjusted R^2	0.264	0.726	0.745	0.801	0.822	0.822
Year FE	✓	✓	✓	✓	✓	✓
Job FE	✓	✓	✓	✓	✓	✓
Major	✓	✓	✓	✓	✓	✓
Degree	✓	✓	✓	✓	✓	✓
Note:						*p<0.1; **p<0.05; ***p<0.01

Notes: This table reports within-job GPA differences across castes at each job search stage. Regressions control for major, degree as well as job and year fixed effects. Bias-adjusted coefficients follow Oster (2019), setting $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$ (assuming unobservables match observables in selection intensity and allowing total explanatory power to exceed observed R^2 by 30%). $^{*}p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

Table OA.4: **Caste Differences in Mean Salary at Each Job Search Stage in the Full Sample**

	Dependent variable:					
	log (Salary)					
	App	Test	Debate	Interview	Offer	Choice
	(1)	(2)	(3)	(4)	(5)	(6)
Disadv. Caste	0.004	−0.009	−0.006	−0.018*	−0.105***	−0.103***
	(0.005)	(0.007)	(0.007)	(0.010)	(0.013)	(0.017)
Disadv. Caste (Oster)	0.005	0.010	0.011	0.012	−0.079***	−0.076***
	(0.005)	(0.008)	(0.008)	(0.011)	(0.014)	(0.019)
Constant	11.032***	10.904***	10.848***	10.758***	10.672***	10.863***
	(0.011)	(0.015)	(0.015)	(0.023)	(0.030)	(0.040)
Observations	4,005	3,618	3,266	3,031	2,527	2,527
Adjusted R^2	0.578	0.491	0.547	0.549	0.488	0.412
Year FE	✓	✓	✓	✓	✓	✓
Other Controls	✓	✓	✓	✓	✓	✓
Job FE						

Notes: This table reports earnings regressions for students still in contention at each job search stage. Each column represents a different stage: (1) application, (2) written test, and subsequent rounds. The dependent variable is log earnings. All regressions control for pre-college skills, GPA, prior labor market experience, major, and degree. Regressions include year fixed effects. Job fixed effects are omitted as salaries are identical within jobs regardless of caste. Bias-adjusted coefficients follow Oster (2019), setting $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$ (assuming unobservables match observables in selection intensity and allowing total explanatory power to exceed observed R^2 by 30%). $^*p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

Table OA.5: **Caste Differences in the Success Rate at Each Job Search Stage Among Client-Facing Jobs**

	Success App (1)	Success Test (2)	Success Debate (3)	Success Interview (4)
Disadv. Caste (OLS)	-0.0066***	-0.0007	-0.0136***	-0.3062***
	(0.0022)	(0.0030)	(0.0046)	(0.0081)
Disadv. Caste (Oster)	0.002	0.006	-0.015***	-0.314***
	(0.002)	(0.004)	(0.005)	(0.009)
Mean Success Rate	0.5770	0.7860	0.2590	0.2600
Observations	117,770	67,927	53,422	13,842
Adjusted R^2	0.61588	0.37467	0.01189	0.24298
Year FE	✓	✓	✓	✓
Job FE	✓	✓	✓	✓
Other Controls	✓	✓	✓	✓

Notes: This table reports caste differences in success rates at each hiring stage for client-facing jobs. The unit of observation is student-job-year, conditional on remaining in contention at each stage. All regressions include controls as well as job and year fixed effects. When firms conduct common screening rounds for multiple positions, each position is treated as a separate job. Bias-adjusted coefficients follow Oster (2019), setting $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$ (assuming unobservables match observables in selection intensity and allowing total explanatory power to exceed observed R^2 by 30%). $^*p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

Table OA.6: **Caste Differences in the Success Rate at Each Job Search Stage Among Non-Client-Facing Jobs**

	Success App (1)	Success Test (2)	Success Debate (3)	Success Interview (4)
Disadv. Caste (OLS)	0.0024	0.0011	0.0070**	-0.2867***
	(0.0018)	(0.0022)	(0.0034)	(0.0060)
Disadv. Caste (Oster)	0.008***	0.014***	0.007**	-0.286***
	(0.002)	(0.002)	(0.003)	(0.007)
Mean Success Rate	0.6140	0.8250	0.2580	0.2620
Observations	187,771	115,223	95,067	24,536
Adjusted R^2	0.60783	0.32634	0.02270	0.24723
Year FE	✓	✓	✓	✓
Job FE	✓	✓	✓	✓
Other Controls	✓	✓	✓	✓

Notes: This table reports caste differences in success rates at each hiring stage for non-client-facing jobs. The unit of observation is student-job-year, conditional on remaining in contention at each stage. All regressions include controls as well as job and year fixed effects. When firms conduct common screening rounds for multiple positions, each position is treated as a separate job. Bias-adjusted coefficients follow Oster (2019), setting $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$ (assuming unobservables match observables in selection intensity and allowing total explanatory power to exceed observed R^2 by 30%). $^*p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

Table OA.7: **Caste Differences in Mean Salary at Each Job Search Stage in the Identifiable Sample**

	Dependent variable:					
	log (Salary)					
	App	Test	Debate	Interview	Offer	Choice
	(1)	(2)	(3)	(4)	(5)	(6)
Disadv. Caste (OLS)	0.012	-0.068^{***}	-0.034^{*}	-0.105^{***}	-0.146^{***}	-0.107^{**}
	(0.014)	(0.019)	(0.018)	(0.027)	(0.033)	(0.049)
Disadv. Caste (Oster)	0.013	-0.049^{***}	-0.016	-0.078^{**}	-0.119^{***}	-0.082
	(0.016)	(0.021)	(0.022)	(0.032)	(0.042)	(0.055)
Constant	10.991^{***}	10.869^{***}	10.807^{***}	10.734^{***}	10.601^{***}	10.834^{***}
	(0.030)	(0.042)	(0.041)	(0.064)	(0.078)	(0.114)
Observations	619	552	500	449	378	378
Adjusted R^2	0.546	0.470	0.537	0.535	0.509	0.362
Year FE	✓	✓	✓	✓	✓	✓
Other Controls	✓	✓	✓	✓	✓	✓
Job FE						

Notes: This table reports earnings regressions for students with identifiable caste surnames at each job search stage. Each column represents a different stage: (1) application, (2) written test, and subsequent rounds. The dependent variable is log earnings. All regressions include controls and year fixed effects. Job fixed effects are omitted as salaries are identical within jobs regardless of caste. Bias-adjusted coefficients follow Oster (2019), setting $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$ (assuming unobservables match observables in selection intensity and allowing total explanatory power to exceed observed R^2 by 30%). $^{*}p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

Table OA.8: **Caste Differences in Mean Salary at Each Job Search Stage in the Non-Identifiable Sample**

	Dependent variable:					
	log (Salary)					
	App	Test	Debate	Interview	Offer	Choice
	(1)	(2)	(3)	(4)	(5)	(6)
Disadv. Caste (OLS)	0.002	0.0004	−0.003	−0.004	−0.100***	−0.106***
	(0.005)	(0.007)	(0.007)	(0.011)	(0.014)	(0.019)
Disadv. Caste (Oster)	0.004	0.019***	0.013	0.024**	−0.075***	−0.076***
	(0.006)	(0.007)	(0.008)	(0.011)	(0.015)	(0.018)
Constant	11.041***	10.913***	10.857***	10.766***	10.685***	10.869***
	(0.012)	(0.017)	(0.016)	(0.024)	(0.033)	(0.043)
Observations	3,386	3,066	2,766	2,582	2,149	2,149
Adjusted R^2	0.584	0.500	0.552	0.556	0.486	0.421
Year FE	✓	✓	✓	✓	✓	✓
Other Controls	✓	✓	✓	✓	✓	✓
Job FE						

Notes: This table reports earnings regressions for students with non-identifiable caste surnames at each job search stage. Each column represents a different stage: (1) application, (2) written test, and subsequent rounds. The dependent variable is log earnings. All regressions include controls and year fixed effects. Job fixed effects are omitted as salaries are identical within jobs regardless of caste. Bias-adjusted coefficients follow Oster (2019), setting $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$ (assuming unobservables match observables in selection intensity and allowing total explanatory power to exceed observed R^2 by 30%). $^*p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

Table OA.9: **Caste Differences in GPA Between Those Who Left the First Job Versus Those Who Did Not**

	Dependent variable:		
	log (GPA)		
	(1)	(2)	(3)
Disadv. Caste	-0.065^{***}	-0.065^{***}	0.010^{***}
	(0.004)	(0.004)	(0.002)
Left Job	-0.072^{***}	-0.072^{***}	-0.004^{**}
	(0.003)	(0.003)	(0.002)
Left Job * Disadv. Caste	0.047^{***}	0.047^{***}	-0.004
	(0.005)	(0.005)	(0.003)
Constant	2.188^{***}	2.182^{***}	2.198^{***}
	(0.003)	(0.005)	(0.028)
Year FE	✓	✓	✓
Major		✓	✓
Degree		✓	✓
Job FE			✓
Observations	2,489	2,489	2,489
Adjusted R^2	0.260	0.265	0.822
Note:		*p<0.1; **p<0.05; ***p<0.01	

Notes: This table reports caste differences in GPA between job stayers and leavers. Column (1) includes year fixed effects only, Column (2) adds major and degree controls, and Column (3) additionally includes job fixed effects. $^{*}p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

Figure OA.1: **Common Support Within Each College GPA Decile**

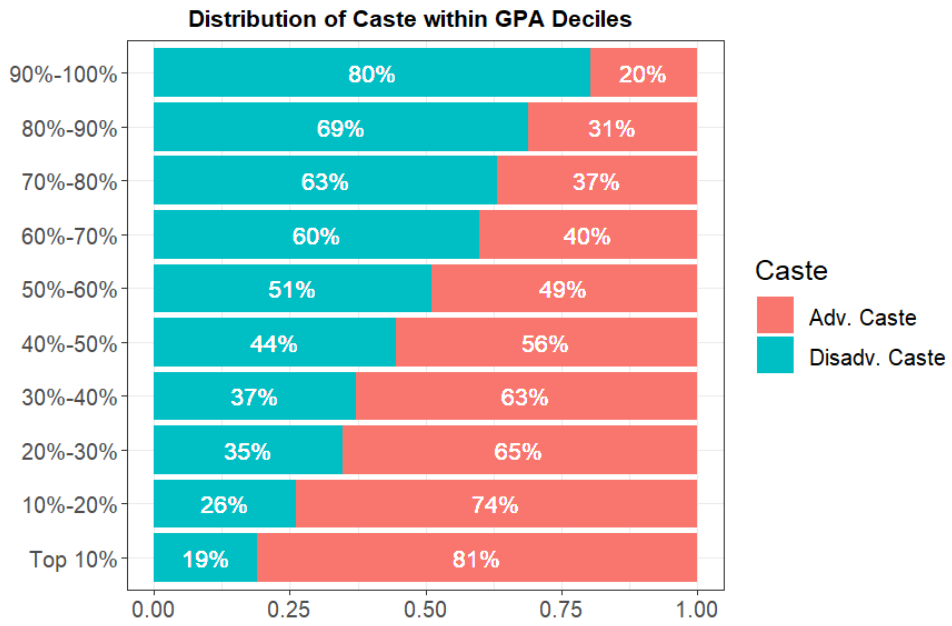

Notes: This figure shows the distribution of advantaged and disadvantaged caste students across GPA deciles. GPAs are normalized within degree-year.

Figure OA.2: **Caste Penalty in the Conditional Success Rate At the Application Reading Round in the Identifiable and the Non-Identifiable Sample**

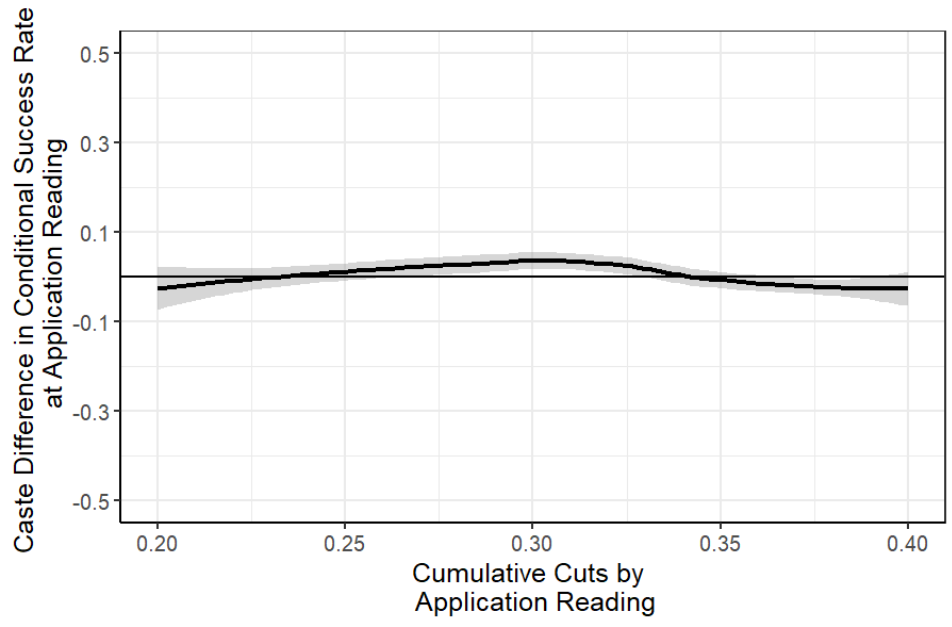

(a) Caste Penalty in Conditional Success Rates versus Cumulative Cuts by Application Reading in the Full Sample

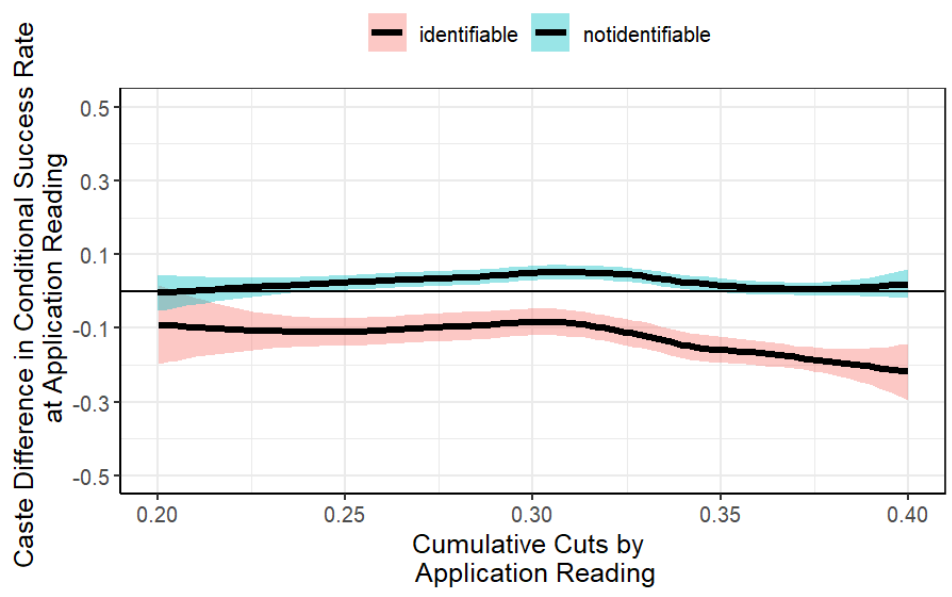

(b) Caste Penalty in Conditional Success Rates versus Cumulative Cuts by Application Reading in the Identifiable and Non-Identifiable Sample

Notes: This figure shows caste differences in success rates at the application stage against cumulative selection cuts. Panel (a) uses the full sample; Panel (b) separates students with identifiable (red) and non-identifiable (blue) caste surnames. The x-axis shows the cumulative proportion of rejected applicants. Y-values represent the disadvantaged caste coefficient from job-specific linear probability models including controls (degree, major, GPA, entrance exam score, cohort year). Black lines show smoothed conditional means with 95% confidence intervals (shaded). The horizontal line at zero indicates no caste difference.

Figure OA.3: **Caste Penalty in the Uncontrolled Success Rate At Each Screening Round**

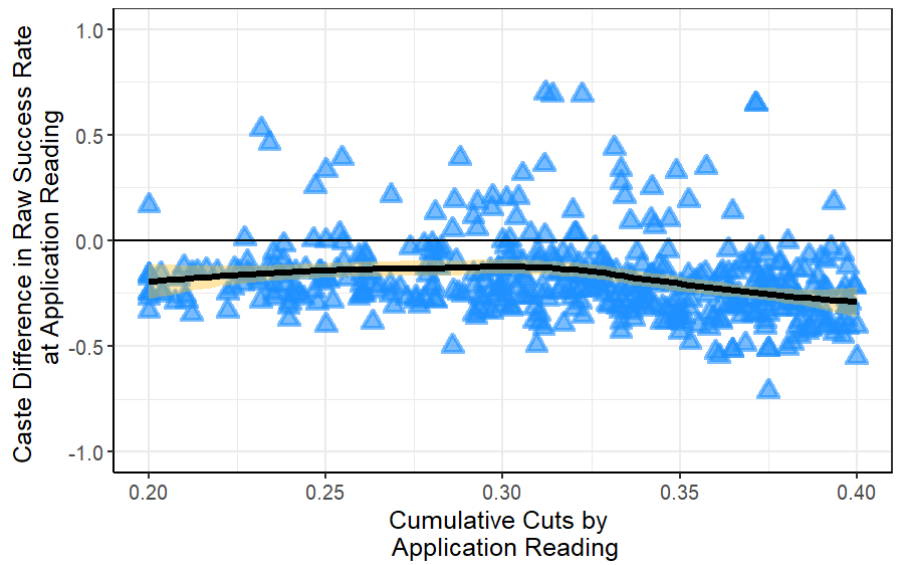

(a) Caste Penalty in Raw Success Rates versus Cumulative Cuts by Application Reading

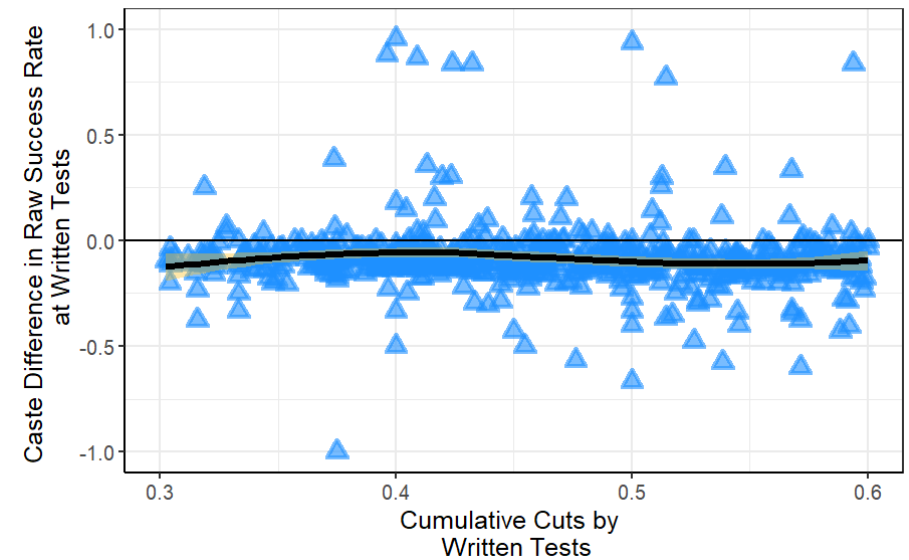

(b) Caste Penalty in Raw Success Rates versus Cumulative Cuts by Written Tests

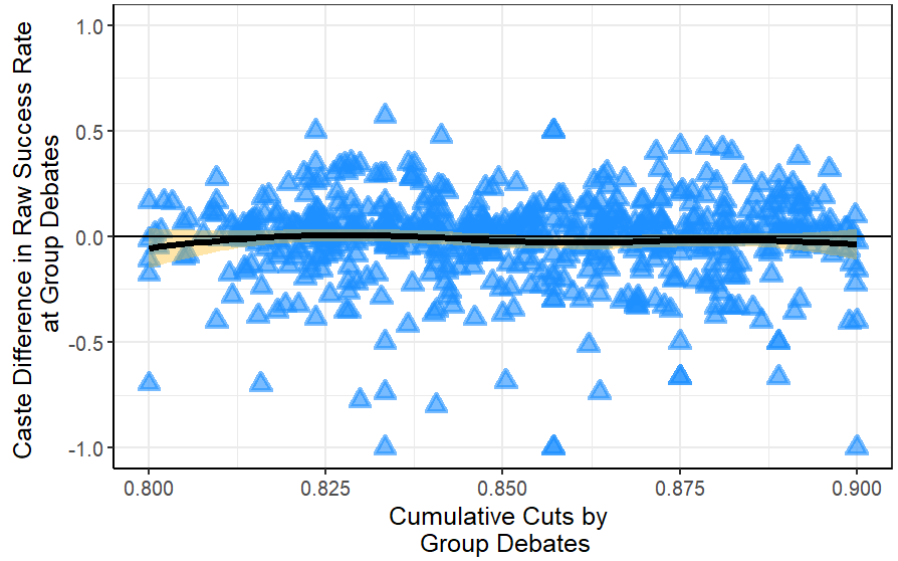

(c) Caste Penalty in Raw Success Rates versus Cumulative Cuts by Group Debates

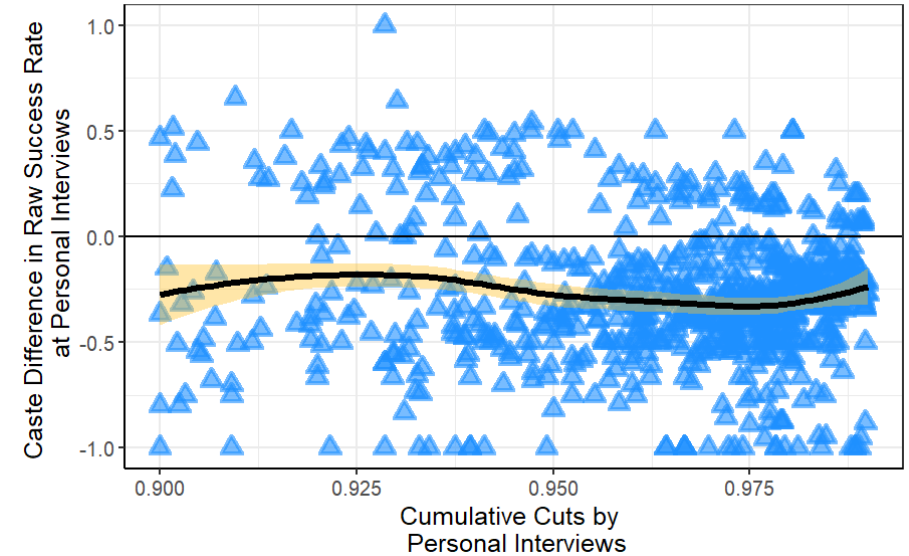

(d) Caste Penalty in Raw Success Rates versus Cumulative Cuts by Personal Interviews

Notes: This figure shows raw (unconditional) caste differences in success rates against cumulative selection cuts at each screening stage. Each blue triangle represents a job. Black lines show smoothed conditional means with 95% confidence intervals (shaded). The x-axis shows the cumulative proportion of rejected applicants at each stage. The y-axis shows the raw caste gap in success rates. The horizontal line at zero indicates no caste difference.