Making the Elite: Class Discrimination at Multinationals

Soumitra Shukla*

May 28, 2025

Abstract

What drives socioeconomic disparities in elite job access, and where do these inequalities emerge? Measuring class discrimination is challenging, as status must often be inferred from subtle markers like accent, education, or networks. India's caste system provides an explicit social class measure, enabling precise identification of how status affects elite job access. Analyzing detailed screening data from over 1,000 jobs at U.S. and European multinationals in India, this paper provides the first systematic quantification of class bias in elite recruitment, identifying both its sources and mechanisms. Personal interviews assessing candidate "fit" drive nearly 90% of the caste earnings penalty (0.16 standard deviations). However, disadvantaged caste hires achieve 20% higher promotion rates, suggesting these disparities cannot be justified by ability differences. To formally distinguish discrimination from unobserved ability differences, I develop a theoretical framework utilizing plausibly exogenous variation in screening selectivity across jobs. The framework predicts that, absent discrimination, the promotion advantage of disadvantaged caste hires should shrink in highly selective jobs. Instead, persistently higher and stable promotion rates of disadvantaged hires across all selectivity tiers are most consistent with discrimination. These patterns persist across both client-facing and non-client-facing roles, suggesting bias rather than statistical or customer-driven motives. My findings reveal how subjective "fit" assessments can undermine both diversity and efficient talent allocation in elite hiring.

*I would like to thank Joseph Altonji, Costas Meghir, and John Eric Humphries for serving on my committee and providing guidance and feedback. I would also like to thank Jason Abaluck, Peter Arcidiacono, Crémieux, Michael Dinerstein, Arpit Gupta, Patrick Kline, Scott Kominers, Eliza McNay, Suanna Oh, Evan Rose, Kamalini Ramadas, Jung Sakong, Rob Valletta, Jon Williams, seminar participants at the Yale labor workshop, the Chicago FED, Chicago Harris, Wharton (BEPP), Penn State, UCSD, SOLE, Southeastern Micro Labor Workshop, System Equitable Growth Conference at the New York Fed, the "Discrimination in the 21st Century Conference" organized by the Becker Friedman Institute and Chicago Booth, MIT Organizational Economics Lunch, and the NBER Race and Stratification Working Group for helpful feedback and support. I used OpenAI's ChatGPT and Anthropic's Claude for light copy-editing.

"People with advantages are loath to believe that they just happen to be people with advantages." — C. Wright Mills, *The Power Elite* (1956)

1 Introduction

Who gains access to top jobs, and how do hiring practices shape inequality in elite recruitment? While extensive research documents labor market discrimination (e.g., Altonji and Blank, 1999), we lack precise evidence on *where*, *how*, and *how much* hiring discrimination occurs in high-stakes, real-world settings like multinational corporations (MNCs).

Measuring discrimination is especially challenging when the relevant marker is social class rather than race, gender, or ethnicity. In many settings, including the U.S. and the U.K., status is often inferred from overlapping cues—accents, school pedigrees, inherited networks, income—complicating precise measurement (Lamont and Lareau, 1988; Bisin and Verdier, 1998; Durlauf, 1999, 2008; Khan, 2011; Ridgeway, 2019; Stansbury and Rodriguez, 2024a). Income itself provides ambiguous signals: high earners with working-class accents face status penalties, while those with elite social networks but modest incomes retain advantages (Bourdieu, 1986). For example, studies in the U.K. show that a high-earning executive who retains a regional working-class accent is often rated as lower status than a lower-paid professional speaking Received Pronunciation or listing a prestigious boarding school (Savage et al., 2013; Savage, 2015; Laurison and Friedman, 2019; Levon et al., 2021).[1] More generally, evaluators choosing among similarly qualified candidates often rely on holistic assessments of who "fits"—referring to candidates' cultural alignment, interpersonal rapport, perceived similarity in values or backgrounds, expectations regarding team dynamics, and other job-specific competencies (Lamont, 1992; Rivera, 2012, 2015).

Holistic evaluations that emphasize "fit" often reinforce class advantages under the guise of comprehensive assessment. Test-optional policies reduce access for high-achieving, low-income students (Sacerdote et al., 2025), legacy preferences at selective U.S. universities favor the wealthy despite weak performance links (Chetty et al., 2023), and "personality" interviews used in selection for India's elite civil service penalize lower-status candidates while adding little predictive value (Bhavnani and Lee, 2021). Such subjective judgments often encode deep-seated cultural expectations (Bisin and Verdier, 2010). Even outside of these specific examples, decades of evidence show holistic assessments predict no better than standardized metrics, *even among experts*, while systematically disadvantaging lower-status applicants (Meehl, 1954; Ægisdóttir et al., 2006; Kuncel and Hezlett, 2007; Murphy et al., 2009; Rivera, 2012; Bol et al.,

[1] Received Pronunciation, sometimes called "BBC English" or "Queen's English," is the historically prestigious accent of English traditionally associated with the British upper classes and elite institutions (Savage et al., 2013).

2014; Kuncel et al., 2014; Laurison and Friedman, 2019; Stansbury and Rodriguez, 2024a). Thus, the very mechanism meant to refine selection often obscures merit.

Establishing that fit assessments lead to social class discrimination requires precise measurement of class itself—a challenge that motivates my focus on the Indian caste system, which provides an unusually comprehensive marker of social stratification. While caste is uniquely Indian, it embodies the broader concept of social class that Bourdieu (1986) conceptualizes as a joint stock of four capital forms: economic (money, property rights), social (durable networks), cultural (embodied traits like accent and etiquette; objectified possessions; institutionalized credentials), and symbolic (prestige from public recognition of other capitals). Unlike the more fluid manifestations of class in other contexts, India's caste system codifies these capital forms into a relatively stable, administratively codified hierarchy (Beteille, 1965, 1969). This codification provides rare precision for measuring class discrimination in elite jobs, which disadvantaged castes remain largely excluded from despite comprising 70% of India's billion-plus population (Bayly, 2008).

Leveraging this measurement advantage, I examine recruitment at leading U.S. and European multinationals in India to make three contributions: First, I quantify social-status discrimination in elite recruitment. Second, I identify personal interviews that assess candidate fit as the primary mechanism revealing and penalizing social status. Third, I show that class bias—not statistical or customer discrimination—drives these disparities.[2] To the best of my knowledge, each of these contributions is novel in the literature.

Personal interviews penalize social status, driving nearly all caste penalties in hiring. Across the 1,050 jobs studied, each position involves four screening rounds progressing from standardized to subjective assessments: application reading, aptitude tests, group debates on socio-emotional skills, and, finally, personal interviews assessing candidate "fit." These jobs all recruit from an elite Indian college, with nearly half the students belonging to disadvantaged castes due to *mandated affirmative action* (Madheswaran, 2008). The unconditional caste earnings gap is 23%, narrowing to 10% after controlling for pre-college skills, academic performance, and other employer-relevant factors. With salaries explicitly *identical across castes* for identical roles and *non-negotiable*, I decompose these gaps by counterfactually assigning posted salaries to the set of jobs each caste group is in contention for at each stage before offers are accepted. Personal interviews alone account for nearly 90% of the conditional earnings penalty (approximately 9% of average earnings, or 0.16 standard deviations) and about 95% of the nearly 30 percentage point caste penalty in hiring success rates, highlighting their role in penalizing social status.

Personal interviews reveal social status. While some surnames indicate caste, most do not convey caste signals in urban-educated settings (e.g., elite colleges typically located in urban

[2] Such discrimination violates Article 15 of the Indian Constitution (Ministry of Law and Justice, Government of India, 2021), limiting mobility and misallocating talent.

centers) due to migration and geographical variation in naming conventions (Mamidi, 2011; Jodkha, 2017; Wiltshire, 2020). Among the 15% of students with distinctly caste-identifiable surnames, penalties emerge immediately at the first screening round (application reading). For the remaining 85% with caste-ambiguous surnames, no early penalty appears—disparities emerge only after personal interviews. This differential timing strongly suggests that the penalty occurs precisely when caste becomes visible to employers, with personal interviews serving as the primary revelation mechanism for social status.

Interviews may provide information, but cannot justify observed disparities. Disadvantaged caste hires consistently outperform advantaged peers by around 20% in promotion rates, even in highly selective roles, strongly indicating disparities stem from discrimination rather than unmeasured skill deficits. To formalize this interpretation, I develop a theoretical framework that leverages plausibly exogenous variation in screening selectivity—the depth in the applicant pool from which each firm selects—across jobs. The framework yields two key implications: first, if disadvantaged candidates genuinely lacked skills, highly selective pre-interview screening stages would show a caste penalty, which they do not; second, if interviews accurately filtered weaker candidates, disadvantaged hires would not consistently outperform advantaged peers across selectivity tiers. As selectivity increases, the average hired candidate's quality approaches marginal (cutoff) quality. If jobs were using the same cutoffs regardless of caste, promotion differences between castes should shrink as job selectivity increases. Thus, persistently higher and stable promotion rates among disadvantaged hires across all selectivity levels suggest discrimination rather than productivity differences.

Analysis of distinctive versus non-distinctive surnames, combined with the selectivity evidence above, strongly indicates that hiring disparities reflect *disparate treatment* (explicit caste-based discrimination). Moreover, disadvantaged caste hires retain their promotion advantage across both client-facing and non-client-facing jobs, suggesting that class bias (either taste-based discrimination or biased beliefs) rather than statistical or customer-driven discrimination drives the caste hiring penalty.

This paper contributes to labor, personnel, and organizational economics by quantifying social-status discrimination in elite hiring and identifying its mechanisms. This paper advances the still-nascent literature on class discrimination, which lags behind research on race, gender, and ethnicity (Durlauf, 2006; Rivera, 2012, 2015). The paper is most closely related to Stansbury and Rodriguez (2024a), who document sizable class gaps in U.S. academic careers; Rivera (2012), who argues that "cultural matching" in elite-firm interviews could engender class disparities; Bhavnani and Lee (2021), who show that disadvantaged caste civil service officers score markedly lower in "personality" interviews, yet still outperform advantaged caste recruits[3]; and

[3]This pattern holds particularly for those who scored high enough to have qualified even without affirmative action quotas.

Zimmerman (2019), Michelman et al. (2022), and Barrios-Fernández et al. (2024), who demonstrate the role of elite education in reproducing class advantage through what Bourdieu (1986) conceptualizes as peer-driven social and cultural capital. No prior work, however, simultaneously quantifies the magnitude of class discrimination, locates it in the hiring pipeline, and isolates its mechanisms.

As traditional social-status markers become less visible across societies, class discrimination operates through *social interactions*—such as interviews—where subtle cues can reveal and penalize social background (Lenski, 1954; Grusky, 2005; Rivera, 2015; Clark, 2015). My findings demonstrate how subjective fit evaluations can preserve discrimination even within ostensibly meritocratic systems, often enabling biases while maintaining plausible deniability (Castilla and Benard, 2010; Reskin, 2012; Castilla, 2015; Kang et al., 2016). These patterns likely extend to other concealable stigmatized identities: LGBTQ+ status and other intersectional attributes (Rivera, 2012, 2015; Eames, 2024; Chow and Knowles, 2016).

This paper traces how personal interviews—a widely adopted practice at multinationals—shape hiring inequality by affording hiring managers the opportunity to discriminate on social status, a topic that has received limited attention in economics. While economic studies extensively document discrimination at early stages like résumé screens (e.g., Kline et al., 2022), few examine bias within interview rooms where subjective evaluations occur. Meta-analyses suggest additional discrimination emerges during face-to-face interactions, as race and gender gaps persist or widen from initial callbacks to final offers (Quillian et al., 2020). By pinpointing when class penalties crystallize and linking them to post-hire performance, this study provides novel evidence on a ubiquitous yet understudied hiring practice while highlighting limitations of interventions like résumé anonymization that may merely shift bias to later stages.

Another central contribution illuminates managers' critical role in perpetuating class discrimination through hiring decisions, revealing both cultural biases and structural agency problems. Recent literature documents substantial managerial influence on firm performance and culture (Turco, 2010; Lazear et al., 2015; Sadun et al., 2017; Hoffman and Tadelis, 2021; Espinosa and Stanton, 2022; Fenizia, 2022; Sadun, 2023; Benson and Shaw, 2024; Minni et al., 2025), particularly through assigning workers to teams and training investments (Minni, 2023; Sadun, 2023; Sadun et al., 2024). This paper uniquely integrates Bourdieu's concept of social class—especially its social and cultural capital dimensions—to illuminate biases in managerial assessments of who "fits" (Bourdieu, 1986). Additionally, this paper highlights agency problems in hiring: the separation of hiring authority from performance accountability and managers' reliance on overconfident private judgments (Hoffman et al., 2018) or productivity-irrelevant cues (Benson et al., 2023). Hiring managers in my setting rarely supervise recruits directly, creating incentives to prioritize perceived "fit" over long-term potential. Such subjective, short-term decisions systematically produce sub-optimal organizational outcomes, as documented by

extensive research (Gibbons and Roberts, 2013; Sadun et al., 2017; Benson and Shaw, 2024; Hoffman and Stanton, 2024).

Finally, this paper reveals how subjective screening practices that provide room for status-based discrimination can erode affirmative action gains, undermining the socioeconomic mobility that elite education is designed to create. Elite colleges like the studied Indian institution, with a talent pool comparable to that at many Ivy League universities (Martellini et al., 2024) and affirmative action polices in admissions (Madheswaran, 2008), serve as gateways to prestigious careers (Jodhka and Naudet, 2019). Although similar affirmative action policies worldwide substantially increase disadvantaged students' access to elite colleges, research indicates elite education disproportionately benefits privileged students, perpetuating social class advantage across generations (Zimmerman, 2019; Barrios-Fernández et al., 2024). Given multinational firms' recruitment practices, my findings highlight the need to integrate affirmative action in college admissions with hiring reforms downstream, emphasizing integrated approaches for socioeconomic mobility—particularly given how social class co-evolves with cultural and institutional factors shaping economic outcomes (Bisin et al., 2023).

2 Institutional Setting and Key Definitions

This section provides essential background on India's caste system as a measure of social status, defines key terms used throughout the analysis, describes the elite college job fair studied in this paper, and explains why this setting is representative of elite hiring in India.

2.1 Caste as Social Status Under Affirmative Action

India's affirmative-action policy framework treats caste as a unified marker of social status, bundling financial resources, social networks, and cultural capital into two administratively codified categories: "advantaged" and "disadvantaged." These policies apply only to public universities and government jobs, not to private multinationals. For historical background on caste classification and reservation policy, see Online Appendix Section A.

2.2 Key Terms: Disadvantaged Castes, Elite Jobs, and Elite Colleges

Below, I define four key terms used throughout the analysis:

a. **Disadvantaged castes.** Following the Indian government's official classifications, the term "disadvantaged castes" encompasses individuals belonging to Scheduled Castes (SCs), Scheduled Tribes (STs), or Other Backward Classes (OBCs).

b. **Job.** A "job" refers specifically to a job designation within a firm. For example, if Google recruits both a product manager and a software engineer, these constitute two distinct jobs. Section 3.6 explains why a job typically corresponds to a firm-designation pair, rather than a particular office location.

c. **Elite jobs.** An "elite entry-level job" refers to a formal entry-level position at multinational corporations (MNCs), major national corporations, or unicorn startups (privately-held startups valued at over $1 billion). Throughout this paper, the terms "elite entry-level jobs" and "elite jobs" are used interchangeably.

d. **Elite colleges.** "Elite colleges" include institutions consistently ranked among the top 30 in their respective fields—such as science, engineering, arts and humanities, social sciences, and business—by *India Today*, an Indian equivalent of the *US News & World Report*. Several elite Indian colleges also place among the top colleges worldwide for graduate earnings, according to a cross-country analysis of college quality in Martellini et al. (2024), underscoring their international standing.

2.3 Structure of the On-Campus Job Fair

The elite college's job fair occurs *entirely on campus* and comprises five phases: (1) pre-placement, (2) student and employer registration, (3) return offers from internships, (4) placement, and (5) post-placement (or the "aftermarket").

a. **Pre-Placement phase.** Firms are invited by the placement office prior to June. Between June and mid-August, recruiters deliver pre-placement talks, gauge student interest, and finalize roles.

b. **Student and employer registration.** Students register for placement by late August, while employers submit detailed registration forms by early September, specifying job roles, compensation, and vacancies, which become fixed thereafter. Compensation is *non-negotiable*, strictly enforced by the placement office through verification of offer letters. Non-compliance can result in restrictions on future campus visits by firms.[4]

c. **Return offers from internships.** By late August, firms must extend pre-placement offers (PPOs) to students who interned with them during the summer, following deadlines set by the placement office. Students accepting PPOs are withdrawn from the main placement process. *Unlike in other contexts such as the U.S., these internship return*

[4]Such rules and timelines are common across elite Indian colleges and are available on their websites. For instance, see Central Placement Cell (Delhi University).

offers play a minor role: fewer than 5% of students accept PPOs and exit the placement process, leaving over 95% to actively participate in the main job fair (Section 4.1).

d. **Placement phase.** Starting in mid-September, students apply for jobs. Employers initially screen applications, followed by written aptitude tests and group debates assessing socio-emotional skills from September to early December. Final-round personal interviews occur between December and January, after which about two-thirds of the total student cohort secure full-time jobs through this standardized placement process.

e. **Post-Placement phase ("aftermarket").** From late January until graduation in July, the placement office operates an "aftermarket" for students without initial job offers, comprising roughly one-third of the cohort. These unplaced students typically pursue entrepreneurship, graduate studies, competitive exam preparation (such as civil services or elite MBAs), or find alternative employment.

Excluding firms and students outside the standard recruitment pipeline. The analysis excludes aftermarket employers—typically smaller startups offering lower salaries with less standardized recruitment—and firms bypassing the four-stage pipeline (e.g., only résumé screening and interviews). These exclusions are necessary given the focus on standardized hiring practices; implications are discussed in Section 4.4.

2.4 Why the Studied Job Fair is Representative

The studied job fair is representative of how elite Indian graduates transition into elite jobs. Nearly all top-ranked Indian colleges use similar standardized placement procedures, elite employers recruit consistently from the same institutional network, elite entry-level roles recruit almost exclusively from these campuses, and student participation rates in campus job fairs are near-universal. For evidence supporting these claims, see Online Appendix Section B.

3 Data Overview

This section describes the data used in the analysis, detailing student and job characteristics, recruitment procedures, and career outcomes such as promotions and tenure. It concludes by discussing relevant data limitations.

3.1 Institutional Source and Completeness of Recruitment Data

Data are obtained directly from the placement office of an elite Indian university (anonymized for confidentiality), which meticulously records each recruitment stage—applications, screening

rounds, interviews, and final job offers. The university's standardized recruitment practices benchmark placement processes at comparable institutions nationwide (Jodhka and Naudet, 2019).

3.2 Centralized Job Portal

Students register through a centralized online job portal managed by the placement office, creating standardized personal profiles including GPA, entrance exam scores, major coursework, and internship experiences. Firms submit detailed profiles specifying compensation, job descriptions, and qualifications (Section 2.3). Both student and job profiles are mutually visible via the portal. Throughout recruitment, employers directly upload candidate shortlists from each screening round onto the portal, ensuring systematic documentation and data consistency.

3.3 Standardized Recruitment Rules and Mandatory Participation Ensuring Data Integrity

Recruitment at elite Indian colleges is tightly standardized: students who enter the campus placement fair are *mandated* by the placement office to participate in all four screening rounds—application screening, written aptitude tests, group debates, and personal interviews (Section 3.8). Offer-letter salaries are likewise verified against the non-negotiable compensation schedules filed by firms (Section 2.3). Such standardized rules enhance transparency and empirical rigor.

3.4 Student Sample: Time Span, Degrees, Demographics, and Caste Representation

The dataset spans four placement cycles at an elite Indian college (exact years anonymized). It includes 4,164 students pursuing four degrees: Bachelor of Technology (B.Tech.), Dual Degree (integrating undergraduate and postgraduate coursework), Master of Technology (M.Tech.), and Master of Science (M.S.).

Admissions based solely on standardized national entrance exams. Admissions are *exclusively* based on standardized national-level exams with caste- and major-specific cutoffs.

Near-equal caste representation via affirmative action. Due to caste-based affirmative action in admissions (Section 2), disadvantaged and advantaged castes are represented at near-

parity *within each major*, and therefore within most degrees (Table 1).[5]

Caste self-reporting. Students self-report caste through official administrative documents, enabling accurate tracking by both the college and researcher.

Predominantly male sample. Approximately 90% of students are male, typical of elite technical colleges in India (cf. Datta, 2017).

Table 1: **Total Students by Caste and Degree**

	Total Students		
	Advantaged Caste	Disadvantaged Caste	Degree Total
B.Tech.	579	710	1289
Dual Degree	622	617	1239
M.Tech.	601	566	1167
M.S.	344	125	469
Caste Total	2146	2018	4164
Fraction	0.52	0.48	1.00

Notes: Table 1 shows student counts by caste and college degree. B.Tech. stands for Bachelor of Technology, M.Tech. for Master of Technology, and M.S. for Master of Science. A Dual degree integrates undergraduate and post-graduate studies, completed one year after the four-year B.Tech.

3.5 Student Characteristics and Caste Differences

Student characteristics include college GPA, pre-college academic skills, and prior labor market experience, highlighting two key patterns:

1. **Significant caste differences in academic performance.** Large caste disparities exist in college GPA and entrance exam (EE) scores (Table 2), although students from both caste groups span all deciles of these metrics (Online Appendix Figure OA.1).

2. **Similar internship patterns across castes.** The dataset includes detailed internship information (e.g., duration, pay, employment sector). Caste is only weakly correlated with internship sector or experience. For instance, among B.Tech. and Dual Degree students, internships across castes show similar sector distributions (25% technology, 37% consulting, and 38% manufacturing). Similar patterns hold for postgraduate degrees (not shown). This likely reflects that return internship offers are rare (Section 2.3), making

[5]In some degrees, particularly the M.S., disadvantaged-caste seats often remain partially unfilled due to eligibility constraints or fewer qualified applicants, consistent with patterns documented in Krishna and Frisancho Robles (2015).

internship choices primarily driven by student sector preferences rather than caste-related pay or opportunities.

Table 2: **Summary Statistics (Students)**

	Adv. Caste	Disadv. Caste	Diff (SD)
		B.Tech.	
Avg. Entrance Exam Score	0.41	-0.33	0.74***
Avg. 10th Grade Score	0.24	-0.19	0.43***
Avg. 12th Grade Score	0.24	-0.20	0.44***
Avg. Overall College GPA	0.51	-0.42	0.93***
		Dual Degree	
Avg. Entrance Exam Score	0.35	-0.35	0.70***
Avg. 10th Grade Score	0.20	-0.20	0.40***
Avg. 12th Grade Score	0.14	-0.14	0.28***
Avg. Overall College GPA	0.43	-0.43	0.86***
		M.Tech.	
Avg. Entrance Exam Score	0.27	-0.30	0.57***
Avg. 10th Grade Score	0.23	-0.15	0.38***
Avg. 12th Grade Score	0.17	-0.10	0.27***
Avg. Overall College GPA	0.34	-0.33	0.67***
		M.S.	
Avg. Entrance Exam Score	-0.02	0.02	-0.04
Avg. 10th Grade Score	0.03	-0.07	0.10
Avg. 12th Grade Score	0.02	-0.06	0.08
Avg. Overall College GPA	0.04	-0.14	0.18*

Notes: Table 2 reports pre-college skills (entrance exam scores, 10th and 12th standard exit exams) and college GPA. All scores are normalized to zero mean and unit standard deviation within each degree and year. Entrance exam scores (originally ranks) are renormalized so higher numbers are better. Caste differences shown in standard deviation units. $^{*}p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$. Adv. Caste and Disadv. Caste denote advantaged and disadvantaged castes.

3.6 Characteristics of Multinational Private-Sector Jobs

This subsection summarizes key job characteristics, including sector distribution, salary structure, recruitment procedures, and managerial involvement.

Sector and employer overview. The dataset includes approximately 1,050 private-sector jobs, mostly at multinational corporations headquartered in the U.S. or Europe (Table 3). Around 75% of roles are non-client-facing (e.g., software engineering), with 25% client-facing (e.g., consulting).

Exclusion of public-sector jobs. Public-sector roles (about 2%; Table 3) are excluded due to distinct salary structures and different recruitment procedures (Seventh Central Pay Commission, 2016.)

Standardized, non-negotiable salaries and benefits. For a given college degree, salaries and amenities are standardized, non-negotiable, and uniform for each job role, independent of individual characteristics like caste, gender or major (or any other candidate-specific factors). Firms typically adopt nationwide salary standards, recruiting identically across elite campuses.[6]

Table 3: **Summary Statistics (Jobs)**

	Count	Fraction	Avg. Salary ($)	Std. Dev. ($)
Private	1011	0.98	60341.38	32591.51
Government	19	0.02	56202.23	20240.01
Headquartered Outside India	620	0.60	64562.74	36038.04
Headquartered in India	410	0.40	53766.04	23946.16
Client Facing	246	0.24	57553.23	28288.23
Non-Client Facing	784	0.76	61115.92	33361.87
All Jobs	1030	1.00	60265.02	32391.89

Notes: Table 3 reports the distribution of jobs between public and private sectors, domestic and foreign headquarters (primarily US or Europe), and client-facing versus non-client-facing roles. Average salaries (USD PPP) are also shown for each group.

Job location and recruitment processes. About 99% of jobs are India-based. Roles abroad are excluded due to different compensation and hiring structures. When hiring for Indian job positions, employers conduct screening rounds for all Indian offices at once. The specific Indian office that a candidate ultimately gets assigned to is decided later and usually depends on the location-specific needs of that employer. For example, Google will conduct a combined recruitment process for its software engineering role regardless of whether the candidate may ultimately get assigned to either the Mumbai or Bangalore office of Google. Therefore, a "job" will almost always mean a firm and job designation pair and not the ultimate (Indian) office location. The standardization of compensation packages is also compatible with how recruitment is conducted for an employer's Indian locations.

[6]Multinationals typically centralize entry-level compensation at headquarters, giving campus recruiters little discretion to vary pay across locations or candidates (e.g., Bloom et al., 2012; Sarsons et al., 2022; Minni, 2023, 2024; Sadun et al., 2024).

Recruitment conducted by Indian hiring managers separately from direct supervisors. Because most jobs are based in India, recruitment is conducted by specialized Indian hiring managers. These managers typically remain distinct from eventual direct supervisors due to logistical constraints posed by placement processes of elite colleges that must be entirely conducted *on-campus* (including written tests that are supervised by recruiters in college-assigned exam rooms) and therefore require in-person presence of recruiters.

Internship return offer policy. Students accepting internship return offers exit the main placement cycle and are excluded from analysis (Section 2.3).

3.7 Quantifying Job Selectivity at Each Recruitment Stage

This subsection summarizes the selectivity of jobs using two key concepts—*conditional cuts*, the fraction of applicants eliminated *at each round*, and *cumulative cuts*, the total fraction eliminated *up to a given round*. Table 4 reports these metrics. For instance, a job eliminating 50% of applicants each round has a 50% conditional cut at each stage, but a cumulative cut of 75% after two rounds. Typically, successful candidates rank within the top 5% of initial applicants. Occasionally, highly selective roles extend no offers.

Table 4: **Variation in Conditional and Cumulative Selectivity Across Jobs**

Cut Type	Round	Q1	Mean	Q2	Q3
Conditional Cut	App. Reading	0.30	0.39	0.39	0.48
Cumulative Cut	App. Reading	0.30	0.39	0.39	0.48
Conditional Cut	Written Test	0.10	0.18	0.15	0.26
Cumulative Cut	Written Test	0.38	0.49	0.49	0.61
Conditional Cut	Group Debate	0.69	0.74	0.74	0.79
Cumulative Cut	Group Debate	0.83	0.87	0.87	0.91
Conditional Cut	Personal Interview	0.61	0.75	0.78	0.90
Cumulative Cut	Personal Interview	0.95	0.96	0.97	0.99

Notes: Table 4 shows the distribution of conditional and cumulative selection cuts across jobs.

3.8 Four Recruitment Stages: From Objective to Subjective Evaluations

Recruitment for all jobs progresses uniformly from relatively objective assessments to subjective evaluations:

Low-cost applications via centralized portal. A single click on the centralized portal forwards the same standardized résumé to any additional job (Section 3.2), making the marginal cost of applying essentially zero; application volumes are therefore virtually identical across castes (Figure 1).

Eligibility criteria restrict applications. Applications are constrained by major, degree, and GPA. For example, software engineering roles often require Computer Science majors, whereas analyst roles typically allow all majors.

Four mandatory screening rounds. All jobs strictly follow these sequential stages:

1. **Application reading.** Employers use standardized criteria and GPA cutoffs to initially filter resumes, typically involving limited subjective assessment.
2. **Written aptitude tests.** Standardized online tests objectively evaluate technical, analytical and critical thinking skills.
3. **Group debates.** Structured debates, conducted in multiple batches, involve 20–30 students divided into two teams, assessing communication, confidence, and teamwork under standardized conditions, including uniform prompts and roughly equal speaking time per candidate.
4. **Personal interviews.** Also sometimes referred to as HR interviews, these final-round interviews last about 30 minutes and emphasize subjective evaluations of interpersonal skills, organizational fit, cultural alignment, and role-specific competencies.

Throughout each stage, employers maintain access to candidates' standardized profiles via the centralized portal (Section 3.2). Additional details for screening rounds and candidate selection, including timeline, are shown in Online Appendix Section D.

Common screening rounds across multiple positions. Employers offering multiple distinct roles frequently administer common screening simultaneously for all applicants (e.g., Nvidia uses a single exam for all advertised profiles).

3.9 Candidate Screening: Summary Statistics

Table 5 summarizes candidate progression through each recruitment stage. On average, jobs receive roughly 209 applications, with approximately 125 advancing to written tests, 102 reaching group debates, 26 participating in personal interviews, and about 6 candidates receiving job offers. This substantial reduction reflects high selectivity at each stage (Table 4). Eligibility constraints (major, degree, GPA) naturally limit the number of applications per student (Section 3.8).

Table 5: **Summary Statistics (Candidate Screening)**

Round	Mean	Median
Application Reading	208.85	140.17
Written Tests	124.91	87.00
Group Debates	101.54	70.00
Personal Interviews	26.07	18.00
Offers	6.49	4.00

Notes: Table 5 reports mean and median counts per job for applications, written tests, group debates, interviews, and offers. Values calculated by averaging student counts across years per job, then averaging across jobs.

3.10 Data on Job Tenure and Promotions

Career outcomes are measured using tenure and promotions, obtained for nearly all (over 98%) placed students using alumni records, also maintained by the placement office and largely curated from an employment-focused platform plus a few other sources (see Online Appendix C for coverage benchmarks). Tenure is calculated using reported employment start and end dates.

Classifying promotions. Promotions are defined as movements to higher job titles within the same firm validated by salary increases and typical career trajectories using external industry-standard benchmarks (such as Glassdoor, levels.fyi) that help distinguish them from lateral moves. Such classifications are particularly reliable for entry-level roles at multinational corporations with globally standardized career ladders and compensation structures (e.g., Bloom et al., 2012; Minni, 2024).

3.11 Data Limitations and Mitigation

Two key limitations include the absence of detailed resumes and screening round performance scores. However, standardized portal data comprehensively capture many relevant student characteristics (e.g., GPA, entrance exam scores, internships), ensuring robust empirical analysis. Strategies addressing the second limitation are discussed in Section 4.7.

4 Penalizing Social Status Through "Fit" Assessments

This section presents six core findings demonstrating how personal interviews penalize social status in entry-level hiring at multinational firms. First, roughly 90% of caste-based disparities in earnings and success rates arise specifically during subjective, "fit"-oriented personal interviews. Second, these disparities are likely conservative. Third, personal interviews are the primary mechanism through which social status is *revealed and penalized.* Fourth, caste penalties consistently emerge at personal interviews across nearly all job positions. Fifth, disadvantaged caste hires exhibit systematically higher promotion rates even in highly selective roles, strongly suggesting hiring disparities result from discrimination rather than unobserved skill differences. Finally, these caste penalties align most closely with *bias* rather than statistical or customer-driven discrimination.

4.1 Main Sample and Reliability of Initial Job Placements

As discussed in Section 2.3, approximately two-thirds of students secure full-time positions through the main placement process. Specifically, fewer than 5% accept return offers from summer internships, while about 60% obtain employment through the standardized four-stage job fair. The remaining one-third participate in the subsequent *aftermarket* hiring phase.[7]

All placed students begin employment in the roles offered. This is further verified through post-graduation surveys conducted by the placement office and cross-validated against external employment platforms, affirming the reliability of placement records as accurate indicators of initial job outcomes.

[7]Of the total 4,164 students, around 159 (4%) accept internship return offers, leaving 4,005 actively participating. Approximately 2,527 students (60% of the total cohort) secure jobs via the four-stage process, totaling 2,686 students (159 internship returns + 2,527 from placement rounds), or about 65% placed through the primary placement process. Consequently, roughly 1,478 students (one-third) enter the aftermarket. See Online Appendix Tables OA.1 and OA.2 for some of these numbers.

4.2 Caste Penalties in Earnings and Hiring Success Rates Emerge at Personal Interviews

Caste disparities in both earnings and hiring success rates predominantly arise after personal interviews. Even uncontrolled regressions show similar patterns (Section 4.7).

Earnings Gap Decomposition. Disadvantaged castes earn significantly lower initial salaries overall: the unconditional earnings gap is approximately 23%, narrowing to roughly 10% after controlling for pre-college skills, college GPA, prior labor market experience, and other factors. This residual 10% gap is about 6,000 (PPP) or approximately 0.18 standard deviations (Table 3).

Because jobs have *fixed, non-negotiable* salaries, caste earnings penalties across screening stages reflect differences in *job access* rather than candidate-specific salary variations (Section 3.6). This enables clear identification of where the caste earnings penalty first appears. Specifically, I quantify these earnings differences using the following regression specification, assigning posted salaries counterfactually to the set of jobs each caste group is in contention for at each stage before final offers are accepted:

$$\log(\text{Avg. Job Salary}_{it}^{\text{Search Stage}}) = \alpha + \beta \times \text{Disadv. Caste}_i + \text{Controls}_i + \gamma_t + \epsilon_{it}, \qquad (1)$$

where $i \in \mathcal{I}$ indexes students, and "Search Stage" $\in$ {Application, Written Tests, Group Debates, Personal Interviews, Offers, Accepted Offers}. Controls include GPA, entrance exam scores, prior internships, major, and degree; γ_t denotes cohort-year fixed effects. For each student, the "average job salary" at a given stage is the mean salary of all roles for which the student remains eligible.

The coefficient β, illustrated in Figure 1, quantifies the caste earnings gap at each recruitment stage. The application reading round shows no statistically significant earnings gap after controls, confirming that broad, low-cost applications (Section 3.8) lead to similar application behavior across castes. The written-test and group-debate rounds likewise reveal no meaningful gaps. Differential caste preferences over offered jobs also contribute negligibly to the total penalty. The earnings penalty emerges sharply at the personal interview stage, accounting for roughly 90% of the total earnings gap observed among graduates in the sample. Full regression results with robustness checks using Oster (2019) bounds to account for selection on unobservables are presented in Online Appendix Table OA.4.[8]

[8] We apply Oster (2019)'s bounding methodology, setting $\delta = 1$ (assuming unobservables match observables in selection intensity) and $R^* = 1.3 \times \tilde{R}$, allowing the total explanatory power to exceed the observed R^2 by 30%. The main findings remain robust under these stringent assumptions and their variations.

Figure 1: **The Caste Earnings Penalty**

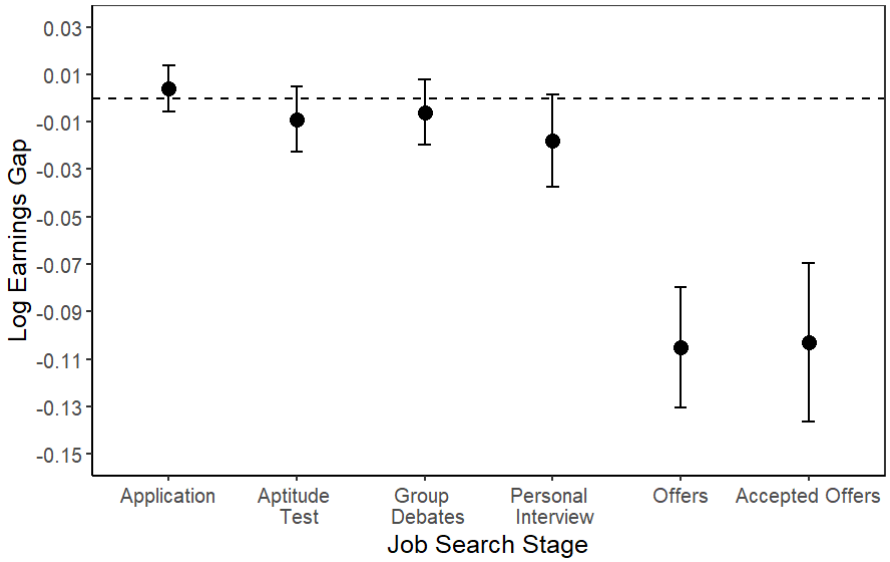

Notes: Figure 1 shows the log earnings gap across castes at each job search stage. Black dots represent coefficients showing the percentage difference in average salary between disadvantaged and advantaged castes, with 95% confidence intervals. Since salaries are identical and non-negotiable for identical jobs, the earnings gap reflects differences in job access at each stage. All regressions include controls.

Success Rate Gaps by Screening Stage. To provide complementary evidence to the earnings analysis, I examine success rates—the probability of progressing from one recruitment stage to the next *within* each job. Because salaries are fixed, a zero earnings gap at stage $t+1$ can occur only if both caste groups have similar odds of advancing at stage t; otherwise they would face different job sets and immediate pay differences. Success-rate regressions therefore furnish a direct consistency check on the earnings decomposition in Figure 1 while also accounting for heterogeneity across jobs.

Table 6 shows negligible caste differences at the application, written-test, and group-debate rounds. The gap opens only after personal interviews, where disadvantaged castes are about 30 percentage points less likely to advance, a drop that explains more than 95% of the overall success-rate disparity and aligns exactly with the interview-driven earnings penalty.[9]

[9]Based on the estimates in Table 6, disadvantaged caste candidates in our sample are about 70% less likely than their advantaged peers to clear the final interview stage. The risk-ratio—defined as the probability of success for disadvantaged castes divided by the probability of success for advantaged castes—is approximately 0.28, indicating disadvantaged candidates have only 28% of the success rate of advantaged candidates. Thorat and Attewell (2007) report an average 50% shortfall in callbacks for disadvantaged groups (risk-ratio ≈ 0.52). Although the point estimates differ, both imply that disadvantaged groups face roughly a one-half to two-thirds reduction in success probability, well within the same ballpark once sampling variation and design differences are taken into account.

Table 6: **Caste Differences in the Success Rate at Each Job Search Stage**

	Success App (1)	Success Test (2)	Success Debate (3)	Success Interview (4)
Disadv. Caste (OLS)	-0.0025*	-0.0002	-0.0005	-0.2938***
	(0.0014)	(0.0018)	(0.0027)	(0.0048)
Disadv. Caste (Oster)	0.009	0.010***	-0.002	-0.296***
	(0.0013)	(0.002)	(0.003)	(0.005)
Mean Success Rate	0.5990	0.8110	0.2580	0.2610
Observations	305,541	183,150	148,489	38,378
Adjusted R^2	0.61069	0.34360	0.01858	0.24450
Year FE	✓	✓	✓	✓
Job FE	✓	✓	✓	✓
Other Controls	✓	✓	✓	✓

Notes: Table 6 reports caste differences in success rates at each job search stage. Each observation is a student-job-year combination for candidates still in contention at that stage. All regressions include job fixed effects and controls. When firms conduct common screening rounds for multiple positions, each position is treated as a separate job. Bias-adjusted coefficients use Oster (2019) with $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$, assuming unobservables match observables in selection intensity and allowing total explanatory power to exceed observed R^2 by 30%. $^*p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

4.3 Robustness and Job-Level Heterogeneity in the Caste Penalty

These findings withstand extensive robustness checks, including a fine-grained *job-level* analysis that treats each employer–position pair as its own fixed effect. Because firms likely vary in written-test content, group debate format, interviewing style, and caste composition, documenting a 30 percentage point caste penalty *within nearly every job* rules out the concern that firm-specific screening methods or other forms of organizational heterogeneity explain the gap. Section 4.7 details these estimates and presents related job-performance evidence based on promotion outcomes.

4.4 Conservative Estimates Due to Exclusion of the Aftermarket

The earnings and success-rate gaps identified above likely represent conservative estimates. The main analysis excludes roughly one-third of students who enter the post-placement *aftermarket* hiring phase (Section 4.1). Disadvantaged caste students entering the aftermarket generally exhibit much *lower* average GPAs than comparable advantaged castes (Online Appendix Table OA.2).[10] Because higher GPA typically correlates with higher-paying and more selective job placements, excluding a subset of lower-GPA disadvantaged students from the main analysis understates the true caste disparities.

[10] *Within* any offered job, disadvantaged hires have slightly higher GPAs than advantaged counterparts, consistent with the caste hiring penalty (Online Appendix Table OA.3).

4.5 Information Observed by Employers at Each Screening Round

Employers continuously access candidates' standardized academic profiles, including GPA, coursework, entrance exam scores, internships, and technical skills, through the centralized online job portal (Section 3.2). While early stages like application screening and written aptitude tests rely primarily on objective metrics (e.g., GPA cutoffs, standardized test scores), group debates and personal interviews introduce direct social interactions.

Group debates involve structured team discussions, allowing employers to assess communication, teamwork, and socio-emotional skills. However, these interactions remain standardized, with limited opportunity for individual background details. In contrast, personal interviews allow deeper, individualized social interactions related to candidates' personal interests, extracurricular activities, background, and interpersonal rapport (Deshpande, 2011; Rivera, 2012, 2015). These social interactions significantly enhance employers' insights into candidates' cultural fit and organizational alignment, making subtle indicators of social status more salient (Bisin and Verdier, 2010). Although we lack data on interviewers' own backgrounds, ethnographic evidence shows that evaluative norms are quickly socialized within elite firms, so managers with diverse origins likely converge on similar notions of fit, reinforcing status-based biases (Turco, 2010). Section 4.6 below provides quantitative evidence supporting our interpretations regarding the role of personal interviews in making social status relatively more salient to employers than previous rounds (qualitative evidence in Online Appendix Section E).

4.6 Interviews as the Main Mechanism for Revealing and Penalizing Social Status

When does caste become visible to employers? This subsection demonstrates quantitatively that personal interviews are where caste is revealed and penalized.

4.6.1 Quantitative Evidence

Although many surnames are caste-indicative, most do not reliably signal caste in elite urban college settings, where students come from diverse regions with varied naming conventions (e.g., Jodkha, 2017; Parmar, 2020).[11] Consequently, only about 15% of students in my sample have clearly identifiable caste surnames. For this subgroup, caste penalties appear earlier—at the application reading round. However, for the remaining 85% with non-identifiable surnames, the caste penalty emerges predominantly during personal interviews (e.g., Rivera, 2012, 2015).

[11] See also Bhupatiraju et al. (2024)—a study of more than one million High Court cases from the urban centre of Patna—which shows that roughly half of petitioners, respondents, and judges use *caste-neutral* surnames (49% of petitioners, 56% of respondents, and 50% of judges), underscoring how common name-based ambiguity is in urban India.

This differential pattern of the emergence of the caste penalty strongly suggests that the caste penalty follows caste revelation. Thus, personal interviews serve as the key mechanism for revealing and penalizing social status.

Caste earnings penalty emerges earlier when surnames reveal caste. Figure 2 replicates Figure 1, separately examining students with and without clearly identifiable caste surnames. Students with identifiable surnames experience a significant penalty at the application reading round, while those without identifiable surnames see penalties emerge only at personal interviews.[12] Corresponding regression results, including Oster's bias-adjusted coefficients, are in Online Appendix Tables OA.7 and OA.8.

Similar patterns in success rates. The same story holds when we look at progression through the screening rounds. Students whose surnames clearly reveal caste suffer large penalties right from the application-reading stage across almost every job (Online Appendix Figure OA.2). By contrast, students whose surnames are *not* caste-identifiable track the overall pattern shown in Table 6: their success rates are essentially equal to advantaged peers until the personal interview, where the penalty appears and drives the final gap.

How are student names associated with caste? Surnames are coded as caste-identifiable versus not by utilizing official government sources, including regularly updated state-specific "schedules" of caste names (the administrative categories "Scheduled Castes" and "Scheduled Tribes" are so named precisely because eligible castes are listed in these government "schedules"; see Puri, 2025). Common advantaged castes surname lists are compiled by the Anthropological Survey of India (Anthropological Survey of India, 1992). Similarly, common disadvantaged castes (SC/ST/OBC) surnames are available through public Election Commission records of reserved constituencies and official state-maintained SC/ST/OBC lists (Election Commission of India, 2024).

Qualitative evidence. Qualitative evidence, consistent with the quantitative patterns above (full details in Online Appendix E), indicates that: (i) visual and auditory cues like skin tone or regional accent do not reliably signal caste in urban, educated settings; and (ii) social interactions such as personal interviews often reveal social status in urban-educated contexts where other signals (e.g., names) are absent.

[12] For example, a purely income-based explanation would predict an early gap in both groups, contrary to what we observe.

Figure 2: **The Caste Earnings Penalty by Subsample**

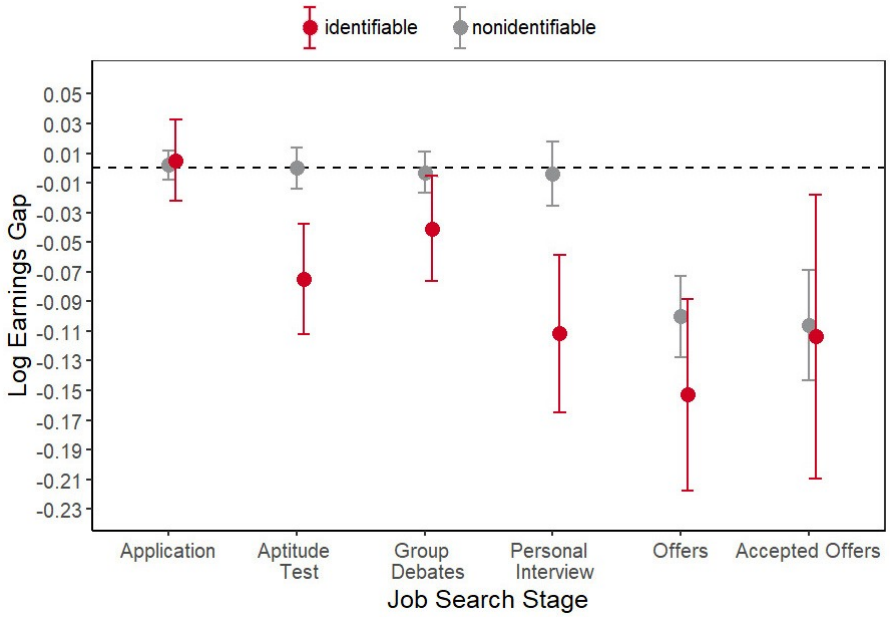

Notes: Figure 2 compares the log earnings gap across castes at each job search stage between students with identifiable (red) and non-identifiable (gray) caste surnames. Both gray and red dots show percentage differences in average salary between disadvantaged and advantaged castes, with 95% confidence intervals. Since salaries are identical and non-negotiable for identical jobs, gaps reflect differences in job access. All regressions include controls.

4.7 Heterogeneity in Caste Penalties, Selectivity, and Promotion Outcomes: Empirical Facts & Stylized Model

We examine *job-level heterogeneity* in the caste penalty to show that disadvantaged-caste applicants experience about a 30 percentage-point lower interview success rate across the *entire* distribution of jobs. Moreover, those who are hired exhibit roughly 20 percent higher promotion rates than their advantaged peers. Since candidate performance scores at each stage are unavailable (Section 3.11), these headline findings raise two critical questions. First, might caste penalties appear earlier in *highly selective* jobs, which recruit primarily from the right tail of the ability distribution? Second, could personal interviews reveal genuine unobserved ability differences—rather than discrimination—that justify hiring disparities across castes?

Key empirical test and intuition. Variation in job selectivity (the depth in the applicant pool a job hires from) offers an empirical test to disentangle these two competing explanations. If disadvantaged castes truly had lower pre-interview performance, such deficits should be particularly evident in highly selective jobs. Yet, pre-interview caste gaps remain negligible even among these selective positions. Furthermore, once hired, disadvantaged castes exhibit higher promotion rates. One might argue this simply reflects inframarginal productivity differences—specifically, a greater share of right-tail performers among disadvantaged hires.

However, examining promotion patterns across jobs of increasing selectivity helps disentangle discrimination from ability differences. If hiring thresholds were identical across castes, disadvantaged hires' promotion advantage should *shrink* as jobs become more selective. Instead, we find this advantage remains stable across all selectivity tiers.

Below, I present two empirical facts leveraging variation in selectivity across jobs followed by a formal model demonstrating that these patterns reflect a systematically stricter hiring threshold for disadvantaged castes.

Fact 1 (Conditional Success Rates Reveal No Early Penalty, but a Large Final Interview Gap). Figure 5 plots, for each job, the conditional success-rate gap (disadvantaged minus advantaged) against the cumulative cut rate at each screening stage (defined in Section 3.7). Negative values indicate lower advancement rates for disadvantaged candidates relative to advantaged candidates. Formally, the gap for job j is the coefficient β_j from the regression:

$$\Pr(\text{Success}_{ijt}^{\text{Job Search Stage}}) = \alpha_j + \beta_j \times \text{Disadv. Caste}_i + \text{Controls}_i + \gamma_t + \epsilon_{ijt}, \tag{2}$$

where $i \in \mathcal{I}$ indexes students and γ_t denotes year fixed effects. Controls include GPA, entrance exam scores, prior internships, major, and degree.

Disadvantaged castes are about 30 percentage points less likely to advance than advantaged castes after personal interviews across all jobs. Before personal interviews—during application reading, written tests, and group debates—the caste gaps remain negligible, even among highly selective jobs. A distinct caste penalty emerges only after personal interviews. *Even uncontrolled regressions demonstrate the same pattern* (see Online Appendix Figure OA.3).

Fact 2 (Disadvantaged Hires Earn Comparably Higher Promotion Rates, Even in Highly Selective Jobs). To further test whether interviews reflect genuine ability or discrimination, I examine post-hire performance outcomes. Alumni records cover 98% of placed graduates, and we use promotion as our performance metric.[13] I estimate caste differences in promotion rates using the following regression across job selectivity tiers:

$$\text{Outcome}_{ijt} = \alpha + \beta \times \text{Disadv. Caste}_i + \text{Controls}_i + k_j + \gamma_t + \epsilon_{ijt}, \tag{3}$$

where $i \in \mathcal{I}$ indexes students, $j \in \mathcal{J}$ indexes jobs, γ_t denotes year fixed effects, and k_j represents job fixed effects. Controls include GPA, entrance exam scores, prior labor market experience, major, and degree. Table 7 presents these results:

- **Columns (1)–(2)** show that, pooling all jobs, disadvantaged hires have approximately

[13] See Section 3.10 for more details and Section 4.1 for details on the placed sample.

a 20% *higher* probability of promotion relative to advantaged peers, despite initial hiring penalties.

- **Columns (3)–(8)** group jobs by *selectivity* (final rejection rates), revealing that disadvantaged hires consistently maintain a stable promotion advantage across the selectivity distribution. For most selectivity tiers, the advantage is approximately 20-25% and statistically significant. In the top 10 percent, the point estimate remains comparable in magnitude to those at other tiers but is imprecise.
- **Columns (9)–(10)** demonstrate consistent patterns in both client-facing and non-client-facing roles, indicating that the promotion advantage is not specific to particular job functions.

Taken together, the coefficients show that disadvantaged hires outperform advantaged peers even after the most stringent selection and regardless of job type.

Promotions as a performance metric. In elite MNCs, entry-level pay is globally fixed, so promotions are the main channel for rewarding productivity; thus they serve as a reliable proxy for job performance (Bloom et al., 2012; Minni, 2024).

Promotion-rate estimates are conservative. When deciding on promotions, local managers often favor subjective traits such as "creativity," "fit,'' and "potential" that typically advantage higher-status candidates (Gorman, 2005; Turco, 2010; Rivera, 2012; Sarsons, 2017; Benson et al., 2019; Sarsons et al., 2021). Caste-aware (or even class-aware) initiatives are scarce: fewer than 10% of major U.S. firms reference social class in DEI statements (Stansbury and Rodriguez, 2024b), and Indian multinationals seldom mention caste (Bapuji et al., 2023). Therefore, the promotion premium for disadvantaged caste hires is likely a lower bound.[14] Appendix Table OA.9 confirms this: advantaged-caste job leavers have weaker GPAs, while promotions correlate positively with GPA.

4.7.1 Explaining the Empirical Facts: A Simple Model of Late-Stage Discrimination

This section introduces a unified theoretical model demonstrating how (i) the absence of early-stage gaps for disadvantaged castes implies no genuine pre-interview ability deficit, and (ii) variation in hiring thresholds across jobs helps distinguish discrimination from genuine ability differences revealed at the interview stage.

Consider a set of jobs $j \in \mathcal{J}$. Each job has up to four sequential rounds $s \in \{1, 2, 3, 4\}$:

[14] In a 2018 U.S. survey, 67% of disadvantaged-caste employees reported workplace mistreatment (California v. Cisco Systems, Inc., 2022).

- Rounds $s = 1, 2, 3$ are pre-interview screening rounds (e.g. application, written tests, group debates)
- Round $s = 4$ is the personal interview

Key notation.

- $j \in \mathcal{J}$: index for jobs.
- $s \in \{1, 2, 3, 4\}$: screening round ($s = 1, 2, 3$ pre-interview; $s = 4$ interview).
- $g \in \{0, 1\}$: caste group (0 = advantaged, 1 = disadvantaged).
- $M_{s,i,g} \sim F_g$: candidate i's score in round s, group g.
- $\tau_{s,j}$ (or $\tau_{4,j}^{(g)}$ at $s = 4$): round-specific cut-off for job j (may differ by g only at $s = 4$).
- $\Delta_{s,j} = P(\text{pass}_s | g = 1) - P(\text{pass}_s | g = 0)$: within-job caste pass-rate gap at round s.

Groups and Screening. There are two castes (groups) $g \in \{0, 1\}$, where $g = 1$ denotes disadvantaged castes and $g = 0$ denotes advantaged castes. At each round s, employers observe a candidate-specific score $M_{s,i,g}$, drawn from a common across-jobs distribution F_g for candidate i in group g. Using residualized scores leaves all results unchanged.[15]

A candidate passes round s if and only if their score satisfies $M_{s,i,g} \geq \tau_{s,j}$ and if they have successfully passed all preceding rounds. At the final interview stage ($s = 4$), the threshold $\tau_{4,j}^{(g)}$ may differ by group g, reflecting the fact that caste becomes salient to the employer for most students at this stage (Section 4.6.1). Thus, a final-stage caste gap is indicative of discrimination if $\tau_{4,j}^{(1)} > \tau_{4,j}^{(0)}$.

Perfect information. For simplicity, we assume employers observe these scores *perfectly* at each stage. This assumption simplifies exposition but is not essential for our key insights: even under imperfect information, variation in screening selectivity and hiring thresholds across jobs allows us to distinguish between discrimination and genuine ability differences (see Online Appendix Section G for details).

[15] We assume the performance distributions for each group are continuous with smooth densities, finite moments, and satisfy standard regularity conditions. These technical assumptions ensure well-behaved pass rates and expectations when calculating promotion probabilities. While these properties underpin our formal analysis, they are not critical to understanding the main conceptual results.

Monotonicity of thresholds. We assume that thresholds are monotonically increasing across screening rounds:

$$\tau_{4,j} \;\geq\; \tau_{3,j} \;\geq\; \tau_{2,j} \;\geq\; \tau_{1,j},$$

implying each successive round is at least as demanding as the previous one.

Assumption 1 (**Common Talent-Pool**)**.** All jobs draw candidates from the same underlying ability distributions F_g for each caste g. Hence any difference in interview cut-offs $\tau_{4,j}$ across jobs reflects genuine variation in job selectivity, not systematic caste-ability differences. Basically, all firms scoop talent from one tub; a smaller scoop (higher cutoff) just digs deeper into the same right tail.

Empirical support.

1. *No observed caste gap pre-interview (strongest test).* Across application reading, tests, and debates—even in the most selective jobs—advancement rates are statistically indistinguishable by caste (Fact 1). These early rounds already probe the right tail of the ability distribution; any true caste-ability gap would surface here first, but we see none, providing the clearest validation of the common-pool assumption.

2. *Quota-driven representation and GPA overlap.* Caste quotas are mandated to equalize representation *within each major*, and the two castes overlap in every GPA decile (Sections 3.4–3.5; Online Appendix Figure OA.1). Thus the right tails of F_0 and F_1 are populated in every field.

3. *Restricted pre-interview sorting.* Near-zero application costs, broad eligibility rules, a standardized portal, and the invisibility of caste for most students at early rounds (Section 4.6.1) leave employers little scope to screen by caste before interviews (Section 3.8).

Implication. Across jobs, the full vector of round-specific thresholds $\boldsymbol{\tau}_j$ varies for reasons unrelated to caste ability; in particular, the common component of the interview cut-off $\tau_{4,j}$ is an *exogenous shifter* of selection stringency, while any *between-caste gap* $\tau_{4,j}^{(1)} - \tau_{4,j}^{(0)}$ signals discrimination.

Assumption 2 (**Light-Tails**)**.** For each caste $g \in \{0, 1\}$, the upper tail of the screening-score distribution decays at least exponentially: there exist constants $C_g > 0$ and $\alpha > 1$ such that

$$\Pr\big(M_{s,i,g} \geq m\big) \;\leq\; C_g\, e^{-\alpha m} \qquad \text{for all sufficiently large } m.$$

All of our performance measures—test scores, written assessments, and interview ratings—are either explicitly bounded or effectively capped in practice; extreme values are therefore rare, making a light-tailed assumption natural in this context.

(A) Ruling Out Systematic Performance Differences Before Interviews

We first rule out systematic caste differences in candidate performance at pre-interview screening rounds.

Proposition 1 (**No Early Gap Implies No Right-Tail Ability Deficit**)**.** Consider a job j that requires up to four rounds of screening, with $s = 1, 2, 3$ the pre-interview stages and $s = 4$ the personal interview. Suppose jobs draw candidates from the same caste–specific ability distributions F_g (Assumption 1) and those distributions have light tails (Assumption 2). Raising the cut-off $\tau_{s,j}$ at any pre-interview stage $s < 4$ therefore selects further into the *common* right tail of F_0 and F_1. If the disadvantaged caste ($g = 1$) were genuinely weaker in that tail, then for a sufficiently high threshold $\tau_{s,j}$ the within-job pass-rate gap

$$\Delta_{s,j} \;=\; P\big(\text{pass}_{s,i,j} \,|\, g = 1\big) \;-\; P\big(\text{pass}_{s,i,j} \,|\, g = 0\big)$$

would become negative. Conversely, observing $\Delta_{s,j} \geq 0$ for *every* job, even when cut-offs are very high, rules out any meaningful right-tail ability deficit for the disadvantaged caste prior to the interview stage.

Sketch of Proof. We illustrate how variation in screening thresholds helps identify systematic performance differences at pre-interview rounds. Figure 3 contrasts two scenarios. Blue curves denote advantaged candidates ($g = 0$); red curves denote disadvantaged candidates ($g = 1$). The shaded mass to the right of each vertical line is the share of candidates who clear a job-specific cut-off $\tau_{s,j}$.

- **Panel (a): Hypothetical right-tail deficit.** Under the light-tail assumption (Assumption 2), raising $\tau_{s,j}$ slides the vertical line rightward into the common tail described in Assumption 1. If the disadvantaged caste had a thinner tail ($C_1 < C_0$ or a larger shape parameter α), its shaded area would shrink *faster*, making $\Delta_{s,j} = P(\text{pass} \mid g = 1) - P(\text{pass} \mid g = 0) < 0$ and ever more negative as selectivity rises.
- **Panel (b): What we observe.** Fact 1 shows $\Delta_{s,j} \approx 0$ at every pre-interview stage—even for the highest cut-offs—so no such deficit appears in the data.

Because, with a common talent pool (Assumption 1), higher $\tau_{s,j}$ merely digs deeper into that shared tail, the persistent non-negative gaps rule out any caste-specific performance deficit before interviews. Thus, any observed late-stage gap at the interview stage ($s = 4$) must stem either from newly revealed ability signals during interviews or from discrimination based explicitly on social status.

Figure 3: **Using Variation in Screening Thresholds to Test for Pre-Interview Performance Differences**

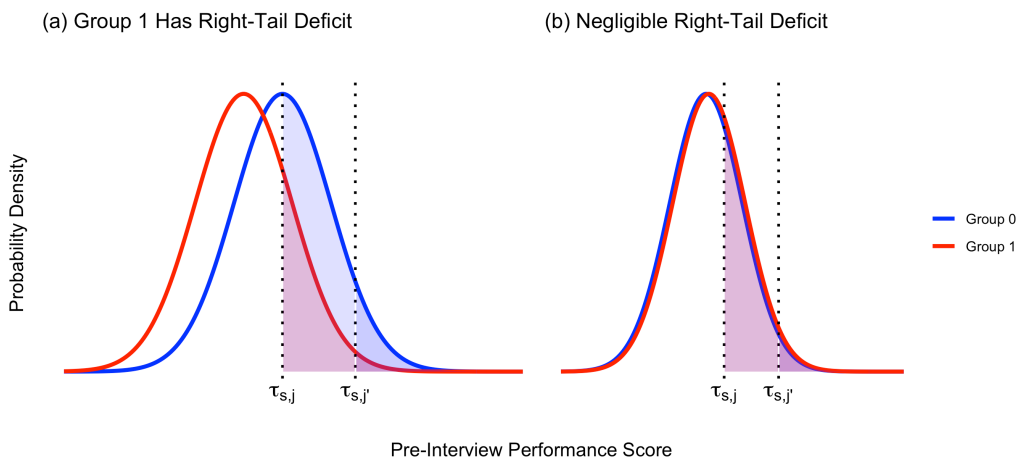

Notes: Figure 3 illustrates how selective screening would reveal systematic pre-interview performance differences. Group 1 denotes disadvantaged castes and Group 0 denotes advantaged castes. Panel (a) shows how a hypothetical right-tail deficit for Group 1 would become more apparent at higher cutoffs. Panel (b) shows the likely distribution of pre-interview performance scores implied by Fact 1, with no meaningful deficit. Shaded regions show candidates who pass the thresholds (blue for Group 0, red for Group 1, purple where they overlap). $\tau_{s,j}$ denotes screening threshold at stage $s < 4$ for job j (j' represents a more selective job).

□

(B) Using Variation in Hiring Cutoffs Across Jobs to Test for Discrimination

Intuition. Across jobs, interview cut-offs $\tau_{4,j}$ vary greatly: prestigious roles set the bar deep in the shared right tail of the talent pool, while less selective jobs admit lower scores. Because the score distribution is light-tailed (Assumption 2), the *average* score of those hired in any group converges to the *marginal* score (the cut-off itself) as $\tau_{4,j}$ moves rightward. Hence, if both castes faced the *same* cut-off within each job, raising $\tau_{4,j}$ would lift the average ability of *both* hired groups toward the same limit, and their post-hire promotion rates would converge. In the data, however, the disadvantaged caste retains roughly a 20% promotion edge even in the most selective jobs (Fact 2). The only way to reconcile this stability with a common talent pool (Assumption 1) is that disadvantaged applicants must clear a uniformly higher bar than their advantaged peers.

Framework. Let job performance be the average of a candidate's four screening scores,

$$\bar{M}_{i,g} \;=\; \tfrac{1}{4}\sum_{s=1}^{4} M_{s,i,g},$$

a monotonic proxy that suffices for our arguments. Candidate i is promoted if $\bar{M}_{i,g} \geq \kappa_j$, where κ_j is the job–specific promotion bar. For hiring, distinguish two cut-offs:

- *Selectivity-tier cut-off* $\tau^{q}_{4,j}$ – the common interview bar applied to *both* castes for all jobs in the top $(1-q)$ selectivity fraction (e.g. $q = 0.8$ means the top 20% most selective jobs).
- *Group-specific cut-off* $\tau^{(g)}_{4,j}$ – the actual bar that group g faces in job j.

Under no discrimination $\tau^{(1)}_{4,j} = \tau^{(0)}_{4,j} = \tau^{q}_{4,j}$; discrimination implies $\tau^{(1)}_{4,j} > \tau^{(0)}_{4,j}$. Let

$$\bar{Y}^{(g)}_{j} \;=\; \mathbb{E}\big[\bar{M}_{i,g} \,\big|\, M_{4,i,g} \geq \tau^{(g)}_{4,j}\big]$$

be the average screening score of hires from caste g, and define the promotion gap in selectivity tier q as

$$\Delta^{\text{Prom}}(q) \;=\; \bar{Y}^{(1)}_{j} - \bar{Y}^{(0)}_{j}.$$

Under the common-pool assumption (Assumption 1), the job-specific interview cutoff $\tau_{4,j}$ plays the role of an *instrument*: it is an exogenous knob that tightens selection stringency (*relevance*) yet affects promotion only through that selection rule (*exclusion*). As $\tau_{4,j}$ moves rightward in more selective jobs, the average ability of hires in each caste converges to the cutoff. Therefore, if both castes face the same threshold, $\Delta^{\text{Prom}}(q) \to 0$; a persistent limit $\delta > 0$ can arise only if disadvantaged applicants clear a uniformly higher bar, $\tau^{(1)}_{4,j} = \tau^{(0)}_{4,j} + \delta$. Proposition 2 formalises this test.

Proposition 2 (Stable Promotion Gaps Across Selectivity Tiers Indicate Discrimination). Fix a selectivity tier q and let $\tau^{q}_{4,j}$ denote the interview cut-off that would apply to *both* castes in the absence of discrimination. Let the actual group-specific cut-offs be $\tau^{(g)}_{4,j}$ and define $\bar{Y}^{(g)}_{j}$ and $\Delta^{\text{Prom}}(q)$ as above.

1. *No discrimination.* If $\tau^{(1)}_{4,j} = \tau^{(0)}_{4,j} = \tau^{q}_{4,j}$ for every job, then under the common-pool and light-tail assumptions (Assumptions 1–2) the mean hired score for each caste converges to the common margin as selectivity tightens, so

$$\Delta^{\text{Prom}}(q) \;\to\; 0 \quad \text{as } q \uparrow 1.$$

2. *Discrimination.* If the disadvantaged caste faces a uniformly higher bar, $\tau^{(1)}_{4,j} = \tau^{(0)}_{4,j} + \delta$ with $\delta > 0$, then the same limit argument gives

$$\Delta^{\text{Prom}}(q) \;\to\; \delta > 0 \quad \text{for all high-selectivity tiers } q.$$

Therefore, a promotion gap that remains positive as jobs grow more selective can only be reconciled with a higher interview threshold for the disadvantaged caste.

Sketch of Proof. Figure 4 contrasts two cases. Blue curves are advantaged ($g = 0$), red curves disadvantaged ($g = 1$); the shaded area right of the vertical line is the share of hires at cut-off $\tau_{4,j}$.

- *Common cut–off (no discrimination).* With $\tau_{4,j}^{(1)} = \tau_{4,j}^{(0)} = \tau$, Assumptions 1–2 imply

$$\bar{Y}_j^{(g)} = \tau + \frac{1}{h(\tau)}, \qquad h(\tau) = \frac{f(\tau)}{1 - F(\tau)} \xrightarrow[\tau \to \infty]{} \infty.$$

  Hence $\bar{Y}_j^{(g)} - \tau \to 0$ for both castes, so $\Delta^{\text{Prom}}(q) \to 0$ as jobs become more selective (derivation in Online Appendix Section F).

- *Higher cut–off for the disadvantaged caste (discrimination).* If $\tau_{4,j}^{(1)} = \tau_{4,j}^{(0)} + \delta$ with $\delta > 0$, the same limit gives

$$\bar{Y}_j^{(1)} - \bar{Y}_j^{(0)} \;\to\; \delta,$$

  so the promotion gap remains positive no matter how high $\tau_{4,j}$ is pushed.

Because these limits hinge only on the cutoff, the conclusion is unaffected by inframarginal hires who sit well above the bar.

Figure 4: **Using Variation in Hiring Thresholds To Test for Discrimination**

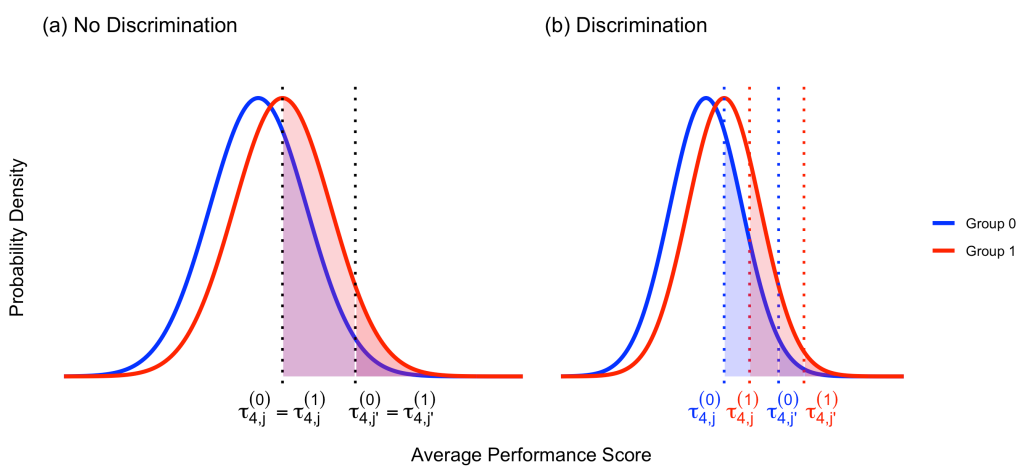

Notes: Figure 4 depicts two scenarios. Group 1 denotes disadvantaged castes and Group 0 denotes advantaged castes. In panel (a), a common cutoff for both groups shrinks their hired-performance gap at higher cutoffs. In panel (b), a constant δ gap in cutoffs preserves a stable performance advantage for Group 1. Shaded regions show candidates who pass the thresholds (blue for Group 0, red for Group 1, purple where they overlap). Here, $\tau_{4,j}^{(g)}$ denotes group g's interview threshold at job j, with j' representing a more selective job.

□

Figure 5: **Caste Penalty in the Conditional Success Rate At Each Screening Round**

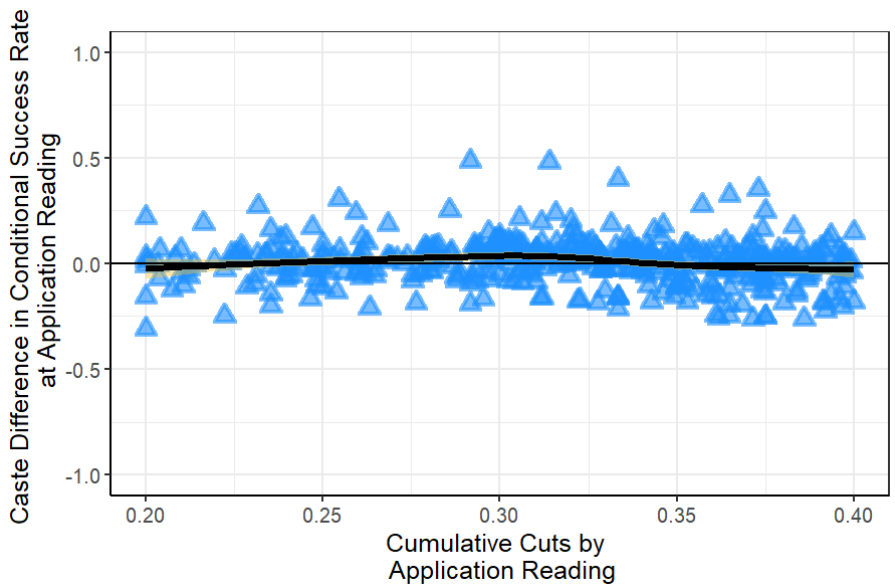

(a) Caste Penalty in Conditional Success Rates versus Cumulative Cuts by Application Reading

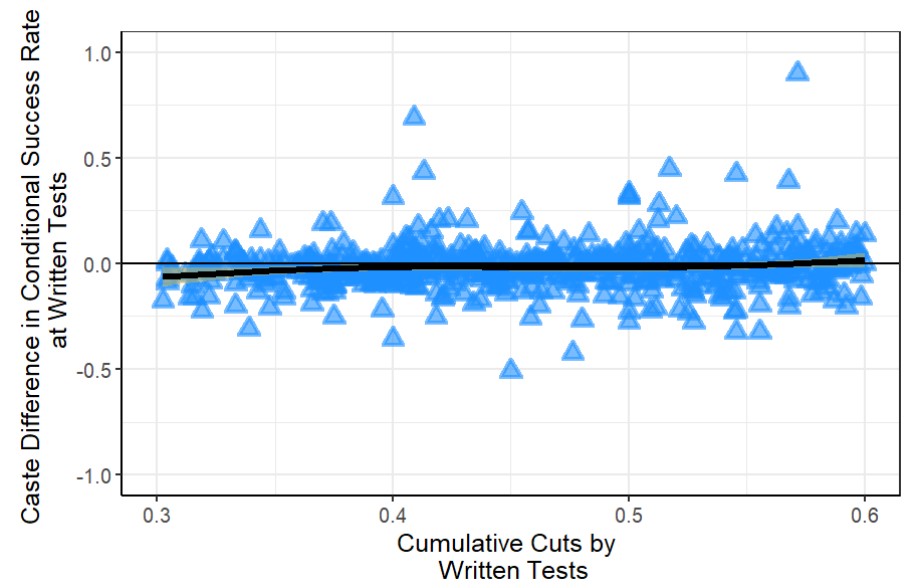

(b) Caste Penalty in Conditional Success Rates versus Cumulative Cuts by Written Tests

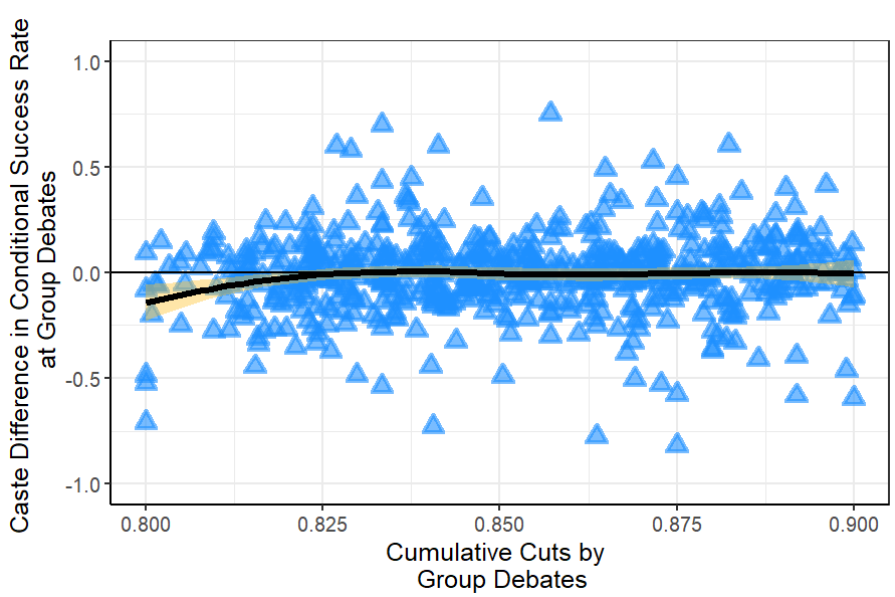

(c) Caste Penalty in Conditional Success Rates versus Cumulative Cuts by Group Debates

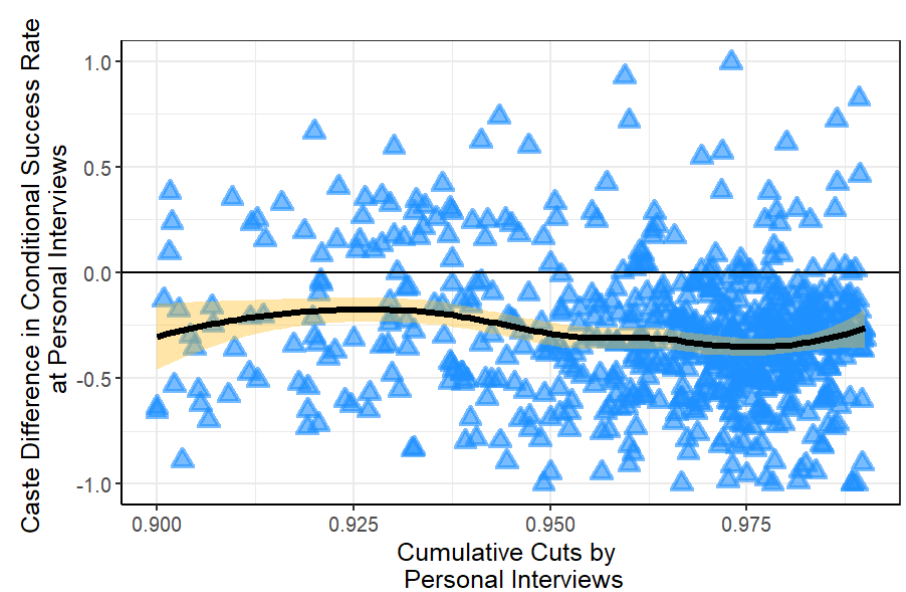

(d) Caste Penalty in Conditional Success Rates versus Cumulative Cuts by Personal Interviews

Notes: Figure 5 plots caste differences in conditional success rates against cumulative cuts at each screening round. Blue triangles represent individual jobs, with smoothed conditional means (black lines) and 95% confidence intervals (yellow). The x-axis shows the cumulative proportion of students eliminated by each round; the y-axis shows the caste coefficient from job-specific regressions with controls (degree, major, GPA, entrance exam scores) and cohort-year fixed effects. The horizontal line at zero indicates no caste difference.

Table 7: **Caste Differences in Promotion Rate**

	Dependent variable:										
	Promoted	Promoted	Promoted at Top 75%	Promoted at Top 50%	Promoted at Top 40%	Promoted at Top 30%	Promoted at Top 20%	Promoted at Top 15%	Promoted at Top 10%	Promoted at Client Facing	Promoted at Non-Client Facing
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)	(9)	(10)	(11)
Disadv. Caste (OLS)	0.137***	0.096***	0.104***	0.112***	0.136***	0.147***	0.190**	0.204**	0.084	0.107**	0.082**
	(0.025)	(0.026)	(0.029)	(0.038)	(0.043)	(0.054)	(0.073)	(0.084)	(0.109)	(0.044)	(0.032)
Disadv. Caste (Oster)	0.142***	0.093***	0.108***	0.122***	0.141***	0.157**	0.196**	0.206**	0.074	0.106***	0.075**
	(0.026)	(0.028)	(0.031)	(0.037)	(0.041)	(0.058)	(0.082)	(0.082)	(0.131)	(0.039)	(0.036)
Tenure Duration		0.084***	0.082***	0.076***	0.074***	0.068***	0.054*	0.079**	0.115**	0.065***	0.097***
		(0.012)	(0.014)	(0.019)	(0.021)	(0.025)	(0.030)	(0.037)	(0.049)	(0.021)	(0.015)
Mean Y	0.684	0.684	0.667	0.668	0.676	0.711	0.711	0.706	0.730	0.687	0.682
Observations	2,489	2,489	2,012	1,246	997	748	502	374	252	835	1,654
Adjusted R^2	0.183	0.205	0.203	0.216	0.211	0.182	0.192	0.189	0.127	0.194	0.214
Job FE	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
Year FE	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
Other Controls		✓	✓	✓	✓	✓	✓	✓	✓	✓	✓

Notes: Table 7 reports caste differences in promotion rates for graduates who obtained jobs through the main placement process. Column (1) includes only job and year fixed effects. Column (2) adds controls for GPA, major, degree, and job tenure. Columns (3)-(8) show results by job selectivity tiers, measured by cumulative applicant elimination rates. Columns (9)-(10) compare client-facing versus non-client-facing roles. Bias-adjusted coefficients use Oster (2019) with $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$, assuming unobservables match observables in selection intensity and allowing total explanatory power to exceed observed R^2 by 30%. $^*p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

4.8 Classification and Preferred Interpretation of the Caste Penalty: Distinguishing from Statistical or Customer Discrimination

Given the evidence that employers directly penalize caste (Section 4.6), the most plausible interpretation is *disparate treatment* or explicit discrimination based on caste. The disparities are most consistent with *bias*. Three facts jointly support this view. First, disadvantaged castes face not only lower hiring rates but penalties that are remarkably consistent across the *entire* job distribution (Fact 1). Second, disadvantaged hires have persistently higher—and stable—promotion rates across all job selectivity tiers (Fact 2). Third, disadvantaged caste hires are more likely to be promoted in both client-facing and non-client-facing jobs (Table 7). Taken together, these findings are most consistent with bias rather than accurate statistical inference about productivity[16] or customer-driven discrimination. Although subjective criteria for assessing "fit" might inadvertently create a disparate impact (e.g., by favoring candidates with particular extracurricular backgrounds), the evidence presented throughout Sections 4.2–4.7 nonetheless points most strongly to *disparate treatment driven by bias*.

4.9 What Happens During Personal Interviews?

We do not directly observe interview transcripts, so we cannot parse the exact conversations. However, prior ethnographic research on elite hiring demonstrates that personal interviews at multinational corporations place substantial weight on subjective markers of "fit" or "culture match" that are closely tied to candidates' socioeconomic backgrounds rather than strictly job-relevant skills (e.g., Rivera, 2012, 2015; Lamont, 2009). A richer qualitative account of how interviewers use cultural cues during such social interactions, complete with illustrative quotations, is provided in Appendix H.

5 Alternative Explanations

Below, I briefly summarize alternative explanations for the observed caste disparities emerging specifically during personal interviews.

1. **Early-stage inclusivity for optics.** Pre-interview pass rates differ by less than one percentage point (Table 6; Figure 5), and surname-identifiable applicants already face a substantial penalty at the initial application screening round (Figure 2; Online Appendix Figure OA.2). There is weak institutional pressure on private employers to promote caste diversity: unlike the U.S., India lacks a federal ombudsman such as the Equal Employment

[16] An accurate statistical discrimination model typically assumes evaluators hold correct beliefs about average group-based productivity and that the underlying productivity distributions are relatively well-behaved (Aigner and Cain, 1977).

Opportunity Commission to regulate private-sector hiring. Moreover, caste, class, or other forms of socioeconomic status indicators are *not* diversity metrics that multinational firms typically monitor, with race and gender being the main focus areas for diversity in such firms (Stansbury and Rodriguez, 2024a; Stansbury, 2024). In fact, as recently as 2018, only 3% of India's top 100 listed firms even collected caste data (BusinessLine, 2018), with many Indian multinational companies explicitly viewing caste-targeted hiring as "anti-merit" (Jodhka and Naudet, 2019).

2. **Delayed discrimination (strategic optics).** Employers have no apparent reason to delay caste discrimination: when caste is identifiable through surnames, penalties emerge immediately at the application screening stage (Figure 2; Online Appendix Figure OA.2). If employers wished to exercise caste discrimination before interviews, the highly subjective group debate rounds—where employers have considerable discretion in candidate evaluation—provide ample opportunity. Yet no caste gaps emerge in these rounds.

3. **Ability gaps (hidden or cumulative).** If disadvantaged candidates truly lagged—either on specific right-tail skills or on *cumulative* pre-interview performance—the deficit would surface first (and most sharply) in the most selective jobs. Yet no caste penalty appears in those rounds (Fact 1), and hired disadvantaged candidates subsequently outperform peers across all jobs (Fact 2).

4. **Socio-emotional skills.** Group debates—explicitly designed to evaluate communication, teamwork, and other socio-emotional skills—reveal no caste differences (Panels (c) of Figures 5–OA.3).

5. **Attrition and external search.** Once students register and submit applications through the centralized portal, virtually no one exits the placement process for non-elimination reasons. Exits occur only through formal elimination decisions at each of the four strictly enforced screening rounds, with the placement office meticulously recording every advance-or-exit decision (Sections 3.1 and 3.3). Return offers from internships are rare (Section 2.3), and the placement office prohibits registered candidates from conducting parallel off-campus searches (cf. Central Placement Cell, Delhi University), leaving little scope for differential outside options to drive the observed gap.

6. **Missing data or skipped rounds.** For students participating in the placement process, the placement office enforces mandatory attendance at every stage and strictly prohibits employers from skipping steps or altering procedures mid-process. Comprehensive exit surveys conducted by the placement office independently corroborate the completeness of official records (Sections 3.1 and 3.3; Section 4.1).

7. **Negotiation and salary differences.** Salaries and benefits are ex-ante standardized, strictly non-negotiable, and verified by the placement office through comparing offer letters with the salaries advertised by firms (Section 2.3; Section 3.6; Section 4.1).

8. **Public-sector competition.** Quota-regulated government jobs account for only about 2% of all campus job listings (Table 3) and thus cannot meaningfully affect the observed caste disparities. All job listings are visible to both candidates and recruiters through a centralized placement portal, ensuring private-sector employers are fully aware of the limited public-sector competition. Moreover, only the most prestigious public-sector positions typically recruit at elite colleges.

9. **Preferences over jobs.** Near-zero application costs lead students to apply broadly (Section 3.8), via the centralized portal (Section 3.2). In addition, caste differences in preferences over amenities can at best explain only a small share of the initial earnings gap and none of the sharp interview-stage penalty (Figure 1).

10. **Preferential promotion for disadvantaged hires.** Given the rarity of caste-focused corporate initiatives (Jodkha, 2017; BusinessLine, 2018; Chakravartty and Subramanian, 2021), and extensive survey evidence pointing to widespread bias *against* disadvantaged caste employees rather than in their favor (California v. Cisco Systems, Inc., 2022)[17], it is unlikely that the promotion results are confounded by policies explicitly favoring disadvantaged castes. Although a few multinational corporations initiated efforts to address caste discrimination in late 2022 and 2023, these changes remain rare (Bapuji et al., 2023).

11. **Alternative job-ladder explanation.** Promotion comparisons analyze identical positions (same job designation, same firm; see Section 2.2) and multinationals typically standardize promotion tracks in the same position (e.g., Bloom et al., 2012).

12. **Negotiation and salary differences.** Salaries and benefits are ex-ante standardized, strictly non-negotiable, and verified by the placement office by comparing offer letters with the salaries advertised by firms (Section 2.3; Section 3.6; Section 4.1).

13. **Public-sector competition.** Quota-regulated government jobs account for only about 2% of all campus job listings (Table 3) and thus cannot meaningfully affect the observed caste disparities. All job listings are visible to both candidates and recruiters through a centralized placement portal, ensuring private-sector employers are fully aware of the

[17] In a 2018 survey of South Asian employees in US workplaces, 67% of disadvantaged caste respondents reported experiencing unfair treatment (California v. Cisco Systems, Inc., 2022).

limited public-sector competition. Moreover, only the most prestigious public-sector positions typically recruit at elite colleges.

14. **Preferences over jobs.** Near-zero application costs lead students to apply broadly (Section 3.8) via the centralized portal (Section 3.2). In addition, caste differences in preferences over amenities can at best explain only a small share of the initial earnings gap and none of the sharp interview-stage penalty (Figure 1).

15. **Preferential promotion for disadvantaged hires.** Given the rarity of caste-focused corporate initiatives (Jodkha, 2017; BusinessLine, 2018; Chakravartty and Subramanian, 2021), and extensive survey evidence pointing to widespread bias *against* disadvantaged caste employees rather than in their favor (California v. Cisco Systems, Inc., 2022)[18], it is unlikely that the promotion results are confounded by policies explicitly favoring disadvantaged castes. Although a few multinational corporations initiated efforts to address caste discrimination in late 2022 and 2023, these changes remain rare (Bapuji et al., 2023).

16. **Alternative job-ladder explanation.** Promotion comparisons analyze identical positions (same job designation, same firm; see Section 2.2) and multinationals typically standardize promotion tracks (among other things) in the same position (e.g., Bloom et al., 2012).

6 External Validity

How generalizable are findings from one elite Indian institution? While caste discrimination might initially appear uniquely Indian, five considerations underscore why these results hold broader implications, not only across India but globally in elite hiring and other high-stakes selection contexts.

A uniquely comprehensive social status label. India's constitutionally recognized caste categories provide a formally codified social hierarchy, offering precise measurement clarity typically unavailable in other contexts where class discrimination must be inferred from ambiguous signals like education, accents, or networks (Rivera, 2015; Laurison and Friedman, 2019). Because caste is explicit rather than inferred, we can pinpoint *when* and *where* "fit"-based heuristics disadvantage lower-status candidates in elite hiring.

[18] In a 2018 survey of South Asian employees in US workplaces, 67% of disadvantaged caste respondents reported experiencing unfair treatment (California v. Cisco Systems, Inc., 2022).

Mirroring global elite hiring pipelines. The hiring process we analyze—résumé reading, preliminary screens (e.g., written tests), and discretionary final interviews—is standard screening procedure at elite jobs worldwide (Rivera, 2015). Multinational corporations use identical evaluation protocols across offices (Bloom et al., 2012; Minni, 2024), while elite employers recruiting through on-campus job fairs at top Indian universities employ similar hiring processes (Section 2.4). By pinpointing class bias at the final interview stage, our evidence speaks directly to mainstream hiring at global consultancies, tech firms, and banks (Rivera, 2012).

A globally competitive talent pool. The students in our sample match their Ivy-Plus counterparts in human-capital quality. Using harmonized alumni earnings metrics for nearly 2,800 institutions around the world, Martellini et al. (2024) place several elite Indian universities alongside leading global universities, including many Ivy League institutions. Because our focal college draws from this applicant pool, the labor market segment we study is comparable to the world's most selective universities.

Subjective "fit" screens reproduce class disparities worldwide. Across settings, discretionary assessments tend to reward higher-status candidates' cultural capital. As explicit status markers fade in diverse urban labor markets (Jodkha, 2017; Clark, 2015), these subjective evaluations become critical gatekeepers. Evidence from U.S. and U.K. professional hiring shows that brief, rapport-heavy interviews amplify class gaps by privileging mannerisms, conversational style, and extracurriculars that signal organizational fit (Rivera, 2012; Laurison and Friedman, 2019; Stansbury and Rodriguez, 2024a). India's elite civil service displays the same pattern: disadvantaged-caste officers who scored high enough to qualify without quotas yet received markedly lower "personality" interview ratings subsequently outperform advantaged peers by 18–20% (Bhavnani and Lee, 2021). U.S. college admissions reveal parallel dynamics, as "holistic" file reads involving personal essays, recommendation letters, extracurricular narratives systematically disadvantage high-achieving, lower-SES applicants (Sacerdote et al., 2025). Together, these findings confirm that status disparities arising at subjective gatekeeper stages are a structural feature of elite selection globally and, in my setting, those disparities are best explained by bias.

Downstream erosion of upward-mobility policies. Policies widening entry at the *front end* of elite pipelines often lose traction *downstream*. In India, caste-based admissions quotas boost college representation, yet discretionary hiring screens at multinationals erode these gains, as this study documents. The U.S. shows parallel patterns: high-income graduates are far more likely to convert prestigious college degrees into elite careers through exclusive networks (Michelman et al., 2022), while low-SES students face persistent gaps in reaching leadership

roles despite academic success (Zimmerman, 2019). These patterns reveal that admissions reforms alone are insufficient.

7 Conclusion

Personal interviews for elite jobs that assess "fit" afford recruiters the opportunity to systematically exclude highly talented candidates based on their social status. This paper traces how such subjective evaluations create barriers that perpetuate inequality while undermining efficient talent allocation. By examining all recruitment stages at the India operations of multinational corporations—from initial applications through final hiring decision—I provide the first large-scale empirical evidence on where and how class-based discrimination emerges in elite hiring and its magnitude. Three findings stand out. First, nearly 90% of caste-based disparities arise during subjective personal interviews, not during earlier, relatively more standardized assessments like résumé screening and written tests. Second, disadvantaged candidates who are hired despite these barriers outperform their advantaged peers by 20% in promotion rates across all jobs, including client- and non-client-facing roles. Third, the hiring penalty is most consistent with bias.

The Indian setting provides unusual clarity for investigating social status discrimination. Although caste is uniquely explicit, the underlying mechanism—subjective evaluations of who "belongs"—operates globally. As explicit class markers fade, interviewers increasingly rely on subtle cues like conversational style, extracurriculars, and shared experiences (Rivera, 2015; Laurison and Friedman, 2019). Dual-process cognition theory explains the psychological mechanisms underlying this pattern: faced with ambiguous hiring decisions, evaluators default to rapid, intuitive "System 1" judgments about cultural similarity rather than slow, deliberative assessment of competence (Kahneman, 2011). These snap impressions, shaped by superficial rapport (Todorov, 2017), implicit stereotypes (Knobe and Nichols, 2008), and immediate affective reactions (Pinker, 2021), predict performance no better than chance yet consistently favor candidates from familiar backgrounds (Highhouse, 2008; Dana et al., 2013; Rivera, 2015).

The problem extends beyond individual bias to structural incentives. Even beyond this study's setting, hiring managers at multinationals rarely supervise the candidates they select (Cappelli, 2019), creating a fundamental agency problem: those making selection decisions bear little to no consequences for poor choices or for overriding more standardized assessments (Hoffman et al., 2018). Consider elite firms where reaching the final interview signals strong qualifications: precisely here, where objective differences are minimal, subjective "culture match" or "fit" assessments disproportionately exclude high-potential candidates from unfamiliar backgrounds. When prestigious firms compete for talent, managers naturally favor candidates who mirror their own backgrounds (Rivera, 2012; Laurison and Friedman, 2019), often from mis-

aligned incentives that reward cultural homogeneity over performance potential (Hoffman and Stanton, 2024; Benson et al., 2023).

These findings expose a critical flaw in diversity efforts: gains from affirmative action in education erode when graduates face subjective hiring screens. Even as elite universities diversify their student bodies, employers inadvertently restore homogeneity by selecting for cultural fit (Rivera, 2012, 2015). This pattern reflects deeper organizational pathologies: hiring managers, rarely accountable for long-term outcomes, prioritize immediate comfort over performance potential (Hoffman and Stanton, 2024; Benson and Shaw, 2024).

This study captures only entry-level hiring and the first few years of career outcomes. While the promotion results provide strong evidence of misallocated talent, they capture only the initial stages of career trajectories. Future research should examine longer-term outcomes tracking whether performance advantages persist after five or ten years, how disadvantaged hires navigate mid-career transitions, and whether early hiring discrimination compounds over time through network effects. Additionally, important questions remain about how these patterns manifest across different organizational settings, educational tiers, and national contexts with varying histories of social stratification.

Interviews can still add genuine value when context and design align. Our results do *not* imply that every personal interview is doomed to bias. Evidence from industrial-organizational psychology shows that well-structured interviews—where questions are job-relevant, asked in a fixed order, and scored against explicit rubrics—can match or exceed the predictive validity of cognitive tests while sharply reducing demographic disparities (e.g., Campion et al., 1997; Levashina et al., 2014). Bias also varies with context: when small teams hire for roles they will personally supervise, interviewers are more accountable for outcomes and tend to weight task competence over social similarity. By contrast, HR gatekeepers or external recruiters, who rarely work with the people they select, have weaker incentives and greater latitude to default to cultural fit heuristics. Seeking some degree of cultural or interpersonal complementarity can facilitate coordination and morale (e.g., Reagans and Zuckerman, 2001), but it becomes costly when those signals crowd out more relevant indicators of capability, thereby misallocating talent. Thus, whether interviews amplify or mute status bias depends on how organizations structure and incentivize them, not on interviews as a selection method *per se*.

What can organizations do? While eliminating interviews is impractical, evidence suggests concrete reforms. Structured behavioral questions reduce bias while preserving valid assessment (Highhouse, 2008). Forcing deliberate pauses before decisions improves judgment quality (Pinker, 2018). Most critically, tracking interviewer ratings against subsequent performance creates accountability—when evaluators know their "fit" assessments will be audited against actual outcomes, superficial biases lose their power.

As hiring increasingly blends human judgment with algorithmic screening, the lessons from

this study grow more urgent. Whether in courts (Albright, 2019), admissions (DeVaul et al., 1987), or employment (Rivera, 2012), whenever discretion enters ostensibly neutral processes, class advantages reassert themselves. Breaking this cycle requires recognizing that our intuitions about who belongs often reflect inherited privilege rather than individual merit. Only by constraining these impulses—through structure, accountability, and continuous monitoring—can organizations tap overlooked talent and help societies achieve genuine mobility.

References

[1] Ægisdóttir, S., M. J. White, P. M. Spengler, A. S. Maugherman, L. A. Anderson, R. S. Cook, C. N. Nichols, G. K. Lampropoulos, B. S. Walker, G. Cohen, and J. D. Rush (2006). The meta-analysis of clinical judgment project: Fifty-six years of accumulated research on clinical versus statistical prediction. *The Counseling Psychologist 34*(3), 341–382.

[2] Aigner, D. J. and G. G. Cain (1977). Statistical theories of discrimination in labor markets. *Industrial and Labor Relations Review 30*(2), 175–187.

[3] Albright, A. (2019). If you give a judge a risk score: Evidence from Kentucky bail decisions. Working paper.

[4] Altonji, J. G. and R. M. Blank (1999). *Race and Gender in the Labor Market. Handbook of Labor Economics*, ed. by O. C. Ashenfelter and D. Card, Elsevier, vol. 3C, chap. 48.

[5] Anthropological Survey of India (1992). *People of India: National Series*. Kolkata: Anthropological Survey of India. K.S. Singh, General Editor.

[6] Baird, M., N. Gahlawat, R. Hood, S. Lara, and P. Ko (2023). Pathways into STEM jobs. Technical report. White paper, November 2023.

[7] Bapuji, H., K. Gupta, S. Chrispal, and T. Roulet (2023). What managers everywhere must know about caste. *MIT Sloan Management Review Winter*. Accessed: 2024-04-22.

[8] Barrios-Fernández, A., C. Neilson, and S. Zimmerman (2024). Elite universities and the intergenerational transmission of human and social capital. *American Economic Review*. conditionally accepted.

[9] Bayly, S. (2008). *Caste, Society and Politics in India from the Eighteenth Century to the Modern Age*. Chicago: Cambridge University Press.

[10] Benson, A., D. Li, and K. Shue (2019). Promotions and the Peter principle. *The Quarterly Journal of Economics 134*(4), 2085–2134.

[11] Benson, A., D. Li, and K. Shue (2023). "Potential" and the gender promotion gap. R&R, *American Economic Review*.

[12] Benson, A. and K. Shaw (2024). What do managers do? *Annual Review of Economics*. Forthcoming.

[13] Beteille, A. (1965). *Caste, Class and Power: Changing Patterns of Stratification in a Tanjore Village*. University of California Press.

[14] Beteille, A. (1969). *Castes: Old and New, Essays in Social Structure and Social Stratification*. Asia Publishing House.

[15] Bhavnani, R. R. and A. Lee (2021). Does affirmative action worsen bureaucratic performance? Evidence from the indian administrative service. *American Journal of Political Science 65*(1), 5–20.

[16] Bhupatiraju, S., D. L. Chen, S. Joshi, and P. Neis (2024). Caste aside? Names, networks and justice in the courts of bihar, india. *European Journal of Empirical Legal Studies 1*(2), 151–178.

[17] Bisin, A., J. Rubin, A. Seror, and T. Verdier (2023). Culture, institutions & the long divergence. *Journal of Economic Growth 28*, 1–40.

[18] Bisin, A. and T. Verdier (1998). On the cultural transmission of preferences for social status. *Journal of Public Economics 70*, 75–97.

[19] Bisin, A. and T. Verdier (2010). The economics of cultural transmission and socialization. In *Handbook of Social Economics*, pp. 339–416. Elsevier.

[20] Bloom, N., R. Sadun, and J. Van Reenan (2012). The organization of firms across countries. *The Quarterly Journal of Economics 127*(4), 1663–1705.

[21] Bol, T., J. Witschge, H. G. Van de Werfhorst, and J. Dronkers (2014). Curricular tracking and central examinations: Counterbalancing the impact of social background on student achievement in 36 countries. *Social Forces 92*(4), 1545–1572.

[22] Bourdieu, P. (1986). The forms of capital. In J. G. Richardson (Ed.), *Handbook of Theory and Research for the Sociology of Education*, pp. 241–258. Greenwood Press.

[23] BusinessLine (2018). An Inclusive India Inc. Business Line, January 13, 2018.

[24] California v. Cisco Systems, Inc. (2022). Dept. of Fair Employment and Housing v. Cisco Systems, Inc. Docket No. H048962.

[25] Campion, M. A., D. K. Palmer, and J. E. Campion (1997). A review of structure in the selection interview. *Personnel Psychology 50*, 655–702.

[26] Cappelli, P. (2019). Your approach to hiring is all wrong. *Harvard Business Review*. Accessed: 2024-08-15.

[27] Castilla, E. J. (2015). Accounting for the gap: A firm study manipulating organizational accountability and transparency in pay decisions. *Organization Science 26*(2), 311–333.

[28] Castilla, E. J. and S. Benard (2010). The paradox of meritocracy in organizations. *Administrative Science Quarterly 55*(4), 543–576.

[29] Chakravartty, P. and A. Subramanian (2021). Why is Caste Inequality Still Legal in America? *The New York Times*, May 25, 2021.

[30] Chetty, R., D. J. Deming, and J. N. Friedman (2023). Diversifying society's leaders? The determinants and causal effects of admission to highly selective private colleges. NBER Working Paper 31492.

[31] Chow, R. M. and E. D. Knowles (2016). Taking race off the table: Agenda setting and support for color-blind public policy. *Personality and Social Psychology Bulletin 42*(1), 25–39.

[32] Clark, G. (2015). *The Son Also Rises: Surnames and the History of Social Mobility*. The Princeton Economic History of the Western World.

[33] Dana, J., R. Dawes, and N. Peterson (2013). Belief in the unstructured interview: The persistence of an illusion. *Judgement and Decision Making 8*(5), 512–520.

[34] Datta, S. (2017). *A History of the Indian University System: Emerging from the Shadows of the Past*. Palgrave Macmillan.

[35] Deshpande, A. (2011). *The Grammar of Caste: Economic Discrimination in Contemporary India*. Oxford: Oxford University Press.

[36] DeVaul, R., F. Jervey, J. Chappell, P. Caver, B. Short, and S. O'Keefe (1987). Medical school performance of initially rejected students. *Journal of the American Medical Association 257*, 47–51.

[37] Durlauf, S. N. (1999). The case "against" social capital. In P. Dasgupta and I. Serageldin (Eds.), *Social Capital: A Multifaceted Perspective*, pp. 84–104. Washington, D.C.: World Bank.

[38] Durlauf, S. N. (2006). Groups, social influences, and inequality: A memberships theory perspective on poverty traps. In S. Bowles, S. N. Durlauf, and K. Hoff (Eds.), *Poverty Traps*. Princeton: Princeton University Press.

[39] Durlauf, S. N. (2008). Memberships and inequality. In D. Castiglione, J. van Deth, and G. Wolleb (Eds.), *Handbook of Social Capital*. Oxford: Oxford University Press.

[40] Eames, T. (2024). Taryn versus taryn (she/her) versus taryn (they/them): A field experiment on pronoun disclosure and hiring discrimination. Working Paper.

[41] Election Commission of India (2024). Reserved constituencies data. Accessed May 25, 2024.

[42] Erickson, B. H. (1996). Culture, class, and connections. *American Journal of Sociology 102*(1), 217–251.

[43] Espinosa, M. and C. T. Stanton (2022, July). Training, communications patterns, and spillovers inside organizations. Working Paper 30224, National Bureau of Economic Research.

[44] Fenizia, A. (2022). Managers and productivity in the public sector. *Econometrica 90*(3), 1063–1084.

[45] Gibbons, R. and J. Roberts (2013). Economic theories of incentives in organizations. *The Handbook of Organizational Economics*.

[46] Gorman, E. H. (2005). Gender stereotypes, same-gender preferences, and organizational variation in the hiring of women: Evidence from law firms. *American Sociological Review 70*(4), 702–728.

[47] Gottfried, J. (2024, January). Americans' social media use. Technical report, Pew Research Center.

[48] Grusky, D. B. (2005). Foundations of a neo-durkheimian class analysis. In E. O. Wright (Ed.), *Approaches to Class Analysis*, pp. 51–81. Cambridge University Press.

[49] Hampole, M., F. Truffa, and A. Wong (2024). Peer effects and the gender gap in corporate leadership: Evidence from MBA students. *Quarterly Journal of Economics*. Revise and Resubmit.

[50] Highet, K. (2023). "RIP English": Race, Class and 'Good English' in India. *Journal of Linguistic Anthropology 33*(2), 184–201.

[51] Highhouse, S. (2008). Stubborn Reliance on Intuition and Subjectivity in Employee Selection. *Industrial and Organizational Psychology: Perspectives on Science and Practice 1*(3), 333–342.

[52] Hoffman, M., L. B. Kahn, and D. Li (2018). Discretion in Hiring. *The Quarterly Journal of Economics 133*(2), 765–800.

[53] Hoffman, M. and C. T. Stanton (2024). People, practices, and productivity: A review of new advances in personnel economics. NBER Working Paper 32849.

[54] Hoffman, M. and S. Tadelis (2021). People management skills, employee attrition, and manager rewards: An empirical analysis. *Journal of Political Economy 129*(1), 243–285.

[55] Human Rights Watch (2001). Caste discrimination: A global concern. HRW Briefing Paper. See ch. 2, p. 2: "Lower-caste communities are almost invariably indistinguishable in physical appearance from higher-caste communities".

[56] Jodhka, S. S. and J. Naudet (2019). *Mapping the Elite: Power, Privilege, and Inequality*. Oxford University Press.

[57] Jodhka, S. S. and K. Newman (2007). In the name of globalisation: Meritocracy, productivity and the hidden language of caste. *Economic and Political Weekly 42*(41), 4125–4132.

[58] Jodkha, S. S. (2017). *Caste in Contemporary India*. Routledge India.

[59] Kahneman, D. (2011). *Thinking, Fast and Slow*. New York: Farrar, Straus and Giroux.

[60] Kang, S. K., K. A. DeCelles, A. Tilcsik, and S. Jun (2016). Whitened resumes: Race and self-presentation in the labor market. *Administrative Science Quarterly 61*(3), 469–502.

[61] Kashyap, R. and F. Verkroost (2021). Analysing global professional gender gaps using LinkedIn advertising data. *EPJ Data Science 10*(39), 1–32.

[62] Khan, S. R. (2011). *Privilege: The Making of an Adolescent Elite at St. Paul's School*. Princeton, NJ: Princeton University Press.

[63] Kline, P., E. Rose, and C. Walters (2022). Systemic discrimination among large U.S. employers. *Quarterly Journal of Economics 137*(4), 1963–2036.

[64] Knobe, J. and S. Nichols (2008). An introduction to experimental philosophy. In *Experimental Philosophy*, pp. 3–34. Oxford: Oxford University Press.

[65] Krishna, K. and V. C. Frisancho Robles (2015). Affirmative Action in Higher Education in India: Targeting, Catch Up, and Mismatch. *Higher Education*.

[66] Kuncel, N. R. and S. A. Hezlett (2007). Standardized tests predict graduate students' success. *Science 315*(5815), 1080–1081.

[67] Kuncel, N. R., R. J. Kochevar, and D. S. Ones (2014). A meta-analysis of letters of recommendation in college and graduate admissions: Reasons for hope. *International Journal of Selection and Assessment 22*(1), 101–107.

[68] Lamont, M. (1992). *Money, Morals, and Manners: The Culture of the French and the American Upper-Middle Class*. Chicago: University of Chicago Press.

[69] Lamont, M. (2009). *How Professors Think: Inside the Curious World of Academic Judgment*. Cambridge, MA: Harvard University Press.

[70] Lamont, M. and A. Lareau (1988). Cultural capital: Allusions, gaps and glissandos in recent theoretical developments. *Sociological Theory 6*(2), 153–168.

[71] Laurison, D. and S. Friedman (2019). *The Class Ceiling: Why it Pays to be Privileged*. Bristol, UK: Policy Press.

[72] Lazear, E. P., K. L. Shaw, and C. T. Stanton (2015). The value of bosses. *Journal of Labor Economics 33*(4), 823–861.

[73] Lenski, G. E. (1954). Status crystallization: A non-vertical dimension of social status. *American Sociological Review 19*(4), 405–413.

[74] Levashina, J., C. J. Hartwell, F. P. Morgeson, and M. A. Campion (2014). The structured employment interview: Narrative and quantitative review of the research literature. *Personnel Psychology 67*(1), 241–293.

[75] Levon, E., D. Sharma, D. J. L. Watt, A. Cardoso, and Y. Ye (2021). Accent bias and perceptions of professional competence in england. *Journal of English Linguistics 49*(4), 355–388.

[76] Madheswaran, S. (2008). Is affirmative action policy for private sector necessary? *Indian Journal of Industrial Relations 44*(2), 164–184.

[77] Mamidi, P. (2011). Strategic manipulation of caste identity signals. Doctoral Dissertation, Oxford University.

[78] Martellini, P., T. Schoellman, and J. Sockin (2024, February). The global distribution of college graduate quality. *Journal of Political Economy 132*(2), 434–483.

[79] Meehl, P. E. (1954). *Clinical versus Statistical Prediction: A Theoretical Analysis and a Review of the Evidence*. University of Minnesota Press.

[80] Michelman, V., J. Price, and S. D. Zimmerman (2022). Old boys' clubs and upward mobility among the educational elite. *The Quarterly Journal of Economics 137*(2), 845–909.

[81] Ministry of Law and Justice, Government of India (2021). The constitution of India. Article 15 explicitly prohibits discrimination on grounds of religion, race, caste, sex, or place of birth.

[82] Minni, V. (2023). Making the invisible hand visible: Managers and the allocation of workers to jobs. R&R, *Quarterly Journal of Economics*.

[83] Minni, V. (2024). Global managers, local workers: Wage setting inside a multinational firm. *AEA Papers & Proceedings 114*, 586–91.

[84] Minni, V., K.-T. Nguyen, and H. Sarsons (2025). Managers and the cultural transmission of gender norms. Working Paper.

[85] Mishra, N. (2015). India and colorism: The finer nuances. *Washington University Global Studies Law Review 14*, 725–750.

[86] Murphy, S. C., D. M. Klieger, M. J. Borneman, and N. R. Kuncel (2009). The predictive power of personal statements in admissions: A meta-analysis and cautionary tale. *Educational and Psychological Measurement 69*(6), 967–982.

[87] Neckerman, K. M. and J. Kirschenman (1991). Hiring strategies, racial bias, and inner-city workers. *Social Problems 38*(4), 433–447.

[88] Oster, E. (2019). Unobservable selection and coefficient stability: Theory and evidence. *Journal of Business & Economic Statistics 37*(2), 187–204.

[89] Parameswaran, R. and K. Cardoza (2015). Melanin on the margins: Advertising and the cultural politics of fair/light/white beauty in india. *Journalism & Communication Monographs 11*(3), 213–274.

[90] Parmar, R. (2020). Transacting caste in modern times: Changing social identity through surnames in urban Gujarat. *Contemporary Voice of Dalit 12*(2), 220–233.

[91] Pinker, S. (2018). *Enlightenment Now: The Case for Reason, Science, Humanism, and Progress*. New York: Viking.

[92] Pinker, S. (2021). *Rationality: What It Is, Why It Seems Scarce, Why It Matters*. New York: Viking.

[93] Puri, S. (2025, February). Capitalism with Indian characteristics. Stigler Center Lecture Series, University of Chicago Booth School of Business. Accessed: 2025-02-11.

[94] Quillian, L., J. J Lee, and M. Oliver (2020). Evidence from Field Experiments in Hiring Shows Substantial Additional Racial Discrimination after the Callback. *Social Forces 99*(2), 732–759.

[95] Reagans, R. and E. W. Zuckerman (2001, July–August). Networks, diversity, and productivity: The social capital of corporate R&D teams. *Organization Science 12*(4), 393–521.

[96] Reskin, B. (2012). The race discrimination system. *Annual Review of Sociology 38*, 17–35.

[97] Ridgeway, C. L. (2019). *Status: Why Is It Everywhere? Why Does It Matter?* New York, NY: Russell Sage Foundation.

[98] Rivera, L. A. (2012). Hiring as cultural matching: The case of elite professional service firms. *American Sociological Review 77*(6), 999–1022.

[99] Rivera, L. A. (2015). *Pedigree: How Elite Students Get Elite Jobs*. Princeton University Press.

[100] Sacerdote, B., D. O. Staiger, and M. Tine (2025, January). How test optional policies in college admissions disproportionately harm high achieving applicants from disadvantaged backgrounds. Working Paper 33389, National Bureau of Economic Research.

[101] Sadun, R. (2023). CEOs and firm performance. NBER Reporter, ISSN 0276-119X, National Bureau of Economic Research (NBER), Cambridge, MA, Iss. 1, pp. 12-16.

[102] Sadun, R., N. Bloom, and J. Van Reenen (2017, September-October). Why do we undervalue competent management? Neither great leadership nor brilliant strategy matters without operational excellence. *Harvard Business Review 95*(5), 120–127. Winner of 59th Annual HBR McKinsey Award.

[103] Sadun, R., J. Tamayo, Jorge A. Ramirez, Diaz, and A. B. S. Neyra-Nazarrett (2024, April). Training within firms. Working Paper 33670, National Bureau of Economic Research.

[104] Sarsons, H. (2017). Recognition for group work: Gender differences in academia. *AEA Papers and Proceedings 107*(5), 141–145.

[105] Sarsons, H., K. Gërxhani, E. Reuben, and A. Schram (2021). Gender differences in recognition for group work. *Journal of Political Economy 129*(1), 101–147.

[106] Sarsons, H., J. Hjort, and X. Li (2022). Across-Country Wage Compression in Multinationals. R&R, *American Economic Review*.

[107] Savage, M. (2015). *Social Class in the 21st Century*. Pelican Books.

[108] Savage, M., F. Devine, N. Cunningham, M. Taylor, Y. Li, J. Hjellbrekke, B. Le Roux, S. Friedman, and A. Miles (2013). A new model of social class? Findings from the bbc's great british class survey experiment. *Sociology 47*(2), 219–250.

[109] Sebring, J. M. (1969). Caste indicators and caste identification of strangers. *Human Organization 28*(3), 199–207.

[110] Sharma, D. (2017). English in India. In J. Filipović and S. C. Deshors (Eds.), *World Englishes: New Theoretical and Methodological Considerations*, pp. 311–332. De Gruyter Mouton.

[111] Stansbury, A. (2024). Social class and career progression in elite professions. *American Economic Journal: Applied Economics Forthcoming.*

[112] Stansbury, A. and K. Rodriguez (2024a). The class gap in career progression: Evidence from US academia. R&R, *Econometrica.*

[113] Stansbury, A. and K. Rodriguez (2024b, December). The class gap in career progression: Evidence from US academia. Presented at the NBER Organizational Economics Working Group, Fall 2024.

[114] Subramanian, A. (2019). *The Caste Of Merit: Engineering Education in India.* Harvard University Press.

[115] Tambe, P., L. Hitt, D. Rock, and E. Brynjolfsson (2021). Digital capital and superstar firms. Technical report, Brookings Institution, Hutchins Center Working Paper #73.

[116] Thorat, S. and P. Attewell (2007, October). The legacy of social exclusion: A correspondence study of job discrimination in india. *Economic and Political Weekly 42*(41), 4141–4145.

[117] Todorov, A. (2017). *Face Value: The Irresistible Influence of First Impressions.* Princeton, NJ: Princeton University Press.

[118] Turco, C. J. (2010). Cultural foundations of tokenism: Evidence from the leveraged buyout industry. *American Sociological Review 75*(6), 894–913.

[119] Wiltshire, C. R. (2020). *Uniformity and Variability in the Indian English Accent.* Cambridge University Press.

[120] Zhang, L., S. Wang, and N. Wilmers (2025). Remote work and careers. SSRN working paper.

[121] Zhang, S. and S. Vucetic (2016). Sampling bias in Linkedin: A case study. In *Proceedings of the 25th International Conference Companion on World Wide Web*, pp. 145–146.

[122] Zimmerman, S. D. (2019). Elite colleges and upward mobility to top jobs and top incomes. *American Economic Review 109*(1), 1–47.

ONLINE APPENDIX

A Caste as Social Status Under Affirmative Action

India's caste system provides a uniquely precise measure of social status, offering clear administrative categories for analyzing class-based inequalities. Affirmative action policies formally recognize caste as a comprehensive marker of historical advantage and disadvantage, codifying how caste membership systematically shapes access to resources, education, and elite employment opportunities. This administrative clarity enables identification of status-based discrimination mechanisms that remain difficult to isolate in Western class analyses.

Jati, Varna, and Modern Caste Classifications. India's modern caste system originates from historical occupational hierarchies. At its most detailed level, the system comprises thousands of hereditary occupational groups known as *Jatis*, historically linked to trades such as merchants, priests, warriors, and laborers. These *Jatis* are grouped into broader categories called *Varnas*: Brahmins (priests), Kshatriyas (warriors), Vaishyas (merchants), Shudras (laborers), and historically marginalized "untouchables" (Dalits). The government maintains official lists classifying communities as "Scheduled Castes" (SCs), "Scheduled Tribes" (STs), and "Other Backward Classes" (OBCs)—groups defined as socially and educationally disadvantaged (Puri, 2025).

This paper adopts two broad categories aligned with India's affirmative action framework: advantaged ("upper") and disadvantaged ("lower") castes. By disadvantaged castes, I refer explicitly to Scheduled Castes, Scheduled Tribes, and Other Backward Classes. This administrative categorization recognizes caste as fundamentally intertwined with economic marginalization and social status, paralleling Western concepts of class (Puri, 2025).

Historical Development of Caste-based Affirmative Action Policies. Affirmative action for disadvantaged communities emerged during British colonial rule, beginning with the Government of India Act of 1919's reserved legislative seats for "depressed classes." The 1935 Government of India Act officially designated these groups as Scheduled Castes. Post-independence, India's 1949 Constitution institutionalized quota-based affirmative action in legislatures, public higher education, and government employment for Scheduled Castes and Scheduled Tribes. Following the Mandal Commission's recommendations in the 1980s, quotas expanded to include Other Backward Classes, raising reservation levels to nearly 50% (Bayly, 2008).

Caste as Proxy for Social Status. Although caste originated as a hereditary hierarchy, Indian policymakers now define caste categories as comprehensive indicators of socioeconomic "backwardness"—encompassing economic, educational, and social disadvantages. This approach recognizes caste as encapsulating multiple dimensions of social status: financial resources, educational credentials, social networks, and cultural capital (cf. Bourdieu, 1986). The Mandal Commission merged lineage-based identities with socioeconomic status, establishing caste as a consolidated marker of social status.

Limited Scope of Affirmative Action in the Private Sector. India's affirmative action policies do not extend to private-sector employment, including the multinational corporations studied here. Private-sector roles typically offer higher wages and career growth but lack reservation mandates, disproportionately excluding disadvantaged castes (Subramanian, 2019). This parallels Western contexts, where firms emphasize racial and gender diversity but rarely address socioeconomic diversity. Fewer than 10% of large US firms track socioeconomic diversity (Stansbury and Rodriguez, 2024b), while fewer than 3% of India's top firms maintain internal caste data (BusinessLine, 2018). Although a few multinational corporations have recently begun addressing caste discrimination (Chakravartty and Subramanian, 2021; Bapuji et al., 2023), these changes occurred in late 2022 or 2023, falling outside the timeframe of the data analyzed in this paper (Section 3.4).

B Representativeness of the Job Fair at the Elite College

The studied on-campus job fair provides a highly representative window into how India's elite-college graduates transition into elite private-sector jobs, supported by four key observations:

1. **Near-identical fairs across elite institutions.** Since their inception in the 1950s, India's premier institutions—including the Indian Institutes of Technology (IITs), Indian Institutes of Management (IIMs), and leading Delhi University colleges—have consciously modeled their practices after prominent Western universities such as MIT, Stanford, and Harvard. Consequently, these institutions share strikingly similar academic calendars, faculty-to-student ratios, curricula, and critically, their standardized placement office procedures. Today, the job fairs conducted by all 23 IITs, 20 IIMs, and top-tier Delhi University colleges closely mirror one another, featuring virtually identical recruitment stages, timelines, and employer pools (e.g., see Central Placement Cell, Delhi University).

2. **Graduates predominantly flow into elite private-sector jobs.** Graduates from India's elite colleges overwhelmingly pursue high-paying entry-level roles in multinational corporations and prominent Indian firms. Public-sector employment remains a distinctly secondary choice, reflecting a long-standing preference among top-tier college graduates for careers in consulting, technology, finance, and multinational management (e.g., Subramanian, 2019; Jodhka and Naudet, 2019).

3. **Elite firms recruit almost exclusively from elite campuses.** The multinational firms and prestigious Indian companies studied here systematically target the same elite Indian institutions year after year. Collectively, these recruiters visit around 30–40 top-ranked colleges annually, each institution typically enrolling more than 1,000 students. This concentrated recruiting strategy ensures a highly consistent candidate pool across years, facilitating a standardized recruitment process and comparable outcomes nationwide. Prior studies further corroborate that elite firms' hiring remains strongly concentrated within these elite campuses (e.g., Subramanian, 2019; Jodhka and Naudet, 2019).

4. **Minimal student withdrawal from the job fair.** Participation in the on-campus job fair is near-universal among graduating students. Less than 5% of final-year students voluntarily deregister from the formal placement process, a proportion consistent across

similar elite institutions (Subramanian, 2019). Such low withdrawal rates underscore the central role of these standardized recruitment fairs in determining employment outcomes for India's most academically competitive graduates.

Collectively, these facts demonstrate that the focal job fair closely mirrors elite hiring pathways across India's top universities.

C Benchmarking Employment-Platform Coverage: External Evidence

Large-scale validations suggest LinkedIn already reaches roughly four-fifths of U.S. bachelor's-degree holders (Zhang et al., 2025). Coverage is markedly higher in technical fields. LinkedIn's own "Pathways into STEM Employment" study reports that 92.7% of 2022 STEM master's graduates— virtually all STEM users with active profiles—interacted with at least one STEM job on the site, a 12-point lift over the bachelor-level average (Baird et al., 2023). Cross-country work buttresses this. In their 92-nation ad-audience study, Kashyap and Verkroost (2021) show that for India the "Information & Communication" sector accounts for 32% of female and 27% of male LinkedIn users but only about 1% of the professional labour force logged by the ILO, an enrichment of more than 25 percentage points that mirrors the global STEM premium they document.

Firm-level evidence points the same way: Tambe et al. (2021) find IT employees make up 5% of LinkedIn's U.S. workforce versus 2% in BLS data, implying near-census coverage of U.S. IT professionals once Pew's 35% overall adoption rate is factored in (Gottfried, 2024). Even outside STEM, MBA follow-ups routinely match 85–90% of graduates from many cohorts to LinkedIn profiles (Hampole et al., 2024). And elite STEM-focused universities often comfortably clear 90%: Stanford's official count lists "more than 220,000 living alumni," yet LinkedIn shows about 314,000 profiles, which after de-duplication and mis-tag removals using standard methods (e.g., Zhang and Vucetic, 2016) suggest a more than 95% hit-rate. Given that our college's talent pool is comparable to many world-leading institutions (Martellini et al., 2024), a well over 90% in-sample match rate at our STEM-heavy campus is entirely plausible. Indeed, the third-party matched dataset achieves a nearly 98% match rate for graduates (modestly higher—about 9%—than other benchmarks that hover around 90%), aided by systematic institutional identifiers that follow predictable patterns within each major-degree-year, which greatly boosts match precision.

D Timeline and Screening Rounds

Every job in the sample follows the same four-stage sequence, eliminating candidates at each step:

1. **Application screening (September).** Recruiters conduct an initial screening, primarily based on GPA, major, degree, and standout résumé items (e.g., internships, competitions).

2. **Written aptitude tests (September–November).** Online technical assessments vary by role: coding exercises for tech/manufacturing jobs, and case-based tests for consulting positions. Firms advertising multiple posts usually administer one common test (e.g., Nvidia uses a single exam for all advertised profiles).

3. **Group debates (September–November).** A non-technical evaluation of soft skills:

   (a) Groups of 30–35 students, split into two teams, debate a general prompt (e.g., "Is social media a boon or a bane?"), while recruiters observe communication, teamwork, and confidence.

   (b) Firms often conduct a single group debate for all advertised roles—for example, Google may simultaneously assess applicants for program manager and software engineer positions.

4. **Personal interviews (December–January).** Final 30-minute "fit" conversations:

   (a) Interviews occur simultaneously in multiple parallel batches, each handled by different members of the recruiting team.

   (b) As with earlier rounds, a single interview panel commonly covers multiple job titles (e.g., Uber may interview data scientist and front-end engineer candidates within the same track).

   (c) Because more than 98% of roles are based in India, interviewers are almost certainly Indian (Section 3.6).

E Qualitative evidence supporting caste revelation at interviews

Qualitative evidence further clarifies why caste becomes most legible to recruiters during personal interviews. As discussed quantitatively in Section 4.6.1, surnames alone rarely disclose caste in urban Indian settings; instead, status cues emerge primarily through richer social interactions. Group debates provide limited scope for caste inference: candidates debate standardized prompts in large groups with minimal individualized questioning, and neither physical appearance nor linguistic cues reliably indicate caste. Anthropological research consistently finds phenotypic traits like skin color are unreliable caste indicators among urban, educated populations (Sebring, 1969; Mishra, 2015; Parameswaran and Cardoza, 2015). Human Rights Watch concurs, stating that "*Lower-caste communities are almost invariably indistinguishable in physical appearance from higher-caste communities*" (Human Rights Watch, 2001, p. 2). Similarly, elite-college varieties of Indian English reflect regional or educational backgrounds rather than caste, as English-medium schooling has expanded significantly beyond upper castes due to affirmative action (Wiltshire, 2020; Sharma, 2017; Highet, 2023). Thus, group debates offer employers minimal caste-identifying information.

In contrast, personal interviews substantially shift the evaluative context. They amplify earlier signals such as schooling pedigree, resume details, and debate performances through deeper discussions. Interviews also elicit richer background information—personal interests, leisure activities, niche extracurricular involvement, and peer networks (e.g., Rivera, 2012, 2015), which

are strongly correlated with caste-linked social and cultural capital, such as networks, resources, and cultural tastes (cf. Bourdieu, 1986). Such characteristics are well-documented caste correlates in elite Indian contexts (cf. Jodhka and Newman, 2007; Deshpande, 2011; Jodkha, 2017). Additionally, caste identities frequently become salient through social interactions in organizational settings, as managers learn subtle cues from conversations about personal background, residential origins, familial occupations, or shared acquaintances (Bapuji et al., 2023). Integrating these nuanced signals with earlier observations enables recruiters to more accurately infer caste, aligning precisely with the timing and intensity of caste penalties documented quantitatively (Section 4.6).

F Derivation of Expected Performance Above Threshold

This appendix establishes the limit result used in Proposition 2,

$$\lim_{\tau\to\infty}\big(\mathbb{E}[M \mid M \geq \tau] - \tau\big) = 0.$$

The proof is self-contained and relies on a single, transparent tail assumption.

Setup. Let M be a real-valued random variable with CDF F, PDF f, survival function $S(a) = 1 - F(a)$, and hazard rate $h(a) = f(a)/S(a)$ (defined wherever $S(a) > 0$).

Assumption A. M has a continuously differentiable density, finite mean, and its hazard rate diverges:

$$\lim_{a\to\infty} h(a) = \infty.$$

This covers all light-tailed, log-concave distributions commonly used in practice (e.g. Normal, Logistic, Weibull with shape > 1).

Proposition 3 (Mean of a Right-Truncated Distribution)**.** For any τ with $S(\tau) > 0$, define the mean residual life

$$m(\tau) \;=\; \mathbb{E}[M - \tau \mid M \geq \tau] \;=\; \frac{\displaystyle\int_\tau^\infty S(u)\,du}{S(\tau)}.$$

Under Assumption A,

$$\lim_{\tau\to\infty} m(\tau) = 0, \quad \text{hence} \quad \mathbb{E}[M \mid M \geq \tau] = \tau + m(\tau) \;\to\; \tau.$$

Proof. Let

$$A(\tau) = \int_\tau^\infty S(u)\,du,$$

so $m(\tau) = A(\tau)/S(\tau)$. Then

$$A'(\tau) = -S(\tau), \qquad S'(\tau) = -f(\tau).$$

Differentiating $m(\tau) = A(\tau)/S(\tau)$ gives

$$m'(\tau) = \frac{A'(\tau)\,S(\tau) \;-\; A(\tau)\,S'(\tau)}{S(\tau)^2} = \frac{-S(\tau)^2 \;+\; f(\tau)\,A(\tau)}{S(\tau)^2} = -1 + \underbrace{\frac{f(\tau)}{S(\tau)}}_{h(\tau)}\,\underbrace{\frac{A(\tau)}{S(\tau)}}_{m(\tau)} = -1 + h(\tau)\,m(\tau).$$

Rearrange to

$$m'(\tau) + 1 \;=\; h(\tau)\,m(\tau) \quad \rightarrow \quad m(\tau) = \frac{1 + m'(\tau)}{h(\tau)}.$$

Since $m(\tau) \geq 0$ and m is non-increasing, $m'(\tau) \leq 0$, so

$$0 \;\leq\; 1 + m'(\tau) \;\leq\; 1.$$

Under Assumption A, $h(\tau) \rightarrow \infty$, hence

$$0 \;\leq\; m(\tau) = \frac{1 + m'(\tau)}{h(\tau)} \;\rightarrow\; 0,$$

which completes the proof. □

Discussion. The key requirement is that the hazard rate diverges. Heavy-tailed distributions with $h(\tau) \rightarrow 0$ (e.g. Pareto) fail this property and yield $m(\tau) \rightarrow \infty$ instead. In our context—where performance scores are bounded or effectively light-tailed—the limit result holds, justifying the asymptotic arguments in Proposition 2.

G Extending the Simple Model of Late-Stage Discrimination with Imperfect Information

In Section 4.7.1, I assumed that employers observe each candidate's true screening-round performance perfectly (i.e., $M_{s,i,g} = A_{i,g}$ at round s). Here, I show that *all* key insights—particularly how variation in job selectivity identifies discrimination versus ability deficits—remain valid even when employers observe *noisy* performance signals.

Setup. As in Section 4.7.1, there are two groups (castes) $g \in \{0, 1\}$, where $g = 1$ denotes disadvantaged castes and $g = 0$ denotes advantaged castes, a continuum of jobs $j \in \mathcal{J}$, and four sequential screening rounds $\{1, 2, 3, 4\}$. Rounds $1, 2, 3$ constitute pre-interview filtering (e.g., applications, written tests, group debates), while round 4 is the final personal interview. Each candidate i from group g has a *latent ability* $A_{i,g}$. Passing round s requires

$$\mathbb{E}\big[M_{s,i,g}\,\big|\mathcal{I}_{s-1}\big] \;\geq\; \tau_{s,j}, \tag{OA.1}$$

where $\mathcal{I}_{s-1}$ denotes the employer's information after observing signals in rounds $1, \ldots, s-1$. Thresholds $\tau_{s,j}$ weakly increase with each stage and may vary across jobs.

Noisy Signals per Round. At each stage s, the employer observes

$$M_{s,i,g} \;=\; A_{i,g} \;+\; \eta_{s,i,g},$$

where $\eta_{s,i,g}$ is a zero-mean noise term. We relax the assumption of independence across candidates and rounds, requiring only that:

(i) $\mathbb{E}[\eta_{s,i,g}] = 0$ for all s, i, g, and

(ii) $\eta_{s,i,1} \overset{d}{=} \eta_{s,i,0}$ for every round s, so that for any subset of candidates (including those who have passed earlier rounds) there is no systematic caste-level difference in noise distributions.

Thus, while noise might correlate across rounds or candidates (e.g., candidate performance may be interdependent across stages), it never introduces systematic group-level biases. Under these conditions, persistent group differences cannot result from noise alone.

Assumptions. Throughout this appendix we maintain the common-pool assumption (Assumption 1): each job draws from the same latent-ability distributions F_g. Screening noise $\eta_{s,i,g}$ is therefore orthogonal to job-level variation in $\tau_{s,j}$. Since screening score distributions are light-tailed (Assumption 2), each F_g has a diverging hazard rate ($h_{F_g}(\tau) \to \infty$). Assuming $\eta_{s,i,g}$ is light-tailed as well, each signal $M_{s,i,g} = A_{i,g} + \eta_{s,i,g}$ inherits that property, ensuring $\mathbb{E}[M - \tau \mid M \geq \tau] \to 0$, ensuring that right-tail truncation leaves negligible residual mass.

Updating Beliefs and Passing Rounds. At each stage s, the employer updates belief about $A_{i,g}$ given $\mathcal{I}_s \equiv \{M_{1,i,g}, \ldots, M_{s,i,g}\}$. Let

$$\mathbb{E}[A_{i,g} \mid \mathcal{I}_s] \quad \text{and} \quad Var(A_{i,g} \mid \mathcal{I}_s)$$

denote the posterior mean and variance of $A_{i,g}$. With each successive stage, beliefs about candidate ability become more precise, and the passing condition can be evaluated using $\mathbb{E}[M_{s,i,g} \mid \mathcal{I}_{s-1}]$.

Final Interview Thresholds and Caste Revelation. As before, let $\tau_{4,j}^{(g)}$ be the *group-specific* final-round threshold at job j. If caste becomes salient at the interview for group 1, a higher threshold ($\tau_{4,j}^{(1)} > \tau_{4,j}^{(0)}$) indicates discrimination. The goal is to demonstrate how variation in thresholds across jobs disentangles true ability differences from discrimination—even with noisy signals.

G.0.1 No Early Caste Penalty Implies No Large Right-Tail Deficit Before Interviews

A substantial right-tail ability deficit for group 1 would yield a particularly large gap in pass rates at early stages for highly selective jobs (with high $\tau_{s,j}$). Formally:

Proposition 4 (No Early Gap $\implies$ No Right-Tail Deficit Under Imperfect Information). Consider any pre-interview stage $s < 4$. Under Assumption 1 (common-pool ability

distributions) and Assumption 2 (screening score distributions are light-tailed), if group 1 truly has fewer high-ability candidates in the top tail, then for sufficiently high $\tau_{s,j}$,

$$P\big(\text{Pass}_{s,i,j} \mid g = 1, \mathcal{I}_{s-1}\big) \;<\; P\big(\text{Pass}_{s,i,j} \mid g = 0, \mathcal{I}_{s-1}\big).$$

Therefore, empirically observing no such gap—even in highly selective jobs—rules out a meaningful right-tail deficit for group 1 before the interview stage.

Sketch. Under Assumption 1, all jobs sample from the same F_g. With noisy signals, passing round s depends on the posterior mean $\mathbb{E}[M_{s,i,g} \mid \mathcal{I}_{s-1}]$. For high $\tau_{s,j}$, employers require a sufficiently high $\mathbb{E}[A_{i,g} \mid \mathcal{I}_{s-1}]$. If group 1 has fewer high-ability individuals, as $\tau_{s,j}$ rises, the pass rate for group 1 falls faster than group 0. This mirrors the perfect-information argument in Proposition 1, and further relies on Assumption 2, that screening score distributions are light-tailed, which ensures that higher cutoffs intensify right-tail gaps. Because noise is symmetric across groups, it cannot conceal large latent ability gaps at high cutoffs. Thus, the absence of early-stage gaps even at very selective jobs implies no meaningful right-tail deficit for group 1. □

G.0.2 Stable Promotion Gaps Across Selectivity Still Imply Discrimination

Post-hire performance Y_i depends on latent ability $A_{i,g}$. Promotion requires $Y_i \geq \kappa_j$. Final-round passing demands $\mathbb{E}[M_{4,i,g} \mid \mathcal{I}_3] \geq \tau_{4,j}^{(g)}$. Suppose top $(1-q)$ fraction have cutoffs $\tau_{4,j}^{q}$. If both groups face identical thresholds, their hired performance converges, and the promotion gap vanishes. However, if $\tau_{4,j}^{(1)} = \tau_{4,j}^{(0)} + \delta$ ($\delta > 0$), group 1's marginal hires have higher posterior means, producing a stable advantage in realized performance.

Proposition 5 (Stable Promotion Gaps Under Imperfect Information Identify Discrimination). Under Assumptions 1-2 and the same logic as in Proposition 2, if group 1 faces a threshold higher by a constant $\delta > 0$, then average hired performance for group 1 remains about δ units higher than group 0, even in highly selective jobs. Hence, persistent promotion-rate gaps at high selectivity levels reject an unobserved-ability explanation in favor of discrimination.

Sketch. The argument parallels Proposition 2. Under imperfect information, stage-4 passage requires $\mathbb{E}[M_{4,i,g} \mid \mathcal{I}_3] \geq \tau_{4,j}^{(g)}$. If $\tau_{4,j}^{(1)} = \tau_{4,j}^{(0)} + \delta$, typical group 1 hires satisfy (and by Proposition 3, the mean-residual life of both ability and noise vanishes):

$$\mathbb{E}[A_{i,1} \mid \mathcal{I}_3] \;\approx\; \tau_{4,j}^{(0)} + \delta \;-\; \mathbb{E}[\eta_{4,i,1} \mid \mathcal{I}_3] \;=\; \tau_{4,j}^{(0)} + \delta \quad (\text{because } \mathbb{E}[\eta_{4,i,1}] = 0).$$

Because noise is identically distributed across groups (mean-zero, no systematic bias), raising group 1's threshold by δ effectively shifts their final hires toward higher regions of the ability distribution. At high selectivity, each group's hired candidates converge around their respective thresholds. Since group 1's threshold is higher by δ, their average hired performance remains about δ above group 0's. Conversely, identical thresholds shrink the performance gap to zero. Regular tail decay ensures marginal hires dominate outcomes at high cutoffs, maintaining a stable δ-sized gap under discrimination.

□

H Cultural Fit in Final-Round Interviews at Elite Multinationals

Final-round interviews at elite multinational firms often hinge less on technical ability than on interviewers' discretionary judgements of *cultural fit*. Building on the concept of *cultural matching* (Rivera, 2012), evaluators favour candidates whose background, tastes, and self-presentation feel "similar to us.'' In Rivera's seminal study of U.S. professional service firms, more than half of interviewers ranked cultural fit as *the single most important* criterion in the final round, ahead of analytical skill or academic pedigree. A law-firm partner captured the prevailing logic: *"In our new associates we are, first and foremost, looking for cultural compatibility. Someone who will . . . fit in''* (Rivera, 2012, p. 1006).

"Would I Want to Be Stuck in an Airport With Them?"

Interviewers often invoke the *airport test*: would I enjoy being stranded with this candidate for hours? A senior banker explained, *"A lot of this job is attitude, not aptitude . . . So you can be the smartest guy ever, but I don't care. I need to be comfortable working everyday with you, then getting stuck in an airport with you, and then going for a beer after. You need chemistry. Not only that the person is smart, but that you like him.''* (Rivera, 2012, p. 1008). Such remarks show how late-stage evaluation often becomes an audition of social rapport rather than pure competence (Lamont, 2009).

To locate common ground quickly, interviewers scour the résumé for shared experiences—college sports, study-abroad destinations, niche hobbies—and open the interview with informal chat. If a spark appears, enthusiasm often translates into advocacy in the hiring committee. Conversely, when overlap is missing, evaluators may withdraw support even when the candidate excels technically (Rivera, 2012, pp. 1013–1019).

Class-Coded Signals

The traits that signal fit are deeply entwined with social class and cultural capital (Bourdieu, 1986). Playing squash or lacrosse can mark a candidate as privileged in one firm and "too snooty'' in another, depending on the organization's self-image. Rivera documents a consulting director who excluded a technically strong candidate because *"we're more rough and tumble''* (Rivera, 2012, p. 1009). Similar dynamics appear in other sectors: Turco's (2010) ethnography of a leveraged-buyout firm showed women struggling to meet an aggressive, sport-centred masculine ideal, while Erickson (1996) finds that shared leisure pursuits foster stronger workplace ties and, ultimately, advancement.

Crucially, professionalism and client-readiness are evaluated separately from fit. As one consultant noted, *"You need someone who speaks in a way that earns your trust, who presents their opinion respectfully but also convincingly . . . But in terms of "fit," it's someone that we want on our case team . . . You want someone that makes you feel comfortable, that you enjoy hanging out with, can maintain a cool head when times are tough and make tough times kind of fun''* (Rivera, 2012, p. 1007). Recruiters believed that while "polish," "presence," and "presentation style" could be taught, fit cannot. Because the finalist pool is already academically and technically elite, these subtle cultural cues decisively influence the offer decision.

Implications for Caste

Our quantitative results show that caste penalties emerge almost exclusively at this final interview stage. The qualitative literature helps explain why. If lower-caste candidates have had less exposure to elite cultural codes—speech patterns, leisure interests, informal networks—they may appear a weaker "fit" despite equal credentials. Interviewers' reliance on self-referential judgements thereby re-inscribes social-status hierarchies under the guise of meritocracy (Neckerman and Kirschenman, 1991; Deshpande, 2011). As Rivera (2012) concludes, employers often *"constructed merit in a manner that validated their own strengths and experiences and perceived similar candidates as better applicants"* (p. 1012), reproducing privilege while believing they are selecting the "best" candidate.

I Tables and Figures

Table OA.1: **Caste Differences in Mean Salary Among Graduates**

	Dependent variable:			
	log (Salary)			
	Without Controls	With Controls	Client Facing	Non-Client Facing
	(1)	(2)	(3)	(4)
Disadv. Caste	-0.231^{***}	-0.103^{***}	-0.107^{***}	-0.089^{***}
	(0.020)	(0.017)	(0.027)	(0.022)
Disadv. Caste (Oster)		-0.076^{***}	-0.095^{***}	-0.072^{***}
		(0.019)	(0.026)	(0.022)
Constant	11.235^{***}	10.863^{***}	10.883^{***}	10.848^{***}
	(0.012)	(0.040)	(0.057)	(0.052)
Observations	2,527	2,527	849	1,678
Adjusted R^2	0.049	0.412	0.549	0.382
Year FE	✓	✓	✓	✓
Other Controls	✓	✓	✓	✓
Job FE				

Notes: This table reports estimates from earnings regressions for graduates who obtained jobs through the college's job fair. The dependent variable is log earnings. Controls include pre-college skills, GPA, prior labor market experience, major, and degree. Regressions include year fixed effects. Column (1) shows estimates without controls, Column (2) adds controls, and Columns (3)-(4) show separate estimates for client-facing and non-client-facing jobs with controls. Job fixed effects are omitted as salaries are identical within jobs regardless of caste. Bias-adjusted coefficients follow Oster (2019), setting $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$ (assuming unobservables match observables in selection intensity and allowing total explanatory power to exceed observed R^2 by 30%). $^*p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

Table OA.2: **Caste Differences in GPA Between Those Omitted From the Earnings Regression Versus Those Who Are Not**

	Dependent variable:	
	log (GPA)	
	(1)	(2)
Disadv. Caste	-0.033^{***}	-0.032^{***}
	(0.004)	(0.004)
Omitted	-0.167^{***}	-0.158^{***}
	(0.005)	(0.005)
Omitted * Disadv. Caste	-0.071^{***}	-0.070^{***}
	(0.007)	(0.007)
Constant	2.147^{***}	2.127^{***}
	(0.004)	(0.008)
Observations	4,164	4,164
Adjusted R^2	0.546	0.559
Year FE	✓	✓
Major		✓
Degree		✓
Job FE		
Note:	*p<0.1; **p<0.05; ***p<0.01	

Notes: This table shows caste differences in GPA between students included in and omitted from the earnings regression. "Omitted" comprises students who either deregistered from placement or received no job offers, typically pursuing aftermarket opportunities, higher education, or entrepreneurship (Sections 2.3 and 4.4). Column (1) includes only year fixed effects; Column (2) adds major and degree controls. Job fixed effects are excluded as they cannot be separately identified from the "omitted" coefficient for students without jobs. $^{*}p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

Table OA.3: **Caste Differences in GPA At Each Job Search Stage**

	Dependent variable:					
	log (GPA)					
	App	Test	Debate	Interview	Offer	Choice
	(1)	(2)	(3)	(4)	(5)	(6)
Disadv. Caste	-0.111^{***}	-0.049^{***}	-0.039^{***}	-0.019^{***}	0.010^{***}	0.009^{***}
	(0.005)	(0.003)	(0.002)	(0.001)	(0.002)	(0.002)
Disadv. Caste (Oster)	-0.104^{***}	-0.035^{***}	-0.034^{***}	-0.014^{***}	0.015^{***}	0.014^{***}
	(0.006)	(0.003)	(0.002)	(0.002)	(0.002)	(0.002)
Constant	2.057^{***}	2.131^{***}	2.163^{***}	2.212^{***}	2.214^{***}	2.199^{***}
	(0.028)	(0.024)	(0.023)	(0.006)	(0.006)	(0.028)
Mean GPA	7.776	7.967	8.17	8.26	8.419	8.419
Observations	4,005	3,618	3,266	3,031	2,527	2,527
Adjusted R^2	0.264	0.726	0.745	0.801	0.822	0.822
Year FE	✓	✓	✓	✓	✓	✓
Job FE	✓	✓	✓	✓	✓	✓
Major	✓	✓	✓	✓	✓	✓
Degree	✓	✓	✓	✓	✓	✓
Note:						*p<0.1; **p<0.05; ***p<0.01

Notes: This table reports within-job GPA differences across castes at each job search stage. Regressions control for major, degree as well as job and year fixed effects. Bias-adjusted coefficients follow Oster (2019), setting $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$ (assuming unobservables match observables in selection intensity and allowing total explanatory power to exceed observed R^2 by 30%). $^{*}p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

Table OA.4: **Caste Differences in Mean Salary at Each Job Search Stage in the Full Sample**

	Dependent variable:					
	log (Salary)					
	App	Test	Debate	Interview	Offer	Choice
	(1)	(2)	(3)	(4)	(5)	(6)
Disadv. Caste	0.004	−0.009	−0.006	−0.018*	−0.105***	−0.103***
	(0.005)	(0.007)	(0.007)	(0.010)	(0.013)	(0.017)
Disadv. Caste (Oster)	0.005	0.010	0.011	0.012	−0.079***	−0.076***
	(0.005)	(0.008)	(0.008)	(0.011)	(0.014)	(0.019)
Constant	11.032***	10.904***	10.848***	10.758***	10.672***	10.863***
	(0.011)	(0.015)	(0.015)	(0.023)	(0.030)	(0.040)
Observations	4,005	3,618	3,266	3,031	2,527	2,527
Adjusted R^2	0.578	0.491	0.547	0.549	0.488	0.412
Year FE	✓	✓	✓	✓	✓	✓
Other Controls	✓	✓	✓	✓	✓	✓
Job FE						

Notes: This table reports earnings regressions for students still in contention at each job search stage. Each column represents a different stage: (1) application, (2) written test, and subsequent rounds. The dependent variable is log earnings. All regressions control for pre-college skills, GPA, prior labor market experience, major, and degree. Regressions include year fixed effects. Job fixed effects are omitted as salaries are identical within jobs regardless of caste. Bias-adjusted coefficients follow Oster (2019), setting $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$ (assuming unobservables match observables in selection intensity and allowing total explanatory power to exceed observed R^2 by 30%). $^*p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

Table OA.5: **Caste Differences in the Success Rate at Each Job Search Stage Among Client-Facing Jobs**

	Success App (1)	Success Test (2)	Success Debate (3)	Success Interview (4)
Disadv. Caste (OLS)	-0.0066***	-0.0007	-0.0136***	-0.3062***
	(0.0022)	(0.0030)	(0.0046)	(0.0081)
Disadv. Caste (Oster)	0.002	0.006	-0.015***	-0.314***
	(0.002)	(0.004)	(0.005)	(0.009)
Mean Success Rate	0.5770	0.7860	0.2590	0.2600
Observations	117,770	67,927	53,422	13,842
Adjusted R^2	0.61588	0.37467	0.01189	0.24298
Year FE	✓	✓	✓	✓
Job FE	✓	✓	✓	✓
Other Controls	✓	✓	✓	✓

Notes: This table reports caste differences in success rates at each hiring stage for client-facing jobs. The unit of observation is student-job-year, conditional on remaining in contention at each stage. All regressions include controls as well as job and year fixed effects. When firms conduct common screening rounds for multiple positions, each position is treated as a separate job. Bias-adjusted coefficients follow Oster (2019), setting $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$ (assuming unobservables match observables in selection intensity and allowing total explanatory power to exceed observed R^2 by 30%). $^*p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

Table OA.6: **Caste Differences in the Success Rate at Each Job Search Stage Among Non-Client-Facing Jobs**

	Success App (1)	Success Test (2)	Success Debate (3)	Success Interview (4)
Disadv. Caste (OLS)	0.0024	0.0011	0.0070**	-0.2867***
	(0.0018)	(0.0022)	(0.0034)	(0.0060)
Disadv. Caste (Oster)	0.008***	0.014***	0.007**	-0.286***
	(0.002)	(0.002)	(0.003)	(0.007)
Mean Success Rate	0.6140	0.8250	0.2580	0.2620
Observations	187,771	115,223	95,067	24,536
Adjusted R^2	0.60783	0.32634	0.02270	0.24723
Year FE	✓	✓	✓	✓
Job FE	✓	✓	✓	✓
Other Controls	✓	✓	✓	✓

Notes: This table reports caste differences in success rates at each hiring stage for non-client-facing jobs. The unit of observation is student-job-year, conditional on remaining in contention at each stage. All regressions include controls as well as job and year fixed effects. When firms conduct common screening rounds for multiple positions, each position is treated as a separate job. Bias-adjusted coefficients follow Oster (2019), setting $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$ (assuming unobservables match observables in selection intensity and allowing total explanatory power to exceed observed R^2 by 30%). $^*p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

Table OA.7: **Caste Differences in Mean Salary at Each Job Search Stage in the Identifiable Sample**

	Dependent variable:					
	log (Salary)					
	App	Test	Debate	Interview	Offer	Choice
	(1)	(2)	(3)	(4)	(5)	(6)
Disadv. Caste (OLS)	0.012	-0.068^{***}	-0.034^{*}	-0.105^{***}	-0.146^{***}	-0.107^{**}
	(0.014)	(0.019)	(0.018)	(0.027)	(0.033)	(0.049)
Disadv. Caste (Oster)	0.013	-0.049^{***}	-0.016	-0.078^{**}	-0.119^{***}	-0.082
	(0.016)	(0.021)	(0.022)	(0.032)	(0.042)	(0.055)
Constant	10.991^{***}	10.869^{***}	10.807^{***}	10.734^{***}	10.601^{***}	10.834^{***}
	(0.030)	(0.042)	(0.041)	(0.064)	(0.078)	(0.114)
Observations	619	552	500	449	378	378
Adjusted R^2	0.546	0.470	0.537	0.535	0.509	0.362
Year FE	✓	✓	✓	✓	✓	✓
Other Controls	✓	✓	✓	✓	✓	✓
Job FE						

Notes: This table reports earnings regressions for students with identifiable caste surnames at each job search stage. Each column represents a different stage: (1) application, (2) written test, and subsequent rounds. The dependent variable is log earnings. All regressions include controls and year fixed effects. Job fixed effects are omitted as salaries are identical within jobs regardless of caste. Bias-adjusted coefficients follow Oster (2019), setting $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$ (assuming unobservables match observables in selection intensity and allowing total explanatory power to exceed observed R^2 by 30%). $^{*}p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

Table OA.8: **Caste Differences in Mean Salary at Each Job Search Stage in the Non-Identifiable Sample**

	Dependent variable:					
			log (Salary)			
	App	Test	Debate	Interview	Offer	Choice
	(1)	(2)	(3)	(4)	(5)	(6)
Disadv. Caste (OLS)	0.002	0.0004	−0.003	−0.004	−0.100***	−0.106***
	(0.005)	(0.007)	(0.007)	(0.011)	(0.014)	(0.019)
Disadv. Caste (Oster)	0.004	0.019***	0.013	0.024**	−0.075***	−0.076***
	(0.006)	(0.007)	(0.008)	(0.011)	(0.015)	(0.018)
Constant	11.041***	10.913***	10.857***	10.766***	10.685***	10.869***
	(0.012)	(0.017)	(0.016)	(0.024)	(0.033)	(0.043)
Observations	3,386	3,066	2,766	2,582	2,149	2,149
Adjusted R^2	0.584	0.500	0.552	0.556	0.486	0.421
Year FE	✓	✓	✓	✓	✓	✓
Other Controls	✓	✓	✓	✓	✓	✓
Job FE						

Notes: This table reports earnings regressions for students with non-identifiable caste surnames at each job search stage. Each column represents a different stage: (1) application, (2) written test, and subsequent rounds. The dependent variable is log earnings. All regressions include controls and year fixed effects. Job fixed effects are omitted as salaries are identical within jobs regardless of caste. Bias-adjusted coefficients follow Oster (2019), setting $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$ (assuming unobservables match observables in selection intensity and allowing total explanatory power to exceed observed R^2 by 30%). $^*p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

Table OA.9: **Caste Differences in GPA Between Those Who Left the First Job Versus Those Who Did Not**

	Dependent variable:		
	log (GPA)		
	(1)	(2)	(3)
Disadv. Caste	-0.065^{***}	-0.065^{***}	0.010^{***}
	(0.004)	(0.004)	(0.002)
Left Job	-0.072^{***}	-0.072^{***}	-0.004^{**}
	(0.003)	(0.003)	(0.002)
Left Job * Disadv. Caste	0.047^{***}	0.047^{***}	-0.004
	(0.005)	(0.005)	(0.003)
Constant	2.188^{***}	2.182^{***}	2.198^{***}
	(0.003)	(0.005)	(0.028)
Year FE	✓	✓	✓
Major		✓	✓
Degree		✓	✓
Job FE			✓
Observations	2,489	2,489	2,489
Adjusted R^2	0.260	0.265	0.822
Note:		*p<0.1; **p<0.05; ***p<0.01	

Notes: This table reports caste differences in GPA between job stayers and leavers. Column (1) includes year fixed effects only, Column (2) adds major and degree controls, and Column (3) additionally includes job fixed effects. $^{*}p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

Figure OA.1: **Common Support Within Each College GPA Decile**

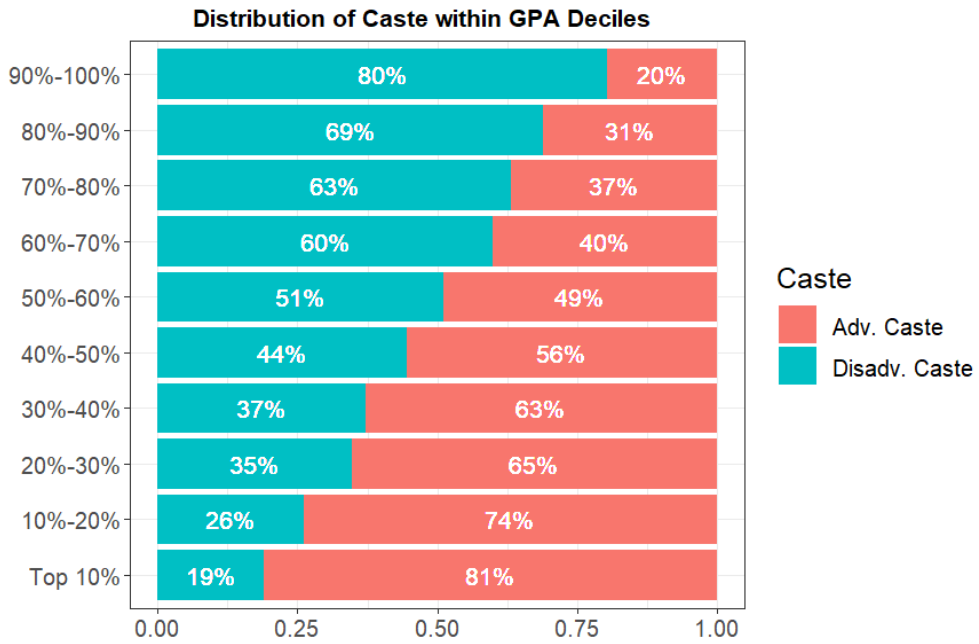

Notes: This figure shows the distribution of advantaged and disadvantaged caste students across GPA deciles. GPAs are normalized within degree-year.

Figure OA.2: **Caste Penalty in the Conditional Success Rate At the Application Reading Round in the Identifiable and the Non-Identifiable Sample**

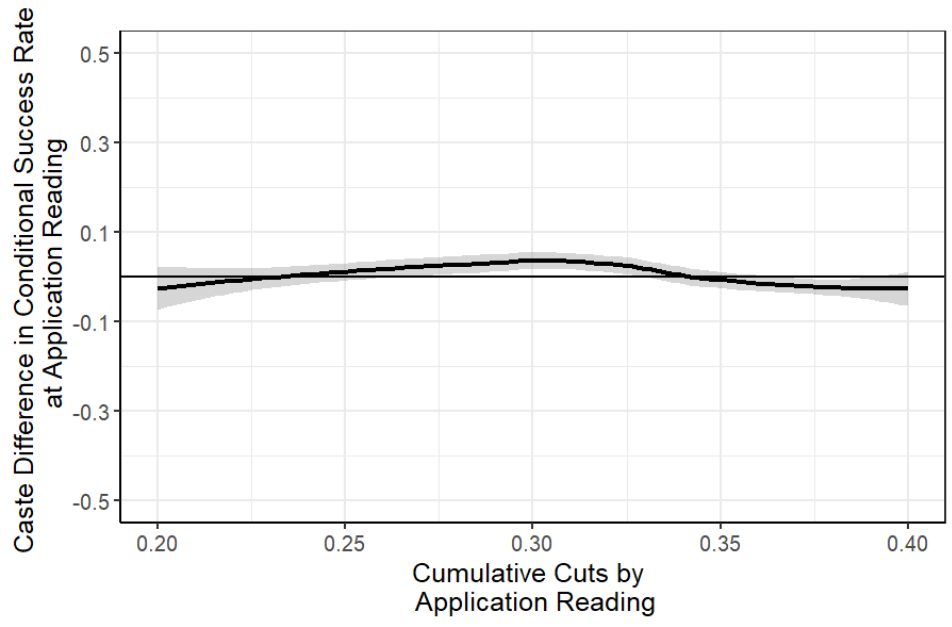

(a) Caste Penalty in Conditional Success Rates versus Cumulative Cuts by Application Reading in the Full Sample

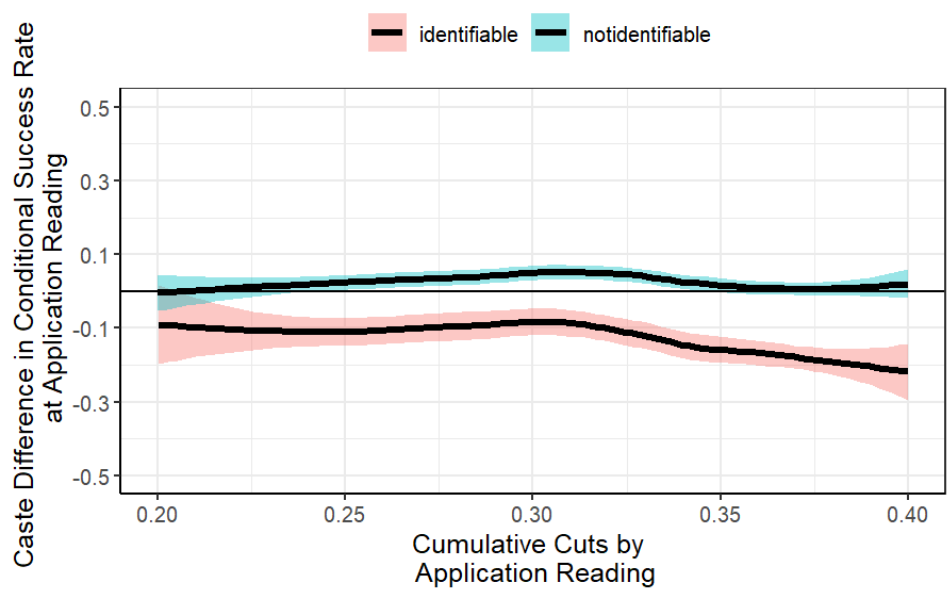

(b) Caste Penalty in Conditional Success Rates versus Cumulative Cuts by Application Reading in the Identifiable and Non-Identifiable Sample

Notes: This figure shows caste differences in success rates at the application stage against cumulative selection cuts. Panel (a) uses the full sample; Panel (b) separates students with identifiable (red) and non-identifiable (blue) caste surnames. The x-axis shows the cumulative proportion of rejected applicants. Y-values represent the disadvantaged caste coefficient from job-specific linear probability models including controls (degree, major, GPA, entrance exam score, cohort year). Black lines show smoothed conditional means with 95% confidence intervals (shaded). The horizontal line at zero indicates no caste difference.

Figure OA.3: **Caste Penalty in the Uncontrolled Success Rate At Each Screening Round**

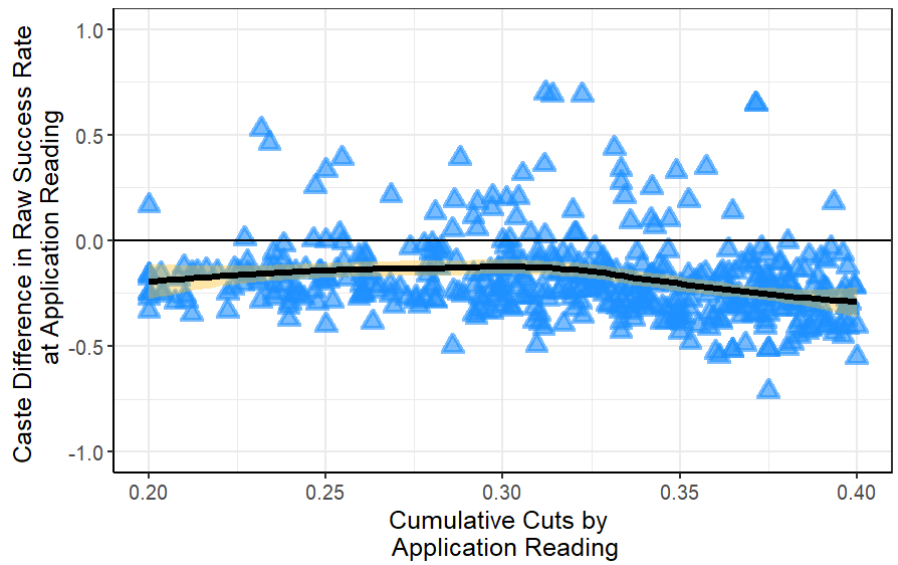

(a) Caste Penalty in Raw Success Rates versus Cumulative Cuts by Application Reading

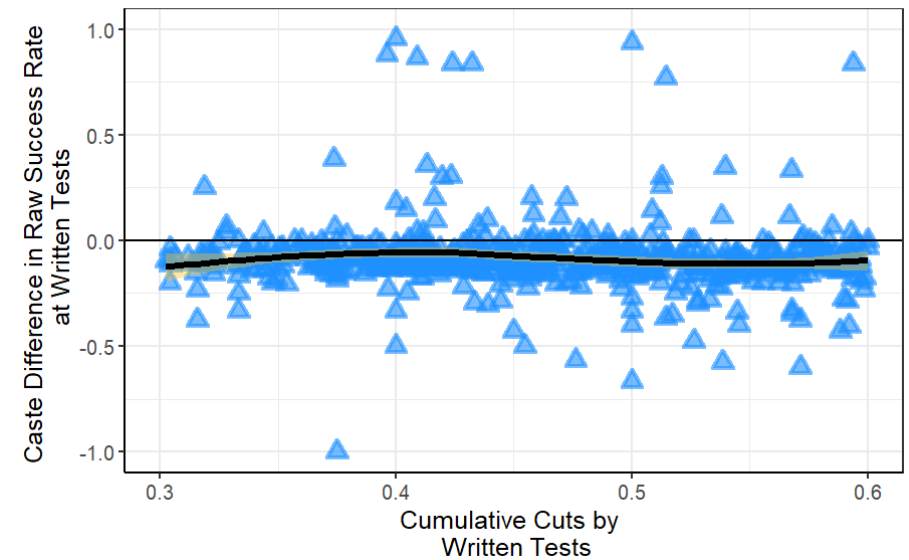

(b) Caste Penalty in Raw Success Rates versus Cumulative Cuts by Written Tests

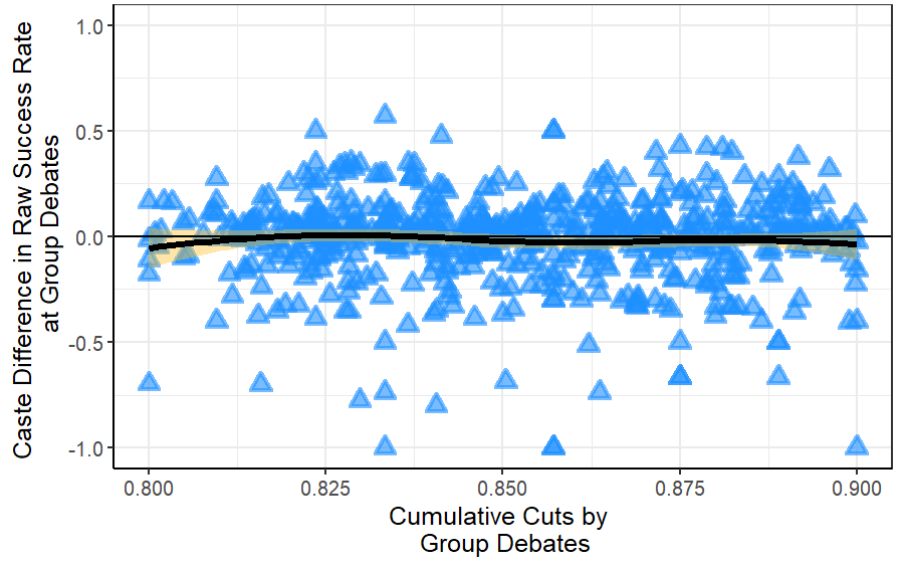

(c) Caste Penalty in Raw Success Rates versus Cumulative Cuts by Group Debates

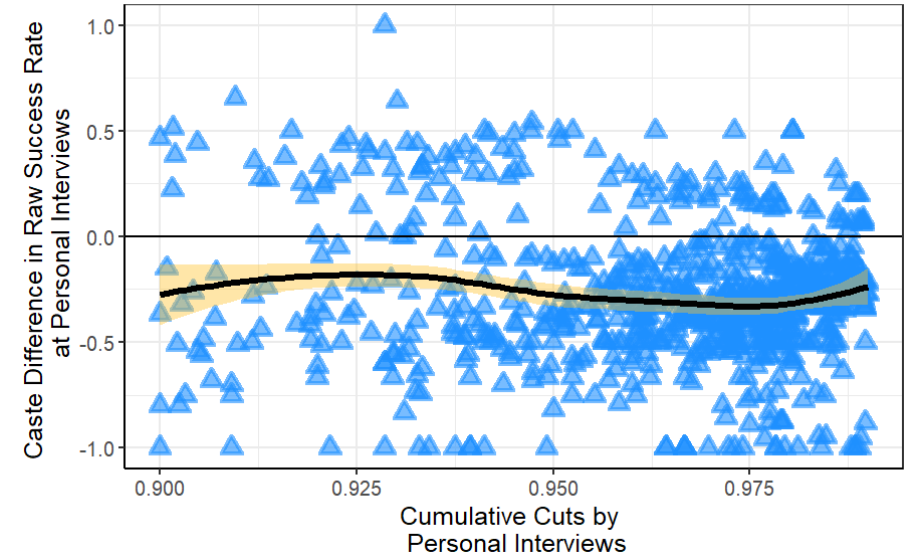

(d) Caste Penalty in Raw Success Rates versus Cumulative Cuts by Personal Interviews

Notes: This figure shows raw (unconditional) caste differences in success rates against cumulative selection cuts at each screening stage. Each blue triangle represents a job. Black lines show smoothed conditional means with 95% confidence intervals (shaded). The x-axis shows the cumulative proportion of rejected applicants at each stage. The y-axis shows the raw caste gap in success rates. The horizontal line at zero indicates no caste difference.