Making the Elite: Class Discrimination at Multinationals

Soumitra Shukla*

April 23, 2025

Abstract

How salient is hiring discrimination in driving socioeconomic gaps in elite job access, and where do these disparities emerge? Measuring discrimination based on social status or class is challenging, especially in Western contexts, where status is indirectly inferred from subtle markers like accent, education, and networks. India's caste system offers a uniquely explicit measure of social class, enabling precise identification of how status affects access to elite positions. Analyzing detailed screening data from over 1,000 hires at US and European multinationals in India, this paper provides the first systematic quantification of class bias in elite recruitment and clearly identifies its source. Subjective personal interviews assessing "fit" drive nearly 90% of caste-based hiring disparities. However, hired disadvantaged castes have approximately 20% higher promotion rates, even in highly selective roles, indicating these disparities are not justified by differences in candidate quality. Given these persistent promotion advantages and the predominantly non-customer-facing roles, results suggest bias rather than statistical or customer-driven discrimination. These findings underscore a critical organizational challenge: how can necessary screening discretion be preserved without reinforcing status-based discrimination that undermines diversity and talent allocation?

*Shukla: Research Affiliate, Department of Economics, Yale University (soumitra.shukla@yale.edu). I would like to thank Joseph Altonji, Costas Meghir, and John Eric Humphries for serving on my committee and providing guidance and feedback. I would also like to thank Jason Abaluck, Alex Albright, Peter Arcidiacono, Crémieux, Michael Dinerstein, Alessandra Fenizia, Arpit Gupta, Alex Imas, Patrick Kline, Scott Kominers, Eliza McNay, Evan Rose, Kamalini Ramadas, Jung Sakong, Rob Valletta, Jon Williams, Seth Zimmerman, seminar participants at the Yale labor workshop, the Chicago FED, Chicago Harris, Federal Reserve Board, Wharton (BEPP), Penn State, UCSD, SOLE, Southeastern Micro Labor Workshop, System Equitable Growth Conference at the New York Fed, the "Discrimination in the 21st Century Conference" organized by the Becker Friedman Institute and Chicago Booth, MIT Organizational Economics Lunch, and the NBER Race and Stratification Working Group for helpful feedback and support.

1 Introduction

Who gains access to elite jobs, and how do hiring practices shape inequality in high-stakes recruitment? Extensive research documents discrimination across labor markets (Altonji and Blank, 1999; Rivera, 2012, 2015; Friedman and Laurison, 2020), yet lacks precise evidence on *where*, *how*, and the *quantitative significance* of discrimination in competitive and real-world contexts such as multinational corporations (MNCs).

Social class or status discrimination poses particular challenges (e.g., Lamont and Lareau, 1988; Durlauf, 1999; Ridgeway, 2019; Stansbury and Rodriguez, 2024). In Western settings, class must be inferred from ambiguous indicators—accents loosely tied to backgrounds, educational credentials obscuring family advantages, and elusive social networks, often transmitted culturally across generations (cf. Bisin and Verdier, 1998). These subtle markers complicate efforts to quantify class discrimination within economics. In elite recruitment contexts, where objective skill differences among candidates are minimal, evaluators disproportionately rely on subtle signals of who "fits," amplifying the role of cultural similarity in hiring decisions (Lamont, 1992, 2009; Rivera, 2012).

Recent meritocratic trends illustrate how well-intentioned policies inadvertently reinforce class advantages. Debates over standardized versus "holistic" assessments highlight this tension. Sacerdote et al. (2025) suggest "test-optional" admissions in U.S. colleges disadvantage high-achieving lower socioeconomic students. Similarly, selective colleges excessively weigh "soft" factors (e.g., legacy and athletic recruitment), favoring high-income students despite limited predictive validity (Chetty et al., 2023). Comparable dynamics occur in India's elite civil service exams, where subjective "personality" interviews disproportionately penalize candidates based on perceived social status, despite limited predictive power for job performance (Bhavnani and Lee, 2021). Indeed, subjective judgment-based methods consistently degrade predictive validity, even among experts (e.g., Meehl, 1954, 1986; Ægisdóttir et al., 2006; Highhouse, 2008; Dana et al., 2013; Kuncel et al., 2013).

Although higher education and public-sector recruitment differ from private-sector hiring, all combine objective metrics with subjective assessments. Removing standardized tests can obscure merit (Kuncel and Hezlett, 2007; Bol et al., 2014); similarly, subjective "fit" evaluations in interviews obscure strong skills, particularly among lower-status candidates. Thus, "fit" often implicitly signals class, disadvantaging less privileged candidates across contexts (Rivera, 2012; Laurison and Friedman, 2019; Stansbury and Rodriguez, 2024).

Dual-process cognition theories (e.g., Kahneman, 2011) help understand how recruiters or admissions committees interpret status markers. Decisions often rely on quick, intuitive evaluations prone to biases rather than reflective judgments. Hiring managers frequently

base choices on rapid impressions of "fit," influenced by superficial rapport (Todorov, 2017), subtle stereotypes (Knobe and Nichols, 2008), overconfidence induced by unstructured interviews (Kausel et al., 2016), and immediate affective reactions (Pinker, 2021), with evaluations often no better than chance (Bergman and Lerner, 2025). These rapid judgments are shaped by deeply internalized cultural norms, reinforcing evaluators' implicit expectations about class-appropriate behaviors (cf. Bisin and Verdier, 2010). Such biases are especially salient when lower socioeconomic markers emerge directly through interpersonal interactions (cf. Imas, 2016; Imas et al., 2021; Imas and Madarász, 2024). Similarly, subjective evaluations via personal statements and recommendation letters often reflect assessors' biases rather than candidates' objective potential (Murphy et al., 2009; Kuncel et al., 2014), obscuring precisely which social identities—signaled by subtle markers such as vocabulary, extracurricular activities, writing style, or stated interests—are penalized (Rivera, 2015).[1]

Extensive global evidence confirms that "fit"-based evaluations foster in-group biases, systematically disadvantaging lower-status candidates by reinforcing preferences for demographically similar individuals—those sharing gender or social backgrounds—in performance assessments and professional recognition (e.g., Gorman, 2005; Sarsons, 2017; Sarsons et al., 2021; Sarsons, 2023; Benson et al., 2023). Even elite institutions struggle to overcome persistent class gaps, as privileged early-career networks shape later outcomes, perpetuating inequality across generations (Rivera, 2015; Laurison and Friedman, 2016, 2019; Stansbury and Schultz, 2023; Asher et al., 2024; Novosad et al., 2024). Elite college admissions disproportionately facilitate upward mobility for already privileged students, magnifying class-based disparities in access to top positions (Zimmerman, 2019; Barrios-Fernández et al., 2024).

While caste is specific to India, underlying discrimination mechanisms are broadly relevant. Caste consolidates economic, social, and cultural capital—financial resources, networks, and symbolic assets like education, accents, or mannerisms respectively—into a single explicit hierarchy (Bourdieu, 1986). India's caste system thus offers a uniquely explicit context for measuring how social status shapes elite job access, even when economic circumstances vary (Beteille, 1965, 1969). By contrast, Western contexts rely on multiple ambiguous indicators to approximate class (cf., Khan, 2011), complicating efforts to isolate and precisely quantify discrimination.

This paper addresses these measurement challenges using recruitment data from leading US and European multinationals hiring in India. Despite affirmative action policies, disadvantaged castes remain underrepresented in elite roles (Madheswaran, 2008). This study makes three contributions: systematically quantifying social-status discrimination in elite re-

[1] Subjective evaluation failures have been extensively documented in other high-stakes contexts like parole decisions (Carroll et al., 1982).

cruitment; identifying personal interviews as the primary mechanism revealing status biases; and providing evidence that class bias—not statistical or customer discrimination—drives entry-level hiring disparities. Crucially, by directly observing the full recruitment process—from initial applications through actual job offers and subsequent career outcomes—this paper significantly advances existing research on hiring discrimination, which typically examines only early-stage outcomes such as callback rates. Given the stakes involved—actual job attainment rather than merely progressing to initial screening—this analysis provides uniquely robust evidence on the economic consequences of discrimination. To my knowledge, each of these contributions is the first of its kind in the literature.

The key innovation demonstrates that personal interviews, often framed as evaluations of candidate "cultural fit," reveal and penalize social status. Detailed data show most caste disparities emerge during final-round interviews. Moreover, hired disadvantaged caste candidates subsequently exhibit significantly higher promotion rates, underscoring systematic undervaluation during interviews.

Elite entry-level positions significantly influence economic power, particularly in India, where disadvantaged castes constitute nearly 70% of the population. The studied elite college job fair exemplifies widespread multinational hiring practices (cf. Bloom et al., 2012). Nearly half the students attending the job fair belong to disadvantaged castes due to mandated affirmative action. This context highlights how affirmative action benefits can erode during hiring stages, as subjective, "fit"-focused interviews reinstate social barriers. Thus, policy attention is required beyond initial educational stages (Yenmez and Sönmez, 2022).

Each of the 1050 jobs studied involves four screening rounds, typically progressing from standardized assessments to more subjective evaluations: application reading, aptitude tests, group debates on socio-emotional skills, and final personal interviews assessing candidate "fit" (Newman and Thorat, 2012; Rivera, 2012, 2015). The unconditional caste earnings gap is 23%, narrowing to 10% after controlling for pre-college skills, academic performance, and other employer-relevant factors. Salaries are explicitly *identical across castes* for identical roles and *non-negotiable*, enabling stage-wise decomposition of earnings gaps by counterfactually assigning identical salaries at each stage before offers are accepted. Personal interviews alone account for nearly 90% of the conditional earnings penalty and about 95% of the caste penalty in hiring success rates, emphasizing their role in highlighting social status through evaluations of extracurricular activities, background, and cultural alignment.

Clear evidence that *caste itself* is penalized emerges among the small subset (around 15%) of students with distinctly caste-identifiable surnames, who face penalties as early as application reading. However, in urban-educated elite college contexts, surnames rarely clearly indicate caste due to migration and diverse naming conventions (cf. Jodkha, 2017;

Wiltshire, 2020). Consequently, the *overall* caste penalty at the application reading round is negligible, with disparities primarily arising in personal interviews, the main channel through which social status is revealed and penalized.

This paper does not argue that personal interviews are useless. Rather, it argues that any *additional diagnostic information* interviews provide cannot *justify* observed caste disparities. Disadvantaged caste hires consistently outperform advantaged peers by about 20% in promotion rates, even in highly selective roles, strongly suggesting disparities stem from discrimination rather than unmeasured skill deficits.

To formalize this interpretation, I develop a theoretical framework based on job selectivity—the depth into the applicant pool each firm selects from. The framework yields two implications: first, if disadvantaged candidates genuinely lacked skills, highly selective pre-interview stages would show a caste penalty, which they do not; second, if interviews accurately filtered weaker candidates, disadvantaged hires would not consistently outperform advantaged peers across selectivity tiers. The logic is straightforward: higher selectivity means average hired candidates' quality approaches the marginal (cutoff) quality. If disadvantaged hires consistently outperform advantaged peers across selectivity levels, this suggests stricter hiring thresholds (discrimination) rather than productivity differences.

Analysis of distinctive versus non-distinctive surnames strongly indicates that hiring disparities result primarily from *disparate treatment*—explicit caste-based discrimination. Persistent promotion advantages for disadvantaged hires and the predominantly non-customer-facing nature of over 75% of these roles further suggest biases—either taste-based discrimination or biased beliefs—rather than statistical or customer-driven discrimination (cf. Benson et al., 2024). Such discrimination explicitly violates Article 15 of the Indian Constitution (Ministry of Law and Justice, Government of India, 2021), severely limiting socioeconomic mobility and impeding optimal talent allocation.

This paper contributes to labor, personnel, and organizational economics by quantifying and identifying sources and mechanisms of social-status discrimination in elite hiring contexts. It uniquely examines each stage of entry-level hiring at multinational firms, pinpointing subjective, discretionary "fit"-focused personal interviews as the primary source of discrimination, accounting for nearly 90% of conditional hiring disparities. This precise quantification advances the nascent literature on class-based discrimination, which remains less developed relative to race and gender (Stansbury and Rodriguez, 2024).

The study is closely related to recent work by Stansbury and Rodriguez (2024), who document substantial class gaps in U.S. academic careers but do not isolate precise institutional sources of these barriers. Similarly, Bhavnani and Lee (2021) conceptually replicate these findings, showing subjective "personality" interviews in India's civil service exams dis-

advantage lower-status candidates without accurately predicting job performance. Although they focus on public-sector recruitment and lack detailed quantification of broader economic consequences, their findings parallel this paper's conclusions about interviews as mechanisms that penalize social status. Rivera (2012) qualitatively identifies "cultural matching" in discretionary hiring evaluations at elite U.S. firms—approximately three-quarters of these firms rank "fit" as their most critical hiring criterion—but does not empirically quantify discrimination's economic scale nor assess its implications for candidate outcomes. Leveraging India's explicit caste measure and detailed multinational recruitment data, this paper provides granular evidence on precisely where and how discrimination occurs, quantifies its magnitude, and distinguishes bias-driven discrimination from statistical or customer discrimination—critical for policy and organizational interventions.

As traditional social-status markers dilute across societies (Lenski, 1954; Grusky, 2005; Clark, 2015), class discrimination likely moves deeper into hiring processes, increasing the need for detailed hiring-stage data. This study addresses this gap, as diffuse markers in Western contexts obscure mechanisms through which status shapes hiring decisions. Moreover, findings have broader implications: subjective interviews could similarly disadvantage candidates based on other concealable stigmatized identities, such as LGBTQ+ status (Eames, 2024, 2025).

Another key contribution highlights managers' critical role in talent allocation, emphasizing their cultural capital. Recent literature underscores managerial influence on firm performance and workplace culture (e.g., Lazear et al., 2015; Hoffman and Tadelis, 2021; Espinosa and Stanton, 2022; Fenizia, 2022; Minni, 2023; Benson and Shaw, 2024; Minni et al., 2025), but this paper uniquely integrates Bourdieu's concept of cultural capital to elucidate biases in managerial assessments of "fit." Bourdieu defines cultural capital as embodied assets (e.g., etiquette, accent, mannerisms, conversational polish), objectified assets (e.g., internships at prestigious organizations), and institutionalized forms (e.g., membership in elite college clubs, alumni status from elite schools), significantly influencing access to elite positions (Bourdieu, 1986). Managers' cultural capital biases evaluations, favoring candidates with similar cultural assets (e.g., Rivera, 2012; Sadun et al., 2022). While many of these "cultural capital" markers might indeed correlate with productivity, my empirical results indicate the subjective information managers glean from interviews, influenced by shared cultural capital, cannot justify observed disparities. Moreover, while we do not have direct information about managers' personal backgrounds, evidence shows organizational socialization processes lead managers from varied backgrounds to converge on similar evaluative standards of cultural capital, reinforcing shared biases (e.g., Turco, 2010; Rivera, 2015). Thus, this study bridges cultural sociology and empirical personnel economics by empirically

demonstrating managerial biases based on cultural alignment.

Additionally, the paper highlights agency problems between hiring managers and supervisors. Hiring managers rarely supervise recruits directly, creating incentives to prioritize short-term outcomes like perceived "fit" over candidates' long-term potential. Empirical evidence, notably disadvantaged candidates' higher promotion rates despite initial undervaluation, underscores this misalignment. Indeed, candidates initially rejected based on subjective criteria frequently perform as well or better than initially accepted peers, a pattern documented in contexts like medical school admissions (DeVaul et al., 1987). Such subjective, short-term-focused decisions by hiring managers systematically produce suboptimal long-term organizational outcomes (e.g., Gibbons and Roberts, 2013; Sadun et al., 2017; Benson and Shaw, 2024; Hoffman and Stanton, 2024). Addressing these agency issues through better-aligned incentives or structured hiring practices emerges as a critical avenue for organizational reform and future research.

A further contribution of this paper is its focus on discretionary evaluations—particularly personal interviews—in elite private-sector screening, such as hiring at multinational corporations, an area receiving limited attention within economics despite significant real-world implications. While economic studies extensively document discrimination at early stages like résumé evaluations (cf. Kline et al., 2022), few empirically examine subjective evaluations in interviews, especially regarding social class. Recent research emphasizes informal, subjective criteria consistently advantage privileged candidates yet correlate poorly with job performance (e.g., Sarsons, 2017; Benson et al., 2023; Cappelli, 2019; Hoffman et al., 2018). By providing direct evidence from elite multinational hiring, this paper deepens our understanding of how subjective interview judgments enable class discrimination. Moreover, it highlights a critical limitation of early-stage interventions like résumé anonymization: reducing initial bias may merely shift discrimination to later, more subjective hiring stages.

This study also contributes to debates on standardized versus "holistic" evaluations common in elite hiring and other high-stakes processes—such as college admissions, faculty hiring, and grant awards. Typically, these processes move from standardized, objective metrics (e.g., GPA, test scores, publications) that reliably predict performance (Bol et al., 2014; Kuncel and Hezlett, 2007), toward subjective evaluations (personal statements, recommendation letters, interviews) with lower predictive validity (Murphy et al., 2009; Borneman et al., 2007; Cooper et al., 2007; Kuncel et al., 2007; Alvero et al., 2021). In my context, screening typically progresses from application reading and aptitude tests toward group debates and finally personal interviews, reflecting a general shift from standardized, objective evaluations toward increasingly subjective assessments. Recent work shows holistic evaluations disproportionately reflect evaluators' biases rather than candidates' potential (Chetty et al., 2023;

Sacerdote et al., 2025; Bol et al., 2014; Checchi and van de Werfhorst, 2018; Barg, 2013). Emerging algorithmic screening tools, intended to increase objectivity, may inadvertently replicate these biases without rigorous validation (Albright, 2024). However, appropriately designed algorithms, such as those leveraging exploration-based approaches, can enhance both efficiency and equity in hiring outcomes (Bergman et al., 2024). By empirically linking subjective evaluations to organizational inefficiencies—demonstrated by higher subsequent promotion rates among initially undervalued candidates—this research emphasizes the need for greater scrutiny and reform of subjective interviewing practices.

Finally, this paper contributes insights into the role elite educational institutions play in socioeconomic mobility and the persistent challenges of discriminatory hiring. Elite colleges like the studied Indian institution, with admission rates more competitive than Harvard or MIT (University World News, 2022), are gateways to prestigious careers. Although affirmative action policies substantially increase disadvantaged students' access, research indicates elite education's socioeconomic benefits disproportionately accrue to already privileged students, perpetuating class inequalities (e.g., Zimmerman, 2019; Barrios-Fernández et al., 2024; Chetty et al., 2023). This study empirically demonstrates that even when disadvantaged students enter highly selective institutions, subjective interview practices in subsequent hiring can erode affirmative action gains, systematically disadvantaging these candidates (cf. Bhavnani and Lee, 2021; Yenmez and Sönmez, 2022). Given multinational firms' widespread recruitment practices, findings highlight the necessity of integrating affirmative action with downstream hiring reforms. Evaluating elite colleges separately from employer hiring practices overlooks critical institutional complementarities (e.g., the alignment of selection criteria or shared networks), emphasizing integrated approaches to enhance socioeconomic mobility—particularly important given the historical co-evolution of social class with cultural and institutional factors shaping long-term economic outcomes (e.g., Bisin et al., 2023).

2 Caste as Social Status Under Affirmative Action

Jati, Varna, and Modern Caste Classifications. India's modern caste system originates from historical occupational hierarchies. At its most detailed level, the system comprises thousands of hereditary occupational groups known as *Jatis*, historically linked to trades like merchants, priests, warriors, and laborers. These *Jatis* are grouped into broader categories called *Varnas*: Brahmins (priests), Kshatriyas (warriors), Vaishyas (merchants), Shudras (laborers), and historically marginalized "untouchables" (Dalits). The British colonial administration explicitly categorized these groups, formalizing classifications in India's

1949 Constitution. The term "Scheduled" literally refers to official lists or schedules maintained by the government, explicitly naming communities as "Scheduled Castes" (SCs), "Scheduled Tribes" (STs), and later "Other Backward Classes" (OBCs)—groups defined as socially and educationally disadvantaged (Puri, 2025).

This paper adopts two broad categories aligned with India's affirmative action framework: advantaged ("upper") and disadvantaged ("lower") castes. By disadvantaged castes, I refer explicitly to Scheduled Castes, Scheduled Tribes, and Other Backward Classes. Though caste may initially seem distinct from Western notions of class, India's policies recognize caste as intertwined with economic marginalization and social status, grouping disadvantaged communities into unified administrative categories (Puri, 2025).

Historical Development of Caste-based Affirmative Action Policies. Affirmative action for disadvantaged communities emerged during British colonial rule. The Government of India Act of 1919 introduced reserved legislative seats for "depressed classes," formally recognizing their socioeconomic disadvantage. This evolved into the Government of India Act of 1935, officially designating these groups as Scheduled Castes. Post-independence, India's 1949 Constitution institutionalized quota-based affirmative action in legislatures, public higher education, and government employment, initially for Scheduled Castes and Scheduled Tribes. In the 1980s, following recommendations by the Mandal Commission, quotas expanded to include Other Backward Classes, collectively raising reservation levels to nearly 50% (Bayly, 2008).

Caste as Proxy for Social Status. Although caste originated as a hereditary hierarchy, Indian policymakers now explicitly define caste categories as comprehensive indicators of socioeconomic "backwardness," encompassing economic, educational, and social disadvantages. This approach recognizes caste as encapsulating multiple dimensions of social status—financial resources, educational credentials, social networks, and cultural capital—that collectively shape socioeconomic opportunities (cf. Bourdieu, 1986). The Mandal Commission explicitly merged lineage-based identities with socioeconomic status, reflecting caste as a consolidated marker of social class. Thus, India's administrative classification mirrors what Western contexts term "lower class."

Limited Scope of Affirmative Action in the Private Sector. India's affirmative action policies, however, do not extend to private-sector employment, including elite positions at multinational corporations studied here. Private-sector roles typically offer higher wages and better career growth but lack reservation mandates, resulting in disadvantaged castes

being disproportionately excluded (Subramanian, 2019). This exclusion parallels Western contexts, where firms often emphasize diversity based on race or gender but rarely prioritize socioeconomic diversity, thus overlooking similar class-based disparities. For instance, fewer than 10% of large US firms explicitly track socioeconomic diversity (Stansbury and Rodriguez, 2024). Similarly, fewer than 3% of India's top firms (e.g., by market capitalization or revenue) explicitly maintain caste data internally (BusinessLine, 2018).[2] Consequently, even robust affirmative action policies in public sectors fail to significantly address structural inequalities in private-sector hiring.

Collectively, India's affirmative action framework highlights caste-based disadvantage as analogous to socioeconomic "backwardness," clearly recognizing how caste reinforces economic barriers and social exclusion. Given that caste distinctly encodes social and economic status, it provides clarity often lacking in Western class measures, making it a uniquely effective metric for analyzing class-based inequalities.

3 Key Definitions and Institutional Setting

This section defines key terms, describes the elite college's job fair studied in this paper, and clarifies why this job fair is representative of how elite Indian college graduates enter elite jobs.

3.1 Key Terms: Disadvantaged Castes, Elite Jobs, and Elite Colleges

Below, I define four key terms used throughout the analysis:

a. **Disadvantaged castes.** Following the Indian government's official classifications, the term "disadvantaged castes" encompasses individuals belonging to Scheduled Castes (SCs), Scheduled Tribes (STs), or Other Backward Classes (OBCs).

b. **Job.** A "job" refers specifically to a job designation within a firm. For example, if Google recruits both a product manager and a software engineer, these constitute two distinct jobs. Section 4.6 explains why a job typically corresponds to a firm-designation pair, rather than a particular office location.

[2] Although a few multinational corporations have recently begun addressing caste discrimination explicitly (Chakravartty and Subramanian, 2021; Bapuji et al., 2023), these changes occurred in late 2022 or 2023, which is after the pre-COVID timeframe of the data analyzed in this paper (Section 4.4).

c. **Elite jobs.** An "elite entry-level job" refers to a formal entry-level position at multinational corporations (MNCs), major national corporations, or unicorn startups (privately-held startups valued at over $1 billion). All elite jobs analyzed in this paper pay salaries within the top 1% of entry-level wages in India (Economic Advisory Council to the Prime Minister, 2022). Throughout this paper, the terms "elite entry-level jobs" and "elite jobs" are used interchangeably.

d. **Elite colleges.** "Elite colleges" include institutions consistently ranked among the top 30 in their respective fields—such as science, engineering, arts and humanities, social sciences, and business—by *India Today*, an Indian equivalent of the *US News & World Report*. Prominent examples include the Indian Institutes of Technology (IITs), Indian Institutes of Management (IIMs), National Institutes of Technology (NITs), and select Delhi University colleges such as St. Stephen's, Shri Ram College of Commerce, and Lady Shri Ram College.

3.2 Structure of the On-Campus Job Fair

The elite college's job fair occurs *entirely on campus* and comprises three phases: (1) pre-placement, (2) placement, and (3) post-placement (or the "aftermarket").

a. **Pre-Placement Phase.** Firms are invited by the placement office before June. Between June and mid-August, recruiters deliver pre-placement talks, gauge student interest, and finalize roles. By late August, firms extend pre-placement offers (PPOs) to students who interned with them during the summer. Students accepting PPOs are withdrawn from the main placement process by the college's placement office.

Significance of return offers from internships. Unlike in the US, PPOs play a minor role; fewer than 5% of students accept these offers, leaving over 95% to participate in the main job fair. Less than 2% choose public-sector employment (Section 4.6).

Student and employer registration. Students register for placement by late August, while employers submit detailed registration forms by early September, specifying job roles, compensation, and vacancies, which become fixed thereafter. Compensation is *non-negotiable*, strictly enforced by the placement office through verification of offer letters. Non-compliance can result in restrictions on future campus visits by firms.[3]

[3]Such rules and timelines are common across elite Indian colleges and are available on their websites. For instance, see Central Placement Cell (Delhi University).

b. **Placement Phase.** Starting in mid-September, students apply for jobs. Employers initially screen applications, followed by written aptitude tests and group debates assessing socio-emotional skills from September to early December. Final-round personal interviews occur between December and January, after which approximately 70% of the cohort secure full-time roles.

c. **Post-Placement Phase ("Aftermarket").** From late January until graduation in July, the placement office operates an "aftermarket" for students without initial job offers—about 30% of the cohort. These unplaced students pursue entrepreneurship, graduate studies, preparation for competitive exams such as civil services or elite MBAs, or find alternative employment.

Excluding students and firms in the aftermarket. Employers in the aftermarket typically differ from those in the main cycle, involving smaller or newer startups offering lower salaries and less standardized recruitment processes. Because this analysis explicitly focuses on standardized hiring, I exclude aftermarket students and employers. Implications of this exclusion are addressed further in Section 5.2.

3.3 Why the Studied Job Fair is Representative

The studied job fair is representative of how elite Indian graduates transition into elite jobs due to four factors (further details and evidence are provided in Online Appendix Section A.1). First, nearly all top-ranked Indian colleges (defined in Section 3.1) use similar standardized procedures, making this fair typical of peer institutions. Second, elite employers—multinationals, major national firms, and unicorn startups—recruit consistently from the same top-tier colleges. These firms visit 30–40 colleges annually, collectively impacting around 50,000 students. Because these employers draw from a common institutional network, their recruitment strategies and job offerings remain consistent.

Third, elite entry-level roles almost exclusively recruit from top-ranked colleges, whose graduates predominantly fill these positions. This mutual exclusivity is documented in prior research (e.g., Subramanian, 2019; Jodhka and Naudet, 2019), creating an integrated elite labor market. Thus, the studied college's employment outcomes directly mirror national trends. Fourth, very few graduates opt out of the formal placement process. In this setting, fewer than 5% of students choose to forego participation. Comparable participation rates at peer institutions confirm the representativeness of the studied fair (e.g., Subramanian, 2019). Most students rely heavily on their college's placement office rather than pursuing off-campus opportunities, further reinforcing representativeness.

4 Detailed Data on Students, Jobs, and Recruitment Process

This section describes the data used in the analysis, detailing student and job characteristics, recruitment procedures, and career outcomes such as promotions and tenure. It concludes by discussing relevant data limitations.

4.1 Institutional Source and Completeness of Recruitment Data

Data were obtained directly from the university's placement office, which meticulously records each stage of recruitment, including applications, screening rounds, personal interviews, and final job offers. The university is one of India's highest ranked institutions and thus qualifies as an elite college (defined in Section 3). Its recruitment practices benchmark placement processes at comparable institutions. The university's identity is anonymized for confidentiality.

4.2 Centralized Job Portal

Students register for the full-time placement process through a centralized online job portal managed by the placement office, creating comprehensive personal profiles that include standardized academic information (such as GPA, entrance exam scores, major coursework, and grades), résumé details (including internships and extracurricular activities), and relevant personal background information. These standardized profiles facilitate uniform screening and evaluation by employers throughout the placement process.

Firms similarly create detailed profiles on the portal by completing standardized employer registration forms, specifying non-negotiable salary structures, job amenities, role descriptions, and required qualifications (Section 3.2). Throughout the recruitment cycle, employers directly upload candidate shortlists from each screening round onto this portal, providing systematic documentation accessible to the placement office for internal verification. This centralized approach ensures comprehensive, consistent, and reliable data collection at every stage, significantly enhancing the empirical rigor and transparency of the recruitment analysis.

4.3 Standardized Recruitment Rules Ensuring Data Integrity

Recruitment fairs at elite Indian colleges follow standardized procedures to ensure data integrity.

Mandatory participation in all screening rounds. Every job conducts four mandatory screening rounds (Section 4.8). The placement office enforces strict compliance, prohibiting firms from informal offers, and requiring student participation at every screening stage for a job they applied to. These rules, common across elite Indian institutions, ensure data consistency.[4]

4.4 Student Sample: Time Span, Degrees Offered, and Caste-Quota Admissions

The dataset spans four placement cycles (years), all occurring prior to the COVID-19 pandemic. Exact years are omitted to protect anonymity further.

Balanced caste representation via affirmative action in admissions. Due to caste-based affirmative action in public college admissions (Section 2), caste representation is approximately balanced across majors. Table 1 summarizes the distribution of the 4,164 students in the dataset, about evenly split between advantaged and disadvantaged castes.

Administrative reporting of student caste. Students self-report caste by submitting official caste certificates, enabling accurate administrative tracking by both the college and the researcher.

Predominantly male student population. About 90% of the sample is male, reflecting the typical demographic composition of elite technical colleges in India (Datta, 2017).

Degrees offered. The dataset includes four degrees (see Table 1): Bachelor of Technology (B.Tech.), Dual Degree, Master of Technology (M.Tech.), and Master of Science (M.S.). The Dual Degree program integrates undergraduate and postgraduate coursework, requiring an additional year beyond the standard four-year B.Tech. curriculum.

Admissions based solely on standardized national entrance exams. Admissions are *exclusively* based on standardized national-level exams with caste- and major-specific cutoffs.

[4] Although exploring the rationale behind these rules is beyond this paper's scope, detailed policies are publicly accessible on most elite colleges' placement websites. For example, see Central Placement Cell (Delhi University).

Table 1: **Total Students by Caste and Degree**

	Total Students		
	Advantaged Caste	Disadvantaged Caste	Degree Total
B.Tech.	579	710	1289
Dual Degree	622	617	1239
M.Tech.	601	566	1167
M.S.	344	125	469
Caste Total	2146	2018	4164
Fraction	0.52	0.48	1.00

Notes: Table 1 shows the total number of students by caste within each college degree, the total number of students by caste, and the fraction of students by caste. B.Tech. stands for Bachelor of Technology, M.Tech. stands for Master of Technology, and M.S. stands for Master of Science. A Dual degree integrates undergraduate and post-graduate studies and is completed a year after the conventional four-year degree (B.Tech.).

4.5 Student Characteristics and Caste-Based Differences

Key characteristics include college GPA, pre-college skills, and prior labor market experience, highlighting two empirical patterns:

1. **Substantial caste differences in academic performance metrics.** Large caste disparities are evident in both college GPA and entrance exam (EE) scores (Table 2), although students from both castes span all college GPA deciles (Online Appendix Figure OA.1) and EE score deciles (not shown).[5]

2. **Weak correlation between caste and previous labor market experience.** The data include details on summer and winter internships covering duration, total pay, employment sector, and startup status. Internship descriptions typically specify eligibility criteria based on factors such as major, degree, and GPA, along with desired skills.

   Caste is only modestly correlated with prior labor market experience. Among B.Tech. and Dual Degree students, internships are similarly distributed across castes: about 25% in technology, 37% in consulting, and 38% in manufacturing. Similar patterns hold for M.Tech. and M.S. students (not shown).

   This weak correlation likely arises because summer internship return offers play a minor role at elite Indian colleges (Section 3.2). Students generally prioritize internships based on preferred sectors or firms rather than internship pay, limiting caste-based differences.

[5] EE scores were originally recorded as ranks but have been renormalized such that higher scores indicate better performance.

Table 2: **Summary Statistics (Students)**

	Adv. Caste	Disadv. Caste	Diff (SD)
		B.Tech.	
Avg. Entrance Exam Score	0.41	-0.33	0.74***
Avg. 10th Grade Score	0.24	-0.19	0.43***
Avg. 12th Grade Score	0.24	-0.20	0.44***
Avg. Overall College GPA	0.51	-0.42	0.93***
		Dual Degree	
Avg. Entrance Exam Score	0.35	-0.35	0.70***
Avg. 10th Grade Score	0.20	-0.20	0.40***
Avg. 12th Grade Score	0.14	-0.14	0.28***
Avg. Overall College GPA	0.43	-0.43	0.86***
		M.Tech.	
Avg. Entrance Exam Score	0.27	-0.30	0.57***
Avg. 10th Grade Score	0.23	-0.15	0.38***
Avg. 12th Grade Score	0.17	-0.10	0.27***
Avg. Overall College GPA	0.34	-0.33	0.67***
		M.S.	
Avg. Entrance Exam Score	-0.02	0.02	-0.04
Avg. 10th Grade Score	0.03	-0.07	0.10
Avg. 12th Grade Score	0.02	-0.06	0.08
Avg. Overall College GPA	0.04	-0.14	0.18*

Notes: Table 2 documents some select summary statistics for students on measures of pre-college skills and college academic performance. Pre-college skill measures include college entrance exam scores and scores in the 10th and 12th standard national-level exit exams. College academic performance is measured by college GPA. All scores are pooled and normalized to have zero mean and unit standard deviation within each degree and year. College entrance exam scores are originally ranks that have been renormalized so that higher numbers are better. The difference across castes is reported in standard deviation units. $^{*}p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$. Adv. Caste stands for advantaged caste and Disadv. Caste stands for disadvantaged caste.

4.6 Characteristics of Multinational Private-Sector Jobs

This subsection summarizes key job characteristics, including sector distribution, salary structure, job location, recruitment processes, and managerial involvement.

Sector and employer overview. The dataset covers approximately 1,050 private-sector jobs, mostly at multinational corporations headquartered in the US or Europe, recruiting primarily for their Indian offices (Table 3). About 75% are non-client-facing roles (e.g., software engineering), while 25% are client-facing (e.g., consulting, managerial positions).

Exclusion of public-sector jobs. Public-sector jobs (fewer than 2%, see Table 3) are excluded due to substantial differences in salary structure and job stability.[6]

Standardized salaries and benefits. Salaries are not match-specific. They are *standardized, non-negotiable*, and *uniform* for each job role. Conditional on college degree, they do not vary by individual characteristics like major, caste or gender. Firms typically adopt nationwide salary standards to simplify administration, as they recruit from multiple elite colleges for identical job profiles.[7] Non-wage benefits such as health insurance, housing assistance, and signing bonuses are similarly uniform.

Job location and recruitment procedures. More than 98% of jobs are based in India, with unified recruitment processes across locations. Thus, a "job" generally refers to a firm-designation combination rather than a particular Indian office location. Less than 2% involve placements abroad and are treated separately due to distinct hiring processes and compensation packages.

Recruitment by Indian hiring managers distinct from direct supervisors. Recruitment is conducted by specialized hiring managers—predominantly Indian professionals—who remain separate from eventual direct supervisors due to logistical constraints posed by campus visits across India.

Summer internships and return offer policies. Students accepting summer internship return offers are excluded from the main placement cycle (Section 3.2).

[6]For further details on public-sector pay and stability, see Seventh Central Pay Commission, 2016.

[7]This "national wage setting" practice has also been documented in elite entry-level hiring in the US (Bloom et al., 2012; Sarsons et al., 2022).

Table 3: **Summary Statistics (Jobs)**

	Count	Fraction	Avg. Salary ($)	Std. Dev. ($)
Private	1011	0.98	60341.38	32591.51
Government	19	0.02	56202.23	20240.01
Headquartered Outside India	620	0.60	64562.74	36038.04
Headquartered in India	410	0.40	53766.04	23946.16
Client Facing	246	0.24	57553.23	28288.23
Non-Client Facing	784	0.76	61115.92	33361.87
All Jobs	1030	1.00	60265.02	32391.89

Notes: Table 3 shows the proportion of jobs in the public versus private sector, the proportion of jobs headquartered outside versus in India, and the average salary breakdown (USD PPP) among client-facing jobs and non-client-facing jobs. Firms headquartered outside India are mostly based in the US or Europe.

4.7 Quantifying Job Selectivity at Each Recruitment Stage

This subsection summarizes the selectivity of jobs and defines two critical concepts—conditional and cumulative cuts—used to measure selectivity at different recruitment stages.

Conditional and cumulative cuts defined. Table 4 summarizes job selectivity using two concepts: *conditional cuts*, the fraction of applicants eliminated *in each round*, and *cumulative cuts*, the total fraction eliminated *up to a given round*. For instance, if a job eliminates 50% of applicants in each of the first two rounds, its conditional cut at each round is 50%, but the cumulative cut after round two totals 75%. Successful candidates typically rank within the top 5% of initial applicants. Some jobs occasionally extend no offers, reflecting highly selective recruitment processes.

Table 4: **Variation in Conditional and Cumulative Selectivity Across Jobs**

Cut Type	Round	Q1	Mean	Q2	Q3
Conditional Cut	App. Reading	0.30	0.39	0.39	0.48
Cumulative Cut	App. Reading	0.30	0.39	0.39	0.48
Conditional Cut	Written Test	0.10	0.18	0.15	0.26
Cumulative Cut	Written Test	0.38	0.49	0.49	0.61
Conditional Cut	Group Debate	0.69	0.74	0.74	0.79
Cumulative Cut	Group Debate	0.83	0.87	0.87	0.91
Conditional Cut	Personal Interview	0.61	0.75	0.78	0.90
Cumulative Cut	Personal Interview	0.95	0.96	0.97	0.99

Notes: Table 4 shows the distribution of both the conditional and cumulative cuts across jobs.

4.8 Four Recruitment Stages: From Objective to Subjective Evaluations

This subsection describes the four screening rounds uniformly employed by all jobs, transitioning from objective assessments to increasingly subjective evaluations.

Low-cost applications via centralized portal. Students apply through a centralized online portal by selecting job listings without additional documents. Employers automatically receive standardized resumes, significantly lowering application costs and ensuring comparable application rates across caste groups.

Eligibility criteria restrict applications. Applications are restricted by major, degree, and GPA. For instance, a software engineering role at Microsoft typically requires Computer Science majors, while analyst roles at Goldman Sachs may accept all majors.

Four mandatory screening rounds. All jobs follow a structured sequence: (1) application screening (early to mid-September), (2) written aptitude tests (late September to October), (3) group debates (October to November), and (4) personal interviews (December to January).

1. **Application screening.** Employers initially screen resumes based on standardized criteria and typically apply GPA cutoffs, though limited subjective evaluation occurs.

2. **Written aptitude tests.** Advancing candidates take standardized tests online assessing technical and analytical skills, producing comparable performance metrics.

3. **Group debates.** Candidates participate in structured debates on common, non-technical prompts. Evaluations consider subjective aspects such as communication and confidence but maintain standardized conditions (e.g., same prompt, equal speaking time per candidate).

4. **Personal interviews.** Final interviews, lasting about 30 minutes, evaluate candidate fit, interpersonal skills, cultural alignment, and job-specific competencies, relying heavily on subjective judgments.

Thus, recruitment transitions from relatively more objective criteria to nuanced interpersonal evaluations.

Common screening rounds across positions. For instance, if Nvidia offers five distinct roles in a given recruitment cycle, it may administer one common technical test simultaneously for all applicants. Additional details are provided in Online Appendix Section A.2.

4.9 Candidate Screening: Summary Statistics

This subsection summarizes the total number of applicants and students advancing through each screening round. Table 5 reports the average (median) number of applications per job as 208.85 (140.17). At each subsequent round, the average (median) number advancing are: written aptitude tests, 124.91 (87.00); group debates, 101.54 (70.00); personal interviews, 26.07 (18.00); and job offers, 6.49 (4.00). The substantial decline at each stage reflects rigorous selectivity (Table 4). Eligibility criteria based on major, degree, and GPA effectively limit the number of applications per student (Section 4.8).

Table 5: **Summary Statistics (Candidate Screening)**

Round	Mean	Median
Application Reading	208.85	140.17
Written Tests	124.91	87.00
Group Debates	101.54	70.00
Personal Interviews	26.07	18.00
Offers	6.49	4.00

Notes: Table 5 shows the mean and median number of applications, written tests conducted, group debates conducted, interviews conducted, and offers extended per job. For a given screening round, these numbers can be calculated by averaging, per job, the number of students across years and then averaging across jobs.

4.10 Data on Job Tenure and Promotions

I measure job performance using two indicators—tenure and promotions—by linking the sample to an employment-focused online platform.[8] I successfully match over 98% of students hired through the college's job fair, cross-referencing student profiles with internship histories, initial job placements, majors, degrees, and graduation years. Tenure, measured in months, is calculated using start and end dates reported on the platform.

Defining and classifying promotions. Promotions are movements to higher job titles within the same firm. Robust classification requires: (1) verification of salary increases associated with new titles using external industry-standard salary benchmarks (platform anonymized), and (2) confirmation that job title progression aligns with typical career trajectories. This classification is particularly reliable for elite, entry-level jobs at multinational firms, where career ladders and salary structures are standardized globally (cf. Bloom et al., 2012; Minni, 2024). For example, advancing from "Analyst" to "Associate" within two years at a consulting firm typically qualifies as a promotion.

Distinguishing vertical career moves. The above approach reliably distinguishes genuine vertical advancements from lateral job changes. Combined with structured compensation practices typical at elite multinationals, these metrics provide robust measures of employee performance and career progression.

[8] Platform name anonymized for confidentiality.

4.11 Data Limitations and Mitigation

There are two primary limitations. First, I lack direct access to student resumes. Nevertheless, all students use a centralized job portal to enter extensive standardized information, ensuring most relevant candidate data are already reflected in my dataset (e.g., entrance exam scores, major, GPA, internships, and job histories) (Section 4.2). Second, I do not observe precise performance scores or evaluator ratings from any screening round. While these limitations restrict direct visibility into intermediate evaluative judgments, I discuss strategies for mitigating their impact and interpreting the results appropriately in Section 5.7.

5 Penalizing Social Status Through "Fit" Assessments

This paper documents five key findings on how personal interviews penalize social status (caste). First, a substantial—and likely conservative—earnings gap exists between castes. Second, about 90% of hiring disparities emerge only after subjective "fit"-focused personal interviews. Third, caste penalties consistently arise during these interviews across nearly all jobs. Fourth, employers explicitly penalize caste once revealed, making interviews the primary channel for discrimination. Finally, disadvantaged caste hires show higher promotion rates even in highly selective roles, suggesting discrimination rather than unobserved skill differences drives hiring inequalities.

5.1 Main Sample and Employment Accuracy

As noted in Section 3.2, about 70% of students secure jobs through the regular placement process by early January, with fewer than 5% accepting summer internship return offers. Consequently, over 95% actively participate in the standardized four-stage job fair, resulting in approximately 2,550 hires (65% of the entire sample).

Reliability of initial job placements. Virtually all hires begin in their offered roles, validating the placement office's data as a reliable measure of actual employment.[9]

5.2 Substantial Earnings Gap Across Castes

Disadvantaged castes earn significantly lower initial salaries. The unconditional earnings gap is 23%, narrowing to about 10% after controlling for pre-college skills, college GPA,

[9] A mandatory post-graduation exit survey (98% response rate), conducted two months after graduation, confirms nearly all hires started in their offered positions without compensation renegotiation. Cross-validation with data from a prominent employment platform further supports placement record accuracy.

prior labor market experience, and other factors. This residual 10% gap—approximately 6,000 (PPP), or roughly 0.18 standard deviations of the private-sector salary distribution (Table 3)—reflects differences in job composition rather than pay within identical roles, given standardized salaries (Section 4.6). Online Appendix Table OA.1 details modest variations by job type.

Students excluded from main earnings analysis. The earnings analysis excludes students who (1) accept summer internship return offers or (2) remain unplaced after the main cycle and enter the "aftermarket" (Section 3.2).

Implications of exclusion. While no fully participating student rejects all offers, approximately 30% of students are excluded due to internships or aftermarket entry. Among these, disadvantaged castes have lower average GPAs compared to advantaged peers (Online Appendix Table OA.2).[10] Because higher GPAs generally correlate with higher pay, omitting disproportionately lower-GPA disadvantaged students likely understates the true earnings gap.

5.3 Caste Earnings Penalty Arises Predominantly During Personal Interviews

Sections 5.3–5.7 demonstrate that personal interviews—emphasizing subjective "fit" criteria such as communication style, personality, and cultural alignment—are the main stage where caste discrimination emerges, accounting for nearly 90% of conditional hiring disparities.

Because jobs have *fixed, non-negotiable salaries* (Section 4.6), caste earnings penalties across screening stages reflect differences in job access rather than candidate-specific salary variations. I quantify these penalties at each stage by counterfactually assigning the posted salaries of available jobs to each caste group before final offers are accepted. Formally, I estimate variations of the following regression:

$$\log(\text{Avg. Job Salary}_{it}^{\text{Search Stage}}) = \alpha + \beta \times \text{Disadv. Caste}_i + \text{Controls}_i + \gamma_t + \epsilon_{it}, \quad (1)$$

where $i \in \mathcal{I}$ indexes students, and "Search Stage" $\in$ {Application, Written Tests, Group Debates, Personal Interviews, Offers, Accepted Offers}. Controls include GPA, entrance exam scores, prior internships, major, and degree; γ_t denotes cohort-year fixed effects. For each student,

[10] *Within* any given job, disadvantaged hires have slightly higher GPAs than advantaged counterparts (Online Appendix Table OA.3).

the "average job salary" at a given stage is the mean salary of all roles for which the student remains eligible.

The coefficient β, illustrated in Figure 1, quantifies the caste earnings gap at each recruitment stage. Full regression results with robustness checks using Oster (2019) bounds, addressing potential omitted-variable biases, are presented in Online Appendix Table OA.4.[11]

Figure 1: **The Caste Earnings Penalty**

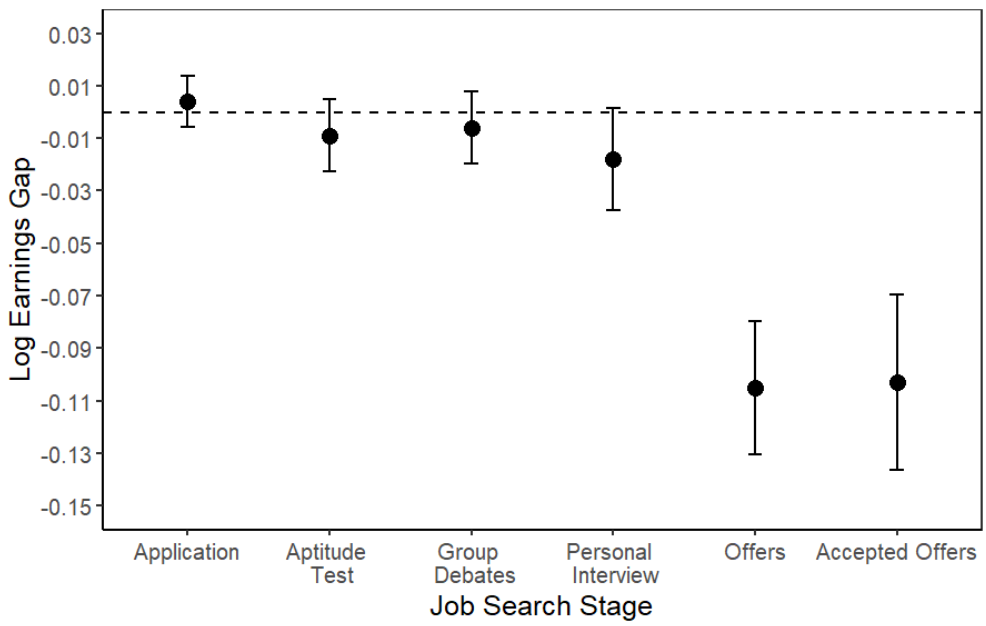

Notes: Figure 1 shows the log earnings gap across castes at successive stages of job search. Each coefficient in the figure is represented by a black dot and reports the percentage difference in the average salary at each job search stage between disadvantaged and advantaged castes. The vertical bars are 95% confidence intervals. These regressions include controls. There are no caste differences in pay for identical jobs and since offered salaries are non-negotiable during the job fair, the caste penalty in mean earnings at each stage can be calculated by comparing the salary differences in the jobs the caste groups are left in contention for after each stage.

Robustness and job-level heterogeneity. To address concerns about aggregate confounding factors, Section 5.7 analyzes heterogeneity in job selectivity. Specifically, I plot job-level caste penalties against cumulative applicant elimination rates ("cumulative selectivity," defined in Section 4.7). Disadvantaged castes face no early-stage penalties—even for highly selective jobs—ruling out pre-interview skill deficits as explanations.

The next subsection further reinforces these results by analyzing caste differences in success rates, confirming the primacy of personal interviews in driving caste disparities.

[11]We apply Oster (2019)'s bounding methodology, setting $\delta = 1$ (assuming unobservables match observables in selection intensity) and $R^* = 1.3 \times \tilde{R}$, allowing the total explanatory power to exceed the observed R^2 by 30%. The main findings remain robust under these stringent assumptions.

5.4 Caste Penalty in Hiring Success Rates Emerges Primarily at the Interview Stage

This subsection analyzes success rates—the probability of advancing through each screening round for a given job—as a complementary measure to the earnings analysis in Section 5.3. Because salaries are standardized, caste disparities in earnings should reflect job composition differences rather than within-job success rate gaps. Thus, if both castes face similar job opportunities at each stage, within-job success rates should not differ significantly by caste until the interview stage. This analysis tests and confirms this logic.

Defining and measuring success rates. A student's success rate at each round is the probability of progressing to the subsequent stage for a given job. Using linear probability models with job and year fixed effects, I estimate caste differences at each stage:

$$\Pr(\text{Success}_{ijt}^{\text{Job Search Stage}}) = \alpha + \beta \times \text{Disadv. Caste}_i + \text{Controls}_i + k_j + \gamma_t + \epsilon_{ijt}, \qquad (2)$$

where i indexes students, j indexes jobs, and the screening stages include application screening, written tests, group debates, personal interviews, job offers, and accepted offers. Controls include GPA, entrance exam scores, prior internships, major, and degree; k_j and γ_t represent job and cohort-year fixed effects, respectively.

Results: Caste penalties emerge during interviews. Table 6 shows negligible caste differences in success rates at the application, written-test, and group-debate stages. However, a significant caste penalty appears at the personal interview stage, accounting for over 95% of the final disparity in job offers. Disadvantaged castes are about 30 percentage points less likely to advance than advantaged castes after personal interviews (mean success rate is 26%). This finding holds for both client-facing and non-client-facing roles (Online Appendix Tables OA.5–OA.6).

Robustness checks. One concern is that early-stage disparities might be obscured if disadvantaged candidates disproportionately fail early rounds only for highly selective roles. Section 5.7 addresses this explicitly, showing no pre-interview caste disparities even in highly selective jobs. Thus, results underscore personal interviews as uniquely generating caste disparities, rather than merely reflecting pre-interview differences in candidate quality.

Having confirmed the centrality of interviews, the subsequent analysis explores precisely *when* and *how* caste signals become salient.

Table 6: **Caste Differences in the Success Rate at Each Job Search Stage**

	Success App (1)	Success Test (2)	Success Debate (3)	Success Interview (4)
Disadv. Caste (OLS)	-0.0025*	-0.0002	-0.0005	-0.2938***
	(0.0014)	(0.0018)	(0.0027)	(0.0048)
Disadv. Caste (Oster)	0.009	0.010***	-0.002	-0.296***
	(0.0013)	(0.002)	(0.003)	(0.005)
Mean Success Rate	0.5990	0.8110	0.2580	0.2610
Observations	305,541	183,150	148,489	38,378
Adjusted R^2	0.61069	0.34360	0.01858	0.24450
Year FE	✓	✓	✓	✓
Job FE	✓	✓	✓	✓
Other Controls	✓	✓	✓	✓

Notes: Table 6 shows the caste difference in the success rate at each stage of job search. Each observation is a student who remains in contention at a given screening round in a given job in a given year. Each column reports a separate linear probability regression and includes controls. These regressions also include job fixed effects. Note that some firms conduct the same screening round (written tests, group debates, interviews) for multiple job positions. However, since these job positions within the same firm represent different jobs, the regressions consider these jobs and the corresponding advance-or-out decisions made by them separately. Bias-adjusted coefficients are calculated using Oster (2019). Following Oster's recommended benchmarks, I set $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$, implying that unobservables are assumed to have the same degree of selection as observables, and the total explanatory power can increase by up to 30% beyond the observed R^2. $^*p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

5.5 Information Observed by Employers at Each Screening Round

This subsection summarizes information available to employers at each recruitment stage. Employers continuously access students' standardized academic records—including GPA, coursework details, entrance exam scores, and selected additional information—entered directly by students into a centralized application portal, complemented by uploaded resumes (Section 4.2).

1. **Application reading.** Employers initially review standardized resumes detailing academic credentials, internships, extracurricular activities, and technical skills. GPA cutoffs are often used to efficiently narrow the applicant pool. Additional standardized details (e.g., entrance exam scores, coursework) from the application portal remain accessible but secondary to resumes at this stage.

2. **Written aptitude tests.** Candidates advance primarily based on their performance on standardized, job-specific online tests. Although resumes remain accessible, evaluations depend predominantly on objective test scores.

3. **Group debates.** Employers directly observe structured group debates, assessing communication, teamwork, and socio-emotional skills. Although academic records and test results remain accessible, debate performance is the primary evaluation criterion.

4. **Personal interviews.** Final personal interviews, lasting roughly 30 minutes per candidate, involve individual assessments. Employers revisit resumes—examining GPA, exam scores, internships, and skills—but primarily emphasize subjective evaluations of cultural fit, interpersonal rapport, job-specific competencies, and organizational alignment.

Personal interviews allow employers to reassess earlier signals in greater depth, enhancing their informational accuracy regarding candidate quality. Crucially, the individualized and subjective nature of these interviews provides additional indicators of cultural fit and interpersonal compatibility through conversations about interests, extracurricular activities, and personal backgrounds. Such in-depth conversations not only amplify previously observed signals but also enable employers to glean richer, subtler insights, potentially clarifying ambiguous indicators of social status. Unlike group debates, which use standardized prompts and structured interactions and primarily assess team dynamics, personal interviews explicitly enable nuanced evaluations, making social status—particularly caste—more salient. Section 5.6 provides quantitative and qualitative evidence supporting this interpretation.

5.6 Interviews as the Main Mechanism for Revealing and Penalizing Social Status

Sections 5.3 and 5.4 raise a critical question: *When does caste become visible to employers?* This subsection quantitatively demonstrates that subjective, fit-focused personal interviews are the primary stage where social status (caste) is revealed and penalized.

5.6.1 Quantitative Evidence

Although many surnames are caste-indicative, most do not reliably signal caste in elite urban college settings, where students come from diverse regions with varied naming conventions (Jodkha, 2017; Parmar, 2020). Consequently, only about 15% of students in my sample have clearly identifiable caste surnames. For this subgroup, caste penalties appear earlier—at the application reading round. However, for the remaining 85% with non-identifiable surnames, the caste penalty emerges predominantly during personal interviews where job-competencies, including candidate background and cultural fit, often become explicit topics

of conversation. Thus, personal interviews, particularly informal "fit" assessments, serve as the key mechanism for uncovering and penalizing social status.

Caste earnings penalty emerges earlier when names reveal caste. Figure 2 replicates Figure 1, separately examining students with and without clearly identifiable caste surnames. I re-estimate Equation 1 separately for each subgroup. Students with identifiable surnames experience a significant penalty starting from application screening, while those without identifiable surnames see penalties appear only at personal interviews. Corresponding regression results, including Oster's bias-adjusted coefficients, are in Online Appendix Tables OA.7 and OA.8.

Figure 2: **The Caste Earnings Penalty by Subsample**

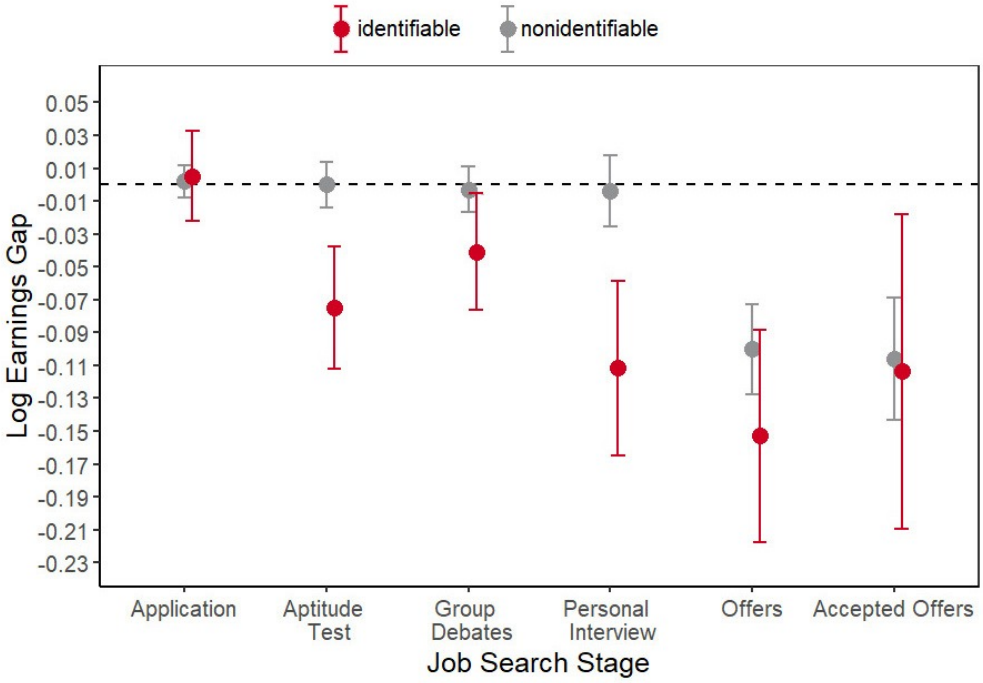

Notes: Figure 2 shows the log earnings gap across castes at successive stages of job search in the identifiable (red) versus the non-identifiable (gray) sample. Each coefficient in the figure is represented by a black dot and reports the percentage difference in the average salary at each job search stage between disadvantaged and advantaged castes. The vertical bars are 95% confidence intervals. These regressions include controls. There are no caste differences in pay for identical jobs and since offered salaries are non-negotiable during the job fair, the caste penalty in mean earnings at each stage can be calculated by comparing the salary differences in the jobs the caste groups are left in contention for after each stage.

Similar patterns in success rates. Comparable patterns emerge in success rates. Students with identifiable caste surnames face substantial penalties during application screening across nearly all jobs, whereas students without identifiable surnames experience penalties predominantly at the interview stage (Online Appendix Figure OA.2).

How are student names associated with caste? To classify surnames, I utilize official government sources, including regularly updated state-specific "schedules" of caste names (Scheduled Castes and Tribes, as described in Section 2). For advantaged castes, I reference surname lists compiled by the Anthropological Survey of India. For disadvantaged castes (SC/ST/OBC), I use Election Commission records of reserved constituencies and official state-maintained SC/ST/OBC lists.

5.6.2 Qualitative evidence supporting caste revelation at interviews

Qualitative evidence further clarifies why caste becomes most legible to recruiters during personal interviews. Group debates provide limited scope for caste inference: candidates debate standardized prompts in large groups with minimal individualized questioning, and neither physical appearance nor linguistic cues reliably indicate caste. Anthropological research consistently finds phenotypic traits like skin color are unreliable caste indicators among urban, educated populations (Sebring, 1969; Mishra, 2015; Parameswaran and Cardoza, 2015). Human Rights Watch concurs, stating that "*Lower-caste communities are almost invariably indistinguishable in physical appearance from higher-caste communities*" (Human Rights Watch, 2001, p. 2). Similarly, elite-college varieties of Indian English reflect regional or educational backgrounds rather than caste, as English-medium schooling has expanded significantly beyond upper castes due to affirmative action (Wiltshire, 2020; Sharma, 2017; Highet, 2023). Thus, group debates offer employers minimal caste-identifying information.

In contrast, personal interviews substantially shift the evaluative context. They amplify earlier signals such as schooling pedigree, resume details, and debate performances through deeper discussions. Interviews also elicit richer background information—personal interests, leisure activities, niche extracurricular involvement, and peer networks (e.g., Rivera, 2012, 2015), which are strongly correlated with caste-linked social and cultural capital, such as networks, resources, and cultural tastes (cf. Bourdieu, 1986). Such characteristics are well-documented caste correlates in elite Indian contexts (cf. Jodhka and Newman, 2007; Deshpande, 2011; Jodkha, 2017). Additionally, caste identities frequently become salient through social interactions in organizational settings, as managers learn subtle cues from conversations about personal background, residential origins, familial occupations, or shared acquaintances (Bapuji et al., 2023). Integrating these nuanced signals with earlier observations enables recruiters to more accurately infer caste, aligning precisely with the timing and intensity of caste penalties documented quantitatively (Section 5.6).

5.7 Heterogeneity in Caste Penalties, Selectivity, and Promotion Outcomes: Empirical Facts & Stylized Model

In Sections 5.3–5.6, we demonstrated that most of the caste penalty—in both earnings and success rates—emerges at the personal interview stage, where caste becomes salient. Because direct measures of candidate performance at each stage are unavailable, unobserved ability differences may remain hidden until later stages (Section 4.11).

This raises two key questions. First, might caste penalties appear earlier in *highly selective* jobs, which recruit primarily from the right tail of the ability distribution? Second, could personal interviews be revealing genuine unobserved ability differences—rather than discrimination—that justify hiring disparities across caste?

Variation in job-level selectivity offers an empirical test. If disadvantaged castes truly had lower pre-interview performance, such deficits should be particularly evident in highly selective jobs. Yet, pre-interview caste gaps remain negligible even among these selective positions.

Furthermore, once hired, disadvantaged castes exhibit higher promotion rates. One might argue this simply reflects inframarginal productivity differences—specifically, a greater share of right-tail performers among disadvantaged hires. However, examining promotion patterns across jobs of increasing selectivity helps disentangle discrimination from ability differences. If hiring thresholds were identical across castes, disadvantaged hires' promotion advantage should shrink as jobs become more selective. Instead, we find this advantage remains stable across all selectivity tiers.

Below, I present two empirical facts leveraging job selectivity variation, followed by a formal model demonstrating that these patterns reflect a systematically stricter hiring threshold for disadvantaged castes. Fact 1 uses job-specific selectivity measures at each screening round, revealing no caste gap in pre-interview success rates across selectivity levels. Fact 2 categorizes jobs by selectivity and documents consistently higher promotion rates for disadvantaged hires throughout the selectivity distribution.

Fact 1 (Conditional Success Rates Reveal No Early Penalty, but a Large Final Interview Gap). Figure 5 plots, for each job, the conditional success-rate gap (disadvantaged minus advantaged) against the cumulative cut rate at each screening stage (defined in Section 4.7). Negative values indicate lower advancement rates for disadvantaged candidates relative to advantaged candidates.

Formally, the gap for job j is the coefficient β_j from the regression:

$$\Pr(\text{Success}_{ijt}^{\text{Job Search Stage}}) = \alpha_j + \beta_j \times \text{Disadv. Caste}_i + \text{Controls}_i + \gamma_t + \epsilon_{ijt}, \tag{3}$$

where $i \in \mathcal{I}$ indexes students and γ_t denotes year fixed effects. Controls include GPA, entrance exam scores, prior internships, major, and degree.

Disadvantaged castes are about 30 percentage points less likely to advance than advantaged castes after personal interviews, consistently across nearly all jobs. Before personal interviews—during application reading, written tests, and group debates—the caste gaps remain negligible, *even* among highly selective jobs that eliminate most applicants. A distinct caste penalty emerges only after personal interviews, underscoring that late-stage subjective "fit" assessments, rather than pre-interview performance differences, drive observed caste disparities.

Robustness. As a robustness check, Online Appendix Figure OA.3 presents the *uncontrolled* success rates. It shows disadvantaged castes initially face a modest penalty during application screening, but this gap nearly disappears before interviews, only to reappear sharply afterward—exactly where the controlled regressions identify intensified caste disparities (Fact 1).

Note that the scatter around the average caste penalty generally increases across successive screening rounds (Figures 5 and OA.3). This reflects the increasing subjectivity of assessments at later stages—particularly during personal interviews—resulting in more heterogeneity in how caste penalties manifest across different job roles.

Fact 2 (Disadvantaged Hires Earn Comparably Higher Promotion Rates, Even in Highly Selective Jobs). The above findings do not fully exclude the possibility that interviews simply reveal genuine productivity-related differences between castes. To distinguish unobserved ability differences from discrimination, I turn to *post-hire* performance outcomes.

Using data from an employment-focused online platform, I track career trajectories for over 98% of the job fair participants, employing promotion rates as a measure of job performance.[12] If disadvantaged castes genuinely had lower unobserved abilities revealed during interviews, their relative performance should decline in more selective jobs, which impose higher—but equal—standards across castes. Instead, disadvantaged hires consistently maintain their promotion advantage even in highly selective positions. This finding aligns with outcome-based tests for discrimination used in healthcare and algorithmic fairness literature

[12]The unlinked sample comprises just 1.5% of graduates. Section 4.10 details the linking procedure.

(Chandra and Staiger, 2010; Abaluck et al., 2016; Chandra and Staiger, 2020), where higher success rates among marginalized groups, conditional on selection, indicate discriminatory screening thresholds.

I estimate caste differences in promotion rates using the following regression across job selectivity tiers:

$$\text{Outcome}_{ijt} = \alpha + \beta \times \text{Disadv. Caste}_i + \text{Controls}_i + k_j + \gamma_t + \epsilon_{ijt}, \tag{4}$$

where $i \in \mathcal{I}$ indexes students, $j \in \mathcal{J}$ indexes jobs, γ_t denotes year fixed effects, and k_j represents job fixed effects. Controls include GPA, entrance exam scores, prior labor market experience, major, and degree.

Table 7 presents these results:

- **Columns (1)–(2)** show that, pooling all jobs, disadvantaged hires have approximately a 20% *higher* probability of promotion relative to advantaged peers, despite initial hiring penalties.
- **Columns (3)–(8)** group jobs by *selectivity* (final rejection rates), revealing that disadvantaged hires maintain a stable promotion advantage of about 20-25% even as selectivity increases—from the top 75% to the top 15% most selective jobs (for example, jobs in the top 20% reject at least 80% of applicants). Thus, disadvantaged hires continue to outperform, even after stringent selection.
- **Columns (9)–(10)** demonstrate consistent patterns in both client-facing and non-client-facing roles, indicating that the promotion advantage is not specific to particular job functions.

Promotion-rate estimates are conservative. Although global promotion grids at top multinationals are nominally standardized (e.g., Bloom et al., 2012), local managers often subjectively evaluate candidates on dimensions such as "leadership," "communication," and "potential" (Benson et al., 2019). Such subjective judgments would, if anything, likely *reduce* the observed promotion probabilities. Additionally, evidence suggests explicit caste-conscious policies favoring disadvantaged castes at multinationals remain exceedingly rare (Bapuji et al., 2023), implying that these results are unlikely contaminated by preferential treatment. Hence, the observed promotion advantage (around 20%) represents a conservative estimate.[13] Online Appendix Table OA.9 corroborates this conservatism: advantaged-caste

[13] A 2018 survey of South Asians in the United States found that 67% of disadvantaged-caste staff reported workplace mistreatment (California v. Cisco Systems, Inc., 2022), underscoring the likely presence of negative biases.

job leavers (during the period for which tenure and promotion data were collected) typically have lower GPAs, whereas promotions correlate (albeit weakly) positively with GPA.

Key Takeaway. Facts 1 and 2 challenge the notion that personal interviews merely reveal latent skill deficits among disadvantaged castes. If disadvantaged candidates genuinely performed worse before interviews, we would observe significantly lower pre-interview success rates—especially in highly selective jobs. Similarly, if interviews accurately filtered out weaker candidates, promotion rates for disadvantaged hires should be comparable or lower. Instead, disadvantaged hires consistently achieve *higher* promotion rates, even in the most selective jobs. Crucially, this promotion advantage does not diminish with job selectivity, contradicting explanations based solely on unobserved ability differences.

Below, I introduce a simple theoretical framework that reconciles these two empirical facts by demonstrating how they are consistent with discrimination in hiring, with personal interviews acting as the primary mechanism translating social status into hiring barriers.

5.7.1 Explaining the Empirical Facts: A Simple Model of Late-Stage Discrimination

This section introduces a unified theoretical model demonstrating how (i) the absence of early-stage gaps for disadvantaged castes implies no genuine pre-interview ability deficit, and (ii) variation in hiring thresholds across jobs helps distinguish discrimination from genuine ability differences revealed at the interview stage.

Consider a set of jobs $j \in \mathcal{J}$. Each job has up to four sequential rounds $s \in \{1, 2, 3, 4\}$:

- Rounds $s = 1, 2, 3$ are pre-interview screening rounds (e.g. application, written tests, group debates)
- Round $s = 4$ is the personal interview

Groups and Screening. There are two castes (groups) $g \in \{0, 1\}$, where $g = 1$ denotes disadvantaged castes and $g = 0$ denotes advantaged castes. At each round s, employers observe a candidate-specific score $M_{s,i,g}$ for candidate i from group g.[14] A candidate passes round s if and only if their score satisfies $M_{s,i,g} \geq \tau_{s,j}$ and if they have successfully passed all preceding rounds.

[14] We assume the performance distributions for each group are continuous with smooth densities, finite moments, and satisfy standard regularity conditions. These technical assumptions ensure well-behaved pass rates and expectations when calculating promotion probabilities. While these properties underpin our formal analysis, they are not critical to understanding the main conceptual results.

For simplicity, we assume employers *perfectly observe* these scores at each stage. This assumption simplifies exposition but is not essential for our key insights: even under imperfect information, variation in screening selectivity and hiring thresholds across jobs allows us to distinguish between discrimination and genuine ability differences (see Online Appendix Section A.3 for details).

At the final interview stage ($s = 4$), the threshold $\tau_{4,j}^{(g)}$ may differ by group g, reflecting the fact that caste becomes salient for approximately 85% of candidates at this stage (Section 5.6.1). Thus, a final-stage caste gap is indicative of discrimination if $\tau_{4,j}^{(1)} > \tau_{4,j}^{(0)}$.

Monotonicity of thresholds. We assume that thresholds are monotonically increasing across screening rounds:

$$\tau_{4,j} \;\geq\; \tau_{3,j} \;\geq\; \tau_{2,j} \;\geq\; \tau_{1,j},$$

implying each successive round is at least as demanding as the previous one.

Identification from Exogenous Variation in Hiring Thresholds Across Jobs. For identification, we assume the underlying ability distributions are common across jobs. This assumption is essential because it ensures differences in hiring thresholds τ_j reflect genuine variations in job selectivity, rather than systematic differences in candidate quality.

To clarify this idea, imagine a large tub of ice cream representing a common talent pool. Each job selects candidates from this pool using a ladle of a specific size: a smaller ladle (higher cutoff) selects fewer candidates but with higher average quality; a larger ladle (lower cutoff) selects more candidates, including those with lower scores. Since all jobs draw from the same pool, ladle size captures selectivity alone.

This assumption is justified empirically by several institutional features:

1. Near-zero application costs and broad eligibility rules prevent sorting into specific jobs by caste (Section 4.8).

2. Admission quotas balance caste representation within majors, effectively ensuring common support across GPA quantiles (Sections 4.4 and 4.5; Online Appendix Figure OA.1).

3. Caste typically becomes visible only during interviews, minimizing the potential for selective screening in earlier stages.

4. Standardized recruitment protocols limit employers' ability to strategically manipulate cutoffs (Section 4).

Thus, *within a given job*, a higher cutoff naturally selects candidates of higher average quality. However, *across jobs*, variations in cutoffs reflect exogenous differences in selectivity, allowing us to use this variation as an instrument in our empirical analysis.

(A) Ruling Out Systematic Performance Differences Before Interviews

We first rule out systematic caste differences in candidate performance at pre-interview screening rounds.

Proposition 1 (No Early Gap Implies No Right-Tail Deficit Before Interviews (i.e., no systematic performance shortfall among highest-performing candidates))**.** Consider a job j with up to four rounds of screening, where the first three ($s = 1, 2, 3$) are pre-interview stages. If group 1 truly performed systematically worse, then for a sufficiently high threshold $\tau_{s,j}$ at some pre-interview stage $s < 4$, the pass-rate gap

$$\Delta_{s,j} \;=\; P(\text{Pass}_{s,i,j}|g = 1) \;-\; P(\text{Pass}_{s,i,j}|g = 0) \;<\; 0$$

would necessarily arise. Therefore, observing that $\Delta_{s,j} \geq 0$—even under high thresholds focusing specifically on the right tail—rules out any meaningful right-tail performance deficit for group 1 before the interview stage.

Sketch of Proof. We illustrate how variation in screening thresholds helps identify systematic performance differences at pre-interview rounds.

Figure 3 contrasts two scenarios. Each panel depicts two groups' performance distributions in pre-interview screening rounds, where blue (solid) represents advantaged candidates ($g = 0$) and red (dashed) represents disadvantaged candidates ($g = 1$). The shaded portion above each vertical dashed line represents candidates passing the screening cutoff for a specific job. The difference between the panels highlights how job selectivity can test for pre-interview performance deficits:

- **Panel (a): Hypothetical Right-Tail Deficit.** Suppose the screening score distribution for each group g, denoted by $\varphi_g(m)$, has a regularly decaying tail; specifically, assume $P(M_{s,i,g} \geq m) \sim C_g\, m^{-\alpha}$ for sufficiently large m, with constants $C_g > 0$ and $\alpha > 1$. This assumption implies that, as screening thresholds rise (reflecting more selective jobs), the pass rate shrinks at a rate determined by the tail density. If group 1's right tail is thinner than group 0 (i.e., a smaller C_1 or larger effective decay rate), the red "passing region" shrinks *more quickly* than the blue, resulting in increasingly negative pass-rate gaps as selectivity increases.

- **Panel (b): No Meaningful Deficit.** In contrast, Fact 1 shows no significant shortfall for disadvantaged candidates (Group 1) at any pre-interview screening stage. Even at high cutoffs for highly selective jobs, both groups maintain similar pass rates. This empirical result holds consistently across application reading, written tests, and group debates.

Thus, observing comparable pass rates across both groups even at the highest thresholds (as in Panel b) rules out systematic pre-interview performance deficits for the disadvantaged caste.

Figure 3: **Using Variation in Screening Thresholds to Test for Pre-Interview Performance Differences**

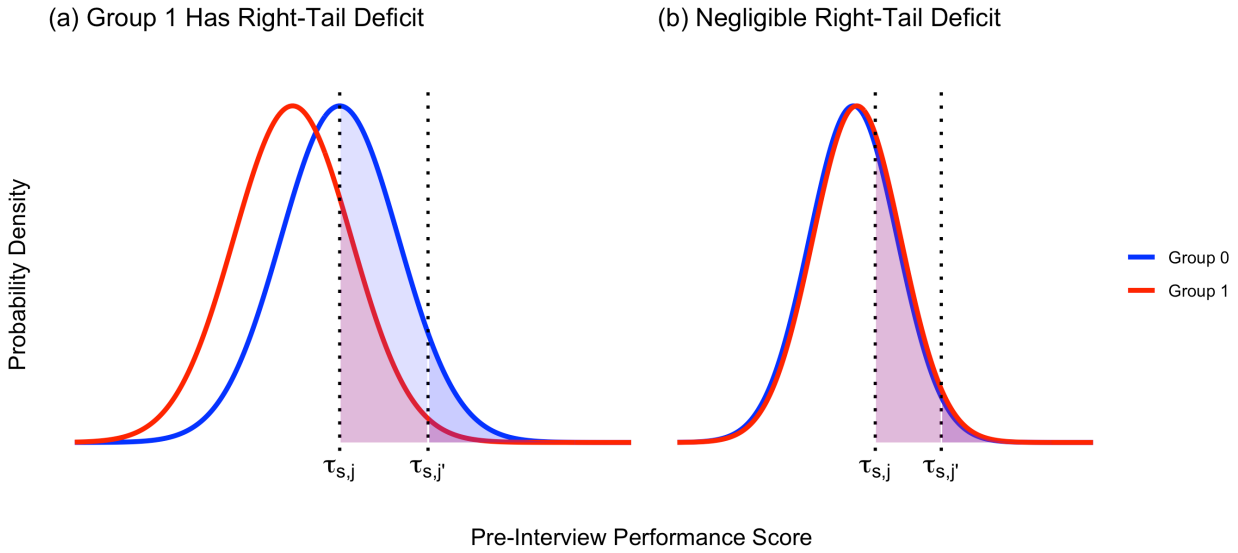

Notes: Figure 3 illustrates how selective screening would reveal systematic pre-interview performance differences. Panel (a) shows how a hypothetical right-tail deficit for Group 1 would become more apparent at higher cutoffs. Panel (b) shows the likely distribution of pre-interview performance scores implied by Fact 1, with no meaningful deficit. $\tau_{s,j}$ denotes screening threshold at stage $s < 4$ for job j (j' represents a more selective job).

□

Key Takeaway. Fact 1 confirms that early-stage caste penalties are negligible before interviews, strongly supporting Proposition 1. Thus, any observed late-stage gap at the interview stage ($s = 4$) must stem either from newly revealed ability signals during interviews or from discrimination based explicitly on social status.

(B) Using Variation in Hiring Thresholds to Test for Discrimination

Intuition. Variation in job-level hiring cutoffs provides a way to test whether the persistent 20% promotion advantage for disadvantaged hires arises from discriminatory hiring standards

or genuine ability differences. If both castes faced identical hiring thresholds, this promotion advantage would diminish as jobs become increasingly selective. The persistence of this advantage instead points toward discrimination arising from systematically stricter hiring thresholds applied to disadvantaged candidates.

Framework. To formalize this intuition, we link pre-hire screening performance to post-hire job performance. Under our maintained assumption of perfect information, we define job performance simply as the average screening score across all four rounds. Specifically, let $\bar{M}_{i,g} = \frac{1}{4}\sum_{s=1}^{4} M_{s,i,g}$ denote candidate i's average screening score from group g.[15] Candidate i is promoted if their average score exceeds a job-specific promotion threshold κ_j.

We distinguish two types of hiring thresholds for jobs ranked by selectivity. For jobs within the top $(1-q)$ fraction by selectivity (e.g., if $q = 0.8$, the top 20% most selective jobs):

- The *selectivity-tier threshold* $\tau^{q}_{4,j}$ captures the common hiring standard at the interview stage across these selectivity-tiered jobs.
- The *group-specific threshold* $\tau^{(g)}_{4,j}$ reflects potentially caste-specific hiring standards within each job.

  Under discrimination, we expect $\tau^{(1)}_{4,j} > \tau^{(0)}_{4,j}$. Without discrimination, thresholds coincide: $\tau^{(1)}_{4,j} = \tau^{(0)}_{4,j} = \tau^{q}_{4,j}$.

Let

$$\bar{Y}^{(g)}_{j} \;=\; \mathbb{E}\big[\bar{M}_{i,g} \,\big|\, M_{4,i,g} \geq \tau^{(g)}_{4,j}\big]$$

represent the average screening score of hired candidates from group g. Define the promotion gap within selectivity tier q as

$$\Delta^{\text{Prom}}(q) \;=\; \bar{Y}^{(1)}_{j} - \bar{Y}^{(0)}_{j}.$$

As hiring cutoffs rise in more selective jobs, the average score among hired candidates approaches the threshold score of marginal hires. If both castes face the same cutoff, their average scores will converge, implying $\Delta^{\text{Prom}}(q) \to 0$. However, if disadvantaged candidates face systematically higher thresholds— $\tau^{(1)}_{4,j} = \tau^{(0)}_{4,j} + \delta,\ \delta > 0$— the promotion gap stabilizes around δ even in the most selective jobs.

[15] Any monotonic link $f(.)$ between screening round performance and job performance works, but we choose to keep things simple.

Proposition 2 (Stable Promotion Gaps Across Selectivity Tiers Indicate Discrimination)**.** Consider a set of jobs J_q within the top $(1-q)$ fraction of selectivity, with associated thresholds $\tau_{4,j}^{q}$. Suppose that the average screening score among hired candidates for group g in these jobs is given by:

$$\bar{Y}_j^{(g)} = \mathbb{E}\left[\bar{M}_{i,g} \mid M_{4,i,g} \geq \tau_{4,j}^{(g)}\right].$$

If no discrimination exists ($\tau_{4,j}^{(1)} = \tau_{4,j}^{(0)}$), then as $q \uparrow$ (as jobs become more selective), average scores of hired candidates converge to the common cutoff, and thus $\Delta^{\text{Prom}}(q) \to 0$.

Conversely, if discrimination is present, with disadvantaged candidates facing a higher cutoff ($\tau_{4,j}^{(1)} = \tau_{4,j}^{(0)} + \delta$ for some $\delta > 0$), then average hired scores for disadvantaged candidates stabilize at approximately δ higher. Consequently, the promotion gap $\Delta^{\text{Prom}}(q)$ remains consistently positive across all levels of selectivity.

Sketch of Proof. Figure 4 illustrates the key intuition. In each scenario, average screening performance for the two groups is shown, with red (Group 1) consistently outperforming blue (Group 0) post-hire due to stricter selection.

- **Panel (a): No Discrimination.** Without discrimination, both groups face a common hiring cutoff: $\tau_{4,j}^{(1)} = \tau_{4,j}^{(0)} = \tau$. Under the regularity condition that each group's screening score distribution decays regularly—specifically, $P(M_{s,i,g} \geq m) \sim C_g\, m^{-\alpha}$ for large m, constants $C_g > 0$, and $\alpha > 1$, we have: $\lim_{\tau\to\infty}\left(\mathbb{E}[\bar{M}_{i,g} \mid M_{4,i,g} \geq \tau] - \tau\right) = 0$. Consequently, as the common cutoff rises in more selective jobs, each group's average hired score converges to the marginal threshold, and the performance gap between groups shrinks to zero.

- **Panel (b): Discrimination.** With discrimination, disadvantaged candidates (Group 1) face systematically higher hiring thresholds than advantaged candidates (Group 0), such that $\tau_{4,j}^{(1)} = \tau_{4,j}^{(0)} + \delta$, $\delta > 0$. Under the same regularity condition stated above, as thresholds become increasingly selective:

$$\lim_{\tau\to\infty}\left(\mathbb{E}[\bar{M}_{i,1} \mid M_{4,i,1} \geq \tau_{4,j}^{(1)}] - \mathbb{E}[\bar{M}_{i,0} \mid M_{4,i,0} \geq \tau_{4,j}^{(0)}]\right) = \delta.$$

  Thus, the average hired score for disadvantaged candidates stabilizes at approximately δ points higher than for advantaged candidates, preserving a consistent performance gap across increasingly selective jobs.

Since a candidate is promoted if $\bar{M}_{i,g}$ exceeds a job-specific threshold κ_j, the asymptotic behavior of the screening score distributions implies that if both groups face the same thresh-

old, the promotion gap $\Delta^{\text{Prom}}(q)$ shrinks toward zero, whereas a constant gap in thresholds (indicative of discrimination) preserves a stable promotion gap.

Figure 4: **Using Variation in Hiring Thresholds To Test for Discrimination**

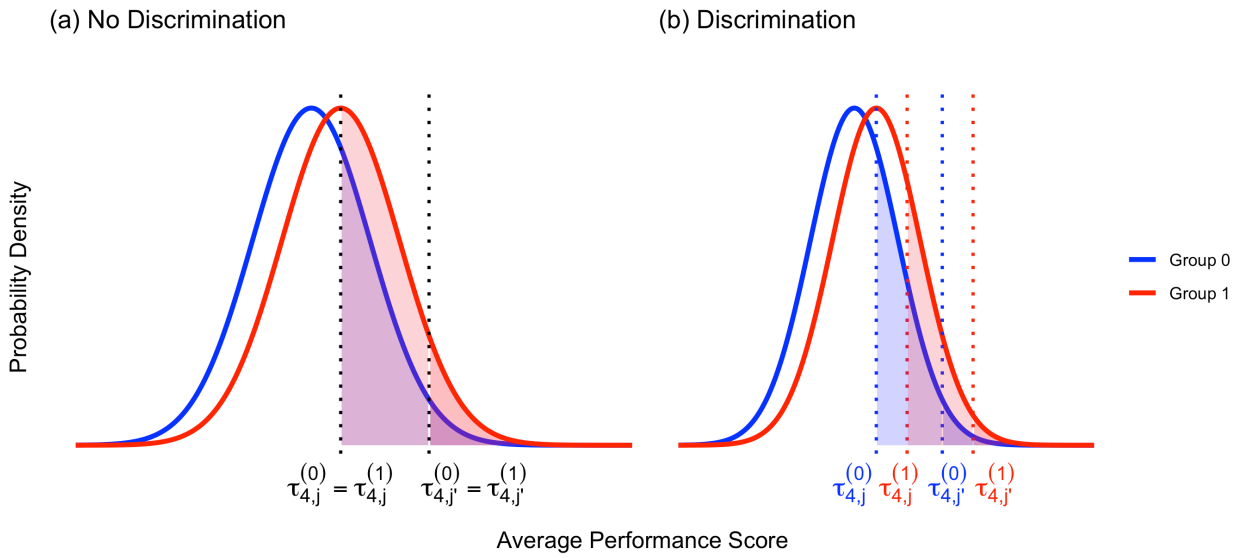

Notes: Figure 4 depicts two scenarios. In panel (a), a common cutoff for both groups shrinks their hired-performance gap at higher cutoffs. In panel (b), a constant δ gap in cutoffs preserves a stable performance advantage for Group 1. Here, $\tau^{(g)}_{4,j}$ denotes group g's interview threshold at job j, with j' representing a more selective job.

□

Not every hire sits on the margin. The approach remains robust even when most hires comfortably exceed hiring thresholds. What matters for identification is how *thresholds* shift with selectivity, enabling clear discrimination tests irrespective of inframarginal hires.

Key Takeaway. Fact 2 strongly supports Proposition 2. The persistent promotion advantage for disadvantaged hires across all selectivity tiers indicates that disadvantaged candidates face consistently higher hiring thresholds, a pattern incompatible with explanations based solely on genuine ability differences.

Summary. Together, the negligible pre-interview caste gaps (Fact 1) and persistent post-interview promotion advantages (Fact 2) provide robust empirical evidence of discrimination at the hiring stage, particularly during personal interviews. These results are consistent with the model predictions of discriminatory hiring standards (Propositions 1 and 2), ruling out unobserved ability differences as the primary explanation.

Figure 5: **Caste Penalty in the Conditional Success Rate At Each Screening Round**

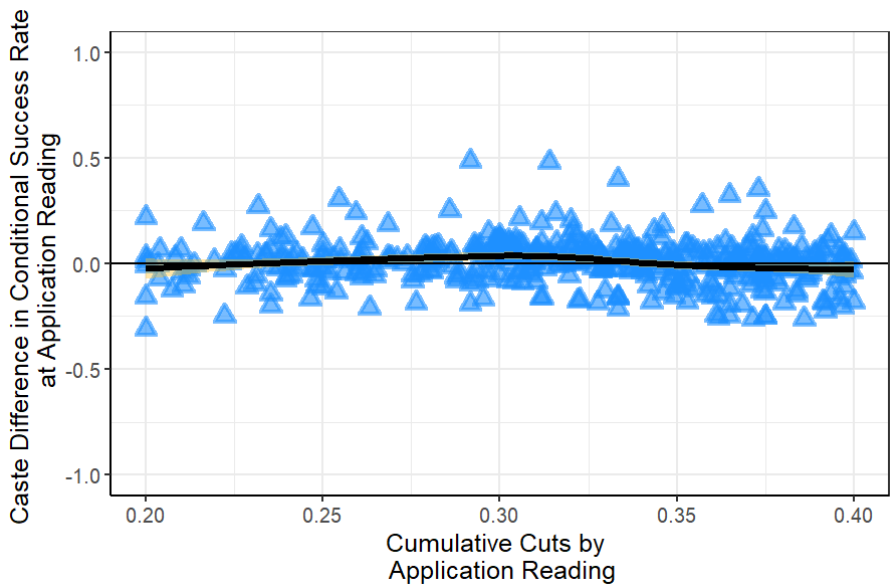

(a) Caste Penalty in Conditional Success Rates versus Cumulative Cuts by Application Reading

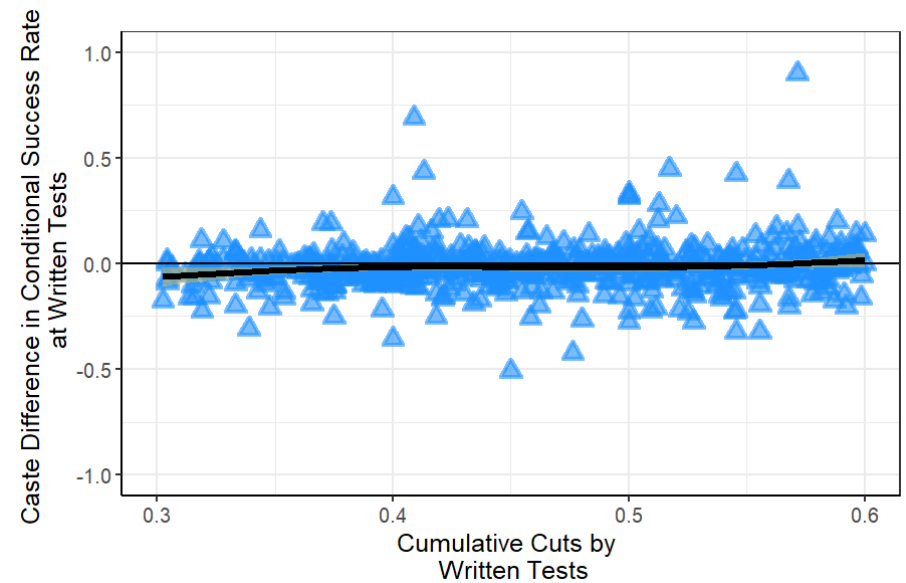

(b) Caste Penalty in Conditional Success Rates versus Cumulative Cuts by Written Tests

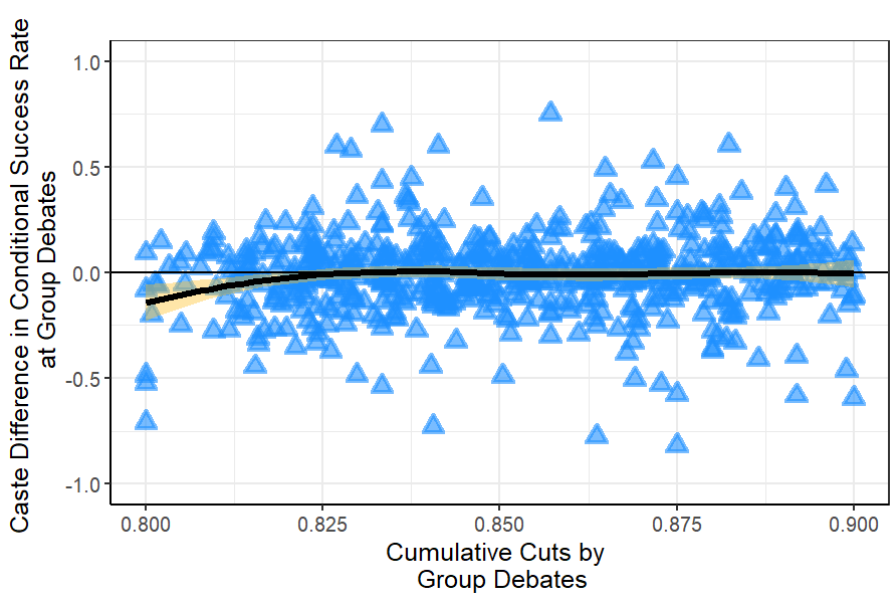

(c) Caste Penalty in Conditional Success Rates versus Cumulative Cuts by Group Debates

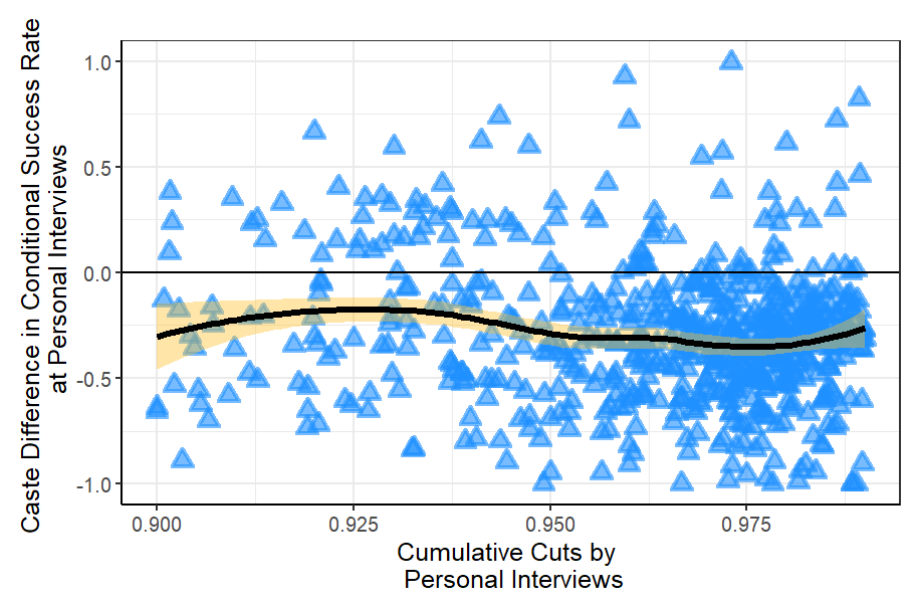

(d) Caste Penalty in Conditional Success Rates versus Cumulative Cuts by Personal Interviews

Notes: Figure 5 shows the caste difference in the conditional (i.e., with controls) success rate against the cumulative cuts made by jobs by the end of each screening round. Each blue triangle denotes a job. Smoothed conditional means (thick black lines) are calculated using the "loess" function in R. The shaded areas in yellow denote 95% confidence intervals. The thin solid horizontal line denotes a y-intercept of 0. The x-axis shows the cumulative proportion of students cut by jobs by the end of each screening round. The y-values are calculated by estimating, *for each job*, a standard linear probability model with controls and cohort-year fixed effects and then plotting the coefficient on disadvantaged caste from these regressions on the y-axis. Each regression includes controls for student characteristics such as degree, major, college GPA, entrance exam score as well as cohort-year fixed effects.

Table 7: **Caste Differences in Promotion Rate**

	Dependent variable:									
	Promoted	Promoted	Promoted at Top 75%	Promoted at Top 50%	Promoted at Top 40%	Promoted at Top 30%	Promoted at Top 20%	Promoted at Top 15%	Promoted at Client Facing	Promoted at Non-Client Facing
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)	(9)	(10)
Disadv. Caste (OLS)	0.137***	0.096***	0.104***	0.112***	0.136***	0.147***	0.190**	0.204**	0.107**	0.082**
	(0.025)	(0.026)	(0.029)	(0.038)	(0.043)	(0.054)	(0.073)	(0.084)	(0.044)	(0.032)
Disadv. Caste (Oster)	0.142***	0.093***	0.108***	0.122***	0.141***	0.157**	0.196**	0.206**	0.106***	0.075**
	(0.026)	(0.028)	(0.031)	(0.037)	(0.041)	(0.058)	(0.082)	(0.082)	(0.039)	(0.036)
Tenure Duration		0.084***	0.082***	0.076***	0.074***	0.068***	0.054*	0.079**	0.065***	0.097***
		(0.012)	(0.014)	(0.019)	(0.021)	(0.025)	(0.030)	(0.037)	(0.021)	(0.015)
Mean Y	0.684	0.684	0.667	0.668	0.676	0.711	0.711	0.706	0.687	0.682
Observations	2,489	2,489	2,012	1,246	997	748	502	374	835	1,654
Adjusted R^2	0.183	0.205	0.203	0.216	0.211	0.182	0.192	0.189	0.194	0.214
Job FE	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
Year FE	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
Other Controls		✓	✓	✓	✓	✓	✓	✓	✓	✓

Notes: Table 7 shows caste differences in job performance measures among students who graduated with full-time jobs through the college's on-campus job fair. All of these students had to succeed in the final personal interview round to obtain jobs. Column (1) reports caste differences in promotion rates only with job and year fixed effects. Column (2) reports the same regression as Column (1) but with all other controls such as GPA, major, degree, and job tenure. Columns (3)-(8) show caste differences in promotion rates with controls among jobs at different quantiles of "job selectivity," proxied by the cumulative proportion of applicants cut by a given job by the end of the final screening round. Columns (9)-(10) show caste differences in promotion rates with controls among client-facing and non-client-facing jobs, respectively. Bias-adjusted coefficients are calculated using Oster (2019). Following Oster's recommended benchmarks, I set $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$, implying that unobservables are assumed to have the same degree of selection as observables, and the total explanatory power can increase by up to 30% beyond the observed R^2. $^*p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

5.8 Classification and Preferred Interpretation of the Caste Penalty: Distinguishing from Statistical or Customer Discrimination

Given the evidence that employers directly penalize caste, the most plausible interpretation is *disparate treatment*—explicit discrimination based on caste. Two facts strongly support this view. First, disadvantaged hires are consistently *more* likely to be promoted across all selectivity tiers (Table 7). Second, over three-quarters of the roles analyzed are non-customer-facing (Table 3). Taken together, these findings are difficult to reconcile with explanations based purely on accurate statistical inference about productivity[16] or customer-driven discrimination; instead, they align closely with *caste-based bias*.

Explicit caste discrimination is prohibited under Article 15 of the Indian Constitution (Ministry of Law and Justice, Government of India, 2021). Although subjective criteria for assessing "fit" might inadvertently create a *disparate impact* (e.g., by favoring candidates with specific extracurricular backgrounds), the comprehensive evidence presented throughout Sections 5.2–5.7 provides the strongest support for the conclusion that observed hiring disparities primarily reflect disparate treatment.

5.9 What Happens During Personal Interviews?

We do not directly observe interview transcripts, so we cannot parse the exact conversations. However, prior ethnographic research on elite hiring demonstrates that final-round interviews at multinational corporations place substantial weight on subjective markers of "cultural fit," which are closely tied to candidates' socioeconomic backgrounds rather than strictly job-relevant skills (e.g., Rivera, 2012, 2015; Lamont, 2009). Similar patterns have been documented in India's civil service examinations, where subjective "personality" interviews disproportionately penalize lower-status candidates despite having limited predictive power for job performance (Bhavnani and Lee, 2021). These external findings strongly support our interpretation that the significant late-stage caste penalty we document arises from discretionary, status-based evaluations rather than from objective assessments of ability.

6 Alternative Explanations

Below, I systematically assess whether alternative explanations can plausibly account for the observed caste gaps that emerge exclusively at the personal interview stage. None fully matches the empirical evidence.

[16] An accurate statistical discrimination model typically assumes evaluators hold correct beliefs about average group-based productivity and that the underlying productivity distributions are relatively well-behaved (cf. Aigner and Cain, 1977).

1. **Early "fairness" to appear inclusive.** If firms intentionally lowered early-stage cut-offs for disadvantaged candidates to appear inclusive, those candidates should experience *significantly higher* conditional success rates in earlier rounds—especially among highly selective jobs. However, Figure 5 reveals parity, not advantage, in pre-interview advancement rates. Additionally, students with caste-identifiable surnames face an early-stage penalty right from application screening (Figure 2), contradicting the notion of early-round inclusivity.

2. **Delayed discrimination saved for interviews.** Employers face no legal or reputational pressure to postpone caste bias; Indian MNCs often view caste-targeted hiring as "anti-merit" (Jodhka and Newman, 2007). Indeed, when caste is obvious through identifiable surnames, penalties emerge at the *application* stage rather than later. Moreover, if employers wished to apply subjective criteria to disadvantage lower-caste candidates earlier, they could easily use the subjective group-debate rounds, which explicitly evaluate soft skills like teamwork or communication. The absence of caste gaps in these group debates indicates that employers do not systematically discriminate early unless caste information is explicitly revealed. Interviews therefore matter primarily because they newly reveal caste for the remaining 85% of candidates.

3. **Hidden right-tail skill gaps.** If disadvantaged castes genuinely had an ability deficit concentrated in the right tail, we would observe systematically lower pre-interview pass rates for disadvantaged candidates in highly selective jobs that explicitly target top performers. Fact 1 clearly demonstrates no such pattern.

4. **Attrition between rounds is negligible.** Among those participating in the placement process, nearly all candidate exits occur through formal elimination decisions made explicitly at each of the four strictly enforced screening rounds (see Section 4.3 and Table 4). Furthermore, the placement office meticulously records every advance-or-exit decision, ensuring that undocumented attrition between rounds is negligible, and exit surveys confirm this comprehensive record-keeping.

5. **Better outside options for advantaged castes.** Candidates accepting return offers (pre-placement job offers from internships) exit the placement process *before* the initial screening rounds, comprising fewer than 5% of each cohort (Section 3.2). Moreover, institutional rules explicitly discourage parallel external job searches by withdrawing placement-office support from candidates who independently pursue job opportunities. Consequently, advantaged-caste students do not disproportionately drop out of the placement process midstream due to superior outside options.

6. **Socio-emotional skill differences.** Group debates—explicitly designed to evaluate communication, teamwork, and other socio-emotional skills—reveal no caste-based differences (Panels (c) of Figures 5–OA.3).

7. **Missing data or skipped rounds.** The placement office enforces mandatory attendance at every stage and strictly prohibits employers from skipping steps or altering procedures mid-process (Section 4.3). Comprehensive exit surveys independently corroborate the completeness and accuracy of the official records (Section 5.1), eliminating concerns about missing data.

8. **Negotiation advantages for advantaged castes.** Salaries (and benefits) are set ex-ante, strictly non-negotiable, and verified by the placement office (Section 3.2). Exit surveys further confirm identical starting salaries (Section 5.1).

9. **Competition from government jobs.** Quota-regulated public-sector positions comprise less than 2% of our sample (Table 3), providing negligible competition for private-sector employers. Additionally, firms recruiting on-campus have full visibility of each other's presence, as all participating companies and their job listings are accessible to candidates and recruiters through a centralized placement portal.

10. **Caste differences in job preferences.** Figure 1 shows that preference-based sorting can at best explain only a small share of the initial earnings gap and none of the sharp interview-stage penalty.

11. **Preferential promotion of disadvantaged hires.** Corporate practices explicitly tracking and actively targeting caste diversity remain exceedingly rare. Under 3% of India's top 100 publicly listed firms systematically track caste representation (BusinessLine, 2018), and fewer than 10% of large U.S. firms maintain socioeconomic diversity metrics (Stansbury and Rodriguez, 2024). Formal or informal caste-based preferential treatment in promotion decisions is uncommon and unlikely. Recent evidence also highlights the continued rarity of explicit caste-focused initiatives at multinational corporations in India (Bapuji et al., 2023). Indeed, extensive survey evidence points to widespread bias *against* disadvantaged caste employees rather than in their favor (California v. Cisco Systems, Inc., 2022). Therefore, the consistently observed 20% promotion advantage for disadvantaged hires reflects systematically stricter hiring thresholds for these candidates rather than reverse favoritism post-hire.

No alternative reconciles the absence of pre-interview gaps with the large, interview-stage penalties and the persistent promotion advantage of disadvantaged hires. The evidence is

most consistent with caste-based disparate treatment during personal interviews.

7 External Validity

How generalizable are findings from one elite Indian institution? While caste discrimination might initially appear uniquely Indian, four considerations underscore why the results have broader implications, not just across India but also for global elite hiring and other high-stakes selection contexts.

1. Caste Provides a Clear Marker to Study Status-Based Discrimination. In most Western hiring contexts, social status or class discrimination must be inferred indirectly through ambiguous indicators such as educational background, accent, or informal networks (e.g., Rivera, 2015; Laurison and Friedman, 2019; Lamont, 1992). In contrast, India's caste system explicitly bundles economic, social, and cultural capital into a single, formally recognized and recorded identity marker (Beteille, 1965; Bourdieu, 1986). This explicitness offers unique measurement clarity, allowing precise identification of the exact stage at which status discrimination emerges.

Moreover, the underlying mechanism identified—discretionary personal interviews assessing subjective "fit"—is by no means uniquely Indian. Extensive evidence from the U.S. and the U.K. confirms that subjective assessments, particularly at late stages of hiring or admissions, disproportionately penalize candidates from lower socioeconomic backgrounds by implicitly assessing their "cultural capital," such as mannerisms, extracurricular backgrounds, conversational style, or perceived alignment with organizational culture (e.g., Rivera, 2012; Laurison and Friedman, 2019). Therefore, while caste is uniquely explicit, the mechanism of status-based discrimination via informal and discretionary evaluations of cultural fit is broadly generalizable.

2. Institutional Representativeness Within India and Globally. The recruitment process studied here is typical of elite Indian higher education. Nearly all top-30 Indian colleges operate similar multi-round placement processes, hosting the same set of multinational corporations (MNCs) recruiting nationwide. These firms employ standardized hiring procedures, unified national salary grids, and centralized hiring teams that visit multiple campuses (Sections 3.3–4.3). Thus, the institutional setting studied accurately mirrors broader recruitment patterns across India's elite higher education and corporate sectors, directly generalizing findings to similar settings throughout India.

Moreover, the multinational corporations examined here replicate identical hiring structures globally, using similarly structured evaluations across their international offices (Bloom et al., 2012). Elite college job fairs in the U.S. and Europe employ comparable procedures—starting from standardized screening and culminating in subjective personal interviews to gauge candidates' "fit" and interpersonal compatibility (e.g., Rivera, 2015; Laurison and Friedman, 2016, 2019). Hence, institutional representativeness extends internationally, particularly within global firms operating in culturally diverse but organizationally standardized settings.

3. Replication of the Underlying Mechanism Across Diverse Contexts. The specific mechanism identified—status discrimination emerging predominantly during subjective final-round interviews—resonates far beyond India's caste context. As explicit status markers become less obvious or deliberately concealed in diverse urban labor markets worldwide (e.g., Jodkha, 2017; Clark, 2015), subjective interpersonal assessments become increasingly critical for identifying subtle status differences. Research across contexts consistently finds these evaluations amplify inequality by favoring candidates with class-aligned cultural capital, reinforcing socioeconomic disparities in access to elite opportunities (Rivera, 2012; Stansbury and Rodriguez, 2024; Laurison and Friedman, 2019).

In U.S. elite college admissions, the shift from standardized test scores to "holistic" evaluations—personal statements, recommendations, and interviews—has been shown to disadvantage high-achieving, lower socioeconomic-status applicants precisely because subjective criteria emphasize subtle class-linked advantages rather than objective academic potential (Kuncel et al., 2013). Similar findings emerge in elite public-sector hiring in India, where subjective "personality" interviews disproportionately disadvantage lower-status candidates without predictive power for subsequent job performance (Bhavnani and Lee, 2021). These cross-context parallels strongly support the broader relevance of the mechanism documented here.

Further reinforcing the broader applicability, dual-process cognitive theories explain why subjective evaluations of "fit" consistently disadvantage candidates lacking the subtle cultural markers of higher social status (Kahneman, 2011; Todorov, 2017). Quick, intuitive evaluations in interviews frequently rely on superficial impressions of rapport, conversational style, and familiarity. These impressions implicitly reflect socioeconomic background and reinforce evaluators' biases toward candidates from similar class backgrounds (e.g., Pinker, 2021; Rivera, 2012).

Such biases are widespread and systematically documented across diverse high-stakes contexts, from hiring and admissions to performance evaluations and promotions in elite firms

(e.g., Gorman, 2005; Sarsons, 2017; Sarsons et al., 2021; Benson et al., 2023). Consequently, although the specific data come from one elite Indian institution, the identified mechanism of subtle, subjective status-based bias during discretionary evaluations applies broadly across elite hiring and selection processes worldwide.

4. Broader Implications for Policy and Organizational Design. Building upon this identified mechanism, this study also highlights a crucial and universal policy implication: affirmative action efforts aimed solely at initial entry points (such as elite college admissions) can be substantially undermined at subsequent selection stages, particularly when hiring practices rely heavily on discretionary, subjective evaluations (Yenmez and Sönmez, 2022). The Indian context exemplifies a widespread international challenge: even if disadvantaged candidates gain access to elite educational institutions, subjective hiring practices by multinational corporations can erode these initial equity gains, perpetuating class disparities in elite occupational attainment.

Given multinational firms' globally standardized hiring practices, these findings underscore the urgent need for systemic reform in structuring and monitoring subjective evaluations, potentially complemented by more objective and transparent assessment practices. The issue identified—subjective "fit"-based evaluation creating class barriers—warrants broader attention wherever socioeconomic diversity in elite careers remains persistently low despite interventions at earlier educational stages (cf. Zimmerman, 2019).

Thus, although the immediate empirical setting focuses on caste in India, the core message holds universal relevance: whenever subjective interpersonal evaluations of "fit" constitute the final gateway into elite roles, hidden markers of social status re-enter, perpetuating inequality and undermining broader efforts to ensure equitable access. This insight applies globally, underscoring the critical need for reforms in elite hiring and selection practices worldwide.

8 Conclusion

By opening the hiring "black box," this paper provides three novel contributions. First, it delivers the first large-scale, stage-by-stage empirical evidence on exactly when class-based bias emerges in elite private-sector hiring, tracing discrimination from initial applications to actual job attainment in a real-world setting. Second, it precisely quantifies the magnitude of discrimination, showing that nearly 90% of the hiring disparites arises exclusively during the discretionary personal-interview stage. Third, it demonstrates that the candidates disadvantaged by these subjective interviews—those from lower-status backgrounds—

subsequently outperform their advantaged peers, securing nearly a 20% higher rate of promotion. Together, these findings provide robust evidence that subjective evaluations of candidate "fit" systematically undervalue talent from disadvantaged groups, likely imposing substantial costs both on marginalized candidates and firms themselves.

The Indian setting provides an unusually clear vantage point for investigating status-based discrimination. While caste offers a uniquely explicit marker of social status, the underlying mechanism identified—subjective and discretionary evaluations of "cultural fit"—is universal. Indeed, as overt class signals—accents, surnames, openly displayed affiliations—fade in global labor markets, hiring managers increasingly rely on subtler, richer dialogue about extracurricular interests, cultural alignment, conversational polish, and life experiences (e.g., Jodkha, 2017; Laurison and Friedman, 2016). These subjective assessments activate implicit biases, as fast, intuitive judgments about rapport or similarity often crowd out deliberate, merit-based evaluations (Kahneman, 2011). Consequently, the same late-stage penalties that disadvantage lower-status candidates in India—such as caste visibility during interviews—recur globally, manifesting as biases against working-class accents in the UK or insufficiently prestigious extracurricular résumés in the US (e.g., Rivera, 2015; Laurison and Friedman, 2019; Stansbury and Rodriguez, 2024).

Crucially, these findings highlight a critical yet overlooked vulnerability in the affirmative action pipeline. Policies that successfully widen university gates for historically disadvantaged groups are systematically undermined by downstream hiring practices, particularly subjective personal interviews. By emphasizing informal cultural alignment and immediate rapport, employers inadvertently re-privilege candidates from familiar socioeconomic backgrounds. The consequences extend beyond equity: firms themselves incur substantial organizational costs. Hiring managers—who are rarely future supervisors—prioritize short-term cultural alignment over candidates' actual long-term potential, thereby discarding high-performing, diverse talent that could substantially boost organizational performance (e.g., Benson and Shaw, 2024; Hoffman and Stanton, 2024).

What can firms do? Eliminating interviews entirely is neither feasible nor advisable; however, reforming interview practices to mitigate bias is both realistic and beneficial. Simple interventions such as structured behavioral questions, blinded scoring rubrics, and systematic interviewer accountability—linking interviewers' assessments to subsequent candidate performance—can significantly reduce subjective biases while preserving the critical human judgment needed to assess intangible qualities. Despite the proven effectiveness of these approaches, fewer than 10% of large U.S. corporations and less than 3% of India's top-100 listed firms systematically track socioeconomic or caste data, forfeiting valuable opportunities to identify and correct these costly, bias-driven hiring failures (Stansbury and Rodriguez, 2024;

BusinessLine, 2018).

The implications extend well beyond traditional hiring. As remote recruitment and algorithmic screening grow prominent, subjective biases may diminish or migrate into digital platforms (Albright, 2024). Without careful auditing, these technologies risk replicating or amplifying existing biases. Extending detailed stage-by-stage audits—like those presented here—across diverse contexts, countries, and evolving hiring methods is thus essential to designing recruitment systems that enhance socioeconomic equity and organizational performance worldwide.

Ultimately, this paper's message transcends its specific empirical context. Wherever subjective assessments form the final hurdle to elite opportunities, subtle markers of social status inevitably resurface, perpetuating inequality. Addressing this persistent barrier requires sustained attention, targeted policy interventions, and institutional reforms. Doing so will foster fairer opportunities and unlock overlooked talent, benefiting organizations and societies globally.

References

[1] Abaluck, J., L. Agha, C. Kabrhel, A. Raja, and A. Venkatesh (2016). The determinants of productivity in medical testing: Intensity and allocation. *American Economic Review 106*(12), 3730–3764.

[2] Ægisdóttir, S., M. J. White, P. M. Spengler, A. S. Maugherman, L. A. Anderson, R. S. Cook, C. N. Nichols, G. K. Lampropoulos, B. S. Walker, G. Cohen, and J. D. Rush (2006). The meta-analysis of clinical judgment project: Fifty-six years of accumulated research on clinical versus statistical prediction. *The Counseling Psychologist 34*(3), 341–382.

[3] Aigner, D. J. and G. G. Cain (1977). Statistical theories of discrimination in labor markets. *Industrial and Labor Relations Review 30*(2), 175–187.

[4] Albright, A. (2024). The hidden effects of algorithmic recommendations. Working paper, Federal Reserve Bank of Minneapolis.

[5] Altonji, J. G. and R. M. Blank (1999). *Race and Gender in the Labor Market. Handbook of Labor Economics*, ed. by O. C. Ashenfelter and D. Card, Elsevier, vol. 3C, chap. 48.

[6] Alvero, A., S. Giebel, B. Gebre-Medhin, A. L. Antonio, M. L. Stevens, and B. W. Domingue (2021). Essay content and style are strongly related to household income and sat scores: Evidence from 60,000 undergraduate applications. *Science Advances 7*(42), eabi9031.

[7] Asher, S., P. Novosad, and C. Rafkin (2024). Intergenerational mobility in india: New measures and estimates across time and social groups. *AEJ: Applied Economics 16*(2), 66–98.

[8] Bapuji, H., K. Gupta, S. Chrispal, and T. Roulet (2023). What managers everywhere must know about caste. *MIT Sloan Management Review Winter*. Accessed: 2024-04-22.

[9] Barg, K. (2013). The influence of students' social background and parental involvement on teachers' school track choices: Reasons and consequences. *European Sociological Review 29*(3), 565–579.

[10] Barrios-Fernández, A., C. Neilson, and S. Zimmerman (2024). Elite universities and the intergenerational transmission of human and social capital. *American Economic Review*. conditionally accepted.

[11] Bayly, S. (2008). *Caste, Society and Politics in India from the Eighteenth Century to the Modern Age*. Chicago: Cambridge University Press.

[12] Benson, A., S. Board, and M. Meyer-ter-Vehn (2024). Discrimination in hiring: Evidence from retail. *Review of Economic Studies 91*(4), 1956–1987.

[13] Benson, A., D. Li, and K. Shue (2019). Promotions and the peter principle. *The Quarterly Journal of Economics 134*(4), 2085–2134.

[14] Benson, A., D. Li, and K. Shue (2023). "Potential" and the gender promotion gap. R&R, *American Economic Review*.

[15] Benson, A. and K. Shaw (2024). What do managers do? *Annual Review of Economics*. Forthcoming.

[16] Bergman, P. and A. Lerner (2025). Are recruiters better than a coin flip at judging resumes? here's the data. https://interviewing.io/blog/are-recruiters-better-than-a-coin-flip-at-judging-resumes. Accessed: 2025-04-21.

[17] Bergman, P., D. Li, and L. R. Raymond (2024). Hiring as exploration. *Review of Economic Studies (forthcoming)*. NBER Working Paper 27736, revised June 2024.

[18] Beteille, A. (1965). *Caste, Class and Power: Changing Patterns of Stratification in a Tanjore Village*. University of California Press.

[19] Beteille, A. (1969). *Castes: Old and New, Essays in Social Structure and Social Stratification*. Asia Publishing House.

[20] Bhavnani, R. R. and A. Lee (2021). Does affirmative action worsen bureaucratic performance? evidence from the indian administrative service. *American Journal of Political Science 65*(1), 5–20.

[21] Bisin, A., J. Rubin, A. Seror, and T. Verdier (2023). Culture, institutions & the long divergence. *Journal of Economic Growth 28*, 1–40.

[22] Bisin, A. and T. Verdier (1998). On the cultural transmission of preferences for social status. *Journal of Public Economics 70*, 75–97.

[23] Bisin, A. and T. Verdier (2010). The economics of cultural transmission and socialization. In *Handbook of Social Economics*, pp. 339–416. Elsevier.

[24] Bloom, N., R. Sadun, and J. Van Reenan (2012). The organization of firms across countries. *The Quarterly Journal of Economics 127*(4), 1663–1705.

[25] Bol, T., J. Witschge, H. G. Van de Werfhorst, and J. Dronkers (2014). Curricular tracking and central examinations: Counterbalancing the impact of social background on student achievement in 36 countries. *Social Forces 92*(4), 1545–1572.

[26] Borneman, M. J., S. R. Cooper, D. M. Klieger, N. R. Kuncel, and N. Brenckman (2007). Admission interviews: Not so nice to meet you. In *Alternative Predictors of Academic Performance, Symposium conducted at the Annual Meeting of the Society for Industrial and Organizational Psychology, New York, NY.*

[27] Bourdieu, P. (1986). The forms of capital. In J. G. Richardson (Ed.), *Handbook of Theory and Research for the Sociology of Education*, pp. 241–258. Greenwood Press.

[28] BusinessLine (2018). An Inclusive India Inc. Business Line, January 13, 2018.

[29] California v. Cisco Systems, Inc. (2022). Dept. of Fair Employment and Housing v. Cisco Systems, Inc. Docket No. H048962.

[30] Cappelli, P. (2019). Your approach to hiring is all wrong. *Harvard Business Review.* Accessed: 2024-08-15.

[31] Carroll, J. S., R. L. Wiener, D. Coates, J. Galegher, and J. J. Alibrio (1982). Evaluation, diagnosis, and prediction in parole decision making. *Law & Society Review 17*, 199–228.

[32] Chakravartty, P. and A. Subramanian (2021). Why is Caste Inequality Still Legal in America? *The New York Times*, May 25, 2021.

[33] Chandra, A. and D. O. Staiger (2010). Identifying provider prejudice in healthcare. Technical Report 16382, National Bureau of Economic Research.

[34] Chandra, A. and D. O. Staiger (2020). Identifying sources of inefficiency in health care. *Quarterly Journal of Economics 135*(2), 785–843.

[35] Checchi, D. and H. G. van de Werfhorst (2018). Policies, skills and earnings: How educational inequality affects earnings inequality. *Socio-Economic Review 16*(1), 137–160.

[36] Chetty, R., D. J. Deming, and J. N. Friedman (2023). Diversifying society's leaders? the determinants and causal effects of admission to highly selective private colleges. NBER Working Paper 31492.

[37] Clark, G. (2015). *The Son Also Rises: Surnames and the History of Social Mobility.* The Princeton Economic History of the Western World.

[38] Cooper, S. R., D. M. Klieger, M. J. Borneman, N. R. Kuncel, and H. Mikeworth (2007). Personal statements: Please keep them to yourself. In *Alternative Predictors of Academic Performance, Symposium conducted at the Annual Meeting of the Society for Industrial and Organizational Psychology, New York, NY.*

[39] Dana, J., R. Dawes, and N. Peterson (2013). Belief in the unstructured interview: The persistence of an illusion. *Judgement and Decision Making 8*(5), 512–520.

[40] Datta, S. (2017). *A History of the Indian University System: Emerging from the Shadows of the Past*. Palgrave Macmillan.

[41] Deshpande, A. (2011). *The Grammar of Caste: Economic Discrimination in Contemporary India*. Oxford: Oxford University Press.

[42] DeVaul, R., F. Jervey, J. Chappell, P. Caver, B. Short, and S. O'Keefe (1987). Medical school performance of initially rejected students. *Journal of the American Medical Association 257*, 47–51.

[43] Durlauf, S. N. (1999). The case "against" social capital. In P. Dasgupta and I. Serageldin (Eds.), *Social Capital: A Multifaceted Perspective*, pp. 84–104. Washington, D.C.: World Bank.

[44] Eames, T. (2024). Taryn versus taryn (she/her) versus taryn (they/them): A field experiment on pronoun disclosure and hiring discrimination. Working Paper.

[45] Eames, T. (2025). How does sex-based occupational discrimination affect nonbinary applicants? *AEA Papers and Proceedings*.

[46] Economic Advisory Council to the Prime Minister (2022). The State of Inequality in India Report. Government of India.

[47] Espinosa, M. and C. T. Stanton (2022, July). Training, communications patterns, and spillovers inside organizations. Working Paper 30224, National Bureau of Economic Research.

[48] Fenizia, A. (2022). Managers and productivity in the public sector. *Econometrica 90*(3), 1063–1084.

[49] Friedman, S. and D. Laurison (2020). *The Class Ceiling: Why it Pays to be Privileged?* Bristol University Press, Policy Press.

[50] Gibbons, R. and J. Roberts (2013). Economic theories of incentives in organizations. *The Handbook of Organizational Economics*.

[51] Gorman, E. H. (2005). Gender stereotypes, same-gender preferences, and organizational variation in the hiring of women: Evidence from law firms. *American Sociological Review 70*(4), 702–728.

[52] Grusky, D. B. (2005). Foundations of a neo-durkheimian class analysis. In E. O. Wright (Ed.), *Approaches to Class Analysis*, pp. 51–81. Cambridge University Press.

[53] Highet, K. (2023). "RIP English": Race, Class and 'Good English' in India. *Journal of Linguistic Anthropology 33*(2), 184–201.

[54] Highhouse, S. (2008). Stubborn Reliance on Intuition and Subjectivity in Employee Selection. *Industrial and Organizational Psychology: Perspectives on Science and Practice 1*(3), 333–342.

[55] Hoffman, M., L. B. Kahn, and D. Li (2018). Discretion in Hiring. *The Quarterly Journal of Economics 133*(2), 765–800.

[56] Hoffman, M. and C. T. Stanton (2024). People, practices, and productivity: A review of new advances in personnel economics. NBER Working Paper 32849.

[57] Hoffman, M. and S. Tadelis (2021). People management skills, employee attrition, and manager rewards: An empirical analysis. *Journal of Political Economy 129*(1), 243–285.

[58] Human Rights Watch (2001). Caste discrimination: A global concern. HRW Briefing Paper. See ch. 2, p. 2: "Lower-caste communities are almost invariably indistinguishable in physical appearance from higher-caste communities".

[59] Imas, A. (2016). The realization effect: Risk-taking after realized versus paper losses. *American Economic Review 106*(8), 2086–2109.

[60] Imas, A., G. Loewenstein, and C. K. Morewedge (2021). Mental money laundering: A motivated violation of fungibility. *Journal of the European Economic Association 19*(4), 2209–2233.

[61] Imas, A. and K. Madarász (2024). Superiority-seeking and the preference for exclusion. *Review of Economic Studies 91*(4), 2347–2386.

[62] Jodhka, S. S. and J. Naudet (2019). *Mapping the Elite: Power, Privilege, and Inequality.* Oxford University Press.

[63] Jodhka, S. S. and K. Newman (2007). In the name of globalisation: Meritocracy, productivity and the hidden language of caste. *Economic and Political Weekly 42*(41), 4125–4132.

[64] Jodkha, S. S. (2017). *Caste in Contemporary India.* Routledge India.

[65] Kahneman, D. (2011). *Thinking, Fast and Slow.* New York: Farrar, Straus and Giroux.

[66] Kausel, E. E., S. S. Culbertson, and H. P. Madrid (2016). Overconfidence in personnel selection: When and why unstructured interview information can hurt hiring decisions. *Organizational Behavior and Human Decision Processes 137*, 27–44.

[67] Khan, S. R. (2011). *Privilege: The Making of an Adolescent Elite at St. Paul's School.* Princeton, NJ: Princeton University Press.

[68] Kline, P., E. Rose, and C. Walters (2022). Systemic discrimination among large u.s. employers. *Quarterly Journal of Economics 137*(4), 1963–2036.

[69] Knobe, J. and S. Nichols (2008). An introduction to experimental philosophy. In *Experimental Philosophy*, pp. 3–34. Oxford: Oxford University Press.

[70] Kuncel, N. R., M. Credé, and L. L. Thomas (2007). A meta-analysis of the predictive validity of the graduate management admission test (gmat) and undergraduate grade point average (ugpa) for graduate student academic performance. *Academy of Management Learning & Education 6*(1), 51–68.

[71] Kuncel, N. R. and S. A. Hezlett (2007). Standardized tests predict graduate students' success. *Science 315*(5815), 1080–1081.

[72] Kuncel, N. R., D. M. Klieger, B. S. Connelly, and D. S. Ones (2013). Mechanical versus clinical data combination in selection and admissions decisions: A meta-analysis. *Journal of Applied Psychology 98*(6), 1060–1072.

[73] Kuncel, N. R., R. J. Kochevar, and D. S. Ones (2014). A meta-analysis of letters of recommendation in college and graduate admissions: Reasons for hope. *International Journal of Selection and Assessment 22*(1), 101–107.

[74] Lamont, M. (1992). *Money, Morals, and Manners: The Culture of the French and the American Upper-Middle Class*. Chicago: University of Chicago Press.

[75] Lamont, M. (2009). *How Professors Think: Inside the Curious World of Academic Judgment*. Cambridge, MA: Harvard University Press.

[76] Lamont, M. and A. Lareau (1988). Cultural capital: Allusions, gaps and glissandos in recent theoretical developments. *Sociological Theory 6*(2), 153–168.

[77] Laurison, D. and S. Friedman (2016). The class pay gap in higher professional and managerial occupations. *American Sociological Review 81*(4), 668–695.

[78] Laurison, D. and S. Friedman (2019). *The Class Ceiling: Why it Pays to be Privileged*. Bristol, UK: Policy Press.

[79] Lazear, E. P., K. L. Shaw, and C. T. Stanton (2015). The value of bosses. *Journal of Labor Economics 33*(4), 823–861.

[80] Lenski, G. E. (1954). Status crystallization: A non-vertical dimension of social status. *American Sociological Review 19*(4), 405–413.

[81] Madheswaran, S. (2008). Is affirmative action policy for private sector necessary? *Indian Journal of Industrial Relations 44*(2), 164–184.

[82] Meehl, P. E. (1954). *Clinical versus Statistical Prediction: A Theoretical Analysis and a Review of the Evidence*. University of Minnesota Press.

[83] Meehl, P. E. (1986). Causes and effects of my disturbing little book. *Journal of Personality Assessment 50*(3), 370–375.

[84] Ministry of Law and Justice, Government of India (2021). The constitution of india. Article 15 explicitly prohibits discrimination on grounds of religion, race, caste, sex, or place of birth.

[85] Minni, V. (2023). Making the invisible hand visible: Managers and the allocation of workers to jobs. R&R, *Quarterly Journal of Economics*.

[86] Minni, V. (2024). Global managers, local workers: Wage setting inside a multinational firm. *AEA Papers & Proceedings 114*, 586–91.

[87] Minni, V., K.-T. Nguyen, and H. Sarsons (2025). Managers and the cultural transmission of gender norms. Working Paper.

[88] Mishra, N. (2015). India and colorism: The finer nuances. *Washington University Global Studies Law Review 14*, 725–750.

[89] Murphy, S. C., D. M. Klieger, M. J. Borneman, and N. R. Kuncel (2009). The predictive power of personal statements in admissions: A meta-analysis and cautionary tale. *Educational and Psychological Measurement 69*(6), 967–982.

[90] Newman, K. and S. Thorat (2012). *Blocked by Caste: Economic Discrimination in Modern India*. Oxford University Press.

[91] Novosad, P., S. Asher, C. Farquharson, and E. Iljazi (2024). Access to opportunity in the sciences: Evidence from the nobel laureates. Working Paper.

[92] Oster, E. (2019). Unobservable selection and coefficient stability: Theory and evidence. *Journal of Business & Economic Statistics 37*(2), 187–204.

[93] Parameswaran, R. and K. Cardoza (2015). Melanin on the margins: Advertising and the cultural politics of fair/light/white beauty in india. *Journalism & Communication Monographs 11*(3), 213–274.

[94] Parmar, R. (2020). Transacting caste in modern times: Changing social identity through surnames in urban Gujarat. *Contemporary Voice of Dalit 12*(2), 220–233.

[95] Pinker, S. (2021). *Rationality: What It Is, Why It Seems Scarce, Why It Matters*. New York: Viking.

[96] Puri, S. (2025, February). Capitalism with Indian characteristics. Stigler Center Lecture Series, University of Chicago Booth School of Business. Accessed: 2025-02-11.

[97] Ridgeway, C. L. (2019). *Status: Why Is It Everywhere? Why Does It Matter?* New York, NY: Russell Sage Foundation.

[98] Rivera, L. A. (2012). Hiring as cultural matching: The case of elite professional service firms. *American Sociological Review 77*(6), 999–1022.

[99] Rivera, L. A. (2015). *Pedigree: How Elite Students Get Elite Jobs*. Princeton University Press.

[100] Sacerdote, B., D. O. Staiger, and M. Tine (2025, January). How test optional policies in college admissions disproportionately harm high achieving applicants from disadvantaged backgrounds. Working Paper 33389, National Bureau of Economic Research. Available at http://www.nber.org/papers/w33389.

[101] Sadun, R., N. Bloom, and J. Van Reenen (2017, September-October). Why do we undervalue competent management? Neither great leadership nor brilliant strategy matters without operational excellence. *Harvard Business Review 95*(5), 120–127. Winner of 59th Annual HBR McKinsey Award.

[102] Sadun, R., J. B. Fuller, S. Hansen, and P. Neal (2022, July-August). The C-suite skills that matter most. *Harvard Business Review 100*(4), 42–50.

[103] Sarsons, H. (2017). Recognition for group work: Gender differences in academia. *AEA Papers and Proceedings 107*(5), 141–145.

[104] Sarsons, H. (2023). Interpreting Signals in the Labor Market: Evidence from Medical Referrals. R&R, *Review of Economic Studies*.

[105] Sarsons, H., K. Gërxhani, E. Reuben, and A. Schram (2021). Gender differences in recognition for group work. *Journal of Political Economy 129*(1), 101–147.

[106] Sarsons, H., J. Hjort, and X. Li (2022). Across-Country Wage Compression in Multinationals. R&R, *American Economic Review*.

[107] Sebring, J. M. (1969). Caste indicators and caste identification of strangers. *Human Organization 28*(3), 199–207.

[108] Sharma, D. (2017). English in india. In J. Filipović and S. C. Deshors (Eds.), *World Englishes: New Theoretical and Methodological Considerations*, pp. 311–332. De Gruyter Mouton.

[109] Stansbury, A. and K. Rodriguez (2024). The class gap in career progression: Evidence from academia. R&R, *Econometrica*.

[110] Stansbury, A. and R. Schultz (2023). The economics profession's socioeconomic diversity problem. *Journal of Economic Perspectives 37*(4), 207–230.

[111] Subramanian, S. (2019). India's elite colleges and the employment divide. *Economic and Political Weekly 54*(2), 42–48.

[112] Todorov, A. (2017). *Face Value: The Irresistible Influence of First Impressions*. Princeton, NJ: Princeton University Press.

[113] Turco, C. J. (2010). Cultural foundations of tokenism: Evidence from the leveraged buyout industry. *American Sociological Review 75*(6), 894–913.

[114] University World News (2022). Thousands sit 'world's toughest' entrance exam, for IITs. University World News, August 27, 2022.

[115] Wiltshire, C. R. (2020). *Uniformity and Variability in the Indian English Accent.* Cambridge University Press.

[116] Yenmez, M. B. and T. Sönmez (2022). Affirmative action in India via vertical, horizontal, and overlapping reservations. *Econometrica 90*(3), 1143–1176.

[117] Zimmerman, S. D. (2019). Elite colleges and upward mobility to top jobs and top incomes. *American Economic Review 109*(1), 1–47.

A ONLINE APPENDIX

A.1 Representativeness of the Job Fair at the Elite College

The studied on-campus job fair provides a reliable window into how India's elite-college graduates enter elite private-sector jobs, for four reasons:

1. **Near-identical fairs across elite colleges.** Since the 1950s, India's top institutions—modeled after MIT and Stanford—have closely replicated each other's academic calendars, faculty-to-student ratios, and placement-office procedures. Today, the 23 Indian Institutes of Technology, 20 Indian Institutes of Management, and leading Delhi University colleges all conduct virtually identical selective job fairs (e.g., Central Placement Cell, Delhi University).

2. **Graduates predominantly flow into elite private-sector jobs.** Most alumni of these institutions join high-paying entry-level positions at multinational corporations rather than opting for the public sector (Subramanian, 2019; Jodhka and Naudet, 2019).

3. **Elite firms hire almost exclusively from elite colleges.** Recruiters in my sample visit roughly 40 similar campuses each year, each enrolling over 1,000 students.[OA.1] See also Subramanian (2019); Jodhka and Naudet (2019).

4. **Few students skip the fair.** Fewer than 5% of final-year students deregister from the placement process (Section 3.2); comparable deregistration rates hold at peer institutions (Subramanian, 2019).

Collectively, these facts demonstrate that the focal job fair closely mirrors elite hiring pathways across India's top universities.

A.2 Timeline and Screening Rounds

Every job in the sample follows the same four-stage sequence, eliminating candidates at each step:

1. **Application screening (September).** Recruiters conduct an initial screening, primarily based on GPA, major, degree, and standout résumé items (e.g., internships, competitions).

2. **Written aptitude tests (September–November).** Online technical assessments vary by role: coding exercises for tech/manufacturing jobs, and case-based tests for consulting positions. Firms advertising multiple posts usually administer one common test (e.g., Nvidia uses a single exam for all advertised profiles).

3. **Group debates (September–November).** A non-technical evaluation of soft skills:

[OA.1] Firm lists were scraped from multiple placement websites; details available upon request.

(a) Groups of 30–35 students, split into two teams, debate a general prompt (e.g., "Is social media a boon or a bane?"), while recruiters observe communication, teamwork, and confidence.

(b) Firms often conduct a single group debate for all advertised roles—for example, Google may simultaneously assess applicants for program manager and software engineer positions.

4. **Personal interviews (December–January).** Final 30-minute "fit" conversations:

   (a) Interviews occur simultaneously in multiple parallel batches, each handled by different members of the recruiting team.

   (b) As with earlier rounds, a single interview panel commonly covers multiple job titles (e.g., Uber may interview data scientist and front-end engineer candidates within the same track).

   (c) Because more than 98% of roles are based in India, interviewers are almost certainly Indian (Section 4.6).

A.3 Extending the Simple Model of Late-Stage Discrimination with Imperfect Information

In Section 5.7.1, I assumed that employers observe each candidate's true screening-round performance perfectly (i.e., $M_{s,i,g} = A_{i,g}$ at round s). Here, I show that *all* key insights—particularly how variation in job selectivity identifies discrimination versus ability deficits—remain valid even when employers observe *noisy* performance signals.

Setup. As in Section 5.7.1, there are two groups (castes) $g \in \{0, 1\}$, a continuum of jobs $j \in \mathcal{J}$, and four sequential screening rounds $\{1, 2, 3, 4\}$. Rounds $1, 2, 3$ constitute pre-interview filtering (e.g., applications, written tests, group debates), while round 4 is the final personal interview. Each candidate i from group g has a *latent ability* $A_{i,g}$. Passing round s requires

$$\mathbb{E}\big[M_{s,i,g} \,\big|\mathcal{I}_{s-1}\big] \;\geq\; \tau_{s,j}, \tag{OA.1}$$

where $\mathcal{I}_{s-1}$ denotes the employer's information after observing signals in rounds $1, \ldots, s-1$. Thresholds $\tau_{s,j}$ weakly increase with each stage and may vary across jobs.

Noisy Signals per Round. At each stage s, the employer observes

$$M_{s,i,g} \;=\; A_{i,g} \;+\; \eta_{s,i,g},$$

where $\eta_{s,i,g}$ is a zero-mean noise term. We relax the assumption of independence across candidates and rounds, requiring only that:

(i) $\mathbb{E}[\eta_{s,i,g}] = 0$ for all s, i, g, and

(ii) $\eta_{s,i,1} \overset{d}{=} \eta_{s,i,0}$, so that for any subset of candidates (including those who have passed earlier rounds), there is no systematic caste-level difference in noise distributions that favors one caste.

Thus, while noise might correlate across rounds or candidates (e.g., candidate performance may be interdependent across stages), it never introduces systematic group-level biases. Under these conditions, persistent group differences cannot result from noise alone.

Updating Beliefs and Passing Rounds. At each stage s, the employer updates belief about $A_{i,g}$ given $\mathcal{I}_s \equiv \{M_{1,i,g}, \ldots, M_{s,i,g}\}$. Let

$$\mathbb{E}[A_{i,g} \mid \mathcal{I}_s] \quad \text{and} \quad Var(A_{i,g} \mid \mathcal{I}_s)$$

denote the posterior mean and variance of $A_{i,g}$. With each successive stage, beliefs about candidate ability become more precise, and the passing condition can be evaluated using $\mathbb{E}[M_{s,i,g} \mid \mathcal{I}_{s-1}]$.

Final Interview Thresholds and Caste Revelation. As before, let $\tau_{4,j}^{(g)}$ be the *group-specific* final-round threshold at job j. If caste becomes salient at the interview for group 1, a higher threshold ($\tau_{4,j}^{(1)} > \tau_{4,j}^{(0)}$) indicates discrimination. The goal is to demonstrate how variation in thresholds across jobs disentangles true ability differences from discrimination—even with noisy signals.

A.3.1 No Early Caste Penalty Implies No Large Right-Tail Deficit Before Interviews

A substantial right-tail ability deficit for group 1 would yield a particularly large gap in pass rates at early stages for highly selective jobs (with high $\tau_{s,j}$). Formally:

Proposition 3 (No Early Gap $\implies$ No Right-Tail Deficit Under Imperfect Information)**.** Consider any pre-interview stage $s < 4$. If group 1 truly has fewer high-ability candidates in the top tail, then for sufficiently high $\tau_{s,j}$,

$$P\big(\text{Pass}_{s,i,j} \mid g = 1, \mathcal{I}_{s-1}\big) \;<\; P\big(\text{Pass}_{s,i,j} \mid g = 0, \mathcal{I}_{s-1}\big).$$

Therefore, empirically observing no such gap—even in highly selective jobs—rules out a meaningful right-tail deficit for group 1 before the interview stage.

Sketch. With noisy signals, passing round s depends on the posterior mean $\mathbb{E}[M_{s,i,g} \mid \mathcal{I}_{s-1}]$. For high $\tau_{s,j}$, employers require a sufficiently high $\mathbb{E}[A_{i,g} \mid \mathcal{I}_{s-1}]$. If group 1 has fewer high-ability individuals, as $\tau_{s,j}$ rises, the pass rate for group 1 falls faster than group 0. This mirrors the perfect-information argument in Proposition 1 and further relies on the *regularly decaying tail* assumption (e.g., $P(M_{s,i,g} \geq m) \sim C_g\, m^{-\alpha}$ for large m), which ensures that higher cutoffs intensify right-tail gaps.

Because noise is symmetric across groups, it cannot conceal large latent ability gaps at high cutoffs. Thus, the absence of early-stage gaps even at very selective jobs implies no meaningful right-tail deficit for group 1. □

A.3.2 Stable Promotion Gaps Across Selectivity Still Imply Discrimination

Post-hire performance Y_i depends on latent ability $A_{i,g}$. Promotion requires $Y_i \geq \kappa_j$. Final-round passing demands $\mathbb{E}[M_{4,i,g} \mid \mathcal{I}_3] \geq \tau_{4,j}^{(g)}$. Suppose top $(1-q)$ fraction have cutoffs $\tau_{4,j}^{q}$. If both groups face identical thresholds, their hired performance converges, and the promotion gap vanishes. However, if, $\tau_{4,j}^{(1)} = \tau_{4,j}^{(0)} + \delta$ $(\delta > 0)$, group 1's marginal hires possess higher posterior means, producing a stable advantage in realized performance.

Proposition 4 (Stable Promotion Gaps Under Imperfect Information Identify Discrimination)**.** Under the regularly decaying tail assumption used in Proposition 2 (i.e., for large m, $P(M_{4,i,g} \geq m) \sim C_g\, m^{-\alpha}$ with $\alpha > 1$), if group 1 faces a threshold higher by a constant $\delta > 0$, then average hired performance for group 1 remains about δ units higher than group 0, even in highly selective jobs. Hence, persistent promotion-rate gaps at high selectivity levels reject an unobserved-ability explanation in favor of discrimination.

Sketch. The argument parallels Proposition 2. Under imperfect information, stage-4 passage requires $\mathbb{E}[M_{4,i,g} \mid \mathcal{I}_3] \geq \tau_{4,j}^{(g)}$. If $\tau_{4,j}^{(1)} = \tau_{4,j}^{(0)} + \delta$, typical group 1 hires satisfy:

$$\mathbb{E}[A_{i,1} \mid \mathcal{I}_3] \;\approx\; \tau_{4,j}^{(0)} + \delta \;-\; \mathbb{E}[\eta_{4,i,1} \mid \mathcal{I}_3].$$

Because noise is identically distributed across groups (mean-zero, no systematic bias), raising group 1's threshold by δ effectively shifts their final hires toward higher regions of the ability distribution. At high selectivity, each group's hired candidates converge around their respective thresholds. Since group 1's threshold is higher by δ, their average hired performance remains about δ above group 0's. Conversely, identical thresholds shrink the performance gap to zero. Regular tail decay ensures marginal hires dominate outcomes at high cutoffs, maintaining a stable δ-sized gap under discrimination.

□

Key Takeaway. Allowing for imperfect information preserves key identifying variations. The absence of pre-interview deficits at high selectivity, combined with stable or increasing performance gaps post-hire, strongly indicates stricter interview thresholds imposed on disadvantaged candidates after caste revelation, rather than hidden ability shortfalls.

A.4 Tables and Figures

Table OA.1: **Caste Differences in Mean Salary Among Graduates**

	Dependent variable:			
	log (Salary)			
	Without Controls	With Controls	Client Facing	Non-Client Facing
	(1)	(2)	(3)	(4)
Disadv. Caste	-0.231^{***}	-0.103^{***}	-0.107^{***}	-0.089^{***}
	(0.020)	(0.017)	(0.027)	(0.022)
Disadv. Caste (Oster)		-0.076^{***}	-0.095^{***}	-0.072^{***}
		(0.019)	(0.026)	(0.022)
Constant	11.235^{***}	10.863^{***}	10.883^{***}	10.848^{***}
	(0.012)	(0.040)	(0.057)	(0.052)
Observations	2,527	2,527	849	1,678
Adjusted R^2	0.049	0.412	0.549	0.382
Year FE	✓	✓	✓	✓
Other Controls	✓	✓	✓	✓
Job FE				

Notes: Online Appendix Table OA.1 includes estimates from an earnings regression run on the sample of all students who graduated with jobs obtained through the college's job fair for full-time jobs. The dependent variable is log earnings. The regressions include detailed controls including measures of pre-college skills, college academic performance, previous labor market experience, and other employer-relevant skills such as major and degree. Column (1) reports estimates from a linear earnings regression without controls. Column (2) reports estimates from a linear earnings regression with controls. Columns (3) and (4) report estimates from a linear earnings regression with controls run separately for client-facing and non-client-facing jobs, respectively. Note that the earnings regressions do not include job fixed effects because there are no caste differences in pay for identical jobs. Bias-adjusted coefficients are calculated using Oster (2019). Following Oster's recommended benchmarks, I set $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$, implying that unobservables are assumed to have the same degree of selection as observables, and the total explanatory power can increase by up to 30% beyond the observed R^2. $^{*}p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

Table OA.2: **Caste Differences in GPA Between Those Omitted From the Earnings Regression Versus Those Who Are Not**

	Dependent variable:	
	log (GPA)	
	(1)	(2)
Disadv. Caste	-0.033^{***}	-0.032^{***}
	(0.004)	(0.004)
Omitted	-0.167^{***}	-0.158^{***}
	(0.005)	(0.005)
Omitted * Disadv. Caste	-0.071^{***}	-0.070^{***}
	(0.007)	(0.007)
Constant	2.147^{***}	2.127^{***}
	(0.004)	(0.008)
Observations	4,164	4,164
Adjusted R^2	0.546	0.559
Year FE	✓	✓
Major		✓
Degree		✓
Job FE		
Note:	*p<0.1; **p<0.05; ***p<0.01	

Notes: Online Appendix Table OA.2 shows the caste differences in GPA between those omitted from the earnings regression versus those who are not. The "omitted" category includes those who deregistered from the placement process and those who registered but did not receive any job offers (Section 5.2). The latter group often participates in the aftermarket and usually pursue higher education or entrepreneurship (Section 3.2). Column (1) only includes year fixed effects. Column (2) adds major and degree as controls. Note that the regressions do not have job fixed effects because these can typically only be included for jobs in my main sample. Therefore, these job fixed effects cannot be separately identified from the coefficient on "omitted."

Table OA.3: **Caste Differences in GPA At Each Job Search Stage**

	Dependent variable:					
	log (GPA)					
	App	Test	Debate	Interview	Offer	Choice
	(1)	(2)	(3)	(4)	(5)	(6)
Disadv. Caste	-0.111^{***}	-0.049^{***}	-0.039^{***}	-0.019^{***}	0.010^{***}	0.009^{***}
	(0.005)	(0.003)	(0.002)	(0.001)	(0.002)	(0.002)
Disadv. Caste (Oster)	-0.104^{***}	-0.035^{***}	-0.034^{***}	-0.014^{***}	0.015^{***}	0.014^{***}
	(0.006)	(0.003)	(0.002)	(0.002)	(0.002)	(0.002)
Constant	2.057^{***}	2.131^{***}	2.163^{***}	2.212^{***}	2.214^{***}	2.199^{***}
	(0.028)	(0.024)	(0.023)	(0.006)	(0.006)	(0.028)
Mean GPA	7.776	7.967	8.17	8.26	8.419	8.419
Observations	4,005	3,618	3,266	3,031	2,527	2,527
Adjusted R^2	0.264	0.726	0.745	0.801	0.822	0.822
Year FE	✓	✓	✓	✓	✓	✓
Job FE	✓	✓	✓	✓	✓	✓
Major	✓	✓	✓	✓	✓	✓
Degree	✓	✓	✓	✓	✓	✓
Note:						*p<0.1; **p<0.05; ***p<0.01

Notes: Online Appendix Table OA.3 compares the *between*-caste and *within*-job GPA differences at each job search stage. These differences are *controlled*, with the linear regressions including standard controls such as major, degree, and year, as well as job fixed effects. Bias-adjusted coefficients are calculated using Oster (2019). Following Oster's recommended benchmarks, I set $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$, implying that unobservables are assumed to have the same degree of selection as observables, and the total explanatory power can increase by up to 30% beyond the observed R^2. $^{*}p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

Table OA.4: **Caste Differences in Mean Salary at Each Job Search Stage in the Full Sample**

	Dependent variable:					
	log (Salary)					
	App	Test	Debate	Interview	Offer	Choice
	(1)	(2)	(3)	(4)	(5)	(6)
Disadv. Caste	0.004	−0.009	−0.006	−0.018*	−0.105***	−0.103***
	(0.005)	(0.007)	(0.007)	(0.010)	(0.013)	(0.017)
Disadv. Caste (Oster)	0.005	0.010	0.011	0.012	−0.079***	−0.076***
	(0.005)	(0.008)	(0.008)	(0.011)	(0.014)	(0.019)
Constant	11.032***	10.904***	10.848***	10.758***	10.672***	10.863***
	(0.011)	(0.015)	(0.015)	(0.023)	(0.030)	(0.040)
Observations	4,005	3,618	3,266	3,031	2,527	2,527
Adjusted R^2	0.578	0.491	0.547	0.549	0.488	0.412
Year FE	✓	✓	✓	✓	✓	✓
Other Controls	✓	✓	✓	✓	✓	✓
Job FE						

Notes: Online Appendix Table OA.4 includes estimates from an earnings regression run on the sample of all students who remain in contention for some job at a given job search stage. For example, column (1) includes the sample of all students who apply for jobs recruiting from the college's job fair. Similarly, column (2) includes the sample of all students who remain in contention for some job at the written test round, and so on. The dependent variable is log earnings. I include detailed controls, including measures of pre-college skills, college academic performance, previous labor market experience, and other employer-relevant skills such as major and degree. Each column is a separate regression and includes these controls. Note that the earnings regressions do not include job fixed effects because there are no caste differences in pay for identical jobs. Bias-adjusted coefficients are calculated using Oster (2019). Following Oster's recommended benchmarks, I set $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$, implying that unobservables are assumed to have the same degree of selection as observables, and the total explanatory power can increase by up to 30% beyond the observed R^2. $^*p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

Table OA.5: **Caste Differences in the Success Rate at Each Job Search Stage Among Client-Facing Jobs**

	Success App (1)	Success Test (2)	Success Debate (3)	Success Interview (4)
Disadv. Caste (OLS)	-0.0066***	-0.0007	-0.0136***	-0.3062***
	(0.0022)	(0.0030)	(0.0046)	(0.0081)
Disadv. Caste (Oster)	0.002	0.006	-0.015***	-0.314***
	(0.002)	(0.004)	(0.005)	(0.009)
Mean Success Rate	0.5770	0.7860	0.2590	0.2600
Observations	117,770	67,927	53,422	13,842
Adjusted R^2	0.61588	0.37467	0.01189	0.24298
Year FE	✓	✓	✓	✓
Job FE	✓	✓	✓	✓
Other Controls	✓	✓	✓	✓

Notes: Online Appendix Table OA.5 shows the caste difference in the success rate at each stage of job search among client-facing jobs. Each observation is a student who remains in contention at a given screening round in a given job in a given year. Each column reports a separate linear probability regression and includes controls. These regressions also include job fixed effects. Note that some firms conduct the same screening round (written tests, group debates, interviews) for multiple job positions. However, since these job positions within the same firm represent different jobs, the regressions consider these jobs and the corresponding advance-or-out decisions made by them separately. Bias-adjusted coefficients are calculated using Oster (2019). Following Oster's recommended benchmarks, I set $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$, implying that unobservables are assumed to have the same degree of selection as observables, and the total explanatory power can increase by up to 30% beyond the observed R^2. $^*p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

Table OA.6: **Caste Differences in the Success Rate at Each Job Search Stage Among Non-Client-Facing Jobs**

	Success App (1)	Success Test (2)	Success Debate (3)	Success Interview (4)
Disadv. Caste (OLS)	0.0024	0.0011	0.0070**	-0.2867***
	(0.0018)	(0.0022)	(0.0034)	(0.0060)
Disadv. Caste (Oster)	0.008***	0.014***	0.007**	-0.286***
	(0.002)	(0.002)	(0.003)	(0.007)
Mean Success Rate	0.6140	0.8250	0.2580	0.2620
Observations	187,771	115,223	95,067	24,536
Adjusted R^2	0.60783	0.32634	0.02270	0.24723
Year FE	✓	✓	✓	✓
Job FE	✓	✓	✓	✓
Other Controls	✓	✓	✓	✓

Notes: Online Appendix Table OA.6 shows the caste difference in the success rate at each stage of job search among non-client-facing jobs. Each observation is a student who remains in contention at a given screening round in a given job in a given year. Each column reports a separate linear probability regression and includes controls. These regressions also include job fixed effects. Note that some firms conduct the same screening round (written tests, group debates, interviews) for multiple job positions. However, since these job positions within the same firm represent different jobs, the regressions consider these jobs and the corresponding advance-or-out decisions made by them separately. Bias-adjusted coefficients are calculated using Oster (2019). Following Oster's recommended benchmarks, I set $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$, implying that unobservables are assumed to have the same degree of selection as observables, and the total explanatory power can increase by up to 30% beyond the observed R^2. $^*p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

Table OA.7: **Caste Differences in Mean Salary at Each Job Search Stage in the Identifiable Sample**

	Dependent variable:					
	log (Salary)					
	App	Test	Debate	Interview	Offer	Choice
	(1)	(2)	(3)	(4)	(5)	(6)
Disadv. Caste (OLS)	0.012	-0.068^{***}	-0.034^{*}	-0.105^{***}	-0.146^{***}	-0.107^{**}
	(0.014)	(0.019)	(0.018)	(0.027)	(0.033)	(0.049)
Disadv. Caste (Oster)	0.013	-0.049^{***}	-0.016	-0.078^{**}	-0.119^{***}	-0.082
	(0.016)	(0.021)	(0.022)	(0.032)	(0.042)	(0.055)
Constant	10.991^{***}	10.869^{***}	10.807^{***}	10.734^{***}	10.601^{***}	10.834^{***}
	(0.030)	(0.042)	(0.041)	(0.064)	(0.078)	(0.114)
Observations	619	552	500	449	378	378
Adjusted R^2	0.546	0.470	0.537	0.535	0.509	0.362
Year FE	✓	✓	✓	✓	✓	✓
Other Controls	✓	✓	✓	✓	✓	✓
Job FE						

Notes: Online Appendix Table OA.7 includes estimates from an earnings regression run on the sample of students with identifiable caste names who remain in contention for some job at a given job search stage. For example, column (1) includes the sample of all students with identifiable caste names who apply for jobs recruiting from the college's job fair. Similarly, column (2) includes the sample of all students with identifiable caste names who remain in contention for some job at the written test round, and so on. The dependent variable is log earnings. Each column is a separate regression and includes controls. Note that the earnings regressions do not include job fixed effects because there are no caste differences in pay for identical jobs. Bias-adjusted coefficients are calculated using Oster (2019). Following Oster's recommended benchmarks, I set $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$, implying that unobservables are assumed to have the same degree of selection as observables, and the total explanatory power can increase by up to 30% beyond the observed R^2. $^{*}p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

Table OA.8: **Caste Differences in Mean Salary at Each Job Search Stage in the Non-Identifiable Sample**

	Dependent variable:					
	log (Salary)					
	App	Test	Debate	Interview	Offer	Choice
	(1)	(2)	(3)	(4)	(5)	(6)
Disadv. Caste (OLS)	0.002	0.0004	−0.003	−0.004	−0.100***	−0.106***
	(0.005)	(0.007)	(0.007)	(0.011)	(0.014)	(0.019)
Disadv. Caste (Oster)	0.004	0.019***	0.013	0.024**	−0.075***	−0.076***
	(0.006)	(0.007)	(0.008)	(0.011)	(0.015)	(0.018)
Constant	11.041***	10.913***	10.857***	10.766***	10.685***	10.869***
	(0.012)	(0.017)	(0.016)	(0.024)	(0.033)	(0.043)
Observations	3,386	3,066	2,766	2,582	2,149	2,149
Adjusted R^2	0.584	0.500	0.552	0.556	0.486	0.421
Year FE	✓	✓	✓	✓	✓	✓
Other Controls	✓	✓	✓	✓	✓	✓
Job FE						

Notes: Online Appendix Table OA.8 includes estimates from an earnings regression run on the sample of students with non-identifiable caste names who remain in contention for some job at a given job search stage. For example, column (1) includes the sample of all students with non-identifiable caste names who apply for jobs recruiting from the college's job fair. Similarly, column (2) includes the sample of all students with non-identifiable caste names who remain in contention for some job at the written test round, and so on. The dependent variable is log earnings. Each column is a separate regression and includes controls. Note that the earnings regressions do not include job fixed effects because there are no caste differences in pay for identical jobs. Bias-adjusted coefficients are calculated using Oster (2019). Following Oster's recommended benchmarks, I set $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$, implying that unobservables are assumed to have the same degree of selection as observables, and the total explanatory power can increase by up to 30% beyond the observed R^2. $^*p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

Table OA.9: **Caste Differences in GPA Between Those Who Left the First Job Versus Those Who Did Not**

	Dependent variable:		
	log (GPA)		
	(1)	(2)	(3)
Disadv. Caste	-0.065^{***}	-0.065^{***}	0.010^{***}
	(0.004)	(0.004)	(0.002)
Left Job	-0.072^{***}	-0.072^{***}	-0.004^{**}
	(0.003)	(0.003)	(0.002)
Left Job * Disadv. Caste	0.047^{***}	0.047^{***}	-0.004
	(0.005)	(0.005)	(0.003)
Constant	2.188^{***}	2.182^{***}	2.198^{***}
	(0.003)	(0.005)	(0.028)
Year FE	✓	✓	✓
Major		✓	✓
Degree		✓	✓
Job FE			✓
Observations	2,489	2,489	2,489
Adjusted R^2	0.260	0.265	0.822
Note:		*p<0.1; **p<0.05; ***p<0.01	

Notes: Online Appendix Table OA.9 shows the caste differences in GPA between those who left the first job versus those who did not. Column (1) only includes year fixed effects. Column (2) adds major and degree as additional controls. Column (3) additionally controls for job fixed effects.

Figure OA.1: **Common Support Within Each College GPA Decile**

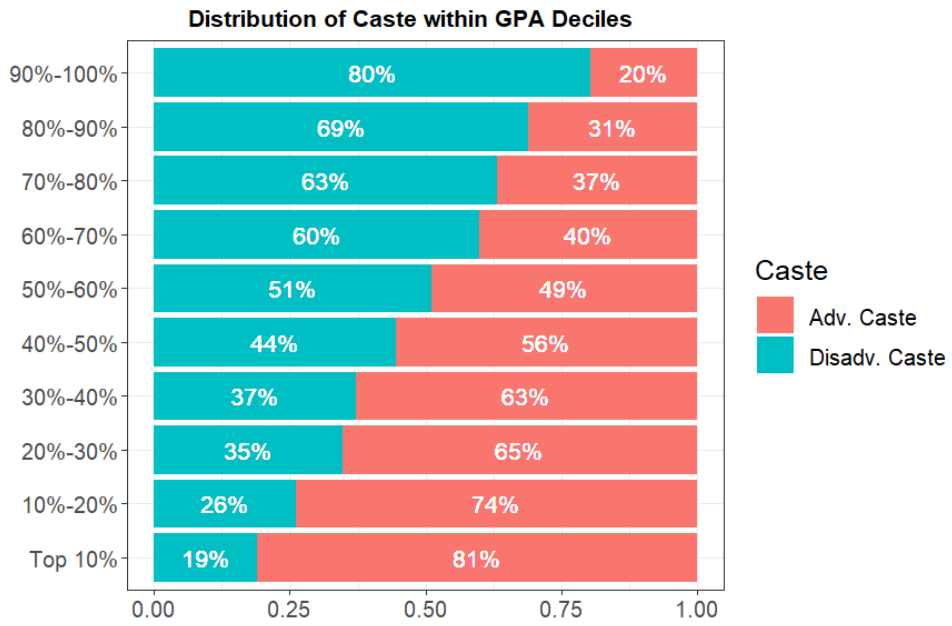

Notes: Online Appendix Figure OA.1 shows that there are students from both disadvantaged and advantaged castes within each college GPA decile. College GPA is normalized within each college degree and year.

Figure OA.2: **Caste Penalty in the Conditional Success Rate At the Application Reading Round in the Identifiable and the Non-Identifiable Sample**

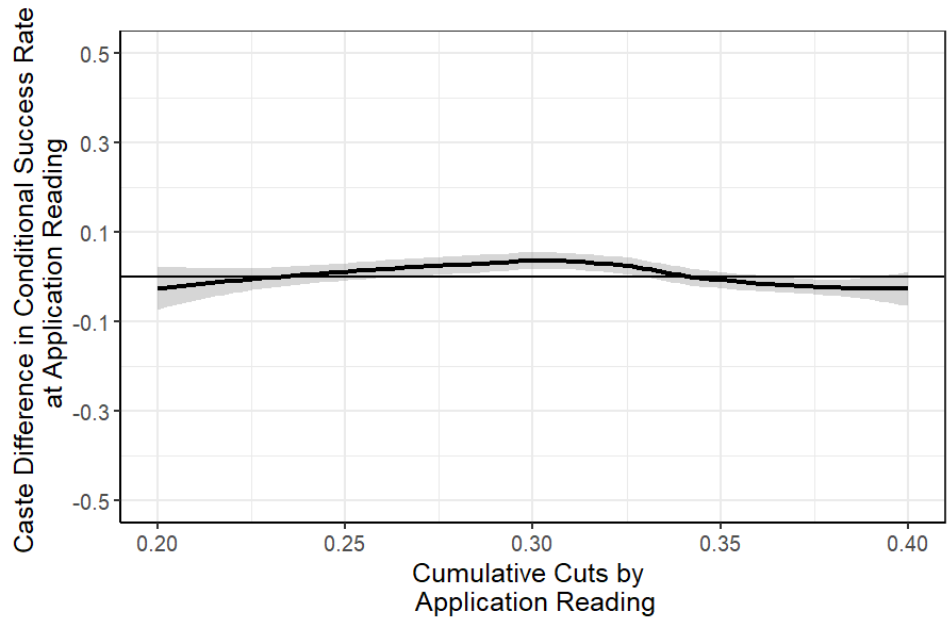

(a) Caste Penalty in Conditional Success Rates versus Cumulative Cuts by Application Reading in the Full Sample

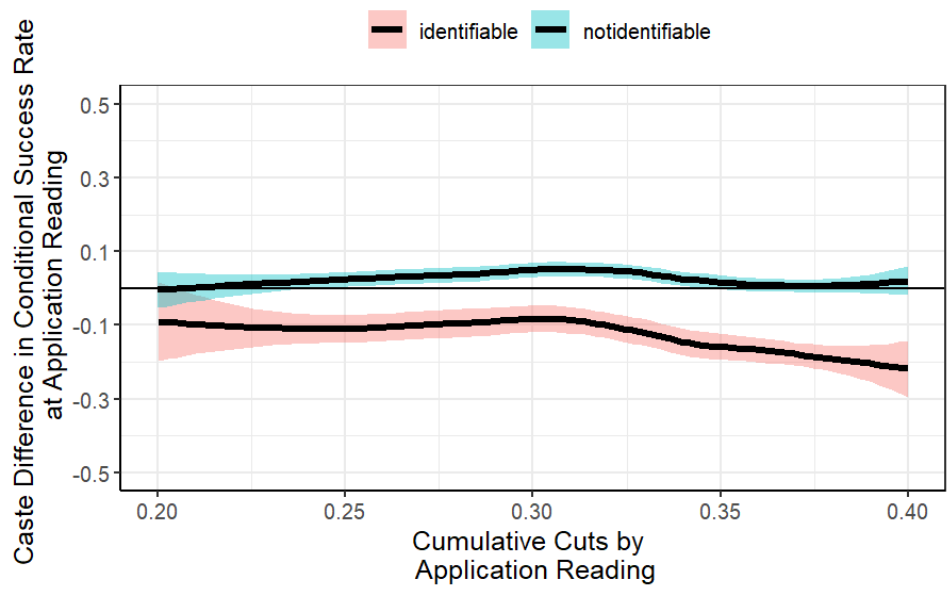

(b) Caste Penalty in Conditional Success Rates versus Cumulative Cuts by Application Reading in the Identifiable and Non-Identifiable Sample

Notes: Figure OA.2 shows the caste difference in the conditional (i.e., with controls) success rate against the cumulative cuts made by jobs by the end of the application reading round. Smoothed conditional means (thick black lines) are calculated using the "loess" function in R. The shaded areas in gray, blue, and red denote 95% confidence intervals. The thin solid horizontal line denotes a y-intercept of 0. Panel (a) is constructed from the full sample. Panel (b) is constructed separately for the sample with identifiable (red) and nonidentifable (blue) caste names. The x-axis shows the cumulative proportion of students cut by jobs by the end of the application reading round. The y-values are calculated by estimating, *for each job*, a standard linear probability model with controls and then plotting the coefficient on disadvantaged caste from these individual regressions on the y-axis. Each regression includes controls for student characteristics such as degree, major, college GPA, entrance exam score, and cohort year.

Figure OA.3: **Caste Penalty in the Uncontrolled Success Rate At Each Screening Round**

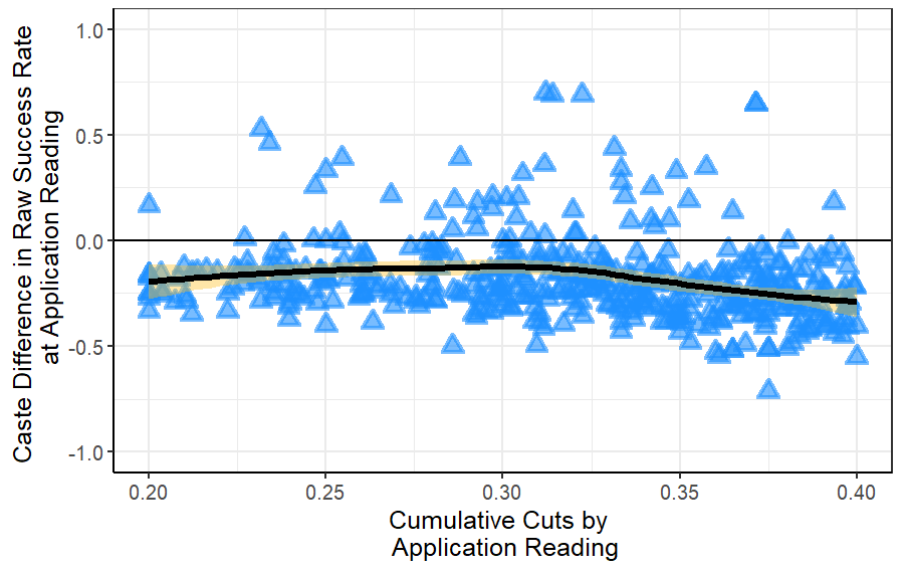

(a) Caste Penalty in Raw Success Rates versus Cumulative Cuts by Application Reading

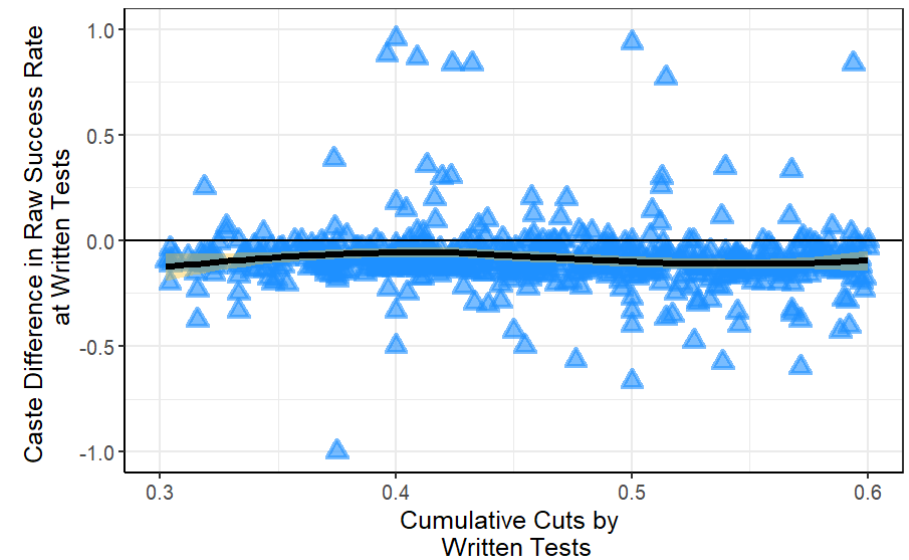

(b) Caste Penalty in Raw Success Rates versus Cumulative Cuts by Written Tests

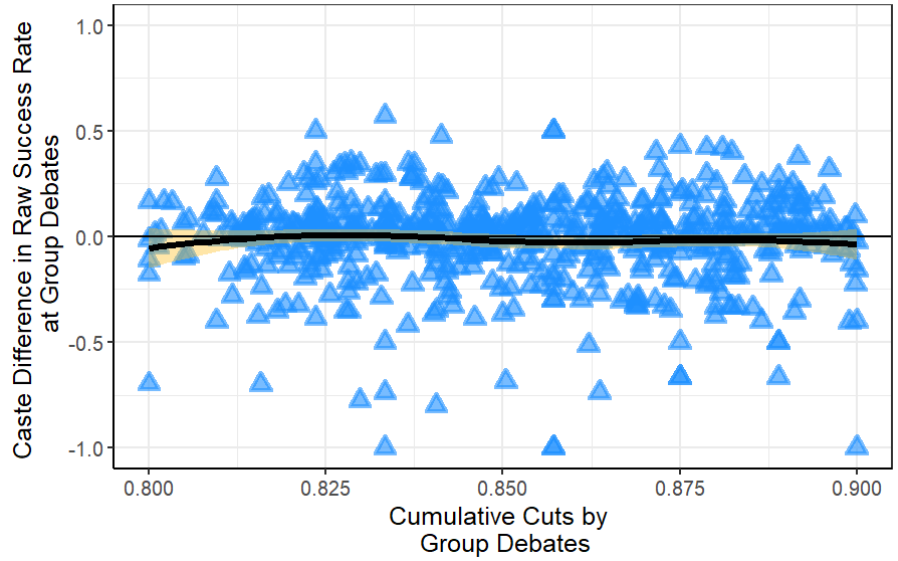

(c) Caste Penalty in Raw Success Rates versus Cumulative Cuts by Group Debates

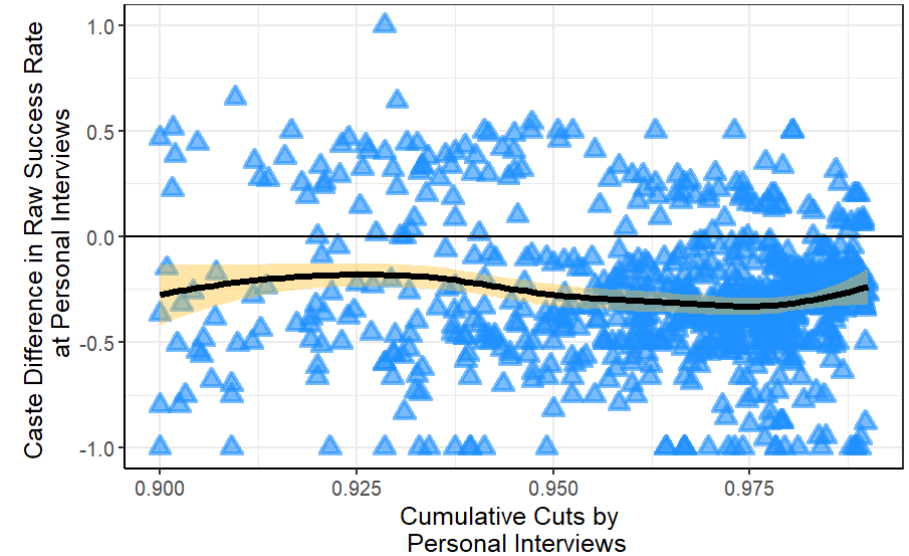

(d) Caste Penalty in Raw Success Rates versus Cumulative Cuts by Personal Interviews

Notes: Figure OA.3 shows the caste difference in the "raw" (i.e., without any controls) success rate against the cumulative cuts made by jobs by the end of each screening round. Each blue triangle denotes a job. Smoothed conditional means (thick black lines) are calculated using the "loess" function in R. The shaded areas in yellow denote 95% confidence intervals. The thin solid horizontal line denotes a y-intercept of 0. The x-axis shows the cumulative proportion of students cut by jobs by the end of each screening round. The y-axis shows the uncontrolled caste penalty in the success rate at each screening round.