Making the Elite: Class Discrimination at Multinationals

Soumitra Shukla*

April 20, 2025

Abstract

How salient is hiring discrimination in driving socioeconomic gaps in elite job access, and where do these disparities emerge? Answering these questions is challenging in the context of social status or class discrimination, particularly in Western contexts, where class must be inferred indirectly from subtle markers such as accent, educational credentials, and social networks. India's caste system offers a uniquely explicit measure of social class, enabling precise identification of how status affects access to elite positions. Using detailed screening data from over 1,000 hires at US and European multinationals in India, this paper offers the first systematic quantification of social-status bias in elite recruitment and clearly identifies its source. Subjective personal interviews assessing "fit" drive nearly 90% of caste-based hiring disparities. Disadvantaged candidates have approximately 20% higher promotion rates, even in highly selective roles, indicating these disparities are not justified by differences in candidate quality. Given these persistent advantages and the predominantly non-customer-facing roles, results suggest explicit bias rather than statistical or customer-driven discrimination. These findings underscore a critical organizational challenge: how can necessary screening discretion be preserved without reinforcing status-based discrimination that undermines diversity and talent allocation?

*Shukla: Research Affiliate, Department of Economics, Yale University (soumitra.shukla@yale.edu). I would like to thank Joseph Altonji, Costas Meghir, and John Eric Humphries for serving on my committee and providing guidance and feedback. I would also like to thank Jason Abaluck, Alex Albright, Peter Arcidiacono, Crémieux, Michael Dinerstein, Alessandra Fenizia, Arpit Gupta, Alex Imas, Patrick Kline, Scott Kominers, Eliza McNay, Evan Rose, Kamalini Ramadas, Jung Sakong, Rob Valletta, Jon Williams, Seth Zimmerman, seminar participants at the Yale labor workshop, the Chicago FED, Chicago Harris, Federal Reserve Board, Wharton (BEPP), Penn State, UCSD, SOLE, Southeastern Micro Labor Workshop, System Equitable Growth Conference at the New York Fed, the "Discrimination in the 21st Century Conference" organized by the Becker Friedman Institute and Chicago Booth, MIT Organizational Economics Lunch, and the NBER Race and Stratification Working Group for helpful feedback and support.

1 Introduction

Who gains access to elite jobs, and how do hiring practices shape inequality in high-stakes recruitment? Extensive research documents discrimination across labor markets (Altonji and Blank, 1999; Rivera, 2012, 2015; Friedman and Laurison, 2020), yet lacks precise evidence on *where*, *how*, and the *quantitative significance* of discrimination in competitive contexts such as multinational corporations (MNCs).

Social class or status discrimination poses a particular challenge (e.g., Lamont and Lareau, 1988; Durlauf, 1999, 2006, 2004, 2008; Ridgeway, 2019; Stansbury and Rodriguez, 2024). In Western settings, social position must be inferred from multiple ambiguous indicators—accents loosely tied to backgrounds, educational credentials obscuring family advantages, and elusive social networks—markers often culturally transmitted across generations (cf. Bisin and Verdier, 1998). These subtle markers complicate efforts to identify how status affects elite job access, making class discrimination particularly understudied within economics. Moreover, in elite recruitment contexts—where objective skill differences among candidates are often minimal—evaluators often rely disproportionately on subtle cultural markers to differentiate candidates, amplifying the role of cultural similarities or differences in hiring decisions (Lamont, 1992, 2009; Rivera, 2012).

Recent trends in meritocratic evaluations illustrate how well-intentioned policies inadvertently reinforce class advantages. For example, debates about standardized versus "holistic" assessments highlight this tension. Evidence from Sacerdote et al. (2025) suggests that "test-optional" policies in U.S. college admissions unintentionally disadvantage high-achieving students from lower socioeconomic backgrounds. Similarly, Chetty et al. (2023) shows selective colleges give excessive weight to "soft" factors (e.g., legacy preference and athletic recruitment) favoring high-income students, despite limited predictive validity for career success. India's elite civil service examinations, among the world's most competitive, exhibit comparable dynamics: subjective "personality" interviews disproportionately penalize high-performing candidates based on perceived social status with little predictive power for subsequent job performance (Bhavnani and Lee, 2021). Indeed, extensive evidence demonstrates that holistic, judgment-based evaluation methods consistently degrade predictive validity, even when employed by knowledgeable experts, thereby systematically undermining objective decision-making (Kuncel et al., 2013).

Although higher education and public-sector recruitment differ from private-sector hiring, all combine objective metrics with subjective assessments. Removing standardized tests can obscure academic merit (Kuncel and Hezlett, 2007; Bol et al., 2014); similarly, subjective "fit" evaluations during interviews may obscure strong underlying skills, particularly

among lower-status candidates. In this sense, "fit" often implicitly signals class or status, disproportionately disadvantaging less privileged candidates. While these dynamics are observed across contexts (Rivera, 2012; Laurison and Friedman, 2019; Stansbury and Rodriguez, 2024), precise measurement of "class" remains challenging.

To understand how decision makers such as recruiters, managers, or admissions committee members interpret class or status markers, dual-process cognition theories (e.g., Kahneman, 2011) are insightful. Decision-making often relies on quick, intuitive evaluations prone to biases rather than reflective judgment. Hiring managers evaluating elite candidates frequently base decisions on rapid impressions of "fit," influenced by superficial rapport (Todorov, 2017), subtle stereotypes (Knobe and Nichols, 2008), and immediate affective reactions (Pinker, 2021). These rapid assessments are often shaped by deeply internalized cultural norms transmitted through socialization, which reinforce evaluators' implicit expectations about class-appropriate behaviors, mannerisms, and attitudes (cf. Bisin and Verdier, 2010). Consequently, such biases become especially salient when negative social markers (e.g., signals of lower socioeconomic background) are realized directly through interpersonal interactions, compared to merely knowing these markers indirectly or abstractly (cf. Imas, 2016; Imas et al., 2021; Imas and Madarász, 2024). Similarly, subjective evaluations through personal statements and recommendation letters—common in admissions committees—often provide limited predictive validity regarding future performance and may disproportionately reflect assessors' biases rather than candidates' objective potential (Murphy et al., 2009; Kuncel et al., 2014). These subjective judgments obscure precisely which social identities—often signaled by subtle markers such as accent, etiquette, extracurricular experiences, vocabulary, or personal interests—are penalized (Rivera, 2015).

Extensive global evidence confirms that such "fit"-based evaluations lead to in-group biases, thereby systematically disadvantaging lower-status candidates, reinforcing preferences for demographically similar candidates—such as those sharing gender or social background with evaluators—in performance assessments and professional recognition (e.g., Gorman, 2005; Sarsons, 2017; Sarsons et al., 2021; Sarsons, 2023; Benson et al., 2023), and sustaining class-based barriers at advanced career stages (e.g., Rivera, 2015; Laurison and Friedman, 2016, 2019). Even elite institutions consistently struggle to overcome persistent class gaps, as privileged early-career networks shape later outcomes, perpetuating inequality across contexts (Stansbury and Schultz, 2023; Asher et al., 2024; Novosad et al., 2024). Admission into elite colleges disproportionately facilitates upward mobility for students from already privileged backgrounds, magnifying class-based disparities in access to top career positions and intergenerational advantages (Zimmerman, 2019; Barrios-Fernández et al., 2024)—patterns shaped by the joint dynamics of culture and institutional persistence across generations (e.g.,

Bisin and Verdier, 2017).

Although caste is specific to India, the underlying discrimination mechanisms are broadly relevant. Caste serves as a uniquely comprehensive measure of social stratification, consolidating economic, social, and cultural capital—concepts described by Bourdieu (1986) as financial resources, networks, and symbolic assets like education, accents, or mannerisms respectively—into a single hierarchy.[1] Even when advantaged caste families are poor or disadvantaged caste families wealthy, caste still shapes perceived social status (Beteille, 1965, 1969). By contrast, Western contexts rely on multiple ambiguous indicators to approximate class (cf., Khan, 2011), complicating efforts to isolate discrimination clearly.

This paper addresses these measurement challenges by analyzing recruitment data from leading US and European multinationals hiring in India, where caste provides a uniquely precise measure of social status. Despite affirmative action policies designed to improve educational outcomes for disadvantaged castes, these groups remain largely excluded from elite positions (Madheswaran, 2008). Leveraging this setting with precisely measurable social status, this study makes three distinct contributions: first, it provides systematic quantification of the salience of discrimination based on social position in elite recruitment; second, it identifies personal interviews as the primary mechanism through which subjective evaluations of candidate "fit" expose and penalize candidates based on social status; and third, it presents robust evidence that class bias—rather than statistical or customer discrimination—drives hiring disparities in entry-level recruitment at elite multinationals. To my knowledge, each of these contributions is the first of its kind in the literature.

The key innovation is demonstrating that personal interviews, often framed as evaluations of candidate "cultural fit," function predominantly as a channel through which social status is revealed and penalized. By "fit," this paper refers both to legitimate skill evaluations and intangible cultural or background alignment assessed by hiring managers. Detailed screening data indicate that the vast majority of conditional caste disparities in hiring emerge during final-round personal interviews. Moreover, disadvantaged castes hired into elite jobs subsequently exhibit significantly higher promotion rates, underscoring that subjective evaluations during interviews systematically undervalue their actual potential.

Elite entry-level jobs are crucial pathways to economic power, particularly in India, where nearly 70% of the population belongs to disadvantaged castes. The studied elite college job fair offers representative insights into hiring practices widely replicated by multinationals across top Indian colleges and their global offices (Bloom et al., 2012). This setting also highlights how earlier affirmative action benefits can erode during hiring stages. Nearly half

[1] For further sociological analyses expanding on Bourdieu's framework, see Lamont and Lareau (1988) and Rivera (2015).

the students attending the job fair belong to disadvantaged castes due to mandated affirmative action. While improper enforcement of affirmative action can disadvantage protected groups (Yenmez and Sönmez, 2022), this study identifies a distinct downstream mechanism: subjective "fit"-focused interviews reinstating social barriers. Thus, even well-designed affirmative action policies require attention to institutional practices at subsequent selection stages.

Each of the 1050 jobs studied has four screening rounds, which generally move from more standardized and objective assessments to less standardized and more subjective evaluations: application reading, written aptitude tests, group debates assessing socio-emotional skills, and final personal interviews evaluating job-relevant attributes, including candidate "fit" (Newman and Thorat, 2012; Rivera, 2012, 2015). The unconditional caste earnings gap is 23%, narrowing to 10% after controlling for pre-college skills, academic performance, and other employer-relevant factors. Because salaries are *identical* across castes for the same jobs and explicitly *non-negotiable*, stage-wise caste penalties reflect differences in job composition, enabling a decomposition of the final earnings penalty by counterfactually assigning the same salary at each stage before offers are accepted.[2] Personal interviews alone account for roughly 90% of the conditional earnings penalty, amounting to nearly 40% of the unconditional caste earnings gap. Similarly, around 95% of the caste penalty in hiring success rates is directly attributable to interviews. Thus, personal interviews assessing extracurricular activities, background, and cultural alignment afford employers the opportunity to highlight social position more saliently than previous stages, making caste disparities particularly pronounced.

The clearest evidence that *caste itself* is penalized emerges among the small subset (around 15%) of students with distinctive caste-identifiable surnames, where disparities arise as early as during the application reading round. However, in urban-educated contexts typical of elite Indian colleges, surnames generally do not reveal caste clearly, due to migration and diverse naming conventions (Jodkha, 2017; Wiltshire, 2020). Consequently, because most candidates do not have clearly caste-identifiable surnames, the *overall* caste penalty at the application-reading stage is negligible. Instead, the penalty primarily emerges during personal interviews, which thus serve as the principal mechanism through which social status is revealed and penalized.

This paper does not argue that personal interviews are useless. Rather, the evidence suggests that any *additional diagnostic information* from interviews *cannot justify* the observed caste disparities. Disadvantaged caste hires consistently outperform their advantaged

[2]Placement offices commonly enforce salary non-negotiation, ensuring offered salaries match initial firm advertisements (Section 3.2).

counterparts by about 20% in promotion rates, even in the most highly selective positions. This enduring promotion edge strongly indicates that hiring disparities arise primarily from discrimination rather than unobserved skill deficits.

To formalize this interpretation, I develop a theoretical framework leveraging job selectivity—how deep into the applicant pool each firm selects. The framework has two implications: first, if disadvantaged candidates genuinely lacked skills, highly selective stages prior to interviews should reveal this deficiency, yet the data show no such pre-interview caste penalty across the full distribution of job selectivity; second, if interviews accurately filtered out weaker candidates, disadvantaged hires would not consistently outperform advantaged peers across all job selectivity tiers. The underlying logic is straightforward: as job selectivity increases, the average quality of those hired approaches the marginal (cutoff) candidate quality. If disadvantaged hires consistently outperform advantaged peers across varying selectivity tiers, it implies either persistent differences in productivity at every margin or, more plausibly, stricter hiring thresholds (i.e., discrimination) for disadvantaged candidates. Given the consistent promotion advantage, differential cutoffs driven by discrimination are the more compelling explanation.

The evidence from the distinct versus non-distinct names analysis strongly indicates that hiring disparities arise primarily from disparate treatment—explicit caste-based discrimination. The persistent promotion advantage of disadvantaged hires across job selectivity levels and the fact that over 75% of these jobs are non-customer-facing strongly suggests that these disparities stem primarily from bias—either taste-based discrimination or biased beliefs—rather than statistical or customer discrimination (cf. Benson et al., 2024). Such discrimination is explicitly illegal under Article 15 of the Indian Constitution (Ministry of Law and Justice, Government of India, 2021). Beyond blatantly violating constitutional mandates, such discrimination severely restricts socioeconomic mobility and hampers the optimal allocation of talent.

This paper makes several contributions to the literatures on labor, personnel, and organizational economics by systematically quantifying and identifying the sources and mechanisms of social status discrimination in high-stakes elite hiring contexts. No other study thoroughly examines each stage of the entry-level hiring process at leading multinationals, pinpointing both the mechanism through which social status emerges as a barrier—subjective, discretionary "fit"-focused personal interviews—and its economic magnitude (nearly 90% of conditional hiring disparities). This precise quantification significantly advances a nascent literature on class-based discrimination, which remains underdeveloped relative to research on race and gender (Stansbury and Rodriguez, 2024).

This paper is most closely related to recent work by Stansbury and Rodriguez (2024),

who document substantial class gaps across academic careers in the U.S. but do not isolate the precise institutional sources of these barriers. Similarly, Bhavnani and Lee (2021) conceptually replicate these findings by demonstrating that subjective 'personality' interviews in India's elite civil service examinations systematically disadvantage lower-status candidates without accurately predicting job performance. Although they primarily focus on public-sector recruitment and do not provide detailed quantification of broader economic consequences, their findings parallel this paper's conclusions regarding interviews as mechanisms through which social status is identified and penalized. Likewise, Rivera (2012) qualitatively identifies "cultural matching" through discretionary hiring evaluations at elite U.S. firms—approximately three-quarters of these firms rank "fit" as their most critical hiring criterion—yet does not empirically quantify the economic scale of discrimination nor assess its implications for candidate outcomes. By leveraging the uniquely explicit measure of social status afforded by India's caste system and detailed stage-wise data from multinational recruitment, this paper bridges these gaps—providing granular evidence on precisely where and how discrimination occurs, quantifying its magnitude, and credibly distinguishing bias-driven discrimination from statistical or customer discrimination, a crucial distinction for accurately targeting policy and organizational interventions.

As longstanding markers of social status become increasingly diluted across societies (Lenski, 1954; Grusky, 2005; Clark, 2015), including urban-educated settings in India, class discrimination likely moves deeper into recruitment processes, increasing the necessity of granular hiring-stage data to reveal discrimination clearly. Thus, this study fills a key methodological gap, as more diffuse markers of social class in Western contexts have previously obscured the specific mechanisms through which status shapes hiring decisions. Furthermore, the findings have broader implications beyond visibly identifiable status markers: subjective interviews may similarly expose and disadvantage candidates based on concealable stigmatized identities, such as LGBTQ+ status (Eames, 2024, 2025).

Another central contribution of this paper is its insight into the role of managers as critical agents in the allocation of talent within organizations, emphasizing the significance of their cultural and social capital. Recent literature highlights that managerial decisions profoundly shape firm performance and workplace culture (e.g., Lazear et al., 2015; Hoffman and Tadelis, 2021; Espinosa and Stanton, 2022; Fenizia, 2022; Minni, 2023; Benson and Shaw, 2024; Minni et al., 2025), yet this paper uniquely integrates Bourdieu's conception of cultural capital to elucidate the biases underlying managerial assessments of "fit." According to Bourdieu, "cultural capital" consists of embodied assets (e.g., etiquette, accent, mannerisms, presentation style, conversational polish), objectified assets (e.g., internships at prestigious organizations), and institutionalized forms (e.g., membership in elite college

clubs, alumni status from elite high schools), which significantly influence status reproduction and access to elite positions (Bourdieu, 1986). Managers' own cultural capital shapes their perceptions and evaluations during hiring, systematically favoring candidates whose cultural assets align closely with their own (e.g., Rivera, 2012; Sadun et al., 2022). While many of these "cultural" markers might indeed correlate with productivity, the empirical results in this study indicate that the extra information managers glean from subjective evaluations of fit—assessments inherently influenced by shared cultural capital and thus rendering social position salient—is insufficient to justify the observed hiring disparities. Moreover, while we do not have direct information about managers' personal backgrounds, sociological research indicates that organizational socialization processes lead managers from varied backgrounds to converge on similar evaluative standards of cultural capital, reinforcing shared biases (e.g., Turco, 2010; Rivera, 2015). By empirically demonstrating how managers' subjective biases based on cultural alignment expose candidates' social status and subsequently facilitate discrimination that excludes objectively high-potential candidates, this study bridges the theoretical insights of cultural sociology with empirical personnel economics.

Additionally, this paper highlights critical agency problems between hiring managers and actual supervisors. Hiring managers seldom become the direct supervisors of the recruits due to logistical constraints—such as extensive travel requirements to multiple colleges over extended periods to conduct recruitment. This separation creates agency problems, as hiring managers might prioritize short-term outcomes such as perceived cultural fit or immediate affective responses rather than candidates' true long-term potential. Essentially, hiring managers may settle for "good enough" candidates rather than selecting the best, since the downstream costs associated with overlooking top talent are borne primarily by future supervisors, not the hiring managers themselves. Empirical evidence presented here, notably the higher promotion rates among disadvantaged candidates initially undervalued by hiring managers, underscores this misalignment. This finding highlights broader inefficiencies in talent allocation, where subjective, short-term-oriented decisions by hiring agents systematically lead to suboptimal long-term outcomes for organizations (e.g., Gibbons and Roberts, 2013; Sadun et al., 2017; Benson and Shaw, 2024; Hoffman and Stanton, 2024). Addressing these agency issues through better-aligned incentives or structured hiring practices thus emerges as an important avenue for organizational reform and future research.

A further contribution of this paper is its focus on discretionary evaluations—particularly personal interviews—in elite private-sector screening, such as hiring at multinational corporations, an aspect that has received limited attention within economics despite its substantial real-world consequences. While the economics literature has extensively documented discrimination at early screening stages such as résumé evaluations (cf. Kline et al., 2022), there is

vanishingly little empirical work examining how subjective evaluations during personal interviews shape hiring outcomes, especially concerning social class. Recent studies emphasize that informal and subjective screening criteria consistently advantage privileged candidates and correlate poorly with actual job performance (e.g., Sarsons, 2017; Benson et al., 2023; Cappelli, 2019; Hoffman et al., 2018). Crucially, by providing direct evidence from high-stakes recruitment processes at elite multinationals, this paper significantly extends our understanding of how subjective judgments during interviews facilitate class-based discrimination. Furthermore, these findings illuminate a critical limitation of résumé anonymization and other early-stage screening interventions: reducing early-stage bias may simply shift class discrimination to later stages of hiring where subjective assessments dominate.

This study also engages with the broader debate about standardized versus "holistic" evaluations prevalent not only in elite hiring but also in other high-stakes selection processes—such as higher education admissions, faculty hiring and promotions, and NSF grant awards. Typically, such screening processes progress from standardized, objective criteria (e.g., GPA cutoffs and standardized test scores or publication records, which reliably predict academic performance and scholarly potential; see Bol et al. 2014; Kuncel and Hezlett 2007) to increasingly subjective evaluations such as personal statements, essays, recommendation letters, and interviews, all of which have substantially lower predictive validity (Murphy et al., 2009; Borneman et al., 2007; Cooper et al., 2007; Kuncel et al., 2007; Alvero et al., 2021). In my setting, group debates—though inherently subjective—still include standardized elements such as uniform prompts and roughly equal speaking time across candidates, making them relatively more objective compared to personal interviews, where evaluators inevitably rely heavily on subjective judgments about cultural alignment and interpersonal rapport. Recent scholarship highlights that holistic evaluations often disproportionately reflect evaluators' biases rather than candidates' true potential (Chetty et al., 2023; Sacerdote et al., 2025; Bol et al., 2014; Checchi and van de Werfhorst, 2018; Barg, 2013). Moreover, emerging algorithmic and AI-driven screening tools, intended to enhance objectivity, could inadvertently replicate these biases if not rigorously validated (Albright, 2024). By empirically linking discretionary evaluations to long-term organizational inefficiencies—demonstrated by higher subsequent promotion rates among candidates initially undervalued in subjective evaluations—this research underscores the need for greater scrutiny and potential reform of discretionary interviewing practices in elite recruitment contexts.

Finally, this paper also contributes insights into the role of elite educational institutions in facilitating socioeconomic mobility and the persistent challenges posed by discriminatory hiring practices. Elite colleges, such as the Indian institution studied here—characterized by admission rates even more competitive than those of Harvard or MIT (University World

News, 2022)—serve as critical gateways to prestigious career opportunities. Although affirmative action policies substantially improve access to these institutions for disadvantaged groups, recent research highlights that the socioeconomic benefits of elite education disproportionately accrue to students from already privileged backgrounds, thereby perpetuating rather than alleviating class-based inequalities (e.g., Zimmerman, 2019; Barrios-Fernández et al., 2024; Chetty et al., 2023). This study demonstrates empirically that even when disadvantaged students gain entry into highly selective institutions, subsequent elite hiring practices—particularly discretionary, subjective evaluations during interviews—can erode affirmative action gains, systematically disadvantaging these candidates in job placement outcomes (cf. Bhavnani and Lee, 2021; Yenmez and Sönmez, 2022). Given that multinational firms typically recruit across several similarly elite institutions during their hiring tours, the findings underscore broader concerns about institutional complementarities: affirmative action policies alone may be insufficient unless complemented by reforms targeting downstream selection processes. Thus, evaluating elite colleges separately from employer hiring practices obscures crucial institutional complementarities, underscoring the necessity of integrated approaches to effectively enhance socioeconomic mobility—especially given the historical co-evolution of social class with cultural and institutional factors that shape long-term economic outcomes (e.g., Bisin et al., 2023).

2 Caste as Social Status Under Affirmative Action

Jati, Varna, and Modern Caste Classifications. The modern caste system in India traces back to historical occupational hierarchies. At its most granular level, the caste system comprises thousands of distinct, hereditary occupational groups known as *Jatis*. These *Jatis* historically mapped onto specific trades, such as merchants, priests, warriors, and agricultural workers, and collectively fall within broader categories called *Varnas*—Brahmins (priests), Kshatriyas (warriors), Vaishyas (merchants), Shudras (laborers), and those historically marginalized as "untouchables" (Dalits). During the colonial era, the British administration began categorizing these groups explicitly, a practice formalized by India's 1949 Constitution through explicit listings known as "schedules." Thus, "Scheduled Castes" (SCs) and "Scheduled Tribes" (STs) refer to specific disadvantaged communities officially listed ("scheduled") by the Constitution itself, alongside "Other Backward Classes" (OBCs), a broader category subsequently defined to include additional socially and educationally disadvantaged groups (Puri, 2025).

This paper uses two broad categories—advantaged ("upper") castes and disadvantaged ("lower") castes—aligned with official classifications employed by India's affirmative action

policies for educational institutions and government employment. Specifically, by disadvantaged castes, I refer to the three constitutionally recognized categories: Scheduled Castes (SCs), Scheduled Tribes (STs), and Other Backward Classes (OBCs). Although "caste" might initially appear distinct from Western notions of "class," India's government has long recognized the tight interplay between caste, economic marginalization, and social status, explicitly grouping historically disadvantaged communities into unified administrative categories (Puri, 2025).

Historical Development of Caste-based Affirmative Action Policies. Early affirmative action for disadvantaged communities emerged during British colonial rule. The Government of India Act of 1919 introduced self-governing institutions and reserved legislative seats for communities termed "depressed classes," explicitly acknowledging their socioeconomic disadvantage. This categorization evolved in the Government of India Act of 1935, which formally designated these groups as Scheduled Castes. After independence, India's 1949 Constitution institutionalized quota-based affirmative action in legislatures, public higher education, and government employment, initially for Scheduled Castes and Scheduled Tribes, and subsequently expanding to Other Backward Classes following the Mandal Commission's recommendations in the 1980s (Bayly, 2008).

Caste as Proxy for Social Status. Although caste originated as a birth-based hierarchy, Indian policymakers now explicitly treat caste categories as comprehensive measures of socioeconomic "backwardness," integrating economic, educational, and social disadvantage. In doing so, policymakers implicitly acknowledge how caste encapsulates various dimensions of social status—financial resources, educational credentials, social networks, and cultural capital—which collectively shape individuals' socioeconomic opportunities (cf. Bourdieu, 1986). The Mandal Commission systematically identified socially disadvantaged groups, leading to a recommended 27% reservation for OBCs in government employment, public-sector undertakings, and educational institutions. Implemented in 1990, these measures raised the overall reservation for constitutionally recognized disadvantaged groups—SCs, STs, and OBCs—to nearly 50% (Bayly, 2008).

The OBC classification explicitly merges lineage-based identity with low socioeconomic status, demonstrating how caste functions as a consolidated marker of social class in India. This administrative approach underscores that what Western contexts label "lower class" often closely aligns with the category of "backward caste" in India.

Limited Scope of Affirmative Action in the Private Sector. India's affirmative action provisions in the public sector do not extend to private-sector employment—where the elite positions at US and European multinational corporations analyzed in this paper are situated. Because private-sector roles typically offer higher wages, better career growth, and lack reservation mandates, disadvantaged castes remain disproportionately excluded from these opportunities (Subramanian, 2019). This gap mirrors similar patterns in Western contexts, where Diversity, Equity, and Inclusion (DEI) efforts often prioritize race and gender, but class-based diversity remains overlooked. Indeed, fewer than 10% of large US firms explicitly track socioeconomic diversity among employees (Stansbury and Rodriguez, 2024), underscoring how class-based disparities, both in India and globally, remain systematically underexamined. Even within India, as recently as 2018, fewer than 3% of the top 100 firms listed on India's premier stock exchange maintained caste data for internal HR purposes (BusinessLine, 2018), reflecting widespread institutional reluctance to acknowledge caste-based disparities explicitly. Thus, even robust affirmative action policies within public education and government employment fail to significantly mitigate deeper structural inequalities in private-sector hiring.

Taken together, India's affirmative action policies underscore that caste-based disadvantage is effectively analogous to socioeconomic "backwardness." These policies clearly recognize how caste reinforces economic barriers and social exclusion. Because caste uniquely and definitively encodes social and economic status, it offers a clarity of measurement that Western conceptions of class frequently lack. Focusing on caste thus captures a historically entrenched, administratively recognized, and practically meaningful measure of social status—making it a uniquely effective tool to analyze class-based inequalities.

3 Key Definitions and Institutional Setting

This section defines key terms, describes the elite college's job fair studied in this paper, and clarifies why this job fair is representative of how elite Indian college graduates enter elite jobs.

3.1 Key Terms: Disadvantaged Castes, Elite Jobs, and Elite Colleges

Below, I define four key terms used throughout the analysis:

a. **Disadvantaged castes.** Following the Indian government's official classifications, the term "disadvantaged castes" encompasses individuals belonging to Scheduled Castes

(SCs), Scheduled Tribes (STs), or Other Backward Classes (OBCs).

b. **Job.** A "job" refers specifically to a job designation within a firm. For example, if Google recruits both a product manager and a software engineer, these constitute two distinct jobs. Section 4.5 explains why a job typically corresponds to a firm-designation pair, rather than a particular office location.

c. **Elite jobs.** An "elite entry-level job" refers to a formal entry-level position at multinational corporations (MNCs), major national corporations, or unicorn startups (privately-held startups valued at over $1 billion). All elite jobs analyzed in this paper pay salaries within the top 1% of entry-level wages in India (Economic Advisory Council to the Prime Minister, 2022). Throughout this paper, the terms "elite entry-level jobs" and "elite jobs" are used interchangeably.

d. **Elite colleges.** "Elite colleges" include institutions consistently ranked among the top 30 in their respective fields—such as science, engineering, arts and humanities, social sciences, and business—by *India Today*, an Indian equivalent of the *US News & World Report*. Prominent examples include the Indian Institutes of Technology (IITs), Indian Institutes of Management (IIMs), National Institutes of Technology (NITs), and select Delhi University colleges such as St. Stephen's, Shri Ram College of Commerce, and Lady Shri Ram College.

3.2 Structure of the On-Campus Job Fair

The elite college's job fair occurs *entirely on campus* and comprises three phases: (1) pre-placement, (2) placement, and (3) post-placement (or the "aftermarket").

a. **Pre-Placement Phase.** The placement office invites firms before June. From June to mid-August, recruiters visit campus to deliver pre-placement talks, gauge student interest, and finalize roles. By late August, firms extend pre-placement offers (PPOs) to students who interned with them during the summer. Students accepting PPOs are *de-registered* from the main job placement process.

Significance of return offers from internships. Unlike in the US, return offers from junior-year internships play only a minor role at elite Indian colleges. Less than 5% of the graduating cohort accepts PPOs, leaving more than 95% to participate in the main job fair. Fewer than 2% choose public-sector jobs (Section 4.5).

Student and employer registration. Students register for the placement process by late August. Employers submit detailed "registration forms" by early September, listing job positions, compensation, and vacancies for the year. Once submitted, these job profiles become fixed.[3] Students cannot negotiate compensation; the placement office enforces advertised pay through verification of offer letters and can restrict future campus visits by firms not complying.

b. **Placement Phase.** Beginning in mid-September, students apply for jobs. Employers initially screen applications, then conduct written aptitude tests and group debates assessing socio-emotional skills. These pre-interview rounds occur from September to early December. From December to January, employers conduct final-round personal interviews and make job offers. By January, approximately 70% of the cohort secures full-time employment.

c. **Post-Placement Phase ("Aftermarket").** From late January to graduation in late July, the placement office operates an "aftermarket" for students who remain without jobs. Roughly 30% of the sample enters this phase.[4] About 80% of these aftermarket participants (25% of the full sample) eventually pursue entrepreneurship, graduate studies (Master's or PhD), or gap years preparing for civil service or elite MBA exams. The remaining 20% (around 5% of the full sample) find employment elsewhere.

Firms in the aftermarket. Firms hiring during the aftermarket differ significantly from those in the regular cycle. Typically offering lower salaries and less standardized recruitment processes, these firms are often smaller or newer startups. Because my analysis specifically examines the standardized hiring processes used during the regular placement cycle, I exclude the aftermarket. The implications of this exclusion are discussed further in Section 5.2.

3.3 Why the Studied Job Fair is Representative

The college's job fair provides a representative window into how elite Indian college graduates transition into elite jobs, based on four key factors (further details and evidence are provided in Online Appendix Section A.1):

[3] Such rules and timelines are common across elite Indian colleges and are publicly available on most of their placement websites. For instance, see Central Placement Cell (Delhi University).

[4] Many students in the aftermarket engage passively, often targeting selective firms or niche startups, having already decided on alternative paths (*e.g.*, graduate studies or civil service preparation).

First, nearly all elite Indian colleges—defined as institutions consistently ranked among the top 30 in their fields (Section 3.1)—conduct similar on-campus job fairs with standardized procedures. Thus, the rules, protocols, and interactions between students and employers at the studied college closely match those at peer elite institutions.

Second, a common set of elite employers—multinational corporations, prominent national firms, and unicorn startups—consistently recruit from the same group of top-tier Indian colleges. These firms typically visit 30–40 colleges annually, each enrolling approximately 1,500 students, collectively influencing around 50,000 students each year. Because these elite employers draw from the same institutional network, their recruitment strategies and job opportunities remain highly consistent. Moreover, because all analyzed positions pay within the top 1% of entry-level Indian salaries (Section 3.1), the studied college's outcomes accurately reflect broader trends in India's elite labor market.

Third, elite entry-level positions almost exclusively recruit graduates from these top-ranked colleges, and graduates from these colleges predominantly join these elite positions. This mutual exclusivity—documented extensively in prior research (Subramanian, 2019; Jodhka and Naudet, 2019)—creates a highly integrated elite labor market ecosystem. Consequently, employment patterns observed at the studied college directly mirror broader national patterns of elite job transitions.

Fourth, very few elite college graduates opt out of the on-campus placement process. In the studied setting, fewer than 5% of students forego participation in formal job placements (Section 3.2). Comparable participation rates at peer institutions confirm the representativeness of the studied fair (Subramanian, 2019). Additionally, most students heavily rely on their college's placement office rather than pursuing off-campus opportunities, further reinforcing the fair's representativeness.

4 Detailed Data on Students, Jobs, and Recruitment Process

This section describes the data used in the analysis, detailing student and job characteristics, recruitment procedures, and longer-term career outcomes such as promotions and tenure. It concludes by briefly discussing relevant data limitations.

4.1 Institutional Source and Completeness of Recruitment Data

Data were obtained directly from the university's placement office, which meticulously records each stage of the recruitment process, including student applications, screening

rounds, personal interviews, and final job offers. The university qualifies as an "elite college" as defined in Section 3.2. It is one of India's highest-ranked institutions, and its recruitment practices serve as a benchmark for placement processes at comparable elite colleges. The identity of the university is anonymized to preserve confidentiality.

4.2 Standardized Recruitment Rules Ensuring Data Integrity

All recruitment fairs at elite Indian colleges adhere to a consistent set of rules applicable to both students and recruiting firms, ensuring standardized procedures and comprehensive data collection.

Mandatory participation in all screening rounds. Every job analyzed in this study conducts all four screening rounds prior to extending offers (see Section 4.7). The placement office strictly enforces compliance: students applying to any job must attend every screening round for which they qualify, while firms are prohibited from making informal or early offers. These widely adopted rules, common across elite Indian institutions, maintain data integrity and completeness. Although exploring the rationale behind these rules is beyond this paper's scope, detailed policies are publicly accessible on most elite colleges' placement websites.[5]

4.3 Student Sample: Time Span, Degrees Offered, and Caste-Quota Admissions

The dataset spans four placement cycles (years), all occurring prior to the COVID-19 pandemic. Exact years are omitted to protect anonymity further.

Balanced caste representation via affirmative action in admissions. Due to India's caste-based affirmative action in public college admissions (discussed previously in Section 2), caste representation within each academic major at the college is approximately balanced. Table 1 summarizes the distribution of the 4,164 students in the dataset, about evenly split between advantaged and disadvantaged castes.

Administrative reporting of student caste. Students self-report their caste to the college by submitting an official "caste certificate," which is administratively recorded for enforcing quota-based admissions. Consequently, this administrative data collection allows the college—and by extension, researchers—to accurately track students' caste categories.

[5]For example, see Central Placement Cell (Delhi University).

Predominantly male student population. About 90% of the sample is male, reflecting the typical demographic composition of elite technical colleges in India (Datta, 2017).

Degrees offered. The dataset includes four degrees (see Table 1): Bachelor of Technology (B.Tech.), Dual Degree, Master of Technology (M.Tech.), and Master of Science (M.S.). The Dual Degree program integrates undergraduate and postgraduate coursework, requiring an additional year beyond the standard four-year B.Tech. curriculum.

Admissions based solely on standardized national entrance exams. Admissions to this elite college occur exclusively through performance on standardized national-level entrance examinations, based on caste- and major-specific cutoff scores.

Table 1: **Total Students by Caste and Degree**

	Total Students		
	Advantaged Caste	Disadvantaged Caste	Degree Total
B.Tech.	579	710	1289
Dual Degree	622	617	1239
M.Tech.	601	566	1167
M.S.	344	125	469
Caste Total	2146	2018	4164
Fraction	0.52	0.48	1.00

Notes: Table 1 shows the total number of students by caste within each college degree, the total number of students by caste, and the fraction of students by caste. B.Tech. stands for Bachelor of Technology, M.Tech. stands for Master of Technology, and M.S. stands for Master of Science. A Dual degree integrates undergraduate and post-graduate studies and is completed a year after the conventional four-year degree (B.Tech.).

4.4 Student Characteristics and Caste-Based Differences

This subsection summarizes three primary student characteristics—college GPA, pre-college skills, and previous labor market experience—and highlights two key empirical patterns that emerge from these data.

College GPA. The dataset contains each student's college GPA. There are substantial caste-based differences in GPA within each degree program (Table 2). Crucially, both advantaged and disadvantaged caste students appear within each GPA decile, ensuring substantial overlap across caste groups (Online Appendix Figure OA.1).

Pre-college skills. Pre-college skills are captured by entrance exam (EE) scores, the sole admissions criterion (Section 4.3). EE scores were originally recorded as ranks but have been

renormalized such that higher scores indicate better performance. Caste differences in EE scores are pronounced (Table 2), though students from both castes are represented across each EE score decile (not shown). This representation is unsurprising, given that both castes appear within every college GPA decile.

Previous labor market experience. The data contain detailed records on students' prior labor market experiences, including summer and winter internships. These records cover internship duration, total pay, employment sector, and whether the internships were with startups. Internship descriptions typically specify eligibility criteria based on factors such as major, degree, and GPA, along with desired skills.

Two clear empirical patterns emerge from these student characteristics:

1. **Substantial caste differences in academic performance metrics.** Large caste disparities are evident in both college GPA and entrance exam scores, although students from both castes span the full range of performance levels (Table 2; Online Appendix Figure OA.1).

2. **Weak correlation between caste and previous labor market experience.** Caste is only modestly correlated with prior labor market experience. For example, among B.Tech. and Dual Degree students, the distribution of internship sectors is similar across castes: approximately 25% interned in technology, 37% in consulting, and 38% in manufacturing. Similar patterns hold for M.Tech. and M.S. students (not shown).

   This weak correlation likely reflects the relatively minor role played by summer internship return offers at elite Indian colleges (Section 3.2). Students tend to prioritize securing internships in preferred sectors or firms rather than focusing primarily on internship pay. Consequently, only modest caste differences emerge naturally from these dynamics.

Table 2: **Summary Statistics (Students)**

	Adv. Caste	Disadv. Caste	Diff (SD)
		B.Tech.	
Avg. Entrance Exam Score	0.41	-0.33	0.74***
Avg. 10th Grade Score	0.24	-0.19	0.43***
Avg. 12th Grade Score	0.24	-0.20	0.44***
Avg. Overall College GPA	0.51	-0.42	0.93***
		Dual Degree	
Avg. Entrance Exam Score	0.35	-0.35	0.70***
Avg. 10th Grade Score	0.20	-0.20	0.40***
Avg. 12th Grade Score	0.14	-0.14	0.28***
Avg. Overall College GPA	0.43	-0.43	0.86***
		M.Tech.	
Avg. Entrance Exam Score	0.27	-0.30	0.57***
Avg. 10th Grade Score	0.23	-0.15	0.38***
Avg. 12th Grade Score	0.17	-0.10	0.27***
Avg. Overall College GPA	0.34	-0.33	0.67***
		M.S.	
Avg. Entrance Exam Score	-0.02	0.02	-0.04
Avg. 10th Grade Score	0.03	-0.07	0.10
Avg. 12th Grade Score	0.02	-0.06	0.08
Avg. Overall College GPA	0.04	-0.14	0.18*

Notes: Table 2 documents some select summary statistics for students on measures of pre-college skills and college academic performance. Pre-college skill measures include college entrance exam scores and scores in the 10th and 12th standard national-level exit exams. College academic performance is measured by college GPA. All scores are pooled and normalized to have zero mean and unit standard deviation within each degree and year. College entrance exam scores are originally ranks that have been renormalized so that higher numbers are better. The difference across castes is reported in standard deviation units. $^{*}p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$. Adv. Caste stands for advantaged caste and Disadv. Caste stands for disadvantaged caste.

4.5 Characteristics of Multinational Private-Sector Jobs

This subsection summarizes key job characteristics, including sector distribution, salary structure, job location, recruitment processes, and managerial involvement.

Sector and employer overview. The dataset covers approximately 1,050 jobs, nearly all of which (about 98%) are in the private sector. Most employers are major multinational corporations, typically headquartered in the US or Europe, and frequently listed in the Fortune 500. These companies recruit for positions located predominantly in their Indian offices (Table 3). Over the four recruitment cycles (Section 4.3), job availability varies: some positions are advertised only once, while others recur annually, with each job appearing on average in two of the four cycles. Approximately 75% of roles are non-client-facing (e.g., software engineering), with the remaining 25% being client-facing roles (e.g., consulting and managerial positions).

Exclusion of public-sector jobs. I exclude public-sector jobs from the analysis for two reasons. First, they represent fewer than 2% of the sample (Table 3). Second, public-sector positions in India differ substantially from private-sector roles, particularly regarding salary structure and job stability.[6]

Standardized salary posting and national wage setting. Job salaries are *standardized, non-negotiable, and job-specific* rather than candidate-specific or match-specific (Section 3.2). Conditional on a candidate's college degree, salaries *do not vary* by major, caste, gender, or other individual factors. Firms typically adopt nationwide salary standards to simplify administration, as they recruit from multiple elite colleges for identical job profiles. This "national wage setting" practice has also been documented in elite entry-level hiring in the US (Bloom et al., 2012; Sarsons et al., 2022).[7]

Detailed job descriptions and standardized non-wage benefits. Before recruitment begins, employers submit comprehensive "registration forms," clearly specifying job responsibilities, required skills, and standardized non-pecuniary benefits—such as health insurance, housing assistance, travel allowances, and signing bonuses. Just like salaries, these non-wage benefits are uniform across all candidates applying for a specific job.

[6]For further details on public-sector pay and stability, see Seventh Central Pay Commission, 2016.

[7]Further details on national wage setting in this context are available upon request.

Job location and recruitment procedures. More than 98% of the jobs analyzed are based in India. Typically, employers conduct unified screening processes for positions across all Indian offices. Assignment to a specific Indian office occurs post-hiring, depending on the employer's location-specific staffing needs. Thus, a "job" generally refers to a firm-designation combination rather than a particular Indian office location. For instance, Google recruits software engineers through a unified screening process, irrespective of the eventual assignment to Mumbai or Bangalore. Salaries and benefits for a given position are also standardized across all Indian locations.

Less than 2% of jobs in the sample involve placements in offices outside India, and these are treated separately. Positions outside India (e.g., Google's Palo Alto office) undergo distinct recruitment processes managed by specialized teams (still primarily comprising Indian recruiters), offering distinct salary packages. Therefore, these non-Indian location roles constitute separate job definitions.

Separation of recruitment teams from direct supervisors. Recruitment, including final interviews, is conducted by specialized hiring managers who often differ from the direct supervisors eventually overseeing hired candidates. Because recruitment occurs entirely on campus at each of the 30–40 elite colleges spread across diverse geographic regions, hiring managers must physically travel over several months (Sections 3.2 and 3.3), making direct supervisor involvement impractical. Informal conversations with employers confirm that dedicated recruitment professionals handle the entire hiring process independently from eventual supervisors, who manage the new hires day-to-day.

Predominantly Indian hiring managers. Because the overwhelming majority of job positions (over 98%) are based in India, recruitment predominantly involves Indian hiring professionals or managers.

Summer internships and return offer policies. Students who accept summer internship return offers by late August are disallowed ("de-registered") from the main placement cycle (Section 3.2). Thus, students applying during the formal placement season starting in September would not have completed junior-year summer internships with the same employer.

Table 3: **Summary Statistics (Jobs)**

	Count	Fraction	Avg. Salary ($)	Std. Dev. ($)
Private	1011	0.98	60341.38	32591.51
Government	19	0.02	56202.23	20240.01
Headquartered Outside India	620	0.60	64562.74	36038.04
Headquartered in India	410	0.40	53766.04	23946.16
Client Facing	246	0.24	57553.23	28288.23
Non-Client Facing	784	0.76	61115.92	33361.87
All Jobs	1030	1.00	60265.02	32391.89

Notes: Table 3 shows the proportion of jobs in the public versus private sector, the proportion of jobs headquartered outside versus in India, and the average salary breakdown (USD PPP) among client-facing jobs and non-client-facing jobs. Firms headquartered outside India are mostly based in the US or Europe.

4.6 Quantifying Job Selectivity at Each Recruitment Stage

This subsection summarizes the selectivity of jobs and defines two critical concepts—conditional and cumulative cuts—used to measure selectivity at different recruitment stages.

Conditional and cumulative cuts defined. Jobs exhibit substantial variation in selectivity across each screening round. Table 4 illustrates this by presenting the distribution of conditional and cumulative cuts. A *conditional cut* refers to the fraction of applicants eliminated *during a specific round*, while a *cumulative cut* measures the total proportion eliminated *by the conclusion of a given round*. For instance, if a job eliminates 50% of applicants in each of the first two rounds, its conditional cut at each round is 50%, but the cumulative cut after round two totals 75

Overall selectivity is high. Successful candidates typically must rank within approximately the top 5% of the initial applicant pool. Moreover, it is not uncommon for some jobs to extend no offers at all, reflecting hiring processes that occasionally fail for various reasons.

Table 4: **Variation in Conditional and Cumulative Selectivity Across Jobs**

Cut Type	Round	Q1	Mean	Q2	Q3
Conditional Cut	App. Reading	0.30	0.39	0.39	0.48
Cumulative Cut	App. Reading	0.30	0.39	0.39	0.48
Conditional Cut	Written Test	0.10	0.18	0.15	0.26
Cumulative Cut	Written Test	0.38	0.49	0.49	0.61
Conditional Cut	Group Debate	0.69	0.74	0.74	0.79
Cumulative Cut	Group Debate	0.83	0.87	0.87	0.91
Conditional Cut	Personal Interview	0.61	0.75	0.78	0.90
Cumulative Cut	Personal Interview	0.95	0.96	0.97	0.99

Notes: Table 4 shows the distribution of both the conditional and cumulative cuts across jobs.

4.7 Four Recruitment Stages: From Objective to Subjective Evaluations

This subsection describes the four screening rounds uniformly employed by all jobs in the sample. The recruitment process generally transitions from standardized and objective assessments toward increasingly subjective evaluations of candidate qualities and interpersonal skills.

Low-cost job applications via centralized portal. Students apply through a centralized online portal, analogous to Job Openings for Economists, by simply selecting desired job listings without additional cover letters or personal statements. Employers automatically receive applicants' standardized resumes through this system. This streamlined application process significantly lowers marginal application costs, enabling comparable application rates across caste groups represented about equally within each major (Section 4.3).

Eligibility criteria restrict applications. Application eligibility is determined by specific combinations of major, degree, and GPA. For instance, a software engineering position at Microsoft typically restricts eligibility to students majoring in Computer Science, while an analyst role at Goldman Sachs might allow students from all majors to apply. Despite low

application costs, these eligibility criteria practically limit the number of positions a student can apply for each year.

Four mandatory screening rounds. All jobs consistently follow a structured sequence of four screening rounds from early September to January: (1) application screening (early to mid-September), (2) written aptitude tests (late September to October), (3) group debates (October to November), and (4) personal interviews (December to January). The recruitment process generally transitions from standardized and objective assessments toward increasingly subjective evaluations.

1. **Application screening.** Employers initially evaluate candidates by reviewing standardized applications and resumes, assessing documented qualifications, credentials, and basic suitability for the position. Based on conversations with employers, this initial stage frequently involves applying a minimum GPA cutoff as a quick, objective filtering mechanism to efficiently narrow down candidate pools—analogous to standardized test score cutoffs employed by admissions committees. While some subjective judgment of resumes inevitably occurs, the primary focus remains on documented and quantifiable academic credentials and experiences.

2. **Written aptitude tests.** Candidates advancing past application screening participate in online written aptitude tests. These tests objectively measure technical and analytical abilities through standardized questions, providing clearly quantifiable and directly comparable performance metrics across candidates.

3. **Group debates.** Candidates who pass the written tests then engage in group debates. Groups of 30–35 students are divided into two teams of roughly equal size (15–20 participants per team) and assigned to debate opposing sides of a common, non-technical prompt. Although group debates inherently involve subjective evaluations—such as assessments of communication, collegiality, confidence, mannerisms, and consensus-building skills—they also include standardized elements: identical prompts, predefined debating positions, structured team compositions, and roughly equal speaking opportunities for all participants. These standardized features help maintain consistency and comparability across candidate evaluations.

4. **Personal interviews.** The final screening stage consists of personal interviews, typically lasting about 30 minutes per candidate. Interviews assess candidate "fit," interpersonal rapport, cultural alignment, and other intangible qualities, alongside direct

evaluations of job-specific competencies. This round relies most heavily on evaluators' subjective judgments.

Thus, the recruitment process naturally evolves from initial objective filters and standardized assessments to nuanced assessments of interpersonal skills and cultural compatibility.

Common screening rounds across job profiles. Firms commonly consolidate screening rounds across multiple advertised positions. For instance, if Nvidia offers five distinct roles in a given recruitment cycle, it may administer one common technical test simultaneously for all applicants. Further details on each screening round are provided in Online Appendix Section A.2.

4.8 Candidate Screening: Summary Statistics

This subsection summarizes key statistics on candidate screening, highlighting the typical volume of applicants and the number of students advancing through each screening round.

Table 5 reports the average (median) number of applications per job as 208.85 (140.17). The average (median) number of candidates advancing per job at each subsequent screening round are: written aptitude tests, 124.91 (87.00); group debates, 101.54 (70.00); personal interviews, 26.07 (18.00); and job offers extended, 6.49 (4.00). The substantial decline at each stage reflects the rigorous selectivity of the hiring process (Table 4). Additionally, because jobs typically restrict applications based on major, degree, and GPA criteria (Section 4.7), these requirements practically limit the number of positions for which any given student can apply.

Table 5: **Summary Statistics (Candidate Screening)**

Round	Mean	Median
Application Reading	208.85	140.17
Written Tests	124.91	87.00
Group Debates	101.54	70.00
Personal Interviews	26.07	18.00
Offers	6.49	4.00

Notes: Table 5 shows the mean and median number of applications, written tests conducted, group debates conducted, interviews conducted, and offers extended per job. For a given screening round, these numbers can be calculated by averaging, per job, the number of students across years and then averaging across jobs.

4.9 Data on Job Tenure and Promotions

I measure two indicators of job performance—tenure and promotion rates—by linking the sample to an employment-focused online platform.[8] I successfully match over 98% of students hired through the college job fair by cross-referencing their names with unique combinations of internship histories, initial job placements, college majors, degrees, and graduation years. Tenure, measured in months, is the total duration an employee spends at a given firm, computed from the start and end dates of each role reported on the platform.

Defining and classifying promotions. Promotions are defined as movements to positions with higher job titles within the same firm. To ensure robust classification, promotions must meet two criteria. First, I verify salary increases associated with new job titles using benchmarks from industry-standard platforms (platform names anonymized for confidentiality). Second, I confirm that the job title progression aligns with typical career advancement trajectories for the given role. Such classifications are particularly reliable in elite, entry-level jobs at multinational corporations, where career ladders and salary structures are highly standardized and globally consistent (Bloom et al., 2012; Minni, 2024). For example, moving from "Analyst" to "Associate" within two years at a major consulting firm constitutes a promotion based on both salary increase and hierarchical advancement.

Distinguishing vertical career moves. This approach clearly identifies genuine vertical advancements rather than lateral job changes. Consequently, these metrics—combined with the structured compensation practices characteristic of elite multinational corporations—minimize subjectivity, providing a robust measure of employee performance and career progression.

4.10 Data Limitations and Mitigation

There are two primary limitations. First, I lack direct access to student resumes. Nevertheless, the placement office mandates a standardized resume format, ensuring most relevant candidate information is already reflected in my dataset (e.g., entrance exam score, major, GPA, internships, and job histories). Second, I do not observe precise performance scores or evaluator ratings from any screening round. While these limitations restrict direct visibility into intermediate evaluative judgments, I discuss strategies for mitigating their impact and interpreting the results appropriately in Section 5.7.

[8] The platform's name is anonymized to maintain confidentiality.

5 Penalizing Social Status Through "Fit" Assessments

This paper documents five key findings on how personal interviews reveal and penalize social status (caste). First, a substantial—and likely conservative—earnings gap exists between castes. Second, about 90% of hiring disparities emerge only after subjective "fit"-focused personal interviews. Third, job-level caste penalties consistently arise during these interviews across nearly all positions. Fourth, employers penalize caste explicitly once it becomes evident, positioning personal interviews as the primary channel for status-based discrimination. Finally, disadvantaged caste hires demonstrate higher promotion rates even in the most selective roles, suggesting discrimination—rather than unobserved skill deficiencies—as the underlying mechanism for hiring inequalities.

5.1 Main Sample and Employment Accuracy

As noted in Section 3.2, about 70% of students secure full-time employment through the regular placement process by early January, with fewer than 5% opting out via summer internship return offers. Consequently, over 95% of the cohort actively participates in the standardized, four-stage on-campus job fair, resulting in approximately 2,550 hires (65% of the entire sample).

Reliability of initial job placements. Virtually all hires begin working in the roles obtained during placement, validating the placement office's data as a reliable measure of actual post-graduation employment.[9]

5.2 Substantial Earnings Gap Across Castes

Disadvantaged castes earn significantly lower initial salaries than their advantaged peers. The unconditional earnings gap is 23%, which narrows to about 10% after controlling for pre-college skills, college academic performance, prior labor market experience, and other relevant factors. This residual 10% gap—approximately $6,000$ (PPP) of mean earnings, or roughly 0.18 standard deviations of the private-sector job salary distribution (Table 3)—arises from differences in job composition, not from differential pay within identical roles, given the standardized compensation at the job fair (Section 4.5). Online Appendix Table OA.1 provides additional detail on modest variations by job type.

[9] A mandatory post-graduation exit survey (98% response rate), conducted two months after graduation, confirms that nearly all hires began in their offered positions without compensation renegotiation. Additional cross-validation with data from a prominent employment platform further supports the accuracy of placement records.

Students excluded from main earnings analysis. Two groups are excluded from the main earnings analysis: (1) students who accept summer internship return offers, and (2) those who remain unplaced after the regular cycle and subsequently participate in the "aftermarket" (Section 3.2).

Implications of excluding these students. No student who participates fully rejects all job offers. However, as noted previously, roughly 30% of the full cohort are excluded from the main sample due to either early acceptance of internships or entry into the aftermarket. Among these omitted students, disadvantaged castes are negatively selected on GPA relative to advantaged peers (Online Appendix Table OA.2).[10] Given that higher GPA generally correlates with higher pay, omitting disproportionately lower-GPA disadvantaged students likely understates the true wage gap.

5.3 Caste Earnings Penalty Arises Predominantly During Personal Interviews

Sections 5.3–5.7 highlight the paper's central contribution: demonstrating that personal interviews—emphasizing subjective "fit" criteria such as communication style, personality, and cultural alignment—are the primary stage where caste discrimination materializes in elite hiring. These interviews alone account for nearly 90% of conditional hiring disparities.

Because each job offers a *fixed, non-negotiable salary* (Section 4.5), differences in the caste earnings penalty across screening stages reflect differences in job composition—i.e., the types of jobs each caste has access to—rather than negotiation or candidate-specific pay variations. This structure allows me to precisely quantify the earnings penalty at each stage by counterfactually assigning the posted salaries of available jobs to each caste group before final offers are accepted. Formally, I estimate variations of the following regression:

$$\log(\text{Avg. Job Salary}_{it}^{\text{Search Stage}}) = \alpha + \beta \times \text{Disadv. Caste}_i + \text{Controls}_i + \gamma_t + \epsilon_{it}, \quad (1)$$

where $i \in \mathcal{I}$ indexes students. "Search Stage" $\in$ {Application, Written Tests, Group Debates, Personal Interviews, Offers, Accepted Offers}. Controls include college GPA, entrance exam scores, prior internships, major, degree; k_j are job fixed effects and γ_t cohort-year fixed effects. For each student, the "average job salary" at a given stage is calculated as the mean salary of all roles for which the student remains eligible at that stage of screening.

[10] *Within* any given job, however, disadvantaged hires tend to have slightly higher GPAs than their advantaged counterparts (Online Appendix Table OA.3).

The coefficient β, illustrated in Figure 1, quantifies the caste earnings gap at each recruitment stage. Full regression results with robustness checks using Oster (2019) bounds, addressing potential omitted-variable biases, are presented in Online Appendix Table OA.4.[11]

Figure 1: **The Caste Earnings Penalty**

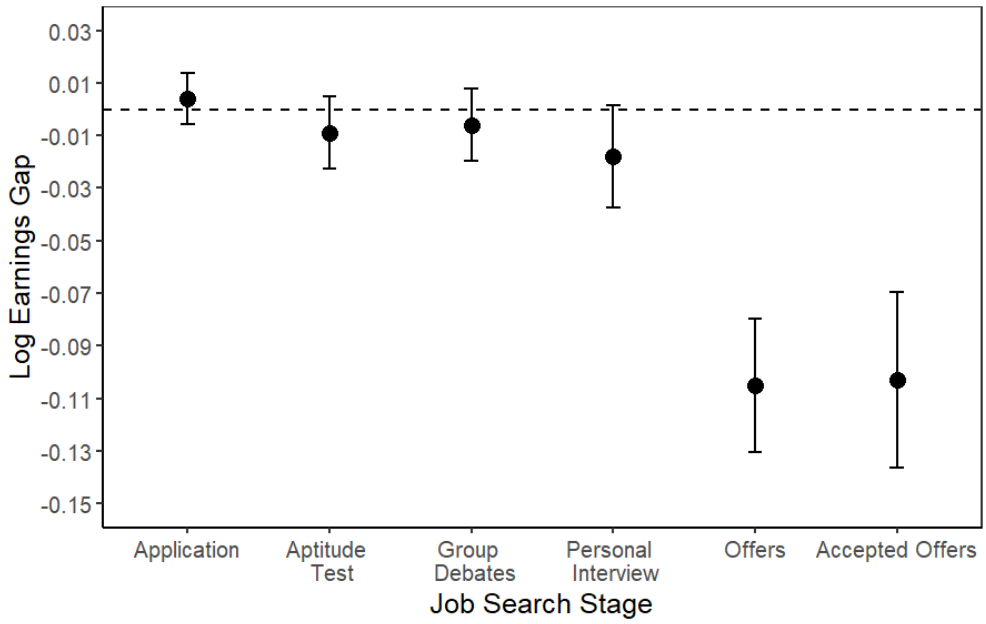

Notes: Figure 1 shows the log earnings gap across castes at successive stages of job search. Each coefficient in the figure is represented by a black dot and reports the percentage difference in the average salary at each job search stage between disadvantaged and advantaged castes. The vertical bars are 95% confidence intervals. These regressions include controls. There are no caste differences in pay for identical jobs and since offered salaries are non-negotiable during the job fair, the caste penalty in mean earnings at each stage can be calculated by comparing the salary differences in the jobs the caste groups are left in contention for after each stage.

Robustness and job-level heterogeneity. While the preceding analysis identifies clear overall trends, potential concerns about aggregate confounding factors remain. To comprehensively address these concerns, Section 5.7 leverages heterogeneity in *job selectivity*. Specifically, I plot *job-level* caste penalties against the cumulative fraction of applicants eliminated by each job at specific screening rounds (defined as "cumulative selectivity" in Section 4.6). This approach reveals that disadvantaged castes experience no early-stage penalties (application screening, written tests, and group debates), even for the most selective jobs, effectively ruling out explanations involving a right-tail skill deficit prior to the interview stage.

[11] We apply Oster (2019)'s bounding methodology, setting $\delta = 1$ (assuming unobservables match observables in selection intensity) and $R^* = 1.3 \times \tilde{R}$, allowing the total explanatory power to exceed the observed R^2 by 30%. The main findings remain robust under these stringent assumptions.

In the next subsection, I show that these late-stage penalties also arise in success rates, reinforcing the primacy of personal interviews in driving caste disparities.

5.4 Caste Penalty in Hiring Success Rates Emerges Primarily at the Interview Stage

This subsection uses success rates—defined as the probability a candidate advances through each screening round for a given job—as an alternative outcome measure, complementing the earnings analysis in Section 5.3. Because salaries are standardized and non-negotiable, caste disparities in earnings should reflect differences in the composition of accessible jobs rather than within-job success rates. Thus, within-job success rates at each screening round should not differ by caste prior to the interview stage if both groups face comparable job opportunities at subsequent stages. This analysis tests and confirms this logic.

Defining and measuring success rates. A student's "success rate" at a screening round is the probability of progressing to the next stage for a given job in a particular year. Using linear probability models with job fixed effects, I estimate caste differences in these success rates at each stage:

$$\Pr(\text{Success}_{ijt}^{\text{Job Search Stage}}) = \alpha + \beta \times \text{Disadv. Caste}_i + \text{Controls}_i + k_j + \gamma_t + \epsilon_{ijt}, \qquad (2)$$

where i indexes students, j indexes jobs, and the screening stage includes application screening, written tests, group debates, personal interviews, job offers, and accepted offers. Controls include college GPA, entrance exam scores, prior internships, major, degree; k_j are job fixed effects and γ_t cohort-year fixed effects. Regressions include year and job fixed effects to account for within-job heterogeneity.

Results: Caste penalties specifically emerge during interviews. Table 6 shows that caste differences in success rates are negligible at the application, written-test, and group-debate stages. Strikingly, these differences only become significant at the personal interview stage. Specifically, over 95% of the final caste penalty in receiving job offers occurs during personal interviews. This result robustly holds across client-facing and non-client-facing roles (see Online Appendix Tables OA.5–OA.6).

Robustness checks. One potential concern is that caste disparities might be hidden at earlier rounds if disadvantaged caste candidates were disproportionately eliminated only

in highly selective positions. Section 5.7 addresses this concern directly, confirming that even in the most selective roles—where early-stage cutoffs are strictest—no caste disparities arise prior to personal interviews. Thus, these results further underscore the unique role of personal interviews in creating caste disparities, rather than merely reflecting pre-interview differences in candidate quality.

Table 6: **Caste Differences in the Success Rate at Each Job Search Stage**

	Success App (1)	Success Test (2)	Success Debate (3)	Success Interview (4)
Disadv. Caste (OLS)	-0.0025*	-0.0002	-0.0005	-0.2938***
	(0.0014)	(0.0018)	(0.0027)	(0.0048)
Disadv. Caste (Oster)	0.009	0.010***	-0.002	-0.296***
	(0.0013)	(0.002)	(0.003)	(0.005)
Mean Success Rate	0.5990	0.8110	0.2580	0.2610
Observations	305,541	183,150	148,489	38,378
Adjusted R^2	0.61069	0.34360	0.01858	0.24450
Year FE	✓	✓	✓	✓
Job FE	✓	✓	✓	✓
Other Controls	✓	✓	✓	✓

Notes: Table 6 shows the caste difference in the success rate at each stage of job search. Each observation is a student who remains in contention at a given screening round in a given job in a given year. Each column reports a separate linear probability regression and includes controls. These regressions also include job fixed effects. Note that some firms conduct the same screening round (written tests, group debates, interviews) for multiple job positions. However, since these job positions within the same firm represent different jobs, the regressions consider these jobs and the corresponding advance-or-out decisions made by them separately. Bias-adjusted coefficients are calculated using Oster (2019). Following Oster's recommended benchmarks, I set $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$, implying that unobservables are assumed to have the same degree of selection as observables, and the total explanatory power can increase by up to 30% beyond the observed R^2. $^*p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

Having confirmed that the interview stage generates the observed caste penalty, the subsequent analysis will explore *when* and *how* caste signals become salient during these interviews.

5.5 Information Observed by Employers at Each Screening Round

This subsection summarizes the information accessible to employers at each recruitment stage. Employers have continuous access to students' standardized academic records—including college GPA, coursework details, entrance exam scores, and selected additional information—entered directly by students into a centralized application portal. This information complements resumes, which students upload to the same portal.

1. **Application reading.** Employers initially review standardized resumes containing academic credentials, internship experiences, extracurricular activities, and technical skills. They frequently apply GPA cutoffs as an objective criterion to efficiently narrow the applicant pool. Employers also have direct access to additional standardized details from the application portal (e.g., entrance exam scores and coursework records), although resumes remain the primary evaluation tool at this stage.

2. **Written aptitude tests.** Advancement is determined primarily by candidates' performance on standardized, job-specific aptitude tests administered online. While resumes remain accessible, evaluations at this stage rely predominantly on objective test scores.

3. **Group debates.** Employers directly observe candidates during structured group debates, assessing communication, teamwork, and socio-emotional competencies. Although earlier academic records and test results remain available, debate performance is the main criterion at this stage.

4. **Personal interviews.** In final personal interviews, lasting approximately thirty minutes per candidate, employers individually assess applicants. While employers review resumes again—including GPA, entrance exam scores, internships, and skills—this stage predominantly emphasizes subjective evaluations of cultural fit, interpersonal rapport, job-specific competencies, and organizational alignment.

Personal interviews allow employers to revisit earlier signals—such as academic credentials and past experiences—in greater detail, thereby enhancing informational accuracy about candidate quality. Crucially, the individualized and subjective nature of interviews elicits additional indicators of cultural alignment and interpersonal compatibility through personalized discussions about interests, extracurricular activities, and background details. Such in-depth conversations not only amplify previously observed signals but also enable employers to glean richer, subtler insights, potentially clarifying ambiguous indicators of social status like caste. Unlike group debates, where standardized prompts and structured interactions restrict personal disclosures and primarily assess team dynamics, personal interviews explicitly facilitate nuanced, individualized evaluations, making social status markers particularly salient. Section 5.6 provides quantitative and qualitative evidence supporting this interpretation.

5.6 Interviews as the Main Mechanism for Revealing and Penalizing Social Status

The evidence in Sections 5.3 and 5.4 raises a critical question: *When does caste become visible to employers?* This subsection demonstrates quantitatively that subjective, fit-focused personal interviews are the main mechanism through which social status (caste) is revealed and penalized.

5.6.1 Quantitative Evidence

While many surnames are indicative of caste, most do not reliably convey it in urban, elite college settings, where students hail from diverse regions and naming traditions (Jodkha, 2017; Parmar, 2020). Consequently, only about 15% of students in my sample have surnames that clearly identify their caste. In that small subgroup, the penalty emerges earlier, during application reading. For the remaining 85% with non-identifiable surnames, the overall caste penalty at the application stage is negligible. Instead, it primarily arises during personal interviews, where students are potentially screened through conversations on many job-relevant characteristics, including candidate background and cultural fit, that also help reveal caste. Thus, personal interviews—through informal "fit" conversations—are the main mechanism for uncovering and penalizing social status.

Caste earnings penalty emerges earlier when names reveal caste. Figure 2 replicates Figure 1, splitting the sample into students with clearly identifiable caste surnames versus those without. I re-estimate Equation 1 separately for each subgroup. Students with identifiable surnames experience a significant penalty beginning at application screening, whereas those without identifiable surnames see disparities emerge only during personal interviews. Corresponding regression results, including Oster's bias-adjusted coefficients, appear in Online Appendix Tables OA.7 and OA.8.

Figure 2: **The Caste Earnings Penalty by Subsample**

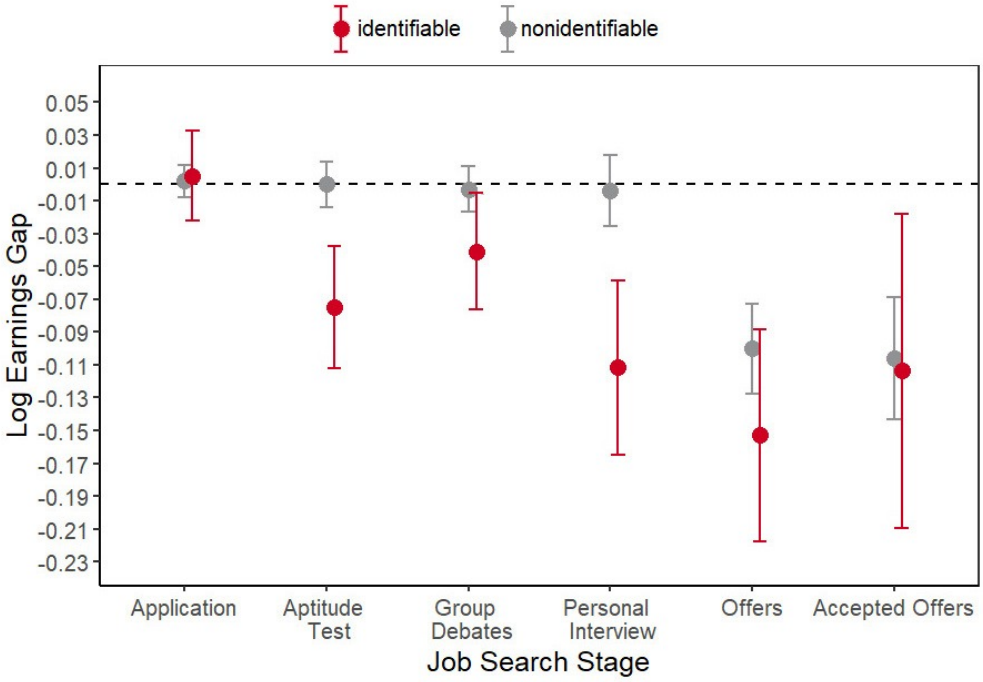

Notes: Figure 2 shows the log earnings gap across castes at successive stages of job search in the identifiable (red) versus the non-identifiable (gray) sample. Each coefficient in the figure is represented by a black dot and reports the percentage difference in the average salary at each job search stage between disadvantaged and advantaged castes. The vertical bars are 95% confidence intervals. These regressions include controls. There are no caste differences in pay for identical jobs and since offered salaries are non-negotiable during the job fair, the caste penalty in mean earnings at each stage can be calculated by comparing the salary differences in the jobs the caste groups are left in contention for after each stage.

Similar patterns in success rates. A comparable pattern appears clearly in success rates. Students with identifiable caste surnames face substantial penalties at the application stage across nearly all jobs, while students without identifiable surnames experience no such early-stage penalty (Online Appendix Figure OA.2).

How are student names associated with caste? To classify surnames as caste-identifiable, I rely on official government sources, including explicitly maintained, periodically updated state-wise "schedules" of community names—hence the terms Scheduled Castes and Tribes (see Section 2). For advantaged castes, I reference surname lists compiled by the Anthropological Survey of India, while for disadvantaged castes (SC/ST/OBC), I draw from Election Commission records of reserved constituencies and official state-maintained SC/ST/OBC lists.

5.6.2 Qualitative evidence supporting caste revelation at interviews

Supporting our quantitative findings, qualitative evidence shows why personal interviews are the first point at which caste becomes legible to recruiters. Group-debate rounds offer little leverage: teams of twenty debate a common prompt with minimal individualized questions, and visual or phonetic cues are weak caste signals. Decades of anthropological work confirm that phenotypic traits such as skin color are unreliable indicators of caste in urban, educated populations (Sebring, 1969; Mishra, 2015; Parameswaran and Cardoza, 2015). A Human Rights Watch study echoes this, noting that "*Lower-caste communities are almost invariably indistinguishable in physical appearance from higher-caste communities*" (Human Rights Watch, 2001, p. 2). Likewise, elite-campus varieties of Indian English encode region and schooling background, not caste; accent features track home state or medium of instruction rather than a student's caste category, and English-medium access has widened well beyond upper-caste groups under affirmative-action expansion (Wiltshire, 2020; Sharma, 2017; Highet, 2023). In short, during debates recruiters confront a large, mixed crowd in which neither looks, surnames, nor English usage reliably betray social origin.

The one-to-one personal interview plausibly shifts the calculus in two ways. First, it amplifies pre-existing cues: recruiters can revisit schooling pedigree, résumé lines, and even debate performance, probing them in greater depth. Second, it adds fresh texture—conversation about background, leisure pursuits, niche extracurriculars, and peer networks—details strongly tied to caste-linked social and cultural capital à la (Bourdieu, 1986)—resources long shown to correlate with caste in elite India (cf. Jodhka and Newman, 2007; Deshpande, 2011; Jodkha, 2017). Triangulating these richer layers with earlier impressions lets recruiters gain nuanced insight into a candidate's social position and infer caste far more accurately, matching the stage at which our quantitative estimates show disparities spike (Section 5.6).

5.7 Heterogeneity in Caste Penalties, Selectivity, and Promotion Outcomes: Empirical Facts & Stylized Model

In Sections 5.3–5.6, we demonstrated that most of the caste penalty in both earnings and success rates emerges at the personal interview stage, when caste becomes salient. Because we lack direct measures of candidate performance at each stage, any unobserved ability differences may remain concealed until later stages (Section 4.10).

This raises two natural questions. First, might these penalties appear earlier in *highly selective* jobs that recruit from the extreme right tail of the ability distribution? Second, do personal interviews reveal unobserved ability that could justify hiring differences across

caste rather than merely reflecting discrimination?

Variation in job-level selectivity provides the test. If disadvantaged castes truly had lower performance scores at pre-interview screening rounds, such deficits would be most evident in highly selective jobs that recruit predominantly from the right tail. Yet, we observe negligible pre-interview caste gaps even in these jobs.

Moreover, once hired, disadvantaged castes exhibit higher promotion rates. One might argue this simply reflects inframarginal productivity differences—that is, more right-tail performers among the hired disadvantaged. However, examining promotion patterns across jobs of increasing selectivity helps distinguish between discrimination and ability differences: while jobs with identical hiring thresholds across castes would generate shrinking promotion advantages as they become more selective, we find these advantages remain persistent and stable across selectivity tiers.

I present two empirical facts that exploit differences in selectivity in across jobs, and then develop a formal model demonstrating why these patterns indicate a stricter hiring threshold for disadvantaged castes. Fact 1 uses job-specific selectivity measures at each screening round and shows no caste gap in the success rate at pre-interview screening rounds across jobs of varying selectivity. Similarly, Fact 2 groups jobs into selectivity bins and shows stable and large promotion rate differences in favor of the hired disadvantaged across the selectivity distribution.

Fact 1 (Conditional Success Rates Reveal No Early Penalty, but a Large Final Interview Gap). Figure 5 plots, job by job, the conditional success-rate gap (disadvantaged-advantaged) against that job's cumulative cut rate at each screening round (defined in Section 4.6). Negative values mean disadvantaged candidates advance less often.

The gap for job j is the coefficient β_j from:

$$\Pr(\text{Success}_{ijt}^{\text{Job Search Stage}}) = \alpha_j + \beta_j \times \text{Disadv. Caste}_i + \text{Controls}_i + \gamma_t + \epsilon_{ijt}, \tag{3}$$

where $i \in \mathcal{I}$ indexes students and γ_t denotes year fixed effects. The controls are characteristics such as college GPA, college entrance exam score, previous internships, major, and degree.

Before personal interviews—across application reading, written tests, and group debates—the figure shows these caste gaps remain near zero, *even* among highly selective jobs that eliminate a large share of applicants. Only at the final interview round does a distinct penalty for disadvantaged castes emerge, underscoring that late-stage "fit" assessments—rather than unobserved performance differences prior to interviews—drive the observed caste disparities.

As a robustness check, Online Appendix Figure OA.3 presents the *uncontrolled* success

rates, which reveal that while disadvantaged castes face an initial penalty at the application reading stage, the penalty nearly disappears prior to personal interviews and then reemerges afterward—precisely where the controlled regressions show the gap intensifying in Fact 1.

Fact 2 (Disadvantaged Hires Earn Comparably Higher Promotion Rates, Even in Highly Selective Jobs). The above patterns do not rule out the possibility that interviews are simply revealing other productivity-related traits in which disadvantaged castes genuinely lag. To distinguish unobserved ability differences from discrimination, we next turn to *post-hire* performance.

Using data from an employment-focused online platform, I track career trajectories for over 98% of the job fair sample and use promotion rates as a measure of job performance.[12] If disadvantaged castes truly had lower ability that was revealed during interviews, the relative performance of the hired disadvantaged should decline in more selective jobs that impose higher—but equal—standards across castes. Instead, as we show below, disadvantaged hires maintain their promotion advantage even in the most selective positions. This pattern aligns with outcome-based tests for discrimination developed in healthcare and algorithmic fairness research (Obermeyer et al., 2019; Chandra and Staiger, 2010, 2020; Abaluck et al., 2016), where higher success rates among marginalized groups, conditional on selection, are suggestive of discriminatory screening thresholds.

To compare caste differences in job performance, I estimate the following linear regression across job selectivity tiers:

$$\text{Outcome}_{ijt} = \alpha + \beta \times \text{Disadv. Caste}_i + \text{Controls}_i + k_j + \gamma_t + \epsilon_{ijt}, \tag{4}$$

where $i \in \mathcal{I}$ indexes students, $j \in \mathcal{J}$ indexes jobs, γ_t denotes year fixed effects, and k_j denotes job fixed effects. The controls are characteristics such as college GPA, college entrance exam score, previous labor market experience, major, and degree.

Table 7 shows the coefficients from these regressions:

- **Columns (1)–(2)** show that, pooling all jobs, disadvantaged hires have about a 20% *higher* probability of promotion relative to advantaged hires, despite having been penalized at the hiring stage.

- **Columns (3)–(8)** group jobs by their *selectivity* (final rejection rates) and reveal that a similar promotion advantage of about 20-25% for disadvantaged hires persists across tiers of increasing stringency—here, the top 75%, 50%, 40%, 30%, 20%, and 15% most selective jobs (e.g., jobs in the top 20% of selectivity reject at least 80% of initial

[12] The unlinked sample comprises just 1.5% of graduates. Section 4.9 details how these links are made.

applicants). In other words, even among candidates who survived highly stringent hiring standards, disadvantaged hires do *not* underperform on promotions.

- **Columns (9)–(10)** demonstrate similar patterns in both client-facing and non–client-facing roles, suggesting that the promotion advantage is not driven by a particular job function.

Promotion-rate estimates are, if anything, conservative. Global promotion grids are nominally standardized (Bloom et al., 2012), but local managers still rate "leadership," "communication," and "potential" subjectively (Benson et al., 2019). Any such bias should *lower*, not raise, disadvantaged candidates' advancement odds, yet we still observe a ∼20% promotion edge. The gap is thus a lower bound.[13] Online Appendix Table OA.9 corroborates this: advantaged–caste leavers are, on average, lower–GPA, whereas promotions correlate (weakly) *positively* with GPA.

Key Takeaway. Facts 1 and 2 challenge the notion that interviews merely catch up to latent skill deficits among disadvantaged castes. A large pre-interview performance deficit should have implied lower success rates for disadvantaged castes before interviews, especially among highly selective jobs. And, if interviews were accurately filtering out less-qualified candidates, we would expect similar or even lower promotion rates among those hired. Instead, disadvantaged hires consistently attain *higher* promotion rates—even in the most selective jobs. The promotion advantage does not shrink or equalize as jobs become more selective, a discrepancy that is difficult to reconcile with explanations based solely on unobserved ability differences.

In what follows, I present a simple theoretical framework to rationalize these two empirical facts, showing how this evidence is consistent with discrimination in hiring—with personal interviews serving as the main mechanism that translates social status into hiring barriers.

5.7.1 Explaining the Empirical Facts: A Simple Model of Late-Stage Discrimination

This section presents a unified model showing how (i) the absence of any early shortfall for disadvantaged castes implies no genuine ability deficit before interviews, and (ii) multiple hiring thresholds across different jobs help distinguish unobserved ability differences revealed at the interview from discrimination.

[13] A 2018 survey of South Asians in the United States found that 67% of disadvantaged–caste staff reported workplace mistreatment (California v. Cisco Systems, Inc., 2022), underscoring the likelihood of negative bias.

We now present a single theoretical model that unifies both these elements.

Consider a set of jobs $j \in \mathcal{J}$. Each job has up to four sequential rounds $s \in \{1, 2, 3, 4\}$:

- Rounds $s = 1, 2, 3$ are pre-interview screening rounds (e.g. application, written tests, group debates)
- Round $s = 4$ is the personal interview

Groups and Screening. There are two castes (groups) $g \in \{0, 1\}$, where $g = 1$ denotes disadvantaged castes and $g = 0$ denotes advantaged castes. At each round s, employers observe a score $M_{s,i,g}$ for candidate i from group g.[14] Passing round s requires $M_{s,i,g} \geq \tau_{s,j}$ and also requires having passed all prior rounds.

For simplicity, we assume employers *perfectly observe* these screening scores. This assumption is not critical for the key insights derived from this model—variation in screening round selectivity and final hiring thresholds across jobs can help distinguish discrimination from ability differences even with imperfect information (see Online Appendix Section A.3 for this extension).

At the final interview ($s = 4$), the threshold $\tau_{4,j}^{(g)}$ may depend on group g since caste becomes salient for nearly 85% of candidates at interviews (Section 5.6.1). Hence, any final-stage gap can reflect discrimination if $\tau_{4,j}^{(1)} > \tau_{4,j}^{(0)}$.

Monotonicity of thresholds. We assume

$$\tau_{4,j} \;\geq\; \tau_{3,j} \;\geq\; \tau_{2,j} \;\geq\; \tau_{1,j},$$

so each round is at least as demanding as the previous one.

Identification from exogenous variation in hiring thresholds across jobs. For our identification strategy, we assume that the underlying ability distributions are common across jobs. This assumption is crucial because it implies that differences in job thresholds τ_j reflect genuine variations in job selectivity, rather than systematic differences in candidate quality.

Imagine one large tub of ice cream representing the common talent pool. Each job dips into this tub with a ladle of a given size: a smaller ladle (higher cutoff) hires fewer

[14] We assume the performance distributions for each group are continuous with smooth densities, finite moments, and satisfy standard regularity conditions. These technical assumptions ensure well-behaved pass rates and expectations when calculating promotion probabilities. While these properties are implicit in our figures and analysis, they are not crucial for understanding the main results.

but on-average stronger candidates; a larger ladle (lower cutoff) hires more, lower-scoring candidates. Because everyone scoops from the same tub, ladle size captures selectivity alone.

1. Near-zero application costs and wide eligibility rules preclude sorting into jobs by caste (Section 4.7).
2. Admission quotas balance caste representation within majors and effectively guarantee common support across GPA quantiles (Section 4.3; Online Appendix Figure OA.1).
3. Caste is typically revealed only at interviews, preventing selective screening in earlier rounds.
4. Standardized recruitment protocols limit employers' ability to strategically manipulate cutoffs (Section 4).

Within a given job, a higher cutoff naturally yields higher-quality hires; across jobs, however, cutoffs vary exogenously with selectivity while the candidate pool remains comparable. We exploit this across-job variation as an instrument for selectivity in the empirical analysis.

(A) Ruling Out Systematic Performance Differences Before Interviews

We rule out systematic caste differences in screening round performance before personal interviews.

Proposition 1 (No Early Gap Implies No Right-Tail Deficit Before Interviews)**.** Consider a job j with up to four rounds of screening, where the first three ($s = 1, 2, 3$) are pre-interview. If group 1 truly performed systematically worse, then for a sufficiently large threshold $\tau_{s,j}$ at some pre-interview stage $s < 4$, the pass-rate gap

$$\Delta_{s,j} \;=\; P(\text{Pass}_{s,i,j}|g = 1) \;-\; P(\text{Pass}_{s,i,j}|g = 0) \;<\; 0$$

must arise. Therefore, if $\Delta_{s,j} \geq 0$ even under high thresholds—essentially focusing on the top tail of each group's performance distribution—this rules out a meaningful right-tail performance deficit for group 1 in earlier rounds.

Sketch of Proof. We illustrate how variation in screening thresholds helps identify any systematic performance differences before interviews.

Figure 3 contrasts two scenarios. In each panel, the blue (solid) and red curves depict two groups' performance distributions in pre-interview screening rounds. Recall that $g = 1$ denotes the disadvantaged group and $g = 0$ the advantaged group. We shade the portion of

each group's distribution above the vertical dashed line, representing the screening cutoff at that job. The contrast between panels reveals why selective jobs provide a powerful test for pre-interview performance differences:

- **Panel (a): Hypothetical Right-Tail Deficit.** Assume that for each group g, the screening score distribution $\varphi_g(m)$ has a regularly decaying tail; for example, for sufficiently large m, $P(M_{s,i,g} \geq m) \sim C_g\, m^{-\alpha}$ for some constants $C_g > 0$ and $\alpha > 1$. This condition implies that as the screening threshold increases (i.e., in more selective jobs), the conditional pass rate declines at a rate determined by the tail density. Consequently, if Group 1 has a thinner right tail than Group 0 (i.e., a smaller C_1 or larger effective decay rate), the red "passing region" would shrink *faster* than the blue, leading to increasingly *negative* pass-rate gaps in more selective jobs.

- **Panel (b): No Meaningful Deficit.** Fact 1 instead shows no systematic shortfall for Group 1 across the pre-interview screening performance distribution. Even at high cutoffs in selective jobs, both groups maintain similar pass rates. This pattern holds across application reading, written tests, and group debates.

The observation of similar pass rates across groups in highly selective jobs (Panel b) thus rules out systematic pre-interview performance differences.

Figure 3: **Using Variation in Screening Thresholds to Test for Pre-Interview Performance Differences**

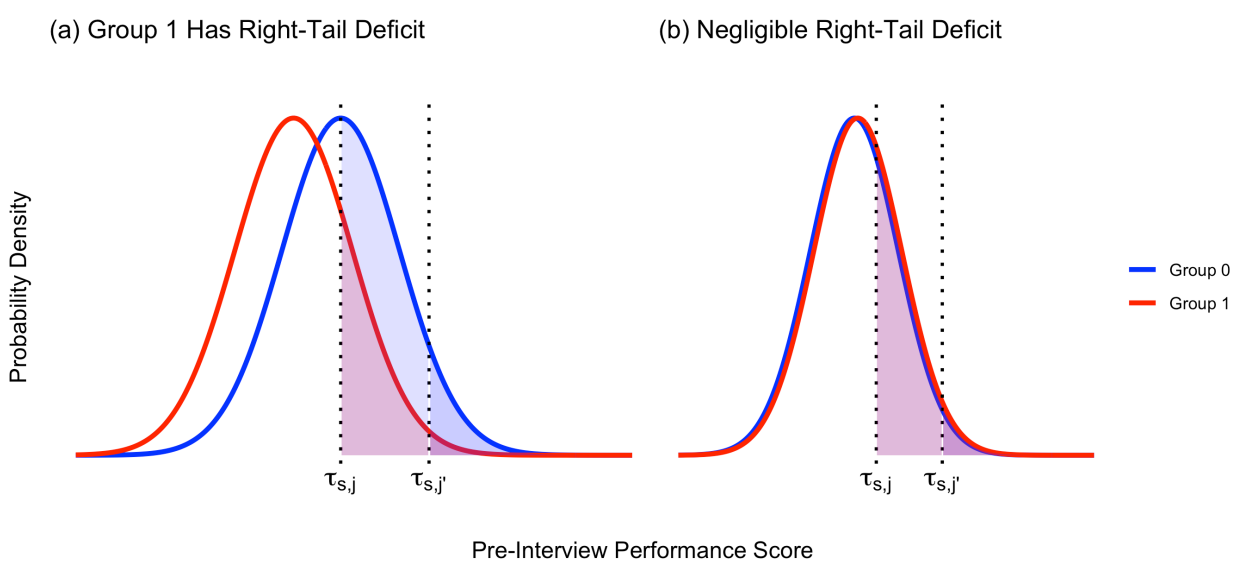

Notes: Figure 3 illustrates how selective screening would reveal systematic pre-interview performance differences. Panel (a) shows how a hypothetical right-tail deficit for Group 1 would become more apparent at higher cutoffs. Panel (b) shows the likely distribution of pre-interview performance scores implied by Fact 1, with no meaningful deficit. $\tau_{s,j}$ denotes screening threshold at stage $s < 4$ for job j (j' represents a more selective job).

□

Key Takeaway. Fact 1 shows that early-stage disadvantages remain negligible before interviews, thereby corroborating Proposition 1. Consequently, any late-stage gap at $s = 4$ must stem either from newly revealed ability signals at the interview or from discrimination.

(B) Using Variation in Hiring Thresholds to Test for Discrimination

Idea. Differences in job-level hiring cutoffs let us test whether the persistent 20% promotion edge for disadvantaged hires reflects stricter screening (i.e., discrimination) or hidden ability gaps. If both castes faced the same cutoff, this edge would shrink to zero in the most selective jobs; its endurance instead points to discrimination.

We begin by linking screening performance to job performance. Maintaining our perfect information assumption, we define job performance as the average screening score of accepted candidates. Let $\bar{M}_{i,g} = \frac{1}{4}\sum_{s=1}^{4} M_{s,i,g}$ denote the average screening score for candidate i from group g (any monotonic link $f(.)$ between screening round performance and job performance works, but we choose to keep things simple). A candidate is promoted if $\bar{M}_{i,g}$ exceeds a job-specific threshold κ_j.

For jobs in the top $(1-q)$ fraction by selectivity (for instance, if $q = 0.8$, then we are conditioning on the top 20% of jobs by selectivity), we distinguish between two types of thresholds: the *selectivity-tier threshold* $\tau^{q}_{4,j}$, which captures the job-level standard for jobs in the top $(1-q)$ fraction, and the *group-specific threshold* $\tau^{(g)}_{4,j}$, which may vary across groups within each job. Under discrimination, $\tau^{(1)}_{4,j} > \tau^{(0)}_{4,j}$; if no discrimination is present, these thresholds coincide (i.e., $\tau^{(1)}_{4,j} = \tau^{(0)}_{4,j} = \tau^{q}_{4,j}$).

Let

$$\bar{Y}^{(g)}_{j} \;=\; \mathbb{E}\big[\bar{M}_{i,g} \,\big|\, M_{4,i,g} \geq \tau^{(g)}_{4,j}\big]$$

denote the mean screening score of candidates hired from group g. Define the promotion gap in selectivity tier q as

$$\Delta^{\text{Prom}}(q) \;=\; \bar{Y}^{(1)}_{j} - \bar{Y}^{(0)}_{j}.$$

As a job's cutoff rises, the average hired score approaches that of the marginal candidate. If both castes share a common cutoff, this convergence forces $\Delta^{\text{Prom}}(q) \to 0$. If disadvantaged candidates face a higher cutoff— $\tau^{(1)}_{4,j} = \tau^{(0)}_{4,j} + \delta,\ \delta > 0$— they are screened on a uniformly stricter margin, so the gap stabilises near δ even in the most selective jobs.

Proposition 2 (Stable Promotion Gaps Across Selectivity Tiers Indicate Discrimination)**.** Let J_q denote the set of jobs in the top $(1-q)$ fraction by selectivity with associated thresholds $\tau^{q}_{4,j}$. Suppose that for jobs in J_q, the average screening score among accepted candidates for

group g is given by

$$\bar{Y}_j^{(g)} = \mathbb{E}\left[\bar{M}_{i,g} \mid M_{4,i,g} \geq \tau_{4,j}^{(g)}\right].$$

If both groups face the same hiring threshold (i.e., $\tau_{4,j}^{(1)} = \tau_{4,j}^{(0)}$), then as $q \uparrow$ (i.e., as jobs become more selective and $\tau_{4,j}^{q}$ increases), the average screening scores of accepted candidates in both groups converge to the common threshold, implying that $\Delta^{\text{Prom}}(q) \to 0$.

Conversely, if discrimination is present so that $\tau_{4,j}^{(1)} = \tau_{4,j}^{(0)} + \delta$ for some $\delta > 0$, then the average screening score for Group 1 converges to a value that is δ higher than that for Group 0, and the promotion gap remains approximately constant at δ across selectivity tiers.

Sketch of Proof. Our aim is to illustrate how variation in hiring thresholds distinguishes discrimination from ability differences.

Figure 4 contrasts two scenarios, with blue and red curves depicting the average screening round performance scores for the two groups. The red curve (Group 1) is shifted rightward relative to the blue curve (Group 0), consistent with both our assumption that average screening round performance is positively correlated with job performance and our empirical finding that disadvantaged castes have higher promotion rates post-hire. The vertical dashed lines represent job-level cutoffs, and the shaded regions above these thresholds denote the hired populations.

- **Panel (a): No Discrimination.** In the absence of discrimination, both groups face a common cutoff, say $\tau_{4,j}^{(1)} = \tau_{4,j}^{(0)} = \tau$. Under the regularity condition that each group's screening score distribution decays regularly (i.e., for sufficiently large m, $P(M_{s,i,g} \geq m) \sim C_g\, m^{-\alpha}$ for some $C_g > 0$ and $\alpha > 1$), it follows that

  $$\lim_{\tau \to \infty} \left(\mathbb{E}[\bar{M}_{i,g} \mid M_{4,i,g} \geq \tau] - \tau\right) = 0.$$

  Hence, as the common cutoff shifts rightward in more selective jobs, the average screening scores of accepted candidates in both groups converge to the cutoff value (i.e., the performance level of the marginal candidate), and the performance gap diminishes toward zero.

- **Panel (b): Discrimination.** If discrimination is present, then $\tau_{4,j}^{(1)} > \tau_{4,j}^{(0)}$ (note that when no discrimination is present, these thresholds coincide with the selectivity-tier threshold, i.e., $\tau_{4,j}^{(1)} = \tau_{4,j}^{(0)} = \tau_{4,j}^{q}$).

  Suppose that Group 1 faces a higher cutoff than Group 0 by a constant $\delta > 0$ (i.e., $\tau_{4,j}^{(1)} = \tau_{4,j}^{(0)} + \delta$). Under the regularity condition that for each group g,

  $$\lim_{\tau \to \infty} \Big(\mathbb{E}[\bar{M}_{i,g} \mid M_{4,i,g} \geq \tau] - \tau\Big) = 0,$$

it follows that as the threshold increases,

$$\lim_{\tau \to \infty} \left(\mathbb{E}[\bar{M}_{i,1} \mid M_{4,i,1} \geq \tau_{4,j}^{(1)}] - \mathbb{E}[\bar{M}_{i,0} \mid M_{4,i,0} \geq \tau_{4,j}^{(0)}] \right) = \delta.$$

In other words, the average screening score for Group 1 converges to a value that is δ higher than that for Group 0, so the performance gap remains approximately constant at δ even as the threshold increases.

Since a candidate is promoted if $\bar{M}_{i,g}$ exceeds a job-specific threshold κ_j, the asymptotic behavior of the screening score distributions implies that if both groups face the same threshold, the promotion gap $\Delta^{\text{Prom}}(q)$ shrinks toward zero, whereas a constant gap in thresholds (indicative of discrimination) preserves a stable promotion gap.

Figure 4: **Using Variation in Hiring Thresholds To Test for Discrimination**

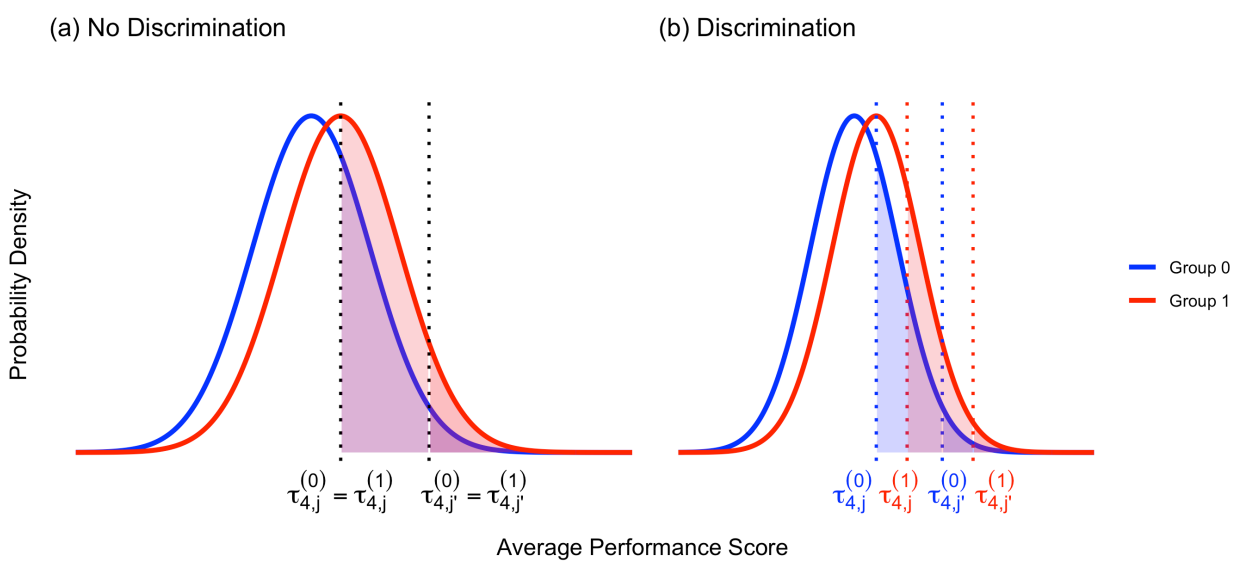

Notes: Figure 4 depicts two scenarios. In panel (a), a common cutoff for both groups shrinks their hired-performance gap at higher cutoffs. In panel (b), a constant δ gap in cutoffs preserves a stable performance advantage for Group 1. Here, $\tau_{4,j}^{(g)}$ denotes group g's interview threshold at job j, with j' representing a more selective job.

□

Not every hire sits on the margin. The test works even when most hires clear the bar comfortably: what matters is how *thresholds* shift across jobs. Selectivity-tier comparisons still reveal whether disadvantaged candidates face a higher bar, separating discrimination from true ability gaps.

Key Takeaway. Fact 2 offers compelling evidence for the discrimination scenario outlined in Proposition 2. The stable promotion gap across selectivity tiers—where disadvantaged

hires maintain a roughly 20% advantage—is inconsistent with an ability-differences explanation, which would predict a diminishing gap at higher thresholds.

Summary. The absence of caste penalties in even the most selective jobs' early rounds rules out meaningful ability deficits before interviews (see Fact 1 and Proposition 1). Moreover, when disparities do appear at interviews, their consistent pattern across jobs distinguishes their origin: the stable 20% promotion advantage observed for disadvantaged castes across all selectivity tiers (Table 7) is precisely what we would expect from discriminatory hiring thresholds (see Fact 2 and Proposition 2), rather than from unobserved ability differences.

Figure 5: **Caste Penalty in the Conditional Success Rate At Each Screening Round**

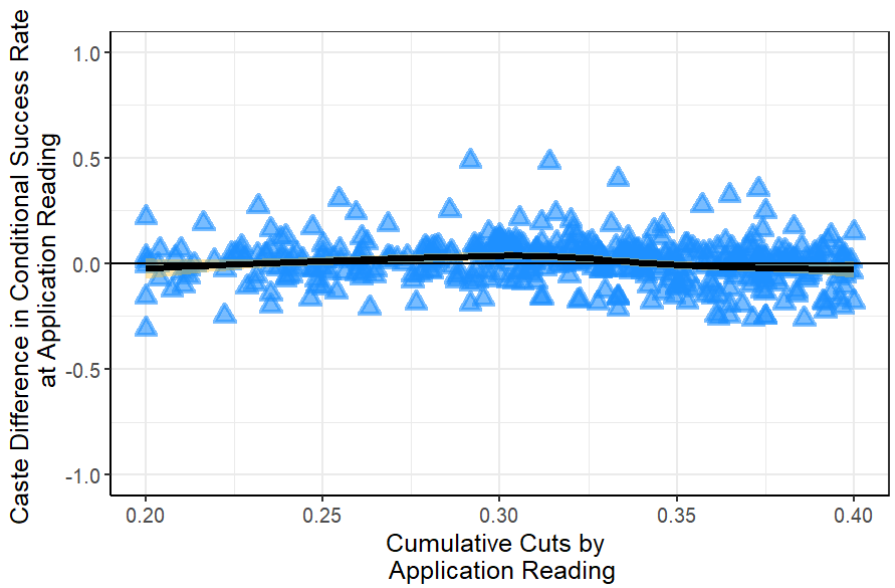

(a) Caste Penalty in Conditional Success Rates versus Cumulative Cuts by Application Reading

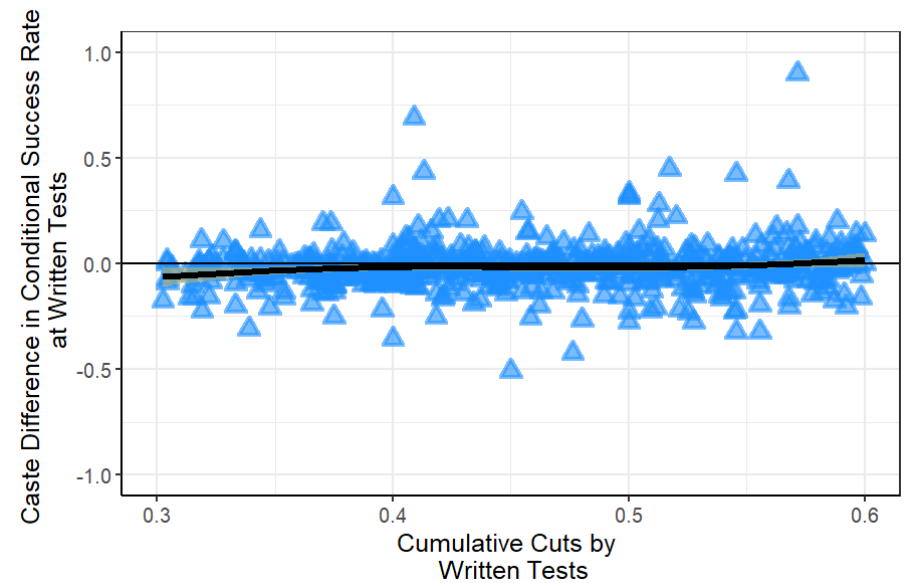

(b) Caste Penalty in Conditional Success Rates versus Cumulative Cuts by Written Tests

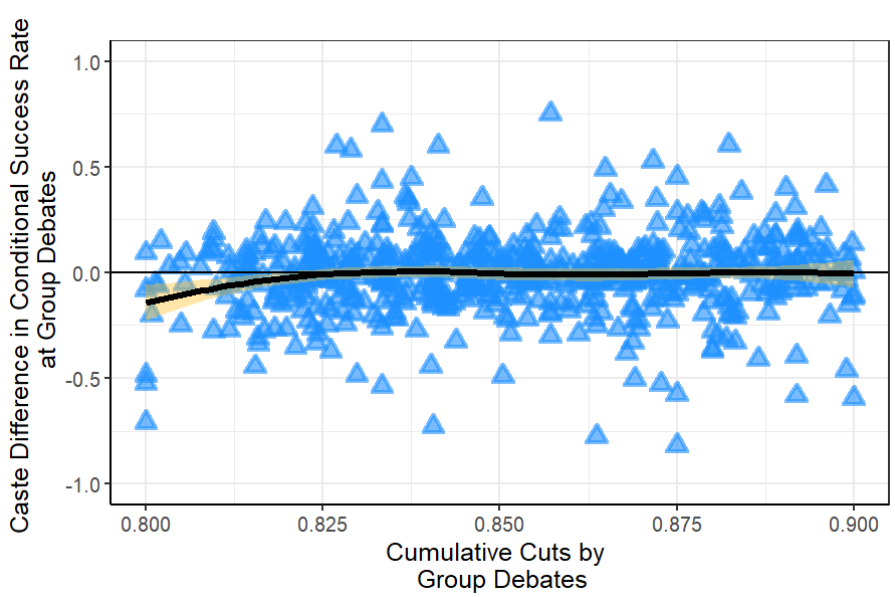

(c) Caste Penalty in Conditional Success Rates versus Cumulative Cuts by Group Debates

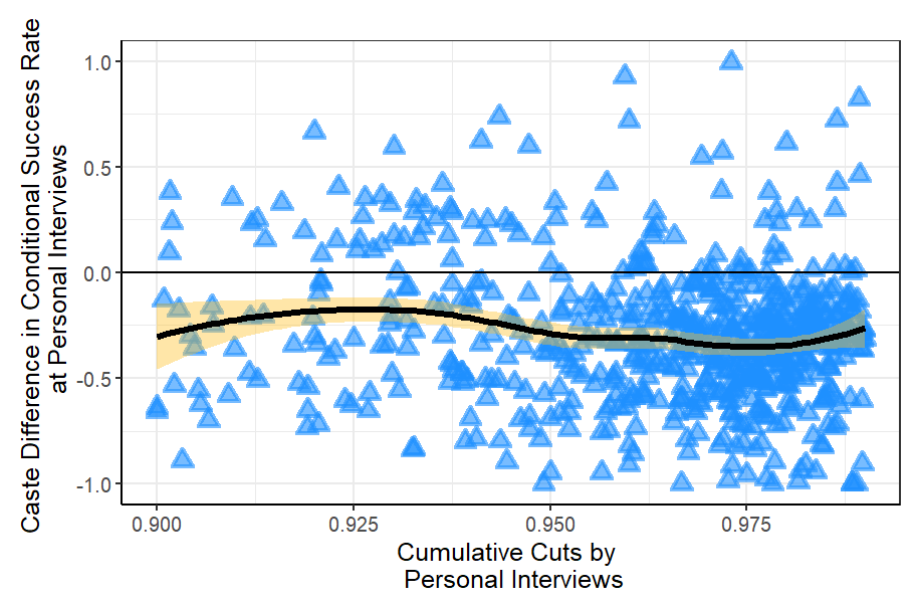

(d) Caste Penalty in Conditional Success Rates versus Cumulative Cuts by Personal Interviews

Notes: Figure 5 shows the caste difference in the conditional (i.e., with controls) success rate against the cumulative cuts made by jobs by the end of each screening round. Each blue triangle denotes a job. Smoothed conditional means (thick black lines) are calculated using the "loess" function in R. The shaded areas in yellow denote 95% confidence intervals. The thin solid horizontal line denotes a y-intercept of 0. The x-axis shows the cumulative proportion of students cut by jobs by the end of each screening round. The y-values are calculated by estimating, *for each job*, a standard linear probability model with controls and cohort-year fixed effects and then plotting the coefficient on disadvantaged caste from these regressions on the y-axis. Each regression includes controls for student characteristics such as degree, major, college GPA, entrance exam score as well as cohort-year fixed effects.

Table 7: **Caste Differences in Promotion Rate**

	Dependent variable:									
	Promoted	Promoted	Promoted at Top 75%	Promoted at Top 50%	Promoted at Top 40%	Promoted at Top 30%	Promoted at Top 20%	Promoted at Top 15%	Promoted at Client Facing	Promoted at Non-Client Facing
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)	(9)	(10)
Disadv. Caste (OLS)	0.137***	0.096***	0.104***	0.112***	0.136***	0.147***	0.190**	0.204**	0.107**	0.082**
	(0.025)	(0.026)	(0.029)	(0.038)	(0.043)	(0.054)	(0.073)	(0.084)	(0.044)	(0.032)
Disadv. Caste (Oster)	0.142***	0.093***	0.108***	0.122***	0.141***	0.157**	0.196**	0.206**	0.106***	0.075**
	(0.026)	(0.028)	(0.031)	(0.037)	(0.041)	(0.058)	(0.082)	(0.082)	(0.039)	(0.036)
Tenure Duration		0.084***	0.082***	0.076***	0.074***	0.068***	0.054*	0.079**	0.065***	0.097***
		(0.012)	(0.014)	(0.019)	(0.021)	(0.025)	(0.030)	(0.037)	(0.021)	(0.015)
Mean Y	0.684	0.684	0.667	0.668	0.676	0.711	0.711	0.706	0.687	0.682
Observations	2,489	2,489	2,012	1,246	997	748	502	374	835	1,654
Adjusted R^2	0.183	0.205	0.203	0.216	0.211	0.182	0.192	0.189	0.194	0.214
Job FE	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
Year FE	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
Other Controls		✓	✓	✓	✓	✓	✓	✓	✓	✓

Notes: Table 7 shows caste differences in job performance measures among students who graduated with full-time jobs through the college's on-campus job fair. All of these students had to succeed in the final personal interview round to obtain jobs. Column (1) reports caste differences in promotion rates only with job and year fixed effects. Column (2) reports the same regression as Column (1) but with all other controls such as GPA, major, degree, and job tenure. Columns (3)-(8) show caste differences in promotion rates with controls among jobs at different quantiles of "job selectivity," proxied by the cumulative proportion of applicants cut by a given job by the end of the final screening round. Columns (9)-(10) show caste differences in promotion rates with controls among client-facing and non-client-facing jobs, respectively. Bias-adjusted coefficients are calculated using Oster (2019). Following Oster's recommended benchmarks, I set $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$, implying that unobservables are assumed to have the same degree of selection as observables, and the total explanatory power can increase by up to 30% beyond the observed R^2. $^*p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

5.8 Classification and Preferred Interpretation of the Caste Penalty: Distinguishing from Statistical or Customer Discrimination

Given the evidence that employers directly penalize caste, the most plausible reading is *disparate treatment*—explicit discrimination based on caste. Two facts support this view. First, disadvantaged hires are *more* likely to be promoted across all selectivity tiers (Table 7). Second, more than three-quarters of the roles are non-customer-facing (Table 3). Together, these points make it hard to attribute the gap to accurate statistical inference or customer preferences; they fit better with caste bias.

Caste discrimination is expressly illegal under Article 15 of the Indian Constitution (Ministry of Law and Justice, Government of India, 2021). It is possible that "fit'' criteria generate a *disparate impact* (e.g., favoring certain extracurricular backgrounds), but the full body of evidence in Sections 5.2–5.7 still points most strongly to disparate treatment.

5.9 What Happens During Personal Interviews?

We do not observe interview transcripts, so we cannot parse the exact conversations. Prior ethnographic work on elite hiring shows that final-round interviews at multinationals heavily weight "cultural-fit'' markers linked to class background rather than job-relevant skill (e.g. Rivera, 2012, 2015; Lamont, 2009). Similar patterns appear in India's civil-service exams, where subjective "personality'' interviews disproportionately penalise lower-status candidates without predicting performance (Bhavnani and Lee, 2021). These external findings reinforce our interpretation that the large, late-stage caste penalty we document is driven by discretionary, status-based assessments rather than objective ability.

6 Alternative Explanations

I examine whether any competing story can account for caste gaps that surface only after personal interviews. None matches the evidence.

1. **Early "fairness" to appear inclusive.** If firms lowered early-round cut-offs for disadvantaged students, those castes should enjoy *higher* conditional success before interviews—especially in the most selective jobs. Figure 5 shows parity, not advantage, and students with caste-identifiable surnames are penalised from the start (Figure 2).

2. **Delayed discrimination saved for interviews.** Employers face no legal or reputational pressure to postpone caste bias; Indian MNCs often view caste-targeted hiring as "anti-merit" (Jodhka and Newman, 2007). When caste is already obvious (identifiable surnames), penalties appear at the *application* stage, not later. Interviews matter because they newly reveal caste for the remaining 85% of students.

3. **Hidden right-tail skill gaps.** A genuine ability deficit would depress disadvantaged pass-rates in early rounds of highly selective jobs. Fact 1 shows no such pattern.

4. **Attrition between rounds is negligible.** Four enforced rounds (Table 4) account for virtually all exits, *and the placement office records every advance-or-out decision*, so unrecorded drop-out is negligible.

5. **Better outside options for advantaged castes.** Return-offer takers exit *before* screening and are under 5% of the cohort (Section 3.2). Parallel job search forfeits placement support and is rare.

6. **Socio-emotional skill differences.** Group debates—designed to test communication and teamwork—show no caste gap (Panels (c) of Figures 5–OA.3).

7. **Missing data or skipped rounds.** The placement office logs all rounds, mandates attendance, and prohibits firms from skipping steps (Section 4.2); exit surveys confirm record accuracy.

8. **Negotiation advantages for advantaged castes.** Pay (and benefits) are posted ex-ante, non-negotiable, and verified by the placement office; exit data confirm identical starting salaries (Section 3.2).

9. **Competition from government jobs.** Quota-regulated public positions make up <2% of our sample (Table 3); private employers face negligible competition.

10. **Caste differences in job preferences.** Figure 1 shows preference-sorting explains only a small share of the initial earnings gap and none of the sharp interview-stage penalty.

11. **Preferential promotion of disadvantaged hires.** Less than 10% of large U.S. firms track socioeconomic diversity, and under 3% of India's top-100 listed firms record caste data (Stansbury and Rodriguez, 2024; BusinessLine, 2018). Formal caste-aware promotion is rare; survey evidence documents bias *against* disadvantaged staff (California v. Cisco Systems, Inc., 2022). The observed 20% promotion edge therefore reflects stricter hiring thresholds, not reverse favoritism.

No alternative reconciles the absence of pre-interview gaps with the large, interview-stage penalties and the persistent promotion advantage of disadvantaged hires. The evidence is most consistent with caste-based disparate treatment during personal interviews.

7 External Validity

How far do results from one elite Indian college travel? Three features make the findings broadly informative for high-stakes recruitment settings.

1. Caste offers a uniquely clean status marker. In Western settings, status signals—accent, schooling, networks—are diffuse. Indian caste bundles economic, social, and cultural capital into a single, officially recorded label (Beteille, 1965; Bourdieu, 1986). That clarity lets us isolate the point—subjective "fit'' interviews—at which status is penalised. The same late-stage conversations screen for hidden hierarchy wherever employers lean on informal rapport (Rivera, 2015; Laurison and Friedman, 2019).

2. The setting is representative inside India (and inside firms). Nearly all top-30 Indian colleges run identical four-round placement fairs; the same 30–40 multinationals visit each campus, using one national salary grid and common hiring teams. The process here therefore mirrors elite recruitment at peer institutions and in these firms' global offices (Sections 3.3–4.2).

3. The mechanism is replicated across contexts. As obvious status cues fade in urban labour markets (Clark, 2015), deep interpersonal interviews become the main detector of social background. Our data show *nearly 90%* of conditional gaps arise in those conversations—paralleling patterns in "test-optional'' U.S. admissions or UK elite hiring (Sacerdote et al., 2025; Chetty et al., 2023; Laurison and Friedman, 2019).

8 Conclusion

By opening the hiring "black box," this paper delivers three firsts. It provides the first large-scale, step-by-step evidence on when class-based bias surfaces in elite private-sector hiring; the first clean quantification of its magnitude (nearly 90% of the gap concentrated in one "fit'' interview); and the first direct link between those hiring decisions and longer-run performance, showing a nearly 20% promotion edge for the very candidate group filtered out. Because caste is an administratively defined and widely recognized marker of social status, the Indian setting offers an unusually precise lens on status-based disadvantage—precision that can inform parallel analyses in settings where class must be inferred from subtler proxies.

Why does bias emerge predominantly at the final conversation? As obvious signals such as surnames or accents fade in global, urban labor markets, recruiters infer social position

from richer dialogue about hobbies, extracurriculars, polish, and background (Jodkha, 2017). Fast, intuitive judgments often crowd out more deliberative, merit-based reasoning (Kahneman, 2011). Parallel late-stage penalties—class accents in the UK, varsity-sport résumés in the US—show the same mechanism at work (Rivera, 2015; Laurison and Friedman, 2019; Stansbury and Rodriguez, 2024) even where caste is absent.

The findings expose a leak in the affirmative-action pipeline: policies that widen university gates are partially undone by private-sector interviews that re-privilege familiar culture. Firms pay the price. Hiring managers—rarely future supervisors—reward short-term rapport, often discarding highly-talented candidates (Benson and Shaw, 2024; Hoffman and Stanton, 2024).

Eliminating interviews is neither feasible nor wise; taming them is. Structured behavioral questions, blinded scoring rubrics, and feedback loops tying interviewers to recruits' eventual productivity can preserve human judgment while curbing bias. Yet fewer than 10% of large U.S. firms and under 3% of India's top-100 track social-class data (Stansbury and Rodriguez, 2024; BusinessLine, 2018), forfeiting the chance to detect these costly "final-hurdle'' failures.

As remote hiring and algorithmic screening spread, conversational bias could shrink—or migrate online (Albright, 2024). Extending stage-by-stage audits across countries and technologies will be vital for designing hiring systems that advance both equity and organizational performance worldwide.

References

[1] Abaluck, J., L. Agha, C. Kabrhel, A. Raja, and A. Venkatesh (2016). The determinants of productivity in medical testing: Intensity and allocation. *American Economic Review 106*(12), 3730–3764.

[2] Albright, A. (2024). The hidden effects of algorithmic recommendations. Working paper, Federal Reserve Bank of Minneapolis.

[3] Altonji, J. G. and R. M. Blank (1999). *Race and Gender in the Labor Market. Handbook of Labor Economics*, ed. by O. C. Ashenfelter and D. Card, Elsevier, vol. 3C, chap. 48.

[4] Alvero, A., S. Giebel, B. Gebre-Medhin, A. L. Antonio, M. L. Stevens, and B. W. Domingue (2021). Essay content and style are strongly related to household income and sat scores: Evidence from 60,000 undergraduate applications. *Science Advances 7*(42), eabi9031.

[5] Asher, S., P. Novosad, and C. Rafkin (2024). Intergenerational mobility in india: New measures and estimates across time and social groups. *AEJ: Applied Economics 16*(2), 66–98.

[6] Barg, K. (2013). The influence of students' social background and parental involvement on teachers' school track choices: Reasons and consequences. *European Sociological Review 29*(3), 565–579.

[7] Barrios-Fernández, A., C. Neilson, and S. Zimmerman (2024). Elite universities and the intergenerational transmission of human and social capital. *American Economic Review*. conditionally accepted.

[8] Bayly, S. (2008). *Caste, Society and Politics in India from the Eighteenth Century to the Modern Age*. Chicago: Cambridge University Press.

[9] Benson, A., S. Board, and M. Meyer-ter-Vehn (2024). Discrimination in hiring: Evidence from retail. *Review of Economic Studies 91*(4), 1956–1987.

[10] Benson, A., D. Li, and K. Shue (2019). Promotions and the peter principle. *The Quarterly Journal of Economics 134*(4), 2085–2134.

[11] Benson, A., D. Li, and K. Shue (2023). "Potential" and the gender promotion gap. R&R, *American Economic Review*.

[12] Benson, A. and K. Shaw (2024). What do managers do? *Annual Review of Economics*. Forthcoming.

[13] Beteille, A. (1965). *Caste, Class and Power: Changing Patterns of Stratification in a Tanjore Village*. University of California Press.

[14] Beteille, A. (1969). *Castes: Old and New, Essays in Social Structure and Social Stratification*. Asia Publishing House.

[15] Bhavnani, R. R. and A. Lee (2021). Does affirmative action worsen bureaucratic performance? evidence from the indian administrative service. *American Journal of Political Science 65*(1), 5–20.

[16] Bisin, A., J. Rubin, A. Seror, and T. Verdier (2023). Culture, institutions & the long divergence. *Journal of Economic Growth 28*, 1–40.

[17] Bisin, A. and T. Verdier (1998). On the cultural transmission of preferences for social status. *Journal of Public Economics 70*, 75–97.

[18] Bisin, A. and T. Verdier (2010). The economics of cultural transmission and socialization. In *Handbook of Social Economics*, pp. 339–416. Elsevier.

[19] Bisin, A. and T. Verdier (2017). The joint dynamics of culture and institutions. *Journal of the European Economic Association 15*(6), 1313–1350.

[20] Bloom, N., R. Sadun, and J. Van Reenan (2012). The organization of firms across countries. *The Quarterly Journal of Economics 127*(4), 1663–1705.

[21] Bol, T., J. Witschge, H. G. Van de Werfhorst, and J. Dronkers (2014). Curricular tracking and central examinations: Counterbalancing the impact of social background on student achievement in 36 countries. *Social Forces 92*(4), 1545–1572.

[22] Borneman, M. J., S. R. Cooper, D. M. Klieger, N. R. Kuncel, and N. Brenckman (2007). Admission interviews: Not so nice to meet you. In *Alternative Predictors of Academic Performance, Symposium conducted at the Annual Meeting of the Society for Industrial and Organizational Psychology, New York, NY.*

[23] Bourdieu, P. (1986). The forms of capital. In J. G. Richardson (Ed.), *Handbook of Theory and Research for the Sociology of Education*, pp. 241–258. Greenwood Press.

[24] BusinessLine (2018). An Inclusive India Inc. Business Line, January 13, 2018.

[25] California v. Cisco Systems, Inc. (2022). Dept. of Fair Employment and Housing v. Cisco Systems, Inc. Docket No. H048962.

[26] Cappelli, P. (2019). Your approach to hiring is all wrong. *Harvard Business Review.* Accessed: 2024-08-15.

[27] Chandra, A. and D. O. Staiger (2010). Identifying provider prejudice in healthcare. Technical Report 16382, National Bureau of Economic Research.

[28] Chandra, A. and D. O. Staiger (2020). Identifying sources of inefficiency in health care. *Quarterly Journal of Economics 135*(2), 785–843.

[29] Checchi, D. and H. G. van de Werfhorst (2018). Policies, skills and earnings: How educational inequality affects earnings inequality. *Socio-Economic Review 16*(1), 137–160.

[30] Chetty, R., D. J. Deming, and J. N. Friedman (2023). Diversifying society's leaders? the determinants and causal effects of admission to highly selective private colleges. NBER Working Paper 31492.

[31] Clark, G. (2015). *The Son Also Rises: Surnames and the History of Social Mobility.* The Princeton Economic History of the Western World.

[32] Cooper, S. R., D. M. Klieger, M. J. Borneman, N. R. Kuncel, and H. Mikeworth (2007). Personal statements: Please keep them to yourself. In *Alternative Predictors of Academic Performance, Symposium conducted at the Annual Meeting of the Society for Industrial and Organizational Psychology, New York, NY.*

[33] Datta, S. (2017). *A History of the Indian University System: Emerging from the Shadows of the Past.* Palgrave Macmillan.

[34] Deshpande, A. (2011). *The Grammar of Caste: Economic Discrimination in Contemporary India.* Oxford: Oxford University Press.

[35] Durlauf, S. N. (1999). The case "against" social capital. In P. Dasgupta and I. Serageldin (Eds.), *Social Capital: A Multifaceted Perspective*, pp. 84–104. Washington, D.C.: World Bank.

[36] Durlauf, S. N. (2004). Neighborhood effects. In J. V. Henderson and J.-F. Thisse (Eds.), *Handbook of Regional and Urban Economics*, Volume 4, pp. 2173–2242. Amsterdam: Elsevier.

[37] Durlauf, S. N. (2006). Groups, social influences, and inequality: A memberships theory perspective on poverty traps. In S. Bowles, S. N. Durlauf, and K. Hoff (Eds.), *Poverty Traps*. Princeton: Princeton University Press.

[38] Durlauf, S. N. (2008). Memberships and inequality. In D. Castiglione, J. van Deth, and G. Wolleb (Eds.), *Handbook of Social Capital*. Oxford: Oxford University Press.

[39] Eames, T. (2024). Taryn versus taryn (she/her) versus taryn (they/them): A field experiment on pronoun disclosure and hiring discrimination. Working Paper.

[40] Eames, T. (2025). How does sex-based occupational discrimination affect nonbinary applicants? *AEA Papers and Proceedings*.

[41] Economic Advisory Council to the Prime Minister (2022). The State of Inequality in India Report. Government of India.

[42] Espinosa, M. and C. T. Stanton (2022, July). Training, communications patterns, and spillovers inside organizations. Working Paper 30224, National Bureau of Economic Research.

[43] Fenizia, A. (2022). Managers and productivity in the public sector. *Econometrica 90*(3), 1063–1084.

[44] Friedman, S. and D. Laurison (2020). *The Class Ceiling: Why it Pays to be Privileged?* Bristol University Press, Policy Press.

[45] Gibbons, R. and J. Roberts (2013). Economic theories of incentives in organizations. *The Handbook of Organizational Economics*.

[46] Gorman, E. H. (2005). Gender stereotypes, same-gender preferences, and organizational variation in the hiring of women: Evidence from law firms. *American Sociological Review 70*(4), 702–728.

[47] Grusky, D. B. (2005). Foundations of a neo-durkheimian class analysis. In E. O. Wright (Ed.), *Approaches to Class Analysis*, pp. 51–81. Cambridge University Press.

[48] Highet, K. (2023). "RIP English": Race, Class and 'Good English' in India. *Journal of Linguistic Anthropology 33*(2), 184–201.

[49] Hoffman, M., L. B. Kahn, and D. Li (2018). Discretion in Hiring. *The Quarterly Journal of Economics 133*(2), 765–800.

[50] Hoffman, M. and C. T. Stanton (2024). People, practices, and productivity: A review of new advances in personnel economics. NBER Working Paper 32849.

[51] Hoffman, M. and S. Tadelis (2021). People management skills, employee attrition, and manager rewards: An empirical analysis. *Journal of Political Economy 129*(1), 243–285.

[52] Human Rights Watch (2001). Caste discrimination: A global concern. HRW Briefing Paper. See ch. 2, p. 2: "Lower-caste communities are almost invariably indistinguishable in physical appearance from higher-caste communities".

[53] Imas, A. (2016). The realization effect: Risk-taking after realized versus paper losses. *American Economic Review 106*(8), 2086–2109.

[54] Imas, A., G. Loewenstein, and C. K. Morewedge (2021). Mental money laundering: A motivated violation of fungibility. *Journal of the European Economic Association 19*(4), 2209–2233.

[55] Imas, A. and K. Madarász (2024). Superiority-seeking and the preference for exclusion. *Review of Economic Studies 91*(4), 2347–2386.

[56] Jodhka, S. S. and J. Naudet (2019). *Mapping the Elite: Power, Privilege, and Inequality.* Oxford University Press.

[57] Jodhka, S. S. and K. Newman (2007). In the name of globalisation: Meritocracy, productivity and the hidden language of caste. *Economic and Political Weekly 42*(41), 4125–4132.

[58] Jodkha, S. S. (2017). *Caste in Contemporary India.* Routledge India.

[59] Kahneman, D. (2011). *Thinking, Fast and Slow.* New York: Farrar, Straus and Giroux.

[60] Khan, S. R. (2011). *Privilege: The Making of an Adolescent Elite at St. Paul's School.* Princeton, NJ: Princeton University Press.

[61] Kline, P., E. Rose, and C. Walters (2022). Systemic discrimination among large u.s. employers. *Quarterly Journal of Economics 137*(4), 1963–2036.

[62] Knobe, J. and S. Nichols (2008). An introduction to experimental philosophy. In *Experimental Philosophy*, pp. 3–34. Oxford: Oxford University Press.

[63] Kuncel, N. R., M. Credé, and L. L. Thomas (2007). A meta-analysis of the predictive validity of the graduate management admission test (gmat) and undergraduate grade point average (ugpa) for graduate student academic performance. *Academy of Management Learning & Education 6*(1), 51–68.

[64] Kuncel, N. R. and S. A. Hezlett (2007). Standardized tests predict graduate students' success. *Science 315*(5815), 1080–1081.

[65] Kuncel, N. R., D. M. Klieger, B. S. Connelly, and D. S. Ones (2013). Mechanical versus clinical data combination in selection and admissions decisions: A meta-analysis. *Journal of Applied Psychology 98*(6), 1060–1072.

[66] Kuncel, N. R., R. J. Kochevar, and D. S. Ones (2014). A meta-analysis of letters of recommendation in college and graduate admissions: Reasons for hope. *International Journal of Selection and Assessment 22*(1), 101–107.

[67] Lamont, M. (1992). *Money, Morals, and Manners: The Culture of the French and the American Upper-Middle Class*. Chicago: University of Chicago Press.

[68] Lamont, M. (2009). *How Professors Think: Inside the Curious World of Academic Judgment*. Cambridge, MA: Harvard University Press.

[69] Lamont, M. and A. Lareau (1988). Cultural capital: Allusions, gaps and glissandos in recent theoretical developments. *Sociological Theory 6*(2), 153–168.

[70] Laurison, D. and S. Friedman (2016). The class pay gap in higher professional and managerial occupations. *American Sociological Review 81*(4), 668–695.

[71] Laurison, D. and S. Friedman (2019). *The Class Ceiling: Why it Pays to be Privileged*. Bristol, UK: Policy Press.

[72] Lazear, E. P., K. L. Shaw, and C. T. Stanton (2015). The value of bosses. *Journal of Labor Economics 33*(4), 823–861.

[73] Lenski, G. E. (1954). Status crystallization: A non-vertical dimension of social status. *American Sociological Review 19*(4), 405–413.

[74] Madheswaran, S. (2008). Is affirmative action policy for private sector necessary? *Indian Journal of Industrial Relations 44*(2), 164–184.

[75] Ministry of Law and Justice, Government of India (2021). The constitution of india. Article 15 explicitly prohibits discrimination on grounds of religion, race, caste, sex, or place of birth.

[76] Minni, V. (2023). Making the invisible hand visible: Managers and the allocation of workers to jobs. R&R, *Quarterly Journal of Economics*.

[77] Minni, V. (2024). Global managers, local workers: Wage setting inside a multinational firm. *AEA Papers & Proceedings 114*, 586–91.

[78] Minni, V., K.-T. Nguyen, and H. Sarsons (2025). Managers and the cultural transmission of gender norms. Working Paper.

[79] Mishra, N. (2015). India and colorism: The finer nuances. *Washington University Global Studies Law Review 14*, 725–750.

[80] Murphy, S. C., D. M. Klieger, M. J. Borneman, and N. R. Kuncel (2009). The predictive power of personal statements in admissions: A meta-analysis and cautionary tale. *Educational and Psychological Measurement 69*(6), 967–982.

[81] Newman, K. and S. Thorat (2012). *Blocked by Caste: Economic Discrimination in Modern India*. Oxford University Press.

[82] Novosad, P., S. Asher, C. Farquharson, and E. Iljazi (2024). Access to opportunity in the sciences: Evidence from the nobel laureates. Working Paper.

[83] Obermeyer, Z., B. Powers, C. Vogeli, and S. Mullainathan (2019). Dissecting racial bias in an algorithm used to manage the health of populations. *Science 366*(6464), 447–453.

[84] Oster, E. (2019). Unobservable selection and coefficient stability: Theory and evidence. *Journal of Business & Economic Statistics 37*(2), 187–204.

[85] Parameswaran, R. and K. Cardoza (2015). Melanin on the margins: Advertising and the cultural politics of fair/light/white beauty in india. *Journalism & Communication Monographs 11*(3), 213–274.

[86] Parmar, R. (2020). Transacting caste in modern times: Changing social identity through surnames in urban Gujarat. *Contemporary Voice of Dalit 12*(2), 220–233.

[87] Pinker, S. (2021). *Rationality: What It Is, Why It Seems Scarce, Why It Matters*. New York: Viking.

[88] Puri, S. (2025, February). Capitalism with Indian characteristics. Stigler Center Lecture Series, University of Chicago Booth School of Business. Accessed: 2025-02-11.

[89] Ridgeway, C. L. (2019). *Status: Why Is It Everywhere? Why Does It Matter?* New York, NY: Russell Sage Foundation.

[90] Rivera, L. A. (2012). Hiring as cultural matching: The case of elite professional service firms. *American Sociological Review 77*(6), 999–1022.

[91] Rivera, L. A. (2015). *Pedigree: How Elite Students Get Elite Jobs*. Princeton University Press.

[92] Sacerdote, B., D. O. Staiger, and M. Tine (2025, January). How test optional policies in college admissions disproportionately harm high achieving applicants from disadvantaged backgrounds. Working Paper 33389, National Bureau of Economic Research. Available at http://www.nber.org/papers/w33389.

[93] Sadun, R., N. Bloom, and J. Van Reenen (2017, September-October). Why do we undervalue competent management? Neither great leadership nor brilliant strategy matters without operational excellence. *Harvard Business Review 95*(5), 120–127. Winner of 59th Annual HBR McKinsey Award.

[94] Sadun, R., J. B. Fuller, S. Hansen, and P. Neal (2022, July-August). The C-suite skills that matter most. *Harvard Business Review 100*(4), 42–50.

[95] Sarsons, H. (2017). Recognition for group work: Gender differences in academia. *AEA Papers and Proceedings 107*(5), 141–145.

[96] Sarsons, H. (2023). Interpreting Signals in the Labor Market: Evidence from Medical Referrals. R&R, *Review of Economic Studies*.

[97] Sarsons, H., K. Gërxhani, E. Reuben, and A. Schram (2021). Gender differences in recognition for group work. *Journal of Political Economy 129*(1), 101–147.

[98] Sarsons, H., J. Hjort, and X. Li (2022). Across-Country Wage Compression in Multinationals. R&R, *American Economic Review*.

[99] Sebring, J. M. (1969). Caste indicators and caste identification of strangers. *Human Organization 28*(3), 199–207.

[100] Sharma, D. (2017). English in india. In J. Filipović and S. C. Deshors (Eds.), *World Englishes: New Theoretical and Methodological Considerations*, pp. 311–332. De Gruyter Mouton.

[101] Stansbury, A. and K. Rodriguez (2024). The class gap in career progression: Evidence from academia. R&R, *Econometrica*.

[102] Stansbury, A. and R. Schultz (2023). The economics profession's socioeconomic diversity problem. *Journal of Economic Perspectives 37*(4), 207–230.

[103] Subramanian, S. (2019). India's elite colleges and the employment divide. *Economic and Political Weekly 54*(2), 42–48.

[104] Todorov, A. (2017). *Face Value: The Irresistible Influence of First Impressions*. Princeton, NJ: Princeton University Press.

[105] Turco, C. J. (2010). Cultural foundations of tokenism: Evidence from the leveraged buyout industry. *American Sociological Review 75*(6), 894–913.

[106] University World News (2022). Thousands sit 'world's toughest' entrance exam, for IITs. University World News, August 27, 2022.

[107] Wiltshire, C. R. (2020). *Uniformity and Variability in the Indian English Accent*. Cambridge University Press.

[108] Yenmez, M. B. and T. Sönmez (2022). Affirmative action in India via vertical, horizontal, and overlapping reservations. *Econometrica 90*(3), 1143–1176.

[109] Zimmerman, S. D. (2019). Elite colleges and upward mobility to top jobs and top incomes. *American Economic Review 109*(1), 1–47.

A ONLINE APPENDIX

A.1 Representativeness of the Job Fair of the Elite College

The studied on-campus job fair is a reliable window on how India's elite-college graduates enter elite private-sector jobs, for four reasons:

1. **Near-identical fairs across elite colleges.** Since the 1950s, India's top institutions—modelled on MIT and Stanford—have copied one another's academic calendars, faculty ratios, and placement offices. Today the 23 Indian Institutes of Technology, 20 Indian Institutes of Management, and leading Delhi University colleges all run virtually the same selective job fairs (Central Placement Cell, Delhi University).

2. **Graduates flow mainly into elite private jobs.** Most alumni of these colleges join high-paying entry-level posts at multinationals rather than the public sector(Subramanian, 2019; Jodhka and Naudet, 2019).

3. **Elite firms hire almost exclusively from elite colleges.** Recruiters in my sample visit roughly 40 similar campuses each year, each enrolling 1,000+ students.[OA.1] SeeSubramanian (2019); Jodhka and Naudet (2019).

4. **Few students skip the fair.** Less than 5% of final-year students deregister from the placement process (Section 3.2); comparable rates hold at peer colleges(Subramanian, 2019).

Together, these facts show that the focal job fair closely mirrors elite hiring pathways across India's top universities.

A.2 Timeline and Screening Rounds

Every job in the sample follows the same four-stage sequence, eliminating candidates at each step.

1. **Application reading (Sept.).** Recruiters make a first cut—largely by GPA, major, degree, and standout résumé items (e.g., internships, competitions).

2. **Written aptitude tests (Sept.–Nov.).** Online "technicals'' vary by role: coding for tech/manufacturing, case-style for consulting. Firms advertising several posts usually administer one common test (e.g., Nvidia's profiles share a single exam).

3. **Group debates (Sept.–Nov.).** Non-technical soft-skill screen.

   (a) Batches of 30–35 students (split into two teams) debate a general prompt such as "Is social media a boon or a bane?" while recruiters observe communication, teamwork, and poise.

[OA.1] Firm lists scraped from multiple placement websites; details available on request.

(b) Firms often run one debate for all their advertised roles—Google, for instance, may assess program-manager and software-engineer applicants together.

4. **Personal interviews (Dec.–Jan.).** Final 30-minute "fit" conversations.

   (a) Conducted in multiple parallel batches, each handled by different members of the recruiting team.

   (b) As with earlier rounds, a single interview panel commonly covers several job titles (e.g., Uber interviews data-scientist and frontend-engineer candidates in one track).

   (c) Because 98% of roles are India-based, interviewers are almost certainly Indian (Section 4.5).

A.3 Extending the Simple Model of Late-Stage Discrimination with Imperfect Information

In Section 5.7.1, we assumed that employers observe each candidate's true screening-round performance perfectly, i.e., $M_{s,i,g} = A_{i,g}$ at round s. We now show that *all* of the key insights—particularly how job-selectivity variation identifies discrimination versus ability deficits—remain valid even when employers observe *noisy* performance signals.

Setup. As in Section 5.7.1, there are two groups (castes) $g \in \{0, 1\}$, a continuum of jobs $j \in \mathcal{J}$, and up to four sequential screening rounds $\{1, 2, 3, 4\}$. Rounds $1, 2, 3$ constitute pre-interview filtering (e.g., applications, written tests, group debates), while round 4 is the final personal interview. Each candidate i from group g has a *latent ability* $A_{i,g}$. Passing round s requires

$$E\left[M_{s,i,g} \,\middle|\, \mathcal{I}_{s-1}\right] \;\geq\; \tau_{s,j}, \tag{OA.1}$$

where $\mathcal{I}_{s-1}$ is the employer's information after observing signals in rounds $1, \ldots, s-1$. As before, each round's threshold $\tau_{s,j}$ is weakly increasing with s and may differ across jobs j.

Noisy Signals per Round. At each stage s, the employer observes

$$M_{s,i,g} \;=\; A_{i,g} \;+\; \eta_{s,i,g},$$

where $\eta_{s,i,g}$ is a zero-mean noise term. We relax the assumption of independence across candidates and rounds, requiring only that:

(i) $\mathrm{E}[\eta_{s,i,g}] = 0$ for all s, i, g, and

(ii) $\eta_{s,i,1} \overset{d}{=} \eta_{s,i,0}$, so that for any subset of candidates (including those who have passed earlier rounds), there is no group-level shift or difference in noise distributions that would systematically favor one caste.

We allow that $(\eta_{s,i,g})$ may be correlated across s or i (e.g., a candidate's performance in different rounds might be interdependent), and may also depend on the conditioning event of "having passed earlier rounds." The key point is that *noise does not differ across castes*

in a way that would introduce persistent group-level biases. Under these assumptions, any stable group differences cannot be attributed to noise alone.

Updating Beliefs and Passing Rounds. At each stage s, the employer forms an updated belief about $A_{i,g}$ given $\mathcal{I}_s \equiv \{M_{1,i,g}, \ldots, M_{s,i,g}\}$. Let

$$E[A_{i,g} \mid \mathcal{I}_s] \quad \text{and} \quad Var(A_{i,g} \mid \mathcal{I}_s)$$

denote the posterior mean and variance of $A_{i,g}$. Intuitively, with more signals $\{M_{1,i,g}, \ldots, M_{s,i,g}\}$, beliefs about each candidate's ability become more precise, and the passing condition (OA.1) can be evaluated in terms of $E[M_{s,i,g} \mid \mathcal{I}_{s-1}]$.

Final Interview Thresholds and Caste Revelation. As before, let $\tau_{4,j}^{(g)}$ be the *group-specific* final-round threshold at job j. If caste becomes salient at the interview for group 1, a higher threshold ($\tau_{4,j}^{(1)} > \tau_{4,j}^{(0)}$) would imply discrimination. Our goal is to see how variation in $\tau_{4,j}$ across jobs still allows us to disentangle true ability differences from discrimination, even when $M_{s,i,g}$ include noise.

A.3.1 No Early Caste Penalty Implies No Large Right-Tail Deficit Before Interviews

If there were a substantial right-tail *ability* deficit for group 1, we would see an especially large gap in pass rates at earlier stages for *highly selective* jobs (i.e., with large $\tau_{s,j}$). Formally:

Proposition 3 (No Early Gap $\implies$ No Right-Tail Deficit Under Imperfect Information)**.** Consider any pre-interview stage $s < 4$. If group 1 truly has fewer high-ability candidates in the top tail, then for sufficiently large $\tau_{s,j}$,

$$P\big(\text{Pass}_{s,i,j} \mid g = 1, \mathcal{I}_{s-1}\big) \;<\; P\big(\text{Pass}_{s,i,j} \mid g = 0, \mathcal{I}_{s-1}\big).$$

Hence, if empirically no such gap arises—even in the most selective jobs—we can rule out a meaningful right-tail deficit for group 1 at the earlier stages.

Sketch. With noisy signals, passing round s depends on the posterior mean $E[M_{s,i,g} \mid \mathcal{I}_{s-1}]$. For large $\tau_{s,j}$, the employer needs $E[A_{i,g} \mid \mathcal{I}_{s-1}]$ to be sufficiently high to clear that threshold. If group 1 truly had fewer high-ability individuals, then as $\tau_{s,j}$ grows, the fraction of group 1 candidates surpassing it falls *faster* than for group 0. This mirrors the perfect-information argument in Proposition 1 and further relies on the *regularly decaying tail* assumption (e.g., $P(M_{s,i,g} \geq m) \sim C_g\, m^{-\alpha}$ for large m), which ensures that higher cutoffs intensify right-tail gaps. Because noise is symmetric across castes, it cannot mask or invert these latent-ability differences at sufficiently high cutoffs. Hence, if we observe no substantial disadvantage for group 1 at early rounds even in highly selective jobs, we can rule out a meaningful right-tail deficit for that group. □

A.3.2 Stable Promotion Gaps Across Selectivity Still Imply Discrimination

Just as in Section 5.7.1, consider post-hire performance Y_i tied to underlying ability $A_{i,g}$. A candidate is promoted if $Y_i \geq \kappa_j$. Let final-round passing require $E[M_{4,i,g} \mid \mathcal{I}_3] \geq \tau_{4,j}^{(g)}$. We then define a *promotion gap* across selectivity tiers. Suppose jobs in the top $(1-q)$ fraction have cutoffs $\tau_{4,j}^{q}$. If both castes face $\tau_{4,j}^{(0)} = \tau_{4,j}^{(1)} = \tau_{4,j}^{q}$, then as $\tau_{4,j}^{q} \to \infty$, average hired performance converges to the *common marginal candidate*, making the promotion gap shrink to zero. By contrast, if $\tau_{4,j}^{(1)} = \tau_{4,j}^{(0)} + \delta$ ($\delta > 0$), group 1's *marginal* hires have strictly higher *posterior* means, yielding a stable advantage in realized performance.

Proposition 4 (Stable Promotion Gaps Under Imperfect Information Identify Discrimination)**.** Under the same regularly decaying tail assumption used in Proposition 2 (e.g. for large m, $P(M_{4,i,g} \geq m) \sim C_g\, m^{-\alpha}$ with $\alpha > 1$), if the final-round threshold for group 1 exceeds that for group 0 by a constant $\delta > 0$, then the average hired performance for group 1 remains approximately δ units higher than group 0, even in highly selective jobs. Consequently, if promotion-rate gaps do not vanish at higher cutoffs, we can reject an unobserved-ability explanation in favor of discrimination.

Sketch. The argument parallels Proposition 2. Under imperfect information, stage-4 passage requires $E[M_{4,i,g} \mid \mathcal{I}_3] \geq \tau_{4,j}^{(g)}$. If group 1 faces a higher threshold $\tau_{4,j}^{(1)} = \tau_{4,j}^{(0)} + \delta$, then a typical group-1 hire must satisfy

$$E[A_{i,1} \mid \mathcal{I}_3] \;\approx\; \tau_{4,j}^{(0)} + \delta \;-\; E[\eta_{4,i,1} \mid \mathcal{I}_3].$$

Because noise is identically distributed across groups (mean-zero, no group-level bias), raising group 1's threshold by δ effectively *shifts* group 1's final hires to a higher region of the ability distribution.

As $\tau_{4,j}$ rises across selectivity tiers, the average performance of each group's hired candidates converges to the respective threshold they must clear. Since group 1's threshold is δ higher, group 1's average hired performance remains about δ above group 0's. By contrast, if both groups share the same threshold, they converge to the same marginal point, so any promotion gap collapses to zero. The regularly decaying tail assumption ensures that these marginal hires dominate average outcomes at high cutoffs, preserving a stable δ-sized gap under discrimination. □

Key Takeaway. Allowing for imperfect information changes the exact pass mechanics but preserves the same identifying variations. If there is *no* sign of a pre-interview deficit at high cutoffs and we see a stable or rising performance gap post-hire, the best explanation remains that group 1 faces a stricter interview threshold once caste is revealed, rather than any hidden ability shortfall.

A.4 Tables and Figures

Table OA.1: **Caste Differences in Mean Salary Among Graduates**

	Dependent variable:			
	log (Salary)			
	Without Controls	With Controls	Client Facing	Non-Client Facing
	(1)	(2)	(3)	(4)
Disadv. Caste	-0.231^{***}	-0.103^{***}	-0.107^{***}	-0.089^{***}
	(0.020)	(0.017)	(0.027)	(0.022)
Disadv. Caste (Oster)		-0.076^{***}	-0.095^{***}	-0.072^{***}
		(0.019)	(0.026)	(0.022)
Constant	11.235^{***}	10.863^{***}	10.883^{***}	10.848^{***}
	(0.012)	(0.040)	(0.057)	(0.052)
Observations	2,527	2,527	849	1,678
Adjusted R^2	0.049	0.412	0.549	0.382
Year FE	✓	✓	✓	✓
Other Controls	✓	✓	✓	✓
Job FE				

Notes: Online Appendix Table OA.1 includes estimates from an earnings regression run on the sample of all students who graduated with jobs obtained through the college's job fair for full-time jobs. The dependent variable is log earnings. The regressions include detailed controls including measures of pre-college skills, college academic performance, previous labor market experience, and other employer-relevant skills such as major and degree. Column (1) reports estimates from a linear earnings regression without controls. Column (2) reports estimates from a linear earnings regression with controls. Columns (3) and (4) report estimates from a linear earnings regression with controls run separately for client-facing and non-client-facing jobs, respectively. Note that the earnings regressions do not include job fixed effects because there are no caste differences in pay for identical jobs. Bias-adjusted coefficients are calculated using Oster (2019). Following Oster's recommended benchmarks, I set $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$, implying that unobservables are assumed to have the same degree of selection as observables, and the total explanatory power can increase by up to 30% beyond the observed R^2. $^{*}p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

Table OA.2: **Caste Differences in GPA Between Those Omitted From the Earnings Regression Versus Those Who Are Not**

	Dependent variable:	
	log (GPA)	
	(1)	(2)
Disadv. Caste	-0.033^{***}	-0.032^{***}
	(0.004)	(0.004)
Omitted	-0.167^{***}	-0.158^{***}
	(0.005)	(0.005)
Omitted * Disadv. Caste	-0.071^{***}	-0.070^{***}
	(0.007)	(0.007)
Constant	2.147^{***}	2.127^{***}
	(0.004)	(0.008)
Observations	4,164	4,164
Adjusted R^2	0.546	0.559
Year FE	✓	✓
Major		✓
Degree		✓
Job FE		
Note:	*p<0.1; **p<0.05; ***p<0.01	

Notes: Online Appendix Table OA.2 shows the caste differences in GPA between those omitted from the earnings regression versus those who are not. The "omitted" category includes those who deregistered from the placement process and those who registered but did not receive any job offers (Section 5.2). The latter group often participates in the aftermarket and usually pursue higher education or entrepreneurship (Section 3.2). Column (1) only includes year fixed effects. Column (2) adds major and degree as controls. Note that the regressions do not have job fixed effects because these can typically only be included for jobs in my main sample. Therefore, these job fixed effects cannot be separately identified from the coefficient on "omitted."

Table OA.3: **Caste Differences in GPA At Each Job Search Stage**

	Dependent variable:					
	log (GPA)					
	App	Test	Debate	Interview	Offer	Choice
	(1)	(2)	(3)	(4)	(5)	(6)
Disadv. Caste	-0.111^{***}	-0.049^{***}	-0.039^{***}	-0.019^{***}	0.010^{***}	0.009^{***}
	(0.005)	(0.003)	(0.002)	(0.001)	(0.002)	(0.002)
Disadv. Caste (Oster)	-0.104^{***}	-0.035^{***}	-0.034^{***}	-0.014^{***}	0.015^{***}	0.014^{***}
	(0.006)	(0.003)	(0.002)	(0.002)	(0.002)	(0.002)
Constant	2.057^{***}	2.131^{***}	2.163^{***}	2.212^{***}	2.214^{***}	2.199^{***}
	(0.028)	(0.024)	(0.023)	(0.006)	(0.006)	(0.028)
Mean GPA	7.776	7.967	8.17	8.26	8.419	8.419
Observations	4,005	3,618	3,266	3,031	2,527	2,527
Adjusted R^2	0.264	0.726	0.745	0.801	0.822	0.822
Year FE	✓	✓	✓	✓	✓	✓
Job FE	✓	✓	✓	✓	✓	✓
Major	✓	✓	✓	✓	✓	✓
Degree	✓	✓	✓	✓	✓	✓
Note:						*p<0.1; **p<0.05; ***p<0.01

Notes: Online Appendix Table OA.3 compares the *between*-caste and *within*-job GPA differences at each job search stage. These differences are *controlled*, with the linear regressions including standard controls such as major, degree, and year, as well as job fixed effects. Bias-adjusted coefficients are calculated using Oster (2019). Following Oster's recommended benchmarks, I set $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$, implying that unobservables are assumed to have the same degree of selection as observables, and the total explanatory power can increase by up to 30% beyond the observed R^2. $^{*}p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

Table OA.4: **Caste Differences in Mean Salary at Each Job Search Stage in the Full Sample**

	Dependent variable:					
	log (Salary)					
	App	Test	Debate	Interview	Offer	Choice
	(1)	(2)	(3)	(4)	(5)	(6)
Disadv. Caste	0.004	−0.009	−0.006	−0.018*	−0.105***	−0.103***
	(0.005)	(0.007)	(0.007)	(0.010)	(0.013)	(0.017)
Disadv. Caste (Oster)	0.005	0.010	0.011	0.012	−0.079***	−0.076***
	(0.005)	(0.008)	(0.008)	(0.011)	(0.014)	(0.019)
Constant	11.032***	10.904***	10.848***	10.758***	10.672***	10.863***
	(0.011)	(0.015)	(0.015)	(0.023)	(0.030)	(0.040)
Observations	4,005	3,618	3,266	3,031	2,527	2,527
Adjusted R^2	0.578	0.491	0.547	0.549	0.488	0.412
Year FE	✓	✓	✓	✓	✓	✓
Other Controls	✓	✓	✓	✓	✓	✓
Job FE						

Notes: Online Appendix Table OA.4 includes estimates from an earnings regression run on the sample of all students who remain in contention for some job at a given job search stage. For example, column (1) includes the sample of all students who apply for jobs recruiting from the college's job fair. Similarly, column (2) includes the sample of all students who remain in contention for some job at the written test round, and so on. The dependent variable is log earnings. I include detailed controls, including measures of pre-college skills, college academic performance, previous labor market experience, and other employer-relevant skills such as major and degree. Each column is a separate regression and includes these controls. Note that the earnings regressions do not include job fixed effects because there are no caste differences in pay for identical jobs. Bias-adjusted coefficients are calculated using Oster (2019). Following Oster's recommended benchmarks, I set $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$, implying that unobservables are assumed to have the same degree of selection as observables, and the total explanatory power can increase by up to 30% beyond the observed R^2. $^*p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

Table OA.5: **Caste Differences in the Success Rate at Each Job Search Stage Among Client-Facing Jobs**

	Success App (1)	Success Test (2)	Success Debate (3)	Success Interview (4)
Disadv. Caste (OLS)	-0.0066***	-0.0007	-0.0136***	-0.3062***
	(0.0022)	(0.0030)	(0.0046)	(0.0081)
Disadv. Caste (Oster)	0.002	0.006	-0.015***	-0.314***
	(0.002)	(0.004)	(0.005)	(0.009)
Mean Success Rate	0.5770	0.7860	0.2590	0.2600
Observations	117,770	67,927	53,422	13,842
Adjusted R^2	0.61588	0.37467	0.01189	0.24298
Year FE	✓	✓	✓	✓
Job FE	✓	✓	✓	✓
Other Controls	✓	✓	✓	✓

Notes: Online Appendix Table OA.5 shows the caste difference in the success rate at each stage of job search among client-facing jobs. Each observation is a student who remains in contention at a given screening round in a given job in a given year. Each column reports a separate linear probability regression and includes controls. These regressions also include job fixed effects. Note that some firms conduct the same screening round (written tests, group debates, interviews) for multiple job positions. However, since these job positions within the same firm represent different jobs, the regressions consider these jobs and the corresponding advance-or-out decisions made by them separately. Bias-adjusted coefficients are calculated using Oster (2019). Following Oster's recommended benchmarks, I set $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$, implying that unobservables are assumed to have the same degree of selection as observables, and the total explanatory power can increase by up to 30% beyond the observed R^2. $^*p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

Table OA.6: **Caste Differences in the Success Rate at Each Job Search Stage Among Non-Client-Facing Jobs**

	Success App (1)	Success Test (2)	Success Debate (3)	Success Interview (4)
Disadv. Caste (OLS)	0.0024	0.0011	0.0070**	-0.2867***
	(0.0018)	(0.0022)	(0.0034)	(0.0060)
Disadv. Caste (Oster)	0.008***	0.014***	0.007**	-0.286***
	(0.002)	(0.002)	(0.003)	(0.007)
Mean Success Rate	0.6140	0.8250	0.2580	0.2620
Observations	187,771	115,223	95,067	24,536
Adjusted R^2	0.60783	0.32634	0.02270	0.24723
Year FE	✓	✓	✓	✓
Job FE	✓	✓	✓	✓
Other Controls	✓	✓	✓	✓

Notes: Online Appendix Table OA.6 shows the caste difference in the success rate at each stage of job search among non-client-facing jobs. Each observation is a student who remains in contention at a given screening round in a given job in a given year. Each column reports a separate linear probability regression and includes controls. These regressions also include job fixed effects. Note that some firms conduct the same screening round (written tests, group debates, interviews) for multiple job positions. However, since these job positions within the same firm represent different jobs, the regressions consider these jobs and the corresponding advance-or-out decisions made by them separately. Bias-adjusted coefficients are calculated using Oster (2019). Following Oster's recommended benchmarks, I set $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$, implying that unobservables are assumed to have the same degree of selection as observables, and the total explanatory power can increase by up to 30% beyond the observed R^2. $^*p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

Table OA.7: **Caste Differences in Mean Salary at Each Job Search Stage in the Identifiable Sample**

	Dependent variable:					
	log (Salary)					
	App	Test	Debate	Interview	Offer	Choice
	(1)	(2)	(3)	(4)	(5)	(6)
Disadv. Caste (OLS)	0.012	-0.068^{***}	-0.034^{*}	-0.105^{***}	-0.146^{***}	-0.107^{**}
	(0.014)	(0.019)	(0.018)	(0.027)	(0.033)	(0.049)
Disadv. Caste (Oster)	0.013	-0.049^{***}	-0.016	-0.078^{**}	-0.119^{***}	-0.082
	(0.016)	(0.021)	(0.022)	(0.032)	(0.042)	(0.055)
Constant	10.991^{***}	10.869^{***}	10.807^{***}	10.734^{***}	10.601^{***}	10.834^{***}
	(0.030)	(0.042)	(0.041)	(0.064)	(0.078)	(0.114)
Observations	619	552	500	449	378	378
Adjusted R^2	0.546	0.470	0.537	0.535	0.509	0.362
Year FE	✓	✓	✓	✓	✓	✓
Other Controls	✓	✓	✓	✓	✓	✓
Job FE						

Notes: Online Appendix Table OA.7 includes estimates from an earnings regression run on the sample of students with identifiable caste names who remain in contention for some job at a given job search stage. For example, column (1) includes the sample of all students with identifiable caste names who apply for jobs recruiting from the college's job fair. Similarly, column (2) includes the sample of all students with identifiable caste names who remain in contention for some job at the written test round, and so on. The dependent variable is log earnings. Each column is a separate regression and includes controls. Note that the earnings regressions do not include job fixed effects because there are no caste differences in pay for identical jobs. Bias-adjusted coefficients are calculated using Oster (2019). Following Oster's recommended benchmarks, I set $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$, implying that unobservables are assumed to have the same degree of selection as observables, and the total explanatory power can increase by up to 30% beyond the observed R^2. $^{*}p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

Table OA.8: **Caste Differences in Mean Salary at Each Job Search Stage in the Non-Identifiable Sample**

	Dependent variable:					
	log (Salary)					
	App	Test	Debate	Interview	Offer	Choice
	(1)	(2)	(3)	(4)	(5)	(6)
Disadv. Caste (OLS)	0.002	0.0004	−0.003	−0.004	−0.100***	−0.106***
	(0.005)	(0.007)	(0.007)	(0.011)	(0.014)	(0.019)
Disadv. Caste (Oster)	0.004	0.019***	0.013	0.024**	−0.075***	−0.076***
	(0.006)	(0.007)	(0.008)	(0.011)	(0.015)	(0.018)
Constant	11.041***	10.913***	10.857***	10.766***	10.685***	10.869***
	(0.012)	(0.017)	(0.016)	(0.024)	(0.033)	(0.043)
Observations	3,386	3,066	2,766	2,582	2,149	2,149
Adjusted R^2	0.584	0.500	0.552	0.556	0.486	0.421
Year FE	✓	✓	✓	✓	✓	✓
Other Controls	✓	✓	✓	✓	✓	✓
Job FE						

Notes: Online Appendix Table OA.8 includes estimates from an earnings regression run on the sample of students with non-identifiable caste names who remain in contention for some job at a given job search stage. For example, column (1) includes the sample of all students with non-identifiable caste names who apply for jobs recruiting from the college's job fair. Similarly, column (2) includes the sample of all students with non-identifiable caste names who remain in contention for some job at the written test round, and so on. The dependent variable is log earnings. Each column is a separate regression and includes controls. Note that the earnings regressions do not include job fixed effects because there are no caste differences in pay for identical jobs. Bias-adjusted coefficients are calculated using Oster (2019). Following Oster's recommended benchmarks, I set $\delta = 1$ and $R^* = 1.3 \times \tilde{R}$, implying that unobservables are assumed to have the same degree of selection as observables, and the total explanatory power can increase by up to 30% beyond the observed R^2. $^*p < 0.1$; $^{**}p < 0.05$; $^{***}p < 0.01$.

Table OA.9: **Caste Differences in GPA Between Those Who Left the First Job Versus Those Who Did Not**

	Dependent variable:		
	log (GPA)		
	(1)	(2)	(3)
Disadv. Caste	-0.065^{***}	-0.065^{***}	0.010^{***}
	(0.004)	(0.004)	(0.002)
Left Job	-0.072^{***}	-0.072^{***}	-0.004^{**}
	(0.003)	(0.003)	(0.002)
Left Job * Disadv. Caste	0.047^{***}	0.047^{***}	-0.004
	(0.005)	(0.005)	(0.003)
Constant	2.188^{***}	2.182^{***}	2.198^{***}
	(0.003)	(0.005)	(0.028)
Year FE	✓	✓	✓
Major		✓	✓
Degree		✓	✓
Job FE			✓
Observations	2,489	2,489	2,489
Adjusted R^2	0.260	0.265	0.822
Note:		*p<0.1; **p<0.05; ***p<0.01	

Notes: Online Appendix Table OA.9 shows the caste differences in GPA between those who left the first job versus those who did not. Column (1) only includes year fixed effects. Column (2) adds major and degree as additional controls. Column (3) additionally controls for job fixed effects.

Figure OA.1: **Common Support Within Each College GPA Decile**

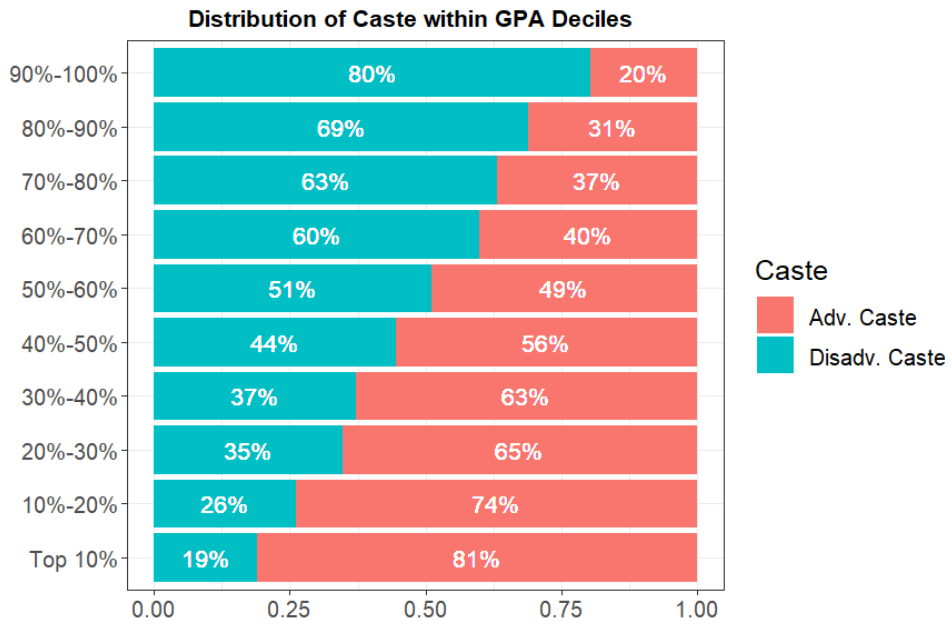

Notes: Online Appendix Figure OA.1 shows that there are students from both disadvantaged and advantaged castes within each college GPA decile. College GPA is normalized within each college degree and year.

Figure OA.2: **Caste Penalty in the Conditional Success Rate At the Application Reading Round in the Identifiable and the Non-Identifiable Sample**

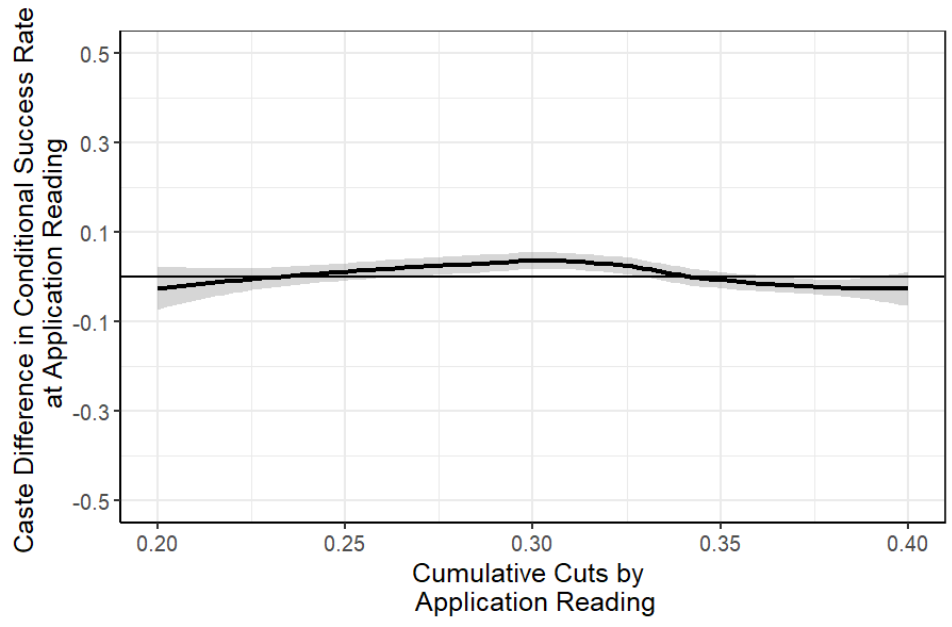

(a) Caste Penalty in Conditional Success Rates versus Cumulative Cuts by Application Reading in the Full Sample

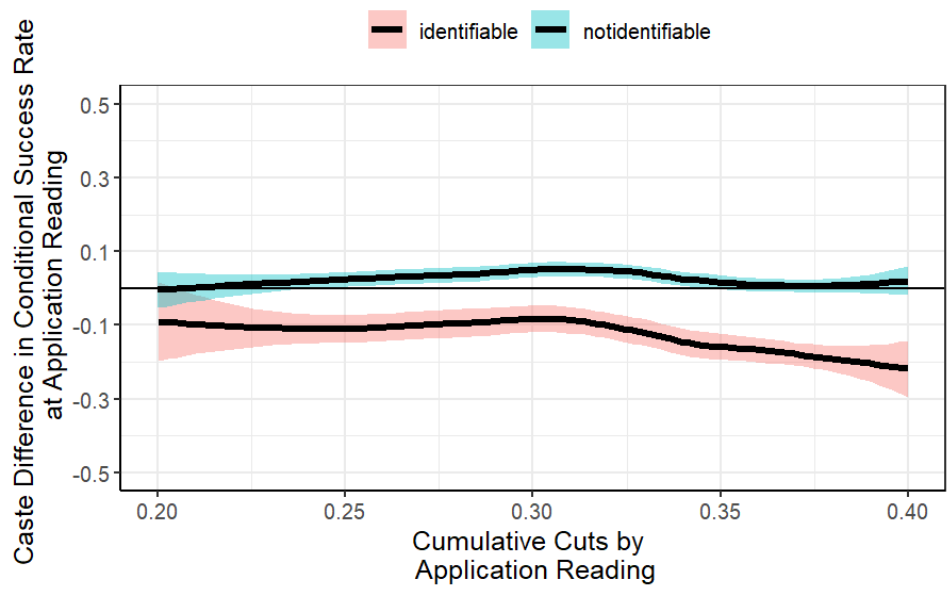

(b) Caste Penalty in Conditional Success Rates versus Cumulative Cuts by Application Reading in the Identifiable and Non-Identifiable Sample

Notes: Figure OA.2 shows the caste difference in the conditional (i.e., with controls) success rate against the cumulative cuts made by jobs by the end of the application reading round. Smoothed conditional means (thick black lines) are calculated using the "loess" function in R. The shaded areas in gray, blue, and red denote 95% confidence intervals. The thin solid horizontal line denotes a y-intercept of 0. Panel (a) is constructed from the full sample. Panel (b) is constructed separately for the sample with identifiable (red) and nonidentifable (blue) caste names. The x-axis shows the cumulative proportion of students cut by jobs by the end of the application reading round. The y-values are calculated by estimating, *for each job*, a standard linear probability model with controls and then plotting the coefficient on disadvantaged caste from these individual regressions on the y-axis. Each regression includes controls for student characteristics such as degree, major, college GPA, entrance exam score, and cohort year.

Figure OA.3: **Caste Penalty in the Uncontrolled Success Rate At Each Screening Round**

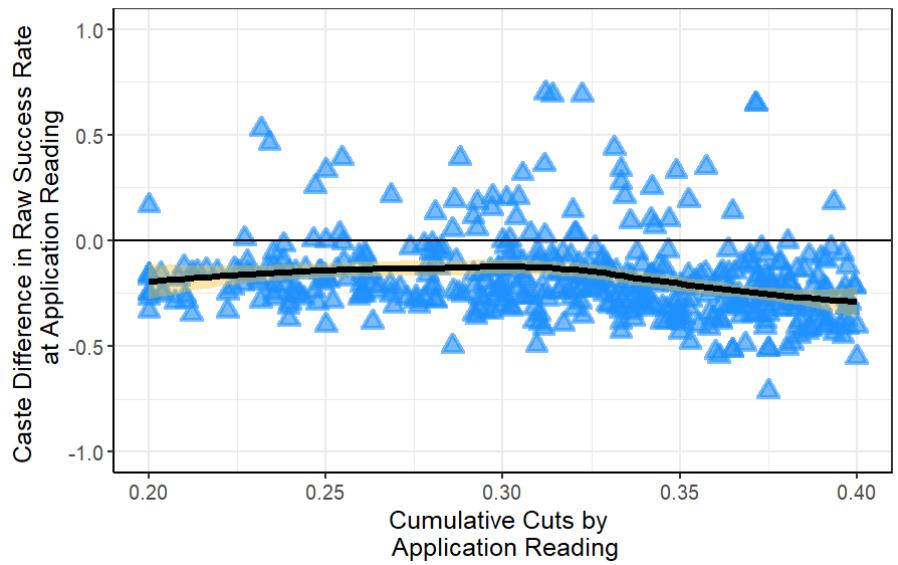

(a) Caste Penalty in Raw Success Rates versus Cumulative Cuts by Application Reading

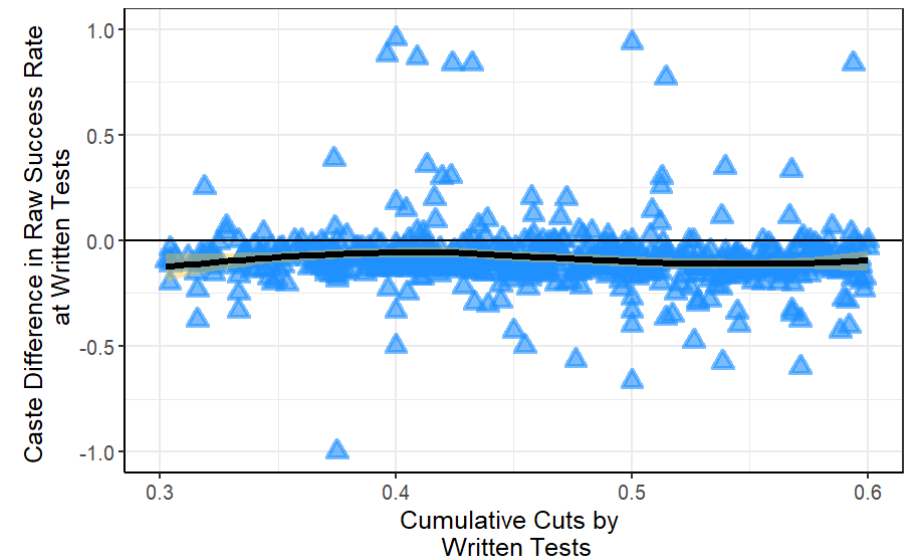

(b) Caste Penalty in Raw Success Rates versus Cumulative Cuts by Written Tests

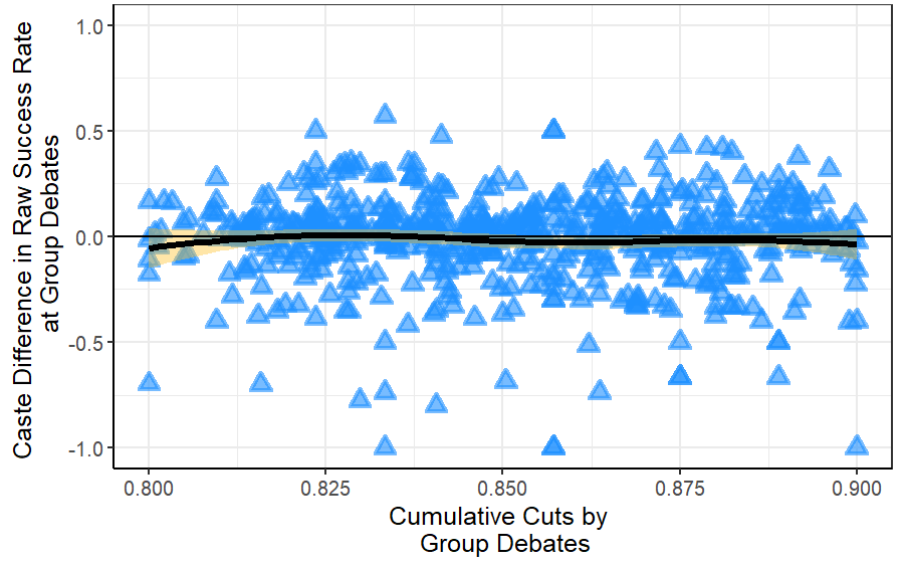

(c) Caste Penalty in Raw Success Rates versus Cumulative Cuts by Group Debates

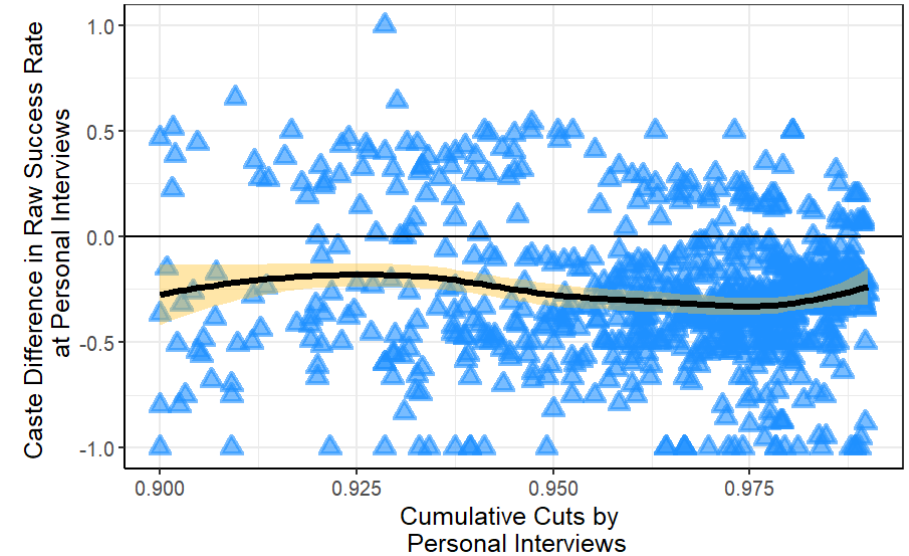

(d) Caste Penalty in Raw Success Rates versus Cumulative Cuts by Personal Interviews

Notes: Figure OA.3 shows the caste difference in the "raw" (i.e., without any controls) success rate against the cumulative cuts made by jobs by the end of each screening round. Each blue triangle denotes a job. Smoothed conditional means (thick black lines) are calculated using the "loess" function in R. The shaded areas in yellow denote 95% confidence intervals. The thin solid horizontal line denotes a y-intercept of 0. The x-axis shows the cumulative proportion of students cut by jobs by the end of each screening round. The y-axis shows the uncontrolled caste penalty in the success rate at each screening round.